
sherpa
Release 1.3

sherpa development team

Aug 16, 2024

CONTENTS

1	Introduction	1
2	Download pdf	3
3	Social groups	5
3.1	WeChat	5
3.2	QQ	5
3.3	Bilibili (B)	6
3.4	YouTube	6
4	Run Next-gen Kaldi in your browser	7
4.1	Visit our Huggingface space	7
4.2	YouTube Video	10
4.3	Other Huggingface spaces	10
5	Pre-trained models	11
5.1	Pre-trained models for different projects	11
5.2	How to download	11
6	sherpa	15
6.1	Installation	15
6.2	Pre-trained models	19
7	sherpa-ncnn	49
7.1	Tutorials	49
7.2	Installation	49
7.3	Python API	64
7.4	WebAssembly	73
7.5	C API	78
7.6	Endpointing	81
7.7	Android	84
7.8	iOS	99
7.9	Pre-trained models	123
7.10	Examples	154
7.11	FAQs	165
8	sherpa-onnx	167
8.1	Tutorials	167
8.2	Installation	167
8.3	Frequently Asked Question (FAQs)	189
8.4	Python	191

8.5	C API	214
8.6	Java API	218
8.7	Javascript API	228
8.8	Kotlin API	229
8.9	Swift API	229
8.10	Go API	230
8.11	C# API	237
8.12	Pascal API	246
8.13	Lazarus	250
8.14	WebAssembly	262
8.15	Android	267
8.16	iOS	281
8.17	Flutter	292
8.18	WebSocket	293
8.19	Hotwords (Contextual biasing)	308
8.20	Keyword spotting	324
8.21	Punctuation	326
8.22	Audio tagging	330
8.23	Spoken language identification	338
8.24	VAD	341
8.25	Pre-trained models	341
8.26	SenseVoice	528
8.27	Speaker Identification	560
8.28	Text-to-speech (TTS)	560
9	Triton	601
9.1	Installation	601
9.2	Triton-server	602
9.3	Triton-client	603
9.4	Perf Analyzer	604
9.5	TensorRT acceleration	605

**CHAPTER
ONE**

INTRODUCTION

`sherpa` is the deployment framework of the Next-gen Kaldi project.

`sherpa` does only one thing, using a pre-trained model to transcribe speech. If you are interested in how to train your own model or fine tune a pre-trained model, please refer to `icefall`.

At present, `sherpa` has the following sub-projects:

- `k2-fsa/sherpa`
- `k2-fsa/sherpa-onnx`
- `k2-fsa/sherpa-ncnn`

The differences are compared below:

	k2-fsa/sherpa	k2-fsa/sherpa-onnx	k2-fsa/sherpa-ncnn
Installation difficulty	hard	easy	easy
NN lib	PyTorch	onnxruntime	ncnn
CPU Support	x86, x86_64	x86, x86_64, arm32, arm64	x86, x86_64, arm32, arm64, **RISC-V**
GPU Support	Yes (with CUDA for NVIDIA GPUs)	Yes	Yes (with Vulkan for ARM GPUs)
OS Support	Linux, Windows, macOS	Linux, Windows, macOS, iOS, Android	Linux, Windows, macOS, iOS, Android
Support batch_size > 1	Yes	Yes	No
Provided APIs	C++, Python	C, C++, Python, C#, Java, Kotlin, Swift, Go, JavaScript, Dart	C, C++, Python, C#, Kotlin, Swift, Go
Supported functions	streaming speech recognition, non-streaming speech recognition	streaming speech recognition, non-streaming speech recognition, text-to-speech, speaker identification, speaker verification, spoken language identification, audio tagging, VAD, keyword spotting,	streaming speech recognition

We also support Triton. Please see [Triton](#).

**CHAPTER
TWO**

DOWNLOAD PDF

We provide a single pdf file containing all the documentation.

Hint: All Chinese related content is not included in the pdf file.

Please download it from the following address:

<https://k2-fsa.github.io/sherpa/sherpa.pdf>

Note: For Chinese users, you can use the following mirror:

<https://hub.nuua.cf/k2-fsa/sherpa/releases/download/doc/sherpa.pdf>

Please always download the latest version.

The pdf file is updated automagically whenever the doc is changed.

SOCIAL GROUPS

3.1 WeChat

If you have a [WeChat](#) account, you can scan the following QR code to join the WeChat group of next-gen Kaldi to get help.



3.2 QQ

The QQ group is also given below:



Hint: The QQ group number is 744602236.

3.3 Bilibili (B)

Please visit <https://search.bilibili.com/all?keyword=%E6%96%B0%E4%B8%80%E4%BB%A3Kaldi> for various demo videos of Next-gen Kaldi.

3.4 YouTube

To get the latest news of next-gen Kaldi, please subscribe the following YouTube channel by Nadira Povey:

https://www.youtube.com/channel/UC_VaumpkmINz1pNkFXAN9mw

CHAPTER
FOUR

RUN NEXT-GEN KALDI IN YOUR BROWSER

This page describes how to try Next-gen Kaldi in your browser.

Hint: You don't need to download or install anything. All you need is a browser.

The server is running on CPU within a docker container provided by [Huggingface](#) and you use a browser to interact with it. The browser can be run on Windows, macOS, Linux, or even on your phone or iPad.

You can upload a file for recognition, record your speech via a microphone from within the browser and submit it for recognition, or even provider an URL to an audio file for speech recognition.


Now let's get started.

4.1 Visit our Huggingface space

Start your browser and visit the following address:

<https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>

and you will see a page like the following screenshot:



Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://hf-mirror.com/spaces/k2-fsa/automatic-speech-recognition>

You can:

1. Select a language for recognition. Currently, we provide pre-trained models from `icefall` for the following languages: Chinese, English, and Chinese+English.
2. After selecting the target language, you can select a pre-trained model corresponding to the language.
3. Select the decoding method. Currently, it provides `greedy_search` and `modified_beam_search`.
4. If you selected `modified_beam_search`, you can choose the number of active paths during the search.
5. Either upload a file or record your speech for recognition.
6. Click the button `Submit for recognition`.
7. Wait for a moment and you will get the recognition results.

The following screenshot shows an example when selecting Chinese+English:

huggingface.co/spaces/k2-fsa/automatic-speech-recognition

Automatic Speech Recognition with Next-gen Kaldi

Language
 Chinese English Chinese+English

Select a model
luomingshuang/icefall_asr_tal-csasr_pruned_transducer_stateless5

Decoding method
 greedy_search modified_beam_search

Number of active paths for modified_beam_search
4

Upload from disk Record from microphone

Record from microphone
0:04 / 0:04

Submit for recognition

Recognized speech from recordings
这个是频繁的啊不认识记下来 FREQUENTLY 频繁地

Wave duration : 4.690 s
Processing time: 0.867 s
RTF: 0.867 / 4.690 = 0.185

In the bottom part of the page, you can find a table of examples. You can click one of them and then click **Submit for recognition**.

huggingface.co/spaces/k2-fsa/automatic-speech-recognition

Upload from disk

Drop Audio Here
- or -
Click to Upload

Submit for recognition

Recognized speech from uploaded file

Examples

Language	Select a model	Decoding method	Number of active paths for modified_beam_search	Upload from disk
English	csukuangfj/icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13	greedy_search	4	./test_wavs/librispeech/1089-134686-0001.wav
English	csukuangfj/icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13	greedy_search	4	./test_wavs/librispeech/1221-135766-0001.wav
English	csukuangfj/icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13	greedy_search	4	./test_wavs/librispeech/1221-135766-0002.wav
English	wgb14/icefall-asr-gigaspeech-pruned-transducer-stateless2	greedy_search	4	./test_wavs/gigaspeech/1-minute-audiobook.opus
English	wgb14/icefall-asr-gigaspeech-pruned-transducer-stateless2	greedy_search	4	./test_wavs/gigaspeech/100-seconds-podcast.opus
English	wgb14/icefall-asr-gigaspeech-pruned-transducer-stateless2	greedy_search	4	./test_wavs/gigaspeech/100-seconds-...

4.2 YouTube Video

We provide the following YouTube video demonstrating how to use <https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>.

Note: To get the latest news of next-gen Kaldi, please subscribe the following YouTube channel by Nadira Povey:

https://www.youtube.com/channel/UC_VaumpkmINz1pNkFXAN9mw

<https://youtu.be/E1N3r9dkKE4>

4.3 Other Huggingface spaces

- ASR + WebAssembly + sherpa-ncnn: Please see *Huggingface Spaces (WebAssembly)*
- TTS: Please see: <https://huggingface.co/spaces/k2-fsa/text-to-speech>

PRE-TRAINED MODELS

5.1 Pre-trained models for different projects

Project	Pretrained models
k2-fsa/sherpa	Click here
k2-fsa/sherpa-onnx	Click here
k2-fsa/sherpa-ncnn	Click here

5.2 How to download

We are hosting our pre-trained models on [Huggingface](#) as git repositories managed by [Git LFS](#).

There are at least two methods for downloading:

- Using `git lfs`
- Using `wget`

In the following, we use the pre-trained model [pkufool/icefall-asr-zipformer-streaming-wenetspeech-20230615 \(Chinese\)](#) as an example.

5.2.1 Using git lfs

Please first install `git-lfs` by following <https://git-lfs.com/>.

Linux

macOS

Windows

```
# apt/deb
sudo apt-get install git-lfs

# yum/rpm
sudo yum install git-lfs
```

Please see <https://github.com/git-lfs/git-lfs/blob/main/INSTALLING.md> for details.

```
brew install git-lfs
```

Please visit <https://gitforwindows.org/> to download and install git-lfs.

Then use the following commands to download pre-trained models:

Linux/macOS

Windows (Powershell)

Windows (cmd)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/pkufool/icefall-asr-zipformer-
˓→streaming-wenetspeech-20230615
cd icefall-asr-zipformer-streaming-wenetspeech-20230615
git lfs pull --include "exp/*chunk-16-left-128.*onnx"
```

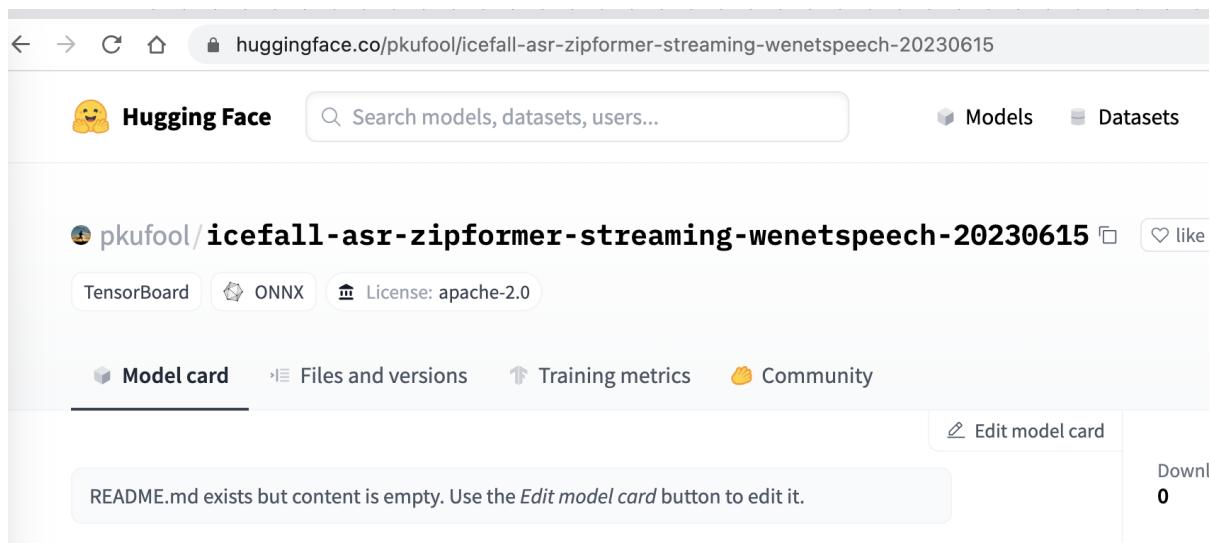
```
$env:GIT_LFS_SKIP_SMUDGE="1"
git clone https://huggingface.co/pkufool/icefall-asr-zipformer-streaming-wenetspeech-
˓→20230615
cd icefall-asr-zipformer-streaming-wenetspeech-20230615
git lfs pull --include "exp/*chunk-16-left-128.*onnx"
```

```
set GIT_LFS_SKIP_SMUDGE="1"
git clone https://huggingface.co/pkufool/icefall-asr-zipformer-streaming-wenetspeech-
˓→20230615
cd icefall-asr-zipformer-streaming-wenetspeech-20230615
git lfs pull --include "exp/*chunk-16-left-128.*onnx"
```

Note: It is very important to set the environment variable GIT_LFS_SKIP_SMUDGE to 1. We don't recommend using git lfs install as it will download many large files that we don't need.

5.2.2 Using wget

First, let us visit the [huggingface git repository](https://huggingface.co/pkufool/icefall-asr-zipformer-streaming-wenetspeech-20230615) of the pre-trained model:



Click **Files and versions** and navigate to the directory containing files for downloading:

The screenshot shows a GitHub repository page for 'icefall-asr-zipformer-streaming-wenetspeech-20230615'. The 'Files and versions' tab is active. A context menu is open over a file named 'decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx'. The menu item 'Copy Link Address' is highlighted with a blue background.

File	Size	Type
decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx	3.12 MB	LFS
decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx	12.4 MB	LFS
decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx	3.12 MB	LFS
decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx	12.4 MB	LFS
decoder-epoch-12-avg-4-chunk-32-left-128.int8.onnx	3.12 MB	LFS
decoder-epoch-12-avg-4-chunk-32-left-128.int8.onnx	12.4 MB	LFS

Right click the arrow that indicates downloading and copy the link address. After that, you can use, for instance, `wget` to download the file with the following command:

```
wget https://huggingface.co/pkufool/icefall-asr-zipformer-streaming-wenetspeech-20230615/
    ↵resolve/main/exp/decoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx
```

Repeat the process until you have downloaded all the required files.

**CHAPTER
SIX**

SHERPA

Hint: During speech recognition, it does not need to access the Internet. Everyting is processed locally on your device.

k2-fsa/sherpa use PyTorch for neural network computation.

<https://k2-fsa.github.io/icefall/model-export/export-with-torch-jit-script.html> for how to export models.

In the following, we describe how to use k2-fsa/sherpa.

6.1 Installation

6.1.1 From pre-compiled wheels

Note: This method supports only Linux and macOS for now. If you want to use Windows, please refer to *From source*.

You can find a list of pre-compiled wheels at the following URLs:

- CPU: <https://k2-fsa.github.io/sherpa/cpu.html>
- CUDA: <https://k2-fsa.github.io/sherpa/cuda.html>

In the following, we demonstrate how to install k2-fsa/sherpa from pre-compiled wheels.

Linux (CPU)

Suppose that we want to install the following wheel

```
https://huggingface.co/csukuangfj/kaldifeat/resolve/main/ubuntu-cpu/k2_sherpa-1.3.  
-dev20230725+cpu.torch2.0.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.whl
```

we can use the following methods:

```
# Before installing k2-fsa/sherpa, we have to install the following dependencies:  
# torch, k2, and kaldifeat
```

```
pip install torch==2.0.1+cpu -f https://download.pytorch.org/whl/torch_stable.html  
pip install k2==1.24.4.dev20231220+cpu.torch2.0.1 -f https://k2-fsa.github.io/k2/cpu.html  
pip install kaldifeat==1.25.3.dev20231221+cpu.torch2.0.1 -f https://csukuangfj.github.io/  
-kaldifeat/cpu.html
```

(continues on next page)

(continued from previous page)

```
# Now we can install k2-fsa/sherpa
pip install k2_sherpa==1.3.dev20230725+cpu.torch2.0.1 -f https://k2-fsa.github.io/sherpa/
˓→cpu.html
```

Please see [Check your installation](#).

macOS (CPU)

Suppose that we want to install the following wheel

```
https://huggingface.co/csukuangfj/kaldifeat/resolve/main/macos/k2_sherpa-1.3.
˓→dev20230725+cpu.torch2.0.1-cp311-cp311-macosx_10_9_x86_64.whl
```

we can use the following methods:

```
# Before installing k2-fsa/sherpa, we have to install the following dependencies:
# torch, k2, and kaldifeat

pip install torch==2.0.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
pip install k2==1.24.4.dev20231220+cpu.torch2.0.1 -f https://k2-fsa.github.io/k2/cpu.html
pip install kaldifeat==1.25.3.dev20231221+cpu.torch2.0.1 -f https://csukuangfj.github.io/
˓→kaldifeat/cpu.html

# Now we can install k2-fsa/sherpa
pip install k2_sherpa==1.3.dev20230725+cpu.torch2.0.1 -f https://k2-fsa.github.io/sherpa/
˓→cpu.html
```

Please see [Check your installation](#).

Linux (CUDA)

Suppose that we want to install the following wheel

```
https://huggingface.co/csukuangfj/kaldifeat/resolve/main/ubuntu-cuda/k2_sherpa-1.3.
˓→dev20230725+cuda11.7.torch2.0.1-cp311-cp311-manylinux_2_17_x86_64.manylinux2014_x86_64.
˓→whl
```

we can use the following methods:

```
# Before installing k2-fsa/sherpa, we have to install the following dependencies:
# torch, k2, and kaldifeat

pip install torch==2.0.1+cu117 -f https://download.pytorch.org/whl/torch_stable.html
pip install k2==1.24.4.dev20231220+cuda11.7.torch2.0.1 -f https://k2-fsa.github.io/k2/
˓→cuda.html
pip install kaldifeat==1.25.3.dev20231221+cuda11.7.torch2.0.1 -f https://csukuangfj.
˓→github.io/kaldifeat/cuda.html

# Now we can install k2-fsa/sherpa
pip install k2_sherpa==1.3.dev20230725+cuda11.7.torch2.0.1 -f https://k2-fsa.github.io/
˓→sherpa/cuda.html
```

(continues on next page)

(continued from previous page)

Please see [Check your installation.](#)

6.1.2 From source

This section describe how to install k2-fsa/sherpa from source.

Install dependencies

Before installing k2-fsa/sherpa from source, we have to install the following dependencies.

- PyTorch
- k2
- kaldifeat

CPU

CUDA

Suppose that we select `torch==2.0.1`. We can use the following commands to install the dependencies:

Linux

macOS

Windows

```
pip install torch==2.0.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
pip install k2==1.24.4.dev20231220+cpu.torch2.0.1 -f https://k2-fsa.github.io/k2/cpu.html
pip install kaldifeat==1.25.3.dev20231221+cpu.torch2.0.1 -f https://csukuangfj.github.io/
↪ kaldifeat/cpu.html
```

```
pip install torch==2.0.1+cpu -f https://download.pytorch.org/whl/torch_stable.html
pip install k2==1.24.4.dev20231220+cpu.torch2.0.1 -f https://k2-fsa.github.io/k2/cpu.html
pip install kaldifeat==1.25.3.dev20231221+cpu.torch2.0.1 -f https://csukuangfj.github.io/
↪ kaldifeat/cpu.html
```

To be done.

Suppose that we select `torch==2.0.1+cu117`. We can use the following commands to install the dependencies:

```
pip install torch==2.0.1+cu117 -f https://download.pytorch.org/whl/torch_
↪ stable.html
pip install k2==1.24.4.dev20231220+cuda11.7.torch2.0.1 -f https://k2-fsa.
↪ github.io/k2/cuda.html
pip install kaldifeat==1.25.3.dev20231221+cuda11.7.torch2.0.1 -f https://
↪ csukuangfj.github.io/kaldifeat/cuda.html
```

Next, please follow <https://k2-fsa.github.io/k2/installation/cuda-cudnn.html> to install CUDA toolkit.

Now we can start to build k2-fsa/sherpa from source.

For general users

You can use the following commands to install k2-fsa/sherpa:

```
# Please make sure you have installed PyTorch, k2, and kaldifeat
# before you continue
#
git clone http://github.com/k2-fsa/sherpa
cd sherpa
python3 -m pip install --verbose .
```

To uninstall k2-fsa/sherpa, please use

```
# Please run it outside of the k2-fsa/sherpa repo
#
pip uninstall k2-sherpa
```

Please see [Check your installation](#).

For developers and advanced users

You can also use the following commands to install k2-fsa/sherpa.

The advantage is that you can have several versions of k2-fsa/sherpa in a single environment.

```
git clone http://github.com/k2-fsa/sherpa
cd sherpa
mkdir build
cd build

# For torch >= 2.0, please use
#
# cmake -DCMAKE_CXX_STANDARD=17 ..
#

cmake ..
make -j

export PATH=$PWD/bin:$PATH
export PYTHONPATH=$PWD/lib:$PWD/..sherpa/python:$PYTHONPATH
```

Please see [Check your installation](#).

6.1.3 Check your installation

To check that you have installed k2-fsa/sherpa successfully, please run:

```
python3 -c "import sherpa; print(sherpa.__file__); print(sherpa.__version__)"

sherpa-online --help
sherpa-offline --help

sherpa-online-microphone --help
```

(continues on next page)

(continued from previous page)

```
sherpa-offline-microphone --help

sherpa-online-websocket-server --help
sherpa-online-websocket-client --help
sherpa-online-websocket-client-microphone --help

sherpa-offline-websocket-server --help
sherpa-offline-websocket-client --help
```

Congratulations! You have installed `k2-fsa/sherpa` successfully. Please refer to [Pre-trained models](#) to download pre-trained models.

Have fun with `k2-fsa/sherpa`!

We suggest that you install `k2-fsa/sherpa` by following [From pre-compiled wheels](#).

6.1.4 Where to get help

If you have any issues about the installation, please create an issue at the following address:

<https://github.com/k2-fsa/sherpa/issues>

6.2 Pre-trained models

Two kinds of end-to-end (E2E) models are supported by `k2-fsa/sherpa`:

- CTC
- Transducer

Hint: For transducer-based models, we only support stateless transducers. To the best of our knowledge, only `icefall` supports that. In other words, only transducer models from `icefall` are currently supported.

For CTC-based models, we support any type of models trained using CTC loss as long as you can export the model via torchscript. Models from the following frameworks are currently supported: `icefall`, `WeNet`, and `torchaudio` (Wav2Vec 2.0). If you have a CTC model and want it to be supported in `k2-fsa/sherpa`, please create an issue at <https://github.com/k2-fsa/sherpa/issues>.

Hint: You can try the pre-trained models in your browser without installing anything. See <https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>.

This page lists all available pre-trained models that you can download.

Hint: We provide pre-trained models for the following languages:

- Arabic
- Chinese
- English
- German

- Tibetan
-

Hint: We provide a colab notebook for you to try offline recognition step by step.

It shows how to install sherpa and use it as offline recognizer, which supports the models from icefall, the WeNet framework and torchaudio.

6.2.1 Offline CTC models

This sections list pre-trained CTC models from the following frameworks:

icefall

Hint: We use the binary `sherpa-offline` below for demonstration. You can replace `sherpa-offline` with `sherpa-offline-websocket-server`.

In this section, we list all pre-trained CTC models from icefall.

icefall-asr-gigaspeech-conformer-ctc (English)

```
# This model is trained using GigaSpeech
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/wgb14/icefall-asr-gigaspeech-
˓→conformer-ctc
cd icefall-asr-gigaspeech-conformer-ctc
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/HLG.pt"
git lfs pull --include "data/lang_bpe_500/tokens.txt"
mkdir test_wavs
cd test_wavs
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1089-
˓→134686-0001.wav
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1221-
˓→135766-0001.wav
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1221-
˓→135766-0002.wav
cd ..

# Decode with H
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

# Decode with HLG
sherpa-offline \
```

(continues on next page)

(continued from previous page)

```
--nn-model=./exp/cpu_jit.pt \
--hlg=./data/lang_bpe_500/HLG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-conformer-ctc-jit-bpe-500-2021-11-09 (English)

```
GIT-lfs_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
↪ librispeech-conformer-ctc-jit-bpe-500-2021-11-09
cd icefall-asr-librispeech-conformer-ctc-jit-bpe-500-2021-11-09

git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/tokens.txt"
git lfs pull --include "data/lang_bpe_500/HLG.pt"
git lfs pull --include "data/lang_bpe_500/HLG_modified.pt"

# Decode with H
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
--use-gpu=false \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

# Decode with HLG
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
--hlg=./data/lang_bpe_500/HLG.pt \
--use-gpu=false \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

# Decode with HLG (modified)
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
--hlg=./data/lang_bpe_500/HLG_modified.pt \
--use-gpu=false \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-tedlium3-conformer-ctc2 (English)

```
# This model is trained using Tedlium3
#
# See https://github.com/k2-fsa/icefall/pull/696
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/videodanchik/icefall-asr-tedlium3-
˓ conformer-ctc2
cd icefall-asr-tedlium3-conformer-ctc2
git lfs pull --include "exp/cpu_jit.pt"

git lfs pull --include "data/lang_bpe/HLG.pt"
git lfs pull --include "data/lang_bpe/tokens.txt"

git lfs pull --include "test_wavs/DanBarber_2010-219.wav"
git lfs pull --include "test_wavs/DanielKahneman_2010-157.wav"
git lfs pull --include "test_wavs/RobertGupta_2010U-15.wav"

# Decode with H
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe/tokens.txt \
./test_wavs/DanBarber_2010-219.wav \
./test_wavs/DanielKahneman_2010-157.wav \
./test_wavs/RobertGupta_2010U-15.wav

# Decode with HLG
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--hlg=./data/lang_bpe/HLG.pt \
--tokens=./data/lang_bpe/tokens.txt \
./test_wavs/DanBarber_2010-219.wav \
./test_wavs/DanielKahneman_2010-157.wav \
./test_wavs/RobertGupta_2010U-15.wav
```

icefall_asr_librispeech_conformer_ctc (English)

```
# This model is trained using LibriSpeech
#
# See https://github.com/k2-fsa/icefall/pull/13
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/pkufool/icefall_asr_librispeech-
˓ conformer_ctc
cd icefall_asr_librispeech_conformer_ctc

git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe/HLG.pt"

# Decode with H
```

(continues on next page)

(continued from previous page)

```
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

# Decode with HLG
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--hlg=./data/lang_bpe/HLG.pt \
--tokens=./data/lang_bpe/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall_asr_aishell_conformer_ctc (Chinese)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/pkufool/icefall_asr_aishell_
→conformer_ctc
cd icefall_asr_aishell_conformer_ctc
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_char/HLG.pt"

# Decode with an H graph
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_char/tokens.txt \
./test_waves/BAC009S0764W0121.wav \
./test_waves/BAC009S0764W0122.wav \
./test_waves/BAC009S0764W0123.wav

# Decode with an HLG graph
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_char/tokens.txt \
--hlg=./data/lang_char/HLG.pt \
./test_waves/BAC009S0764W0121.wav \
./test_waves/BAC009S0764W0122.wav \
./test_waves/BAC009S0764W0123.wav
```

icefall-asr-mgb2-conformer_ctc-2022-27-06 (Arabic)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/AmirHussein/icefall-asr-mgb2-
˓→conformer_ctc-2022-27-06
cd icefall-asr-mgb2-conformer_ctc-2022-27-06
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_5000/HLG.pt"
git lfs pull --include "data/lang_bpe_5000/tokens.txt"

# Decode with an H graph
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_5000/tokens.txt \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0053813:0054281.wav \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0051454:0052244.wav \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0052244:0053004.wav

# Decode with an HLG graph
sherpa-offline \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_5000/tokens.txt \
--hlg=./data/lang_bpe_5000/HLG.pt \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0053813:0054281.wav \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0051454:0052244.wav \
./test_wavs/94D37D38-B203-4FC0-9F3A-538F5C174920_spk-0001_seg-0052244:0053004.wav
```

WeNet

This section lists models from WeNet.

wenet-english-model (English)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/wenet-english-model
cd wenet-english-model
git lfs pull --include "final.zip"

sherpa-offline \
--normalize-samples=false \
--modified=true \
--nn-model=./final.zip \
--tokens=./units.txt \
--use-gpu=false \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

wenet-chinese-model (Chinese)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/wenet-chinese-model
cd wenet-chinese-model
git lfs pull --include "final.zip"

sherpa-offline \
--normalize-samples=false \
--modified=true \
--nn-model=./final.zip \
--tokens=./units.txt \
./test_wavs/BAC009S0764W0121.wav \
./test_wavs/BAC009S0764W0122.wav \
./test_wavs/BAC009S0764W0123.wav \
./test_wavs/DEV_T0000000000.wav \
./test_wavs/DEV_T0000000001.wav \
./test_wavs/DEV_T0000000002.wav
```

torchaudio

This section lists models from torchaudio.

wav2vec2_asr_base (English)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio
cd wav2vec2.0-torchaudio

# Note: There are other kinds of models fine-tuned with different
# amount of data. We use a model that is fine-tuned with 10 minutes of data.

git lfs pull --include "wav2vec2_asr_base_10m.pt"

sherpa-offline \
--nn-model=wav2vec2_asr_base_10m.pt \
--tokens=tokens.txt \
--use-gpu=false \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

voxpupuli_asr_base (German)

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio
cd wav2vec2.0-torchaudio
git lfs pull --include "voxpupuli_asr_base_10k_de.pt"

sherpa-offline \
--nn-model=voxpupuli_asr_base_10k_de.pt \
```

(continues on next page)

(continued from previous page)

```
--tokens=tokens-de.txt \
--use-gpu=false \
./test_wavs/20120315-0900-PLENARY-14-de_20120315.wav \
./test_wavs/20170517-0900-PLENARY-16-de_20170517.wav
```

NeMo

This section lists models from NeMo.

sherpa-nemo-ctc-en-citrinet-512 (English)

This model is converted from

https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_citrinet_512

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-en-
↪citrinet-512
cd sherpa-nemo-ctc-en-citrinet-512
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
--tokens=./tokens.txt \
--use-gpu=false \
--modified=false \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 kuangfangjun root 142M Mar  9 21:23 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-zh-citrinet-512 (Chinese)

This model is converted from

https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_zh_citrinet_512

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-zh-
↪citrinet-512
cd sherpa-nemo-ctc-zh-citrinet-512
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
```

(continues on next page)

(continued from previous page)

```
--tokens=./tokens.txt \
--use-gpu=false \
--modified=true \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 kuangfangjun root 153M Mar 10 15:07 model.pt
```

Hint: Since the vocabulary size of this model is very large, i.e, 5207, we use `--modified=true` to use a modified CTC topology

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-zh-citrinet-1024-gamma-0-25 (Chinese)

This model is converted from

https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_zh_citrinet_1024_gamma_0_25

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-zh-
˓→citrinet-1024-gamma-0-25
cd sherpa-nemo-ctc-zh-citrinet-1024-gamma-0-25
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
--tokens=./tokens.txt \
--use-gpu=false \
--modified=true \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 kuangfangjun root 557M Mar 10 16:29 model.pt
```

Hint: Since the vocabulary size of this model is very large, i.e, 5207, we use `--modified=true` to use a modified CTC topology

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-de-citrinet-1024 (German)

This model is converted from

```
https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt\_de\_citrinet\_1024
```

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-de-
˓→citrinet-1024
cd sherpa-nemo-ctc-de-citrinet-1024
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
--tokens=./tokens.txt \
--use-gpu=false \
--modified=false \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 kuangfangjun root 541M Mar 10 16:55 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-en-conformer-small (English)

This model is converted from

```
https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt\_en\_conformer\_ctc\_small
```

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-en-
˓→conformer-small
cd sherpa-nemo-ctc-en-conformer-small
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
--tokens=./tokens.txt \
--use-gpu=false \
--modified=false \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 fangjun staff 82M Mar 10 19:55 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-en-conformer-medium (English)

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_ctc_medium

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-en-  
conformer-medium  
cd sherpa-nemo-ctc-en-conformer-medium  
git lfs pull --include "model.pt"  
  
sherpa-offline \  
--nn-model=./model.pt \  
--tokens=./tokens.txt \  
--use-gpu=false \  
--modified=false \  
--nemo-normalize=per_feature \  
./test_wavs/0.wav \  
./test_wavs/1.wav \  
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 fangjun staff 152M Mar 10 20:26 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-en-conformer-large (English)

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_ctc_large

Hint: The vocabulary size is 129

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-en-  
conformer-large  
cd sherpa-nemo-ctc-en-conformer-large  
git lfs pull --include "model.pt"  
  
sherpa-offline \  
--nn-model=./model.pt \  
--tokens=./tokens.txt \  
--use-gpu=false \  
--modified=false \  
--nemo-normalize=per_feature \  
--
```

(continues on next page)

(continued from previous page)

```
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 fangjun staff 508M Mar 10 20:44 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

sherpa-nemo-ctc-de-conformer-large (German)

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_de_conformer_ctc_large

Hint: The vocabulary size is 129

```
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-nemo-ctc-de-
˓→conformer-large
cd sherpa-nemo-ctc-de-conformer-large
git lfs pull --include "model.pt"

sherpa-offline \
--nn-model=./model.pt \
--tokens=./tokens.txt \
--use-gpu=false \
--modified=false \
--nemo-normalize=per_feature \
./test_wavs/0.wav \
./test_wavs/1.wav \
./test_wavs/2.wav
```

```
ls -lh model.pt
-rw-r--r-- 1 fangjun staff 508M Mar 10 21:34 model.pt
```

Caution: It is of paramount importance to specify `--nemo-normalize=per_feature`.

How to convert NeMo models to sherpa

This section describes how to export NeMo pre-trained CTC models to sherpa.

You can find a list of pre-trained models from NeMo by visiting:

https://catalog.ngc.nvidia.com/orgs/nvidia/collections/nemo_asr.

Let us take `stt_en_conformer_ctc_small` as an example.

You can use the following code to obtain `model.pt` and `tokens.txt`:

```
import nemo.collections.asr as nemo_asr
m = nemo_asr.models.EncDecCTCModelBPE.from_pretrained('stt_en_conformer_ctc_small')
m.export("model.pt")

with open('tokens.txt', 'w', encoding='utf-8') as f:
    f.write("<blk> 0\n")
    for i, s in enumerate(m.decoder.vocabulary):
        f.write(f"{s} {i+1}\n")
```

One thing to note is that the blank token has the largest token ID in NeMo. However, it is always `0` in sherpa. During network computation, we shift the last column of the `log_prob` tensor to the first column so that it matches the convention about using 0 for the blank in sherpa.

You can find the exported `model.pt` and `tokens.txt` by visiting

<https://huggingface.co/csukuangfj/sherpa-nemo-ctc-en-conformer-small>

6.2.2 Offline transducer models

Hint: We use the binary `sherpa-offline` below for demonstration. You can replace `sherpa-offline` with `sherpa-offline-websocket-server`.

Hint: Please visit <https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition> to try the pre-trained models in your browser. You don't need to install anything.

icefall

This sections lists models trained using `icefall`.

English

icefall-asr-librispeech-zipformer-2023-05-15

```
# This model is trained using LibriSpeech with zipformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/1058
#
```

(continues on next page)

(continued from previous page)

```
# normal-scaled model, number of model parameters: 65549011, i.e., 65.55 M
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/Zengwei/icefall-asr-librispeech-
zipformer-2023-05-15
cd icefall-asr-librispeech-zipformer-2023-05-15

git lfs pull --include "exp/jit_script.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/jit_script.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/jit_script.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/1089-134686-0001.wav \
    ./test_wavs/1221-135766-0001.wav \
    ./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-zipformer-small-2023-05-16

```
# This model is trained using LibriSpeech with zipformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/1058
#
# small-scaled model, number of model parameters: 23285615, i.e., 23.3 M
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/Zengwei/icefall-asr-librispeech-
zipformer-small-2023-05-16
cd icefall-asr-librispeech-zipformer-small-2023-05-16

git lfs pull --include "exp/jit_script.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/jit_script.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
```

(continues on next page)

(continued from previous page)

```

./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
--decoding-method=fast_beam_search \
--nn-model=./exp/jit_script.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

```

icefall-asr-librispeech-zipformer-large-2023-05-16

```

# This model is trained using LibriSpeech with zipformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/1058
#
# large-scaled model, number of model parameters: 148439574, i.e., 148.4 M
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/Zengwei/icefall-asr-librispeech-
↪zipformer-large-2023-05-16
cd icefall-asr-librispeech-zipformer-large-2023-05-16

git lfs pull --include "exp/jit_script.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/jit_script.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
--decoding-method=fast_beam_search \
--nn-model=./exp/jit_script.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav

```

icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04

```
# This model is trained using GigaSpeech + LibriSpeech + Common Voice 13.0 with zipformer
#
# See https://github.com/k2-fsa/icefall/pull/1010
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/yfyeung/icefall-asr-multidataset-
˓→pruned_transducer_stateless7-2023-05-04
cd icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04
git lfs pull --include "exp/cpu_jit-epoch-30-avg-4.pt"
cd exp
ln -s cpu_jit-epoch-30-avg-4.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done
```

icefall-asr-librispeech-pruned-transducer-stateless8-2022-12-02

```
# This model is trained using GigaSpeech + LibriSpeech with zipformer
#
# See https://github.com/k2-fsa/icefall/pull/728
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/WeijiZhuang/icefall-asr-
˓→librispeech-pruned-transducer-stateless8-2022-12-02
cd icefall-asr-librispeech-pruned-transducer-stateless8-2022-12-02
git lfs pull --include "exp/cpu_jit-torch-1.10.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

cd exp
rm cpu_jit.pt
ln -sv cpu_jit-torch-1.10.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done
```

(continues on next page)

(continued from previous page)

```
sherpa-offline \
--decoding-method=fast_beam_search \
--nn-model=./exp/cpu_jit.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-pruned-transducer-stateless8-2022-11-14

```
# This model is trained using GigaSpeech + LibriSpeech with zipformer
#
# See https://github.com/k2-fsa/icefall/pull/675
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
librispeech-pruned-transducer-stateless8-2022-11-14
cd icefall-asr-librispeech-pruned-transducer-stateless8-2022-11-14
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
--decoding-method=fast_beam_search \
--nn-model=./exp/cpu_jit.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-pruned-transducer-stateless7-2022-11-11

```
# This model is trained using LibriSpeech with zipformer
#
# See https://github.com/k2-fsa/icefall/pull/672
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
librispeech-pruned-transducer-stateless7-2022-11-11
```

(continues on next page)

(continued from previous page)

```
cd icefall-asr-librispeech-pruned-transducer-stateless7-2022-11-11
git lfs pull --include "exp/cpu_jit-torch-1.10.0.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"
cd exp
ln -s cpu_jit-torch-1.10.0.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/1089-134686-0001.wav \
    ./test_wavs/1221-135766-0001.wav \
    ./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13

```
# This model is trained using LibriSpeech + GigaSpeech
#
# See https://github.com/k2-fsa/icefall/pull/363
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
librispeech-pruned-transducer-stateless3-2022-05-13
cd icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

sherpa-offline \
    --decoding-method=fast_beam_search \
```

(continues on next page)

(continued from previous page)

```
--nn-model=./exp/cpu_jit.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-gigaspeech-pruned-transducer-stateless2

```
# This model is trained using GigaSpeech
#
# See https://github.com/k2-fsa/icefall/pull/318
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/wgb14/icefall-asr-gigaspeech-
˓→pruned-transducer-stateless2
cd icefall-asr-gigaspeech-pruned-transducer-stateless2
git lfs pull --include "exp/cpu_jit-iter-3488000-avg-15.pt"
git lfs pull --include "data/lang_bpe_500/bpe.model"

cd ../exp
ln -s cpu_jit-iter-3488000-avg-15.pt cpu_jit.pt
cd ..

# Since this repo does not provide tokens.txt, we generate it from bpe.model
# by ourselves
/path/to/sherpa/scripts/bpe_model_to_tokens.py ./data/lang_bpe_500/bpe.model > ./data/
˓→lang_bpe_500/tokens.txt

mkdir test_wavs
cd test_wavs
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1089-
˓→134686-0001.wav
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1221-
˓→135766-0001.wav
wget https://huggingface.co/csukuangfj/wav2vec2.0-torchaudio/resolve/main/test_wavs/1221-
˓→135766-0002.wav

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done
```

Chinese

icefall_asr_wenetspeech_pruned_transducer_stateless2

```
# This models is trained using WenetSpeech
#
# See https://github.com/k2-fsa/icefall/pull/349
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/luomingshuang/icefall_asr_
-wenetspeech_pruned_transducer_stateless2

cd icefall_asr_wenetspeech_pruned_transducer_stateless2
git lfs pull --include "exp/cpu_jit_epoch_10_avg_2_torch_1.7.1.pt"
git lfs pull --include "data/lang_char/LG.pt"
cd exp
ln -s cpu_jit_epoch_10_avg_2_torch_1.7.1.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char/tokens.txt \
        ./test_wavs/DEV_T0000000000.wav \
        ./test_wavs/DEV_T0000000001.wav \
        ./test_wavs/DEV_T0000000002.wav
done

sherpa-offline \
    --decoding-method=$m \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_char/LG.pt \
    --tokens=./data/lang_char/tokens.txt \
    ./test_wavs/DEV_T0000000000.wav \
    ./test_wavs/DEV_T0000000001.wav \
    ./test_wavs/DEV_T0000000002.wav
```

icefall_asr_aiddatatang-200zh_pruned_transducer_stateless2

```
# This models is trained using aiddatatang_200zh
#
# See https://github.com/k2-fsa/icefall/pull/355
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/luomingshuang/icefall_asr_
-aiddatatang-200zh_pruned_transducer_stateless2
cd icefall_asr_aiddatatang-200zh_pruned_transducer_stateless2
git lfs pull --include "exp/cpu_jit_torch.1.7.1.pt"

cd exp
ln -sv cpu_jit_torch.1.7.1.pt cpu_jit.pt
```

(continues on next page)

(continued from previous page)

```
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char/tokens.txt \
        ./test_wavs/T0055G0036S0002.wav \
        ./test_wavs/T0055G0036S0003.wav \
        ./test_wavs/T0055G0036S0004.wav
done
```

icefall-asr-alimeeting-pruned-transducer-stateless7

```
# This models is trained using alimeeting (https://www.openslr.org/119/)
#
# See https://github.com/k2-fsa/icefall/pull/751
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/desh2608/icefall-asr-alimeeting-
˓pruned-transducer-stateless7
cd icefall-asr-alimeeting-pruned-transducer-stateless7

git lfs pull --include "exp/cpu_jit.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char/tokens.txt \
        ./test_wavs/165.wav \
        ./test_wavs/74.wav \
        ./test_wavs/209.wav
done
```

Chinese + English**icefall_asr_tal-csasr_pruned_transducer_stateless5**

```
# This models is trained using TAL_CSASR dataset from
# https://ai.100tal.com/dataset
# where each utterance contains both English and Chinese.
#
# See https://github.com/k2-fsa/icefall/pull/428
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/luomingshuang/icefall_asr_tal-
˓csasr_pruned_transducer_stateless5
cd icefall_asr_tal-csasr_pruned_transducer_stateless5
git lfs pull --include "exp/cpu_jit.pt"
```

(continues on next page)

(continued from previous page)

```
for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char/tokens.txt \
        ./test_wavs/210_36476_210_8341_1_1533271973_7057520_132.wav \
        ./test_wavs/210_36476_210_8341_1_1533271973_7057520_138.wav \
        ./test_wavs/210_36476_210_8341_1_1533271973_7057520_145.wav \
        ./test_wavs/210_36476_210_8341_1_1533271973_7057520_148.wav
done
```

Tibetan

icefall-asr-xbmu-amdo31-pruned-transducer-stateless7-2022-12-02

```

# This model is trained using the XBMU-AMDO31 corpus
#
# See https://github.com/k2-fsa/icefall/pull/706
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/syzym/icefall-asr-xbmu-amdo31-
˓pruned-transducer-stateless7-2022-12-02
cd icefall-asr-xbmu-amdo31-pruned-transducer-stateless7-2022-12-02
git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/a_0_cacm-A70_31116.wav \
        ./test_wavs/a_0_cacm-A70_31117.wav \
        ./test_wavs/a_0_cacm-A70_31118.wav
done

sherpa-offline \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/a_0_cacm-A70_31116.wav \
    ./test_wavs/a_0_cacm-A70_31117.wav \
    ./test_wavs/a_0_cacm-A70_31118.wav

```

icefall-asr-xbmu-amdo31-pruned-transducer-stateless5-2022-11-29

```
# This model is trained using the XBMU-AMDO31 corpus
#
# See https://github.com/k2-fsa/icefall/pull/706
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/syzym/icefall-asr-xbmu-amdo31-
˓→pruned-transducer-stateless5-2022-11-29
cd icefall-asr-xbmu-amdo31-pruned-transducer-stateless5-2022-11-29
git lfs pull --include "data/lang_bpe_500/LG.pt"
git lfs pull --include "data/lang_bpe_500/tokens.txt"
git lfs pull --include "exp/cpu_jit-epoch-28-avg-23-torch-1.10.0.pt"
git lfs pull --include "test_wavs/a_0_cacm-A70_31116.wav"
git lfs pull --include "test_wavs/a_0_cacm-A70_31117.wav"
git lfs pull --include "test_wavs/a_0_cacm-A70_31118.wav"

cd exp
rm cpu_jit.pt
ln -sv cpu_jit-epoch-28-avg-23-torch-1.10.0.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-offline \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/a_0_cacm-A70_31116.wav \
        ./test_wavs/a_0_cacm-A70_31117.wav \
        ./test_wavs/a_0_cacm-A70_31118.wav
done

sherpa-offline \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/a_0_cacm-A70_31116.wav \
    ./test_wavs/a_0_cacm-A70_31117.wav \
    ./test_wavs/a_0_cacm-A70_31118.wav
```

6.2.3 Online transducer models

Hint: We use the binary `sherpa-online` below for demonstration. You can replace `sherpa-online` with `sherpa-online-websocket-server` and `sherpa-online-microphone`.

Hint: At present, only streaming transducer models from `icefall` are supported.

icefall

This sections lists models trained using icefall.

English

icefall-asr-librispeech-streaming-zipformer-2023-05-17

icefall-asr-librispeech-pruned-transducer-stateless7-streaming-2022-12-29

(continues on next page)

(continued from previous page)

```

git lfs pull --include "exp/cpu_jit.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

# For fast_beam_search with LG
sherpa-online \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/1089-134686-0001.wav \
    ./test_wavs/1221-135766-0001.wav \
    ./test_wavs/1221-135766-0002.wav

```

icefall-asr-librispeech-conv-emformer-transducer-stateless2-2022-07-05

```

# This model is trained using LibriSpeech with ConvEmformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/440
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/Zengwei/icefall-asr-librispeech-
˓→conv-emformer-transducer-stateless2-2022-07-05
cd icefall-asr-librispeech-conv-emformer-transducer-stateless2-2022-07-05

git lfs pull --include "exp/cpu-jit-epoch-30-avg-10-torch-1.10.0.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"
cd exp
ln -sv cpu-jit-epoch-30-avg-10-torch-1.10.0.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

# For fast_beam_search with LG

```

(continues on next page)

(continued from previous page)

```
./build/bin/sherpa-online \
--decoding-method=fast_beam_search \
--nn-model=./exp/cpu_jit.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-lstm-transducer-stateless2-2022-09-03

```
# This model is trained using LibriSpeech with LSTM transducer
#
# See https://github.com/k2-fsa/icefall/pull/558
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
↳ librispeech-lstm-transducer-stateless2-2022-09-03
cd icefall-asr-librispeech-lstm-transducer-stateless2-2022-09-03

git lfs pull --include "exp/encoder_jit_trace-iter-468000-avg-16.pt"
git lfs pull --include "exp/decoder_jit_trace-iter-468000-avg-16.pt"
git lfs pull --include "exp/joiner_jit_trace-iter-468000-avg-16.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"

cd exp
ln -sv encoder_jit_trace-iter-468000-avg-16.pt encoder_jit_trace.pt
ln -sv decoder_jit_trace-iter-468000-avg-16.pt decoder_jit_trace.pt
ln -sv joiner_jit_trace-iter-468000-avg-16.pt joiner_jit_trace.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --encoder-model=./exp/encoder_jit_trace.pt \
        --decoder-model=./exp/decoder_jit_trace.pt \
        --joiner-model=./exp/joiner_jit_trace.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

# For fast_beam_search with LG
sherpa-online \
    --decoding-method=fast_beam_search \
    --encoder-model=./exp/encoder_jit_trace.pt \
    --decoder-model=./exp/decoder_jit_trace.pt \
    --joiner-model=./exp/joiner_jit_trace.pt \
    --lg=./data/lang_bpe_500/LG.pt \
```

(continues on next page)

(continued from previous page)

```
--tokens=./data/lang_bpe_500/tokens.txt \
./test_wavs/1089-134686-0001.wav \
./test_wavs/1221-135766-0001.wav \
./test_wavs/1221-135766-0002.wav
```

icefall-asr-librispeech-pruned-stateless-emformer-rnnt2-2022-06-01

```
# This model is trained using LibriSpeech with Emformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/390
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
→librispeech-pruned-stateless-emformer-rnnt2-2022-06-01
cd icefall-asr-librispeech-pruned-stateless-emformer-rnnt2-2022-06-01

git lfs pull --include "exp/cpu_jit-epoch-39-avg-6-use-averaged-model-1.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"
cd exp
ln -sv cpu_jit-epoch-39-avg-6-use-averaged-model-1.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_bpe_500/tokens.txt \
        ./test_wavs/1089-134686-0001.wav \
        ./test_wavs/1221-135766-0001.wav \
        ./test_wavs/1221-135766-0002.wav
done

# For fast_beam_search with LG

sherpa-online \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_bpe_500/LG.pt \
    --tokens=./data/lang_bpe_500/tokens.txt \
    ./test_wavs/1089-134686-0001.wav \
    ./test_wavs/1221-135766-0001.wav \
    ./test_wavs/1221-135766-0002.wav
```

icefall/librispeech.streaming_pruned_transducer_stateless4_20220625

```

# This model is trained using LibriSpeech with Conformer transducer
#
# See https://github.com/k2-fsa/icefall/pull/440
#
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/pkufool/icefall_librispeech_
˓→streaming_pruned_transducer_stateless4_20220625
cd icefall_librispeech_streaming_pruned_transducer_stateless4_20220625

git lfs pull --include "exp/cpu_jit-epoch-25-avg-3.pt"
git lfs pull --include "data/lang_bpe_500/LG.pt"
cd exp
ln -sv cpu_jit-epoch-25-avg-3.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
sherpa-online \
--decoding-method=$m \
--nn-model=./exp/cpu_jit.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_waves/1089-134686-0001.wav \
./test_waves/1221-135766-0001.wav \
./test_waves/1221-135766-0002.wav
done

# For fast_beam_search with LG

sherpa-online \
--decoding-method=fast_beam_search \
--nn-model=./exp/cpu_jit.pt \
--lg=./data/lang_bpe_500/LG.pt \
--tokens=./data/lang_bpe_500/tokens.txt \
./test_waves/1089-134686-0001.wav \
./test_waves/1221-135766-0001.wav \
./test_waves/1221-135766-0002.wav

```

Chinese

`icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming`

(continues on next page)

(continued from previous page)

```

git lfs pull --include "data/lang_char/LG.pt"
cd exp
ln -sv cpu_jit_epoch_7_avg_1_torch.1.7.1.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char/tokens.txt \
        ./test_wavs/DEV_T000000000000.wav \
        ./test_wavs/DEV_T000000000001.wav \
        ./test_wavs/DEV_T000000000002.wav
done

# For fast_beam_search with LG

sherpa-online \
    --decoding-method=fast_beam_search \
    --nn-model=./exp/cpu_jit.pt \
    --lg=./data/lang_char/LG.pt \
    --tokens=./data/lang_char/tokens.txt \
    ./test_wavs/DEV_T000000000000.wav \
    ./test_wavs/DEV_T000000000001.wav \
    ./test_wavs/DEV_T000000000002.wav

```

Chinese + English (all-in-one)

pfluo/k2fsa-zipformer-chinese-english-mixed

It is a streaming zipformer model

```

# This model supports both Chinese and English
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/pfluo/k2fsa-zipformer-chinese-
↪english-mixed
cd k2fsa-zipformer-chinese-english-mixed
git lfs pull --include "exp/cpu_jit.pt"

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char_bpe/tokens.txt \
        ./test_wavs/0.wav \
        ./test_wavs/1.wav \
        ./test_wavs/2.wav \
        ./test_wavs/3.wav \
        ./test_wavs/4.wav
done

```

icefall-asr-conv-emformer-transducer-stateless2-zh

It is a ConvEmformer model

```
# This model supports both Chinese and English
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/ptrnnull/icefall-asr-conv-emformer-
→transducer-stateless2-zh
cd icefall-asr-conv-emformer-transducer-stateless2-zh
git lfs pull --include "exp/cpu_jit-epoch-11-avg-1.pt"
cd exp
ln -sv cpu_jit-epoch-11-avg-1.pt cpu_jit.pt
cd ..

for m in greedy_search modified_beam_search fast_beam_search; do
    sherpa-online \
        --decoding-method=$m \
        --nn-model=./exp/cpu_jit.pt \
        --tokens=./data/lang_char_bpe/tokens.txt \
        ./test_wavs/0.wav \
        ./test_wavs/1.wav \
        ./test_wavs/2.wav \
        ./test_wavs/3.wav \
        ./test_wavs/4.wav
done
```

SHERPA-NCNN

Hint: During speech recognition, it does not need to access the Internet. Everything is processed locally on your device.

We support using `ncnn` to replace PyTorch for neural network computation. The code is put in a separate repository `sherpa-ncnn`

`sherpa-ncnn` is self-contained and everything can be compiled from source.

Please refer to <https://k2-fsa.github.io/icefall/model-export/export-ncnn.html> for how to export models to `ncnn` format.

In the following, we describe how to build `sherpa-ncnn` for Linux, macOS, Windows, embedded systems, Android, and iOS.

Also, we show how to use it for speech recognition with pre-trained models.

7.1 Tutorials

This page contains links to tutorials written by our users.

Caution: The tutorials are not necessarily written in English.

7.2 Installation

Hint: Please refer to [Python API](#) for its usage with Python.

In this section, we describe how to install `sherpa-ncnn` for the following platforms:

7.2.1 Installation videos

This section presents some videos about how to install and use `sherpa-ncnn`.

Window (64-bit)

The following video shows how to install and use `sherpa-ncnn` on 64-bit Windows.

Thanks to <https://space.bilibili.com/7990701> for his contribution.

Caution: It is in Chinese.

7.2.2 Linux

This page describes how to build `sherpa-ncnn` on Linux.

Hint: You can follow this section if you want to build `sherpa-ncnn` directly on your board.

Hint: For the Python API, please refer to [Python API](#).

All you need is to run:

x86/x86_64

32-bit ARM

64-bit ARM

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
make -j6
```

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build
cmake \
-DCMAKE_BUILD_TYPE=Release \
-DCMAKE_C_FLAGS="-march=armv7-a -mfloat-abi=hard -mfpu=neon" \
-DCMAKE_CXX_FLAGS="-march=armv7-a -mfloat-abi=hard -mfpu=neon" \
..
make -j6
```

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
```

(continues on next page)

(continued from previous page)

```
cd build
cmake \
-DCMAKE_BUILD_TYPE=Release \
-DCMAKE_C_FLAGS="-march=armv8-a" \
-DCMAKE_CXX_FLAGS="-march=armv8-a" \
..
make -j6
```

After building, you will find two executables inside the bin directory:

```
$ ls -lh bin/
total 13M
-rwxr-xr-x 1 kuangfangjun root 6.5M Dec 18 11:31 sherpa-ncnn
-rwxr-xr-x 1 kuangfangjun root 6.5M Dec 18 11:31 sherpa-ncnn-microphone
```

That's it!

Please read [Pre-trained models](#) for usages about the generated binaries.

Read below if you want to learn more.

You can strip the binaries by

```
$ strip bin/sherpa-ncnn
$ strip bin/sherpa-ncnn-microphone
```

After stripping, the file size of each binary is:

```
$ ls -lh bin/
total 12M
-rwxr-xr-x 1 kuangfangjun root 5.8M Dec 18 11:35 sherpa-ncnn
-rwxr-xr-x 1 kuangfangjun root 5.8M Dec 18 11:36 sherpa-ncnn-microphone
```

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained.

You can use the following commands to check that and you will find they depend only on system libraries.

```
$ readelf -d bin/sherpa-ncnn

Dynamic section at offset 0x5c0650 contains 34 entries:
  Tag          Type           Name/Value
 0x0000000000000001 (NEEDED)   Shared library: [libgomp.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)   Shared library: [libstdc++.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libmvec.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libgcc_s.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libc.so.6]
 0x0000000000000001d (RUNPATH) Library runpath: [$ORIGIN:]


$ readelf -d bin/sherpa-ncnn-microphone

Dynamic section at offset 0x5c45d0 contains 34 entries:
  Tag          Type           Name/Value
```

(continues on next page)

(continued from previous page)

0x0000000000000001 (NEEDED)	Shared library: [libpthread.so.0]
0x0000000000000001 (NEEDED)	Shared library: [libgomp.so.1]
0x0000000000000001 (NEEDED)	Shared library: [libstdc++.so.6]
0x0000000000000001 (NEEDED)	Shared library: [libm.so.6]
0x0000000000000001 (NEEDED)	Shared library: [libmvec.so.1]
0x0000000000000001 (NEEDED)	Shared library: [libgcc_s.so.1]
0x0000000000000001 (NEEDED)	Shared library: [libc.so.6]
0x0000000000000001d (RUNPATH)	Library runpath: [\$ORIGIN:]

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

7.2.3 macOS

This page describes how to build `sherpa-ncnn` on macOS.

Hint: For the Python API, please refer to [Python API](#).

All you need is to run:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
make -j6
```

After building, you will find two executables inside the `bin` directory:

```
$ ls -lh bin/
total 24232
-rwxr-xr-x  1 fangjun  staff   5.9M Dec 18 12:39 sherpa-ncnn
-rwxr-xr-x  1 fangjun  staff   6.0M Dec 18 12:39 sherpa-ncnn-microphone
```

That's it!

Please read [Pre-trained models](#) for usages about the generated binaries.

Read below if you want to learn more.

You can strip the binaries by

```
$ strip bin/sherpa-ncnn
$ strip bin/sherpa-ncnn-microphone
```

After stripping, the file size of each binary is:

```
$ ls -lh bin/
total 23000
-rwxr-xr-x  1 fangjun  staff   5.6M Dec 18 12:40 sherpa-ncnn
-rwxr-xr-x  1 fangjun  staff   5.6M Dec 18 12:40 sherpa-ncnn-microphone
```

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained. You can use the following commands to check that and you will find they depend only on system libraries.

```
$ otool -L bin/sherpa-ncnn
bin/sherpa-ncnn:
    /usr/local/opt/libomp/lib/libomp.dylib (compatibility version 5.0.0, current version 5.0.0)
        /usr/lib/libc++.1.dylib (compatibility version 1.0.0, current version 902.1.0)
            /usr/lib/libSystem.B.dylib (compatibility version 1.0.0, current version 1281.100.1)

$ otool -L bin/sherpa-ncnn-microphone
bin/sherpa-ncnn-microphone:
    /System/Library/Frameworks/CoreAudio.framework/Versions/A/CoreAudio (compatibility version 1.0.0, current version 1.0.0)
        /System/Library/Frameworks/AudioToolbox.framework/Versions/A/AudioToolbox (compatibility version 1.0.0, current version 1000.0.0)
            /System/Library/Frameworks/AudioUnit.framework/Versions/A/AudioUnit (compatibility version 1.0.0, current version 1.0.0)
                /System/Library/Frameworks/CoreFoundation.framework/Versions/A/CoreFoundation (compatibility version 150.0.0, current version 1677.104.0)
                    /System/Library/Frameworks/CoreServices.framework/Versions/A/CoreServices (compatibility version 1.0.0, current version 1069.24.0)
                        /usr/lib/libSystem.B.dylib (compatibility version 1.0.0, current version 1281.100.1)
                            /usr/local/opt/libomp/lib/libomp.dylib (compatibility version 5.0.0, current version 5.0.0)
                                /usr/lib/libc++.1.dylib (compatibility version 1.0.0, current version 902.1.0)
```

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

7.2.4 Windows

This page describes how to build sherpa-ncnn on Windows.

Hint: For the Python API, please refer to [Python API](#).

Hint: MinGW is known not to work. Please install Visual Studio before you continue.

64-bit Windows (x64)

All you need is to run:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
cmake --build . --config Release
```

It will generate two executables inside ./bin/Release/:

- **sherpa-ncnn.exe**: For decoding a single wave file.
- **sherpa-ncnn-microphone.exe**: For real-time speech recognition from a microphone

That's it!

Please read *Pre-trained models* for usages about the generated binaries.

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

32-bit Windows (x86)

All you need is to run:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build

# Please select one toolset among VS 2015, 2017, 2019, and 2022 below
# We use VS 2022 as an example.

# For Visual Studio 2015
# cmake -T v140,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2017
# cmake -T v141,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2019
# cmake -T v142,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2022
cmake -T v143,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

cmake --build . --config Release -- -m:6
```

It will generate two executables inside ./bin/Release/:

- **sherpa-ncnn.exe**: For decoding a single wave file.
- **sherpa-ncnn-microphone.exe**: For real-time speech recognition from a microphone

That's it!

Please read *Pre-trained models* for usages about the generated binaries.

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

7.2.5 Embedded Linux (arm)

This page describes how to build `sherpa-ncnn` for embedded Linux (arm, 32-bit) with cross-compiling on an x86 machine with Ubuntu OS.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Hint: This page is for cross-compiling.

Install toolchain

The first step is to install a toolchain for cross-compiling.

Warning: You can use any toolchain that is suitable for your platform. The toolchain we use below is just an example.

Visit <https://developer.arm.com/tools-and-software/open-source-software/developer-tools/gnu-toolchain/gnu-a/downloads/8-3-2019-03> to download the toolchain:

We are going to download `gcc-arm-8.3-2019.03-x86_64-arm-linux-gnueabihf.tar.xz`, which has been uploaded to <https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains>.

Assume you want to install it in the folder `$HOME/software`:

```
mkdir -p $HOME/software
cd $HOME/software
wget https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-8.3-
˓→2019.03-x86_64-arm-linux-gnueabihf.tar.xz
tar xvf gcc-arm-8.3-2019.03-x86_64-arm-linux-gnueabihf.tar.xz
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-arm-8.3-2019.03-x86_64-arm-linux-gnueabihf/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
arm-linux-gnueabihf-gcc --version
```

which should print the following log:

```
arm-linux-gnueabihf-gcc (GNU Toolchain for the A-profile Architecture 8.3-2019.03 (arm-  
↪rel-8.36)) 8.3.0  
Copyright (C) 2018 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions. There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Congratulations! You have successfully installed a toolchain for cross-compiling `sherpa-ncnn`.

Build `sherpa-ncnn`

Finally, let us build `sherpa-ncnn`.

```
git clone https://github.com/k2-fsa/sherpa-ncnn  
cd sherpa-ncnn  
../build-arm-linux-gnueabihf.sh
```

After building, you will get two binaries:

```
$ ls -lh build-arm-linux-gnueabihf/install/bin/  
  
total 6.6M  
-rwxr-xr-x 1 kuangfangjun root 2.2M Jan 14 21:46 sherpa-ncnn  
-rwxr-xr-x 1 kuangfangjun root 2.2M Jan 14 21:46 sherpa-ncnn-alsa
```

That's it!

Hint:

- `sherpa-ncnn` is for decoding a single file
 - `sherpa-ncnn-alsa` is for real-time speech recognition by reading the microphone with [ALSA](#)
-

Caution: We recommend that you use `sherpa-ncnn-alsa` on embedded systems such as Raspberry pi.

You need to provide a `device_name` when invoking `sherpa-ncnn-alsa`. We describe below how to find the device name for your microphone.

Run the following command:

```
arecord -l
```

to list all available microphones for recording. If it complains that `arecord`: command not found, please use `sudo apt-get install alsa-utils` to install it.

If the above command gives the following output:

```
**** List of CAPTURE Hardware Devices ****  
card 0: Audio [Axera Audio], device 0: 49ac000.i2s_mst-es8328-hifi-analog  
↪es8328-hifi-analog-0 []  
Subdevices: 1/1  
Subdevice #0: subdevice #0
```

In this case, I only have 1 microphone. It is card `0` and that card has only device `0`. To select card `0` and device `0` on that card, we need to pass `plughw:0,0` to `sherpa-ncnn-alsa`. (Note: It has the format `plughw:card_number,device_index`.)

For instance, you have to use

```
# Note: We use int8 models for encoder and joiner below.
./bin/sherpa-ncnn-alsa \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/tokens.txt \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-
    ↪ pnnx.ncnn.int8.param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-
    ↪ pnnx.ncnn.int8.bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-
    ↪ pnnx.ncnn.param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-
    ↪ pnnx.ncnn.bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-
    ↪ pnnx.ncnn.int8.param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-
    ↪ pnnx.ncnn.int8.bin \
    "plughw:0,0"
```

Please change the card number and also the device index on the selected card accordingly in your own situation. Otherwise, you won't be able to record with your microphone.

Please read [Pre-trained models](#) for usages about the generated binaries.

Read below if you want to learn more.

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained.

You can use the following commands to check that and you will find they depend only on system libraries.

```
$ readelf -d build-arm-linux-gnueabihf/install/bin/sherpa-ncnn

Dynamic section at offset 0x1c7ee8 contains 30 entries:
  Tag          Type                 Name/Value
  0x00000001 (NEEDED)           Shared library: [libstdc++.so.6]
  0x00000001 (NEEDED)           Shared library: [libm.so.6]
  0x00000001 (NEEDED)           Shared library: [libgcc_s.so.1]
  0x00000001 (NEEDED)           Shared library: [libpthread.so.0]
  0x00000001 (NEEDED)           Shared library: [libc.so.6]
  0x0000000f (RPATH)            Library rpath: [$ORIGIN]

$ readelf -d build-arm-linux-gnueabihf/install/bin/sherpa-ncnn-alsa

Dynamic section at offset 0x22ded8 contains 32 entries:
  Tag          Type                 Name/Value
  0x00000001 (NEEDED)           Shared library: [libasound.so.2]
  0x00000001 (NEEDED)           Shared library: [libgomp.so.1]
  0x00000001 (NEEDED)           Shared library: [libpthread.so.0]
  0x00000001 (NEEDED)           Shared library: [libstdc++.so.6]
  0x00000001 (NEEDED)           Shared library: [libm.so.6]
```

(continues on next page)

(continued from previous page)

0x00000001 (NEEDED)	Shared library: [libgcc_s.so.1]
0x00000001 (NEEDED)	Shared library: [libc.so.6]
0x0000000f (RPATH)	Library rpath: [\${ORIGIN}]

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

7.2.6 Embedded Linux (aarch64)

This page describes how to build `sherpa-ncnn` for embedded Linux (aarch64, 64-bit) with cross-compiling on an x86 machine with Ubuntu OS.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-ncnn` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Hint: This page is for cross-compiling.

Install toolchain

The first step is to install a toolchain for cross-compiling.

Warning: You can use any toolchain that is suitable for your platform. The toolchain we use below is just an example.

Visit <https://releases.linaro.org/components/toolchain/binaries/latest-7/aarch64-linux-gnu/> to download the toolchain.

We are going to download `gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz`, which has been uploaded to <https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains>.

Assume you want to install it in the folder `$HOME/software`:

```
mkdir -p $HOME/software
cd $HOME/software
wget https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-linaro-7.
˓→5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz

# For users from China
# huggingface,
```

(continues on next page)

(continued from previous page)

```
# wget https://hf-mirror.com/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-linaro-7.
→5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz
tar xvf gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
aarch64-linux-gnu-gcc --version
```

which should print the following log:

```
aarch64-linux-gnu-gcc (Linaro GCC 7.5-2019.12) 7.5.0
Copyright (C) 2017 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Congratulations! You have successfully installed a toolchain for cross-compiling `sherpa-ncnn`.

Build `sherpa-ncnn`

Finally, let us build `sherpa-ncnn`.

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
./build-aarch64-linux-gnu.sh
```

After building, you will get two binaries:

```
$ ls -lh build-aarch64-linux-gnu/install/bin/
total 10M
-rwxr-xr-x 1 kuangfangjun root 3.4M Jan 13 21:16 sherpa-ncnn
-rwxr-xr-x 1 kuangfangjun root 3.4M Jan 13 21:16 sherpa-ncnn-alsa
```

That's it!

Hint:

- `sherpa-ncnn` is for decoding a single file
- `sherpa-ncnn-alsa` is for real-time speech recognition by reading the microphone with `ALSA`

sherpa-ncnn-alsa

Caution: We recommend that you use `sherpa-ncnn-alsa` on embedded systems such as Raspberry pi.

You need to provide a `device_name` when invoking `sherpa-ncnn-alsa`. We describe below how to find the device name for your microphone.

Run the following command:

```
arecord -l
```

to list all available microphones for recording. If it complains that `arecord: command not found`, please use `sudo apt-get install alsa-utils` to install it.

If the above command gives the following output:

```
**** List of CAPTURE Hardware Devices ****
card 3: UACDemoV10 [UACDemoV1.0], device 0: USB Audio [USB Audio]
  Subdevices: 1/1
  Subdevice #0: subdevice #0
```

In this case, I only have 1 microphone. It is card 3 and that card has only device 0. To select card 3 and device 0 on that card, we need to pass `plughw:3,0` to `sherpa-ncnn-alsa`. (Note: It has the format `plughw:card_number,device_index`.)

For instance, you have to use

```
./bin/sherpa-ncnn-alsa \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/tokens.txt \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-
  ↪ pnx.ncnn.param \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-
  ↪ pnx.ncnn.bin \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-
  ↪ pnx.ncnn.param \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-
  ↪ pnx.ncnn.bin \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-
  ↪ pnx.ncnn.param \
  ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-
  ↪ pnx.ncnn.bin \
  "plughw:3,0"
```

Please change the card number and also the device index on the selected card accordingly in your own situation. Otherwise, you won't be able to record with your microphone.

Please read [Pre-trained models](#) for usages about the generated binaries.

Hint: If you want to select a pre-trained model for Raspberry that can be run on real-time, we recommend you to use [marcoyang/sherpa-ncnn-conv-emformer-transducer-small-2023-01-09 \(English\)](#).

Read below if you want to learn more.

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained.

You can use the following commands to check that and you will find they depend only on system libraries.

```
$ readelf -d build-aarch64-linux-gnu/install/bin/sherpa-ncnn

Dynamic section at offset 0x302a80 contains 30 entries:
  Tag      Type           Name/Value
 0x0000000000000001 (NEEDED)  Shared library: [libgomp.so.1]
 0x0000000000000001 (NEEDED)  Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)  Shared library: [libstdc++.so.6]
 0x0000000000000001 (NEEDED)  Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)  Shared library: [libgcc_s.so.1]
 0x0000000000000001 (NEEDED)  Shared library: [libc.so.6]
 0x000000000000000f (RPATH)   Library rpath: [$ORIGIN]

$ readelf -d build-aarch64-linux-gnu/install/bin/sherpa-ncnn-alsa

Dynamic section at offset 0x34ea48 contains 31 entries:
  Tag      Type           Name/Value
 0x0000000000000001 (NEEDED)  Shared library: [libasound.so.2]
 0x0000000000000001 (NEEDED)  Shared library: [libgomp.so.1]
 0x0000000000000001 (NEEDED)  Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)  Shared library: [libstdc++.so.6]
 0x0000000000000001 (NEEDED)  Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)  Shared library: [libgcc_s.so.1]
 0x0000000000000001 (NEEDED)  Shared library: [libc.so.6]
 0x000000000000000f (RPATH)   Library rpath: [$ORIGIN]
```

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

7.2.7 Embedded Linux (riscv64)

This page describes how to build `sherpa-ncnn` for embedded Linux (RISC-V, 64-bit) with cross-compiling on an x64 machine with Ubuntu OS.

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows/macOS or you don't want to setup your local environment for cross-compiling, please use the above colab notebook.

Install toolchain

The first step is to install a toolchain for cross-compiling.

```
sudo apt-get install gcc-riscv64-linux-gnu
sudo apt-get install g++-riscv64-linux-gnu
```

To check that you have installed the toolchain successfully, please run

```
$ riscv64-linux-gnu-gcc --version
riscv64-linux-gnu-gcc (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0
```

(continues on next page)

(continued from previous page)

```
Copyright (C) 2017 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions. There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

```
$ riscv64-linux-gnu-g++ --version  
riscv64-linux-gnu-g++ (Ubuntu 7.5.0-3ubuntu1~18.04) 7.5.0  
Copyright (C) 2017 Free Software Foundation, Inc.  
This is free software; see the source for copying conditions. There is NO  
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Build sherpa-ncnn

Next, let us build sherpa-ncnn.

```
git clone https://github.com/k2-fsa/sherpa-ncnn  
cd sherpa-ncnn  
./build-riscv64-linux-gnu.sh
```

After building, you will get two binaries:

```
$ ls -lh build-riscv64-linux-gnu/install/bin/  
total 3.8M  
-rwxr-xr-x 1 kuangfangjun root 1.9M May 23 22:12 sherpa-ncnn  
-rwxr-xr-x 1 kuangfangjun root 1.9M May 23 22:12 sherpa-ncnn-alsa
```

That's it!

Hint:

- sherpa-ncnn is for decoding a single file
 - sherpa-ncnn-alsa is for real-time speech recognition by reading the microphone with [ALSA](#)
-

Please read [Pre-trained models](#) for usages about the generated binaries.

Hint: If you want to select a pre-trained model for [VisionFive 2](#) that can be run on real-time, we recommend you to use [csukuangfj/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16](#) (*Bilingual, Chinese + English*).

You can use the following command with the above model:

```
./sherpa-ncnn \  
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.  
  ↵txt \  
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/  
  ↵encoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/  
  ↵encoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/  
  ↵decoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/  
  ↵decoder_jit_trace-pnnx.ncnn.bin \  
  ↵
```

(continues on next page)

(continued from previous page)

```
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_
↳ jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_
↳ jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/
↳ 5.wav \
4 \
greedy_search
```

Read below if you want to learn more.

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained.

You can use the following commands to check that and you will find they depend only on system libraries.

```
$ readelf -d build-riscv64-linux-gnu/install/bin/sherpa-ncnn

Dynamic section at offset 0x1d6dc0 contains 31 entries:
  Tag          Type           Name/Value
 0x0000000000000001 (NEEDED)   Shared library: [libgomp.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)   Shared library: [libstdc++.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libgcc_s.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libc.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [ld-linux-riscv64-
↳ lp64d.so.1]
 0x0000000000000001d (RUNPATH) Library runpath: [$ORIGIN]
 0x00000000000000020 (PREINIT_ARRAY) 0x1e18e0
 0x00000000000000021 (PREINIT_ARRAYSZ) 0x8

$ readelf -d build-riscv64-linux-gnu/install/bin/sherpa-ncnn-alsa

Dynamic section at offset 0x1d3db0 contains 32 entries:
  Tag          Type           Name/Value
 0x0000000000000001 (NEEDED)   Shared library: [libasound.so.2]
 0x0000000000000001 (NEEDED)   Shared library: [libgomp.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)   Shared library: [libstdc++.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libgcc_s.so.1]
 0x0000000000000001 (NEEDED)   Shared library: [libc.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [ld-linux-riscv64-
↳ lp64d.so.1]
 0x0000000000000001d (RUNPATH) Library runpath: [$ORIGIN]
 0x00000000000000020 (PREINIT_ARRAY) 0x1de8c8
 0x00000000000000021 (PREINIT_ARRAYSZ) 0x8
```

Please create an issue at <https://github.com/k2-fsa/sherpa-ncnn/issues> if you have any problems.

If you want to build an Android app, please refer to *Android*. If you want to build an iOS app, please refer to *iOS*.

7.3 Python API

Hint: It is known to work for Python >= 3.6 on Linux, macOS, and Windows.

In this section, we describe

1. How to install the Python package `sherpa-ncnn`
2. How to use `sherpa-ncnn` Python API for real-time speech recognition with a microphone
3. How to use `sherpa-ncnn` Python API to recognize a single file

7.3.1 Installation

You can use 1 of the 4 methods below to install the Python package `sherpa-ncnn`:

Method 1

Hint: This method supports x86_64, arm64 (e.g., Mac M1, 64-bit Raspberry Pi), and arm32 (e.g., 32-bit Raspberry Pi).

```
pip install sherpa-ncnn
```

If you use Method 1, it will install pre-compiled libraries. The disadvantage is that it may not be optimized for your platform, while the advantage is that you don't need to install `cmake` or a C++ compiler.

For the following methods, you have to first install:

- `cmake`, which can be installed using `pip install cmake`
- A C++ compiler, e.g., GCC on Linux and macOS, Visual Studio on Windows

Method 2

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
python3 setup.py install
```

Method 3

```
pip install git+https://github.com/k2-fsa/sherpa-ncnn
```

Method 4 (For developers and embedded boards)

x86/x86_64

32-bit ARM

64-bit ARM

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build

cmake \
-D SHERPA_NCNN_ENABLE_PYTHON=ON \
-D SHERPA_NCNN_ENABLE_PORTAUDIO=OFF \
-D BUILD_SHARED_LIBS=ON \
..

make -j6

export PYTHONPATH=$PWD/lib:$PWD/../../sherpa-ncnn/python:$PYTHONPATH
```

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build

cmake \
-D SHERPA_NCNN_ENABLE_PYTHON=ON \
-D SHERPA_NCNN_ENABLE_PORTAUDIO=OFF \
-D BUILD_SHARED_LIBS=ON \
-DCMAKE_C_FLAGS="-march=armv7-a -mfloating-abi=hard -mfpu=neon" \
-DCMAKE_CXX_FLAGS="-march=armv7-a -mfloating-abi=hard -mfpu=neon" \
..

make -j6

export PYTHONPATH=$PWD/lib:$PWD/../../sherpa-ncnn/python:$PYTHONPATH
```

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build
cd build

cmake \
-D SHERPA_NCNN_ENABLE_PYTHON=ON \
-D SHERPA_NCNN_ENABLE_PORTAUDIO=OFF \
-D BUILD_SHARED_LIBS=ON \
-DCMAKE_C_FLAGS="-march=armv8-a" \
-DCMAKE_CXX_FLAGS="-march=armv8-a" \
..

make -j6
```

(continues on next page)

(continued from previous page)

```
export PYTHONPATH=$PWD/lib:$PWD/./sherpa-ncnn/python:$PYTHONPATH
```

Let us check whether `sherpa-ncnn` was installed successfully:

```
python3 -c "import sherpa_ncnn; print(sherpa_ncnn.__file__)"  
python3 -c "import _sherpa_ncnn; print(_sherpa_ncnn.__file__)"
```

They should print the location of `sherpa_ncnn` and `_sherpa_ncnn`.

Hint: If you use Method 1, Method 2, and Method 3, you can also use

```
python3 -c "import sherpa_ncnn; print(sherpa_ncnn.__version__)"
```

It should print the version of `sherpa-ncnn`, e.g., 1.1.

Next, we describe how to use `sherpa-ncnn` Python API for speech recognition:

- (1) Real-time speech recognition with a microphone
- (2) Recognize a file

7.3.2 Real-time recognition with a microphone

The following Python code shows how to use `sherpa-ncnn` Python API for real-time speech recognition with a microphone.

Hint: We use `sounddevice` for recording. Please run `pip install sounddevice` before you run the code below.

Note: You can download the code from

<https://github.com/k2-fsa/sherpa-ncnn/blob/master/python-api-examples/speech-recognition-from-microphone.py>

Listing 7.1: Real-time speech recognition with a microphone using `sherpa-ncnn` Python API

```
import sys  
  
try:  
    import sounddevice as sd  
except ImportError as e:  
    print("Please install sounddevice first. You can use")  
    print()  
    print("  pip install sounddevice")  
    print()  
    print("to install it")  
    sys.exit(-1)
```

(continues on next page)

(continued from previous page)

```

import sherpa_ncnn

def create_recognizer():
    # Please replace the model files if needed.
    # See https://k2-fsa.github.io/sherpa/ncnn/pretrained_models/index.html
    # for download links.
    recognizer = sherpa_ncnn.Recognizer(
        tokens="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt",
        encoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_"
        ↪trace-pnnx.ncnn.param",
        encoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-"
        ↪pnnx.ncnn.bin",
        decoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_"
        ↪trace-pnnx.ncnn.param",
        decoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-"
        ↪pnnx.ncnn.bin",
        joiner_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-"
        ↪pnnx.ncnn.param",
        joiner_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-"
        ↪pnnx.ncnn.bin",
        num_threads=4,
    )
    return recognizer

def main():
    print("Started! Please speak")
    recognizer = create_recognizer()
    sample_rate = recognizer.sample_rate
    samples_per_read = int(0.1 * sample_rate) # 0.1 second = 100 ms
    last_result = ""
    with sd.InputStream(
        channels=1, dtype="float32", samplerate=sample_rate
    ) as s:
        while True:
            samples, _ = s.read(samples_per_read) # a blocking read
            samples = samples.reshape(-1)
            recognizer.accept_waveform(sample_rate, samples)
            result = recognizer.text
            if last_result != result:
                last_result = result
                print(result)

if __name__ == "__main__":
    devices = sd.query_devices()
    print(devices)
    default_input_device_idx = sd.default.device[0]
    print(f'Use default device: {devices[default_input_device_idx]["name"]}')

try:

```

(continues on next page)

(continued from previous page)

main()

Code explanation:**1. Import the required packages**

```
try:
    import sounddevice as sd
except ImportError as e:
    print("Please install sounddevice first. You can use")
    print()
    print("  pip install sounddevice")
    print()
    print("to install it")
    sys.exit(-1)

import sherpa_ncnn
```

Two packages are imported:

- `sounddevice`, for recording with a microphone
- `sherpa-ncnn`, for real-time speech recognition

2. Create the recognizer

```
def create_recognizer():
    # Please replace the model files if needed.
    # See https://k2-fsa.github.io/sherpa/ncnn/pretrained_models/index.html
    # for download links.
    recognizer = sherpa_ncnn.Recognizer(
        tokens=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt",
        encoder_param=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_"
        ↪trace-pnnx.ncnn.param",
        encoder_bin=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-"
        ↪pnnx.ncnn.bin",
        decoder_param=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_"
        ↪trace-pnnx.ncnn.param",
        decoder_bin=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-"
        ↪pnnx.ncnn.bin",
        joiner_param=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-"
        ↪pnnx.ncnn.param",
        joiner_bin=". ./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-"
        ↪pnnx.ncnn.bin",
        num_threads=4,
    )
    return recognizer

def main():
    print("Started! Please speak")
    recognizer = create_recognizer()
```

We use the model [csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06](https://huggingface.co/csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06) (*Chinese + English*) as an example, which is able to recognize both English and Chinese. You can replace it with other pre-trained models.

Please refer to [Pre-trained models](#) for more models.

Hint: The above example uses a `float16` encoder and joiner. You can also use the following code to switch to 8-bit (i.e., `int8`) quantized encoder and joiner.

```
recognizer = sherpa_ncnn.Recognizer(
    tokens="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt",
    encoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_"
    ↴jit_trace-pnnx.ncnn.int8.param",
    encoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_"
    ↴trace-pnnx.ncnn.int8.bin",
    decoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_"
    ↴jit_trace-pnnx.ncnn.param",
    decoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_"
    ↴trace-pnnx.ncnn.bin",
    joiner_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_"
    ↴jit_trace-pnnx.ncnn.int8.param",
    joiner_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_"
    ↴trace-pnnx.ncnn.int8.bin",
    num_threads=4,
)
```

3. Start recording

```
sample_rate = recognizer.sample_rate
with sd.InputStream(
```

Note that:

- We set channel to 1 since the model supports only a single channel
- We use dtype `float32` so that the resulting audio samples are normalized to the range [-1, 1].
- The sampling rate has to be `recognizer.sample_rate`, which is 16 kHz for all models at present.

4. Read audio samples from the microphone

```
samples_per_read = int(0.1 * sample_rate) # 0.1 second = 100 ms
) as s:
    while True:
```

Note that:

- It reads 100 ms of audio samples at a time. You can choose a larger value, e.g., 200 ms.
- No queue or callback is used. Instead, we use a blocking read here.
- The `samples` array is reshaped to a 1-D array

5. Invoke the recognizer with audio samples

```
samples, _ = s.read(samples_per_read) # a blocking read
```

Note that:

- `samples` has to be a 1-D tensor and should be normalized to the range [-1, 1].
- Upon accepting the audio samples, the recognizer starts the decoding automatically. There is no separate call for decoding.

6. Get the recognition result

```
samples = samples.reshape(-1)
recognizer.accept_waveform(sample_rate, samples)
result = recognizer.text
if last_result != result:
```

We use `recognizer.text` to get the recognition result. To avoid unnecessary output, we compare whether there is new result in `recognizer.text` and don't print to the console if there is nothing new recognized.

That's it!

Summary

In summary, you need to:

1. Create the recognizer
2. Start recording
3. Read audio samples
4. Call `recognizer.accept_waveform(sample_rate, samples)`
5. Call `recognizer.text` to get the recognition result

The following is a YouTube video for demonstration.

<https://youtu.be/74SxVueROok>

Hint: If you don't have access to YouTube, please see the following video from bilibili:

Note:

<https://github.com/k2-fsa/sherpa-ncnn/blob/master/python-api-examples/speech-recognition-from-microphone-with-endpoint-detection.py> supports endpoint detection.

Please see the following video for its usage:

7.3.3 Recognize a file

The following Python code shows how to use sherpa-ncnn Python API to recognize a wave file.

Caution: The sampling rate of the wave file has to be 16 kHz. Also, it should contain only a single channel and samples should be 16-bit (i.e., int16) encoded.

Note: You can download the code from

<https://github.com/k2-fsa/sherpa-ncnn/blob/master/python-api-examples/decode-file.py>

Listing 7.2: Decode a file with sherpa-ncnn Python API

```
import wave

import numpy as np
import sherpa_ncnn

def main():
    recognizer = sherpa_ncnn.Recognizer(
        tokens="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt",
        encoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_
        ↵trace-pnnx.ncnn.param",
        encoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-
        ↵pnnx.ncnn.bin",
        decoder_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_
        ↵trace-pnnx.ncnn.param",
        decoder_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-
        ↵pnnx.ncnn.bin",
        joiner_param="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-
        ↵pnnx.ncnn.param",
        joiner_bin="./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-
        ↵pnnx.ncnn.bin",
        num_threads=4,
    )

    filename = (
        "./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/1.wav"
    )
    with wave.open(filename) as f:
        assert f.getframerate() == recognizer.sample_rate, (
            f.getframerate(),
            recognizer.sample_rate,
        )
        assert f.getnchannels() == 1, f.getnchannels()
        assert f.getsampwidth() == 2, f.getsampwidth() # it is in bytes
        num_samples = f.getnframes()
        samples = f.readframes(num_samples)
        samples_int16 = np.frombuffer(samples, dtype=np.int16)
        samples_float32 = samples_int16.astype(np.float32)
```

(continues on next page)

(continued from previous page)

```

        samples_float32 = samples_float32 / 32768

    recognizer.accept_waveform(recognizer.sample_rate, samples_float32)

    tail_paddings = np.zeros(
        int(recognizer.sample_rate * 0.5), dtype=np.float32
    )
    recognizer.accept_waveform(recognizer.sample_rate, tail_paddings)

    recognizer.input_finished()

    print(recognizer.text)

if __name__ == "__main__":
    main()

```

We use the model [csukuangji/sherpa-ncnn-conv-emformer-transducer-2022-12-06](https://huggingface.co/csukuangji/sherpa-ncnn-conv-emformer-transducer-2022-12-06) (*Chinese + English*) as an example, which is able to recognize both English and Chinese. You can replace it with other pre-trained models.

Please refer to [Pre-trained models](#) for more models.

Hint: The above example uses a float16 encoder and joiner. You can also use the following code to switch to 8-bit (i.e., int8) quantized encoder and joiner.

```

recognizer = sherpa_ncnn.Recognizer(
    tokens=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt",
    encoder_param=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_"
    ↪jit_trace-pnnx.ncnn.int8.param",
    encoder_bin=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_"
    ↪trace-pnnx.ncnn.int8.bin",
    decoder_param=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_"
    ↪jit_trace-pnnx.ncnn.param",
    decoder_bin=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_"
    ↪trace-pnnx.ncnn.bin",
    joiner_param=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_"
    ↪trace-pnnx.ncnn.int8.param",
    joiner_bin=".sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_"
    ↪trace-pnnx.ncnn.int8.bin",
    num_threads=4,
)

```

7.4 WebAssembly

In this section, we describe how to build `sherpa-ncnn` for WebAssembly so that you can run real-time speech recognition with `WebAssembly`.

Please follow the steps below to build and run `sherpa-ncnn` for `WebAssembly`.

7.4.1 Install Emscripten

We need to compile the C/C++ files in `sherpa-ncnn` with the help of `emscripten`.

Please refer to https://emscripten.org/docs/getting_started/downloads for detailed installation instructions.

The following is an example to show you how to install it on Linux/macOS.

```
git clone https://github.com/emscripten-core/emsdk.git
cd emsdk
git pull
./emsdk install latest
./emsdk activate latest
source ./emsdk_env.sh
```

To check that you have installed `emscripten` successfully, please run:

```
emcc -v
```

The above command should print something like below:

```
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
(e967e20b4727956a30592165a3c1cde5c67fa0a8)
shared:INFO: (Emscripten: Running sanity checks)
(py38) fangjuns-MacBook-Pro:open-source fangjun$ emcc -v
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
(e967e20b4727956a30592165a3c1cde5c67fa0a8)
clang version 18.0.0 (https://github.com/llvm/llvm-project
a54545ba6514802178cf7cf1c1dd9f7efbf3cde7)
Target: wasm32-unknown-emscripten
Thread model: posix
InstalledDir: /Users/fangjun/open-source/emsdk/upstream/bin
```

Congratulations! You have successfully installed `emscripten`.

7.4.2 Build

After installing `emscripten`, we can build `sherpa-ncnn` for `WebAssembly` now.

Please use the following command to build it:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn

cd wasm/assets
wget -q https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
streaming-zipformer-bilingual-zh-en-2023-02-13.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13.tar.bz2
mv -v sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/*pnnx.ncnn.param .
mv -v sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/*pnnx.ncnn.bin .
mv -v sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/tokens.txt .
rm -rf sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13
rm -v sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13.tar.bz2
cd ../.

./build-wasm-simd.sh
```

Hint: You can visit <https://github.com/k2-fsa/sherpa-ncnn/releases/tag/models> to download a different model.

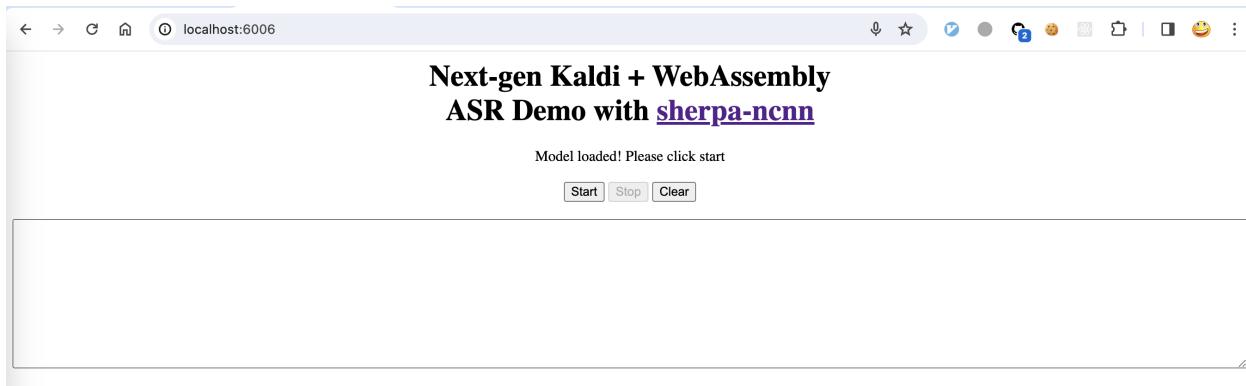
After building, you should see the following output:

```
Install the project...
-- Install configuration: "Release"
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/lib/
  ↵libkaldi-native-fbank-core.a
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/lib/
  ↵libncnn.a
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/.sherpa-
  ↵ncnn.pc
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/lib/
  ↵libsherpa-ncnn-core.a
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/lib/
  ↵libsherpa-ncnn-c-api.a
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/include/
  ↵sherpa-ncnn/c-api/c-api.h
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵sherpa-ncnn-wasm-main.js
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵sherpa-ncnn.js
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵app.js
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵index.html
-- Up-to-date: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵sherpa-ncnn-wasm-main.js
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵sherpa-ncnn-wasm-main.wasm
-- Installing: /Users/fangjun/open-source/sherpa-ncnn/build-wasm-simd/install/bin/wasm/
  ↵sherpa-ncnn-wasm-main.data
+ ls -lh install/bin/wasm
total 280152
-rw-r--r-- 1 fangjun staff  9.0K Feb  6 15:42 app.js
-rw-r--r-- 1 fangjun staff 936B Feb  6 15:42 index.html
-rw-r--r-- 1 fangjun staff 135M Feb  6 17:06 sherpa-ncnn-wasm-main.data
-rw-r--r-- 1 fangjun staff  79K Feb  6 17:06 sherpa-ncnn-wasm-main.js
-rw-r--r-- 1 fangjun staff 1.7M Feb  6 17:06 sherpa-ncnn-wasm-main.wasm
-rw-r--r-- 1 fangjun staff  6.9K Feb  6 15:42 sherpa-ncnn.js
```

Now you can use the following command to run it:

```
cd build-wasm-simd/install/bin/wasm/
python3 -m http.server 6006
```

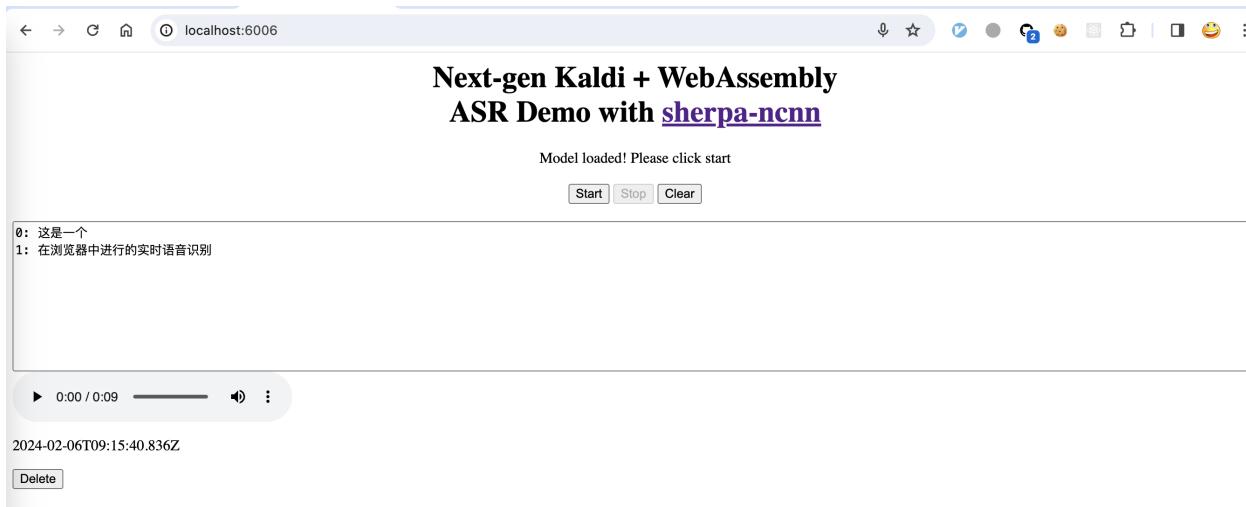
Start your browser and visit <http://localhost:6006/>; you should see the following page:



Now click start and speak! You should see the recognition results in the text box.

Warning: We are using a bilingual model (Chinese + English) in the above example, which means you can only speak Chinese or English in this case.

A screenshot is given below:



Congratulations! You have successfully run real-time speech recognition with [WebAssembly](#) in your browser.

7.4.3 Use pre-built WebAssembly library

In this section, we describe how to use the pre-built WebAssembly library of sherpa-ncnn for real-time speech recognition.

Note: Note that the pre-built library used in this section uses a bilingual model (Chinese + English), which is from [csukuangfj/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13](https://github.com/csukuangfj/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13) (*Bilingual, Chinese + English*).

Download

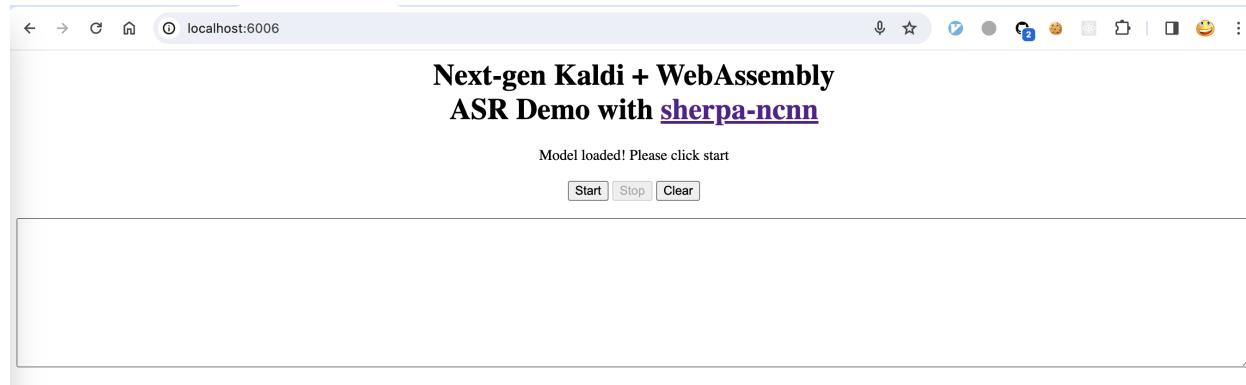
Please use the following command to download the pre-built library for version v2.1.7, which is the latest release as of 2024.02.06.

Hint: Please always use the latest release. You can visit <https://github.com/k2-fsa/sherpa-ncnn/releases> to find the latest release.

```
wget -q https://github.com/k2-fsa/sherpa-ncnn/releases/download/v2.1.7/sherpa-ncnn-wasm-
˓→simd-v2.1.7.tar.bz2
tar xvf sherpa-ncnn-wasm-simd-v2.1.7.tar.bz2
rm sherpa-ncnn-wasm-simd-v2.1.7.tar.bz2
cd sherpa-ncnn-wasm-simd-v2.1.7

python3 -m http.server 6006
```

Start your browser and visit <http://localhost:6006/>; you should see the following page:

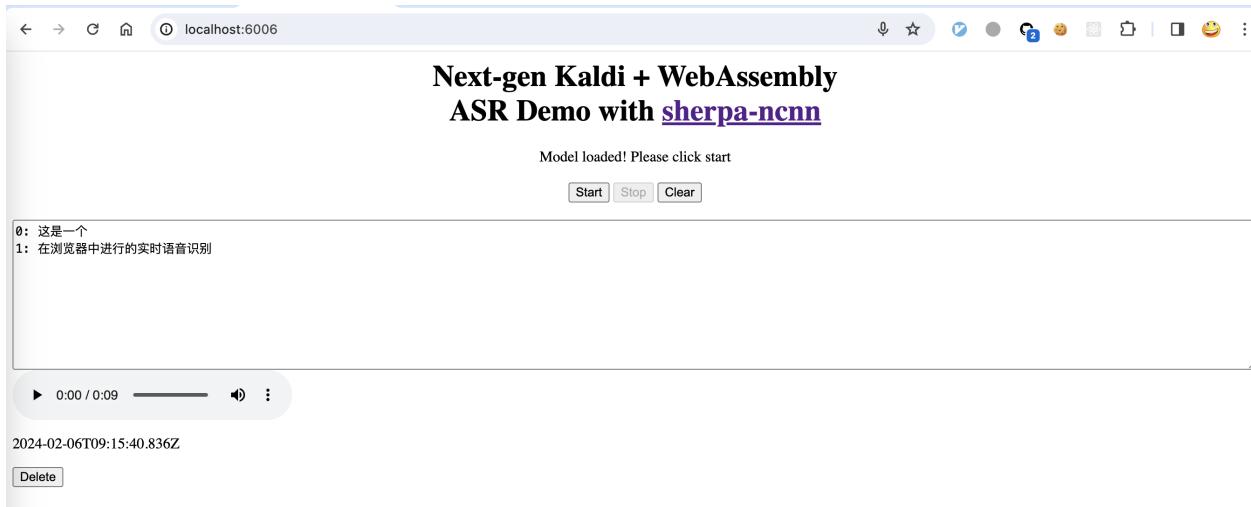


Now click start and speak! You should see the recognition results in the text box.

Warning: We are using a bilingual model (Chinese + English) in the above example, which means you can only speak Chinese or English in this case.

A screenshot is given below:

Congratulations! You have successfully run real-time speech recognition with WebAssembly in your browser.



7.4.4 Huggingface Spaces (WebAssembly)

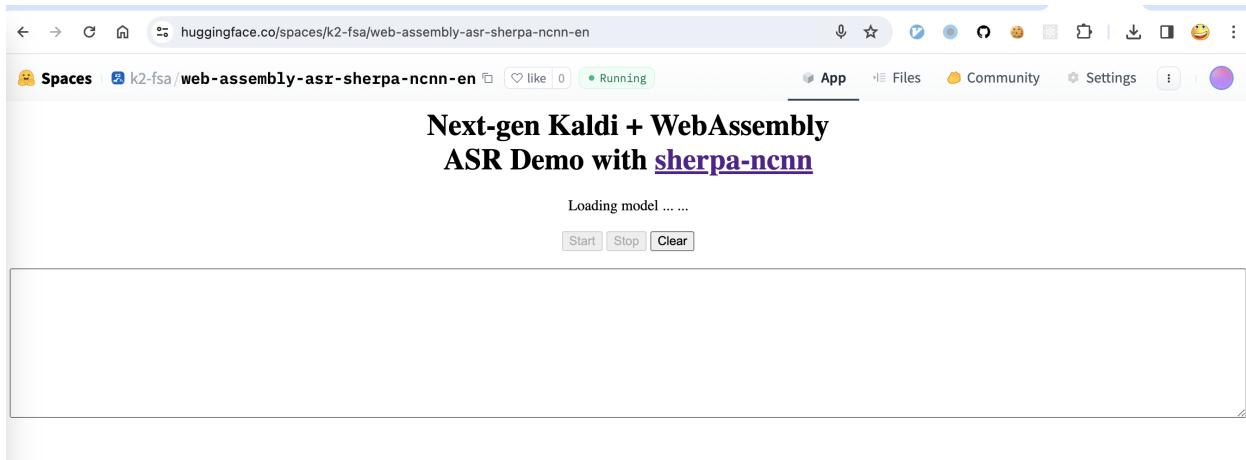
We provide two Huggingface spaces so that you can try real-time speech recognition with WebAssembly in your browser.

English only

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-ncnn-en>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-ncnn-en/summary>



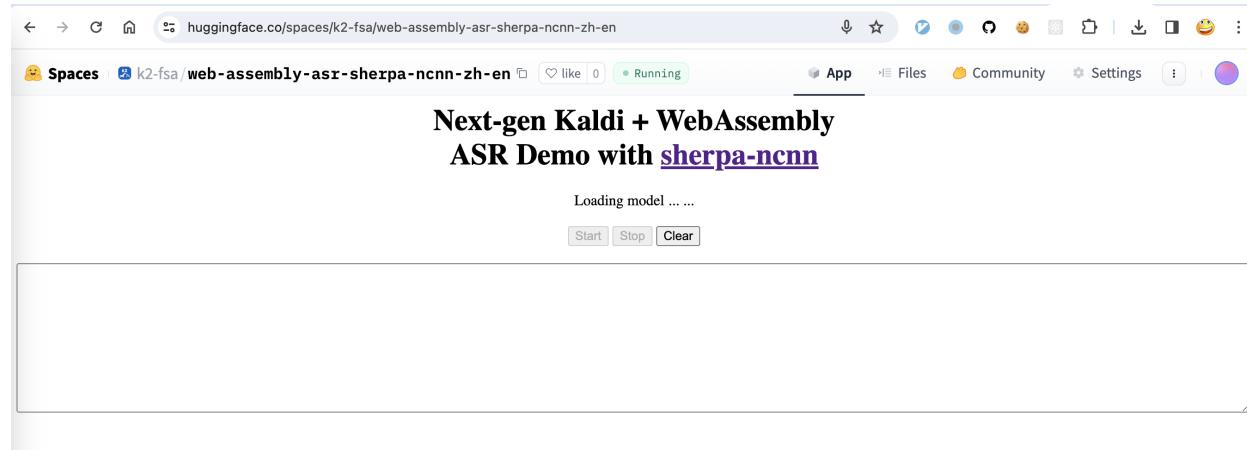
Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-ncnn/blob/master/.github/workflows/wasm-simd-hf-space-en.yaml>

Chinese + English

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-ncnn-zh-en>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-ncnn-zh-en/summary>



Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-ncnn/blob/master/.github/workflows/wasm-simd-hf-space-zh-en.yaml>

7.5 C API

In this section, we describe how to use the C API of `sherpa-ncnn`.

Specifically, we will describe:

- How to generate required files
- How to use `pkg-config` with `sherpa-ncnn`

7.5.1 Generate required files

Before using the C API of `sherpa-ncnn`, we need to first build required libraries. You can choose either to build static libraries or shared libraries.

Build shared libraries

Assume that we want to put library files and header files in the directory `/tmp/sherpa-ncnn/shared`:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build-shared
cd build-shared

cmake \
-DSHERPA_NCNN_ENABLE_C_API=ON \
-DCMAKE_BUILD_TYPE=Release \
-DBUILD_SHARED_LIBS=ON \
-DCMAKE_INSTALL_PREFIX=/tmp/sherpa-ncnn/shared \
..

make -j6
make install
```

You should find the following files inside `/tmp/sherpa-ncnn/shared`:

macOS

Linux

```
$ tree /tmp/sherpa-ncnn/shared/
/tmp/sherpa-ncnn/shared/
├── bin
│   └── sherpa-ncnn
│       └── sherpa-ncnn-microphone
├── include
│   └── sherpa-ncnn
│       └── c-api
│           └── c-api.h
└── lib
    ├── libkaldi-native-fbank-core.dylib
    ├── libncnn.dylib
    ├── libsherpa-ncnn-c-api.dylib
    └── libsherpa-ncnn-core.dylib
    sherpa-ncnn.pc

5 directories, 8 files
```

```
$ tree /tmp/sherpa-ncnn/shared/
/tmp/sherpa-ncnn/shared/
├── bin
│   └── sherpa-ncnn
│       └── sherpa-ncnn-microphone
├── include
│   └── sherpa-ncnn
│       └── c-api
│           └── c-api.h
└── lib
    └── libkaldi-native-fbank-core.so
```

(continues on next page)

(continued from previous page)

```
└── libncnn.so
    ├── libsherpa-ncnn-c-api.so
    └── libsherpa-ncnn-core.so
    └── sherpa-ncnn.pc
```

5 directories, 8 files

Build static libraries

Assume that we want to put library files and header files in the directory `/tmp/sherpa-ncnn/static`:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
cd sherpa-ncnn
mkdir build-static
cd build-static

cmake \
-DSHERPA_NCNN_ENABLE_C_API=ON \
-DCMAKE_BUILD_TYPE=Release \
-DBUILD_SHARED_LIBS=OFF \
-DCMAKE_INSTALL_PREFIX=/tmp/sherpa-ncnn/static \
..

make -j6
make install
```

You should find the following files in `/tmp/sherpa-ncnn/static`:

```
$ tree /tmp/sherpa-ncnn/static/
/tmp/sherpa-ncnn/static/
├── bin
│   └── sherpa-ncnn
│       └── sherpa-ncnn-microphone
├── include
│   └── sherpa-ncnn
│       └── c-api
│           └── c-api.h
└── lib
    ├── libkaldi-native-fbank-core.a
    ├── libncnn.a
    ├── libsherpa-ncnn-c-api.a
    └── libsherpa-ncnn-core.a
    └── sherpa-ncnn.pc
```

5 directories, 8 files

7.5.2 Build decode-file-c-api.c with generated files

To build the following file:

```
https://github.com/k2-fsa/sherpa-ncnn/blob/master/c-api-examples/decode-file-c-api.c
```

We can use:

static link

dynamic link

```
export PKG_CONFIG_PATH=/tmp/sherpa-ncnn/static:$PKG_CONFIG_PATH

cd ./c-api-examples
gcc -o decode-file-c-api $(pkg-config --cflags sherpa-ncnn) ./decode-file-c-api.c $(pkg-
config --libs sherpa-ncnn)
```

```
export PKG_CONFIG_PATH=/tmp/sherpa-ncnn/shared:$PKG_CONFIG_PATH

cd ./c-api-examples
gcc -o decode-file-c-api $(pkg-config --cflags sherpa-ncnn) ./decode-file-c-api.c $(pkg-
config --libs sherpa-ncnn)
```

7.6 Endpointing

We have three rules for endpoint detection. If any of them is activated, we assume an endpoint is detected.

Note: We borrow the implementation from

https://kaldi-asr.org/doc/structkaldi_1_1OnlineEndpointRule.html

7.6.1 Rule 1

In Rule 1, we count the duration of trailing silence. If it is larger than a user specified value, Rule 1 is activated. The following is an example, which uses 2.4 seconds as the threshold.

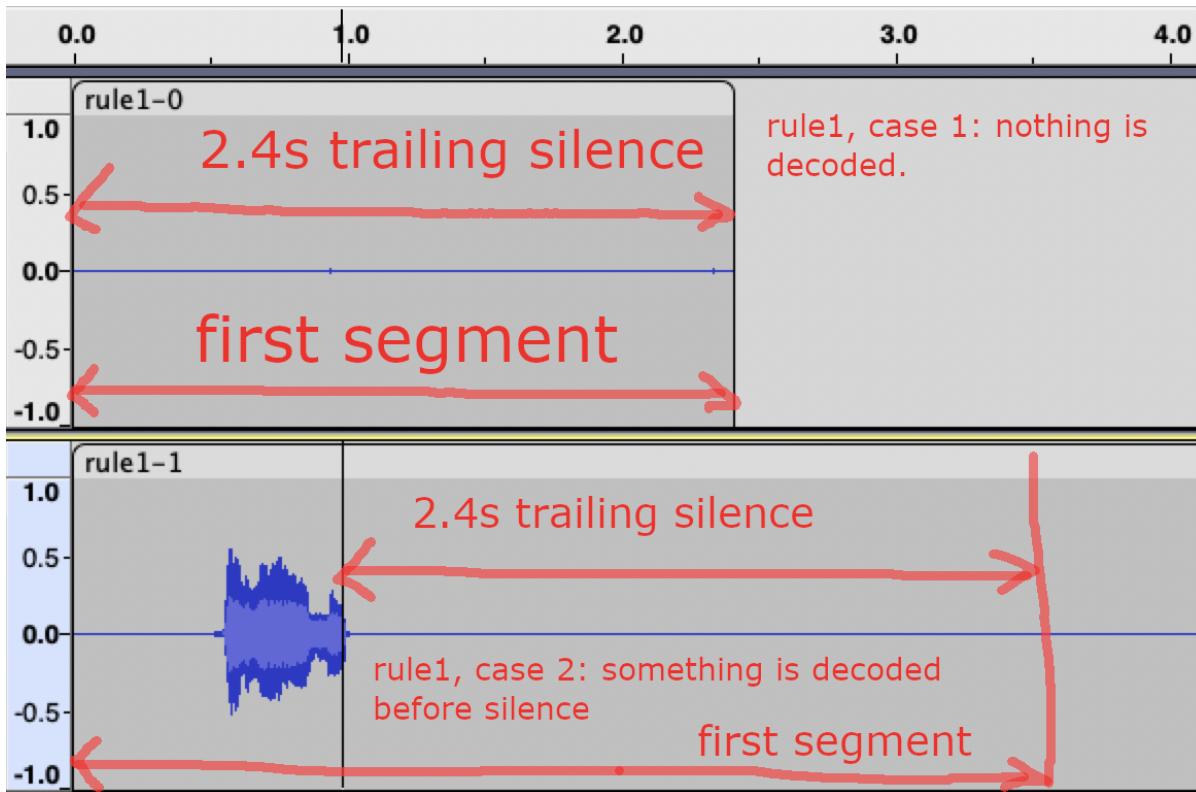
Two cases are given:

- (1) In the first case, nothing has been decoded when the duration of trailing silence reaches 2.4 seconds.
- (2) In the second case, we first decode something before the duration of trailing silence reaches 2.4 seconds.

In both cases, Rule 1 is activated.

Hint: In the Python API, you can specify `rule1_min_trailing_silence` while constructing an instance of `sherpa_ncnn.Recognizer`.

In the C++ API, you can specify `rule1.min_trailing_silence` when creating `EndpointConfig`.



7.6.2 Rule 2

In Rule 2, we require that it has to first decode something before we count the trailing silence. In the following example, after decoding something, Rule 2 is activated when the duration of trailing silence is larger than the user specified value 1.2 seconds.

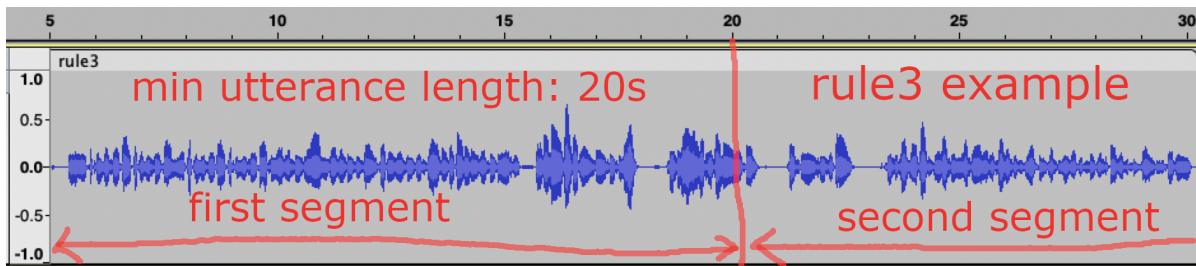


Hint: In the Python API, you can specify `rule2_min_trailing_silence` while constructing an instance of `sherpa_ncnn.Recognizer`.

In the C++ API, you can specify `rule2.min_trailing_silence` when creating `EndpointConfig`.

7.6.3 Rule 3

Rule 3 is activated when the utterance length in seconds is larger than a given value. In the following example, Rule 3 is activated after the first segment reaches a given value, which is 20 seconds in this case.



Hint: In the Python API, you can specify `rule3_min_utterance_length` while constructing an instance of `sherpa_ncnn.Recognizer`.

In the C++ API, you can specify `rule3.min_utterance_length` when creating `EndpointConfig`.

Note: If you want to deactivate this rule, please provide a very large value for `rule3_min_utterance_length` or `rule3.min_utterance_length`.

7.6.4 Demo

Multilingual (Chinese + English)

The following video demonstrates using the Python API of `sherpa-ncnn` for real-time speech recognition with endpointing.

Hint: The code is available at

<https://github.com/k2-fsa/sherpa-ncnn/blob/master/python-api-examples/speech-recognition-from-microphone-with-endpoint-detection.py>

7.6.5 FAQs

How to compute duration of silence

For each frame to be decoded, we can output either a blank or a non-blank token. We record the number of contiguous blanks that has been decoded so far. In the current default setting, each frame is 10 ms. Thus, we can get the duration of trailing silence by counting the number of contiguous trailing blanks.

Note: If a model uses a subsampling factor of 4, the time resolution becomes $10 * 4 = 40$ ms.

7.7 Android

In this section, we describe how to build an Android app for **real-time** speech recognition with `sherpa-ncnn`. We also provide real-time speech recognition video demos.

Hint: During speech recognition, it does not need to access the Internet. Everything is processed locally on your phone.

7.7.1 Video demos

In this page, we list some videos about using `sherpa-ncnn` for real-time speech recognition on Android.

Hint: You can find pre-built APK packages used by the following videos at:

<https://huggingface.co/csukuangfj/sherpa-ncnn-apk/tree/main>

- CPU versions require Android ≥ 5.0
 - GPU versions with Vulkan require Android ≥ 7.0
-

Note: You can also find the latest APK for each release at

<https://github.com/k2-fsa/sherpa-ncnn/releases>

Video 1: Chinese

Video 2: Chinese + English

Video 3: Chinese with background noise

Video 4: Chinese poem with background music

7.7.2 Build `sherpa-ncnn` for Android

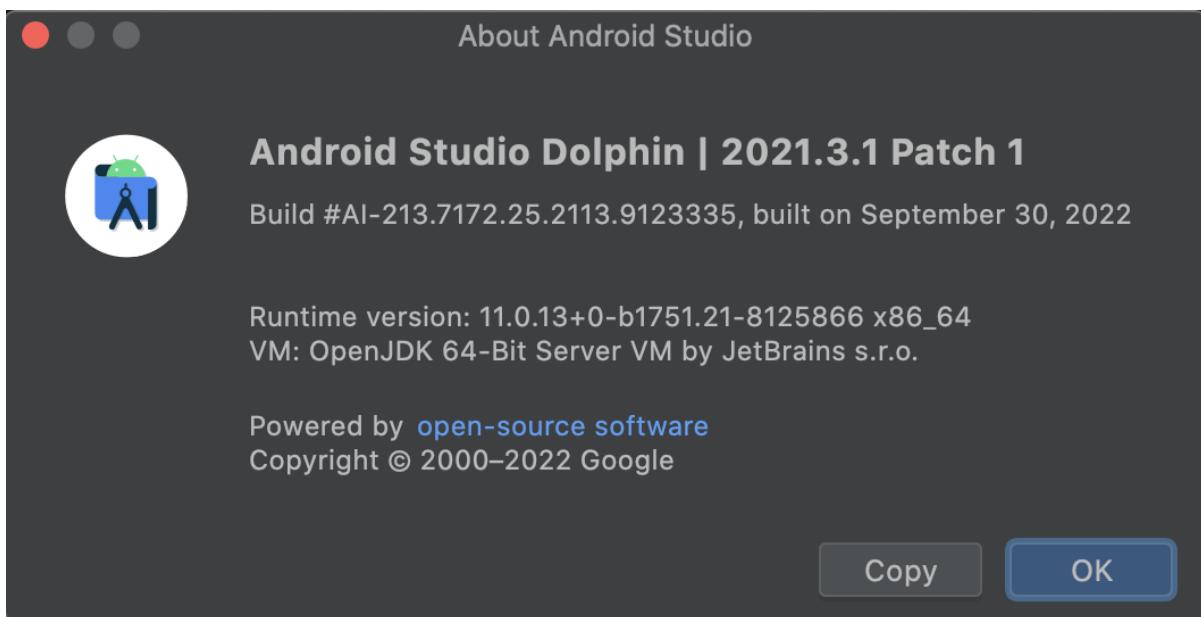
Install Android Studio

The first step is to download and install Android Studio.

Please refer to <https://developer.android.com/studio> for how to install Android Studio.

Hint: Any recent version of Android Studio should work fine. Also, you can use the default settings of Android Studio during installation.

For reference, we post the version we are using below:



Download sherpa-ncnn

Next, download the source code of sherpa-ncnn:

```
git clone https://github.com/k2-fsa/sherpa-ncnn
```

Install NDK

Step 1, start Android Studio.

Step 2, Open `sherpa-ncnn/android/SherpaNcnn`.

Step 3, Select Tools -> SDK Manager.

Step 4, Install NDK.

In the following, we assume Android SDK location was set to `/Users/fangjun/software/my-android`. You can change it accordingly below.

After installing NDK, you can find it in

```
/Users/fangjun/software/my-android/ndk/22.1.7171670
```

Warning: If you selected a different version of NDK, please replace `22.1.7171670` accordingly.

Next, let us set the environment variable `ANDROID_NDK` for later use.

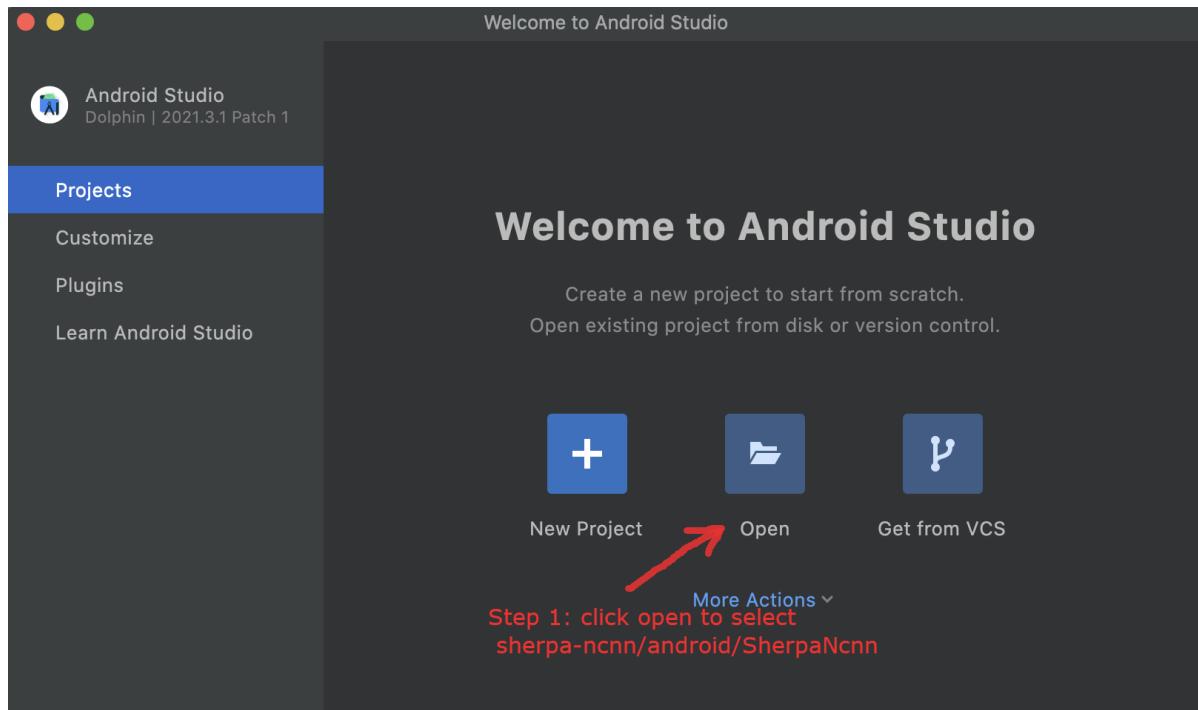


Fig. 7.1: Step 1: Click Open to select sherpa-ncnn/android/SherpaNcnn




Fig. 7.2: Step 2: Open SherpaNcnn.




Fig. 7.3: Step 3: Select Tools -> SDK Manager.

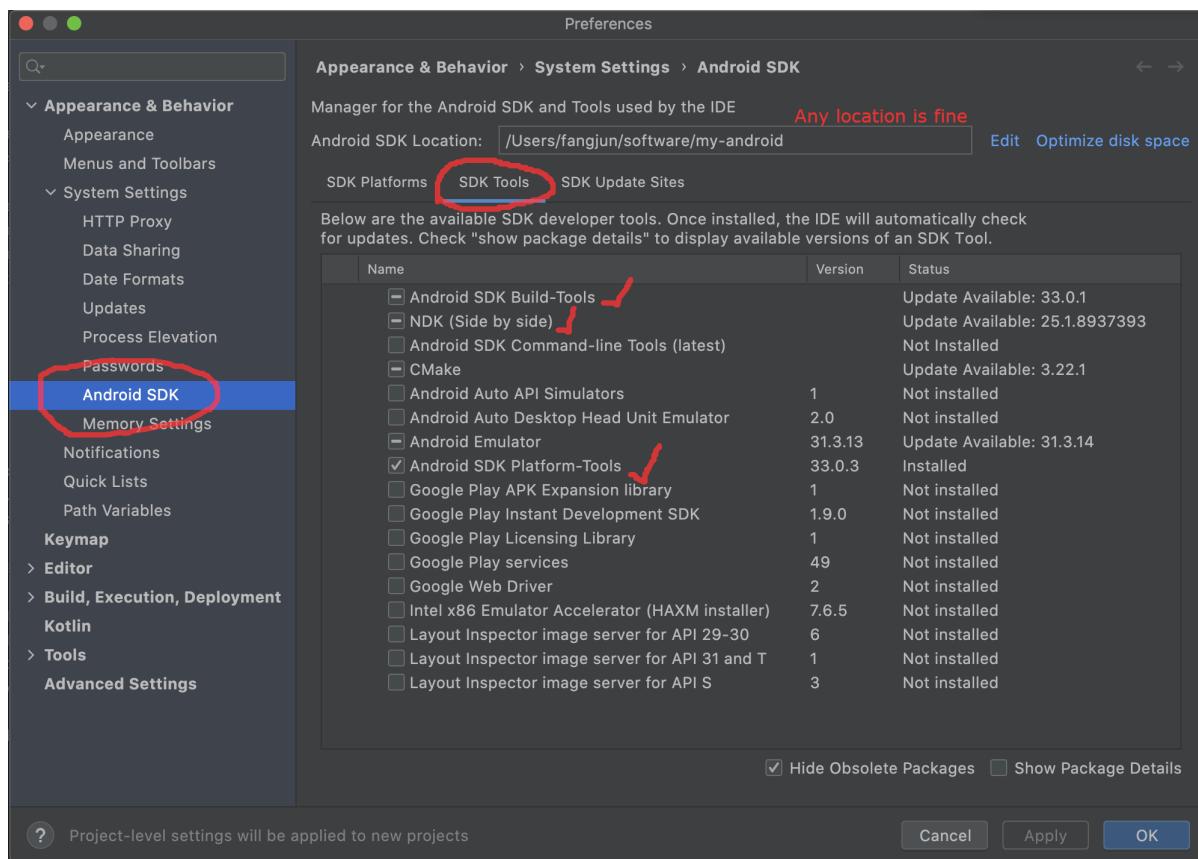


Fig. 7.4: Step 4: Install NDK.

```
export ANDROID_NDK=/Users/fangjun/software/my-android/ndk/22.1.7171670
```

Note: Note from <https://github.com/Tencent/ncnn/wiki/how-to-build#build-for-android>

(Important) remove the hardcoded debug flag in Android NDK to fix the android-ndk issue: <https://github.com/android/ndk/issues/243>

1. open \$ANDROID_NDK/build/cmake/android.toolchain.cmake for ndk < r23 or \$ANDROID_NDK/build/cmake/android-legacy.toolchain.cmake for ndk >= r23

2. delete the line containing “-g”

```
list(APPEND ANDROID_COMPILER_FLAGS
-g
-DANDROID
```

Caution: If you don't delete the line containin -g above, the generated library libncnn.so can be as large as 21 MB or even larger!

Build sherpa-ncnn (C++ code)

After installing NDK, it is time to build the C++ code of sherpa-ncnn.

In the following, we show how to build `sherpa-ncnn` for the following Android ABIs:

- arm64-v8a
- armeabi-v7a
- x86_64
- x86

Caution: You only need to select one and only one ABI. `arm64-v8a` is probably the most common one.

If you want to test the app on an emulator, you probably need `x86_64`.

Hint: Building scripts for this section are for macOS and Linux. If you are using Windows or if you don't want to build the shared libraries by yourself, you can download pre-compiled shared libraries for this section by visiting

<https://github.com/k2-fsa/sherpa-ncnn/releases>

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows or you don't want to setup your local environment to build the C++ libraries, please use the above colab notebook.

Build for arm64-v8a

```
cd sherpa-ncnn # Go to the root repo  
./build-android-arm64-v8a.sh
```

After building, you will find the following shared libraries:

```
$ ls -lh build-android-arm64-v8a/install/lib/lib*.so  
-rwxr-xr-x 1 fangjun staff 848K Dec 18 16:49 build-android-arm64-v8a/install/lib/  
↳ libkaldi-native-fbank-core.so  
-rwxr-xr-x 1 fangjun staff 3.4M Dec 18 16:49 build-android-arm64-v8a/install/lib/  
↳ libncnn.so  
-rwxr-xr-x 1 fangjun staff 195K Dec 18 16:49 build-android-arm64-v8a/install/lib/  
↳ libsherpa-ncnn-core.so  
-rwxr-xr-x 1 fangjun staff 19K Dec 18 16:49 build-android-arm64-v8a/install/lib/  
↳ libsherpa-ncnn-jni.so
```

Please copy them to android/SherpaNcnn/app/src/main/jniLibs/arm64-v8a/:

```
$ cp build-android-arm64-v8a/install/lib/lib*.so android/SherpaNcnn/app/src/main/  
↳ jniLibs/arm64-v8a/
```

You should see the following screen shot after running the above copy cp command.

Note: If you have Android >= 7.0 and want to run `sherpa-ncnn` on GPU, please replace `./build-android-arm64-v8a.sh` with `./build-android-arm64-v8a-with-vulkan.sh` and replace `build-android-arm64-v8a/install/lib/lib*.so` with `./build-android-arm64-v8a-with-vulkan/install/lib/lib*.so`. That is all you need to do and you don't need to change any code.

Also, you need to install Vulkan sdk. Please see <https://github.com/k2-fsa/sherpa-ncnn/blob/master/install-vulkan-macos.md> for details.

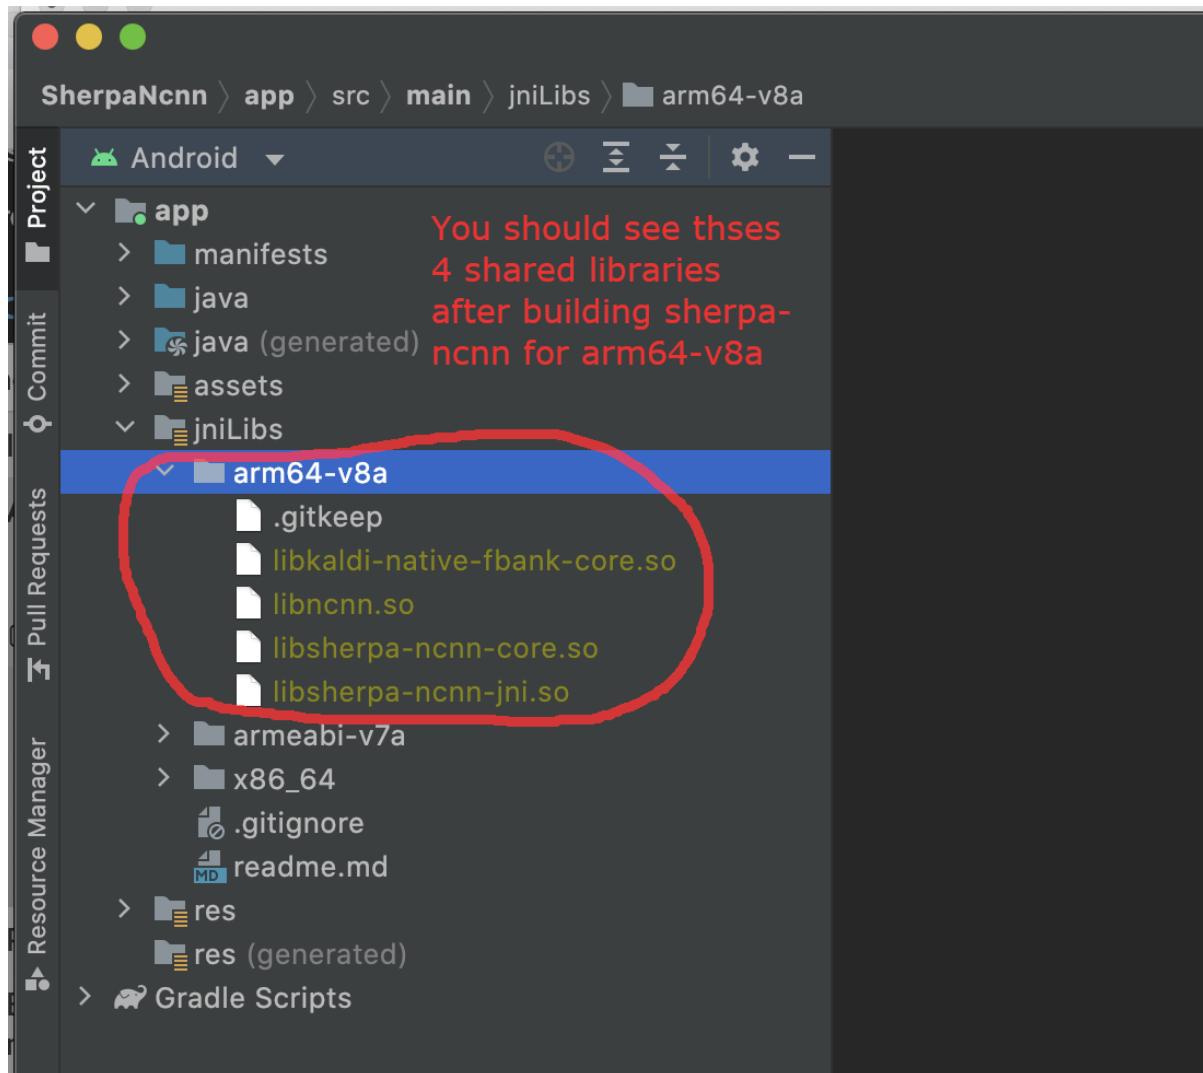
Build for armeabi-v7a

```
cd sherpa-ncnn # Go to the root repo  
./build-android-armv7-eabi.sh
```

After building, you will find the following shared libraries:

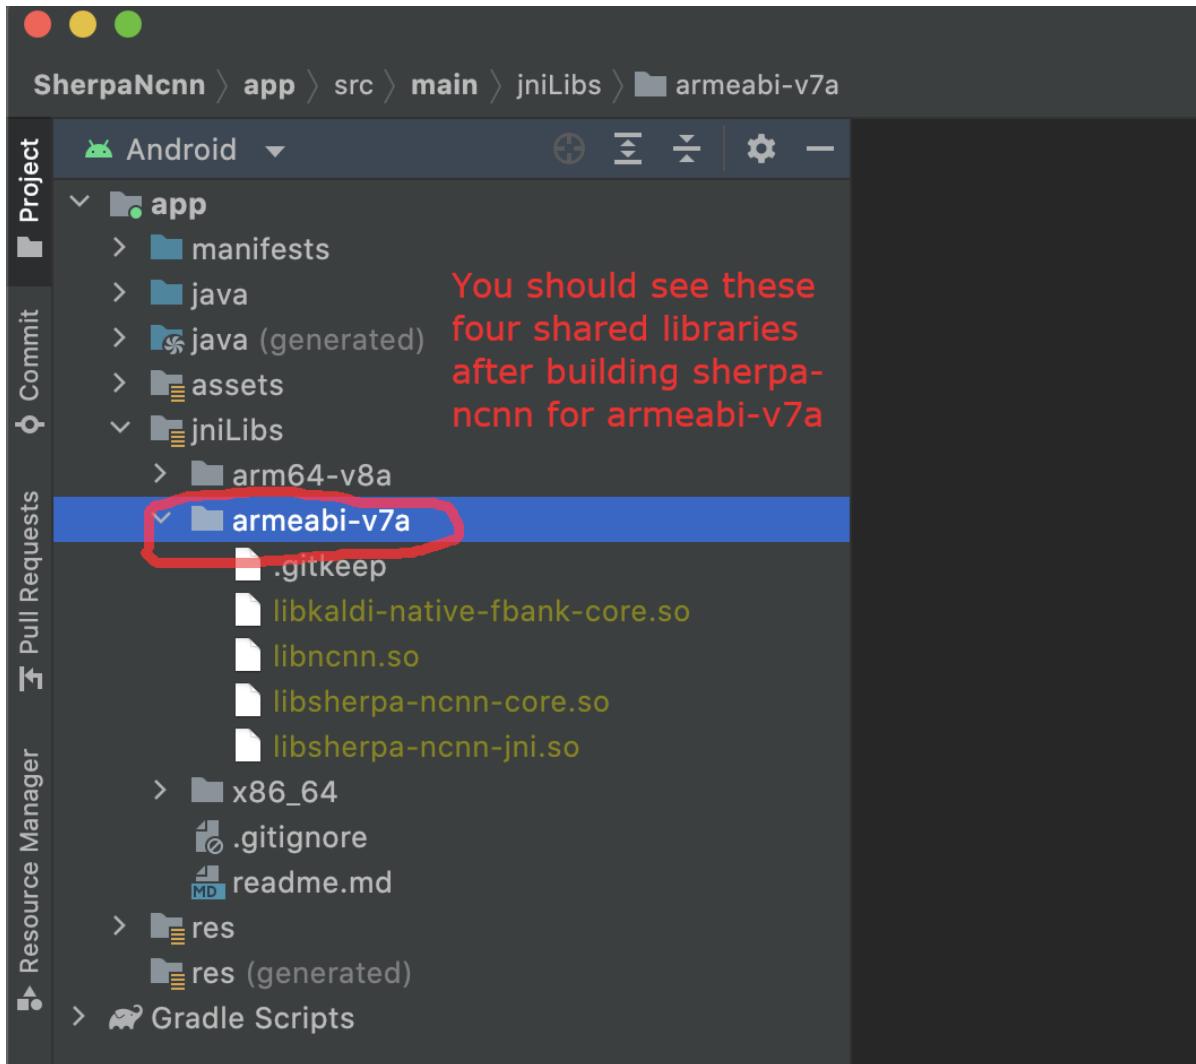
```
$ ls -lh build-android-armv7-eabi/install/lib/lib*.so  
-rwxr-xr-x 1 fangjun staff 513K Dec 18 17:04 build-android-armv7-eabi/install/lib/  
↳ libkaldi-native-fbank-core.so  
-rwxr-xr-x 1 fangjun staff 1.9M Dec 18 17:04 build-android-armv7-eabi/install/lib/  
↳ libncnn.so  
-rwxr-xr-x 1 fangjun staff 163K Dec 18 17:04 build-android-armv7-eabi/install/lib/  
↳ libsherpa-ncnn-core.so  
-rwxr-xr-x 1 fangjun staff 28K Dec 18 17:04 build-android-armv7-eabi/install/lib/  
↳ libsherpa-ncnn-jni.so
```

Please copy them to android/SherpaNcnn/app/src/main/jniLibs/armeabi-v7a/:



```
cp build-android-armv7-eabi/install/lib/lib*.so android/SherpaNcnn/app/src/main/jniLibs/  
↳ armeabi-v7a/
```

You should see the following screen shot after running the above copy cp command.



Build for x86_64

```
cd sherpa-ncnn # Go to the root repo  
./build-android-x86-64.sh
```

After building, you will find the following shared libraries:

```
$ ls -lh build-android-x86-64/install/lib/lib*.so  
-rwxr-xr-x 1 fangjun staff 901K Dec 18 17:14 build-android-x86-64/install/lib/  
↳ libkaldi-native-fbank-core.so  
-rwxr-xr-x 1 fangjun staff 6.9M Dec 18 17:14 build-android-x86-64/install/lib/  
↳ libncnn.so
```

(continues on next page)

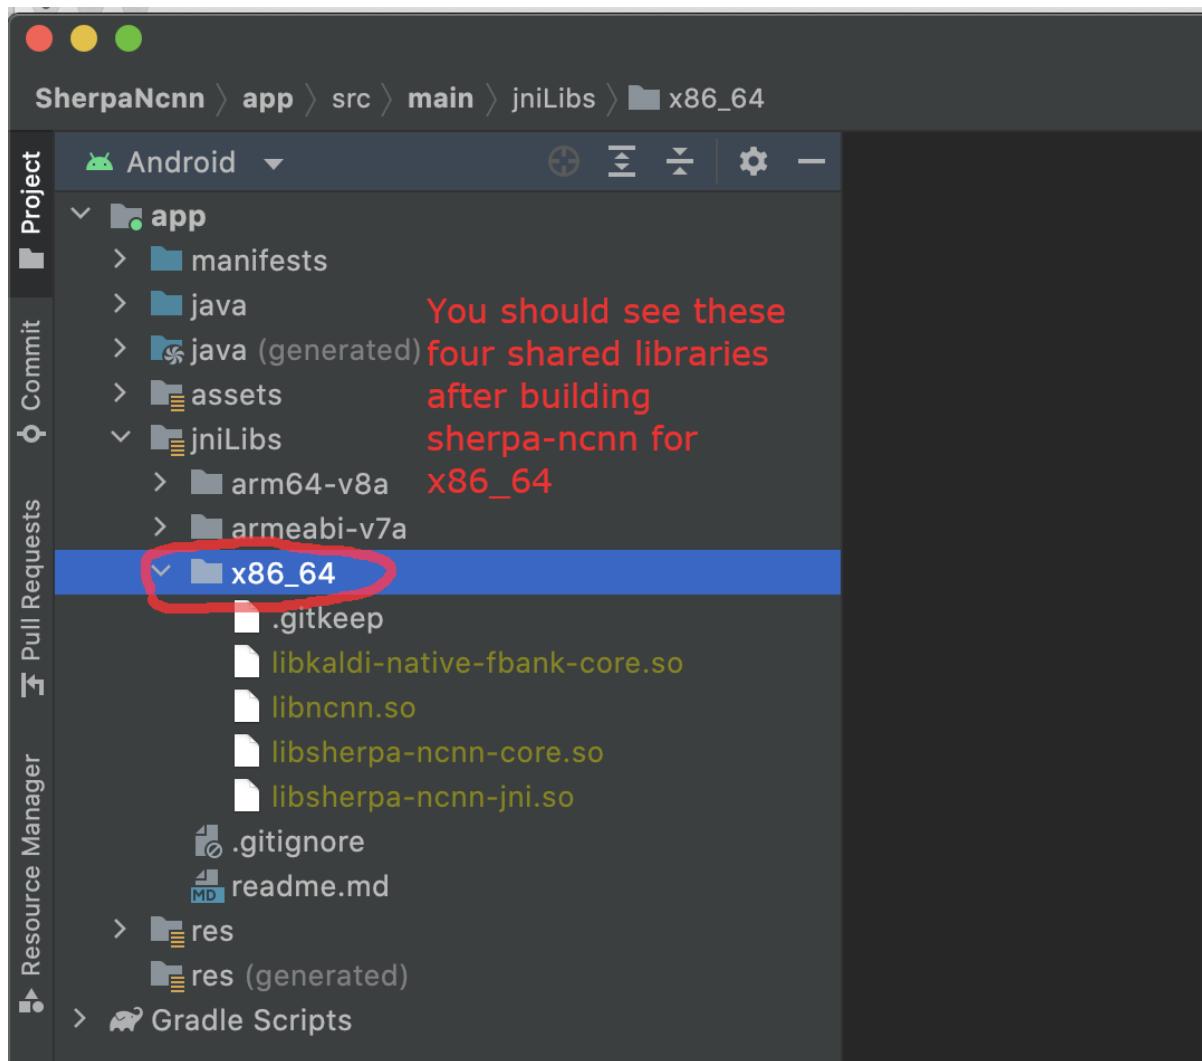
(continued from previous page)

```
-rwxr-xr-x 1 fangjun staff 208K Dec 18 17:14 build-android-x86-64/install/lib/
  ↳ libsherpa-ncnn-core.so
-rwxr-xr-x 1 fangjun staff 19K Dec 18 17:14 build-android-x86-64/install/lib/
  ↳ libsherpa-ncnn-jni.so
```

Please copy them to android/SherpaNcnn/app/src/main/jniLibs/x86_64/:

```
cp build-android-x86-64/install/lib/lib*.so android/SherpaNcnn/app/src/main/jniLibs/x86_
  ↳ 64/
```

You should see the following screen shot after running the above copy cp command.



Build for x86

```
cd sherpa-ncnn # Go to the root repo  
./build-android-x86.sh
```

Download pre-trained models

Please read *Pre-trained models* for all available pre-trained models.

In the following, we use a pre-trained model from <https://huggingface.co/csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06>, which supports both Chinese and English.

Hint: The model is trained using icefall and the original torchscript model is from <https://huggingface.co/ptrnnull/icefall-asr-conv-emformer-transducer-stateless2-zh>.

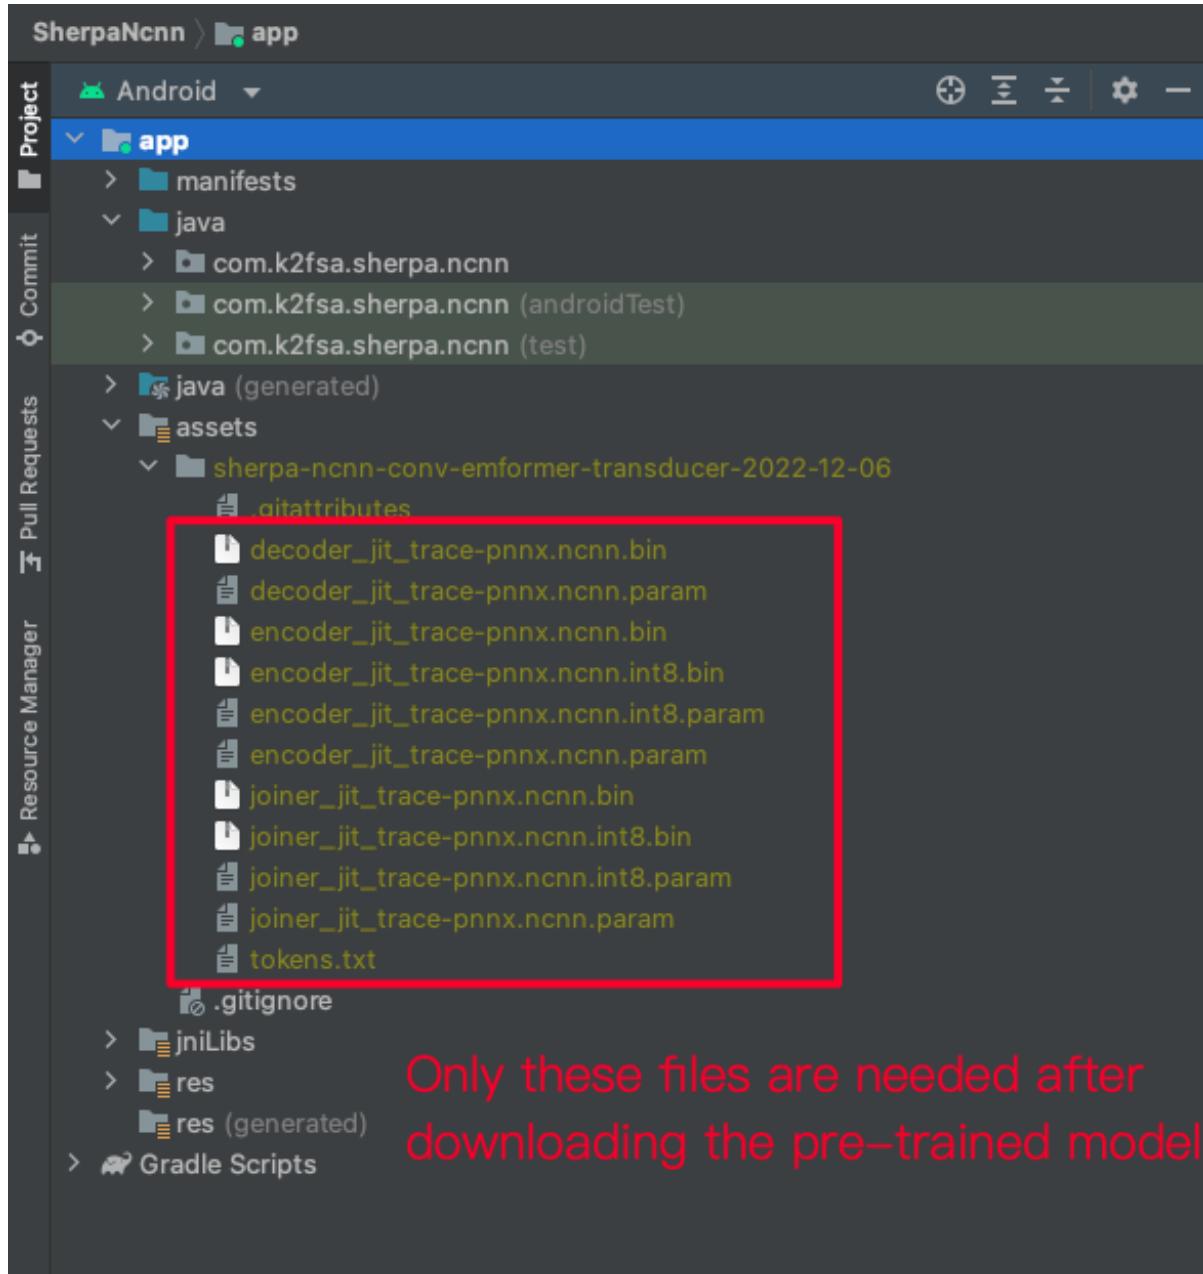
Use the following command to download the pre-trained model and place it into android/SherpaNcnn/app/src/main/assets/:

```
cd android/SherpaNcnn/app/src/main/assets/  
  
sudo apt-get install git-lfs  
  
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/sherpa-ncnn-conv-  
~emformer-transducer-2022-12-06  
cd sherpa-ncnn-conv-emformer-transducer-2022-12-06  
git lfs pull --include "*.bin"  
  
# Now, remove extra files to reduce the file size of the generated apk  
rm -rf .git test_wavs scripts/  
rm export-for-ncnn.sh *.png README.md
```

In the end, you should have the following files:

```
$ ls -lh  
total 525224  
-rw-r--r-- 1 fangjun staff 5.9M Dec 18 17:40 decoder_jit_trace-pnnx.ncnn.bin  
-rw-r--r-- 1 fangjun staff 439B Dec 18 17:39 decoder_jit_trace-pnnx.ncnn.param  
-rw-r--r-- 1 fangjun staff 141M Dec 18 17:40 encoder_jit_trace-pnnx.ncnn.bin  
-rw-r--r-- 1 fangjun staff 99M Dec 18 17:40 encoder_jit_trace-pnnx.ncnn.int8.bin  
-rw-r--r-- 1 fangjun staff 78K Dec 18 17:40 encoder_jit_trace-pnnx.ncnn.int8.param  
-rw-r--r-- 1 fangjun staff 79K Dec 18 17:39 encoder_jit_trace-pnnx.ncnn.param  
-rw-r--r-- 1 fangjun staff 6.9M Dec 18 17:40 joiner_jit_trace-pnnx.ncnn.bin  
-rw-r--r-- 1 fangjun staff 3.5M Dec 18 17:40 joiner_jit_trace-pnnx.ncnn.int8.bin  
-rw-r--r-- 1 fangjun staff 498B Dec 18 17:40 joiner_jit_trace-pnnx.ncnn.int8.param  
-rw-r--r-- 1 fangjun staff 490B Dec 18 17:39 joiner_jit_trace-pnnx.ncnn.param  
-rw-r--r-- 1 fangjun staff 53K Dec 18 17:39 tokens.txt  
  
$ du -h -d1 .  
256M .
```

You should see the following screen shot after downloading the pre-trained model:



Hint: If you select a different pre-trained model, make sure that you also change the corresponding code listed in the following screen shot:

```
val numBytes = AudioRecord.getMinBufferSize(sampleRateInHz, channelConfig, audioFormat)
Log.i(
    TAG,
    msg: "Buffer size in milliseconds: ${numBytes * 1000.0f / sampleRateInHz}"
)

audioRecord = AudioRecord(
    audioSource,
    sampleRateInHz,
    channelConfig,
    audioFormat,
    numBytes
) // numBytes * 2 // a sample has two bytes as we are using 16-bit PCM
return true
}

private fun initModel() {
    val featureConfig = getFeatureExtractorConfig(
        sampleRate = 16000.0f,
        featureDim = 80
    )
    //Please change the argument "type" if you use a different model.
    val modelConfig = getModelConfig(type = 1, useGPU = false)!!
    val decoderConfig = getDecoderConfig(decoderType = "greedy_search", numActivePaths = 4)
    val config = RecognizerConfig(
        featureConfig = featureConfig,
        modelConfig = modelConfig,
        decoderConfig = decoderConfig,
        enableEndpoint = true,
        ruleMinTrailingSilence = 2.0f,
        ruleMinIntraSilence = 0.8f,
        ruleMinUtteranceLength = 20.0f,
    )
    model = SherpaNcnn(
        assetManager = application.assets,
        config = config,
    )
}
}
```

Generate APK

Finally, it is time to build `sherpa-ncnn` to generate an APK package.

Select Build -> Make Project, as shown in the following screen shot.

You can find the generated APK in `android/SherpaNcnn/app/build/outputs/apk/debug/app-debug.apk`:

```
$ ls -lh android/SherpaNcnn/app/build/outputs/apk/debug/app-debug.apk
-rw-r--r-- 1 fangjun staff 152M Dec 18 17:53 android/SherpaNcnn/app/build/outputs/
apk/debug/app-debug.apk
```

Congratulations! You have successfully built an APK for Android.

Read below to learn more.

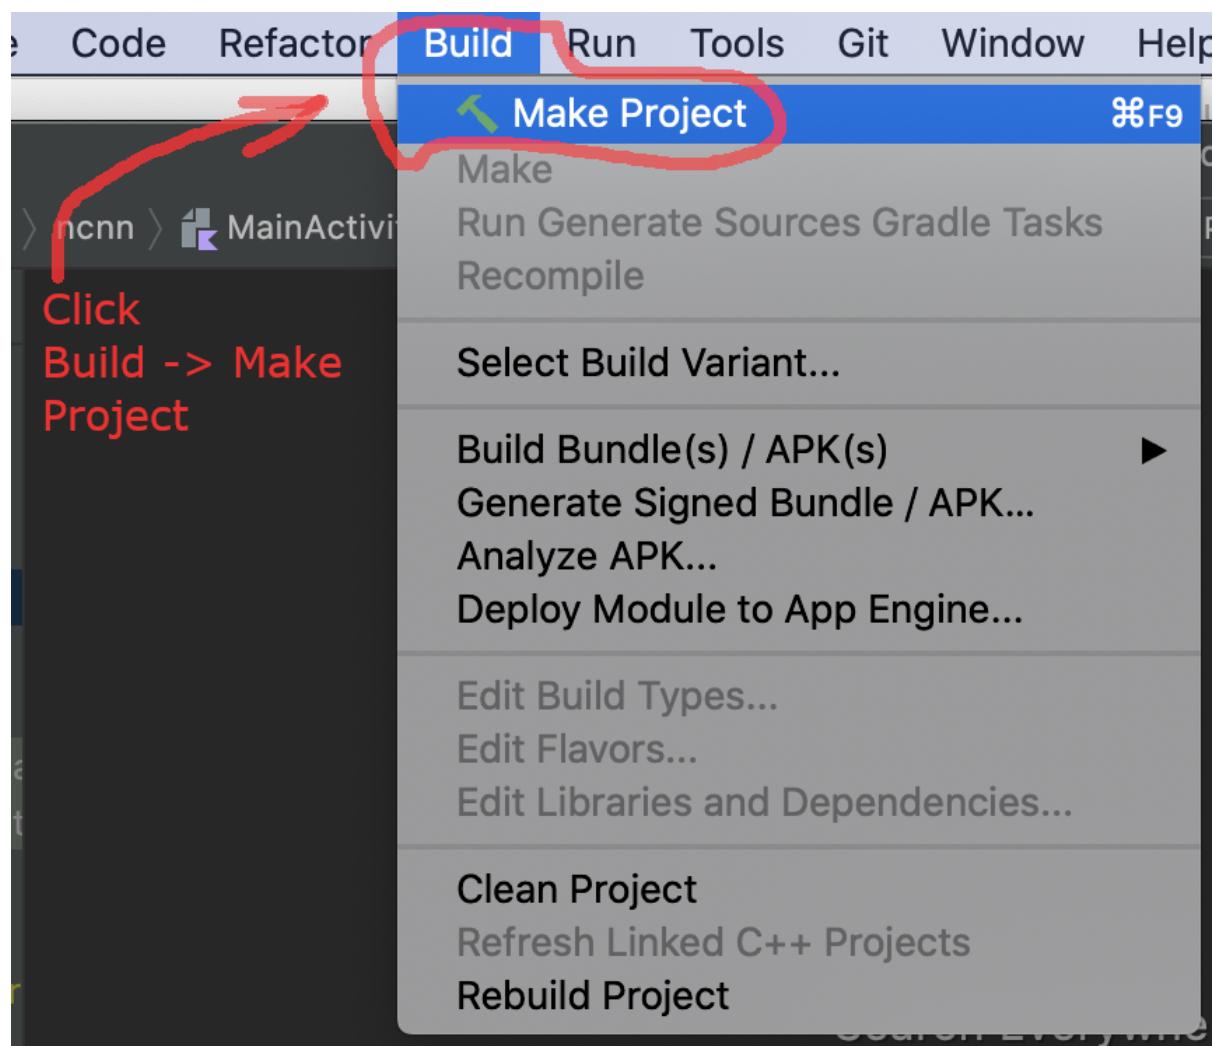
Analyze the APK

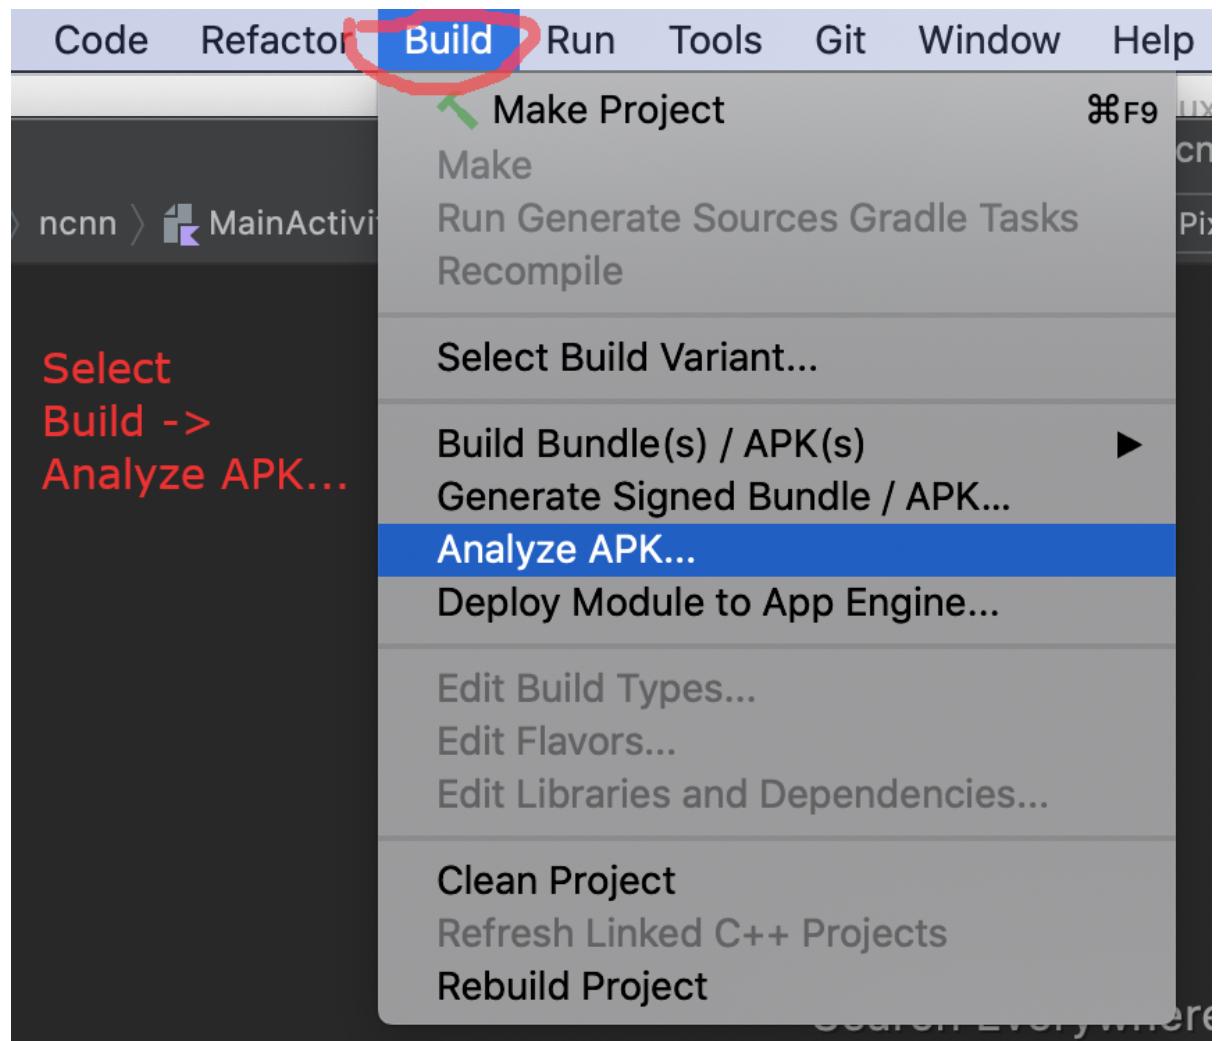
Select Build -> Analyze APK ... in the above screen shot, in the popped-up dialog select the generated APK `app-debug.apk`, and you will see the following screen shot:


You can see from the above screen shot that most part of the APK is occupied by the pre-trained model, while the runtime, including the shared libraries, is only 1.7 MB.

Hint: We have pre-built APKs that can be downloaded from <https://huggingface.co/csukuangfj/sherpa-ncnn-apk>

Please refer to demo videos about using the above APKs: [Video demos](#).







7.8 iOS

In this section, we describe how to build an iOS app for real-time speech recognition with `sherpa-ncnn` and run it within a simulator on your Mac, run it on you iPhone or iPad.

We also provide video demos for real-time speech recognition.

Hint: During speech recognition, it does not need to access the Internet. Everyting is processed locally on your iPhone or iPad.

7.8.1 Video demos

In this page, we list some videos about using `sherpa-ncnn` for real-time speech recognition on iPhone and iPad.

Video 1: Chinese + English on iPhone 14 Pro (simulator)

Video 2: Chinese + English on iPad 11 Pro (simulator)

7.8.2 Build `sherpa-ncnn` for iOS

This section describes how to build `sherpa-ncnn` for iPhone and iPad.

Requirement

Warning: The minimum deployment requires the iOS version >= 13.0.

Before we continue, please make sure the following requirements are satisfied:

- macOS. It won't work on Windows or Linux.
- Xcode. The version 14.2 (14C18) is known to work. Other versions may also work.
- CMake. CMake 3.25.1 is known to work. Other versions may also work.
- (Optional) iPhone or iPad. This is for testing the app on your device. If you don't have a device, you can still run the app within a simulator on your Mac.

Caution:

If you get the following error:

```
CMake Error at toolchains/ios.toolchain.cmake:544 (get_filename_component):
  get_filename_component called with incorrect number of arguments
Call Stack (most recent call first):
  /usr/local/Cellar/cmake/3.29.0/share/cmake/Modules/CMakeDetermineSystem.
  ↵cmake:146 (include)
  CMakeLists.txt:2 (project)
```

please run:

```
sudo xcode-select --install
sudo xcodebuild -license
```

And then delete the build directory ./build-ios and re-build.

Please see also <https://github.com/k2-fsa/sherpa-onnx/issues/702>.

Download sherpa-ncnn

First, let us download the source code of **sherpa-ncnn**.

Note: In the following, I will download **sherpa-ncnn** to **\$HOME/open-source**, i.e., **/Users/fangjun/open-source**, on my Mac.

You can put it anywhere as you like.

```
mkdir -p $HOME/open-source
cd $HOME/open-source
git clone https://github.com/k2-fsa/sherpa-ncnn
```

Build sherpa-ncnn (in commandline, C++ Part)

After downloading `sherpa-ncnn`, let us build the C++ part of `sherpa-ncnn`.

```
cd $HOME/open-source/sherpa-ncnn/
./build-ios.sh
```

It will generate a directory `$HOME/open-source/sherpa-ncnn/build-ios`, which we have already pre-configured for you in Xcode.

Hint: You don't have to look at the generated files in `$HOME/open-source/sherpa-ncnn/build-ios` to build an app. We have pre-configured it for you.

If you are eager to learn more about the generated files or want to use `sherpa-ncnn` in your own iOS project, please have a look at [For the more curious](#).

Build sherpa-ncnn (in Xcode)

Use the following command to open `sherpa-ncnn` in Xcode:

```
cd $HOME/open-source/sherpa-ncnn/ios-swift/SherpaNcnn
open SherpaNcnn.xcodeproj
```

It will start Xcode and you will see the following screenshot:




Fig. 7.5: Screenshot after running the command `open SherpaNcnn.xcodeproj`

Please select Product -> Build to build the project. See the screenshot below:

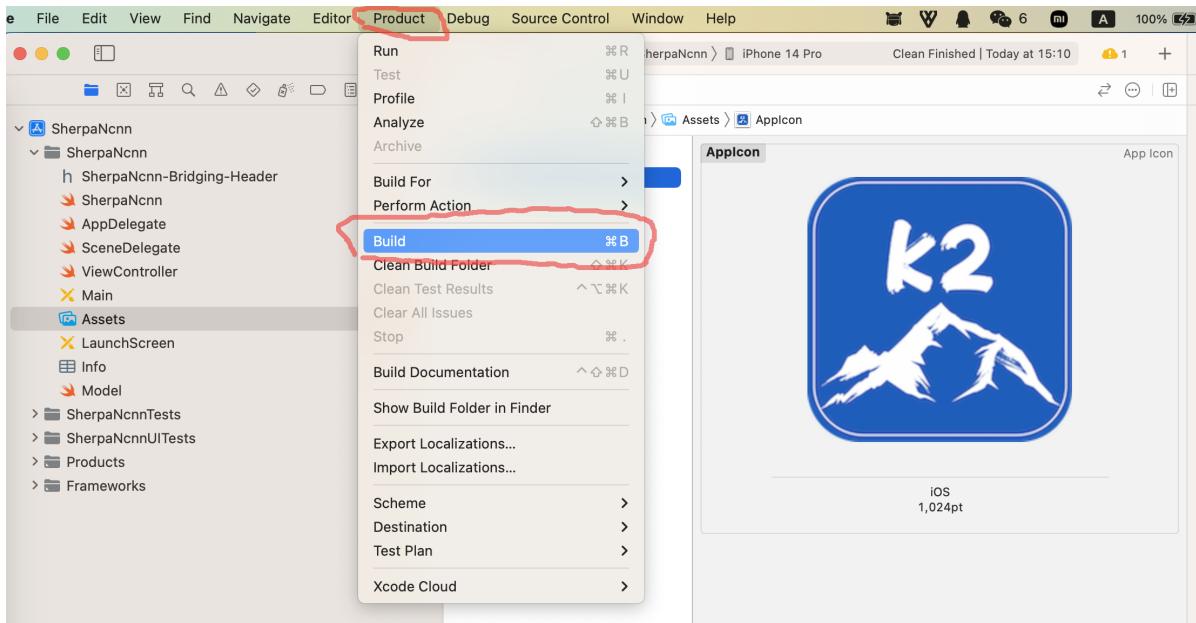


Fig. 7.6: Screenshot for selecting Product -> Build

After finishing the build, you should see the following screenshot:



Fig. 7.7: Screenshot after finishing the build.

Congratulations! You have successfully built the project. Let us run the project by selecting Product -> Run, which is shown in the following screenshot:

Please wait for a few seconds before Xcode starts the simulator.

Unfortunately, it will throw the following error:

The reason for the above error is that we have not provided the pre-trained model yet.

The file `ViewController.swift` pre-selects the pre-trained model to be `csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06 (Chinese + English)`, shown in the screenshot below:

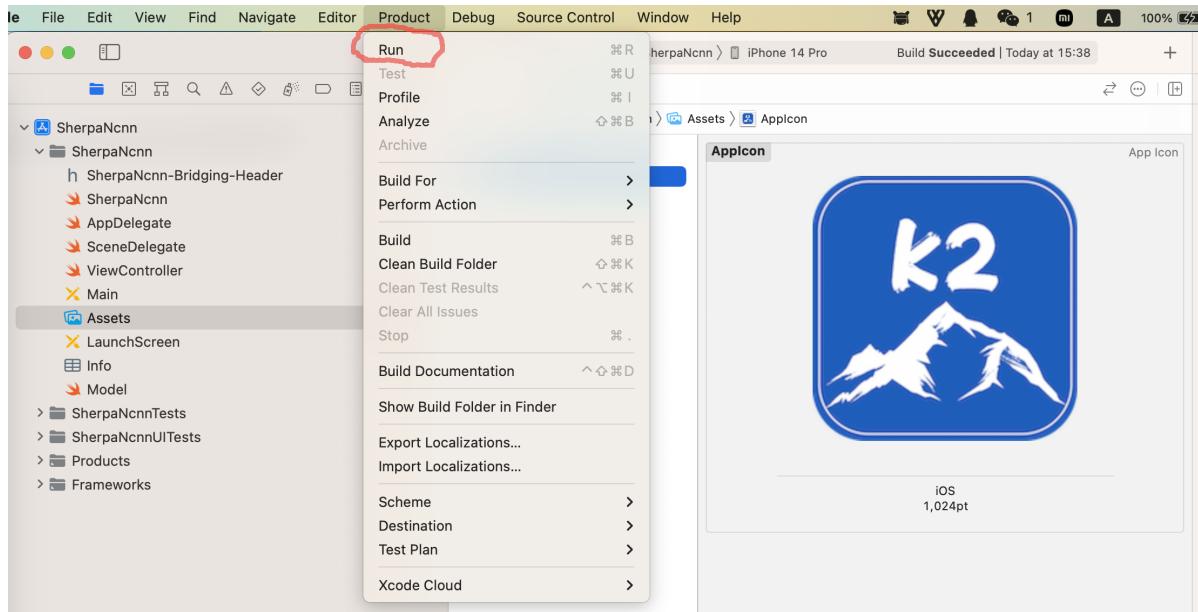


Fig. 7.8: Screenshot for Product -> Run.

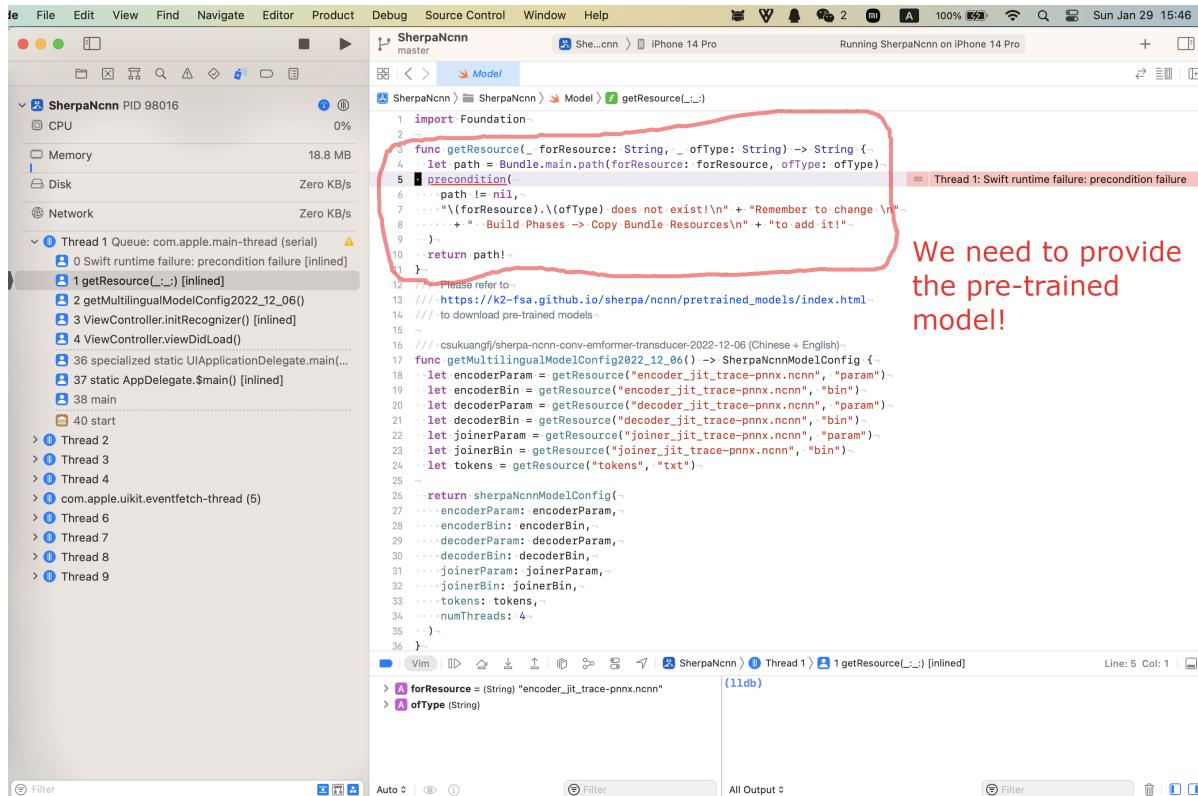


Fig. 7.9: Screenshot for the error

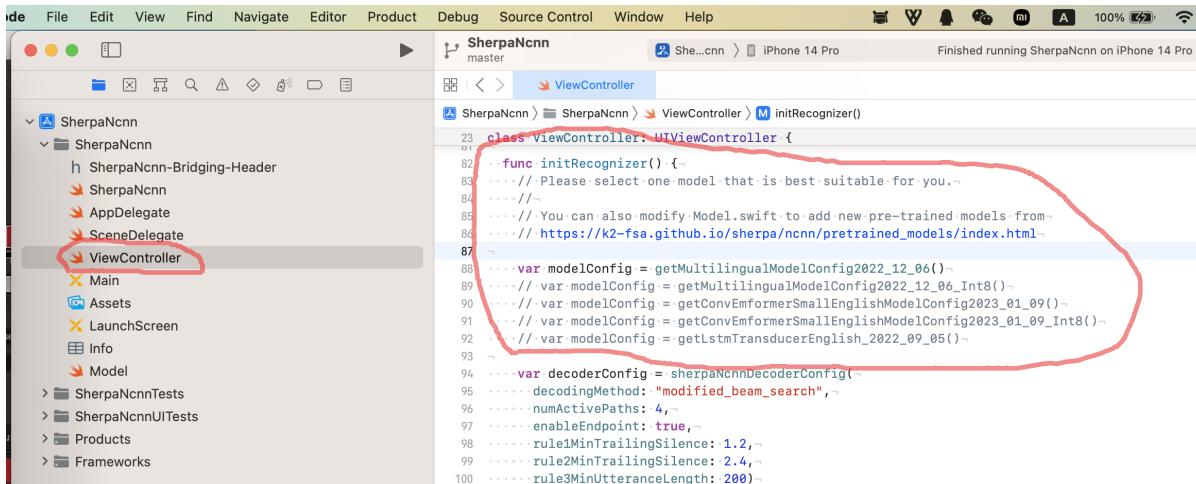


Fig. 7.10: Screenshot for the pre-selected pre-trained model

Let us add the pre-trained model [csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06 \(Chinese + English\)](https://huggingface.co/csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06) to Xcode. Please follow [csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06 \(Chinese + English\)](https://huggingface.co/csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06) to download it from [huggingface](#). You can download it to any directory as you like.

Please right click the project SherpaNcnn and select Add Files to "SherpaNcnn"... in the popup menu, as is shown in the screenshot below:




Fig. 7.11: Screenshot for adding files to SherpaNcnn

In the popup dialog, switch to the folder where you just downloaded the pre-trained model.

In the screenshot below, it is the folder /Users/fangjun/open-source/icefall-models/sherpa-ncnn-conv-emformer-transducer-2022-12-06:

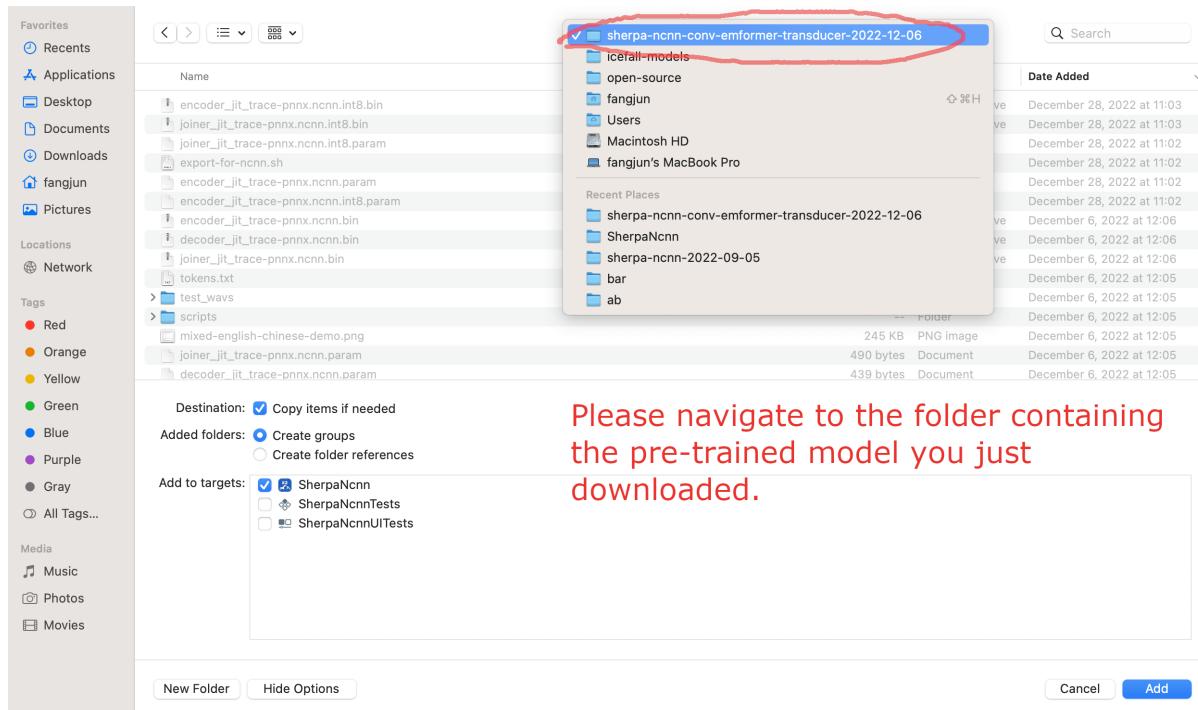


Fig. 7.12: Screenshot for navigating to the folder containing the downloaded pre-trained

Select required files and click the button Add:

After adding pre-trained model files to Xcode, you should see the following screenshot:

At this point, you should be able to select the menu **Product** -> **Run** to run the project and you should finally see the following screenshot:

Click the button to start recording! A screenshot is given below:

Congratulations! You have finally succeeded in running **sherpa-ncnn** with iOS, though it is in a simulator.

Please read below if you want to run **sherpa-ncnn** on your iPhone or iPad.




Fig. 7.13: Screenshot for selecting required files




Fig. 7.14: Screenshot after add pre-trained model files

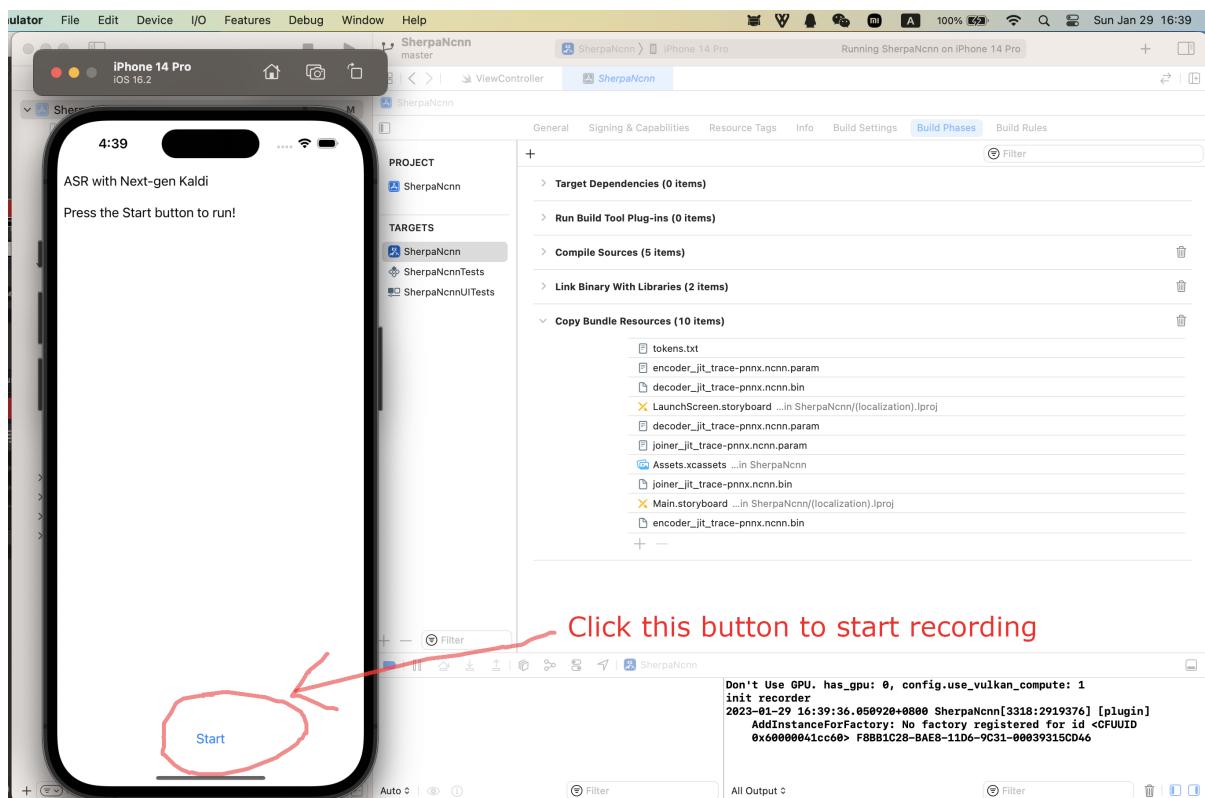


Fig. 7.15: Screenshot for a successful run.

Run sherpa-ncnn on your iPhone/iPad

First, please make sure the iOS version of your iPhone/iPad is ≥ 13.0 .

Click the menu **Xcode** \rightarrow **Settings...**, as is shown in the following screenshot:

In the popup dialog, please select **Account** and click **+** to add your Apple ID, as is shown in the following screenshots.

After adding your Apple ID, please connect your iPhone or iPad to your Mac and select your device in Xcode. The following screenshot is an example to select my iPhone.

Now your Xcode should look like below after selecting a device:

Please select **Product** \rightarrow **Run** again to run **sherpa-ncnn** on your selected device, as is shown in the following screenshot:

After a successful build, check your iPhone/iPad and you should see the following screenshot:

To fix that, please select **Settings** \rightarrow **General** \rightarrow **Device Management** on your device

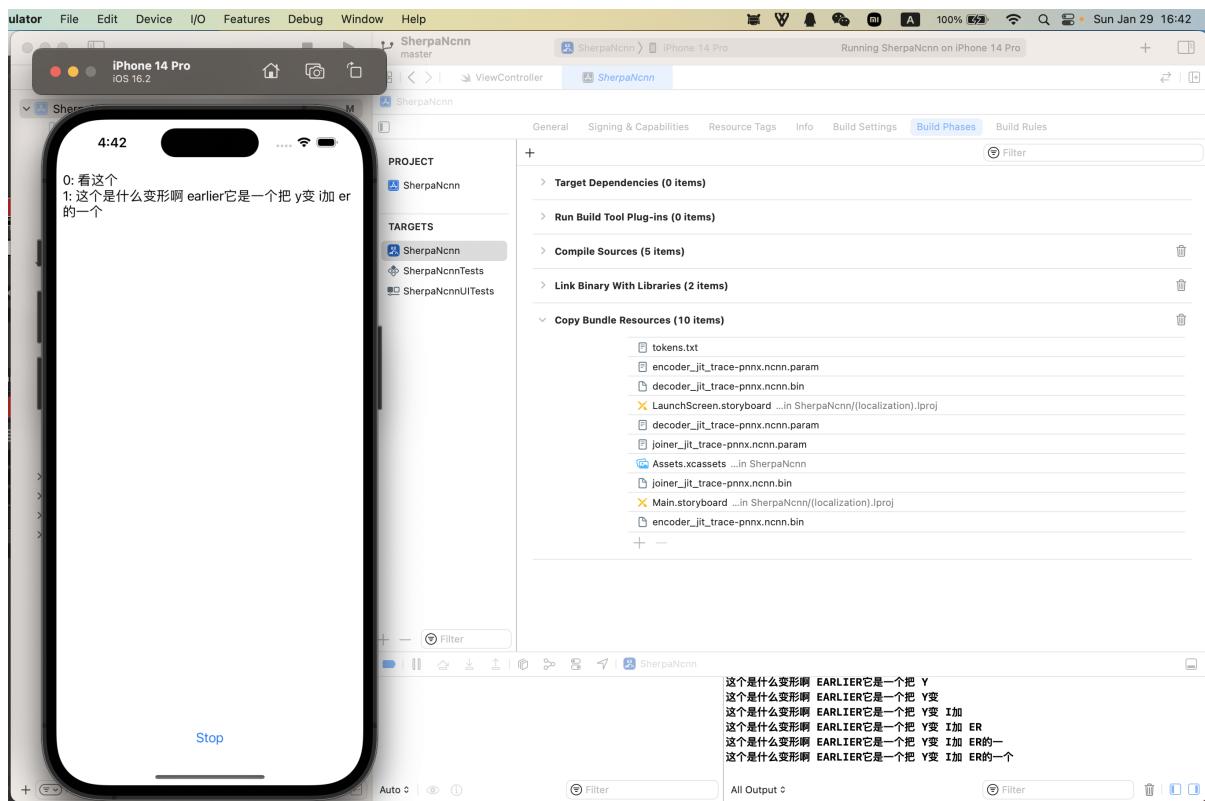


Fig. 7.16: Screenshot for recording and recognition.

Please click Apple Development: csukuangfj... and click Trust "Apple Development: csukuangfj@g..." in the subsequent dialog, as is shown below:

At this point, you should be able to run the app on your device. The following is a screenshot about running it on my iPhone:

Congratulations! You have successfully run `sherpa-ncnn` on your device!

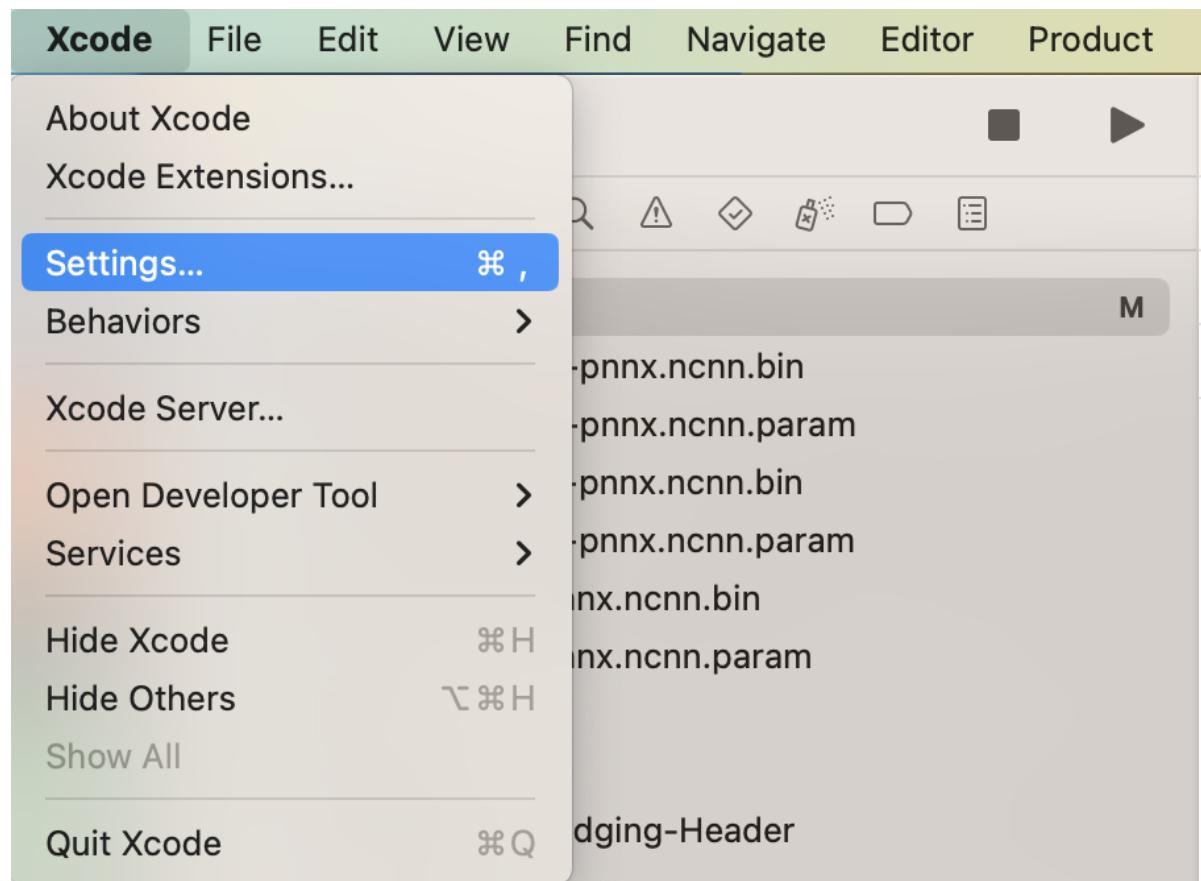


Fig. 7.17: Screenshot for Xcode -> Settings...

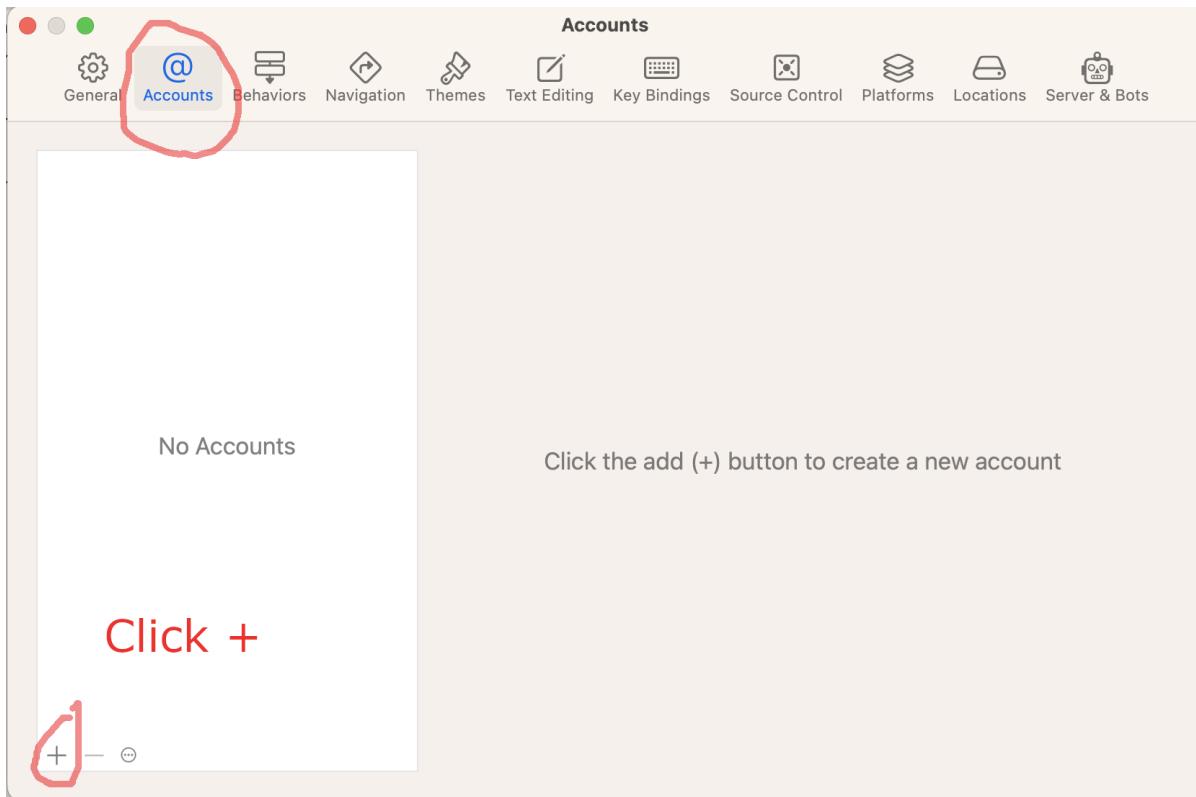


Fig. 7.18: Screenshot for selecting Account and click +.

7.8.3 For the more curious

This section is for those who want to learn more about how to use `sherpa-ncnn` in an iOS project.

Files generated by running `./build-ios.sh`

After running:

```
./build-ios.sh
```

You may be curious about the generated files.

Hint: Please have a look at `./build-ios.sh` so that you know what it does for you.

The above command generates files inside the directory `./build-ios`:

```
sherpa-ncnn fangjun$ ls -lh build-ios
total 1912
drwxr-xr-x  6 fangjun  staff  192B Feb 26 16:48 build
drwxr-xr-x  4 fangjun  staff  128B Feb 26 16:46 install
drwxr-xr-x 15 fangjun  staff  480B Feb 14 18:09 openmp-11.0.0.src
-rw-r--r--  1 fangjun  staff  952K Dec  8 2021 openmp-11.0.0.src.tar.xz
drwxr-xr-x  6 fangjun  staff  192B Feb 26 16:44 openmp.xcframework
drwxr-xr-x  6 fangjun  staff  192B Feb 26 16:48 sherpa-ncnn.xcframework
```

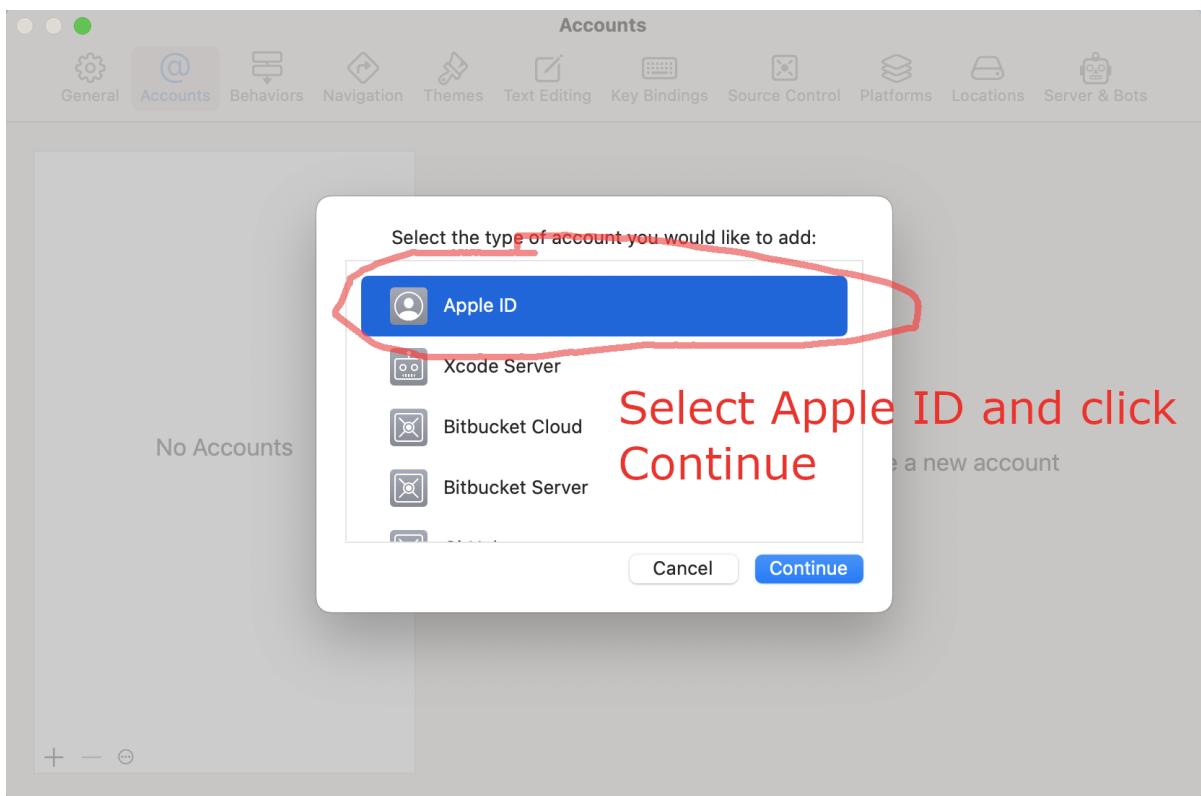


Fig. 7.19: Screenshot for selecting Apple ID and click Continue

What is interesting here is the two framework folders `openmp.xcframework` and `sherpa-ncnn.xcframework`. All other folders can be safely removed. We only need the two framework folders.

In the following, we describe the content in these two framework folders.

`openmp.xcframework`

```
$ tree build-ios/openmp.xcframework/
build-ios/openmp.xcframework/
├── Headers
│   └── omp.h
├── Info.plist
└── ios-arm64
    └── libomp.a
```

3 directories, 4 files

Explanation:

- `omp.h`: The header file, which is used by `ncnn`
- `Info.plist`: A file that is dedicated for framework on macOS/iOS
- `ios-arm64/libomp.a`: A static library for iOS device, e.g., for iPhone

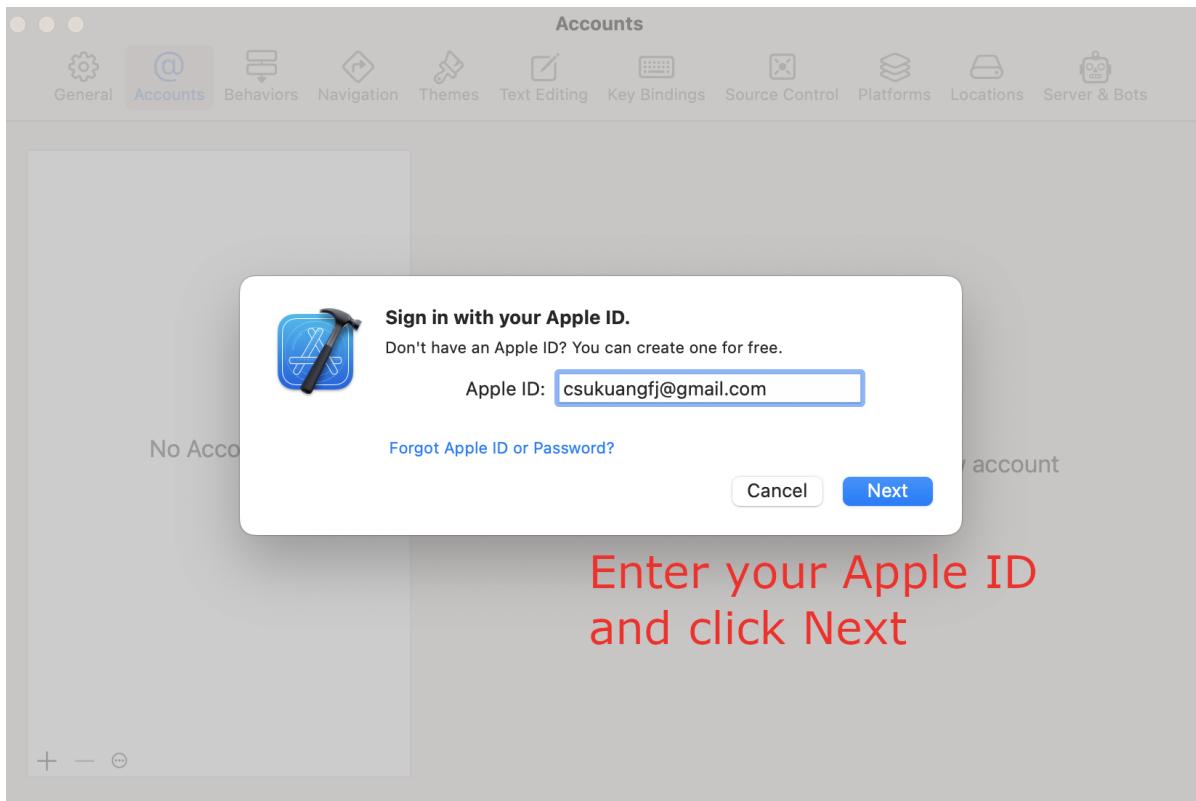


Fig. 7.20: Screenshot for adding your Apple ID and click Next

- `ios-arm64_x86_64-simulator/libomp.a`: A static library for iOS simulators, including simulators for Intel chips and Apple Silicon (e.g., M1)

```
sherpa-ncnn fangjun$ file build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a
build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a: Mach-O universal binary with 2 architectures: [x86_64:current ar archive random library] [arm64:current ar archive random library]
build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a (for architecture x86_64):      current ar archive random library
build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a (for architecture arm64):      current ar archive random library

sherpa-ncnn fangjun$ lipo -info build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a
Architectures in the fat file: build-ios/openmp.xcframework/ios-arm64_x86_64-simulator/libomp.a are: x86_64 arm64

sherpa-ncnn fangjun$ file build-ios/openmp.xcframework/ios-arm64/libomp.a
build-ios/openmp.xcframework/ios-arm64/libomp.a: current ar archive random library

sherpa-ncnn fangjun$ lipo -info build-ios/openmp.xcframework/ios-arm64/libomp.a
Non-fat file: build-ios/openmp.xcframework/ios-arm64/libomp.a is architecture: arm64
```

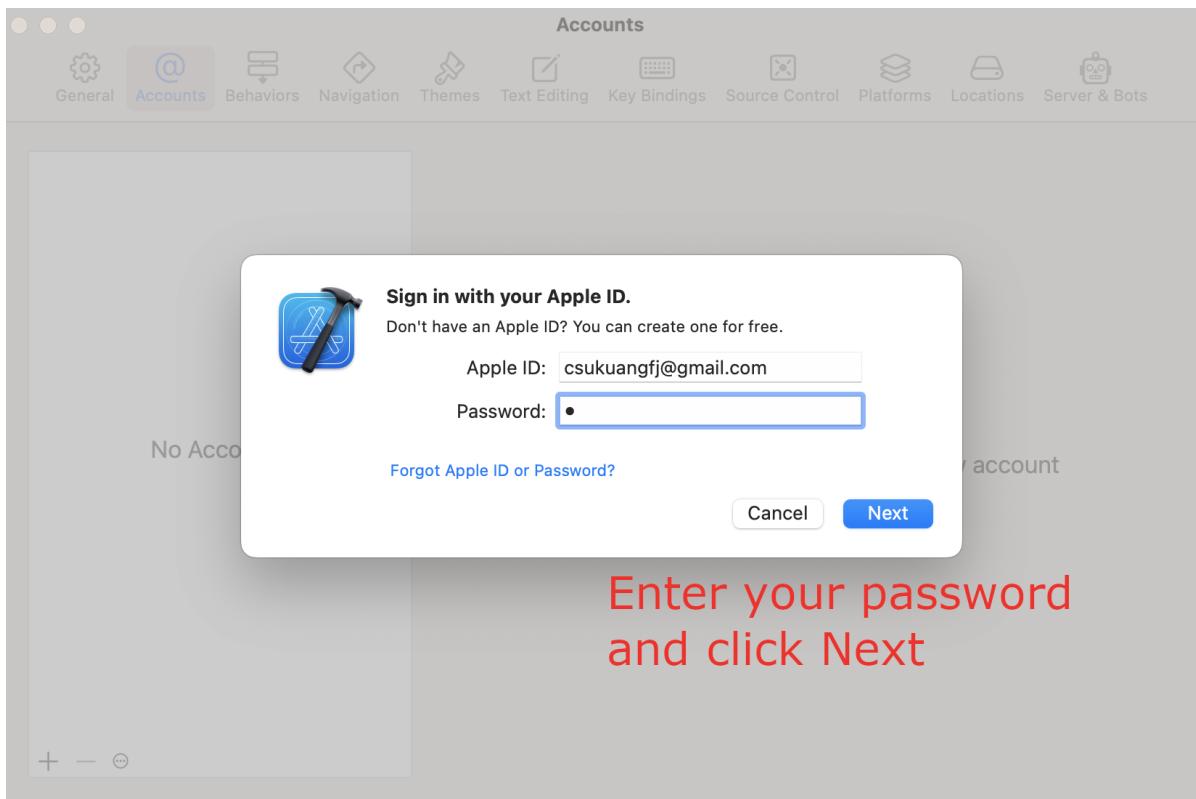


Fig. 7.21: Screenshot for entering your password and click Next

sherpa-ncnn.xcframework

```
sherpa-ncnn fangjun$ tree build-ios/sherpa-ncnn.xcframework/
build-ios/sherpa-ncnn.xcframework/
├── Headers
│   └── sherpa-ncnn
│       └── c-api
│           └── c-api.h
├── Info.plist
└── ios-arm64
    └── sherpa-ncnn.a
    └── ios-arm64_x86_64-simulator
        └── sherpa-ncnn.a

5 directories, 4 files
```

Explanation:

- **c-api.h:** The header file, which is copied from <https://github.com/k2-fsa/sherpa-ncnn/blob/master/sherpa-ncnn/c-api/c-api.h>
- **Info.plist:** A file that is dedicated for framework on macOS/iOS
- **ios-arm64/sherpa-ncnn.a:** A static library for iOS, e.g., iPhone
- **ios-arm64_x86_64-simulator/sherpa-ncnn.a:** A static library for simulators, including simulators for Intel chips and Apple Silicon (e.g., M1)

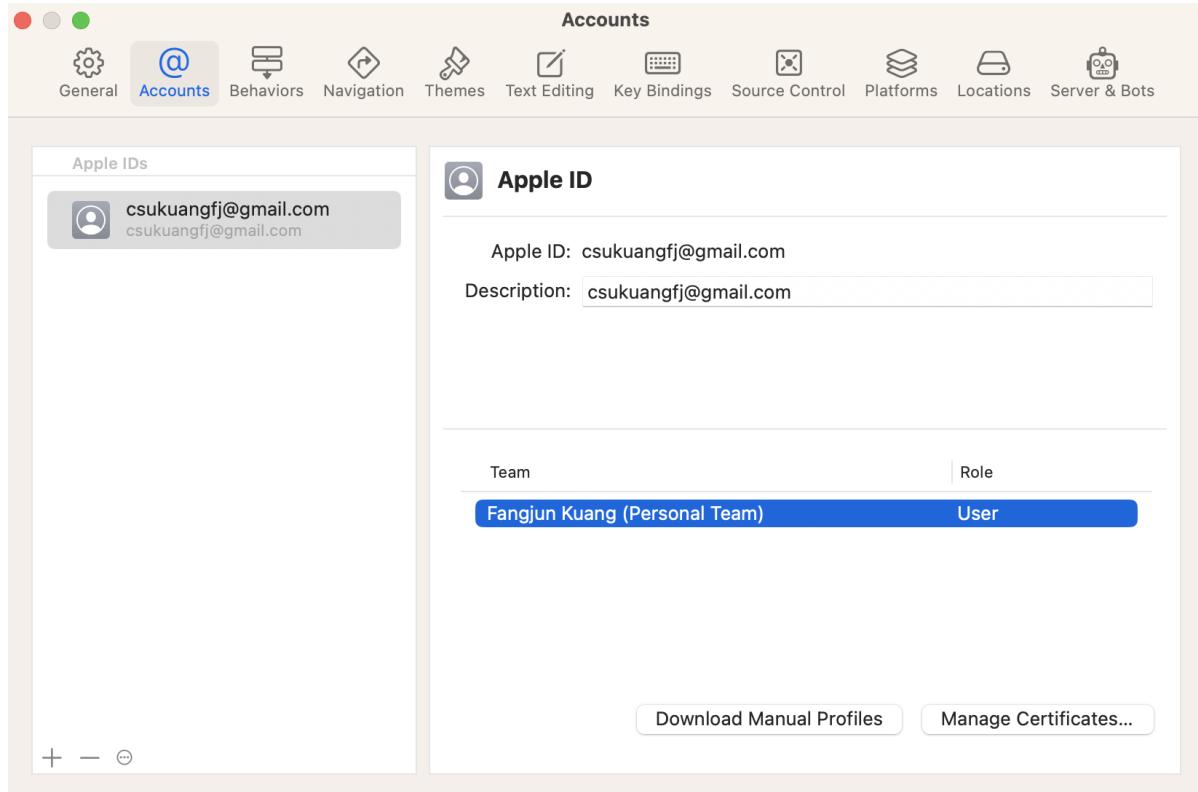


Fig. 7.22: Screenshot after adding your Apple ID

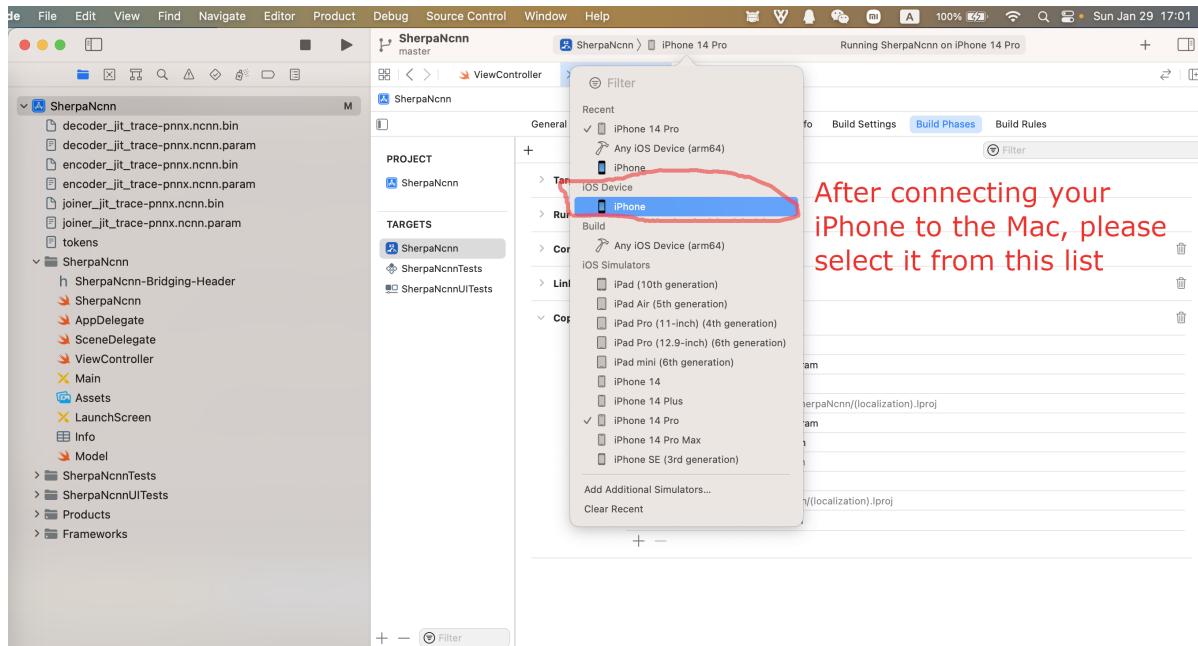


Fig. 7.23: Screenshot for selecting your device

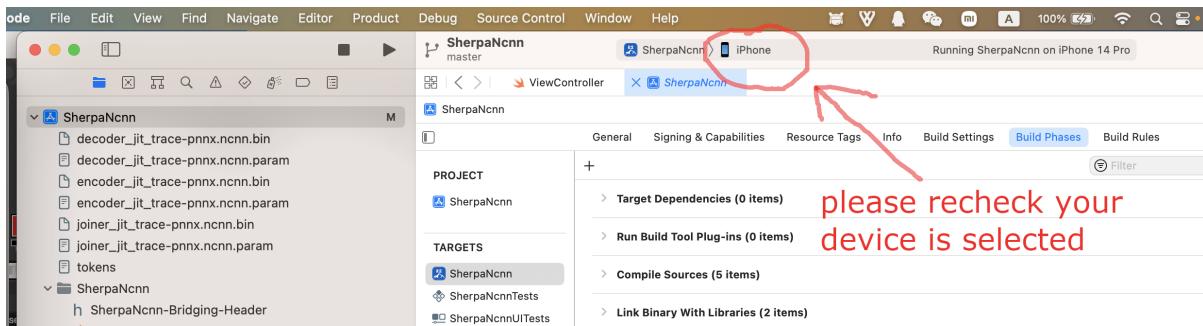


Fig. 7.24: Screenshot after selecting your device

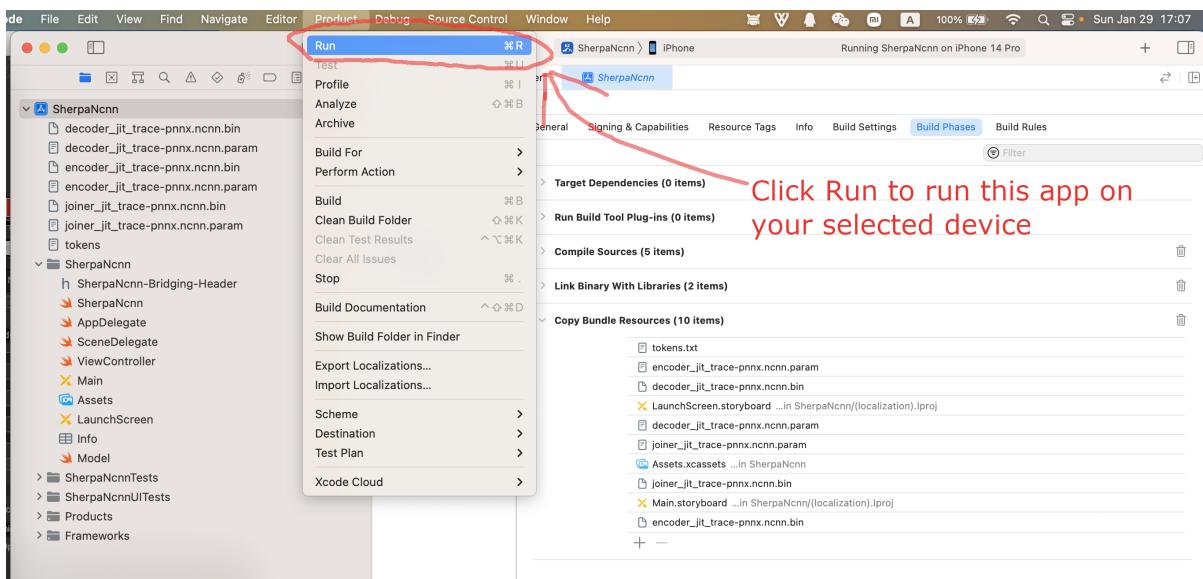


Fig. 7.25: Screenshot for selecting Product -> Run

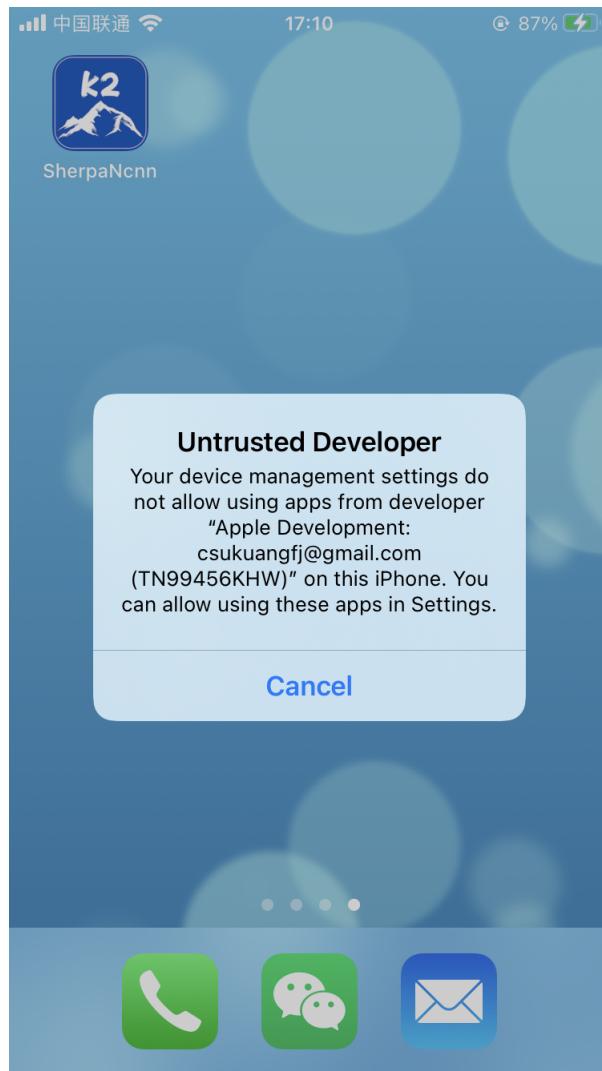


Fig. 7.26: Screenshot for running sherpa-ncnn on your device

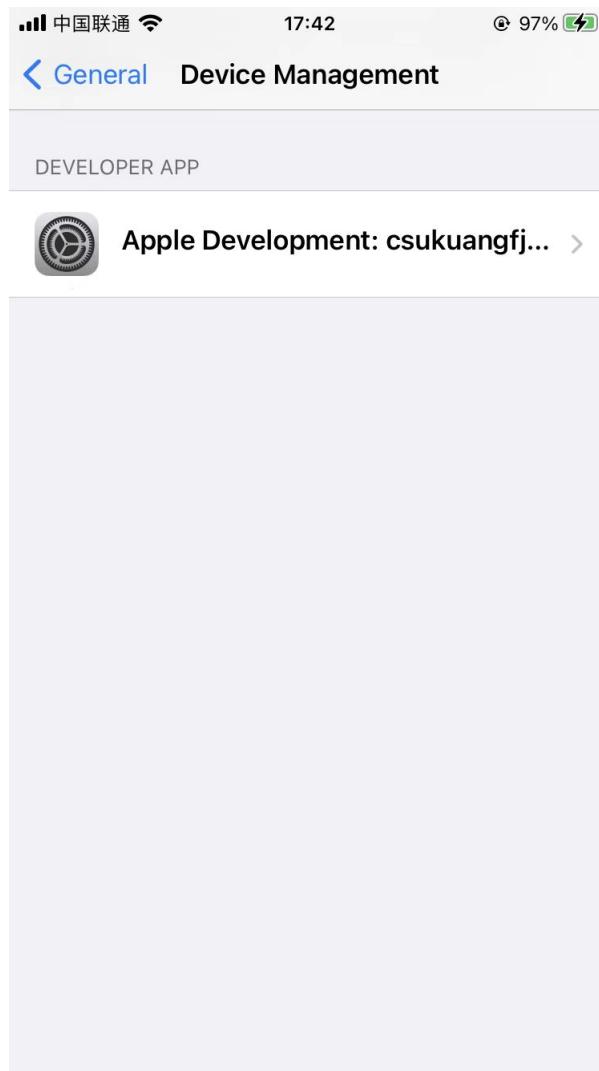


Fig. 7.27: Screenshot for selecting *Settings -> General -> Device Management* on your device

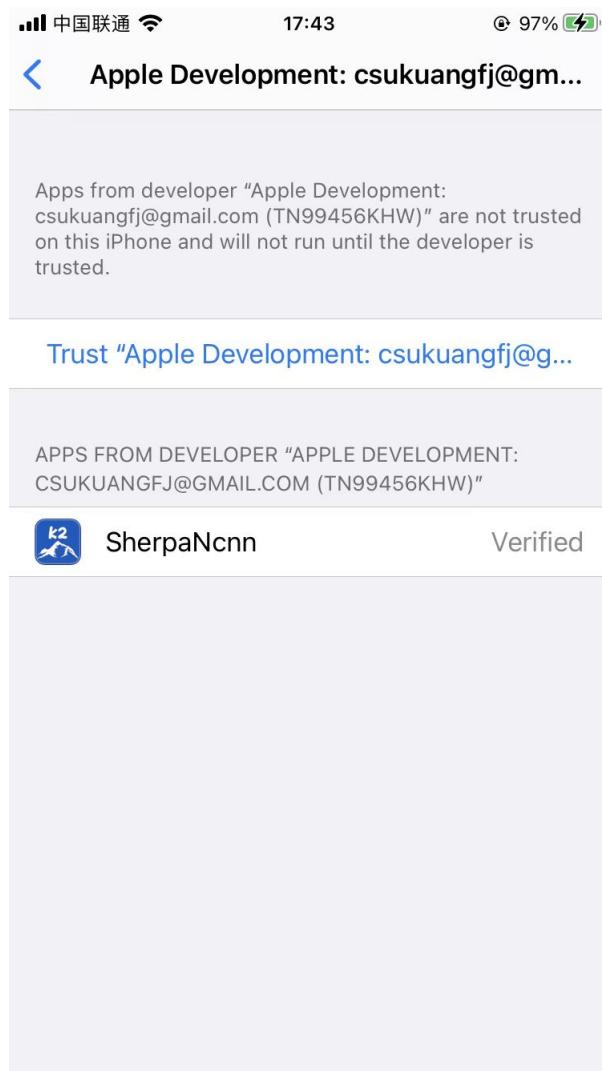


Fig. 7.28: Screenshot for “Trust “Apple Development: csukuangfj@g...””

中国联通 18:41 97%

ASR with Next-gen Kaldi

Press the Start button to run!

Click Start to
run it!



Fig. 7.29: Screenshot for running sherpa-ncnn on iPhone

```

sherpa-ncnn fangjun$ file build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-simulator/
sherpa-ncnn.a
build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-simulator/sherpa-ncnn.a: Mach-O
universal binary with 2 architectures: [x86_64:current ar archive] [arm64]
build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-simulator/sherpa-ncnn.a (for
architecture x86_64): current ar archive
build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-simulator/sherpa-ncnn.a (for
architecture arm64): current ar archive

sherpa-ncnn fangjun$ lipo -info build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-
simulator/sherpa-ncnn.a
Architectures in the fat file: build-ios/sherpa-ncnn.xcframework/ios-arm64_x86_64-
simulator/sherpa-ncnn.a are: x86_64 arm64

sherpa-ncnn fangjun$ file build-ios/sherpa-ncnn.xcframework/ios-arm64/sherpa-ncnn.a
build-ios/sherpa-ncnn.xcframework/ios-arm64/sherpa-ncnn.a: current ar archive

sherpa-ncnn fangjun$ lipo -info build-ios/sherpa-ncnn.xcframework/ios-arm64/sherpa-ncnn.a
Non-fat file: build-ios/sherpa-ncnn.xcframework/ios-arm64/sherpa-ncnn.a is architecture:
arm64

```

How to use files generated by ./build-ios.sh in Xcode

In this section, we describe how to use `openmp.xcframework` and `sherpa-ncnn.xcframework` in Xcode.

The underlying implementation of `sherpa-ncnn` is in C++. It also provides C API.

To use C API in Xcode with Swift, we have to write a bridging header.

We provide a bridging header for you: `SherpaNcnn-Bridging-Header.h`. All you need is to add this file to your iOS project and click Build Settings -> Swift Compiler - General and set Objective-C Bridging Header to `$(PROJECT_DIR)/../../swift-api-examples/SherpaNcnn-Bridging-Header.h`. See the screenshot below for reference:

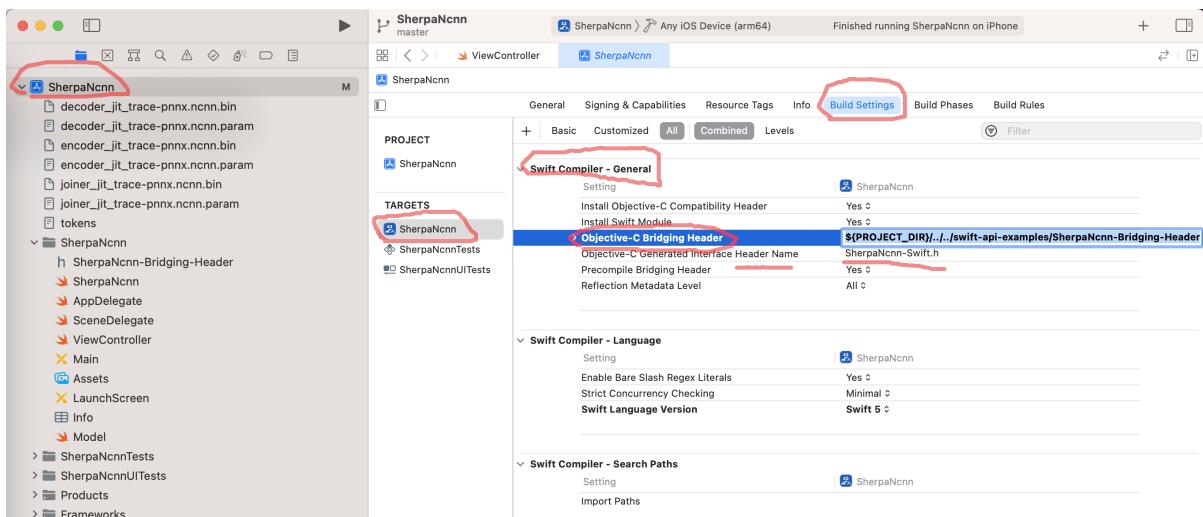


Fig. 7.30: Screenshot for setting the bridging header

We list the content of the bridging header below for reference:

```
#ifndef SWIFT_API_EXAMPLES_SHERPANCNN_BRIDGING_HEADER_H_
#define SWIFT_API_EXAMPLES_SHERPANCNN_BRIDGING_HEADER_H_

#import "sherpa-ncnn/c-api/c-api.h"

#endif // SWIFT_API_EXAMPLES_SHERPANCNN_BRIDGING_HEADER_H_
```

After adding the bridging header to your iOS project, Xcode will complain it cannot find `sherpa-ncnn/c-api/c-api.h`. The fix is to add the path `build-ios/sherpa-ncnn.xcframework/Headers` to Header Search Paths by changing Build Settings -> Search Paths -> Header Search Paths, as is shown in the following screenshot:

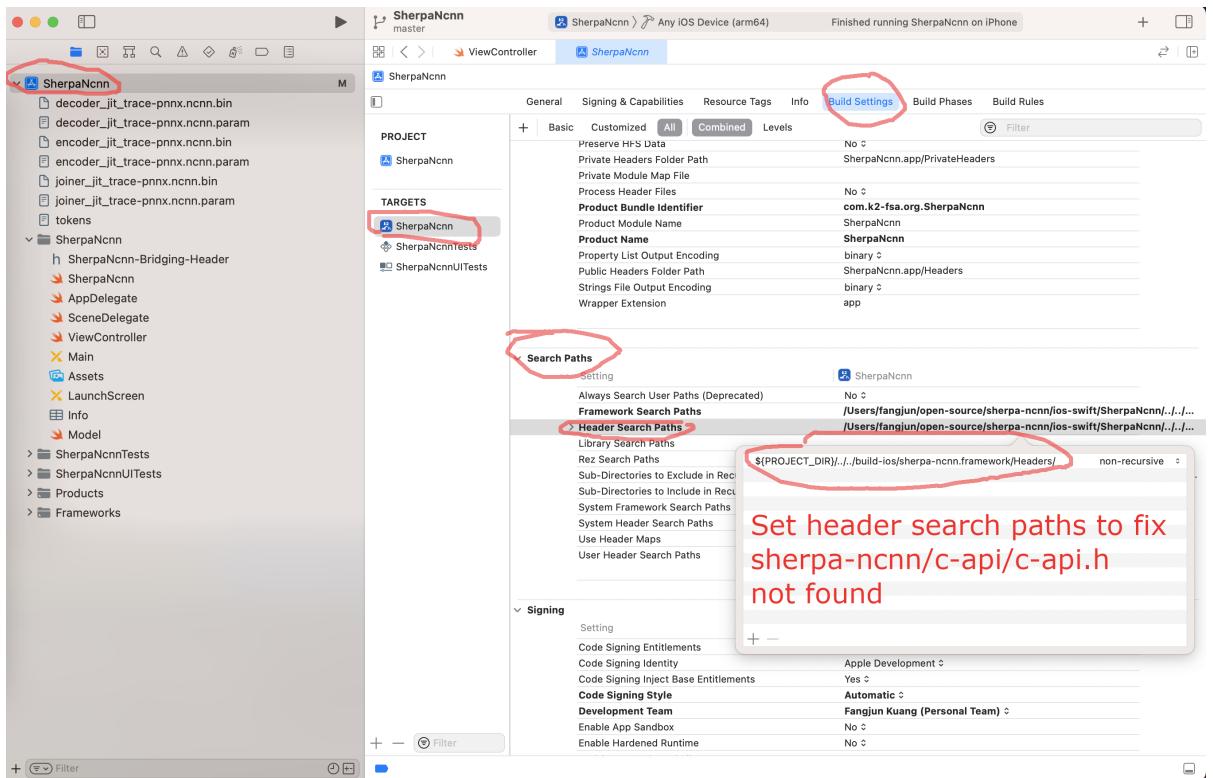


Fig. 7.31: Screenshot for setting the header search paths

Hint: Instead of using an absolute path, we use `$(PROJECT_DIR)/../../build-ios/sherpa-ncnn.xcframework/Headers/`

For instance, my `sherpa-ncnn` is downloaded to `/Users/fangjun/open-source/sherpa-ncnn` and the path to `sherpa-ncnn.xcframework` is `/Users/fangjun/open-source/sherpa-ncnn/build-ios/sherpa-ncnn.xcframework`.

The value of `PROJECT_DIR` is `/Users/fangjun/open-source/sherpa-ncnn/ios-swift/SherpaNcnn`, so we can use `$(PROJECT_DIR)/../../build-ios/sherpa-ncnn.xcframework/Headers/`.

Also note that `PROJECT_DIR` is a pre-defined variable in Xcode.

Please also add `SherpaNcnn.swift` to your iOS project, which is a utility to make it easier to access functions from the bridging header.

The next thing is to add `openmp.xcframework` and `sherpa-ncnn.xcframework` as dependencies to your iOS project. Select Build Phases -> Link Binary with Libraries and then click + to add `sherpa-ncnn.xcframework` and `openmp.xcframework`. See the screenshot below for reference:

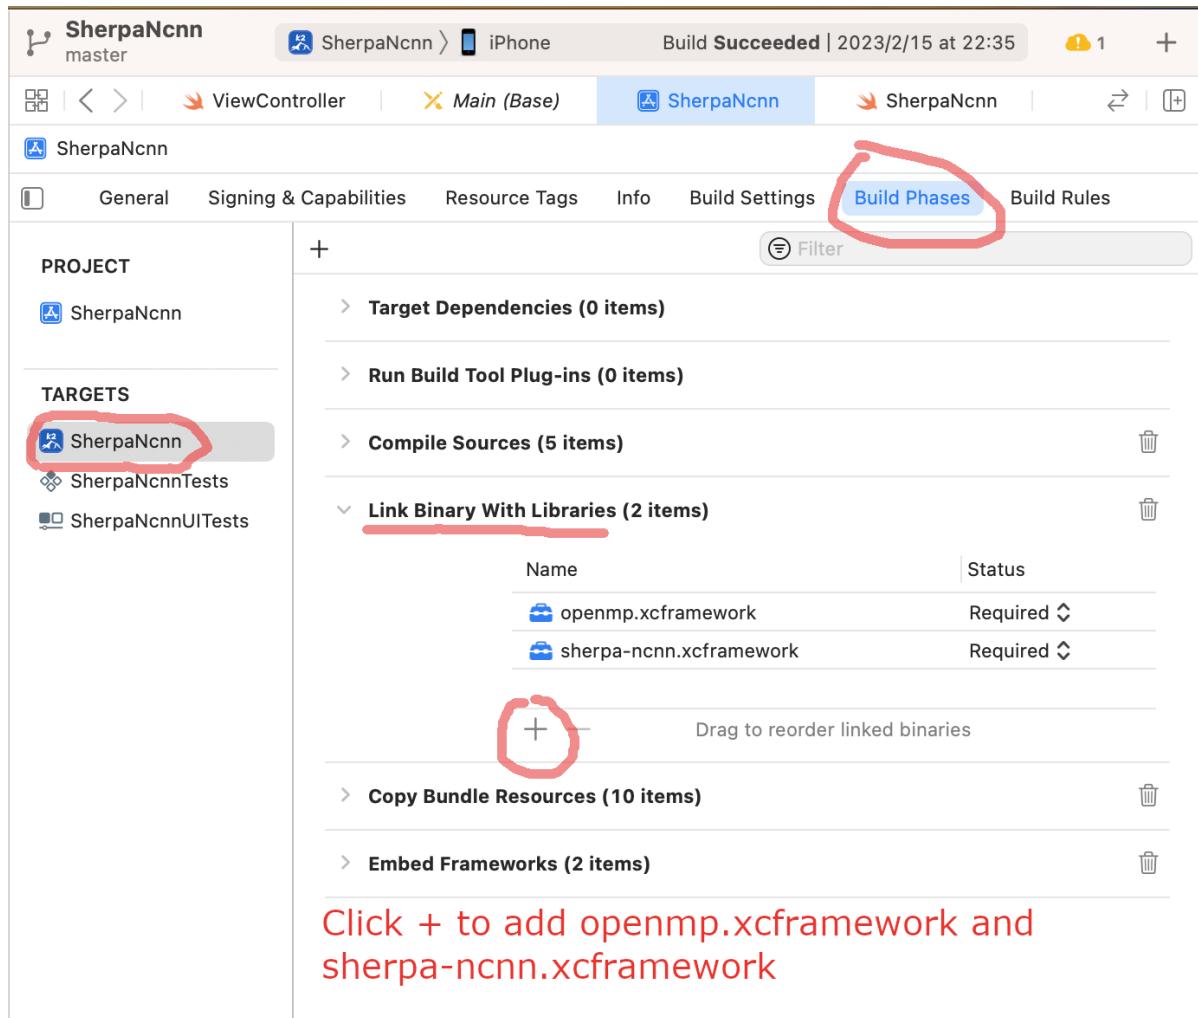


Fig. 7.32: Screenshot for adding a framework to your project

Hint: After clicking +, please select Add Other... -> Add Files..., and then add the path to `sherpa-ncnn.xcframework`.

Repeat the step for `openmp.xcframework`.

See the screenshot below for reference:

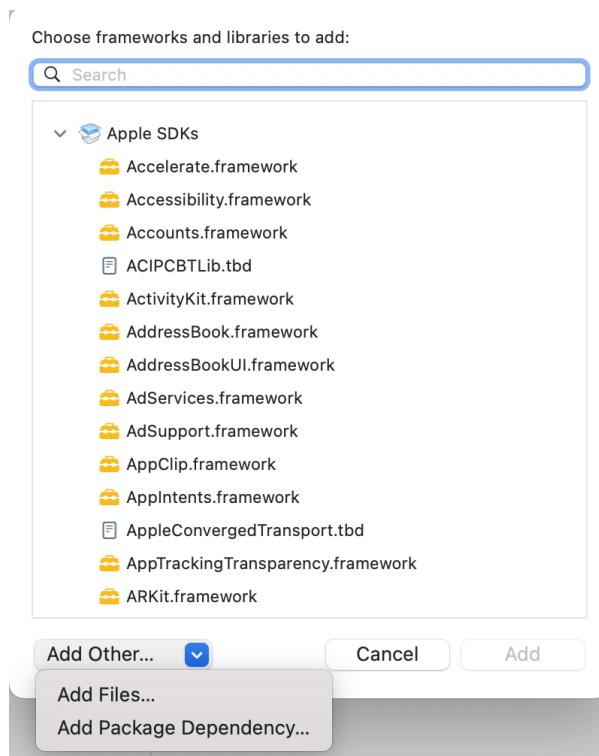


Fig. 7.33: Screenshot for adding a framework

One more thing you need to do after adding the framework is to setup the framework search path. Click **Build Settings** -> **Search Paths** -> **Framework Search Paths** and add the path to `build-ios/`. See the screenshot below:

If you encounter link errors about the C++ standard library, please add `-lc++` to link against `libc++` by clicking **Build Settings** -> **Linking** -> **Other Linker Flags** and adding `-lc++`. See the screenshot below for reference:

That is all you need to add `sherpa-ncnn` to your iOS project.

7.9 Pre-trained models

In this section, we describe how to download and use all available pre-trained models.

Hint: Please install `git-lfs` before you continue.

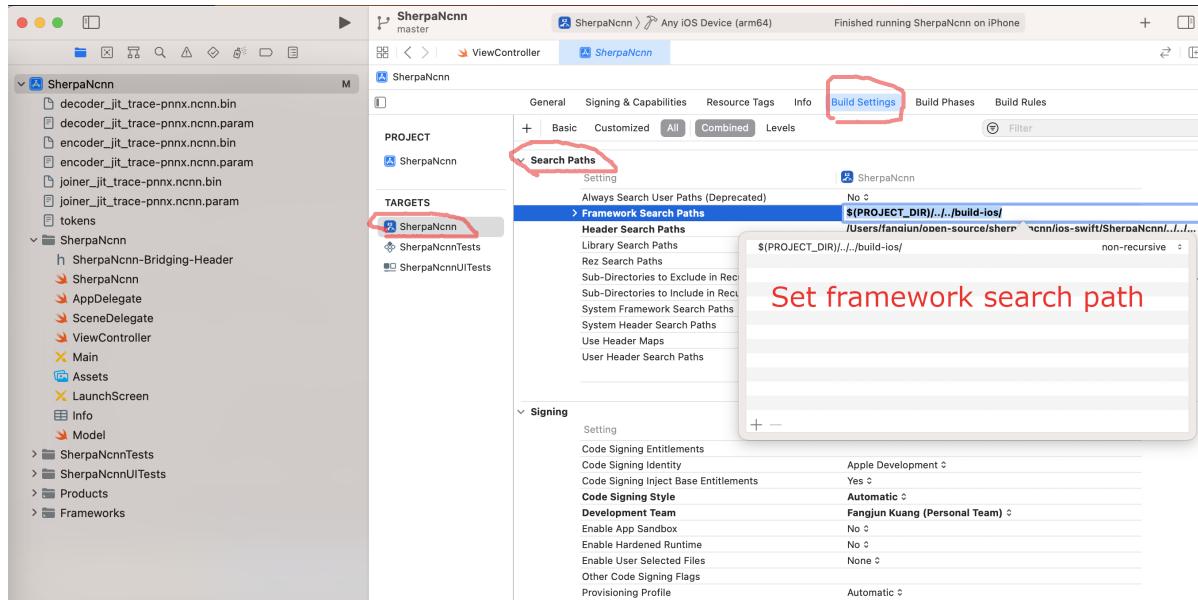


Fig. 7.34: Screenshot for setting framework search paths

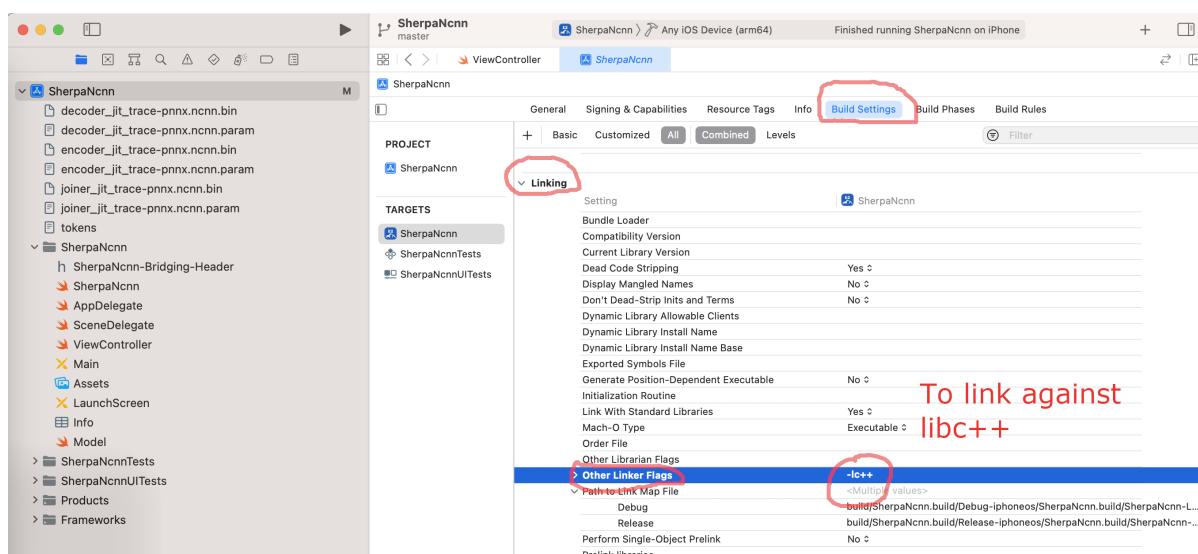


Fig. 7.35: Screenshot for adding -lc++ to linker flags

7.9.1 Small models

In this section, we list models with fewer parameters that are suitable for resource constrained embedded systems.

Hint: If you are using Raspberry Pi 4, this section is not so helpful for you since all models in `sherpa-ncnn` are able to run in real-time on it.

This page is especially useful for systems with less resource than Raspberry Pi 4.

- [marcoyang/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 \(Chinese\)](#)
- [marcoyang/sherpa-ncnn-streaming-zipformer-20M-2023-02-17 \(English\)](#)
- [marcoyang/sherpa-ncnn-conv-emformer-transducer-small-2023-01-09 \(English\)](#)
- [cskuangfj/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16 \(Bilingual, Chinese + English\)](#)
- [marcoyang/sherpa-ncnn-lstm-transducer-small-2023-02-13 \(Bilingual, Chinese + English\)](#)

7.9.2 Zipformer-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-ncnn` before you read this section.

marcoyang/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 (Chinese)

This model is a streaming Zipformer model which has around 14 million parameters. It is trained on the [WenetSpeech](#) corpus so it supports only Chinese.

You can find the training code at https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless7_streaming

In the following, we describe how to download it and use it with `sherpa-ncnn`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
˓→streaming-zipformer-zh-14M-2023-02-23.tar.bz2
tar xvf sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23.tar.bz2
```

Decode a single wave file

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/tokens.txt \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/encoder_jit_trace-pnnx.ncnn.
    ↪ param \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/encoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/decoder_jit_trace-pnnx.ncnn.
    ↪ param \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/decoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/joiner_jit_trace-pnnx.ncnn.param
    ↪ \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/joiner_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/test_wavs/0.wav \
    2 \
    $method
done
```

You should see the following output:

```
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param="pruned_transducer_stateless7_streaming/exp-small-L/encoder_
↪ jit_trace-pnnx.ncnn.param", encoder_bin="pruned_transducer_stateless7_streaming/exp-
↪ small-L/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="pruned_transducer_stateless7_
↪ streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="pruned_
↪ transducer_stateless7_streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.bin", joiner_
↪ param="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_trace-pnnx.ncnn.
↪ param", joiner_bin="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_
↪ trace-pnnx.ncnn.bin", tokens="data/lang_char/tokens.txt", encoder num_threads=4,_
↪ decoder num_threads=4, joiner num_threads=4)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,_
↪ endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
↪ trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
↪ nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
↪ rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
↪ utterance_length=20)))
wav filename: ./test_wavs_zh/0.wav
wav duration (s): 5.6115
Started!
Done!
Recognition result for ./test_wavs_zh/0.wav

Elapsed seconds: 0.678 s
Real time factor (RTF): 0.678 / 5.611 = 0.121
```

(continues on next page)

(continued from previous page)

```

Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param="pruned_transducer_stateless7_streaming/exp-small-L/encoder_
↪jit_trace-pnnx.ncnn.param", encoder_bin="pruned_transducer_stateless7_streaming/exp-
↪small-L/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="pruned_transducer_stateless7_
↪streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="pruned_
↪transducer_stateless7_streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.bin", joiner_
↪param="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_trace-pnnx.ncnn.
↪param", joiner_bin="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_
↪trace-pnnx.ncnn.bin", tokens="data/lang_char/tokens.txt", encoder num_threads=4,_
↪decoder num_threads=4, joiner num_threads=4)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,_
↪endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
↪trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
↪nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
↪rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
↪utterance_length=20)))
wav filename: ./test_wavs_zh/1.wav
wav duration (s): 5.15306
Started!
Done!
Recognition result for ./test_wavs_zh/1.wav

Elapsed seconds: 0.676 s
Real time factor (RTF): 0.676 / 5.153 = 0.131

Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param="pruned_transducer_stateless7_streaming/exp-small-L/encoder_
↪jit_trace-pnnx.ncnn.param", encoder_bin="pruned_transducer_stateless7_streaming/exp-
↪small-L/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="pruned_transducer_stateless7_
↪streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="pruned_
↪transducer_stateless7_streaming/exp-small-L/decoder_jit_trace-pnnx.ncnn.bin", joiner_
↪param="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_trace-pnnx.ncnn.
↪param", joiner_bin="pruned_transducer_stateless7_streaming/exp-small-L/joiner_jit_
↪trace-pnnx.ncnn.bin", tokens="data/lang_char/tokens.txt", encoder num_threads=4,_
↪decoder num_threads=4, joiner num_threads=4)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,_
↪endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
↪trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
↪nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
↪rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
↪utterance_length=20)))
wav filename: ./test_wavs_zh/2.wav
wav duration (s): 4.52431
Started!
Done!
Recognition result for ./test_wavs_zh/2.wav

Elapsed seconds: 0.592 s
Real time factor (RTF): 0.592 / 4.524 = 0.131

```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

`CHCP 65001`

in your commandline.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn  
  
./build/bin/sherpa-ncnn-microphone \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/tokens.txt \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/encoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/encoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/decoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/decoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/joiner_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/joiner_jit_trace-pnnx.ncnn.bin \  
  2 \  
  greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use `sherpa-ncnn-alsa` to do real-time speech recognition with your microphone if `sherpa-ncnn-microphone` does not work for you.

marcoyang/sherpa-ncnn-streaming-zipformer-20M-2023-02-17 (English)

This model is a streaming Zipformer model converted from

<https://huggingface.co/desh2608/icefall-asr-librispeech-pruned-transducer-stateless7-streaming-small>

which has around 20 million parameters. It is trained on the LibriSpeech corpus so it supports only English. The word-error-rates(%) on `test-clean` is 3.88.

You can find the training code at https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless7_streaming

In the following, we describe how to download it and use it with `sherpa-ncnn`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
tar streaming-zipformer-20M-2023-02-17.tar.bz2
tar xvf sherpa-ncnn-streaming-zipformer-20M-2023-02-17.tar.bz2
```

Decode a single wave file

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn
for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/tokens.txt \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/encoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/encoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/decoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/decoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/test_wavs/0.wav \
    2 \
    $method
done
```

You should see the following output:

```
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param=".sherpa-ncnn-streaming-zipformer-20M-2023-02-17/encoder_jit_"
            "trace-pnnx.ncnn.param", encoder_bin=".sherpa-ncnn-streaming-zipformer-20M-2023-02-17/"
            "encoder_jit_trace-pnnx.ncnn.bin", decoder_param=".sherpa-ncnn-streaming-zipformer-20M-"
            "2023-02-17/decoder_jit_trace-pnnx.ncnn.param", decoder_bin=".sherpa-ncnn-streaming-"
            "zipformer-20M-2023-02-17/decoder_jit_trace-pnnx.ncnn.bin", joiner_param=".sherpa-ncnn-"
            "streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./"
            "sherpa-ncnn-streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.bin", tokens=
            "./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/tokens.txt", encoder_num_threads=2,_
            "decoder num_threads=2, joiner num_threads=2)
DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False,_
            "endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_"
            "trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_"
            "nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
            "rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_"
            "utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/test_wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
wav duration (s): 6.625
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/test_wavs/0.wav
  AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
  ↵ QUARTER OF THE BRAFFLELS
Elapsed seconds: 0.472 s
Real time factor (RTF): 0.472 / 6.625 = 0.071
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/tokens.txt \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/encoder_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/encoder_jit_trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/decoder_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/decoder_jit_trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-20M-2023-02-17/joiner_jit_trace-pnnx.ncnn.bin \
  2 \
  greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use `sherpa-ncnn-alsa` to do real-time speech recognition with your microphone if `sherpa-ncnn-microphone` does not work for you.

csukuangfj/sherpa-ncnn-streaming-zipformer-en-2023-02-13 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-pruned-transducer-stateless7-streaming-2022-12-29>

which supports only English as it is trained on the LibriSpeech corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless7_streaming

In the following, we describe how to download it and use it with `sherpa-ncnn`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn

wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
˓→streaming-zipformer-en-2023-02-13.tar.bz2
tar xvf sherpa-ncnn-streaming-zipformer-en-2023-02-13.tar.bz2
```

Decode a single wave file

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
./build/bin/sherpa-ncnn \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/tokens.txt \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/test_wavs/1221-135766-0002.wav \
2 \
$method
done
```

You should see the following output:

```
ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_
˓→trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-en-2023-02-13/
˓→encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-ncnn-streaming-zipformer-en-
˓→2023-02-13/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./sherpa-ncnn-streaming-
˓→zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./sherpa-ncnn-
˓→streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./
˓→sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin", tokens=
˓→"./sherpa-ncnn-streaming-zipformer-en-2023-02-13/tokens.txt", encoder_num_threads=2,_
˓→decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False,_
˓→endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
˓→trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
˓→nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
˓→rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
˓→utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-en-2023-02-13/test_wavs/1221-135766-0002.
˓→wav
wav duration (s): 4.825
```

(continues on next page)

(continued from previous page)

```
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-en-2023-02-13/test_wavs/1221-
↪135766-0002.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
Elapsed seconds: 0.569 s
Real time factor (RTF): 0.569 / 4.825 = 0.118
ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_"
↪trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-en-2023-02-13/
↪encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-ncnn-streaming-zipformer-en-
↪2023-02-13/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./sherpa-ncnn-streaming-
↪zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./sherpa-ncnn-
↪streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./
↪sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin", tokens=
↪"./sherpa-ncnn-streaming-zipformer-en-2023-02-13/tokens.txt", encoder num_threads=2,
↪decoder num_threads=2, joiner num_threads=2)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,
↪endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
↪trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
↪nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),
↪rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
↪utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-en-2023-02-13/test_wavs/1221-135766-0002.
↪wav
wav duration (s): 4.825
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-en-2023-02-13/test_wavs/1221-
↪135766-0002.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
Elapsed seconds: 0.554 s
Real time factor (RTF): 0.554 / 4.825 = 0.115
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/tokens.txt \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin \
2 \
greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-ncnn-alsa* to do real-time speech recognition with your microphone if *sherpa-ncnn-microphone* does not work for you.

cskuangfj/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13 (Bilingual, Chinese + English)

This model is converted from

<https://huggingface.co/pfluo/k2fsa-zipformer-chinese-english-mixed>

which supports both Chinese and English. The model is contributed by the community and is trained on tens of thousands of some internal dataset.

In the following, we describe how to download it and use it with *sherpa-ncnn*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
˓→streaming-zipformer-bilingual-zh-en-2023-02-13.tar.bz2
tar xvf sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13.tar.bz2
```

Decode a single wave file

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/tokens.txt \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.
    ↵ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.
    ↵ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.
    ↵ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.
    ↵ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.
    ↵ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.
    ↵ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/test_wavs/1.wav \
    ↵\
```

(continues on next page)

(continued from previous page)

```
$method
done
```

You should see the following output:

```
ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/
˓encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-
˓bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-
˓ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.param",
˓decoder_bin="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_
˓trace-pnnx.ncnn.bin", joiner_param="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-
˓2023-02-13/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./sherpa-ncnn-streaming-
˓zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin", tokens="./sherpa-
˓ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/tokens.txt", encoder num_threads=2,
˓decoder num_threads=2, joiner num_threads=2)
DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False,
˓endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
˓trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
˓nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),
˓rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
˓utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/test_wavs/1.
˓wav
wav duration (s): 5.1
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/test_
˓wavs/1.wav
ALWAYS ALWAYS
Elapsed seconds: 0.598 s
Real time factor (RTF): 0.598 / 5.100 = 0.117
ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/
˓encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-
˓bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-
˓ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.ncnn.param",
˓decoder_bin="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_
˓trace-pnnx.ncnn.bin", joiner_param="./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-
˓2023-02-13/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./sherpa-ncnn-streaming-
˓zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin", tokens="./sherpa-
˓ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/tokens.txt", encoder num_threads=2,
˓decoder num_threads=2, joiner num_threads=2)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,
˓endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
˓trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
˓nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),
˓rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
˓utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/test_wavs/1.
˓wav
wav duration (s): 5.1
Started!
Done!
```

(continues on next page)

(continued from previous page)

```
Recognition result for ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/test_
↳ wav/1.wav
ALWAYS ALWAYS
Elapsed seconds: 0.943 s
Real time factor (RTF): 0.943 / 5.100 = 0.185
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/tokens.txt \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.
↳ ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/encoder_jit_trace-pnnx.
↳ ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.
↳ ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/decoder_jit_trace-pnnx.
↳ ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.
↳ ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13/joiner_jit_trace-pnnx.
↳ ncnn.bin \
  2 \
  greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use `sherpa-ncnn-alsa` to do real-time speech recognition with your microphone if `sherpa-ncnn-microphone` does not work for you.

[csukuangfj/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16](https://huggingface.co/csukuangfj/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16) (Bilingual, Chinese + English)

This model is converted from

<https://huggingface.co/csukuangfj/k2fsa-zipformer-bilingual-zh-en-t>

which supports both Chinese and English. The model is contributed by the community and is trained on tens of thousands of some internal dataset.

In the following, we describe how to download it and use it with `sherpa-ncnn`.

Note: Unlike [csukuangfj/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13](https://huggingface.co/csukuangfj/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13) (Bilingual, Chinese + English), this model is much smaller.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn  
  
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-  
→ streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2  
tar xvf sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2
```

Decode a single wave file

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn  
  
for method in greedy_search modified_beam_search; do  
    ./build/bin/sherpa-ncnn \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-  
→ pnnx.ncnn.param \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-  
→ pnnx.ncnn.bin \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-  
→ pnnx.ncnn.param \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-  
→ pnnx.ncnn.bin \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-  
→ pnnx.ncnn.param \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-  
→ pnnx.ncnn.bin \  
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/1.wav \  
    2 \  
done
```

(continues on next page)

(continued from previous page)

\$method
done

You should see the following output:

```
ModelConfig(encoder_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-pnnx.ncnn.param", encoder_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-pnnx.ncnn.param", decoder_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-pnnx.ncnn.bin", joiner_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-pnnx.ncnn.param", joiner_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-pnnx.ncnn.bin", tokens=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt", encoder_num_threads=2, decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False, endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/1.wav
wav duration (s): 5.1
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/1.wav
ALWAYS
ModelConfig(encoder_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-pnnx.ncnn.param", encoder_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-pnnx.ncnn.param", decoder_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-pnnx.ncnn.bin", joiner_param=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-pnnx.ncnn.param", joiner_bin=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-pnnx.ncnn.bin", tokens=".sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt", encoder_num_threads=2, decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False, endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)))
wav filename: ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/1.wav
wav duration (s): 5.1
Started!
Done!
```

(continues on next page)

(continued from previous page)

```
Recognition result for ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-
˓→16/test_wavs/1.wav
ALWAYS
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-
˓→pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder_jit_trace-
˓→pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-
˓→pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder_jit_trace-
˓→pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-
˓→pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner_jit_trace-
˓→pnnx.ncnn.bin \
  2 \
  greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use `sherpa-ncnn-alsa` to do real-time speech recognition with your microphone if `sherpa-ncnn-microphone` does not work for you.

A faster model of `sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16`

We provide a second version of the model that is exported with `--decode-chunk-len=96` instead of 32.

Hint: Please see the model export script at

<https://huggingface.co/csukuangfj/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/blob/main/96/export-for-ncnn-bilingual-small.sh>

if you are interested.

Note: You can also find a third version with folder 64.

The advantage of using this model is that it runs much faster, while the downside is that you will see some delay before you see the recognition result after you speak.

To decode a file, please use:

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/tokens.txt \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/encoder_jit_
    ↵trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/encoder_jit_
    ↵trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/decoder_jit_
    ↵trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/decoder_jit_
    ↵trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/joiner_jit_
    ↵trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/96/joiner_jit_
    ↵trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/1.wav \
    2 \
    $method
done
```

shaojieli/sherpa-ncnn-streaming-zipformer-fr-2023-04-14

This model is converted from

<https://huggingface.co/shaojieli/icefall-asr-commonvoice-fr-pruned-transducer-stateless7-streaming-2023-04-02>

which supports only French as it is trained on the [CommonVoice](#) corpus. In the following, we describe how to download it and use it with [sherpa-ncnn](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn

wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-
↪streaming-zipformer-fr-2023-04-14.tar.bz2
tar xvf sherpa-ncnn-streaming-zipformer-fr-2023-04-14.tar.bz2
```

To decode a file, please use:

```
cd /path/to/sherpa-ncnn
for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/tokens.txt \
        ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.bin \
```

(continues on next page)

(continued from previous page)

```

./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_19364697.
wav \
2 \
$method
done

```

You should see the following output:

```

RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80),
    model_config=ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-fr-2023-04-
    14/encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-
    fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-ncnn-streaming-
    zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./sherpa-ncnn-
    streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./
    sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.param",
    joiner_bin="./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.
    bin", tokens="./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/tokens.txt", encoder_num_
    threads=2, decoder_num_threads=2, joiner_num_threads=2), decoder_
    config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_
    config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
    silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, u
    min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_
    nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
    endpoint=False)
wav filename: ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_-
19364697.wav
wav duration (s): 7.128
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_
voice_fr_19364697.wav
text: CE SITE CONTIENT QUATRE TOMBEAUX DE LA DYNASTIE ASHÉMÉNIDE ET SEPT DES SASSANDIDES
timestamps: 0.96 1.44 1.52 1.76 1.96 2.08 2.28 2.56 2.64 2.76 2.8 2.96 3.04 3.2 3.28 3.4.
3.48 3.72 3.8 4 4.16 4.24 4.32 4.44 4.6 4.68 4.92 5.2 5.52 5.84 6.04 6.12 6.24 6.56 6.
68
Elapsed seconds: 1.082 s
Real time factor (RTF): 1.082 / 7.128 = 0.152
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80),
    model_config=ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-fr-2023-04-
    14/encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-
    fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-ncnn-streaming-
    zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./sherpa-ncnn-
    streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./
    sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.param",
    joiner_bin="./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.
    bin", tokens="./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/tokens.txt", encoder_num_
    threads=2, decoder_num_threads=2, joiner_num_threads=2), decoder_
    config=DecoderConfig(method="modified_beam_search", num_active_paths=4), endpoint_
    config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
    silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, u
    min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_
    nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
    endpoint=False)

```

(continued from previous page)

```
wav filename: ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_
↪19364697.wav
wav duration (s): 7.128
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_
↪voice_fr_19364697.wav
text: CE SITE CONTIENT QUATRE TOMBEAUX DE LA DYNASTIE ASHÉMÉNIDE ET SEPT DES SASSANDIDES
timestamps: 0.96 1.44 1.52 1.76 1.96 2.08 2.28 2.56 2.64 2.76 2.8 2.96 3.04 3.2 3.28 3.4_
↪3.48 3.72 3.8 4 4.16 4.24 4.32 4.44 4.6 4.68 4.92 5.2 5.52 5.84 6.04 6.12 6.24 6.56 6.
↪68
Elapsed seconds: 0.812 s
Real time factor (RTF): 0.812 / 7.128 = 0.114
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-ncnn
./build/bin/sherpa-ncnn-microphone \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/tokens.txt \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/encoder_jit_trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/decoder_jit_trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/joiner_jit_trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_19364697.wav_
↪ \
  2 \
greedy_search
```

Hint: If your system is Linux (including embedded Linux), you can also use `sherpa-ncnn-alsa` to do real-time speech recognition with your microphone if `sherpa-ncnn-microphone` does not work for you.

7.9.3 LSTM-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-ncnn` before you read this section.

marcoyang/sherpa-ncnn-lstm-transducer-small-2023-02-13 (Bilingual, Chinese + English)

This model is a small version of `lstm-transducer` trained in `icefall`.

It only has 13.3 million parameters and can be deployed on embedded devices for real-time speech recognition. You can find the models in fp16 format at <https://huggingface.co/marcoyang/sherpa-ncnn-lstm-transducer-small-2023-02-13>.

The model is trained on a bi-lingual dataset `tal_csasr` (Chinese + English), so it can be used for both Chinese and English.

In the following, we show you how to download it and deploy it with `sherpa-ncnn`.

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn  
  
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-lstm-  
→transducer-small-2023-02-13.tar.bz2  
tar xvf sherpa-ncnn-lstm-transducer-small-2023-02-13.tar.bz2
```

Note: Please refer to *Embedded Linux (arm)* for how to compile `sherpa-ncnn` for a 32-bit ARM platform.

Decode a single wave file with `./build/bin/sherpa-ncnn`

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn  
  
. ./build/bin/sherpa-ncnn \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/tokens.txt \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/encoder_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/encoder_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/decoder_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/decoder_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/joiner_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/joiner_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-lstm-transducer-small-2023-02-13/test_wavs/0.wav
```

Note: The default option uses 4 threads and `greedy_search` for decoding.

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

csukuangfi/sherpa-ncnn-2022-09-05 (English)

This is a model trained using the GigaSpeech and the LibriSpeech dataset.

Please see <https://github.com/k2-fsa/icefall/pull/558> for how the model is trained.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/lstm_transducer_stateless2

In the following, we describe how to download it and use it with `sherpa-ncnn`.

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-2022-09-
↪05.tar.bz2
tar xvf sherpa-ncnn-2022-09-05.tar.bz2
```

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
    ./build/bin/sherpa-ncnn \
        ./sherpa-ncnn-2022-09-05/tokens.txt \
        ./sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-2022-09-05/joiner_jit_trace-pnnx.ncnn.param \
        ./sherpa-ncnn-2022-09-05/joiner_jit_trace-pnnx.ncnn.bin \
        ./sherpa-ncnn-2022-09-05/test_wavs/1089-134686-0001.wav \
    2 \
    $method
done
```

You should see the following output:

```
ModelConfig(encoder_param="../sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.param",
↪encoder_bin="../sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=
↪"../sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="../sherpa-
↪ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="../sherpa-ncnn-2022-09-
↪05/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="../sherpa-ncnn-2022-09-05/joiner_jit_
↪trace-pnnx.ncnn.bin", tokens="../sherpa-ncnn-2022-09-05/tokens.txt", encoder_num_
↪threads=2, decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False,
↪endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
↪trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
↪nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),
↪rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
↪utterance_length=20)))
```

(continues on next page)

(continued from previous page)

```
wav filename: ./sherpa-ncnn-2022-09-05/test_wavs/1089-134686-0001.wav
wav duration (s): 6.625
Started!
Done!
Recognition result for ./sherpa-ncnn-2022-09-05/test_wavs/1089-134686-0001.wav
  AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
  ↵ QUARTER OF THE BROTHELS
ModelConfig(encoder_param="../sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.param",_
  ↵ encoder_bin="../sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=_
  ↵ "../sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="../sherpa-_
  ↵ -ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="../sherpa-ncnn-2022-09-_
  ↵ 05/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="../sherpa-ncnn-2022-09-05/joiner_jit_-
  ↵ trace-pnnx.ncnn.bin", tokens="../sherpa-ncnn-2022-09-05/tokens.txt", encoder_num_
  ↵ threads=2, decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,_
  ↵ endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
  ↵ trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
  ↵ nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0),_
  ↵ rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
  ↵ utterance_length=20)))
wav filename: ./sherpa-ncnn-2022-09-05/test_wavs/1089-134686-0001.wav
wav duration (s): 6.625
Started!
Done!
Recognition result for ./sherpa-ncnn-2022-09-05/test_wavs/1089-134686-0001.wav
  AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
  ↵ QUARTER OF THE BROTHELS
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
  ./.sherpa-ncnn-2022-09-05/tokens.txt \
  ./.sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.param \
  ./.sherpa-ncnn-2022-09-05/encoder_jit_trace-pnnx.ncnn.bin \
  ./.sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.param \
  ./.sherpa-ncnn-2022-09-05/decoder_jit_trace-pnnx.ncnn.bin \
  ./.sherpa-ncnn-2022-09-05/joiner_jit_trace-pnnx.ncnn.param \
  ./.sherpa-ncnn-2022-09-05/joiner_jit_trace-pnnx.ncnn.bin \
  2 \
  greedy_search
```

Note: Please use `./build/bin/Release/sherpa-ncnn-microphone.exe` for Windows.

It will print something like below:

```
Number of threads: 4
num devices: 4
```

(continues on next page)

(continued from previous page)

```
Use default device: 2
  Name: MacBook Pro Microphone
  Max input channels: 1
Started
```

Speak and it will show you the recognition result in real-time.

You can find a demo below:

<https://youtu.be/m6ynSxycpX0>

csukuangfj/sherpa-ncnn-2022-09-30 (Chinese)

This is a model trained using the [WenetSpeech](#) dataset.

Please see <https://github.com/k2-fsa/icefall/pull/595> for how the model is trained.

In the following, we describe how to download it and use it with [sherpa-ncnn](#).

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn

wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-2022-09-
˓→30.tar.bz2
tar xvf sherpa-ncnn-2022-09-30.tar.bz2
```

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

for method in greedy_search modified_beam_search; do
./build/bin/sherpa-ncnn \
./sherpa-ncnn-2022-09-30/tokens.txt \
./sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-2022-09-30/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-2022-09-30/joiner_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-2022-09-30/test_wavs/0.wav \
2 \
$method
done
```

You should see the following output:

```
ModelConfig(encoder_param=".sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.param",
˓→encoder_bin=".sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=
˓→".sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.param", decoder_bin=".sherpa-
˓→ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.bin", joiner_param=".sherpa-ncnn-2022-09-
˓→30/joiner_jit_trace-pnnx.ncnn.param", joiner_bin=".sherpa-ncnn-2022-09-30/joiner_jit_
˓→trace-pnnx.ncnn.bin", tokens=".sherpa-ncnn-2022-09-30/tokens.txt", encoder_num_
˓→threads=2, decoder_num_threads=2, joiner_num_threads=2)
```

(continues on next page)

(continued from previous page)

```

DecoderConfig(method="greedy_search", num_active_paths=4, enable_endpoint=False,
    ↵ endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
    ↵ trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
    ↵ nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), ↵
    ↵ rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
    ↵ utterance_length=20)))
wav filename: ./sherpa-ncnn-2022-09-30/test_wavs/0.wav
wav duration (s): 5.61462
Started!
Done!
Recognition result for ./sherpa-ncnn-2022-09-30/test_wavs/0.wav

ModelConfig(encoder_param="../sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.param",
    ↵ encoder_bin="../sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.bin", decoder_param=
    ↵ "../sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="../sherpa-
    ↵ -ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.bin", joiner_param="../sherpa-ncnn-2022-09-
    ↵ -30/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="../sherpa-ncnn-2022-09-30/joiner_jit-
    ↵ -trace-pnnx.ncnn.bin", tokens="../sherpa-ncnn-2022-09-30/tokens.txt", encoder_num_
    ↵ threads=2, decoder_num_threads=2, joiner_num_threads=2)
DecoderConfig(method="modified_beam_search", num_active_paths=4, enable_endpoint=False,
    ↵ endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
    ↵ trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
    ↵ nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), ↵
    ↵ rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
    ↵ utterance_length=20)))
wav filename: ./sherpa-ncnn-2022-09-30/test_wavs/0.wav
wav duration (s): 5.61462
Started!
Done!
Recognition result for ./sherpa-ncnn-2022-09-30/test_wavs/0.wav

```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

```

cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn-microphone \
    ./.sherpa-ncnn-2022-09-30/tokens.txt \
    ./.sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.param \
    ./.sherpa-ncnn-2022-09-30/encoder_jit_trace-pnnx.ncnn.bin \
    ./.sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.param \
    ./.sherpa-ncnn-2022-09-30/decoder_jit_trace-pnnx.ncnn.bin \
    ./.sherpa-ncnn-2022-09-30/joiner_jit_trace-pnnx.ncnn.param \
    ./.sherpa-ncnn-2022-09-30/joiner_jit_trace-pnnx.ncnn.bin \
    2 \
    greedy_search

```

Note: Please use `./build/bin/Release/sherpa-ncnn-microphone.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You can find a demo below:

```
https://youtu.be/bbQfoRT75oM
```

7.9.4 Conv-Emformer-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-ncnn` before you read this section.

[marcoyang/sherpa-ncnn-conv-emformer-transducer-small-2023-01-09 \(English\)](#)

This model is a small version of `conv-emformer-transducer` trained in `icefall`.

It only has 8.8 million parameters and can be deployed on embedded devices for real-time speech recognition. You can find the models in `fp16` and `int8` format at <https://huggingface.co/marcoyang/sherpa-ncnn-conv-emformer-transducer-small-2023-01-09>.

This model is trained using `LibriSpeech` and thus it supports only English.

In the following, we show you how to download it and deploy it with `sherpa-ncnn` on an embedded device, whose CPU is `RV1126` (Quad core ARM Cortex-A7)

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-conv-
      ↵emformer-transducer-small-2023-01-09.tar.bz2
tar xvf sherpa-ncnn-conv-emformer-transducer-small-2023-01-09.tar.bz2
```

Note: Please refer to [Embedded Linux \(arm\)](#) for how to compile `sherpa-ncnn` for a 32-bit ARM platform. In the following, we test the pre-trained model on an embedded device, whose CPU is `RV1126` (Quad core ARM Cortex-A7).

Decode a single wave file with ./build/bin/sherpa-ncnn

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/tokens.txt \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-pnnx.ncnn.
    ↵param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-pnnx.ncnn.
    ↵bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-pnnx.ncnn.
    ↵param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-pnnx.ncnn.
    ↵bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-pnnx.ncnn.
    ↵param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-pnnx.ncnn.bin \
    ↵\
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/test_wavs/1089-134686-0001.wav
    ↵\
```

The outputs are shown below. The CPU used for decoding is RV1126 (Quad core ARM Cortex-A7).

```
[root@RV1126_RV1109:/asr/small]# time ./sherpa-ncnn tokens.txt encoder_jit_trace
-pnnx.ncnn.param encoder_jit_trace-pnnx.ncnn.bin decoder_jit_trace-pnnx.ncnn.par
am decoder_jit_trace-pnnx.ncnn.bin joiner_jit_trace-pnnx.ncnn.param joiner_jit_t
race-pnnx.ncnn.bin ./test_wavs/1089-134686-0001.wav 4 greedy_search
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param="encoder_jit_trace-pnnx.ncnn.param", encoder_bin="encoder_jit_trace-pnnx.ncnn.bin", decoder_param="decoder_jit_
_trace-pnnx.ncnn.param", decoder_bin="decoder_jit_trace-pnnx.ncnn.bin", joiner_param="joiner_jit_trace-pnnx.ncnn.param", joiner_bin="joi
ner_jit_trace-pnnx.ncnn.bin", tokens="tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4)
wav filename: ./test_wavs/1089-134686-0001.wav
wav duration (s): 6.625
Recognition result for ./test_wavs/1089-134686-0001.wav
    AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLES
real    0m 3.26s
user    0m 10.75s
sys     0m 0.23s
```

Note: The default option uses 4 threads and greedy_search for decoding.

Note: Please use ./build/bin/Release/sherpa-ncnn.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

Decode a single wave file with ./build/bin/sherpa-ncnn (with int8 quantization)

Note: We also support int8 quantization to compresss the model and speed up inference. Currently, only encoder and joiner are quantized.

To decode the int8-quantized model, use the following command:

```
cd /path/to/sherpa-ncnn

./build/bin/sherpa-ncnn \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/tokens.txt \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-pnnx.ncnn.
    ↵int8.param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/encoder_jit_trace-pnnx.ncnn.
    ↵int8.bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-pnnx.ncnn.
    ↵param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/decoder_jit_trace-pnnx.ncnn.
    ↵bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-pnnx.ncnn.
    ↵int8.param \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/joiner_jit_trace-pnnx.ncnn.
    ↵int8.bin \
    ./sherpa-ncnn-conv-emformer-transducer-small-2023-01-09/test_wavs/1089-134686-0001.wav
    ↵\
```

The outputs are shown below. The CPU used for decoding is RV1126 (Quad core ARM Cortex-A7).

```
[root@RV1126_RV1109:/asr/small]# time ./sherpa-ncnn tokens.txt encoder_jit_trace
-pnnx.ncnn.int8.param encoder_jit_trace-pnnx.ncnn.int8.bin decoder_jit_trace-pnn
x.ncnn.param decoder_jit_trace-pnnx.ncnn.bin joiner_jit_trace-pnnx.ncnn.int8.par
am joiner_jit_trace-pnnx.ncnn.int8.bin ./test_wavs/1089-134686-0001.wav 4 greedy
_search
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
ModelConfig(encoder_param="encoder_jit_trace-pnnx.ncnn.int8.param", encoder_bin="encoder_jit_trace-pnnx.ncnn.int8.bin", decoder_param="d
ecoder_jit_trace-pnnx.ncnn.param", decoder_bin="decoder_jit_trace-pnnx.ncnn.bin", joiner_param="joiner_jit_trace-pnnx.ncnn.int8.param",
joiner_bin="joiner_jit_trace-pnnx.ncnn.int8.bin", tokens="tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4)
wav filename: ./test_wavs/1089-134686-0001.wav
wav duration (s): 6.625
Recognition result for ./test_wavs/1089-134686-0001.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLES
real    0m 2.44s
user    0m 7.89s
sys     0m 0.23s
```

Compared to the original model in fp16 format, the decoding speed is significantly improved. The decoding time is changed from 3.26 s to 2.44 s.

Note: When the model's weights are quantized to float16, it is converted to float32 during computation.

When the model's weights are quantized to int8, it is using int8 during computation.

Hint: Even if we use only 1 thread for the int8 model, the resulting real time factor (RTF) is still less than 1.

csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-06 (Chinese + English)

This model is converted from <https://huggingface.co/ptrnull/icefall-asr-conv-emformer-transducer-stateless2-zh>, which supports both Chinese and English.

Hint: If you want to train your own model that is able to support both Chinese and English, please refer to our training code:

https://github.com/k2-fsa/icefall/tree/master/egs/tal_csasr/ASR

You can also try the pre-trained models in your browser without installing anything by visiting:

<https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>

In the following, we describe how to download and use it with `sherpa-ncnn`.

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn  
  
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-conv-  
-emformer-transducer-2022-12-06.tar.bz2  
tar xvf sherpa-ncnn-conv-emformer-transducer-2022-12-06.tar.bz2
```

Decode a single wave file with `./build/bin/sherpa-ncnn`

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn  
  
. ./build/bin/sherpa-ncnn \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-pnnx.ncnn.param \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-pnnx.ncnn.bin \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/0.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/1.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/2.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/3.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/4.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/5.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/6.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/7.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/8.wav \  
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/test_wavs/9.wav
```

Note: Please use `./build/bin/Release/sherpa-ncnn.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

Real-time speech recognition from a microphone with build/bin/sherpa-ncnn-microphone

```
cd /path/to/sherpa-ncnn
./build/bin/sherpa-ncnn-microphone \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/tokens.txt \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-conv-emformer-transducer-2022-12-06/joiner_jit_trace-pnnx.ncnn.bin
```

Note: Please use ./build/bin/Release/sherpa-ncnn-microphone.exe for Windows.

It will print something like below:

```
Number of threads: 4
num devices: 4
Use default device: 2
  Name: MacBook Pro Microphone
  Max input channels: 1
Started
```

Speak and it will show you the recognition result in real-time.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-08 (Chinese)

Hint: This is a very small model that can be run in real-time on embedded systems.

This model is trained using [WenetSpeech](#) dataset and it supports only Chinese.

In the following, we describe how to download and use it with [sherpa-ncnn](#).

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn

wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-conv-
      -emformer-transducer-2022-12-08.tar.bz2
tar xvf sherpa-ncnn-conv-emformer-transducer-2022-12-08.tar.bz2
```

Decode a single wave file with ./build/bin/sherpa-ncnn

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn  
  
./build/bin/sherpa-ncnn \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/tokens.txt \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/encoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/encoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/decoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/decoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/joiner_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/joiner_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-ncnn.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

Real-time speech recognition from a microphone with build/bin/sherpa-ncnn-microphone

```
cd /path/to/sherpa-ncnn  
./build/bin/sherpa-ncnn-microphone \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/tokens.txt \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/encoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/encoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/decoder_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/decoder_jit_trace-pnnx.ncnn.bin \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/joiner_jit_trace-pnnx.ncnn.param \  
  ./sherpa-ncnn-conv-emformer-transducer-2022-12-08/joiner_jit_trace-pnnx.ncnn.bin
```

Note: Please use ./build/bin/Release/sherpa-ncnn-microphone.exe for Windows.

It will print something like below:

```
Number of threads: 4  
num devices: 4  
Use default device: 2  
  Name: MacBook Pro Microphone  
  Max input channels: 1  
Started
```

Speak and it will show you the recognition result in real-time.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

csukuangfj/sherpa-ncnn-conv-emformer-transducer-2022-12-04 (English)

This model is trained using [GigaSpeech](#) and [LibriSpeech](#). It supports only English.

In the following, we describe how to download and use it with [sherpa-ncnn](#).

Please use the following commands to download it.

```
cd /path/to/sherpa-ncnn
wget https://github.com/k2-fsa/sherpa-ncnn/releases/download/models/sherpa-ncnn-conv-
      ↵emformer-transducer-2022-12-04.tar.bz2
tar xvf sherpa-ncnn-conv-emformer-transducer-2022-12-04.tar.bz2
```

Decode a single wave file with ./build/bin/sherpa-ncnn

Hint: It supports decoding only wave files with a single channel and the sampling rate should be 16 kHz.

```
cd /path/to/sherpa-ncnn
./build/bin/sherpa-ncnn \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/tokens.txt \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/encoder_jit_trace-pnnx.ncnn.param \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/encoder_jit_trace-pnnx.ncnn.bin \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/decoder_jit_trace-pnnx.ncnn.param \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/decoder_jit_trace-pnnx.ncnn.bin \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/joiner_jit_trace-pnnx.ncnn.param \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/joiner_jit_trace-pnnx.ncnn.bin \
    ./sherpa-ncnn-conv-emformer-transducer-2022-12-04/test_wavs/1089-134686-0001.wav
```

Note: Please use ./build/bin/Release/sherpa-ncnn.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

Real-time speech recognition from a microphone with build/bin/sherpa-ncnn-microphone

```
cd /path/to/sherpa-ncnn
./build/bin/sherpa-ncnn-microphone \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/tokens.txt \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/joiner_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-conv-emformer-transducer-2022-12-04/joiner_jit_trace-pnnx.ncnn.param
```

Note: Please use ./build/bin/Release/sherpa-ncnn-microphone.exe for Windows.

It will print something like below:

```
Number of threads: 4
num devices: 4
Use default device: 2
  Name: MacBook Pro Microphone
  Max input channels: 1
Started
```

Speak and it will show you the recognition result in real-time.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

7.10 Examples

In this section, we describe some usage examples of `sherpa-ncnn` on various boards.

7.10.1 Raspberry Pi 3B E14

This page posts some screenshots of running `sherpa-ncnn` on Raspberry Pi 3B E14.

Hint: You can find pre-compiled binaries used in this example at

```
https://huggingface.co/csukuangfj/sherpa-ncnn-pre-compiled-binaries/tree/main/arm32
```

Board info

OS release

```
pi@raspberrypi: ~
File Edit Tabs Help
pi@raspberrypi:~ $ cat /etc/os-release
PRETTY_NAME="Raspbian GNU/Linux 10 (buster)"
NAME="Raspbian GNU/Linux"
VERSION_ID="10"
VERSION="10 (buster)"
VERSION_CODENAME=buster
ID=raspbian
ID_LIKE=debian
HOME_URL="http://www.raspbian.org/"
SUPPORT_URL="http://www.raspbian.org/RaspbianForums"
BUG_REPORT_URL="http://www.raspbian.org/RaspbianBugs"
pi@raspberrypi:~ $
```

lscpu

```
pi@raspberrypi:~ $ lscpu
Architecture:          armv7l
Byte Order:            Little Endian
CPU(s):                4
On-line CPU(s) list:  0-3
Thread(s) per core:   1
Core(s) per socket:   4
Socket(s):             1
Vendor ID:             ARM
Model:                 4
Model name:            Cortex-A53
Stepping:               r0p4
CPU max MHz:           1200.0000
CPU min MHz:           600.0000
BogoMIPS:              57.60
Flags:                 half thumb fastmult vfp edsp neon vfpv3 tls vfpv4 idiva idivt vfpd32 lpae evtstrm crc32
```

cpuinfo

```
File Edit Tabs Help
pi@raspberrypi:~ $ cat /proc/cpuinfo
processor : 0
model name : ARMv7 Processor rev 4 (v7l)
BogoMIPS : 38.40
Features : half thumb fastmult vfp edsp neon vfpv3 tls vfpv4 idiva idivt
vfpd32 lpae evtstrm crc32
CPU implementer : 0x41
CPU architecture: 7
CPU variant : 0x0
CPU part : 0xd03
CPU revision : 4

processor : 1
model name : ARMv7 Processor rev 4 (v7l)
BogoMIPS : 38.40
Features : half thumb fastmult vfp edsp neon vfpv3 tls vfpv4 idiva idivt
vfpd32 lpae evtstrm crc32
CPU implementer : 0x41
CPU architecture: 7
CPU variant : 0x0
CPU part : 0xd03
CPU revision : 4

processor : 2
model name : ARMv7 Processor rev 4 (v7l)
BogoMIPS : 38.40
Features : half thumb fastmult vfp edsp neon vfpv3 tls vfpv4 idiva idivt
vfpd32 lpae evtstrm crc32
CPU implementer : 0x41
CPU architecture: 7
CPU variant : 0x0
CPU part : 0xd03
CPU revision : 4

processor : 3
model name : ARMv7 Processor rev 4 (v7l)
BogoMIPS : 38.40
Features : half thumb fastmult vfp edsp neon vfpv3 tls vfpv4 idiva idivt
vfpd32 lpae evtstrm crc32
CPU implementer : 0x41
CPU architecture: 7
CPU variant : 0x0
CPU part : 0xd03
CPU revision : 4

Hardware : BCM2835
Revision : a02082
Serial : 0000000e3495009
pi@raspberrypi:~ $
```

RTF (1 thread)

```
pi@raspberrypi: ~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23
File Edit Tabs Help
dpointConfig(rule1=EndpointRule(must_contain_nonsense=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsense=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsense=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 5.6115
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 对我做了介绍那么我想说的是大家如果对我的研究感兴趣
timestamps:
Elapsed seconds: 2.520 s
Real time factor (RTF): 2.520 / 5.611 = 0.449
pi@raspberrypi:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 $ ./sherpa-ncnn ./tokens.txt ./encoder_jit_trace-pnnx.ncnn.param ./encoder_jit_trace-pnnx.ncnn.bin ./decoder_jit_trace-pnnx.ncnn.param ./decoder_jit_trace-pnnx.ncnn.bin ./joiner_jit_trace-pnnx.ncnn.param ./joiner_jit_trace-pnnx.ncnn.bin ./test_wavs/0.wav
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt"), encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=dpointConfig(rule1=EndpointRule(must_contain_nonsense=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsense=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsense=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 5.6115
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 对我做了介绍那么我想说的是大家如果对我的研究感兴趣
timestamps:
Elapsed seconds: 2.021 s
Real time factor (RTF): 2.021 / 5.611 = 0.360
pi@raspberrypi:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 $
```

RTF (2 threads)

```
pi@raspberrypi: ~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23
File Edit Tabs Help
Please refer to
https://k2-fsa.github.io/sherpa/ncnn/pretrained\_models/index.html
for a list of pre-trained models to download.

pi@raspberrypi:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 $ ./sherpa-ncnn \
> ./tokens.txt \
> ./encoder_jit_trace-pnnx.ncnn.param \
> ./encoder_jit_trace-pnnx.ncnn.bin \
> ./decoder_jit_trace-pnnx.ncnn.param \
> ./decoder_jit_trace-pnnx.ncnn.bin \
> ./joiner_jit_trace-pnnx.ncnn.param \
> ./joiner_jit_trace-pnnx.ncnn.bin \
> ./test_wavs/0.wav \
> 2
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt"), encoder_num_threads=2, decoder_num_threads=2, joiner_num_threads=2), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=dpointConfig(rule1=EndpointRule(must_contain_nonsense=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsense=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsense=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 5.6115
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 对我做了介绍那么我想说的是大家如果对我的研究感兴趣
timestamps:
Elapsed seconds: 2.520 s
Real time factor (RTF): 2.520 / 5.611 = 0.449
pi@raspberrypi:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23 $
```

7.10.2 Jetson Nano

This page posts some screenshots of running sherpa-ncnn on Jetson Nano.

Hint: You can find pre-compiled binaries used in this example at

<https://huggingface.co/csukuangfj/sherpa-ncnn-pre-compiled-binaries/tree/main/aarch64>

Board info

```
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$ lscpu
Architecture:          aarch64
Byte Order:            Little Endian
CPU(s):                4
On-line CPU(s) list:  0-3
Thread(s) per core:   1
Core(s) per socket:   4
Socket(s):             1
Vendor ID:             ARM
Model:                 1
Model name:            Cortex-A57
Stepping:               r1p1
CPU max MHz:           1479.0000
CPU min MHz:           102.0000
BogoMIPS:              38.40
L1d cache:             32K
L1i cache:             48K
L2 cache:              2048K
Flags:                 fp asimd evtstrm aes pmull sha1 sha2 crc32
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$ lsb_release -a
No LSB modules are available.
Distributor ID:        Ubuntu
Description:            Ubuntu 18.04.6 LTS
Release:               18.04
Codename:              bionic
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$ uname -a
Linux nikki-desktop 4.9.253-tegra #1 SMP PREEMPT Sat Feb 19 08:59:22 PST 2022 aarch64 aarch64 aarch64 GNU/Linux
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$
```

RTF (4 threads)

```
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$ ./sherpa-ncnn ./tokens.txt ./encoder_jit_trace-pnnx.ncnn.param ./encoder_jit_trace-pnnx.ncnn.bin ./decoder_jit_trace-pnnx.ncnn.param ./decoder_jit_trace-pnnx.ncnn.bin ./joiner_jit_trace-pnnx.ncnn.param ./joiner_jit_trace-pnnx.ncnn.bin ./test_wavs/0.wav 4
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 5.6115
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 对我做了介绍那么我想说的是大家如果对我的研究感兴趣
timestamps:
Elapsed seconds: 0.845 s
Real time factor (RTF): 0.845 / 5.611 = 0.151
nikki@nikki-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23$
```

7.10.3 Jetson NX

This page posts some screenshots of running sherpa-ncnn on Jetson NX.

Hint: You can find pre-compiled binaries used in this example at

<https://huggingface.co/csukuangfj/sherpa-ncnn-pre-compiled-binaries/tree/main/aarch64>

Board info

```
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ lscpu
Architecture:           aarch64
Byte Order:             Little Endian
CPU(s):                6
On-line CPU(s) list:   0-5
Thread(s) per core:    1
Core(s) per socket:    2
Socket(s):              3
Vendor ID:              Nvidia
Model:                 0
Model name:             ARMv8 Processor rev 0 (v8l)
Stepping:               0x0
CPU max MHz:            1907.2000
CPU min MHz:            115.2000
BogoMIPS:               62.50
L1i cache:              64K
L1i cache:              128K
L2 cache:               2048K
L3 cache:               4996K
Flags:                  fp asimd evtstrm aes pmull sha1 sha2 crc32 atomics fphp asimdhfp
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ lsb_release -a
No LSB modules are available.
Distributor ID:        Ubuntu
Description:             Ubuntu 18.04.6 LTS
Release:                18.04
Codename:               bionic
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ uname -a
Linux e2ai-desktop 4.9.253-tegra #1 SMP PREEMPT Sat Feb 19 08:58:27 PST 2022 aarch64 aarch64 aarch64 GNU/Linux
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ █
```

RTF (2 threads)

```
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ ./sherpa-ncnn \
>   ./tokens.txt \
>   ./encoder_jit_trace-pnnx.ncnn.param \
>   ./encoder_jit_trace-pnnx.ncnn.bin \
>   ./decoder_jit_trace-pnnx.ncnn.param \
>   ./decoder_jit_trace-pnnx.ncnn.bin \
>   ./joiner_jit_trace-pnnx.ncnn.param \
>   ./joiner_jit_trace-pnnx.ncnn.bin \
>   ./test_wavs/0.wav \
2>   ./decoder_jit_trace-pnnx.ncnn.param \
>   ./decoder_jit_trace-pnnx.ncnn.bin \
>   ./joiner_jit_trace-pnnx.ncnn.param \
>   ./joiner_jit_trace-pnnx.ncnn.bin \
>   ./test_wavs/0.wav \
>   2
Disable fp16 for zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt", encoder_num_threads=2, decoder_num_threads=2, joiner_num_threads=2), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 10.0531
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 昨天是 MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW是星期三
timestamps:
Elapsed seconds: 3.661 s
Real time factor (RTF): 3.661 / 10.053 = 0.364
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ █
```

RTF (4 threads)

```
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ ./sherpa-ncnn ./tokens.txt ./encoder_jit_trace-pnnx.ncnn.param ./encoder_jit_trace-pnnx.ncnn.bin ./decoder_jit_trace-pnnx.ncnn.param ./decoder_jit_trace-pnnx.ncnn.bin ./joiner_jit_trace-pnnx.ncnn.param ./joiner_jit_trace-pnnx.ncnn.bin ./test_wavs/0.wav 4
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4), decoder_config=DecoderConfig(method="greedy_search", num_active_patterns=4), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 10.0531
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 昨天是 MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW是星期三
timestamps:
Elapsed seconds: 3.163 s
Real time factor (RTF): 3.163 / 10.053 = 0.315
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ █
```

RTF (6 threads)

```
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ ./sherpa-ncnn ./tokens.txt ./encoder_jit_trace-pnnx.ncnn.param ./encoder_jit_trace-pnnx.ncnn.bin ./decoder_jit_trace-pnnx.ncnn.param ./decoder_jit_trace-pnnx.ncnn.bin ./joiner_jit_trace-pnnx.ncnn.param ./joiner_jit_trace-pnnx.ncnn.bin ./test_wavs/0.wav 6
Disable fp16 for Zipformer encoder
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80, max_feature_vectors=-1), model_config=ModelConfig(encoder_param="./encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./encoder_jit_trace-pnnx.ncnn.bin", decoder_param="./decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./decoder_jit_trace-pnnx.ncnn.bin", joiner_param="./joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./joiner_jit_trace-pnnx.ncnn.bin", tokens="./tokens.txt", encoder_num_threads=6, decoder_num_threads=6, joiner_num_threads=6), decoder_config=DecoderConfig(method="greedy_search", num_active_patterns=6), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./test_wavs/0.wav
wav duration (s): 10.0531
Started!
Done!
Recognition result for ./test_wavs/0.wav
text: 昨天是 MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW是星期三
timestamps:
Elapsed seconds: 3.854 s
Real time factor (RTF): 3.854 / 10.053 = 0.383
e2ai@e2ai-desktop:~/Downloads/sherpa-ncnn/sherpa-ncnn-streaming-zipformer-bilingual-zh-en-2023-02-13$ █
```

7.10.4 VisionFive 2

This page describes how to run `sherpa-ncnn` on `VisionFive2`, which is a 64-bit RISC-V board with 4 CPUs.

Hint: You can find pre-compiled binaries used in this example at

<https://huggingface.co/csukuangfj/sherpa-ncnn-pre-compiled-binaries/tree/main/riscv64>

Caution: The latest debian image from https://doc-en.rvspace.org/VisionFive2/Quick_Start_Guide/VisionFive2_QSG/flashing_with_mac_linux.html does not work since it does not support USB devices.

That is, you cannot use USB microphones on the board with the above debian image.

Note: We have compiled <https://github.com/starfive-tech/VisionFive2> and the resulting `sdcard.img` is available at <https://huggingface.co/csukuangfj/visionfive2-sd-card-img>.

Please use this image for testing. It supports USB microphones.

The username for this image is **root** and the password is **starfive**.

Board info

```
on the cross build 3
root@buildroot:~# lscpu
Architecture:          riscv64
Byte Order:            Little Endian
CPU(s):                4
On-line CPU(s) list: 0-3
root@buildroot:~# cat /proc/cpuinfo
processor      : 0
hart          : 1
isa           : rv64imafdc
mmu           : sv39
isa-ext       :
uarch         : sifive,u74-mc

processor      : 1
hart          : 2
isa           : rv64imafdc
mmu           : sv39
isa-ext       :
uarch         : sifive,u74-mc

processor      : 2
hart          : 3
isa           : rv64imafdc
mmu           : sv39
isa-ext       :
uarch         : sifive,u74-mc

processor      : 3
hart          : 4
isa           : rv64imafdc
mmu           : sv39
isa-ext       :
uarch         : sifive,u74-mc
```

RTF (4 threads)

We use [csukuangji/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16](https://github.com/csukuangji/sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16) (Bilingual, Chinese + English) for testing. The RTF is given below:

```

root@buildroot:next-gen-kaldi# cat run-sherpa-ncnn.sh
#!/usr/bin/env bash

./sherpa-ncnn \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav \
4 \
greedy_search

root@buildroot:next-gen-kaldi# ./run-sherpa-ncnn.sh
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
RecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80), model_config=ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_trace-pnnx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnnx.ncnn.bin", decoder_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnnx.ncnn.param", decoder_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnnx.ncnn.param", joiner_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnnx.ncnn.param", joiner_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnnx.ncnn.param", tokens="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.4, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=False)
wav filename: ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
wav duration (s): 10.0531
Started!
Done!
Recognition result for ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
text: 昨天是 MONDAY TODAY IS 李八二 TODAY AFTER TOMORROW 是星期三
timestamps: 0.64 1.08 1.6 2.08 2.24 2.36 4.04 4.28 5.12 5.48 5.68 6.12 6.8 6.96 7.28 7.92 8.04 8.2 8.28 9.04 9.36 9.6 10.04
Elapsed seconds: 9.334 s
Real time factor (RTF): 9.334 / 10.053 = 0.928

```

You can see that the RTF is less than 1, which means it is able to perform streaming (i.e., real-time) speech recognition. The following posts the commands used for testing so that you can copy and paste them if you want to test it by yourself.

```

./sherpa-ncnn \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_\
trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_\
trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_\
trace-pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_\
trace-pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-\
pnnx.ncnn.param \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-\
pnnx.ncnn.bin \
./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav \
4 \
greedy_search

```

Real-time speech recognition with a microphone

Since the board does not have microphones, we use a USB microphone for testing.

Caution: We use the image from <https://huggingface.co/csukuangfj/visionfive2-sd-card-img/tree/main>, which provides support for USB microphones.

After connecting a USB microphone to the board, use the following command to check it:

```
root@buildroot:next-gen-kaldi# arecord -l
**** List of CAPTURE Hardware Devices ****
card 2: UACDemoV10 [UACDemoV1.0], device 0: USB Audio [USB Audio]
  Subdevices: 1/1
  Subdevice #0: subdevice #0
```

The output shows Card 2 and device 0, so the device name is hw:2,0.

The command to start the program for real-time speech recognition is

```
./sherpa-ncnn-alsa \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_
→ trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_
→ trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_
→ trace-pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_
→ trace-pnnx.ncnn.bin \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-
→ pnnx.ncnn.param \
  ./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-
→ pnnx.ncnn.bin \
  hw:2,0 \
  4 \
  greedy_search
```

A screenshot is given below:

```

root@buildroot:next-gen-kaldi# ./run-sherpa-ncnn-alsa.sh
./sherpa-ncnn-alsa: /usr/lib/libasound.so.2: no version information available (required by ./sherpa-ncnn-alsa)
./sherpa-ncnn-alsa: /usr/lib/libasound.so.2: no version information available (required by ./sherpa-ncnn-alsa)
RecognizerConfig{feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80), model_config=ModelConfig(encoder_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_trace-pnmx.ncnn.param", encoder_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/encoder_jit_trace-pnmx.ncnn.param", decoder_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnmx.ncnn.param", decoder_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/decoder_jit_trace-pnmx.ncnn.param", joiner_param="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/64/joiner_jit_trace-pnmx.ncnn.param", joiner_bin="./sherpa-ncnn-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt", encoder_num_threads=4, decoder_num_threads=4, joiner_num_threads=4), decoder_config=DecoderConfig(method="greedy_search", num_active_paths=4), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=300)), enable_endpoint=True)
Don't Use GPU. has_gpu: 0, config.use_vulkan_compute: 1
Failed to set sample rate to 16000
Current sample rate is 48000
Creating a resampler:
    in_sample_rate: 48000
    output_sample_rate: 16000
Recording started!
Use recording device: hw:2,0
0:星光二上的语音识别
1:这是一块基于 risk five
2:的开发版

```

7.11 FAQs

7.11.1 Where to get help

If you have any questions, please create an issue at <https://github.com/k2-fsa/sherpa-ncnn>

We also have active social groups:

- : Kaldi
- Kaldi, ,
- QQ 744602236

7.11.2 No default input device found

If you are using Linux and if `sherpa-ncnn-microphone` throws the following error:

```

Num device: 0
No default input device found.

```

Please consider using `sherpa-ncnn-alsa` to replace `sherpa-ncnn-microphone`. If you cannot find `sherpa-ncnn-alsa` in `./build/bin`, please run the following commands:

```

cd /path/to/sherpa-ncnn
sudo apt-get install alsa-utils libasound2-dev
cd build
rm CMakeCache.txt # Important, remove the cmake cache file
make -j

```

After the above commands, you should see a binary file `./build/bin/sherpa-ncnn-alsa`.

Please follow `sherpa-ncnn-alsa` to use `sherpa-ncnn-alsa`.

SHERPA-ONNX

Hint: During speech recognition, it does not need to access the Internet. Everything is processed locally on your device.

We support using `onnx` with `onnxruntime` to replace `PyTorch` for neural network computation. The code is put in a separate repository `sherpa-onnx`.

`sherpa-onnx` is self-contained and everything can be compiled from source.

Please refer to <https://k2-fsa.github.io/icefall/model-export/export-onnx.html> for how to export models to `onnx` format.

In the following, we describe how to build `sherpa-onnx` for Linux, macOS, Windows, embedded systems, Android, and iOS.

Also, we show how to use it for speech recognition with pre-trained models.

8.1 Tutorials

This page contains links to tutorials written by our users.

Caution: The tutorials are not necessarily written in English.

8.2 Installation

In this section, we describe how to install `sherpa-onnx` on various platforms.

Requirements:

- CMake >= 3.13
- A compiler that supports at least C++14

Hint: You can use `pip install cmake` to install the latest cmake.

8.2.1 Linux

This page describes how to build `sherpa-onnx` on Linux.

All you need is to run:

CPU

Nvidia GPU (CUDA)

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
make -j6
```

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release -DBUILD_SHARED_LIBS=ON -DSHERPA_ONNX_ENABLE_GPU=ON ..
make -j6
```

Hint: You need to install CUDA toolkit. Otherwise, you would get errors at runtime.

You can refer to <https://k2-fsa.github.io/k2/installation/cuda-cudnn.html> to install CUDA toolkit.

After building, you will find an executable `sherpa-onnx` inside the `bin` directory.

That's it!

Please refer to *Pre-trained models* for a list of pre-trained models.

8.2.2 macOS

This page describes how to build `sherpa-onnx` on macOS.

Hint: It supports both Intel and Apple Silicon (e.g., M1).

All you need is to run:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
make -j6
```

After building, you will find an executable `sherpa-onnx` inside the `bin` directory.

That's it!

Please refer to *Pre-trained models* for a list of pre-trained models.

8.2.3 Windows

This page describes how to build `sherpa-onnx` on Windows.

Hint: MinGW is known not to work. Please install Visual Studio before you continue.

Note: You can download pre-compiled binaries for both 32-bit and 64-bit Windows from the following URL <https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main>.

Please always download the latest version.

URLs to download the version 1.9.12 is given below.

64-bit Windows (static lib)	https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/win64/sherpa-onnx-v1.9.12-win-x64-static.tar.bz2
64-bit Windows (shared lib)	https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/win64/sherpa-onnx-v1.9.12-win-x64-shared.tar.bz2
32-bit Windows (static lib)	https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/win32/sherpa-onnx-v1.9.12-win-x86-static.tar.bz2
32-bit Windows (shared lib)	https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/win32/sherpa-onnx-v1.9.12-win-x86-shared.tar.bz2

If you cannot access `huggingface.co`, then please replace `huggingface.co` with `hf-mirror.com`.

64-bit Windows (x64)

All you need is to run:

CPU

Nvidia GPU (CUDA)

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release ..
cmake --build . --config Release
```

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build
cmake -DCMAKE_BUILD_TYPE=Release -DBUILD_SHARED_LIBS=ON -DSHERPA_ONNX_ENABLE_GPU=ON ..
cmake --build . --config Release
```

Hint: You need to install CUDA toolkit. Otherwise, you would get errors at runtime.

After building, you will find an executable `sherpa-onnx.exe` inside the `bin/Release` directory.

That's it!

Please refer to *Pre-trained models* for a list of pre-trained models.

32-bit Windows (x86)

Hint: It does not support NVIDIA GPU for Win32/x86.

All you need is to run:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build

# Please select one toolset among VS 2015, 2017, 2019, and 2022 below
# We use VS 2022 as an example.

# For Visual Studio 2015
# cmake -T v140,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2017
# cmake -T v141,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2019
# cmake -T v142,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

# For Visual Studio 2022
cmake -T v143,host=x64 -A Win32 -D CMAKE_BUILD_TYPE=Release ..

cmake --build . --config Release
```

After building, you will find an executable `sherpa-onnx.exe` inside the `bin/Release` directory.

That's it!

Please refer to *Pre-trained models* for a list of pre-trained models.

Hint: By default, it builds static libraries of `sherpa-onnx`. To get dynamic/shared libraries, please pass `-DBUILD_SHARED_LIBS=ON` to `cmake`. That is, use

```
cmake -DCMAKE_BUILD_TYPE=Release -DBUILD_SHARED_LIBS=ON ..
```

8.2.4 Embedded Linux (aarch64)

This page describes how to build `sherpa-onnx` for embedded Linux (aarch64, 64-bit) with cross-compiling on an x64 machine with Ubuntu OS.

Warning: By cross-compiling we mean that you do the compilation on a `x86_64` machine. And you copy the generated binaries from a `x86_64` machine and run them on an `aarch64` machine.

If you want to compile `sherpa-onnx` on an `aarch64` machine directly, please see [Linux](#).

Note: You can download pre-compiled binaries for `aarch64` from the following URL <https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/aarch64>

Please always download the latest version.

Example command to download the version 1.9.12:

```
# binaries built with shared libraries
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/aarch64/
↳sherpa-onnx-v1.9.12-linux-aarch64-shared.tar.bz2

# binaries built with static link
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/aarch64/
↳sherpa-onnx-v1.9.12-linux-aarch64-static.tar.bz2

# For users from China
# huggingface,

# binaries built with shared libraries
wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/aarch64/
↳sherpa-onnx-v1.9.12-linux-aarch64-shared.tar.bz2

# binaries built with static link
wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/aarch64/
↳sherpa-onnx-v1.9.12-linux-aarch64-static.tar.bz2
```

Hint: We provide two colab notebooks for you to try this section step by step.

Build with shared libraries	Build with static libraries

If you are using Windows/macOS or you don't want to setup your local environment for cross-compiling, please use the above colab notebooks.

Install toolchain

The first step is to install a toolchain for cross-compiling.

Warning: You can use any toolchain that is suitable for your platform. The toolchain we use below is just an example.

Visit <https://releases.linaro.org/components/toolchain/binaries/latest-7/aarch64-linux-gnu/>

We are going to download `gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz`, which has been uploaded to <https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains>.

Assume you want to install it in the folder `$HOME/software`:

```
mkdir -p $HOME/software
cd $HOME/software
wget https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-linaro-7.
˓→5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz

# For users from China
# huggingface,
# wget https://hf-mirror.com/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-linaro-7.
˓→5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz

tar xvf gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu.tar.xz
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-linaro-7.5.0-2019.12-x86_64_aarch64-linux-gnu/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
aarch64-linux-gnu-gcc --version
```

which should print the following log:

```
aarch64-linux-gnu-gcc (Linaro GCC 7.5-2019.12) 7.5.0
Copyright (C) 2017 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Congratulations! You have successfully installed a toolchain for cross-compiling `sherpa-onnx`.

Build sherpa-onnx

Finally, let us build `sherpa-onnx`.

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
export BUILD_SHARED_LIBS=ON
./build-aarch64-linux-gnu.sh
```

After building, you will get two binaries:

```
sherpa-onnx$ ls -lh build-aarch64-linux-gnu/install/bin/
total 378K
-rwxr-xr-x 1 kuangfangjun root 187K Feb 21 21:55 sherpa-onnx
-rwxr-xr-x 1 kuangfangjun root 191K Feb 21 21:55 sherpa-onnx-alsa
```

Note: Please also copy the `onnxruntime` lib to your embedded systems and put it into the same directory as `sherpa-onnx` and `sherpa-onnx-alsa`.

```
sherpa-onnx$ ls -lh build-aarch64-linux-gnu/install/lib/*onnxruntime*
lrw-r--r-- 1 kuangfangjun root 24 Feb 21 21:38 build-aarch64-linux-gnu/install/lib/
↳ libonnxruntime.so -> libonnxruntime.so.1.14.0
-rw-r--r-- 1 kuangfangjun root 15M Feb 21 21:38 build-aarch64-linux-gnu/install/lib/
↳ libonnxruntime.so.1.14.0
```

That's it!

Hint:

- `sherpa-onnx` is for decoding a single file
- `sherpa-onnx-alsa` is for real-time speech recognition by reading the microphone with [ALSA](#)

sherpa-onnx-alsa

Caution: We recommend that you use `sherpa-onnx-alsa` on embedded systems such as Raspberry pi.

You need to provide a `device_name` when invoking `sherpa-onnx-alsa`. We describe below how to find the device name for your microphone.

Run the following command:

```
arecord -l
```

to list all available microphones for recording. If it complains that `arecord: command not found`, please use `sudo apt-get install alsa-utils` to install it.

If the above command gives the following output:

```
**** List of CAPTURE Hardware Devices ****
card 3: UACDemoV10 [UACDemoV1.0], device 0: USB Audio [USB Audio]
Subdevices: 1/1
Subdevice #0: subdevice #0
```

In this case, I only have 1 microphone. It is card 3 and that card has only device 0. To select card 3 and device 0 on that card, we need to pass `plughw:3,0` to `sherpa-onnx-alsa`. (Note: It has the format `plughw:card_number,device_index`.)

For instance, you have to use

```
./sherpa-onnx-alsa \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/
↪encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/
↪decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/
↪joiner-epoch-99-avg-1.onnx \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/
↪tokens.txt \
plughw:3,0
```

Please change the card number and also the device index on the selected card accordingly in your own situation. Otherwise, you won't be able to record with your microphone.

Please read *Pre-trained models* for usages about the generated binaries.

Hint: If you want to select a pre-trained model for Raspberry that can be run on real-time, we recommend you to use *Zipformer-transducer-based Models*.

Please create an issue at <https://github.com/k2-fsa/sherpa-onnx/issues> if you have any problems.

How to build static libraries and static linked binaries

If you want to build static libraries and static linked binaries, please first download a cross compile toolchain with GCC ≥ 9.0 . The following is an example:

```
mkdir -p $HOME/software
cd $HOME/software
wget -q https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-10.
↪3-2021.07-x86_64-aarch64-none-linux-gnu.tar.xz

# For users from China
# huggingface,
# wget -q https://hf-mirror.com/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-
↪10.3-2021.07-x86_64-aarch64-none-linux-gnu.tar.xz

tar xf gcc-arm-10.3-2021.07-x86_64-aarch64-none-linux-gnu.tar.xz
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-arm-10.3-2021.07-x86_64-aarch64-none-linux-gnu/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
aarch64-none-linux-gnu-gcc --version
```

which should print the following log:

```
aarch64-none-linux-gnu-gcc (GNU Toolchain for the A-profile Architecture 10.3-2021.07
↪(arm-10.29)) 10.3.1 20210621
Copyright (C) 2020 Free Software Foundation, Inc.
```

(continues on next page)

(continued from previous page)

This **is** free software; see the source **for** copying conditions. There **is** NO warranty; **not** even **for** MERCHANTABILITY **or** FITNESS FOR A PARTICULAR PURPOSE.

Now you can build static libraries and static linked binaries with the following commands:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
export BUILD_SHARED_LIBS=OFF
./build-aarch64-linux-gnu.sh
```

You can use the following commands to check that the generated binaries are indeed static linked:

```
$ cd build-aarch64-linux-gnu/bin

$ ldd sherpa-onnx-alsa
    not a dynamic executable

$ readelf -d sherpa-onnx-alsa

Dynamic section at offset 0xed9950 contains 30 entries:
  Tag          Type           Name/Value
 0x0000000000000001 (NEEDED)   Shared library: [libasound.so.2]
 0x0000000000000001 (NEEDED)   Shared library: [libdl.so.2]
 0x0000000000000001 (NEEDED)   Shared library: [libm.so.6]
 0x0000000000000001 (NEEDED)   Shared library: [libpthread.so.0]
 0x0000000000000001 (NEEDED)   Shared library: [libc.so.6]
 0x000000000000000f (RPATH)     Library rpath: [$ORIGIN:/star-fj/fangjun/open-
->source/sherpa-onnx/build-aarch64-linux-gnu/_deps/onnxruntime-sr
c/lib:]
 0x000000000000000c (INIT)      0x404218
```

8.2.5 Embedded Linux (arm)

This page describes how to build `sherpa-onnx` for embedded Linux (arm, 32-bit) with **cross-compiling** on an x86 machine with Ubuntu OS.

Caution: If you want to build `sherpa-onnx` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-onnx` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Caution: If you want to build `sherpa-onnx` directly on your board, please don't use this document. Refer to [Linux](#) instead.

Hint: This page is for cross-compiling.

Note: You can download pre-compiled binaries for 32-bit ARM from the following URL <https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/arm32>

Please always download the latest version.

Example command to download the version 1.9.12:

```
# binaries built with shared libraries
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/arm32/
↳sherpa-onnx-v1.9.12-linux-arm-gnueabihf-shared.tar.bz2

# binaries built with static link
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/arm32/
↳sherpa-onnx-v1.9.12-linux-arm-gnueabihf-static.tar.bz2

# For users from China
# huggingface,

# binaries built with shared libraries
wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/arm32/
↳sherpa-onnx-v1.9.12-linux-arm-gnueabihf-shared.tar.bz2

# binaries built with static link
wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/arm32/
↳sherpa-onnx-v1.9.12-linux-arm-gnueabihf-static.tar.bz2
```

Hint: We provide two colab notebooks for you to try this section step by step.

Build with shared libraries	Build with static libraries

If you are using Windows/macOS or you don't want to setup your local environment for cross-compiling, please use the above colab notebooks.

Install toolchain

The first step is to install a toolchain for cross-compiling.

Warning: You can use any toolchain that is suitable for your platform. The toolchain we use below is just an example.

Visit <https://developer.arm.com/downloads/-/arm-gnu-toolchain-downloads> to download the toolchain:

We are going to download `gcc-arm-10.3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz`, which has been uploaded to <https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains>.

Assume you want to install it in the folder \$HOME/software:

```
mkdir -p $HOME/software
cd $HOME/software
wget -q https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-10.
→3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz

# For users from China
# huggingface,
# wget -q https://hf-mirror.com/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-
→10.3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz

tar xf gcc-arm-10.3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-arm-10.3-2021.07-x86_64-arm-none-linux-gnueabihf/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
arm-none-linux-gnueabihf-gcc --version
```

which should print the following log:

```
arm-none-linux-gnueabihf-gcc (GNU Toolchain for the A-profile Architecture 10.3-2021.07
→(arm-10.29)) 10.3.1 20210621
Copyright (C) 2020 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Congratulations! You have successfully installed a toolchain for cross-compiling `sherpa-onnx`.

Build `sherpa-onnx`

Finally, let us build `sherpa-onnx`.

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
export BUILD_SHARED_LIBS=ON
./build-arm-linux-gnueabihf.sh
```

After building, you will get the following binaries:

```
$ ls -lh build-arm-linux-gnueabihf/install/bin/
total 1.2M
-rwxr-xr-x 1 kuangfangjun root 395K Jul  7 16:28 sherpa-onnx
-rwxr-xr-x 1 kuangfangjun root 391K Jul  7 16:28 sherpa-onnx-alsa
-rwxr-xr-x 1 kuangfangjun root 351K Jul  7 16:28 sherpa-onnx-offline
```

That's it!

Hint:

- `sherpa-onnx` is for decoding a single file using a streaming model

- `sherpa-onnx-offline` is for decoding a single file using a non-streaming model
- `sherpa-onnx-alsa` is for real-time speech recognition using a streaming model by reading the microphone with [ALSA](#)

Caution: We recommend that you use `sherpa-onnx-alsa` on embedded systems such as Raspberry pi.

You need to provide a `device_name` when invoking `sherpa-onnx-alsa`. We describe below how to find the device name for your microphone.

Run the following command:

```
arecord -l
```

to list all available microphones for recording. If it complains that `arecord: command not found`, please use `sudo apt-get install alsa-utils` to install it.

If the above command gives the following output:

```
**** List of CAPTURE Hardware Devices ****
card 0: Audio [Axera Audio], device 0: 49ac000.i2s_mst-es8328-hifi-analog
  ↳ es8328-hifi-analog-0 []
    Subdevices: 1/1
    Subdevice #0: subdevice #0
```

In this case, I only have 1 microphone. It is card 0 and that card has only device 0. To select card 0 and device 0 on that card, we need to pass `plughw:0,0` to `sherpa-onnx-alsa`. (Note: It has the format `plughw:card_number,device_index`.)

For instance, you have to use

```
# Note: We use int8 models below.
./bin/sherpa-onnx-alsa \
  ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
  ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-
  ↳ chunk-16-left-64.int8.onnx \
  ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-
  ↳ chunk-16-left-64.int8.onnx \
  ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-
  ↳ chunk-16-left-64.int8.onnx \
  "plughw:0,0"
```

Please change the card number and also the device index on the selected card accordingly in your own situation. Otherwise, you won't be able to record with your microphone.

Please read [Pre-trained models](#) for usages about the generated binaries.

Read below if you want to learn more.

Hint: By default, all external dependencies are statically linked. That means, the generated binaries are self-contained (except that it requires the onnxruntime shared library at runtime).

You can use the following commands to check that and you will find they depend only on system libraries.

```
$ readelf -d build-arm-linux-gnueabihf/install/bin/sherpa-onnx
```

(continues on next page)

(continued from previous page)

```

Dynamic section at offset 0x61ee8 contains 30 entries:
  Tag      Type           Name/Value
  0x00000001 (NEEDED)   Shared library: [libbonnxruntime.so.
↪ 1.14.0]
  0x00000001 (NEEDED)   Shared library: [libstdc++.so.6]
  0x00000001 (NEEDED)   Shared library: [libm.so.6]
  0x00000001 (NEEDED)   Shared library: [libgcc_s.so.1]
  0x00000001 (NEEDED)   Shared library: [libc.so.6]
  0x0000000f (RPATH)    Library rpath: [$ORIGIN:$ORIGIN/../
↪ lib:$ORIGIN/../../../../sherpa_onda/lib]

$ readelf -d build-arm-linux-gnueabihf/install/bin/sherpa-onda-alsa

Dynamic section at offset 0x60ee0 contains 31 entries:
  Tag      Type           Name/Value
  0x00000001 (NEEDED)   Shared library: [libasound.so.2]
  0x00000001 (NEEDED)   Shared library: [libbonnxruntime.so.
↪ 1.14.0]
  0x00000001 (NEEDED)   Shared library: [libstdc++.so.6]
  0x00000001 (NEEDED)   Shared library: [libm.so.6]
  0x00000001 (NEEDED)   Shared library: [libgcc_s.so.1]
  0x00000001 (NEEDED)   Shared library: [libc.so.6]
  0x0000000f (RPATH)    Library rpath: [$ORIGIN]

```

Please create an issue at <https://github.com/k2-fsa/sherpa-onda/issues> if you have any problems.

How to build static libraries and static linked binaries

If you want to build static libraries and static linked binaries, please first download a cross compile toolchain with GCC >= 9.0. The following is an example:

```

mkdir -p $HOME/software
cd $HOME/software
wget -q https://huggingface.co/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-10.
↪ 3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz

# For users from China
# huggingface,
wget -q https://hf-mirror.com/csukuangfj/sherpa-ncnn-toolchains/resolve/main/gcc-arm-10.
↪ 3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz

tar xf gcc-arm-10.3-2021.07-x86_64-arm-none-linux-gnueabihf.tar.xz

```

Next, we need to set the following environment variable:

```
export PATH=$HOME/software/gcc-arm-10.3-2021.07-x86_64-arm-none-linux-gnueabihf/bin:$PATH
```

To check that we have installed the cross-compiling toolchain successfully, please run:

```
arm-none-linux-gnueabihf-gcc --version
```

which should print the following log:

```
arm-none-linux-gnueabihf-gcc (GNU Toolchain for the A-profile Architecture 10.3-2021.07
 ↵(arm-10.29)) 10.3.1 20210621
Copyright (C) 2020 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Now you can build static libraries and static linked binaries with the following commands:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
export BUILD_SHARED_LIBS=OFF
./build-arm-linux-gnueabihf.sh
```

You can use the following commands to check that the generated binaries are indeed static linked:

```
$ cd build-arm-linux-gnueabihf/bin

$ ldd sherpa-onnx-alsa
    not a dynamic executable

$ readelf -d sherpa-onnx-alsa

Dynamic section at offset 0xa68eb4 contains 31 entries:
  Tag          Type           Name/Value
 0x00000001 (NEEDED)      Shared library: [libasound.so.2]
 0x00000001 (NEEDED)      Shared library: [libdl.so.2]
 0x00000001 (NEEDED)      Shared library: [libm.so.6]
 0x00000001 (NEEDED)      Shared library: [libpthread.so.0]
 0x00000001 (NEEDED)      Shared library: [libc.so.6]
 0x00000001 (NEEDED)      Shared library: [ld-linux-armhf.so.3]
 0x0000000f (RPATH)       Library rpath: [$ORIGIN:/star-fj/fangjun/open-
 ↵-source/sherpa-onnx/build-arm-linux-gnueabihf/_deps/espeak_ng-src/lib:/star-fj/fangjun/
 ↵-open-source/sherpa-onnx/build-arm-linux-gnueabihf/_deps/onnxruntime-src/lib:]
 0x0000000c (INIT)        0x13550
```

8.2.6 Embedded Linux (riscv64)

This page describes how to build `sherpa-onnx` for embedded Linux (RISC-V, 64-bit) with cross-compiling on an x64 machine with Ubuntu OS. It also demonstrates how to use `qemu` to run the compiled binaries.

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows/macOS or you don't want to setup your local environment for cross-compiling, please use the above colab notebook.

Note: You can download pre-compiled binaries for `riscv64` from the following URL <https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/riscv64>

Please always download the latest version.

Example command to download the version 1.9.12:

```
# binaries built with shared libraries
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/riscv64/
↳sherpa-onnx-v1.9.12-linux-riscv64-shared.tar.bz2

# For users from China
# huggingface,

# binaries built with shared libraries
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/riscv64/
↳sherpa-onnx-v1.9.12-linux-riscv64-shared.tar.bz2
```

Install toolchain

The first step is to install a toolchain for cross-compiling.

```
mkdir -p $HOME/toolchain

wget -q https://occ-oss-prod.oss-cn-hangzhou.aliyuncs.com/resource//1663142514282/
↳Xuantie-900-gcc-linux-5.10.4-glibc-x86_64-V2.6.1-20220906.tar.gz

tar xf ./Xuantie-900-gcc-linux-5.10.4-glibc-x86_64-V2.6.1-20220906.tar.gz --strip-
↳components 1 -C $HOME/toolchain
```

Next, we need to set the following environment variable:

```
export PATH=$HOME/toolchain/bin:$PATH
```

To check that you have installed the toolchain successfully, please run

```
$ riscv64-unknown-linux-gnu-gcc --version

riscv64-unknown-linux-gnu-gcc (Xuantie-900 linux-5.10.4 glibc gcc Toolchain V2.6.1 B-
↳20220906) 10.2.0
Copyright (C) 2020 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

$ riscv64-unknown-linux-gnu-g++ --version

riscv64-unknown-linux-gnu-g++ (Xuantie-900 linux-5.10.4 glibc gcc Toolchain V2.6.1 B-
↳20220906) 10.2.0
Copyright (C) 2020 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.
```

Build sherpa-onnx

Next, let us build `sherpa-onnx`.

Hint: Currently, only shared libraries are supported. We will support static linking in the future.

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
./build-riscv64-linux-gnu.sh
```

After building, you will get the following files

```
$ ls -lh build-riscv64-linux-gnu/install/bin
$ echo "---"
$ ls -lh build-riscv64-linux-gnu/install/lib

total 292K
-rwxr-xr-x 1 root root 23K Mar 20 09:41 sherpa-onnx
-rwxr-xr-x 1 root root 27K Mar 20 09:41 sherpa-onnx-alsa
-rwxr-xr-x 1 root root 31K Mar 20 09:41 sherpa-onnx-alsa-offline
-rwxr-xr-x 1 root root 40K Mar 20 09:41 sherpa-onnx-alsa-offline-speaker-identification
-rwxr-xr-x 1 root root 23K Mar 20 09:41 sherpa-onnx-keyword-spotter
-rwxr-xr-x 1 root root 27K Mar 20 09:41 sherpa-onnx-keyword-spotter-alsa
-rwxr-xr-x 1 root root 23K Mar 20 09:41 sherpa-onnx-offline
-rwxr-xr-x 1 root root 39K Mar 20 09:41 sherpa-onnx-offline-parallel
-rwxr-xr-x 1 root root 19K Mar 20 09:41 sherpa-onnx-offline-tts
-rwxr-xr-x 1 root root 31K Mar 20 09:41 sherpa-onnx-offline-tts-play-alsa
---
total 30M
-rw-r--r-- 1 root root 256K Mar 20 09:41 libespeak-ng.so
-rw-r--r-- 1 root root 71K Mar 20 09:41 libkaldi-decoder-core.so
-rw-r--r-- 1 root root 67K Mar 20 09:41 libkaldi-native-fbank-core.so
-rw-r--r-- 1 root root 13M Mar 20 09:35 libonnxruntime.so
-rw-r--r-- 1 root root 13M Mar 20 09:35 libonnxruntime.so.1.14.1
lrwxrwxrwx 1 root root 23 Mar 20 09:41 libpiper_phonemize.so -> libpiper_phonemize.
so.1
lrwxrwxrwx 1 root root 27 Mar 20 09:41 libpiper_phonemize.so.1 -> libpiper_phonemize.
so.1.2.0
-rw-r--r-- 1 root root 395K Mar 20 09:41 libpiper_phonemize.so.1.2.0
-rw-r--r-- 1 root root 1.3M Mar 20 09:41 libsherpa-onnx-core.so
lrwxrwxrwx 1 root root 23 Mar 20 09:41 libsherpa-onnx-fst.so -> libsherpa-onnx-fst.
so.6
-rw-r--r-- 1 root root 1.4M Mar 20 09:41 libsherpa-onnx-fst.so.6
-rw-r--r-- 1 root root 752K Mar 20 09:41 libsherpa-onnx-kaldifst-core.so
-rw-r--r-- 1 root root 202K Mar 20 09:41 libucd.so
drwxr-xr-x 2 root root 4.0K Mar 20 09:41 pkgconfig
```

```
$ file build-riscv64-linux-gnu/install/bin/sherpa-onnx

build-riscv64-linux-gnu/install/bin/sherpa-onnx: ELF 64-bit LSB executable, UCB RISC-V, ↵
RVC, double-float ABI, version 1 (GNU/Linux), dynamically linked, interpreter /lib/ld- ↵
linux-riscv64-lp64d.so.1, for GNU/Linux 4.15.0, stripped
```

```
$ readelf -d build-riscv64-linux-gnu/install/bin/sherpa-onnx
```

```
$ find $HOME/toolchain/ -name ld-linux-riscv64-lp64d.so.1

Dynamic section at offset 0x4d40 contains 39 entries:
  Tag          Type           Name/Value
  0x0000000000000001 (NEEDED)   Shared library: [libsherpa-onnx-core.so]
  0x0000000000000001 (NEEDED)   Shared library: [libkaldi-native-fbank-core.
↪so]
  0x0000000000000001 (NEEDED)   Shared library: [libkaldi-decoder-core.so]
  0x0000000000000001 (NEEDED)   Shared library: [libsherpa-onnx-kaldifst-
↪core.so]
  0x0000000000000001 (NEEDED)   Shared library: [libsherpa-onnx-fst.so.6]
  0x0000000000000001 (NEEDED)   Shared library: [libpiper_phonemize.so.1]
  0x0000000000000001 (NEEDED)   Shared library: [libonnxruntime.so.1.14.1]
  0x0000000000000001 (NEEDED)   Shared library: [libespeak-ng.so]
  0x0000000000000001 (NEEDED)   Shared library: [libucd.so]
  0x0000000000000001 (NEEDED)   Shared library: [libstdc++.so.6]
  0x0000000000000001 (NEEDED)   Shared library: [libm.so.6]
  0x0000000000000001 (NEEDED)   Shared library: [libgcc_s.so.1]
  0x0000000000000001 (NEEDED)   Shared library: [libpthread.so.0]
  0x0000000000000001 (NEEDED)   Shared library: [libc.so.6]
  0x000000000000000f (RPATH)    Library rpath: [$ORIGIN:$ORIGIN/../lib:
↪$ORIGIN/../../sherpa_onnx/lib]
  0x0000000000000020 (PREINIT_ARRAY) 0x15d20
  0x0000000000000021 (PREINIT_ARRAYSZ) 8 (bytes)
  0x0000000000000019 (INIT_ARRAY) 0x15d28
  0x000000000000001b (INIT_ARRAYSZ) 16 (bytes)
  0x000000000000001a (FINI_ARRAY) 0x15d38
  0x000000000000001c (FINI_ARRAYSZ) 8 (bytes)
  0x0000000000000004 (HASH) 0x10280
  0x0000000000ffffef5 (GNU_HASH) 0x10418
  0x0000000000000005 (STRTAB) 0x10bd8
  0x0000000000000006 (SYMTAB) 0x105f0
  0x000000000000000a (STRSZ) 3652 (bytes)
  0x000000000000000b (SYMENT) 24 (bytes)
  0x0000000000000015 (DEBUG) 0x0
  0x0000000000000003 (PLTGOT) 0x16000
  0x0000000000000002 (PLTRELSZ) 1056 (bytes)
  0x0000000000000014 (PLTREL) RELA
  0x0000000000000017 (JMPREL) 0x11bb0
  0x0000000000000007 (RELA) 0x11b80
  0x0000000000000008 (RELASZ) 1104 (bytes)
  0x0000000000000009 (RELAENT) 24 (bytes)
  0x0000000000fffffe (VERNEED) 0x11aa0
  0x0000000000fffffe (VERNEEDNUM) 4
  0x0000000000fffff0 (VERSYM) 0x11a1c
  0x0000000000000000 (NULL) 0x0

/root/toolchain/sysroot/lib/ld-linux-riscv64-lp64d.so.1
```

That's it!

Please create an issue at <https://github.com/k2-fsa/sherpa-onnx/issues> if you have any problems.

Read more if you want to run the binaries with qemu.

qemu

Hint: This subsection works only on x64 Linux.

Caution: Please don't use any other methods to install `qemu-riscv64`. Only the method listed in this subsection is known to work.

Please use the following command to download the `qemu-riscv64` binary.

```
mkdir -p $HOME/qemu

mkdir -p /tmp
cd /tmp
wget -q https://files.pythonhosted.org/packages/21/f4/
↳ 733f29c435987e8bb264a6504c7a4ea4c04d0d431b38a818ab63eef082b9/xuantie_qemu-20230825-py3-
↳ none-manylinux1_x86_64.whl

unzip xuantie_qemu-20230825-py3-none-manylinux1_x86_64.whl
cp -v ./qemu/qemu-riscv64 $HOME/qemu

export PATH=$HOME/qemu:$PATH
```

To check that we have installed `qemu-riscv64` successfully, please run:

```
qemu-riscv64 -h
```

which should give the following output:

```
usage: qemu-riscv64 [options] program [arguments...]
Linux CPU emulator (compiled for riscv64 emulation)

Options and associated environment variables:

Argument          Env-variable      Description
-h                print this help
-help
-g port           QEMU_GDB        wait gdb connection to 'port'
-L path           QEMU_LD_PREFIX  set the elf interpreter prefix to 'path'
-s size           QEMU_STACK_SIZE set the stack size to 'size' bytes
-cpu model       QEMU_CPU        select CPU (-cpu help for list)
-E var=value     QEMU_SET_ENV   sets targets environment variable (see below)
-U var           QEMU_UNSET_ENV  unsets targets environment variable (see below)
-@ argv@         QEMU_ARGV@     forces target process argv[0] to be 'argv@'
-r uname         QEMU_UNAME     set qemu uname release string to 'uname'
-B address       QEMU_GUEST_BASE set guest_base address to 'address'
-R size          QEMU_RESERVED_VA reserve 'size' bytes for guest virtual address
-space
-d item[,...]    QEMU_LOG       enable logging of specified items (use '-d help' for
a list of items)
```

(continues on next page)

(continued from previous page)

-dfilter	range[,...]	QEMU_DFILTER	filter logging based on address range
-D	logfile	QEMU_LOG_FILENAME	write logs to ' logfile ' (default stderr)
-p	pagesize	QEMU_PAGESIZE	set the host page size to ' pagesize '
-singlestep		QEMU_SINGLESTEP	run in singlestep mode
-strace		QEMU_STRACE	log system calls
-pctrace		QEMU_PCTRACE	log pctrace
-seed		QEMU_RAND_SEED	Seed for pseudo-random number generator
-trace		QEMU_TRACE	[enable=<pattern>][,events=<file>][,file=<file>]
-csky-extend		CSKY_EXTEND	[tb_trace=<on off>][,jcount_start=<addr>][,jcount_end=<addr>][vdsp=<vdsp>][exit_addr=<addr>][denormal=<on off>]
-CPF		CSKY_PROFILING	
-csky-trace		CSKY_TRACE	[port=<port>][,tb_trace=<on off>][,mem_trace=<on off>][,auto_trace=<on off>][,start=addr][,exit=addr]
-plugin		QEMU_PLUGIN	[file=<file>][,arg=<string>]
-version		QEMU_VERSION	display version information and exit

Defaults:

QEMU_LD_PREFIX = /usr/gnemu/qemu-riscv64

QEMU_STACK_SIZE = 8388608 byte

You can use **-E** **and** **-U** options **or** the QEMU_SET_ENV **and** QEMU_UNSET_ENV environment variables to **set** **and** **unset** environment variables **for** the target process.

It **is** possible to provide several variables by separating them by commas **in** getsubopt(3) style. Additionally it **is** possible to provide the **-E** **and** **-U** options multiple times.

The following lines are equivalent:

```
-E vari=val2 -E var2=val2 -U LD_PRELOAD -U LD_DEBUG
-E vari=val2,var2=val2 -U LD_PRELOAD,LD_DEBUG
QEMU_SET_ENV=vari=val2,var2=val2 QEMU_UNSET_ENV=LD_PRELOAD,LD_DEBUG
```

Note that **if** you provide several changes to a single variable the last change will stay **in** effect.

See <<https://qemu.org/contribute/report-a-bug>> **for** how to report bugs.
More information on the QEMU project at <<https://qemu.org>>.

We describe below how to use qemu-riscv64 to run speech-to-text and text-to-speech.

Run speech-to-text with qemu

We use [csukuangfj/sherpa-onnx-streaming-zipformer-en-20M-2023-02-17 \(English\)](https://github.com/csukuangfj/sherpa-onnx-streaming-zipformer-en-20M-2023-02-17) as the test model.

Note: You can select any model from *Pre-trained models*.

Please use the following command to download the model:

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→ streaming-zipformer-en-20M-2023-02-17.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf sherpa-onnx-streaming-zipformer-en-20M-2023-02-17.tar.bz2  
rm sherpa-onnx-streaming-zipformer-en-20M-2023-02-17.tar.bz2
```

Now you can use the following command to run it with `qemu-riscv64`:

```
cd /path/to/sherpa-onnx  
  
export PATH=$HOME/qemu:$PATH  
  
qemu-riscv64 build-riscv64-linux-gnu/install/bin/sherpa-onnx \  
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \  
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.\  
onnx \  
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.\  
onnx \  
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.\  
onnx \  
../sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
```

It will throw the following error:

```
qemu-riscv64: Could not open '/lib/ld-linux-riscv64-lp64d.so.1': No such file or  
directory
```

Please use the following command instead:

```
cd /path/to/sherpa-onnx  
  
export PATH=$HOME/qemu:$PATH  
export QEMU_LD_PREFIX=$HOME/toolchain/sysroot  
  
qemu-riscv64 build-riscv64-linux-gnu/install/bin/sherpa-onnx \  
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \  
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.\  
onnx \  
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.\  
onnx \  
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.\  
onnx \  
../sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
```

It will throw a second error:

```
build-riscv64-linux-gnu/install/bin/sherpa-onnx: error while loading shared libraries:  
ld-linux-riscv64xthead-lp64d.so.1: cannot open shared object file: No such file or  
directory
```

Please use the following command instead:

```
cd /path/to/sherpa-onnx  
  
export PATH=$HOME/qemu:$PATH  
export QEMU_LD_PREFIX=$HOME/toolchain/sysroot
```

(continues on next page)

(continued from previous page)

```
export LD_LIBRARY_PATH=$HOME/toolchain/sysroot/lib:$LD_LIBRARY_PATH

qemu-riscv64 build-riscv64-linux-gnu/install/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.
˓→onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.
˓→onnx \
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
```

Finally, it prints the following output:

```
/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 build-riscv64-linux-gnu/
˓→install/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
˓→tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-
˓→99-avg-1.onnx --decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-
˓→epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
˓→joiner-epoch-99-avg-1.onnx ./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_
˓→wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
˓→"./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.onnx", u
˓→decoder=".sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
˓→onnx", joiner=".sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-
˓→1.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), wenet_
˓→ctc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4), zipformer2_
˓→ctc=OnlineZipformer2CtcModelConfig(model=""), tokens=".sherpa-onnx-streaming-
˓→zipformer-en-20M-2023-02-17/tokens.txt", num_threads=1, debug=False, provider="cpu", u
˓→model_type=""), lm_config=OnlineLMConfig(model="", scale=0.5), endpoint_
˓→config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓→silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, u
˓→min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
˓→nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
˓→endpoint=True, max_active_paths=4, hotwords_score=1.5, hotwords_file="", decoding_
˓→method="greedy_search", blank_penalty=0)
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
Elapsed seconds: 70, Real time factor (RTF): 11
THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLELS
{ "text": " THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE
˓→BRAFFLELS", "tokens": [ " THE", " YE", " LL", " OW", " LA", " M", " P", " S", " WOULD", " u
˓→LIGHT", " UP", " HE", " RE", " AND", " THERE", " THE", " S", " QUA", " LI", " D", " ", " QUA
˓→", " R", " TER", " OF", " THE", " B", " RA", " FF", " L", " EL", " S" ], "timestamps": [ 2.04,
˓→2.16, 2.28, 2.36, 2.52, 2.64, 2.68, 2.76, 2.92, 3.08, 3.40, 3.60, 3.72, 3.88, 4.12, 4.
˓→48, 4.64, 4.68, 4.84, 4.96, 5.16, 5.20, 5.32, 5.36, 5.60, 5.72, 5.92, 5.96, 6.08, 6.24,
˓→6.36, 6.60 ], "ys_probs": [ -0.454799, -0.521409, -0.345871, -0.001244, -0.240359, -0.
˓→013972, -0.010445, -0.051701, -0.000371, -0.171570, -0.002205, -0.026703, -0.006903, -
˓→0.021168, -0.011662, -0.001059, -0.0005089, -0.000273, -0.575480, -0.024973, -0.159344,
˓→-0.000042, -0.011082, -0.187136, -0.004002, -0.292751, -0.084873, -0.241302, -0.543844,
˓→-0.428164, -0.853198, -0.093776 ], "lm_probs": [ ], "context_scores": [ ], "segment
˓→": 0, "start_time": 0.00, "is_final": false}
```

(continues on next page)

(continued from previous page)

Hint: As you can see, the RTF is 11, indicating that it is very slow to run the model with the qemu simulator. Running on a real RISC-V board should be much faster.

Run text-to-speech with qemu

Please visit <https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models> to download a text-to-speech model. We use the following model `vits-piper-en_US-amy-low.tar.bz2`:

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-piper-en_US-
amy-low.tar.bz2
tar xf vits-piper-en_US-amy-low.tar.bz2
rm vits-piper-en_US-amy-low.tar.bz2
```

After downloading the model, we can use the following command to run it:

```
cd /path/to/sherpa-onnx

export PATH=$HOME/qemu:$PATH
export QEMU_LD_PREFIX=$HOME/toolchain/sysroot
export LD_LIBRARY_PATH=$HOME/toolchain/sysroot/lib:$LD_LIBRARY_PATH

qemu-riscv64 build-riscv64-linux-gnu/install/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-piper-en_US-amy-low/en_US-amy-low.onnx \
--vits-tokens=./vits-piper-en_US-amy-low/tokens.txt \
--vits-data-dir=./vits-piper-en_US-amy-low/espeak-ng-data \
--output-filename=./a-test.wav \
"Friends fell out often because life was changing so fast. The easiest thing in the_
world was to lose touch with someone."
```

The log of the above command is given below:

```
/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 build-riscv64-linux-gnu/
install/bin/sherpa-onnx-offline-tts --vits-model=./vits-piper-en_US-amy-low/en_US-amy-
low.onnx --vits-tokens=./vits-piper-en_US-amy-low/tokens.txt --vits-data-dir=./vits-
piper-en_US-amy-low/espeak-ng-data --output-filename=./a-test.wav 'Friends fell out_
often because life was changing so fast. The easiest thing in the world was to lose_
touch with someone.'
```

Elapsed seconds: 270.745 s
Audio duration: 7.904 s
Real-time factor (RTF): 270.745/7.904 = 34.254
The text is: Friends fell out often because life was changing so fast. The easiest thing_
in the world was to lose touch with someone.. Speaker ID: 0
Saved to ./a-test.wav successfully!

If you want to build an Android app, please refer to [Android](#). If you want to build an iOS app, please refer to [iOS](#).

8.3 Frequently Asked Question (FAQs)

This page contains frequently asked questions for `sherpa-onnx`.

8.3.1

,

,

8.3.2 Cannot open shared library `libasound_module_conf_pulse.so`

The detailed errors are given below:

```
Cannot open shared library libasound_module_conf_pulse.so
(/usr/lib64/alsa-lib/libasound_module_conf_pulse.so: cannot open shared object file: No such file or directory)
ALSA lib pcm.c:2722:(snd_pcm_open_noupdate) Unknown PCM
```

If you use Linux and get the above error when trying to use the microphone, please do the following:

1. Locate where is the file `libasound_module_conf_pulse.so` on your system

```
find / -name libasound_module_conf_pulse.so 2>/dev/null
```

2. If the above search command prints:

```
/usr/lib/x86_64-linux-gnu/alsa-lib/libasound_module_conf_pulse.so
/usr/lib/i386-linux-gnu/alsa-lib/libasound_module_conf_pulse.so
```

3. Please run:

```
sudo mkdir -p /usr/lib64/alsa-lib
sudo ln -s /usr/lib/x86_64-linux-gnu/alsa-lib/libasound_module_conf_pulse.so /usr/lib64/alsa-lib
```

4. Now your issue should be fixed.

8.3.3 TTS

Please see [How to enable UTF-8 on Windows](#). You need to use UTF-8 encoding for your system.

8.3.4 ./gitcompile: line 89: libtoolize: command not found

If you are using Linux and get the following error:

```
./gitcompile: line 89: libtoolize: command not found
```

Please run:

```
sudo apt-get install libtool
```

8.3.5 OSError: PortAudio library not found

If you have the following error on Linux (Ubuntu),

```
Traceback (most recent call last):
  File "/mnt/sdb/shared/sherpa-onnx./python-api-examples/vad-microphone.py", line 8, in 
    <module>
    import sounddevice as sd
  File "/mnt/sdb/shared/py311/lib/python3.11/site-packages/sounddevice.py", line 71, in 
    <module>
    raise OSError('PortAudio library not found')
OSError: PortAudio library not found
```

Then please run:

```
sudo apt-get install libportaudio2
```

and then re-try.

8.3.6 imports github.com/k2-fsa/sherpa-onnx-go-linux: build constraints exclude all Go files

If you have the following output when running go build:

```
[root@VM-0-3-centos non-streaming-decode-files]# go build
package non-streaming-decode-files
imports github.com/k2-fsa/sherpa-onnx-go/sherpa_onnx
imports github.com/k2-fsa/sherpa-onnx-go-linux: build constraints exclude all Go files
  in /root/go/pkg/mod/github.com/k2-fsa/sherpa-onnx-go-linux@v1.9.21
```

Please first run:

```
go env -w CGO_ENABLED=1
```

And then re-run go build.

8.4 Python

In this section, we describe how to install the Python package `sherpa-onnx`.

8.4.1 Install the Python Package

You can select one of the following methods to install the Python package.

Method 1 (From pre-compiled wheels, CPU only)

Hint: This method supports the following platforms:

- Linux (x64, aarch64, armv7l),
- macOS (x64, arm64)
- Windows (x64, x86)

Note that this method installs a CPU-only version of `sherpa-onnx`.

```
pip install sherpa-onnx
```

To check you have installed `sherpa-onnx` successfully, please run

```
python3 -c "import sherpa_onnx; print(sherpa_onnx.__file__)"  
which sherpa-onnx  
sherpa-onnx --help  
ls -lh $(dirname $(which sherpa-onnx))/sherpa-onnx*
```

Hint: You can find previous releases at <https://k2-fsa.github.io/sherpa/onnx/cpu.html>

For Chinese users and users who have no access to huggingface, please visit <https://k2-fsa.github.io/sherpa/onnx/cpu-cn.html>.

You can use:

```
pip install sherpa-onnx -f https://k2-fsa.github.io/sherpa/onnx/cpu.html
```

or:

```
# For Chinese users  
pip install sherpa-onnx -f https://k2-fsa.github.io/sherpa/onnx/cpu-cn.html
```

Method 2 (From pre-compiled wheels, CPU + CUDA)

Note: This method installs a version of `sherpa-onnx` supporting both CUDA and CPU. You need to pass the argument `provider=cuda` to use NVIDIA GPU, which always uses GPU 0. Otherwise, it uses CPU by default.

Please use the environment variable `CUDA_VISIBLE_DEVICES` to control which GPU is mapped to GPU 0.

By default, `provider` is set to `cpu`.

Remember to follow <https://onnxruntime.ai/docs/execution-providers/CUDA-ExecutionProvider.html#requirements> to install CUDA 11.8.

If you have issues about installing CUDA 11.8, please have a look at <https://k2-fsa.github.io/k2/installation/cuda-cudnn.html#cuda-11-8>.

Note that you don't need to have `sudo` permission to install CUDA 11.8

This approach supports only Linux x64 and Windows x64.

Please use the following command to install CUDA-enabled `sherpa-onnx`:

```
# We use 1.10.16 here for demonstration.  
#  
# Please visit https://k2-fsa.github.io/sherpa/onnx/cuda.html  
# to find available versions  
  
pip install sherpa-onnx==1.10.16+cuda -f https://k2-fsa.github.io/sherpa/onnx/cuda.html  
  
# For Chinese users, please use  
# pip install sherpa-onnx==1.10.16+cuda -f https://k2-fsa.github.io/sherpa/onnx/cuda-cn.  
#html
```

The installation logs are given below:

```
Looking in links: https://k2-fsa.github.io/sherpa/onnx/cuda.html  
Collecting sherpa-onnx==1.10.16+cuda  
  Downloading https://huggingface.co/csukuangfj/sherpa-onnx-wheels/resolve/main/cuda/1.  
10.16/sherpa_onnx-1.10.16%2Bcuda-cp310-cp310-linux_x86_64.whl (183.3 MB)  
    183.3/183.3 MB 4.4 MB/s eta 0:00:00  
Installing collected packages: sherpa-onnx  
Successfully installed sherpa-onnx-1.10.16+cuda
```

To check that you have installed `sherpa-onnx` successfully, please run:

```
python3 -c "import sherpa_onnx; print(sherpa_onnx.__version__)"
```

which should print something like below:

```
1.10.16+cuda
```

Method 3 (From source)

CPU

Nvidia GPU (CUDA)

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
python3 setup.py install
```

```
git clone https://github.com/k2-fsa/sherpa-onnx
export SHERPA_ONNX_CMAKE_ARGS="-DSHERPA_ONNX_ENABLE_GPU=ON"
cd sherpa-onnx
python3 setup.py install
```

Method 4 (For developers)

CPU

Nvidia GPU (CUDA)

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build

cmake \
-DSHERPA_ONNX_ENABLE_PYTHON=ON \
-DBUILD_SHARED_LIBS=ON \
-DSHERPA_ONNX_ENABLE_CHECK=OFF \
-DSHERPA_ONNX_ENABLE_PORTAUDIO=OFF \
-DSHERPA_ONNX_ENABLE_C_API=OFF \
-DSHERPA_ONNX_ENABLE_WEBSOCKET=OFF \
..

make -j
export PYTHONPATH=$PWD/../../sherpa-onnx/python/:$PWD/lib:$PYTHONPATH
```

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build
cd build

cmake \
-DSHERPA_ONNX_ENABLE_PYTHON=ON \
-DBUILD_SHARED_LIBS=ON \
-DSHERPA_ONNX_ENABLE_CHECK=OFF \
-DSHERPA_ONNX_ENABLE_PORTAUDIO=OFF \
-DSHERPA_ONNX_ENABLE_C_API=OFF \
-DSHERPA_ONNX_ENABLE_WEBSOCKET=OFF \
-DSHERPA_ONNX_ENABLE_GPU=ON \
..
```

(continues on next page)

(continued from previous page)

```
make -j
export PYTHONPATH=$PWD/../../sherpa-onnx/python/:$PWD/lib:$PYTHONPATH
```

Hint: You need to install CUDA toolkit. Otherwise, you would get errors at runtime.

You can refer to <https://k2-fsa.github.io/k2/installation/cuda-cudnn.html> to install CUDA toolkit.

Check your installation

To check that `sherpa-onnx` has been successfully installed, please use:

```
python3 -c "import sherpa_onnx; print(sherpa_onnx.__file__)"
```

It should print some output like below:

```
/Users/fangjun/py38/lib/python3.8/site-packages/sherpa_onnx/__init__.py
```

Please refer to:

<https://github.com/k2-fsa/sherpa-onnx/tree/master/python-api-examples>

for usages.

Please refer to *Pre-trained models* for a list of pre-trained models.

8.4.2 Decode files

In this section, we demonstrate how to use the Python API of `sherpa-onnx` to decode files.

Hint: We only support WAVE files of single channel and each sample should have 16-bit, while the sample rate of the file can be arbitrary and it does not need to be 16 kHz

Streaming zipformer

We use `csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20` (*Bilingual, Chinese + English*) as an example below.

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/online-decode-files.py \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/0.wav \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.wav
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/2.wav \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/3.wav \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/8k.wav
```

Hint: `online-decode-files.py` is from <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/online-decode-files.py>

Note: You can replace `encoder-epoch-99-avg-1.onnx` with `encoder-epoch-99-avg-1.int8.onnx` to use int8 models for decoding.

The output is given below:

```
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Started!
Done!
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/0.wav
MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW
-----
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.wav
ALWAYS ALWAYS
-----
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/2.wav
FREQUENTLY
-----
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/3.wav
ES
-----
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/8k.wav
IN TIME
-----

num_threads: 1
decoding_method: greedy_search
Wave duration: 17.640 s
Elapsed time: 3.907 s
Real time factor (RTF): 3.907/17.640 = 0.221
```

Non-streaming zipformer

We use [csukuangfj/sherpa-onnx-zipformer-en-2023-04-01 \(English\)](https://github.com/csukuangfj/sherpa-onnx-zipformer-en-2023-04-01) as an example below.

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-decode-files.py \
--tokens=./sherpa-onnx-zipformer-en-2023-04-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
```

Hint: offline-decode-files.py is from <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/offline-decode-files.py>

Note: You can replace encoder-epoch-99-avg-1.onnx with encoder-epoch-99-avg-1.int8.onnx to use int8 models for decoding.

The output is given below:

```
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Started!
Done!
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
↳ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY_
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONoured BOSOM TO CONNECT HER PARENT FOR EVER_
↳ WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num_threads: 1
decoding_method: greedy_search
Wave duration: 4.825 s
Elapsed time: 2.567 s
Real time factor (RTF): 2.567/4.825 = 0.532
```

Non-streaming paraformer

We use [csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28](https://github.com/csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28) (*Chinese + English*) as an example below.

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-decode-files.py \
--tokens=./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2023-03-28/model.onnx \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

Note: You can replace `model.onnx` with `model.int8.onnx` to use `int8` models for decoding.

The output is given below:

```
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Started!
Done!
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav

-----
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav

-----
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav

-----
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav

-----
num_threads: 1
decoding_method: greedy_search
Wave duration: 4.204 s
Elapsed time: 1.663 s
Real time factor (RTF): 1.663/4.204 = 0.396
```

8.4.3 Real-time speech recognition from a microphone

In this section, we demonstrate how to use the Python API of `sherpa-onnx` for real-time speech recognition with a microphone.

With endpoint detection

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/speech-recognition-from-microphone-with-endpoint-detection.py \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-avg-1.onnx
```

Hint: `speech-recognition-from-microphone-with-endpoint-detection.py` is from <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/speech-recognition-from-microphone-with-endpoint-detection.py>

In the above demo, the model files are from <https://github.com/csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20> (*Bilingual, Chinese + English*).

Without endpoint detection

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/speech-recognition-from-microphone.py \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-avg-1.onnx
```

Hint: `speech-recognition-from-microphone.py` is from <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/speech-recognition-from-microphone.py>

In the above demo, the model files are from <https://github.com/csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20> (*Bilingual, Chinese + English*).

8.4.4 Speech recognition from URLs

sherpa-onnx also supports decoding from URLs.

Hint: Only streaming models are currently supported. Please modify the code for non-streaming models on need.

All types of URLs supported by `ffmpeg` are supported.

The following table lists some example URLs.

Type	Example
RTMP	<code>rtmp://localhost/live/livestream</code>
OPUS file	<code>https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition/resolve/main/test_wavs/wenetspeech/DEV_T0000000000.opus</code>
WAVE file	<code>https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition/resolve/main/test_wavs/aishell2/ID0012W0030.wav</code>
Local WAVE file	<code>file:///Users/fangjun/open-source/sherpa-onnx/a.wav</code>

Before you continue, please install `ffmpeg` first.

For instance, you can use `sudo apt-get install ffmpeg` for Ubuntu and `brew install ffmpeg` for macOS.

We use the model `csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20` (*Bilingual, Chinese + English*) for demonstration in the following examples.

Decode a URL

This example shows you how to decode a URL pointing to a file.

Hint: The file does not need to be a WAVE file. It can be a file of any format supported by `ffmpeg`.

```
python3 ./python-api-examples/speech-recognition-from-url.py \
--encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--url https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition/resolve/main/
˓→test_wavs/librispeech/1089-134686-0001.wav
```

RTMP

In this example, we use `ffmpeg` to capture a microphone and push the audio stream to a server using RTMP, and then we start `sherpa-onnx` to pull the audio stream from the server for recognition.

Install the server

We will use `srs` as the server. Let us first install `srs` from source:

```
git clone -b develop https://github.com/ossrs/srs.git
cd srs/trunk
./configure
make

# Check that we have compiled srs successfully
./objs/srs --help

# Note: ./objs/srs is statically linked and depends only on system libraries.
```

Start the server

```
ulimit -HSn 10000

# switch to the directory srs/trunk
./objs/srs -c conf/srs.conf
```

The above command gives the following output:

```
srs(51047,0x7ff8451198c0) malloc: nano zone abandoned due to inability to preallocate
↳ reserved vm space.
Asan: Please setup the env MallocNanoZone=0 to disable the warning, see https://
↳ stackoverflow.com/a/70209891/17679565
[2023-07-05 12:19:23.017][INFO][51047][78gw8v44] XCORE-SRS/6.0.55(Bee)
[2023-07-05 12:19:23.021][INFO][51047][78gw8v44] config parse complete
[2023-07-05 12:19:23.021][INFO][51047][78gw8v44] you can check log by: tail -n 30 -f ./
↳ objs/srs.log
[2023-07-05 12:19:23.021][INFO][51047][78gw8v44] please check SRS by: ./etc/init.d/srs
↳ status
```

To check the status of `srs`, use

```
./etc/init.d/srs status
```

which gives the following output:

```
SRS(pid 51548) is running. [ OK ]
```

Hint: If you fail to start the `srs` server, please check the log file `./objs/srs.log` for a fix.

Start ffmpeg to push audio stream

First, let us list available recording devices on the current computer with the following command:

```
ffmpeg -hide_banner -f avfoundation -list_devices true -i ""
```

It gives the following output on my computer:

```
[AVFoundation indev @ 0x7f9f41904840] AVFoundation video devices:  
[AVFoundation indev @ 0x7f9f41904840] [0] FaceTime HD Camera (Built-in)  
[AVFoundation indev @ 0x7f9f41904840] [1] Capture screen 0  
[AVFoundation indev @ 0x7f9f41904840] AVFoundation audio devices:  
[AVFoundation indev @ 0x7f9f41904840] [0] Background Music  
[AVFoundation indev @ 0x7f9f41904840] [1] MacBook Pro Microphone  
[AVFoundation indev @ 0x7f9f41904840] [2] Background Music (UI Sounds)  
[AVFoundation indev @ 0x7f9f41904840] [3] WeMeet Audio Device  
: Input/output error
```

We will use the device [1] MacBook Pro Microphone. Note that its index is 1, so we will use `-i ":1"` in the following command to start recording and push the recorded audio stream to the server under the address `rtmp://localhost/live/livestream`.

Hint: The default TCP port for RTMP is 1935.

```
ffmpeg -hide_banner -f avfoundation -i ":1" -acodec aac -ab 64k -ar 16000 -ac 1 -f flv  
-r rtmp://localhost/live/livestream
```

The above command gives the following output:

```
Input #0, avfoundation, from ':1':  
  Duration: N/A, start: 830938.803938, bitrate: 1536 kb/s  
    Stream #0:0: Audio: pcm_f32le, 48000 Hz, mono, flt, 1536 kb/s  
Stream mapping:  
  Stream #0:0 -> #0:0 (pcm_f32le (native) -> aac (native))  
Press [q] to stop, [?] for help  
Output #0, flv, to 'rtmp://localhost/live/livestream':  
  Metadata:  
    encoder : Lavf60.3.100  
  Stream #0:0: Audio: aac (LC) ([10][0][0][0] / 0x000A), 16000 Hz, mono, fltp, 64 kb/s  
  Metadata:  
    encoder : Lavc60.3.100 aac  
size= 64kB time=00:00:08.39 bitrate= 62.3kbit/s speed=0.977x
```

Start sherpa-onnx to pull audio stream

Now we can start `sherpa-onnx` to pull audio stream from `rtmp://localhost/live/livestream` for speech recognition.

```
python3 ./python-api-examples/speech-recognition-from-url.py \
--encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--url rtmp://localhost/live/livestream
```

You should see the recognition result printed to the console as you speak.

Hint: You can replace `localhost` with your server IP and start `sherpa-onnx` on many computers at the same time to pull audio stream from the address `rtmp://your_server_ip/live/livestream`.

8.4.5 Streaming WebSocket Server

This section describes how to use the Python streaming WebSocket server of `sherpa-onnx` for speech recognition.

Hint: The server supports multiple clients connecting at the same time.

The code for the streaming server can be found at

https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/streaming_server.py

Start the server

Hint:

If you don't use a X.509 certificate, due to security reasons imposed by the browser, you are only allowed to use the domain `localhost` to access the server if you want to access the microphone in the browser. That is, you can only use

`http://localhost:port`

to access the server. You cannot use `http://0.0.0.0:port`, or `http://127.0.0.1:port`, or `http://public_ip:port`.

You can use the following command to generate a self-signed certificate:

```
cd python-api-examples/web
./generate-certificate.py
```

The above commands will generate 3 files. You only need to use the file `cert.pem`. When starting the server, you pass the following argument:

```
--certificate=./python-api-examples/web/cert.pem
```

Please refer to *Pre-trained models* to download a streaming model before you continue.

We will use [csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26 \(English\)](https://github.com/csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26) as an example.

First, let us download it:

```
cd /path/to/sherpa-onnx/
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
      -streaming-zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
cd /path/to/sherpa-onnx/
python3 ./python-api-examples/streaming_server.py \
    --encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
    ↵16-left-128.onnx \
    --decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
    ↵16-left-128.onnx \
    --joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
    ↵16-left-128.onnx \
    --tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
    --port 6006
```

It will print the following logs:

```
2023-08-11 16:29:51,522 INFO [streaming_server.py:678] {'encoder': './sherpa-onnx-
      -streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.onnx',
      'decoder': './sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-
      -chunk-16-left-128.onnx', 'joiner': './sherpa-onnx-streaming-zipformer-en-2023-06-26/
      -joiner-epoch-99-avg-1-chunk-16-left-128.onnx', 'tokens': './sherpa-onnx-streaming-
      -zipformer-en-2023-06-26/tokens.txt', 'sample_rate': 16000, 'feat_dim': 80, 'provider':
      'cpu', 'decoding_method': 'greedy_search', 'num_active_paths': 4, 'use_endpoint': 1,
      'rule1_min_trailing_silence': 2.4, 'rule2_min_trailing_silence': 1.2, 'rule3_min_
      -utterance_length': 20, 'port': 6006, 'nn_pool_size': 1, 'max_batch_size': 50, 'max_
      -wait_ms': 10, 'max_message_size': 1048576, 'max_queue_size': 32, 'max_active_
      -connections': 500, 'num_threads': 2, 'certificate': None, 'doc_root': './python-api-
      -examples/web'}
2023-08-11 16:29:57,476 INFO [streaming_server.py:520] No certificate provided
2023-08-11 16:29:57,480 INFO [server.py:707] server listening on 0.0.0.0:6006
2023-08-11 16:29:57,480 INFO [server.py:707] server listening on [:]:6006
2023-08-11 16:29:57,480 INFO [streaming_server.py:546] Please visit one of the following
      ↵addresses:
      ↵http://localhost:6006

Since you are not providing a certificate, you cannot use your microphone from within
      ↵the browser using public IP addresses. Only localhost can be used. You also cannot use
      ↵0.0.0.0 or 127.0.0.1
```

We can use the following two methods to interact with the server:

- Use Python API
- Use a browser by accessing <http://localhost:6006>

We describe each method below in details.

Use Python API

We provide two Python example files:

Description	URL
Send a file for decoding	https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/online-websocket-client-decode-file.py
Send audio samples from a microphone for decoding	https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/speech-recognition-from-microphone.py

Send a file for decoding

Hint: The example file supports only *.wav files with a single channel and the each sample should be of type `int16_t`. The sample rate does not need to be 16000 Hz, e.g., it can be 48000 Hz, 8000 Hz or some other value.

We use the following command to send a file for decoding:

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/online-websocket-client-decode-file.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
```

It should give the following output:

```
2023-08-11 16:37:03,877 INFO [online-websocket-client-decode-file.py:133] Sending ./
↳sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
2023-08-11 16:37:03,931 INFO [online-websocket-client-decode-file.py:115] {"text": "", 
↳"segment": 0}
2023-08-11 16:37:04,012 INFO [online-websocket-client-decode-file.py:115] {"text": "", 
↳"segment": 0}
2023-08-11 16:37:04,128 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER
↳", "segment": 0}
2023-08-11 16:37:04,170 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY", "segment": 0}
2023-08-11 16:37:04,228 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY", "segment": 0}
2023-08-11 16:37:04,331 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY NIGHTFA", "segment": 0}
2023-08-11 16:37:04,373 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY NIGHTFALL THE", "segment": 0}
2023-08-11 16:37:04,433 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY NIGHTFALL THE YELLOW LA", "segment": 0}
2023-08-11 16:37:04,535 INFO [online-websocket-client-decode-file.py:115] {"text": 
↳"AFTER EARLY NIGHTFALL THE YELLOW LAMPS", "segment": 0} (continues on next page)
```

(continued from previous page)

```

2023-08-11 16:37:04,576 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT", "segment": 0}
2023-08-11 16:37:04,645 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP", "segment": 0}
2023-08-11 16:37:04,685 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE", "segment": 0}
2023-08-11 16:37:04,755 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE", "segment": 0}
2023-08-11 16:37:04,847 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE", "segment": 0}
2023-08-11 16:37:04,887 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUA", "segment": 0}
2023-08-11 16:37:04,958 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID", "segment": 0}
2023-08-11 16:37:05,057 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}
2023-08-11 16:37:05,095 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}
2023-08-11 16:37:05,164 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}
2023-08-11 16:37:05,268 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}
2023-08-11 16:37:05,369 INFO [online-websocket-client-decode-file.py:115] {"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}
2023-08-11 16:37:05,370 INFO [online-websocket-client-decode-file.py:154]
Final result is:
{"text": "AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUAR", "segment": 0}

```

Send audio samples from a microphone for decoding

We use the following command to run the script:

```

cd /path/to/sherpa-onnx

python3 ./python-api-examples/online-websocket-client-microphone.py \
--server-addr localhost \
--server-port 6006

```

It should give you the following output:

```

{'server_addr': 'localhost', 'server_port': 6006}
Started! Please Speak
0 Background Music, Core Audio (2 in, 2 out)

```

(continues on next page)

(continued from previous page)

```
1 Background Music (UI Sounds), Core Audio (2 in, 2 out)
> 2 MacBook Pro Microphone, Core Audio (1 in, 0 out)
< 3 MacBook Pro Speakers, Core Audio (0 in, 2 out)
  4 WeMeet Audio Device, Core Audio (2 in, 2 out)
Use default device: MacBook Pro Microphone
```

```
Started! Please speak
```

If you speak, you will see the recognition result returned by the server.

Use a browser

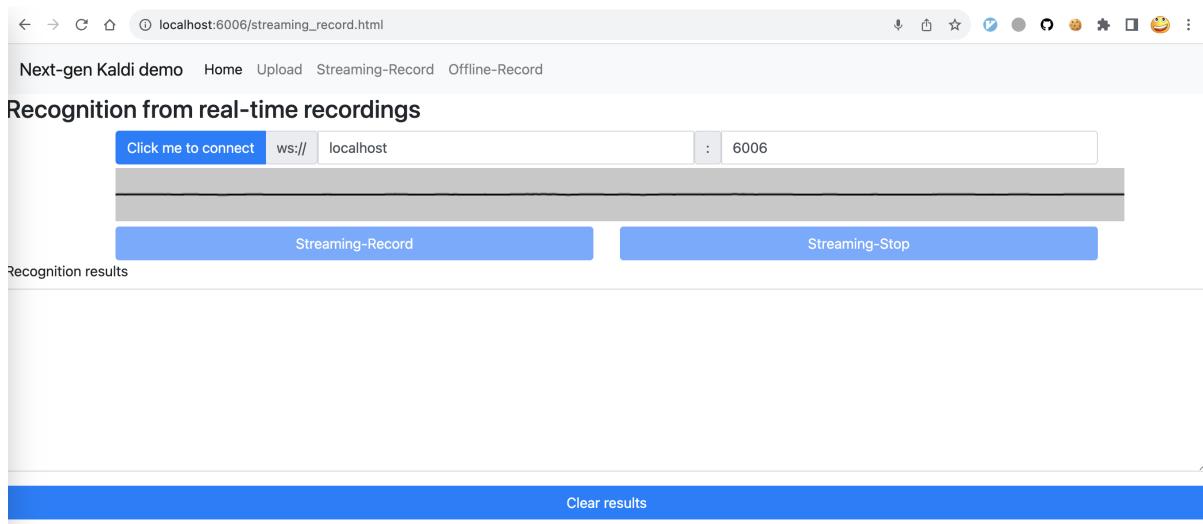
Start your browser and visit the following address:

<http://localhost:6006/>

You should see a page like below:

The screenshot shows a web browser window with the URL `localhost:6006` in the address bar. The page content includes a navigation bar with links for Home, Upload, Streaming-Record, and Offline-Record. Below the navigation bar, there are three main sections: **Upload**, **Streaming_Record**, and **Offline_Record**. Each section has a brief description and a link to the corresponding GitHub repository: <https://github.com/k2-fsa/sherpa-onnx>.

Click **Streaming-Record** and you will see the following page:



Click the button `Click me to connect` to connect to the server and then you can click the `Streaming-Record` button to start recording. You should see the decoded results as you speak.

colab

We provide a colab notebook for you to try the Python streaming websocket server example of `sherpa-onnx`.

8.4.6 Non-Streaming WebSocket Server

This section describes how to use the Python non-streaming WebSocket server of `sherpa-onnx` for speech recognition.

Hint: The server supports multiple clients connecting at the same time.

The code for the non-streaming server can be found at

https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/non_streaming_server.py

Please refer to *Pre-trained models* to download a non-streaming model before you continue.

We use the following types of models for demonstration.

Description	URL
Non-streaming transducer	csukuangfj/sherpa-onnx-zipformer-en-2023-06-26 (English)
Non-streaming paraformer	csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28 (Chinese + English)
Non-streaming CTC model from NeMo	stt_en_conformer_ctc_medium
Non-streaming Whisper tiny.en	tiny.en

Non-streaming transducer

Start the server

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
      -zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-zipformer-en-2023-06-26.tar.bz2

python3 ./python-api-examples/non_streaming_server.py \
    --encoder ./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
    --decoder ./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
    --joiner ./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
    --tokens ./sherpa-onnx-zipformer-en-2023-06-26/tokens.txt \
    --port 6006
```

Start the client

Decode multiple files in parallel

```
python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py \
    --server-addr localhost \
    --server-port 6006 \
    ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav \
    ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav \
    ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:19:26,000 INFO [offline-websocket-client-decode-files-paralell.py:139] {
    ↵ 'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-
      -zipformer-en-2023-06-26/test_wavs/0.wav', './sherpa-onnx-zipformer-en-2023-06-26/test_
      -wavs/1.wav', './sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav']}
2023-08-11 18:19:26,034 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
2023-08-11 18:19:26,058 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav
2023-08-11 18:19:26,205 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
2023-08-11 18:19:26,262 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵ sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
2023-08-11 18:19:26,609 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵ Sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY ↵
    ↵ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH ↵
    ↵ THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
2023-08-11 18:19:26,773 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵ sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID ↵
    ↵ QUARTER OF THE BROTHELS
```

(continues on next page)

(continued from previous page)

Decode multiple files sequentially

```
python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:20:36,677 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID ↵
↳ QUARTER OF THE BROTHELS
2023-08-11 18:20:36,861 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY ↵
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH ↵
↳ THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
2023-08-11 18:20:37,375 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTION HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
```

Non-streaming paraformer

Start the server

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↳ paraformer-zh-2023-03-28.tar.bz2
tar xvf sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2
rm sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2

python3 ./python-api-examples/non_streaming_server.py \
--paraformer ./sherpa-onnx-paraformer-zh-2023-03-28/model.int8.onnx \
--tokens ./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
--port 6006
```

Start the client

Decode multiple files in parallel

```
python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:22:54,189 INFO [offline-websocket-client-decode-files-paralell.py:113] \
↳ Sending ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav
2023-08-11 18:22:54,233 INFO [offline-websocket-client-decode-files-paralell.py:113] \
↳ Sending ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav
2023-08-11 18:22:54,275 INFO [offline-websocket-client-decode-files-paralell.py:113] \
↳ Sending ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
2023-08-11 18:22:54,295 INFO [offline-websocket-client-decode-files-paralell.py:113] \
↳ Sending ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav
2023-08-11 18:22:54,380 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳ sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav

2023-08-11 18:22:54,673 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳ sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav

2023-08-11 18:22:54,673 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳ sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav

2023-08-11 18:22:54,674 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳ sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav
```

Decode multiple files sequentially

```
python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:24:32,678 INFO [offline-websocket-client-decode-files-sequential.py:141] { \
↳ 'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx- \
↳ paraformer-zh-2023-03-28/test_wavs/0.wav', './sherpa-onnx-paraformer-zh-2023-03-28/ \
↳ test_wavs/1.wav', './sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav', './sherpa- \
↳ onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav']}
2023-08-11 18:24:32,709 INFO [offline-websocket-client-decode-files-sequential.py:114] \
↳ Sending ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
2023-08-11 18:24:32,883 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/1.wav

2023-08-11 18:24:33,042 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/2.wav

2023-08-11 18:24:33,175 INFO [offline-websocket-client-decode-files-sequential.py:114] ↵
↳ Sending ./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/8k.wav
```

Non-streaming CTC model from NeMo

Start the server

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-
↳ ctc-en-conformer-medium.tar.bz2
tar xvf sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
rm sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2

python3 ./python-api-examples/non_streaming_server.py \
--nemo-ctc ./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx \
--tokens ./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt \
--port 6006
```

Start the client

Decode multiple files in parallel

```
python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:31:32,432 INFO [offline-websocket-client-decode-files-paralell.py:139] {
↳ 'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-nemo-
↳ ctc-en-conformer-medium/test_wavs/0.wav', './sherpa-onnx-nemo-ctc-en-conformer-medium/
↳ test_wavs/1.wav', './sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav']}
2023-08-11 18:31:32,462 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
↳ Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
2023-08-11 18:31:32,513 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
↳ Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
2023-08-11 18:31:32,533 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
↳ Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav
```

(continues on next page)

(continued from previous page)

```
2023-08-11 18:31:32,670 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
↳quarter of the brothels
2023-08-11 18:31:32,741 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
yet these thoughts affected hester pryne less with hope than apprehension
2023-08-11 18:31:33,117 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
↳sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav
god as a direct consequence of the sin which man thus punished had given her a lovely_
↳child whose place was on that same dishonored bosom to connect her parent for ever_
↳with the race and descent of mortals and to be finally a blessed soul in heaven
```

Decode multiple files sequentially

```
python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:33:14,520 INFO [offline-websocket-client-decode-files-sequential.py:141] {
↳'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-nemo-
↳ctc-en-conformer-medium/test_wavs/0.wav', './sherpa-onnx-nemo-ctc-en-conformer-medium/
↳test_wavs/1.wav', './sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav']}
2023-08-11 18:33:14,547 INFO [offline-websocket-client-decode-files-sequential.py:114]_
↳Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
↳quarter of the brothels
2023-08-11 18:33:14,716 INFO [offline-websocket-client-decode-files-sequential.py:114]_
↳Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav
god as a direct consequence of the sin which man thus punished had given her a lovely_
↳child whose place was on that same dishonored bosom to connect her parent for ever_
↳with the race and descent of mortals and to be finally a blessed soul in heaven
2023-08-11 18:33:15,218 INFO [offline-websocket-client-decode-files-sequential.py:114]_
↳Sending ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
yet these thoughts affected hester pryne less with hope than apprehension
```

Non-streaming Whisper tiny.en

Start the server

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↳whisper-tiny.en.tar.bz2
tar xvf sherpa-onnx-whisper-tiny.en.tar.bz2
rm sherpa-onnx-whisper-tiny.en.tar.bz2
```

(continues on next page)

(continued from previous page)

```
python3 ./python-api-examples/non_streaming_server.py \
--whisper-encoder=./sherpa-onnx-whisper-tiny.en/tiny.en-encoder.onnx \
--whisper-decoder=./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.onnx \
--tokens=./sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt \
--port 6006
```

Start the client

Decode multiple files in parallel

```
python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:35:28,866 INFO [offline-websocket-client-decode-files-paralell.py:139] {
    ↵'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-
    ↵whisper-tiny.en/test_wavs/0.wav', './sherpa-onnx-whisper-tiny.en/test_wavs/1.wav', './
    ↵sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav']}
2023-08-11 18:35:28,894 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav
2023-08-11 18:35:28,947 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav
2023-08-11 18:35:29,082 INFO [offline-websocket-client-decode-files-paralell.py:113] ↵
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
2023-08-11 18:35:29,754 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵sherpa-onnx-whisper-tiny.en/test_wavs/0.wav
    After early nightfall, the yellow lamps would light up here and there, the squalid,
    ↵quarter of the brothels.
2023-08-11 18:35:30,276 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
    Yet these thoughts affected Hester Prin less with hope than apprehension.
2023-08-11 18:35:31,592 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
    ↵sherpa-onnx-whisper-tiny.en/test_wavs/1.wav
    God, as a direct consequence of the sin which man thus punished, had given her a lovely,
    ↵child, whose place was on that same dishonored bosom to connect her parent forever,
    ↵with the race and descent of mortals, and to be finally a blessed soul in heaven.
```

Decode multiple files sequentially

```
python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
```

You should see the following output:

```
2023-08-11 18:36:42,148 INFO [offline-websocket-client-decode-files-sequential.py:141] {
    ↵'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-
    ↵whisper-tiny.en/test_wavs/0.wav', './sherpa-onnx-whisper-tiny.en/test_wavs/1.wav', './
    ↵sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav']}
2023-08-11 18:36:42,176 INFO [offline-websocket-client-decode-files-sequential.py:114]
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav
    After early nightfall, the yellow lamps would light up here and there, the squalid
    ↵quarter of the brothels.
2023-08-11 18:36:42,926 INFO [offline-websocket-client-decode-files-sequential.py:114]
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav
    God, as a direct consequence of the sin which man thus punished, had given her a lovely
    ↵child, whose place was on that same dishonored bosom to connect her parent forever
    ↵with the race and descent of mortals, and to be finally a blessed soul in heaven.
2023-08-11 18:36:44,314 INFO [offline-websocket-client-decode-files-sequential.py:114]
    ↵Sending ./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
    Yet these thoughts affected Hester Prin less with hope than apprehension.
```

colab

We provide a colab notebook for you to try the Python non-streaming websocket server example of [sherpa-onnx](#).

8.5 C API

In this section, we describe how to use the C API of [sherpa-onnx](#).

Specifically, we will describe:

- How to generate required files
- How to use `pkg-config` with [sherpa-onnx](#)

You can find the implementation at

- <https://github.com/k2-fsa/sherpa-onnx/blob/master/sherpa-onnx/c-api/c-api.h>
- <https://github.com/k2-fsa/sherpa-onnx/blob/master/sherpa-onnx/c-api/c-api.cc>

8.5.1 Generate required files

Before using the C API of [sherpa-onnx](#), we need to first build required libraries. You can choose either to build static libraries or shared libraries.

Build shared libraries

Assume that we want to put library files and header files in the directory `/tmp/sherpa-onnx/shared`:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build-shared
cd build-shared

cmake \
-DSHERPA_ONNX_ENABLE_C_API=ON \
-DCMAKE_BUILD_TYPE=Release \
-DBUILD_SHARED_LIBS=ON \
-DCMAKE_INSTALL_PREFIX=/tmp/sherpa-onnx/shared \
..

make -j6
make install
```

You should find the following files inside `/tmp/sherpa-onnx/shared`:

macOS

Linux

```
$ tree /tmp/sherpa-onnx/shared/
/tmp/sherpa-onnx/shared/
    bin
        ├── sherpa-onnx
        ├── sherpa-onnx-microphone
        ├── sherpa-onnx-microphone-offline
        ├── sherpa-onnx-offline
        ├── sherpa-onnx-offline-websocket-server
        ├── sherpa-onnx-online-websocket-client
        └── sherpa-onnx-online-websocket-server
    include
        ├── cargs.h
        └── sherpa-onnx
            └── c-api
                └── c-api.h
    lib
        ├── cargs.h
        ├── libcargs.dylib
        ├── libkaldi-native-fbank-core.dylib
        ├── libonnxruntime.1.15.1.dylib
        ├── libonnxruntime.dylib -> libonnxruntime.1.15.1.dylib
        ├── libsherpa-onnx-c-api.dylib
        ├── libsherpa-onnx-core.dylib
        └── libsherpa-onnx-portaudio.dylib
    sherpa-onnx.pc
5 directories, 18 files
```

```
$ tree /tmp/sherpa-onnx/shared/
/tmp/sherpa-onnx/shared/
|-- bin
|   |-- sherpa-onnx
|   |-- sherpa-onnx-alsa
|   |-- sherpa-onnx-microphone
|   |-- sherpa-onnx-microphone-offline
|   |-- sherpa-onnx-offline
|   |-- sherpa-onnx-offline-websocket-server
|   |-- sherpa-onnx-online-websocket-client
|   '-- sherpa-onnx-online-websocket-server
|-- include
|   |-- cargs.h
|   '-- sherpa-onnx
|       '-- c-api
|           '-- c-api.h
|-- lib
|   |-- cargs.h
|   |-- libcargs.so
|   |-- libkaldi-native-fbank-core.so
|   |-- libonnxruntime.so -> libonnxruntime.so.1.15.1
|   |-- libonnxruntime.so.1.15.1
|   |-- libsherpa-onnx-c-api.so
|   |-- libsherpa-onnx-core.so
|   '-- libsherpa-onnx-portaudio.so
`-- sherpa-onnx.pc

5 directories, 19 files
```

Build static libraries

Assume that we want to put library files and header files in the directory `/tmp/sherpa-onnx/static`:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
mkdir build-static
cd build-static

cmake \
-DSHERPA_ONNX_ENABLE_C_API=ON \
-DCMAKE_BUILD_TYPE=Release \
-DBUILD_SHARED_LIBS=OFF \
-DCMAKE_INSTALL_PREFIX=/tmp/sherpa-onnx/static \
.

make -j6
make install
```

You should find the following files in `/tmp/sherpa-onnx/static`:

macOS

Linux

```
$ tree /tmp/sherpa-onnx/static/
/tmp/sherpa-onnx//static/
├── bin
│   ├── sherpa-onnx
│   ├── sherpa-onnx-microphone
│   ├── sherpa-onnx-microphone-offline
│   ├── sherpa-onnx-offline
│   ├── sherpa-onnx-offline-websocket-server
│   ├── sherpa-onnx-online-websocket-client
│   └── sherpa-onnx-online-websocket-server
├── include
│   ├── cargs.h
│   └── sherpa-onnx
│       └── c-api
│           └── c-api.h
└── lib
    ├── cargs.h
    ├── libcargs.a
    ├── libkaldi-native-fbank-core.a
    ├── libonnxruntime.1.15.1.dylib
    ├── libonnxruntime.dylib -> libonnxruntime.1.15.1.dylib
    ├── libsherpa-onnx-c-api.a
    ├── libsherpa-onnx-core.a
    └── libsherpa-onnx-portaudio_static.a
    sherpa-onnx.pc
```

5 directories, 18 files

```
$ tree /tmp/sherpa-onnx/static/
/tmp/sherpa-onnx/static/
|-- bin
|   |-- sherpa-onnx
|   |-- sherpa-onnx-alsa
|   |-- sherpa-onnx-microphone
|   |-- sherpa-onnx-microphone-offline
|   |-- sherpa-onnx-offline
|   |-- sherpa-onnx-offline-websocket-server
|   |-- sherpa-onnx-online-websocket-client
|   `-- sherpa-onnx-online-websocket-server
|-- include
|   |-- cargs.h
|   `-- sherpa-onnx
|       '-- c-api
|           '-- c-api.h
`-- lib
    |-- cargs.h
    |-- libcargs.a
    |-- libkaldi-native-fbank-core.a
    '-- libonnxruntime.so -> libonnxruntime.so.1.15.1
```

(continues on next page)

(continued from previous page)

```
|   |-- libonnxruntime.so.1.15.1
|   |-- libsherpa-onnx-c-api.a
|   |-- libsherpa-onnx-core.a
|   '-- libsherpa-onnx-portaudio_static.a
`-- sherpa-onnx.pc

5 directories, 19 files
```

8.5.2 Build decode-file-c-api.c with generated files

To build the following file:

```
https://github.com/k2-fsa/sherpa-onnx/blob/master/c-api-examples/decode-file-c-api.c
```

We can use:

static link

dynamic link

```
export PKG_CONFIG_PATH=/tmp/sherpa-onnx/static:$PKG_CONFIG_PATH

cd ./c-api-examples
gcc -o decode-file-c-api $(pkg-config --cflags sherpa-onnx) ./decode-file-c-api.c $(pkg-
config --libs sherpa-onnx)

./decode-file-c-api --help
```

```
export PKG_CONFIG_PATH=/tmp/sherpa-onnx/shared:$PKG_CONFIG_PATH

cd ./c-api-examples
gcc -o decode-file-c-api $(pkg-config --cflags sherpa-onnx) ./decode-file-c-api.c $(pkg-
config --libs sherpa-onnx)

./decode-file-c-api --help
```

8.5.3 colab

We provide a colab notebook for you to try the C API of `sherpa-onnx`.

8.6 Java API

In this section, we describe how to use the Java API of `sherpa-onnx`.

The core part of `sherpa-onnx` is written in C++. We have provided `JNI` interface for `sherpa-onnx` so that you can use it in Java.

Before using the Java API of `sherpa-onnx`, you have to build the `JNI` interface.

Hint: We provide pre-built `JNI` libs. Please see

<https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/jni>

For Chinese users, please use

<https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/tree/main/jni>

8.6.1 Build JNI interface (macOS)

In the following, we describe how to build the JNI interface for macOS. It is applicable for both macOS x64 and arm64.

For Linux users, please refer to [Build JNI interface \(Linux\)](#)

Hint: For Windows users, you have to modify the commands by yourself.

Setup the environment

Make sure you have the following two items ready:

- a working C/C++ compiler that supports C++17
- you are able to run java and javac commands in your terminal.

To check your environment, please run:

```
gcc --version
java -version
javac -version
```

The above three commands print the following output on my computer. You don't need to use the exact versions as I am using.

```
# output of gcc --version

Apple clang version 14.0.0 (clang-1400.0.29.202)
Target: x86_64-apple-darwin22.2.0
Thread model: posix
InstalledDir: /Applications/Xcode.app/Contents/Developer/Toolchains/XcodeDefault.xctoolchain/usr/bin

# output of java -version

openjdk version "19.0.1" 2022-10-18
OpenJDK Runtime Environment (build 19.0.1+10-21)
OpenJDK 64-Bit Server VM (build 19.0.1+10-21, mixed mode, sharing)

# output of javac -version

javac 19.0.1
```

Build sherpa-onnx

Please use the following commands to build `sherpa-onnx`:

```
git clone https://github.com/k2-fsa/sherpa-onnx  
cd sherpa-onnx  
mkdir build  
cd build  
cmake \  
  -DSHERPA_ONNX_ENABLE_PYTHON=OFF \  
  -DSHERPA_ONNX_ENABLE_TESTS=OFF \  
  -DSHERPA_ONNX_ENABLE_CHECK=OFF \  
  -DBUILD_SHARED_LIBS=ON \  
  -DSHERPA_ONNX_ENABLE_PORTAUDIO=OFF \  
  -DSHERPA_ONNX_ENABLE_JNI=ON \  
  ..  
make -j4  
ls -lh lib
```

You should see the following output for `ls -lh lib`:

```
total 11576  
-rwxr-xr-x 1 fangjun staff    33K May 15 11:24 libcargs.dylib  
-rwxr-xr-x 1 fangjun staff  350K May 15 11:24 libespeak-ng.dylib  
-rwxr-xr-x 1 fangjun staff  471K May 15 11:24 libkaldi-decoder-core.dylib  
-rwxr-xr-x 1 fangjun staff 113K May 15 11:24 libkaldi-native-fbank-core.dylib  
-rwxr-xr-x 1 fangjun staff 423K May 15 11:24 libpiper_phonemize.1.2.0.dylib  
lrwxr-xr-x 1 fangjun staff   30B May 15 11:24 libpiper_phonemize.1.dylib -> libpiper_<  
phonemize.1.2.0.dylib  
lrwxr-xr-x 1 fangjun staff   26B May 15 11:24 libpiper_phonemize.dylib -> libpiper_<  
phonemize.1.dylib  
-rwxr-xr-x 1 fangjun staff 113K May 15 11:25 libsherpa-onnx-c-api.dylib  
-rwxr-xr-x 1 fangjun staff 1.8M May 15 11:24 libsherpa-onnx-core.dylib  
-rwxr-xr-x 1 fangjun staff 1.5M May 15 11:24 libsherpa-onnx-fst.6.dylib  
lrwxr-xr-x 1 fangjun staff   26B May 15 11:24 libsherpa-onnx-fst.dylib -> libsherpa-<  
onnx-fst.6.dylib  
-rwxr-xr-x 1 fangjun staff   34K May 15 11:24 libsherpa-onnx-fstfar.7.dylib  
lrwxr-xr-x 1 fangjun staff   29B May 15 11:24 libsherpa-onnx-fstfar.dylib -> libsherpa-<  
onnx-fstfar.7.dylib  
-rwxr-xr-x 1 fangjun staff 127K May 15 11:25 libsherpa-onnx-jni.dylib  
-rwxr-xr-x 1 fangjun staff 933K May 15 11:24 libsherpa-onnx-kalidfst-core.dylib  
-rwxr-xr-x 1 fangjun staff 187K May 15 11:24 libucd.dylib
```

Note that all these `*.dylib` files are required.

`libsherpa-onnx-jni.dylib` contains the JNI interface for `sherpa-onnx`.

Download pre-built JNI libs

If you don't want to build JNI libs by yourself, please download pre-built JNI libs from

<https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/jni>

For Chinese users, please use

<https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/tree/main/jni>

Please always use the latest version. In the following, we describe how to download the version 1.10.2.

Intel CPU (x86_64)

Apple Silicon (arm64)

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2-osx-x86_64-jni.tar.bz2

# For Chinese users
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2-osx-x86_64-jni.tar.bz2

tar xf sherpa-onnx-v1.10.2-osx-x86_64-jni.tar.bz2
rm sherpa-onnx-v1.10.2-osx-x86_64-jni.tar.bz2
```

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2-osx-arm64-jni.tar.bz2

# For Chinese users
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2-osx-arm64-jni.tar.bz2

tar xf sherpa-onnx-v1.10.2-osx-arm64-jni.tar.bz2
rm sherpa-onnx-v1.10.2-osx-arm64-jni.tar.bz2
```

After downloading, you should see the following files:

```
# For x86_64
ls -lh sherpa-onnx-v1.10.2-osx-x86_64-jni/lib/

-rwxr-xr-x 1 fangjun staff 33K Jun 25 11:39 libcargs.dylib
-rwxr-xr-x 1 fangjun staff 335K Jun 25 11:39 libespeak-ng.dylib
-rwxr-xr-x 1 fangjun staff 484K Jun 25 11:39 libkaldi-decoder-core.dylib
-rwxr-xr-x 1 fangjun staff 121K Jun 25 11:39 libkaldi-native-fbank-core.dylib
-rw-r--r-- 1 fangjun staff 25M Jun 25 11:39 libonnxruntime.1.17.1.dylib
lrwxr-xr-x 1 fangjun staff 27B Jun 25 11:39 libonnxruntime.dylib -> libonnxruntime.1.17.1.dylib
-rwxr-xr-x 1 fangjun staff 427K Jun 25 11:39 libpiper_phonemize.1.2.0.dylib
lrwxr-xr-x 1 fangjun staff 30B Jun 25 11:39 libpiper_phonemize.1.dylib -> libpiper_phonemize.1.2.0.dylib
lrwxr-xr-x 1 fangjun staff 26B Jun 25 11:39 libpiper_phonemize.dylib -> libpiper_phonemize.1.dylib
-rwxr-xr-x 1 fangjun staff 115K Jun 25 11:39 libsherpa-onnx-c-api.dylib
-rwxr-xr-x 1 fangjun staff 1.9M Jun 25 11:39 libsherpa-onnx-core.dylib
-rwxr-xr-x 1 fangjun staff 1.5M Jun 25 11:39 libsherpa-onnx-fst.dylib
```

(continues on next page)

(continued from previous page)

-rwxr-xr-x	1	fangjun	staff	40K	Jun 25 11:39	libsherpa-onnx-fstfar.dylib
-rwxr-xr-x	1	fangjun	staff	134K	Jun 25 11:39	libsherpa-onnx-jni.dylib
-rwxr-xr-x	1	fangjun	staff	963K	Jun 25 11:39	libsherpa-onnx-kaldifst-core.dylib
-rwxr-xr-x	1	fangjun	staff	117K	Jun 25 11:39	libsherpa-onnx-portaudio.dylib
-rwxr-xr-x	1	fangjun	staff	150K	Jun 25 11:39	libssentencepiece_core.dylib
-rwxr-xr-x	1	fangjun	staff	187K	Jun 25 11:39	libucd.dylib
<hr/>						
# For arm64						
ls -lh sherpa-onnx-v1.10.2-osx-arm64-jni/lib/						
-rwxr-xr-x	1	fangjun	staff	49K	Jun 25 11:39	libcargs.dylib
-rwxr-xr-x	1	fangjun	staff	318K	Jun 25 11:39	libespeak-ng.dylib
-rwxr-xr-x	1	fangjun	staff	436K	Jun 25 11:39	libkaldi-decoder-core.dylib
-rwxr-xr-x	1	fangjun	staff	122K	Jun 25 11:39	libkaldi-native-fbank-core.dylib
-rw-r--r--	1	fangjun	staff	23M	Jun 25 11:39	libonnxruntime.1.17.1.dylib
lrwxr-xr-x	1	fangjun	staff	27B	Jun 25 11:39	libonnxruntime.dylib -> libonnxruntime. → 1.17.1.dylib
-rwxr-xr-x	1	fangjun	staff	430K	Jun 25 11:39	libpiper_phonemize.1.2.0.dylib
lrwxr-xr-x	1	fangjun	staff	30B	Jun 25 11:39	libpiper_phonemize.1.dylib -> libpiper_ → phonemize.1.2.0.dylib
lrwxr-xr-x	1	fangjun	staff	26B	Jun 25 11:39	libpiper_phonemize.dylib -> libpiper_ → phonemize.1.dylib
-rwxr-xr-x	1	fangjun	staff	132K	Jun 25 11:39	libsherpa-onnx-c-api.dylib
-rwxr-xr-x	1	fangjun	staff	1.7M	Jun 25 11:39	libsherpa-onnx-core.dylib
-rwxr-xr-x	1	fangjun	staff	1.4M	Jun 25 11:39	libsherpa-onnx-fst.dylib
-rwxr-xr-x	1	fangjun	staff	56K	Jun 25 11:39	libsherpa-onnx-fstfar.dylib
-rwxr-xr-x	1	fangjun	staff	151K	Jun 25 11:39	libsherpa-onnx-jni.dylib
-rwxr-xr-x	1	fangjun	staff	859K	Jun 25 11:39	libsherpa-onnx-kaldifst-core.dylib
-rwxr-xr-x	1	fangjun	staff	118K	Jun 25 11:39	libsherpa-onnx-portaudio.dylib
-rwxr-xr-x	1	fangjun	staff	149K	Jun 25 11:39	libssentencepiece_core.dylib
-rwxr-xr-x	1	fangjun	staff	188K	Jun 25 11:39	libucd.dylib

8.6.2 Build JNI interface (Linux)

In the following, we describe how to build the JNI interface for Linux. It is applicable for both Linux x64 and arm64.

For macOS users, please refer to [Build JNI interface \(macOS\)](#)

Hint: For Windows users, you have to modify the commands by yourself.

Setup the environment

Make sure you have the following two items ready:

- a working C/C++ compiler that supports C++17
- you are able to run `java` and `javac` commands in your terminal.

To check your environment, please run:

```
gcc --version
java -version
javac -version
```

The above three commands print the following output on my computer. You don't need to use the exact versions as I am using.

```
# output of gcc --version

gcc (Ubuntu 12.3.0-1ubuntu1~23.04) 12.3.0
Copyright (C) 2022 Free Software Foundation, Inc.
This is free software; see the source for copying conditions. There is NO
warranty; not even for MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE.

# output of java -version

java version "17.0.11" 2024-04-16 LTS
Java(TM) SE Runtime Environment (build 17.0.11+7-LTS-207)
Java HotSpot(TM) 64-Bit Server VM (build 17.0.11+7-LTS-207, mixed mode, sharing)

# output of javac -version

javac 17.0.11
```

Build sherpa-onnx

Please use the following commands to build `sherpa-onnx`:

```
git clone https://github.com/k2-fsa/sherpa-onnx

cd sherpa-onnx

mkdir build

cd build

cmake \
-DSHERPA_ONNX_ENABLE_PYTHON=OFF \
-DSHERPA_ONNX_ENABLE_TESTS=OFF \
-DSHERPA_ONNX_ENABLE_CHECK=OFF \
-DBUILD_SHARED_LIBS=ON \
-DSHERPA_ONNX_ENABLE_PORTAUDIO=OFF \
-DSHERPA_ONNX_ENABLE_JNI=ON \
..
```

(continues on next page)

(continued from previous page)

make -j4

ls -lh lib

You should see the following output for ls -lh lib:

```
total 7.7M
-rwxrwxr-x 1 fangjun fangjun 16K May 15 03:53 libcargs.so
-rwxrwxr-x 1 fangjun fangjun 396K May 15 03:53 libespeak-ng.so
-rwxrwxr-x 1 fangjun fangjun 652K May 15 03:53 libkaldi-decoder-core.so
-rwxrwxr-x 1 fangjun fangjun 108K May 15 03:53 libkaldi-native-fbank-core.so
lrwxrwxrwx 1 fangjun fangjun 23 May 15 03:53 libpiper_phonemize.so -> libpiper_
↳phonemize.so.1
lrwxrwxrwx 1 fangjun fangjun 27 May 15 03:53 libpiper_phonemize.so.1 -> libpiper_
↳phonemize.so.1.2.0
-rwxrwxr-x 1 fangjun fangjun 450K May 15 03:53 libpiper_phonemize.so.1.2.0
-rwxrwxr-x 1 fangjun fangjun 107K May 15 03:55 libsherpa-onnx-c-api.so
-rwxrwxr-x 1 fangjun fangjun 2.4M May 15 03:54 libsherpa-onnx-core.so
lrwxrwxrwx 1 fangjun fangjun 26 May 15 03:53 libsherpa-onnx-fstfar.so -> libsherpa-
↳onnx-fstfar.so.7
-rwxrwxr-x 1 fangjun fangjun 18K May 15 03:53 libsherpa-onnx-fstfar.so.7
lrwxrwxrwx 1 fangjun fangjun 23 May 15 03:53 libsherpa-onnx-fst.so -> libsherpa-onnx-
↳fst.so.6
-rwxrwxr-x 1 fangjun fangjun 2.1M May 15 03:53 libsherpa-onnx-fst.so.6
-rwxrwxr-x 1 fangjun fangjun 134K May 15 03:55 libsherpa-onnx-jni.so
-rwxrwxr-x 1 fangjun fangjun 1.2M May 15 03:53 libsherpa-onnx-kaldifst-core.so
-rwxrwxr-x 1 fangjun fangjun 229K May 15 03:53 libucd.so
```

Note that all these *.so files are required.

libsherpa-onnx-jni.so contains the JNI interface for sherpa-onnx.

Download pre-built JNI libs

If you don't want to build JNI libs by yourself, please download pre-built JNI libs from

<https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/jni>

For Chinese users, please use

<https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/tree/main/jni>

Please always use the latest version. In the following, we describe how to download the version 1.10.2.

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.
↳10.2-linux-x64-jni.tar.bz2

# For Chinese users
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.
↳10.2-linux-x64-jni.tar.bz2

tar xf sherpa-onnx-v1.10.2-linux-x64-jni.tar.bz2
rm sherpa-onnx-v1.10.2-linux-x64-jni.tar.bz2
```

You should find the following files:

```
ls -lh sherpa-onnx-v1.10.2-linux-x64-jni/lib/

-rwxr-xr-x 1 fangjun staff 16K Jun 25 11:41 libcargs.so
-rwxr-xr-x 1 fangjun staff 383K Jun 25 11:41 libespeak-ng.so
-rwxr-xr-x 1 fangjun staff 648K Jun 25 11:41 libkaldi-decoder-core.so
-rwxr-xr-x 1 fangjun staff 214K Jun 25 11:41 libkaldi-native-fbank-core.so
-rw-r--r-- 1 fangjun staff 15M Jun 25 11:40 libonnxruntime.so
-rw-r--r-- 1 fangjun staff 15M Jun 25 11:40 libonnxruntime.so.1.17.1
lrwxr-xr-x 1 fangjun staff 23B Jun 25 11:44 libpiper_phonemize.so -> libpiper_
↳ phonemize.so.1
lrwxr-xr-x 1 fangjun staff 27B Jun 25 11:44 libpiper_phonemize.so.1 -> libpiper_
↳ phonemize.so.1.2.0
-rwxr-xr-x 1 fangjun staff 535K Jun 25 11:41 libpiper_phonemize.so.1.2.0
-rwxr-xr-x 1 fangjun staff 128K Jun 25 11:44 libsherpa-onnx-c-api.so
-rwxr-xr-x 1 fangjun staff 2.5M Jun 25 11:43 libsherpa-onnx-core.so
-rwxr-xr-x 1 fangjun staff 2.1M Jun 25 11:41 libsherpa-onnx-fst.so
-rwxr-xr-x 1 fangjun staff 18K Jun 25 11:41 libsherpa-onnx-fstfar.so
-rwxr-xr-x 1 fangjun staff 155K Jun 25 11:44 libsherpa-onnx-jni.so
-rwxr-xr-x 1 fangjun staff 1.1M Jun 25 11:41 libsherpa-onnx-kalidifst-core.so
-rwxr-xr-x 1 fangjun staff 81K Jun 25 11:41 libsherpa-onnx-portaudio.so
-rwxr-xr-x 1 fangjun staff 275K Jun 25 11:41 libssentencepiece_core.so
-rwxr-xr-x 1 fangjun staff 224K Jun 25 11:40 libucd.so
```

8.6.3 Build JNI interface (Windows)

If you want to build JNI libs by yourself, please see <https://github.com/k2-fsa/sherpa-onnx/issues/882>.

Hint: The PDFs in the above link are in Chinese.

If you want to download pre-built JNI libs, please see below.

Download pre-built JNI libs

If you don't want to build JNI libs by yourself, please download pre-built JNI libs from

<https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/jni>

For Chinese users, please use

<https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/tree/main/jni>

Please always use the latest version. In the following, we describe how to download the version 1.10.2.

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.
↳ 10.2-win-x64-jni.tar.bz2

# For Chinese users
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.
↳ 10.2-win-x64-jni.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xf sherpa-onnx-v1.10.2-win-x64-jni.tar.bz2  
rm sherpa-onnx-v1.10.2-win-x64-jni.tar.bz2
```

You should find the following files:

```
ls -lh sherpa-onnx-v1.10.2-win-x64-jni/lib/  
  
-rwxr-xr-x 1 fangjun staff 12K Jun 25 11:43 cargs.dll  
-rw-r--r-- 1 fangjun staff 3.1K Jun 25 11:43 cargs.lib  
-rwxr-xr-x 1 fangjun staff 271K Jun 25 11:43 espeak-ng.dll  
-rw-r--r-- 1 fangjun staff 49K Jun 25 11:43 espeak-ng.lib  
-rwxr-xr-x 1 fangjun staff 533K Jun 25 11:44 kaldidecoder-core.dll  
-rw-r--r-- 1 fangjun staff 1.1M Jun 25 11:44 kaldidecoder-core.lib  
-rwxr-xr-x 1 fangjun staff 109K Jun 25 11:43 kaldinative-fbank-core.dll  
-rw-r--r-- 1 fangjun staff 103K Jun 25 11:43 kaldinative-fbank-core.lib  
-rwxr-xr-x 1 fangjun staff 10M Jun 25 11:43 onnxruntime.dll  
-rwxr-xr-x 1 fangjun staff 22K Jun 25 11:43 onnxruntime_providers_shared.dll  
-rwxr-xr-x 1 fangjun staff 439K Jun 25 11:43 piper_phonemize.dll  
-rw-r--r-- 1 fangjun staff 212K Jun 25 11:43 piper_phonemize.lib  
-rwxr-xr-x 1 fangjun staff 76K Jun 25 11:47 sherpa-onnx-c-api.dll  
-rw-r--r-- 1 fangjun staff 63K Jun 25 11:47 sherpa-onnx-c-api.lib  
-rwxr-xr-x 1 fangjun staff 1.7M Jun 25 11:47 sherpa-onnx-core.dll  
-rw-r--r-- 1 fangjun staff 2.5M Jun 25 11:47 sherpa-onnx-core.lib  
-rwxr-xr-x 1 fangjun staff 1.4M Jun 25 11:43 sherpa-onnx-fst.dll  
-rw-r--r-- 1 fangjun staff 3.8M Jun 25 11:43 sherpa-onnx-fst.lib  
-rwxr-xr-x 1 fangjun staff 30K Jun 25 11:43 sherpa-onnx-fstfar.dll  
-rw-r--r-- 1 fangjun staff 24K Jun 25 11:43 sherpa-onnx-fstfar.lib  
-rwxr-xr-x 1 fangjun staff 109K Jun 25 11:47 sherpa-onnx-jni.dll  
-rw-r--r-- 1 fangjun staff 92K Jun 25 11:47 sherpa-onnx-jni.lib  
-rwxr-xr-x 1 fangjun staff 1.1M Jun 25 11:44 sherpa-onnx-kaldifst-core.dll  
-rw-r--r-- 1 fangjun staff 2.8M Jun 25 11:44 sherpa-onnx-kaldifst-core.lib  
-rwxr-xr-x 1 fangjun staff 173K Jun 25 11:43 sherpa-onnx-portaudio.dll  
-rw-r--r-- 1 fangjun staff 43K Jun 25 11:43 sherpa-onnx-portaudio.lib  
-rwxr-xr-x 1 fangjun staff 135K Jun 25 11:43 ssentencepiece_core.dll  
-rw-r--r-- 1 fangjun staff 174K Jun 25 11:43 ssentencepiece_core.lib  
-rwxr-xr-x 1 fangjun staff 168K Jun 25 11:43 ucd.dll  
-rw-r--r-- 1 fangjun staff 5.9K Jun 25 11:43 ucd.lib
```

Hint: Only *.dll files are needed during runtime.

8.6.4 Build the jar package

Note: Please see the end of this page for how to download pre-built jar.

```
cd sherpa-onnx/sherpa-onnx/java-api  
ls -lh
```

You should see the following output:

```
(py311) fangjun@ubuntu23-04:/mnt/sdb/shared/sherpa-onnx/sherpa-onnx/java-api$ ls -lh
total 8.0K
-rw-rw-r-- 1 fangjun fangjun 2.5K May  8 06:17 Makefile
drwxrwxr-x 3 fangjun fangjun 4.0K Mar  1 04:29 src
```

Please run the following command in the directory `sherpa-onnx/java-api`:

```
make
```

You should see the following output after running `make`:

```
(py311) fangjun@ubuntu23-04:/mnt/sdb/shared/sherpa-onnx/sherpa-onnx/java-api$ ls -lh
total 12K
drwxrwxr-x 3 fangjun fangjun 4.0K May 15 03:59 build
-rw-rw-r-- 1 fangjun fangjun 2.5K May  8 06:17 Makefile
drwxrwxr-x 3 fangjun fangjun 4.0K Mar  1 04:29 src
(py311) fangjun@ubuntu23-04:/mnt/sdb/shared/sherpa-onnx/sherpa-onnx/java-api$ ls -lh
~/build/
total 60K
drwxrwxr-x 3 fangjun fangjun 4.0K May 15 03:58 com
-rw-rw-r-- 1 fangjun fangjun 53K May 15 03:59 sherpa-onnx.jar
```

Congratulations! You have generated `sherpa-onnx.jar` successfully.

Hint: You can find the Java API source files at

<https://github.com/k2-fsa/sherpa-onnx/tree/master/sherpa-onnx/java-api/src/com/k2fsa/sherpa/onnx>

Download pre-built jar

If you don't want to build jar by yourself, you can download pre-built jar from from

<https://huggingface.co/csukuangfj/sherpa-onnx-libs/tree/main/jni>

For Chinese users, please use

<https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/tree/main/jni>

Please always use the latest version. In the following, we describe how to download the version 1.10.2.

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2.jar
# For Chinese users
# wget https://hf-mirror.com/csukuangfj/sherpa-onnx-libs/resolve/main/jni/sherpa-onnx-v1.10.2.jar
```

8.6.5 Examples

Please see <https://github.com/k2-fsa/sherpa-onnx/tree/master/java-api-examples>

You can find detailed instructions at

<https://github.com/k2-fsa/sherpa-onnx/blob/master/java-api-examples/README.md>

8.7 Javascript API

For using Javascript in the browser, please see our [WebAssembly](#) doc.

This section describes how to use `sherpa-onnx` in Node with Javascript API.

8.7.1 Install

We provide npm packages for `sherpa-onnx`.

It can be found at

<https://www.npmjs.com/package/sherpa-onnx-node>

Hint: It requires Node>=v16.

Please always use the latest version.

To install it, please run:

```
npm install sherpa-onnx-node
```

It supports the following platforms:

- Linux x64
- Linux arm64
- macOS x64
- macOS arm64
- Windows x64

Hint: You don't need to pre-install anything in order to install `sherpa-onnx-node`.

That is, you don't need to install a C/C++ compiler. You don't need to install Python. You don't need to install CMake, etc.

8.7.2 Examples

Please see <https://github.com/k2-fsa/sherpa-onnx/tree/master/nodejs-addon-examples>

You can find detailed instructions at

<https://github.com/k2-fsa/sherpa-onnx/tree/master/nodejs-addon-examples/README.md>

8.8 Kotlin API

In this section, we describe how to use the **Kotlin API** of `sherpa-onnx`.

The core part of `sherpa-onnx` is written in C++. We have provided **JNI** interface for `sherpa-onnx` so that you can use it in **Kotlin**.

Before using the **Kotlin API** of `sherpa-onnx`, you have to build the **JNI** interface.

8.8.1 Build JNI interface

For macOS users, please refer to [Build JNI interface \(macOS\)](#).

For Linux users, please refer to [Build JNI interface \(Linux\)](#).

Hint: For Windows users, you have to modify the commands by yourself.

8.8.2 Examples

Please see <https://github.com/k2-fsa/sherpa-onnx/tree/master/kotlin-api-examples>

You can find detailed instructions at

<https://github.com/k2-fsa/sherpa-onnx/blob/master/kotlin-api-examples/run.sh>

8.9 Swift API

8.9.1 Build

Please use the following script to build `sherpa-onnx` for Swift API:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/build-swift-macos.sh>

The following is an example command:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx
./build-swift-macos.sh
```

8.9.2 Examples

Please see <https://github.com/k2-fsa/sherpa-onnx/tree/master/swift-api-examples>

Each example has a corresponding shell script. Please use the shell script to run it.

For instance,

- to run the text-to-speech example, please use <https://github.com/k2-fsa/sherpa-onnx/blob/master/swift-api-examples/run-tts.sh>
- to run the speech-to-text example, please use <https://github.com/k2-fsa/sherpa-onnx/blob/master/swift-api-examples/run-decode-file.sh>

8.10 Go API

In this section, we describe how to use the [Go API](#) of [sherpa-onnx](#).

The [Go API](#) of [sherpa-onnx](#) supports both streaming and non-streaming speech recognition.

The following table lists some [Go API](#) examples:

Description	URL
Decode a file with non-streaming models	https://github.com/k2-fsa/sherpa-onnx/tree/master/go-api-examples/non-streaming-decode-files
Decode a file with streaming models	https://github.com/k2-fsa/sherpa-onnx/tree/master/go-api-examples/streaming-decode-files
Real-time speech recognition from a microphone	https://github.com/k2-fsa/sherpa-onnx/tree/master/go-api-examples/real-time-speech-recognition-from-microphone

One thing to note is that we have provided pre-built libraries for [Go](#) so that you don't need to build [sherpa-onnx](#) by yourself when using the [Go API](#).

To make supporting multiple platforms easier, we split the [Go API](#) of [sherpa-onnx](#) into multiple packages, as listed in the following table:

OS	Package name	Supported Arch	Doc
Linux	sherpa-onnx-go-linux	<code>x86_64, aarch64, arm</code>	https://pkg.go.dev/github.com/k2-fsa/sherpa-onnx-go-linux
macOS	sherpa-onnx-go-macos	<code>x86_64, aarch64</code>	https://pkg.go.dev/github.com/k2-fsa/sherpa-onnx-go-macos
Windows	sherpa-onnx-go-windows	<code>x86_64, x86</code>	https://pkg.go.dev/github.com/k2-fsa/sherpa-onnx-go-windows

To simplify the usage, we have provided a single [Go](#) package for [sherpa-onnx](#) that supports multiple operating systems. It can be found at

<https://github.com/k2-fsa/sherpa-onnx-go>

Hint: Such a design is inspired by the following article:

[Cross platform Go modules for giants.](#)

You can use the following `import` to import [sherpa-onnx-go](#) into your [Go](#) project:

```
import (
    sherpa "github.com/k2-fsa/sherpa-onnx-go/sherpa_onnx"
)
```

In the following, we describe how to run our provided Go API examples.

Note: Before you continue, please make sure you have installed Go. If not, please follow <https://go.dev/doc/install> to install Go.

Hint: You need to enable `cgo` to build `sherpa-onnx-go`.

8.10.1 Decode files with non-streaming models

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/go-api-examples/non-streaming-decode-files
go mod tidy
go build
./non-streaming-decode-files --help
```

You will find the following output:

```
Usage of ./non-streaming-decode-files:
  --debug int            Whether to show debug message
  --decoder string       Path to the decoder model
  --decoding-method string Decoding method. Possible values: greedy_search, ↵
  ↵modified_beam_search (default "greedy_search")
  --encoder string        Path to the encoder model
  --joiner string         Path to the joiner model
  --lm-model string       Optional. Path to the LM model
  --lm-scale float32      Optional. Scale for the LM model (default 1)
  --max-active-paths int  Used only when --decoding-method is modified_beam_
  ↵search (default 4)
  --model-type string     Optional. Used for loading the model in a faster way
  --nemo-ctc string       Path to the NeMo CTC model
  --num-threads int        Number of threads for computing (default 1)
  --paraformer string      Path to the paraformer model
  --provider string        Provider to use (default "cpu")
  --tokens string          Path to the tokens file
pflag: help requested
```

Congratulations! You have successfully built your first Go API example for speech recognition.

Note: If you are using Windows and don't see any output after running `./non-streaming-decode-files --help`, please copy `*.dll` from https://github.com/k2-fsa/sherpa-onnx-go-windows/tree/master/lib/x86_64-pc-windows-gnu (for Win64) or <https://github.com/k2-fsa/sherpa-onnx-go-windows/tree/master/lib/i686-pc-windows-gnu> (for Win32) to the directory `sherpa-onnx/go-api-examples/non-streaming-decode-files`.

Now let us refer to [Pre-trained models](#) to download a non-streaming model.

We give several examples below for demonstration.

Non-streaming transducer

We will use [csukuangfj/sherpa-onnx-zipformer-en-2023-06-26 \(English\)](#) as an example.

First, let us download it:

```
cd sherpa-onnx/go-api-examples/non-streaming-decode-files
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
./non-streaming-decode-files \
--encoder ./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder ./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner ./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
--tokens ./sherpa-onnx-zipformer-en-2023-06-26/tokens.txt \
--model-type transducer \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
```

It should give you the following output:

```
2023/08/10 14:52:48.723098 Reading ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
2023/08/10 14:52:48.741042 Initializing recognizer (may take several seconds)
2023/08/10 14:52:51.998848 Recognizer created!
2023/08/10 14:52:51.998870 Start decoding!
2023/08/10 14:52:52.258818 Decoding done!
2023/08/10 14:52:52.258847 after early nightfall the yellow lamps would light up here.
    ↵and there the squalid quarter of the brothels
2023/08/10 14:52:52.258952 Wave duration: 6.625 seconds
```

Non-streaming paraformer

We will use [csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28 \(Chinese + English\)](#) as an example.

First, let us download it:

```
cd sherpa-onnx/go-api-examples/non-streaming-decode-files
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
paraformer-zh-2023-03-28.tar.bz2
tar xvf sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2
rm sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2
```

Now we can use:

```
./non-streaming-decode-files \
--paraformer ./sherpa-onnx-paraformer-zh-2023-03-28/model.int8.onnx \
--tokens ./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
```

(continues on next page)

(continued from previous page)

```
--model-type paraformer \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav
```

It should give you the following output:

```
2023/08/10 15:07:10.745412 Reading ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav
2023/08/10 15:07:10.758414 Initializing recognizer (may take several seconds)
2023/08/10 15:07:13.992424 Recognizer created!
2023/08/10 15:07:13.992441 Start decoding!
2023/08/10 15:07:14.382157 Decoding done!
2023/08/10 15:07:14.382847
2023/08/10 15:07:14.382898 Wave duration: 5.614625 seconds
```

Non-streaming CTC model from NeMo

We will use *stt_en_conformer_ctc_medium* as an example.

First, let us download it:

```
cd sherpa-onnx/go-api-examples/non-streaming-decode-files
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-
↪ctc-en-conformer-medium.tar.bz2
tar xvf sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
rm sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
```

Now we can use:

```
./non-streaming-decode-files \
--nemo-ctc ./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx \
--tokens ./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt \
--model-type nemo_ctc \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
```

It should give you the following output:

```
2023/08/10 15:11:48.667693 Reading ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/
↪0.wav
2023/08/10 15:11:48.680855 Initializing recognizer (may take several seconds)
2023/08/10 15:11:51.900852 Recognizer created!
2023/08/10 15:11:51.900869 Start decoding!
2023/08/10 15:11:52.125605 Decoding done!
2023/08/10 15:11:52.125630 after early nightfall the yellow lamps would light up here_
↪and there the squalid quarter of the brothels
2023/08/10 15:11:52.125645 Wave duration: 6.625 seconds
```

8.10.2 Decode files with streaming models

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/go-api-examples/streaming-decode-files
go mod tidy
go build
./streaming-decode-files --help
```

You will find the following output:

```
Usage of ./streaming-decode-files:
  --debug int            Whether to show debug message
  --decoder string       Path to the decoder model
  --decoding-method string Decoding method. Possible values: greedy_search, u
  ↵modified_beam_search (default "greedy_search")
    --encoder string      Path to the encoder model
    --joiner string        Path to the joiner model
    --max-active-paths int Used only when --decoding-method is modified_beam_
  ↵search (default 4)
    --model-type string   Optional. Used for loading the model in a faster way
    --num-threads int     Number of threads for computing (default 1)
    --provider string     Provider to use (default "cpu")
    --tokens string        Path to the tokens file
pflag: help requested
```

Note: If you are using Windows and don't see any output after running `./streaming-decode-files --help`, please copy `*.dll` from https://github.com/k2-fsa/sherpa-onnx-go-windows/tree/master/lib/x86_64-pc-windows-gnu (for Win64) or <https://github.com/k2-fsa/sherpa-onnx-go-windows/tree/master/lib/i686-pc-windows-gnu> (for Win32) to the directory `sherpa-onnx/go-api-examples/streaming-decode-files`.

Now let us refer to [Pre-trained models](#) to download a streaming model.

We give one example below for demonstration.

Streaming transducer

We will use [csukuangji/sherpa-onnx-streaming-zipformer-en-2023-06-26 \(English\)](https://github.com/csukuangji/sherpa-onnx-streaming-zipformer-en-2023-06-26) as an example.

First, let us download it:

```
cd sherpa-onnx/go-api-examples/streaming-decode-files
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
  ↵streaming-zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
./streaming-decode-files \
  --encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
  ↵16-left-128.onnx \
```

(continues on next page)

(continued from previous page)

```
--decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
↪16-left-128.onnx \
--joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
↪16-left-128.onnx \
--tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--model-type zipformer2 \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
```

It should give you the following output:

```
2023/08/10 15:17:00.226228 Reading ./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_
↪wavs/0.wav
2023/08/10 15:17:00.241024 Initializing recognizer (may take several seconds)
2023/08/10 15:17:03.352697 Recognizer created!
2023/08/10 15:17:03.352711 Start decoding!
2023/08/10 15:17:04.057130 Decoding done!
2023/08/10 15:17:04.057215 after early nightfall the yellow lamps would light up here.
↪and there the squalid quarter of the brothels
2023/08/10 15:17:04.057235 Wave duration: 6.625 seconds
```

8.10.3 Real-time speech recognition from microphone

Hint: You need to install portaudio for this example.

```
# for macOS
brew install portaudio
export PKG_CONFIG_PATH=/usr/local/Cellar/portaudio/19.7.0

# for Ubuntu
sudo apt-get install libasound-dev portaudio19-dev libportaudio2 libportaudiocpp0
```

To check that you have installed portaudio successfully, please run:

```
pkg-config --cflags --libs portaudio-2.0
```

It should give you something like below:

```
# for macOS
-I/usr/local/Cellar/portaudio/19.7.0/include -L/usr/local/Cellar/portaudio/19.
↪7.0/lib -lportaudio -framework CoreAudio -framework AudioToolbox -framework
↪AudioUnit -framework CoreFoundation -framework CoreServices

# for Ubuntu
-pthread -lportaudio -lasound -lm -lpthread
```

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/go-api-examples/real-time-speech-recognition-from-microphone
go mod tidy
```

(continues on next page)

(continued from previous page)

```
go build
./real-time-speech-recognition-from-microphone --help
```

You will find the following output:

```
Select default input device: MacBook Pro Microphone
Usage of ./real-time-speech-recognition-from-microphone:
    --debug int                                Whether to show debug message
    --decoder string                            Path to the decoder model
    --decoding-method string                   Decoding method. Possible values: greedy_
↳ search, modified_beam_search (default "greedy_search")
    --enable-endpoint int                     Whether to enable endpoint (default 1)
    --encoder string                          Path to the encoder model
    --joiner string                           Path to the joiner model
    --max-active-paths int                  Used only when --decoding-method is_
↳ modified_beam_search (default 4)
    --model-type string                      Optional. Used for loading the model in a_
↳ faster way
    --num-threads int                        Number of threads for computing (default 1)
    --provider string                         Provider to use (default "cpu")
    --rule1-min-trailing-silence float32     Threshold for rule1 (default 2.4)
    --rule2-min-trailing-silence float32     Threshold for rule2 (default 1.2)
    --rule3-min-utterance-length float32      Threshold for rule3 (default 20)
    --tokens string                           Path to the tokens file
pflag: help requested
```

Now let us refer to [Pre-trained models](#) to download a streaming model.

We give one example below for demonstration.

Streaming transducer

We will use [csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26 \(English\)](#) as an example.

First, let us download it:

```
cd sherpa-onnx/go-api-examples/real-time-speech-recognition-from-microphone
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↳ streaming-zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
./real-time-speech-recognition-from-microphone \
    --encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
↳ 16-left-128.onnx \
    --decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
↳ 16-left-128.onnx \
    --joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
↳ 16-left-128.onnx \
    --tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
    --model-type zipformer2
```

It should give you the following output:

```
Select default input device: MacBook Pro Microphone
2023/08/10 15:22:00 Initializing recognizer (may take several seconds)
2023/08/10 15:22:03 Recognizer created!
Started! Please speak
0: this is the first test
1: this is the second
```

8.10.4 colab

We provide a colab notebook for you to try the Go API examples of `sherpa-onnx`.

8.11 C# API

In this section, we describe how to use the C# API examples of `sherpa-onnx`.

The C# API of `sherpa-onnx` supports both streaming and non-streaming speech recognition.

The following table lists some C# API examples:

Description	URL
Decode a file with non-streaming models	https://github.com/k2-fsa/sherpa-onnx/tree/master/dotnet-examples/offline-decode-files
Decode a file with streaming models	https://github.com/k2-fsa/sherpa-onnx/tree/master/dotnet-examples/online-decode-files
Real-time speech recognition from a microphone	https://github.com/k2-fsa/sherpa-onnx/tree/master/dotnet-examples/speech-recognition-from-microphone

You can find the implementation in the following files:

- API for **streaming** speech recognition
<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/dotnet/online.cs>
- API for **non-streaming** speech recognition
<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/dotnet/offline.cs>

We also provide a nuget package for `sherpa-onnx`:

<https://www.nuget.org/packages/org.k2fsa.sherpa.onnx>

You can use the following statement in your `csproj` file to introduce the dependency on `sherpa-onnx`:

```
<PackageReference Include="org.k2fsa.sherpa.onnx" Version="*" />
```

One thing to note is that we have provided pre-built libraries for C# so that you don't need to build `sherpa-onnx` by yourself when using the C# API.

In the following, we describe how to run our provided C# API examples.

Note: Before you continue, please make sure you have installed .Net. If not, please follow <https://dotnet.microsoft.com/en-us/download> to install .Net.

Hint: .Net supports Windows, macOS, and Linux.

8.11.1 Decode files with non-streaming models

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/dotnet-examples/offline-decode-files/
dotnet build -c Release
./bin/Release/net6.0/offline-decode-files --help
```

You will find the following output:

```
# Zipformer

dotnet run \
--tokens=./sherpa-onnx-zipformer-en-2023-04-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.onnx \
--files ./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav

Please refer to
https://k2-fsa.github.io/sherpa/onnx/pretrained\_models/offline-transducer/index.html
to download pre-trained non-streaming zipformer models.

# Paraformer

dotnet run \
--tokens=./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2023-03-28/model.onnx \
--files ./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav

Please refer to
https://k2-fsa.github.io/sherpa/onnx/pretrained\_models/offline-paraformer/index.html
to download pre-trained paraformer models

# NeMo CTC

dotnet run \
--tokens=./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt \
--nemo-ctc=./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx \
--num-threads=1 \
--files ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
```

(continues on next page)

(continued from previous page)

Please refer to
https://k2-fsa.github.io/sherpa/onnx/pretrained_models/offline-ctc/index.html
 to download pre-trained paraformer models

Copyright (c) 2023 Xiaomi Corporation

--tokens	Path to tokens.txt
--encoder	Path to encoder.onnx. Used only for transducer models
--decoder	Path to decoder.onnx. Used only for transducer models
--joiner	Path to joiner.onnx. Used only for transducer models
--paraformer	Path to model.onnx. Used only for paraformer models
--nemo-ctc	Path to model.onnx. Used only for NeMo CTC models
--num-threads	(Default: 1) Number of threads for computation
--decoding-method	(Default: greedy_search) Valid decoding methods are: greedy_search, modified_beam_search
--max-active-paths	(Default: 4) Used only when --decoding--method is modified_beam_search. It specifies number of active paths to keep during the search
--files	Required. Audio files for decoding
--help	Display this help screen.
--version	Display version information.

Now let us refer to *Pre-trained models* to download a non-streaming model.

We give several examples below for demonstration.

Non-streaming transducer

We will use *csukuangfj/sherpa-onné-zipformer-en-2023-06-26 (English)* as an example.

First, let us download it:

```
cd sherpa-onné/dotnet-examples/offline-decode-files
wget https://github.com/k2-fsa/sherpa-onné/releases/download/asr-models/sherpa-onné-
zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onné-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onné-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
dotnet run -c Release \
--encoder ./sherpa-onné-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder ./sherpa-onné-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner ./sherpa-onné-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
--tokens ./sherpa-onné-zipformer-en-2023-06-26/tokens.txt \
--files ./sherpa-onné-zipformer-en-2023-06-26/test_wavs/0.wav \
./sherpa-onné-zipformer-en-2023-06-26/test_wavs/1.wav \
./sherpa-onné-zipformer-en-2023-06-26/test_wavs/8k.wav
```

It should give you the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/offline-stream.  
↳ cc:AcceptWaveformImpl:117 Creating a resampler:  
    in_sample_rate: 8000  
    output_sample_rate: 16000  
  
-----  
. ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav  
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID  
↳ QUARTER OF THE BROTHELS  
  
-----  
. ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav  
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY  
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH  
↳ THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN  
  
-----  
. ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav  
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION  
-----
```

Non-streaming paraformer

We will use [csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28 \(Chinese + English\)](https://csukuangfj.github.io/sherpa-onnx-sherpa-onnx-paraformer-zh-2023-03-28.html) as an example.

First, let us download it:

```
cd sherpa-onnx/dotnet-examples/offline-decode-files  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
↳ paraformer-zh-2023-03-28.tar.bz2  
tar xvf sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2  
rm sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2
```

Now we can use:

```
dotnet run -c Release \  
--paraformer ./sherpa-onnx-paraformer-zh-2023-03-28/model.int8.onnx \  
--tokens ./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \  
--files ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \  
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \  
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

It should give you the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/offline-stream.  
↳ cc:AcceptWaveformImpl:117 Creating a resampler:  
    in_sample_rate: 8000  
    output_sample_rate: 16000  
  
-----  
. ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav  
  
-----  
. ./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav
```

(continues on next page)

(continued from previous page)

```
-----  
./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/8k.wav  
-----
```

Non-streaming CTC model from NeMo

We will use `stt_en_conformer_ctc_medium` as an example.

First, let us download it:

```
cd sherpa-onnx/dotnet-examples/offline-decode-files
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-
→ctc-en-conformer-medium.tar.bz2
tar xvf sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
rm sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
```

Now we can use:

```
dotnet run -c Release \
--nemo-ctc ./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx \
--tokens ./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt \
--files ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
```

It should give you the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
→cc:AcceptWaveformImpl:117 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

-----
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
→quarter of the brothels
-----

./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav
god as a direct consequence of the sin which man thus punished had given her a lovely_
→child whose place was on that same dishonored bosom to connect her parent for ever_
→with the race and descent of mortals and to be finally a blessed soul in heaven
-----

./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
yet these thoughts affected hester pryne less with hope than apprehension
-----
```

8.11.2 Decode files with streaming models

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/dotnet-examples/online-decode-files
dotnet build -c Release
./bin/Release/net6.0/online-decode-files --help
```

You will find the following output:

```
dotnet run \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-
˓→99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-
˓→99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-
˓→avg-1.onnx \
--num-threads=2 \
--decoding-method=modified_beam_search \
--debug=false \
--files ./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/0.wav \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.wav
```

Please refer to

https://k2-fsa.github.io/sherpa/onnx/pretrained_models/online-transducer/index.html
to download pre-trained streaming models.

Copyright (c) 2023 Xiaomi Corporation

--tokens	Required. Path to tokens.txt
--provider	(Default: cpu) Provider, e.g., cpu, coreml
--encoder	Required. Path to encoder.onnx
--decoder	Required. Path to decoder.onnx
--joiner	Required. Path to joiner.onnx
--num-threads	(Default: 1) Number of threads for computation
--decoding-method	(Default: greedy_search) Valid decoding methods are: greedy_search, modified_beam_search
--debug	(Default: false) True to show model info during loading
--sample-rate	(Default: 16000) Sample rate of the data used to train the model
--max-active-paths	(Default: 4) Used only when --decoding--method is modified_beam_search. It specifies number of active paths to keep during the search
--enable-endpoint	(Default: false) True to enable endpoint detection.
--rule1-min-trailing-silence	(Default: 2.4) An endpoint is detected if trailing silence in seconds is larger than this value even if nothing has been decoded. Used only when --enable-endpoint

(continues on next page)

(continued from previous page)

--rule2-min-trailing-silence	is true. (Default: 1.2) An endpoint is detected if trailing silence in seconds is larger than this value after something that is not blank has been decoded. Used only when --enable-endpoint is true.
--rule3-min-utterance-length	(Default: 20) An endpoint is detected if the utterance in seconds is larger than this value. Used only when --enable-endpoint is true.
--files	Required. Audio files for decoding
--help	Display this help screen.
--version	Display version information.

Now let us refer to *Pre-trained models* to download a streaming model.

We give one example below for demonstration.

Streaming transducer

We will use *cskuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26 (English)* as an example.

First, let us download it:

```
cd sherpa-onnx/dotnet-examples/online-decode-files/
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
dotnet run -c Release \
--encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
˓→16-left-128.onnx \
--tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--files ./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/8k.wav
```

You will find the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/features.
˓→cc:AcceptWaveform:76 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

-----
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
↳ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY_
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER_
↳ WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTIONATELY HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
```

8.11.3 Real-time speech recognition from microphone

First, let us build the example:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx/dotnet-examples/speech-recognition-from-microphone
dotnet build -c Release
./bin/Release/net6.0/speech-recognition-from-microphone --help
```

You will find the following output:

```
dotnet run -c Release \
--tokens ./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_char/tokens.
↳ .txt \
--encoder ./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-
↳ avg-4-chunk-16-left-128.onnx \
--decoder ./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-
↳ avg-4-chunk-16-left-128.onnx \
--joiner ./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-
↳ avg-4-chunk-16-left-128.onnx \
```

Please refer to
https://k2-fsa.github.io/sherpa/onnx/pretrained_models/online-transducer/index.html
to download pre-trained streaming models.

Copyright (c) 2023 Xiaomi Corporation

--tokens	Required. Path to tokens.txt
--provider	(Default: cpu) Provider, e.g., cpu, coreml
--encoder	Required. Path to encoder.onnx
--decoder	Required. Path to decoder.onnx
--joiner	Required. Path to joiner.onnx
--num-threads	(Default: 1) Number of threads for computation
--decoding-method	(Default: greedy_search) Valid decoding methods are: greedy_search, modified_beam_search
--debug	(Default: false) True to show model info during loading
--sample-rate	(Default: 16000) Sample rate of the data used

(continues on next page)

(continued from previous page)

--max-active-paths	to train the model (Default: 4) Used only when --decoding--method is modified_beam_search. It specifies number of active paths to keep during the search
--enable-endpoint	(Default: true) True to enable endpoint detection.
--rule1-min-trailing-silence	(Default: 2.4) An endpoint is detected if trailing silence in seconds is larger than this value even if nothing has been decoded. Used only when --enable-endpoint is true.
--rule2-min-trailing-silence	(Default: 0.8) An endpoint is detected if trailing silence in seconds is larger than this value after something that is not blank has been decoded. Used only when --enable-endpoint is true.
--rule3-min-utterance-length	(Default: 20) An endpoint is detected if the utterance in seconds is larger than this value. Used only when --enable-endpoint is true.
--help	Display this help screen.
--version	Display version information.

Now let us refer to [Pre-trained models](#) to download a streaming model.

We give one example below for demonstration.

Streaming transducer

We will use [csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26 \(English\)](#) as an example.

First, let us download it:

```
cd sherpa-onnx/dotnet-examples/speech-recognition-from-microphone
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→streaming-zipformer-en-2023-06-26.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Now we can use:

```
dotnet run -c Release \
--encoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
→16-left-128.onnx \
--decoder ./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
→16-left-128.onnx \
--joiner ./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
→16-left-128.onnx \
--tokens ./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt
```

You will find the following output:

```
PortAudio V19.7.0-devel, revision 147dd722548358763a8b649b3e4b41dfffbcfbb6
Number of devices: 5
Device 0
    Name: Background Music
    Max input channels: 2
    Default sample rate: 44100
Device 1
    Name: Background Music (UI Sounds)
    Max input channels: 2
    Default sample rate: 44100
Device 2
    Name: MacBook Pro Microphone
    Max input channels: 1
    Default sample rate: 48000
Device 3
    Name: MacBook Pro Speakers
    Max input channels: 0
    Default sample rate: 48000
Device 4
    Name: WeMeet Audio Device
    Max input channels: 2
    Default sample rate: 48000

Use default device 2 (MacBook Pro Microphone)
StreamParameters [
    device=2
    channelCount=1
    sampleFormat=Float32
    suggestedLatency=0.03452083333333334
    hostApiSpecificStreamInfo?=[False]
]
Started! Please speak

0: THIS IS A TEST
1: THIS IS A SECOND TEST
```

8.11.4 colab

We provide a colab notebook for you to try the C# API examples of sherpa-onnx.

8.12 Pascal API

We provide APIs for Object Pascal.

In other words, you can develop the following types of applications using Object Pascal:

- Voice activity detection
- Streaming speech recognition (i.e., real-time speech recognition)
- Non-streaming speech recognition

on Windows, Linux, and macOS.

Hint: For macOS, both Apple Silicon (i.e., macOS arm64, M1/M2/M3) and Intel chips are supported.

Note: We will support text-to-speech, audio tagging, keyword spotting, speaker recognition, speech identification, and spoken language identification with object pascal later.

In the following, we describe how to use the object pascal API to decode files.

We use macOS below as an example. You can adapt it for Linux and Windows.

Hint: We support both static link and dynamic link; the example below uses dynamic link. You can pass `-DBUILD_SHARED_LIBS=OFF` to `cmake` if you want to use static link.

On the Windows platform, it supports only dynamic link though.

8.12.1 Install free pascal

Please visit https://wiki.freepascal.org/Installing_the_Free_Pascal_Compiler for installation.

To check that you have installed `fpc` successfully, please run:

```
fpc -h
```

which should print the usage information of `fpc`.

8.12.2 Build sherpa-onnx

```
mkdir -p $HOME/open-source
cd $HOME/open-source
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

mkdir build
cd build

cmake \
-DBUILD_SHARED_LIBS=ON \
-DCMAKE_INSTALL_PREFIX=../install \
..

make
make install

ls -lh install/lib
```

You should get the following two shared library files:

```
(py38) fangjuns-MacBook-Pro:build fangjun$ ls -lh install/lib/
total 59696
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 25M Aug 14 14:09 libonnxruntime.1.17.1.dylib  
lrwxr-xr-x 1 fangjun staff 27B Aug 14 14:18 libonnxruntime.dylib -> libonnxruntime.  
→ 1.17.1.dylib  
-rwxr-xr-x 1 fangjun staff 3.9M Aug 15 15:01 libsherpa-onnx-c-api.dylib
```

8.12.3 Non-streaming speech recognition from files

We use the `Whisper` `tiny.en` model below as an example.

Hint: We have hardcoded the model filenames in the code.

```
cd $HOME/open-source/sherpa-onnx  
  
cd pascal-api-examples/non-streaming-asr/  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
→ whisper-tiny.en.tar.bz2  
tar xvf sherpa-onnx-whisper-tiny.en.tar.bz2  
rm sherpa-onnx-whisper-tiny.en.tar.bz2  
  
fpc \  
-dSHERPA_ONNX_USE_SHARED_LIBS \  
-Fu$HOME/open-source/sherpa-onnx/sherpa-onnx/pascal-api \  
-Fl$HOME/open-source/sherpa-onnx/build/install/lib \  
./whisper.pas  
  
# It will generate a file ./whisper
```

The output logs of the above `fpc` command are given below:

```
Free Pascal Compiler version 3.2.2 [2021/05/16] for x86_64  
Copyright (c) 1993-2021 by Florian Klaempfl and others  
Target OS: Darwin for x86_64  
Compiling ./whisper.pas  
Compiling /Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/pascal-api/sherpa_onnx.pas  
Assembling sherpa_onnx  
Assembling whisper  
Linking whisper  
ld: warning: dylib (/Users/fangjun/open-source/sherpa-onnx/build/install/lib//libsherpa-  
→ onnx-c-api.dylib) was built for newer macOS version (10.14) tha  
n being linked (10.8)  
1530 lines compiled, 3.8 sec
```

Explanation of the options for the `fpc` command:

- `-dSHERPA_ONNX_USE_SHARED_LIBS`

It defines a symbol `SHERPA_ONNX_USE_SHARED_LIBS`, which means we want to use dynamic link in the code. If you omit it, it will use static link. Please search for the string `SHERPA_ONNX_USE_SHARED_LIBS` in the file https://github.com/k2-fsa/sherpa-onnx/blob/master/sherpa-onnx/pascal-api/sherpa_onnx.pas if you want to learn more.

- -Fu\$HOME/open-source/sherpa-onnx/pascal-api

It specifies the unit search path. Our `sherpa_onnx.pas` is inside the directory `$HOME/open-source/sherpa-onnx/pascal-api` and we have to tell `fpc` where to find it.

- -Fl\$HOME/sherpa-onnx/build/install/lib

It tells `fpc` where to look for `libsherpa-onnx-c-api.dylib`.

After running the above `fpc` command, we will find an executable file `whisper` in the current directory, i.e., `$HOME/open-source/sherpa-onnx/pascal-api-examples/non-streaming-asr/whisper`:

```
(py38) fangjuns-MacBook-Pro:non-streaming-asr fangjun$ ls -lh ./whisper
-rwxr-xr-x 1 fangjun staff 2.3M Aug 16 12:13 ./whisper
```

If we run it:

```
(py38) fangjuns-MacBook-Pro:non-streaming-asr fangjun$ ./whisper
dyld[23162]: Library not loaded: @rpath/libsherpa-onnx-c-api.dylib
  Referenced from: <3AE58F60-4925-335D-89A5-B30FD7D97D7E> /Users/fangjun/open-source/
  ↵sherpa-onnx/pascal-api-examples/non-streaming-asr/whisper
  Reason: tried: '/Users/fangjun/py38/lib/python3.8/site-packages/libsherpa-onnx-c-api.
  ↵dylib' (no such file), '/usr/local/Cellar/ghostscript/9.55.0/lib/libsherpa-onnx-c-api.
  ↵dylib' (no such file), '/Users/fangjun/py38/lib/python3.8/site-packages/libsherpa-onnx-
  ↵c-api.dylib' (no such file), '/usr/local/Cellar/ghostscript/9.55.0/lib/libsherpa-onnx-
  ↵c-api.dylib' (no such file), '/libsherpa-onnx-c-api.dylib' (no such file), '/System/
  ↵Volumes/Preboot/Cryptexes/OS@rpath/libsherpa-onnx-c-api.dylib' (no such file), '/usr/
  ↵local/lib/libsherpa-onnx-c-api.dylib' (no such file), '/usr/lib/libsherpa-onnx-c-api.
  ↵dylib' (no such file, not in dyld cache)
Abort trap: 6
```

You can see it cannot find `libsherpa-onnx-c-api.dylib`.

At the compilation time, we have used `-Fl$HOME/sherpa-onnx/build/install/lib` to tell the compiler `fpc` where to find `libsherpa-onnx-c-api.dylib`.

At the runtime, we also need to do something to tell the executable where to look for `libsherpa-onnx-c-api.dylib`.

The following command does exactly that:

```
(py38) fangjuns-MacBook-Pro:non-streaming-asr fangjun$ export DYLD_LIBRARY_PATH=$HOME/
  ↵open-source/sherpa-onnx/build/install/lib:$DYLD_LIBRARY_PATH
(py38) fangjuns-MacBook-Pro:non-streaming-asr fangjun$ ./whisper
TSherpaOnnxOfflineRecognizerResult(Text := After early nightfall, the yellow lamps,
  ↵would light up here and there the squalid quarter of the brothels., Tokens := [ After,
  ↵early, night, fall, ,, the, yellow, lamps, would, light, up, here, and,
  ↵there, the, squ, alid, quarter, of, the, bro, the, ls, .], Timestamps := [])
NumThreads 1
Elapsed 0.803 s
Wave duration 6.625 s
RTF = 0.803/6.625 = 0.121
```

Hint: If you are using Linux, please replace `DYLD_LIBRARY_PATH` with `LD_LIBRARY_PATH`.

Congratulations! You have successfully managed to use the object pascal API with Whisper for speech recognition!

You can find more examples at:

<https://github.com/k2-fsa/sherpa-onnx/tree/master/pascal-api-examples>

8.12.4 Colab notebook

We provide a colab notebook for you to try this section step by step.

8.13 Lazarus

We also provide examples for developing with <https://www.lazarus-ide.org/> using Object Pascal.

We provide support for the following platforms and architectures:

- Linux-x64
- Windows-x64
- macOS-x64
- macOS-arm64

8.13.1 Pre-built APPs using Lazarus

This page lists some pre-built APPs using Lazarus with Object Pascal.

Hint: It runs locally on CPU without accessing the network.

All you need is to download it, unzip it, and then double click it to run it.

No need to install anything!

CPU

	URL
Generate subtitles ()	https://k2-fsa.github.io/sherpa/onnx/lazarus/download-generated-subtitles.html

8.13.2 Generate subtitles

This page describes how to run the code in the following directory:

https://github.com/k2-fsa/sherpa-onnx/tree/master/lazarus-examples/generate_subtitles

Hint: Before you continue, we assume you have installed [Lazarus](#).

Screenshots on different platforms

The same code can be compiled without any modifications for different operating systems and architectures.

That is **WOMA**,

Write once, compile anywhere.

The following screenshots give an example about that.

Linux x64 screenshot

Windows x64 screenshot

macOS x64 screenshot

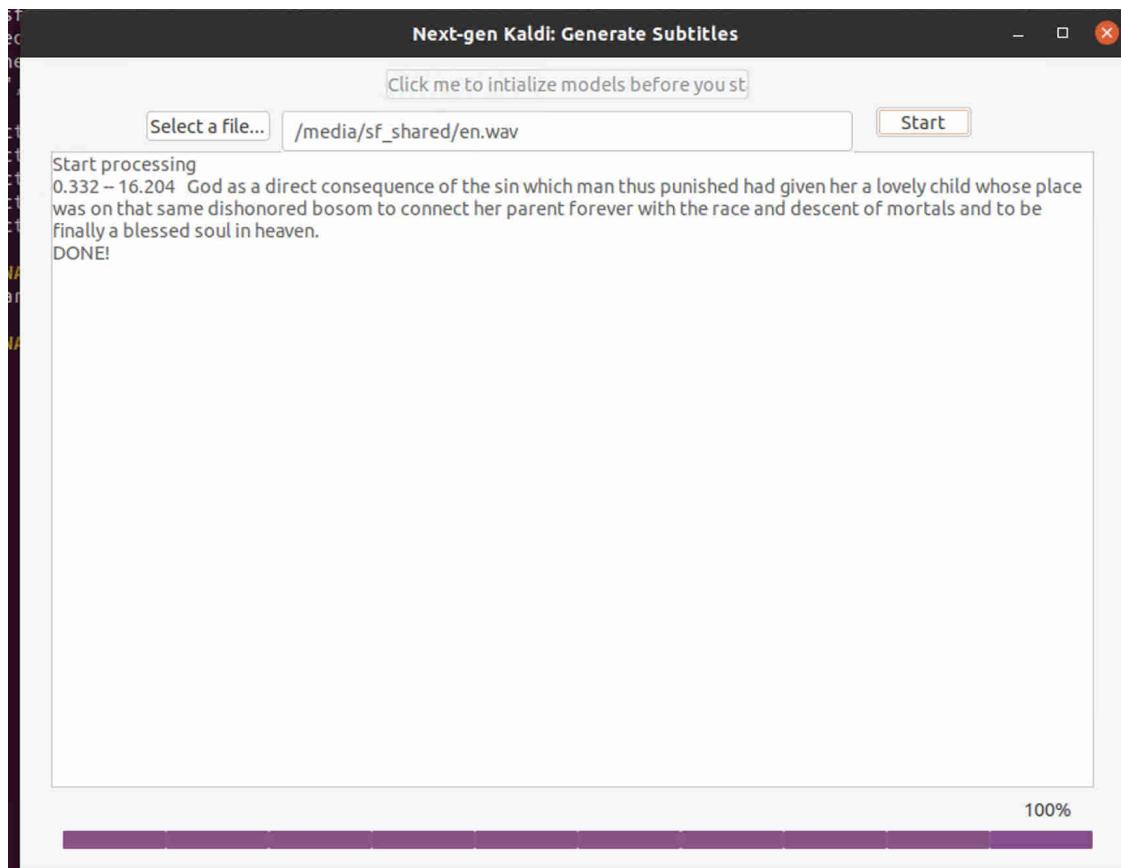


Fig. 8.1: Linux-x64 screenshot

Get sherpa-onnx libraries

sherpa-onnx is implemented in C++. To use it with Object Pascal, we have to get either the static library or the dynamic library for **sherpa-onnx**.

To achieve that, you can either build **sherpa-onnx** from source or download pre-built libraries from

<https://github.com/k2-fsa/sherpa-onnx/releases>

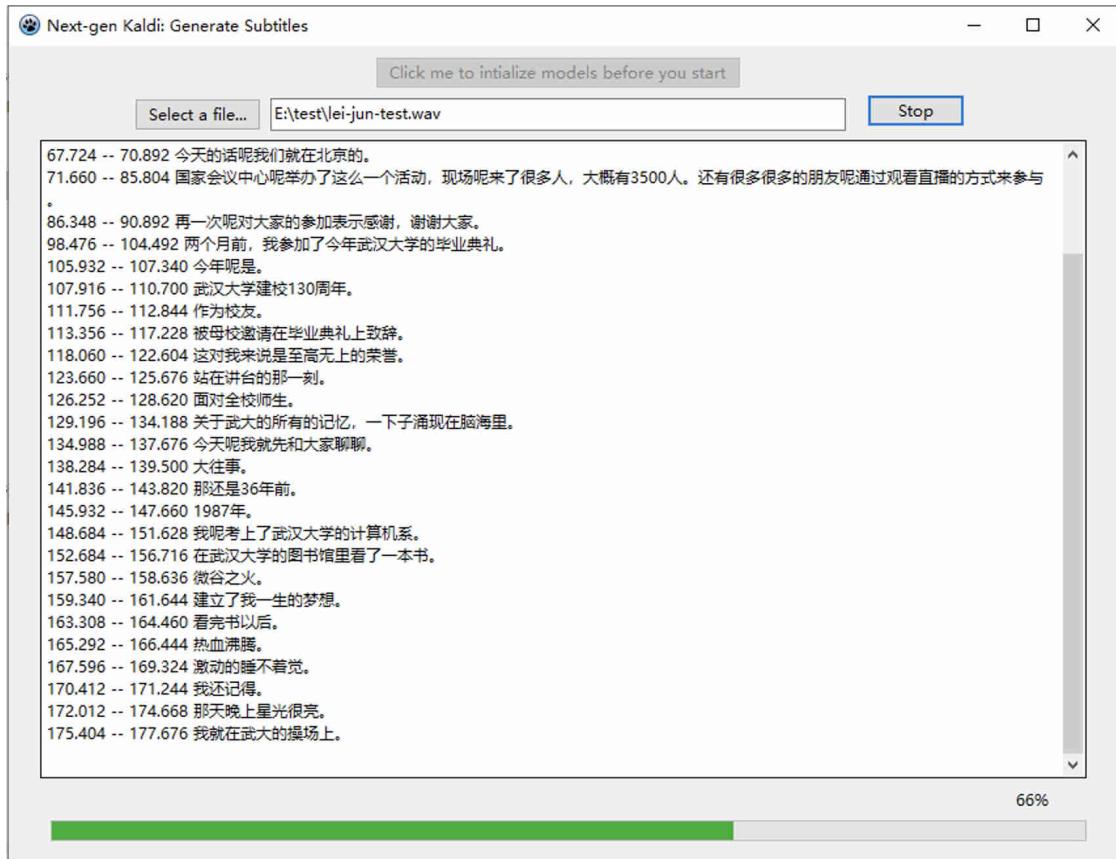


Fig. 8.2: Windows-x64 screenshot

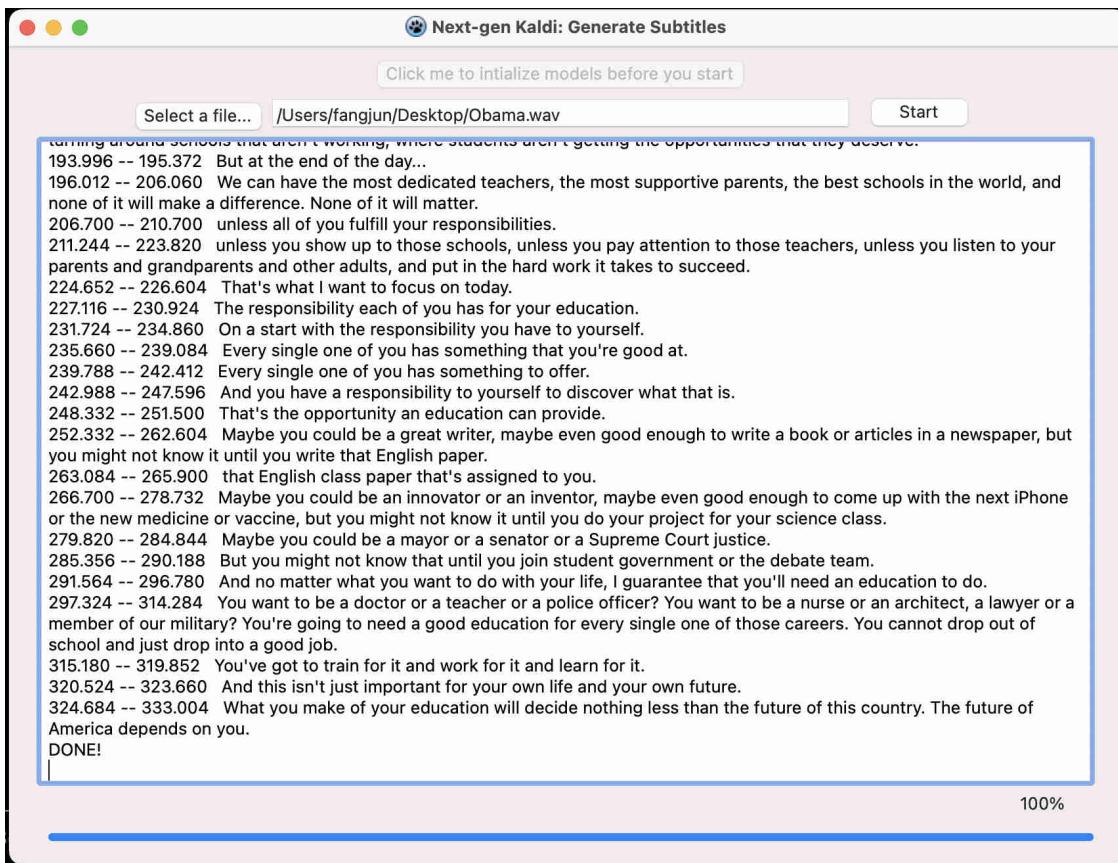


Fig. 8.3: macOS-x64 screenshot

1. Build sherpa-onnx from source

The following code builds shared libraries for sherpa-onnx:

```
mkdir -p $HOME/open-source/
cd $HOME/open-source/
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

# The build directory must be named as "build"
# for shared libraries

mkdir build
cd build

cmake \
-DBUILD_SHARED_LIBS=ON \
-DCMAKE_BUILD_TYPE=Release \
-DCMAKE_INSTALL_PREFIX=../install \
..
cmake --build . --target install --config Release
```

The following code builds static libraries for sherpa-onnx:

```
mkdir -p $HOME/open-source/
cd $HOME/open-source/
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

# The build directory must be named as "build-static"
# for shared libraries

mkdir build-static
cd build-static

cmake \
-DBUILD_SHARED_LIBS=OFF \
-DCMAKE_BUILD_TYPE=Release \
-DCMAKE_INSTALL_PREFIX=../install \
..
cmake --build . --target install --config Release
```

Caution:

- For building shared libraries, the build directory must be `build`.
- For building static libraries, the build directory must be `build-static`.

If you want to learn why there are such constraints, please search for `build-static` in the file `generate_subtitles.lpi`

2. Download pre-built libraries

If you don't want to build `sherpa-onnx` from source, please download pre-built libraries from <https://github.com/k2-fsa/sherpa-onnx/releases>.

We suggest that you always use the latest release.

	Required dynamic library files	Required static library files
Windows	<code>sherpa-onnx-c-api.dll</code> , <code>onnxruntime.dll</code>	N/A (We only support dynamic linking with <code>sherpa-onnx</code> in Lazarus on Windows)
Linux	<ul style="list-style-type: none"> • <code>libsherpa-onnx-c-api.so</code> • <code>libonnxruntime.so</code> 	<ul style="list-style-type: none"> • <code>libsherpa-onnx-c-api.a</code> • <code>libsherpa-onnx-core.a</code> • <code>libkaldi-decoder-core.a</code> • <code>libsherpa-onnx-kaldifst-core.a</code> • <code>libsherpa-onnx-fstfar.a</code> • <code>libsherpa-onnx-fst.a</code> • <code>libkaldi-native-fbank-core.a</code> • <code>libpiper_phonemize</code> • <code>liblipespeak-ng.a</code> • <code>libucd.a</code> • <code>liblibonnxruntime.a</code> • <code>libssentencepiece_core.a</code>
macOS	<ul style="list-style-type: none"> • <code>libsherpa-onnx-c-api.dylib</code> • <code>libonnxruntime.1.17.1.dylib</code> 	<ul style="list-style-type: none"> • <code>libsherpa-onnx-c-api.a</code> • <code>libsherpa-onnx-core.a</code> • <code>libkaldi-decoder-core.a</code> • <code>libsherpa-onnx-kaldifst-core.a</code> • <code>libsherpa-onnx-fstfar.a</code> • <code>libsherpa-onnx-fst.a</code> • <code>libkaldi-native-fbank-core.a</code> • <code>libpiper_phonemize</code> • <code>liblipespeak-ng.a</code> • <code>libucd.a</code> • <code>liblibonnxruntime.a</code> • <code>libssentencepiece_core.a</code>

If you download shared libraries, please create a build directory inside the `sherpa-onnx` project directory and put the library files into `build/install/lib`. An example on my macOS is given below:

```
(py38) fangjuns-MacBook-Pro:sherpa-onnx fangjun$ pwd  
/Users/fangjun/open-source/sherpa-onnx  
(py38) fangjuns-MacBook-Pro:sherpa-onnx fangjun$ ls -lh build/install/lib  
total 59696  
-rw-r--r-- 1 fangjun staff 25M Aug 14 14:09 libonnxruntime.1.17.1.dylib  
lrwxr-xr-x 1 fangjun staff 27B Aug 14 14:18 libonnxruntime.dylib -> libonnxruntime.  
→1.17.1.dylib  
-rwxr-xr-x 1 fangjun staff 3.9M Aug 15 15:01 libsherpa-onnx-c-api.dylib
```

If you download static libraries, please create a `build-static` directory inside the `sherpa-onnx` project directory and put the library files into `build-static/install/lib`. An example on my macOS is given below:

```
(py38) fangjuns-MacBook-Pro:sherpa-onnx fangjun$ pwd  
/Users/fangjun/open-source/sherpa-onnx  
(py38) fangjuns-MacBook-Pro:sherpa-onnx fangjun$ ls -lh build-static/install/lib  
total 138176  
-rw-r--r-- 1 fangjun staff 438K Aug 15 15:03 libespeak-ng.a  
-rw-r--r-- 1 fangjun staff 726K Aug 15 15:03 libkaldi-decoder-core.a  
-rw-r--r-- 1 fangjun staff 198K Aug 15 15:03 libkaldi-native-fbank-core.a  
-rw-r--r-- 1 fangjun staff 56M Aug 14 14:25 libonnxruntime.a  
-rw-r--r-- 1 fangjun staff 421K Aug 15 15:03 libpiper_phonemize.a  
-rw-r--r-- 1 fangjun staff 87K Aug 15 15:03 libsherpa-onnx-c-api.a  
-rw-r--r-- 1 fangjun staff 5.7M Aug 15 15:03 libsherpa-onnx-core.a  
-rw-r--r-- 1 fangjun staff 2.3M Aug 15 15:03 libsherpa-onnx-fst.a  
-rw-r--r-- 1 fangjun staff 30K Aug 15 15:03 libsherpa-onnx-fstfar.a  
-rw-r--r-- 1 fangjun staff 1.6M Aug 15 15:03 libsherpa-onnx-kaldifst-core.a  
-rw-r--r-- 1 fangjun staff 131K Aug 15 15:03 libsherpa-onnx-portaudio_static.a  
-rw-r--r-- 1 fangjun staff 147K Aug 15 15:03 libssentencepiece_core.a  
-rw-r--r-- 1 fangjun staff 197K Aug 15 15:03 libucd.a
```

Build the `generate_subtitles` project

Now you can start Lazarus and open `generate_subtitles.lpi`.

Click the menu Run -> Compile. It should be able to build the project without any errors.

Hint: Please ignore warnings, if there are any.

After building, you should find the following files inside the directory `generate_subtitles`:

macOS

Windows

Linux

```
(py38) fangjuns-MacBook-Pro:generate_subtitles fangjun$ pwd  
/Users/fangjun/open-source/sherpa-onnx/lazarus-examples/generate_subtitles  
(py38) fangjuns-MacBook-Pro:generate_subtitles fangjun$ ls -lh generate_subtitles.  
→generate_subtitles.app/  
-rwxr-xr-x 1 fangjun staff 25M Aug 15 20:44 generate_subtitles
```

(continues on next page)

(continued from previous page)

```
generate_subtitles.app/:
total 0
drwxr-xr-x 6 fangjun staff 192B Aug 14 23:01 Contents
```

```
fangjun@M-0LQSDCC2RV398 C:\Users\fangjun\open-source\sherpa-onnx\lazarus-examples\
└─generate_subtitles>dir generate_subtitles.exe
Volume in drive C is
Volume Serial Number is 8E17-A21F

Directory of C:\Users\fangjun\open-source\sherpa-onnx\lazarus-examples\generate_
└─subtitles

08/15/2024 09:39 PM      2,897,408 generate_subtitles.exe
    1 File(s)   2,897,408 bytes
    0 Dir(s)  38,524,039,168 bytes free
```

```
cd lazarus-examples/generate_subtitles
ls -lh generate_subtitles

-rwxr-xr-x 1 runner docker 3.1M Aug 16 03:37 generate_subtitles
```

Now you can start the generated executable `generate_subtitles` and you should get the screenshot like the one listed at the start of this section.

If you get any issues about `shared libraries not found`, please copy the shared library files from `build/install/lib` to the directory `lazarus-examples/generate_subtitles` or you can set the environment variable `DYLD_LIBRARY_PATH` (for macOS) and `LD_LIBRARY_PATH` (for Linux).

Download models

The generated executable expects that there are model files located in the same directory.

Download the VAD model

```
cd lazarus-examples/generate_subtitles
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx
```

and put `silero_vad.onnx` into `lazarus-examples/generate_subtitles`.

Hint: If you are using macOS, please put it into `lazarus-examples/generate_subtitles/generate_subtitles.app/Contents/Resources/`

Download an speech recognition model

The executable expects a non-streaming speech recognition model. Currently, we have supported the following types of models

- Whisper
- Zipformer transducer
- NeMo transducer
- SenseVoice
- Paraformer
- TeleSpeech CTC

You can download them from <https://github.com/k2-fsa/sherpa-onnx/releases/tag/asr-models>

Note that you have to rename the model files after downloading.

	Expected filenames
Whisper	<ul style="list-style-type: none">• tokens.txt• whisper-encoder.onnx• whisper-decoder.onnx
Zipformer transducer	<ul style="list-style-type: none">• tokens.txt• transducer-encoder.onnx• transducer-decoder.onnx• transducer-joiner.onnx
NeMo transducer	<ul style="list-style-type: none">• tokens.txt• nemo-transducer-encoder.onnx• nemo-transducer-decoder.onnx• nemo-transducer-joiner.onnx
SenseVoice	<ul style="list-style-type: none">• tokens.txt• sense-voice.onnx
Paraformer	<ul style="list-style-type: none">• tokens.txt• paraformer.onnx
TeleSpeech	<ul style="list-style-type: none">• tokens.txt• telespeech.onnx

We give several examples below.

1. Wisper

```
cd lazarus-examples/generate_subtitles
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-whisper-tiny.en.tar.bz2

tar xvf sherpa-onnx-whisper-tiny.en.tar.bz2
rm sherpa-onnx-whisper-tiny.en.tar.bz2

cd sherpa-onnx-whisper-tiny.en

mv -v tiny.en-encoder.int8.onnx ../whisper-encoder.onnx
mv -v tiny.en-decoder.int8.onnx ../whisper-decoder.onnx
mv -v tiny.en-tokens.txt ../tokens.txt

cd ..
rm -rf sherpa-onnx-whisper-tiny.en
```

You can replace `tiny.en` with other types of Whisper models, e.g., `tiny`, `base`, etc.

2. Zipformer transducer

We give two examples below for Zipformer transducer

Example 1

```
cd lazarus-examples/generate_subtitles
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-
˓→asr-zipformer-wenetspeech-20230615.tar.bz2
tar xvf icefall-asr-zipformer-wenetspeech-20230615.tar.bz2
rm icefall-asr-zipformer-wenetspeech-20230615.tar.bz2

cd icefall-asr-zipformer-wenetspeech-20230615

mv -v data/lang_char/tokens.txt ../

mv -v exp/encoder-epoch-12-avg-4.int8.onnx ../transducer-encoder.onnx
mv -v exp/decoder-epoch-12-avg-4.onnx ../transducer-decoder.onnx
mv -v exp/joiner-epoch-12-avg-4.int8.onnx ../transducer-joiner.onnx

cd ..
rm icefall-asr-zipformer-wenetspeech-20230615
```

Example 2

```
cd lazarus-examples/generate_subtitles
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-zipformer-ja-reazonspeech-2024-08-01.tar.bz2
tar xvf sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01.tar.bz2
rm sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01.tar.bz2

cd sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01
```

(continues on next page)

(continued from previous page)

```
mv ./tokens.txt ../  
mv encoder-epoch-99-avg-1.int8.onnx ../transducer-encoder.onnx  
mv decoder-epoch-99-avg-1.onnx ../transducer-decoder.onnx  
mv joiner-epoch-99-avg-1.int8.onnx ../transducer-joiner.onnx  
  
cd ../  
  
rm -rf sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01
```

3. NeMo transducer

```
cd lazarus-examples/generate_subtitles
curl -SL -o https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-nemo-fast-conformer-transducer-be-de-en-es-fr-hr-it-pl-ru-uk-20k.tar.bz2
tar xvf sherpa-onnx-nemo-fast-conformer-transducer-be-de-en-es-fr-hr-it-pl-ru-uk-20k.tar.
˓→bz2
rm sherpa-onnx-nemo-fast-conformer-transducer-be-de-en-es-fr-hr-it-pl-ru-uk-20k.tar.bz2

cd sherpa-onnx-nemo-fast-conformer-transducer-be-de-en-es-fr-hr-it-pl-ru-uk-20k

mv tokens.txt ..
mv encoder.onnx ../nemo-transducer-encoder.onnx
mv decoder.onnx ../nemo-transducer-decoder.onnx
mv joiner.onnx ../nemo-transducer-joiner.onnx

cd ..
rm -rf sherpa-onnx-nemo-fast-conformer-transducer-be-de-en-es-fr-hr-it-pl-ru-uk-20k
```

4. SenseVoice

```
cd lazarus-examples/generate_subtitles
curl -SL -0 https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

cd sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17
mv tokens.txt ../
mv model.int8.onnx ../

cd ../
rm -rf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17
```

5. Paraformer

```
cd lazarus-examples/generate_subtitles
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-paraformer-zh-2023-09-14.tar.bz2

tar xvf sherpa-onnx-paraformer-zh-2023-09-14.tar.bz2
rm sherpa-onnx-paraformer-zh-2023-09-14.tar.bz2

cd sherpa-onnx-paraformer-zh-2023-09-14

mv tokens.txt ../
mv model.int8.onnx ../

cd ../
rm -rf sherpa-onnx-paraformer-zh-2023-09-14
```

6. TeleSpeech

```
cd lazarus-examples/generate_subtitles
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-telespeech-ctc-int8-zh-2024-06-04.tar.bz2

tar xvf sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04.tar.bz2
rm sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04.tar.bz2

cd sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04

mv tokens.txt ../
mv model.int8.onnx ../

cd ../
rm -rf sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04
```

For the more curious

If you want to find out how we generate the APPs in <<https://k2-fsa.github.io/sherpa/onnx/lazarus/download-generated-subtitles.html>>, please have a look at

- <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/lazarus/generate-subtitles.py>
- <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/lazarus/build-generate-subtitles.sh.in>
- <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/lazarus.yaml>

8.14 WebAssembly

In this section, we describe how to build `sherpa-onnx` for WebAssembly so that you can run real-time speech recognition with [WebAssembly](#).

Please follow the steps below to build and run `sherpa-onnx` for [WebAssembly](#).

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows or you don't want to setup your local environment to build WebAssembly support, please use the above colab notebook.

8.14.1 Install Emscripten

We need to compile the C/C++ files in `sherpa-onnx` with the help of `emscripten`.

Please refer to https://emscripten.org/docs/getting_started/downloads for detailed installation instructions.

The following is an example to show you how to install it on Linux/macOS.

```
git clone https://github.com/emscripten-core/emsdk.git
cd emsdk
git pull
./emsdk install latest
./emsdk activate latest
source ./emsdk_env.sh
```

To check that you have installed `emscripten` successfully, please run:

```
emcc -v
```

The above command should print something like below:

```
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
(e967e20b4727956a30592165a3c1cde5c67fa0a8)
shared:INFO: (Emscripten: Running sanity checks)
(py38) fangjuns-MacBook-Pro:open-source fangjun$ emcc -v
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
(e967e20b4727956a30592165a3c1cde5c67fa0a8)
clang version 18.0.0 (https://github.com/llvm/llvm-project
(a54545ba6514802178cf7cf1c1dd9f7efbf3cde7)
Target: wasm32-unknown-emscripten
Thread model: posix
InstalledDir: /Users/fangjun/open-source/emsdk/upstream/bin
```

Congratulations! You have successfully installed `emscripten`.

8.14.2 Build

After installing emscripten, we can build `sherpa-onnx` for WebAssembly now.

Please use the following command to build it:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

cd wasm/asr/assets

wget -q https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
rm sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2

# Note it is not an error that we rename encoder.int8.onnx to encoder.onnx

mv sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-99-avg-1.
˓→int8.onnx encoder.onnx
mv sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-99-avg-1.
˓→onnx decoder.onnx
mv sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-avg-1.int8.
˓→onnx joiner.onnx
mv sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt ./
rm -rf sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/

cd ../../..
./build-wasm-simd-asr.sh
```

Hint: You can visit <https://github.com/k2-fsa/sherpa-onnx/releases/tag/asr-models> to download a different model.

If you want to use a streaming Paraformer model, please see <https://github.com/k2-fsa/sherpa-onnx/blob/master/wasm/asr/assets/README.md#paraformer>

After building, you should see the following output:

```
Install the project...
-- Install configuration: "Release"
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/sherpa-onnx-wasm-asr-main.js
-- Up-to-date: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/sherpa-onnx-wasm-asr-main.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/index.html
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/sherpa-onnx.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/app.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/sherpa-onnx-wasm-asr-main.wasm
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-asr/install/bin/
˓→wasm/asr/sherpa-onnx-wasm-asr-main.data
```

(continues on next page)

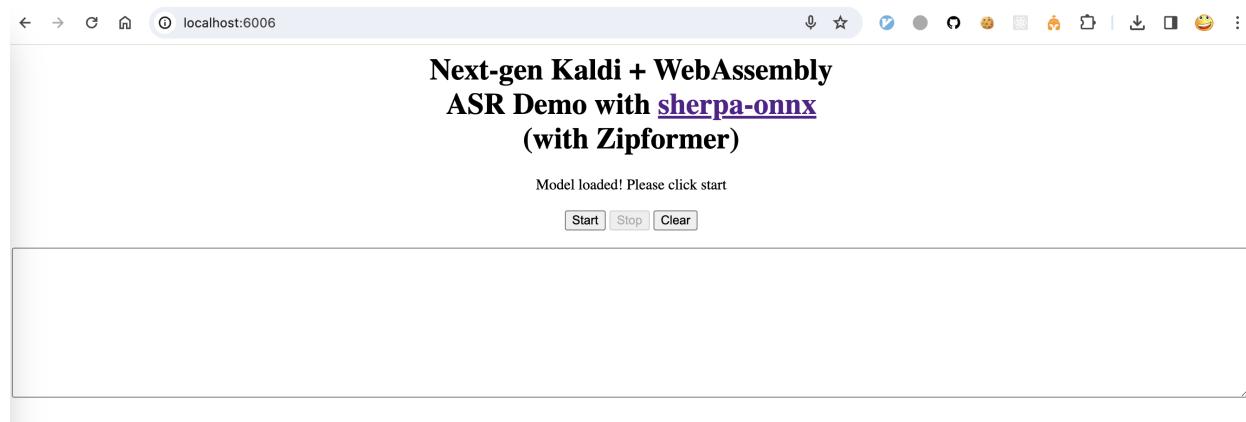
(continued from previous page)

```
+ ls -lh install/bin/wasm/asr
total 440080
-rw-r--r-- 1 fangjun staff 9.0K Feb 23 17:39 app.js
-rw-r--r-- 1 fangjun staff 978B Feb 23 17:39 index.html
-rw-r--r-- 1 fangjun staff 199M Feb 23 18:34 sherpa-onnx-wasm-asr-main.data
-rw-r--r-- 1 fangjun staff 90K Feb 23 18:38 sherpa-onnx-wasm-asr-main.js
-rw-r--r-- 1 fangjun staff 10M Feb 23 18:38 sherpa-onnx-wasm-asr-main.wasm
-rw-r--r-- 1 fangjun staff 9.1K Feb 23 17:39 sherpa-onnx.js
```

Now you can use the following command to run it:

```
cd build-wasm-simd-asr/install/bin/wasm/asr
python3 -m http.server 6006
```

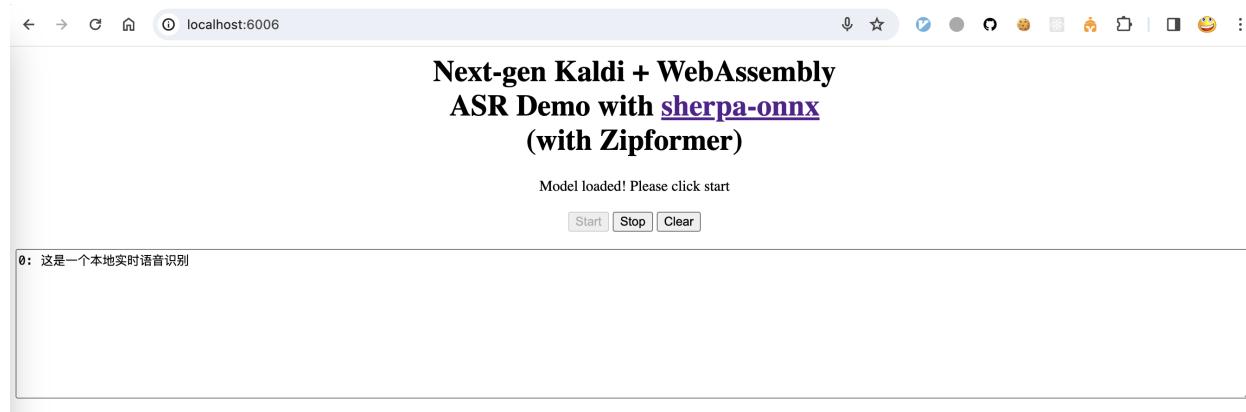
Start your browser and visit <http://localhost:6006>; you should see the following page:



Now click start and speak! You should see the recognition results in the text box.

Warning: We are using a bilingual model (Chinese + English) in the above example, which means you can only speak Chinese or English in this case.

A screenshot is given below:



Congratulations! You have successfully run real-time speech recognition with [WebAssembly](#) in your browser.

8.14.3 Huggingface Spaces (WebAssembly)


We provide four Huggingface spaces so that you can try real-time speech recognition with WebAssembly in your browser.

English only (Zipformer)

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-onnx-en>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-onnx-en/summary>



Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-en-asr-zipformer.yaml>

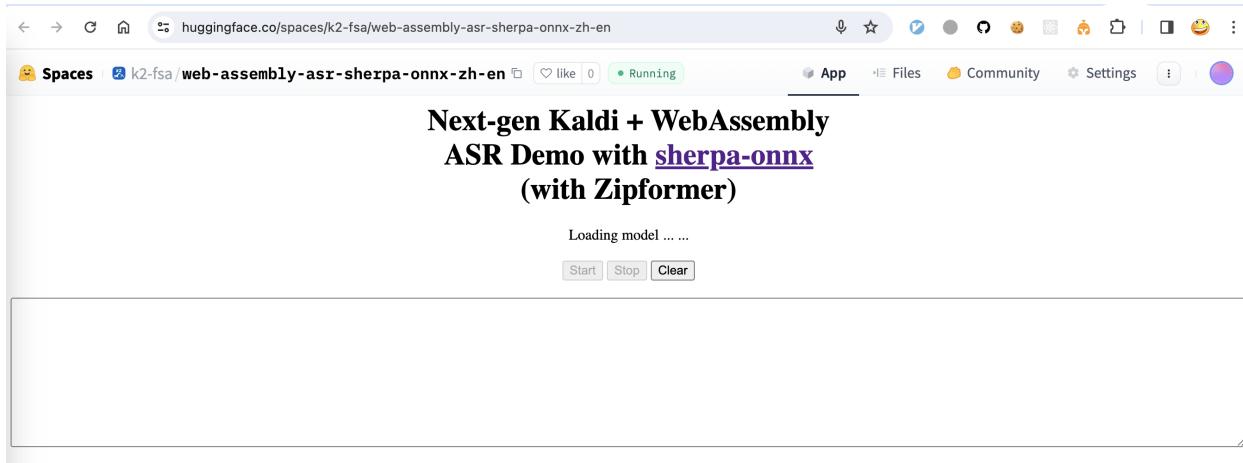
Chinese + English (Zipformer)

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-onnx-zh-en>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-onnx-zh-en/summary>

Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-zh-en-asr-zipformer.yaml>




Chinese + English (Paraformer)

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-onnx-zh-en-paraformer>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-onnx-zh-en-paraformer/summary>



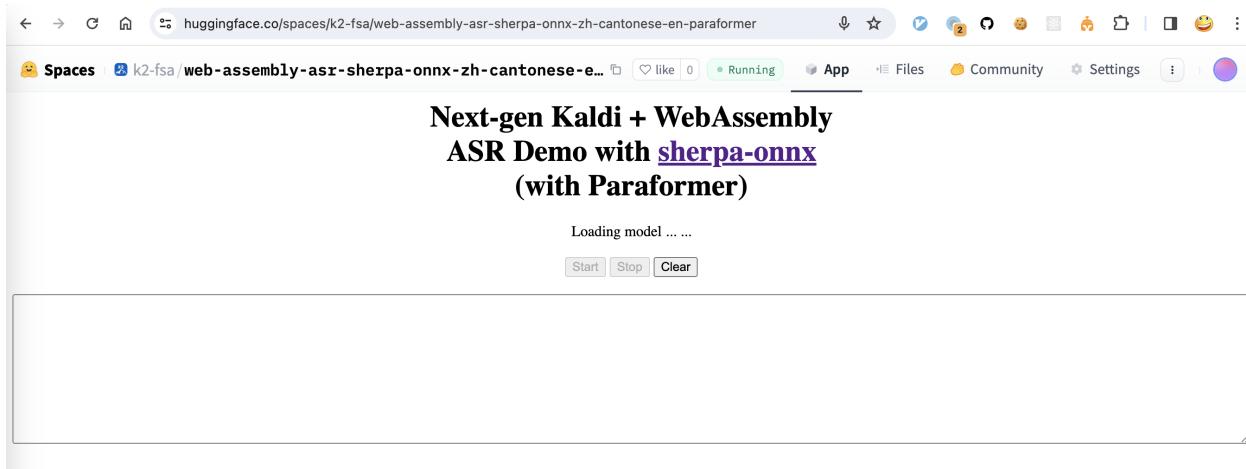
Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-zh-en-asr-paraformer.yaml>

Chinese + English + Cantonese (Paraformer)

<https://huggingface.co/spaces/k2-fsa/web-assembly-asr-sherpa-onnx-zh-cantonese-en-paraformer>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-asr-sherpa-onnx-zh-cantonese-en-paraformer/>
summary



Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-zh-cantonese-en-asr-paraformer.yaml>

8.15 Android

In this section, we describe how to build an Android app for with `sherpa-onnx`.

Hint: For real-time speech recognition, it does not need to access the Internet. Everything is processed locally on your phone.

8.15.1 Pre-built APKs

Links for pre-built APKs can be found in the following table:

Hint: It runs locally, without internet connection.

	URL
Streaming speech recognition	https://k2-fsa.github.io/sherpa/onnx/android/apk.html
Text-to-speech engine	https://k2-fsa.github.io/sherpa/onnx/tts/apk-engine.html
Voice activity detection (VAD)	https://k2-fsa.github.io/sherpa/onnx/vad/apk.html
VAD + non-streaming speech recognition	https://k2-fsa.github.io/sherpa/onnx/vad/apk-asr.html
Two-pass speech recognition	https://k2-fsa.github.io/sherpa/onnx/android/apk-2pass.html
Audio tagging	https://k2-fsa.github.io/sherpa/onnx/audio-tagging/apk.html
Audio tagging (WearOS)	https://k2-fsa.github.io/sherpa/onnx/audio-tagging/apk-wearos.html
Speaker identification	https://k2-fsa.github.io/sherpa/onnx/speaker-identification/apk.html
Spoken language identification	https://k2-fsa.github.io/sherpa/onnx/spoken-language-identification/apk.html
Keyword spotting	https://k2-fsa.github.io/sherpa/onnx/kws/apk.html

8.15.2 Build sherpa-onnéx for Android

You can use this section for both **speech-to-text** (STT, ASR) and **text-to-speech** (TTS).

Hint: The build scripts mentioned in this section run on both Linux and macOS.

If you are using Windows or if you don't want to build the shared libraries, you can download pre-built shared libraries by visiting the release page <https://github.com/k2-fsa/sherpa-onnéx/releases/>

For instance, for the release v1.10.19, you can visit <https://github.com/k2-fsa/sherpa-onnéx/releases/tag/v1.10.19> and download the file `sherpa-onnéx-v1.10.19-android.tar.bz2` using the following command:

```
wget https://github.com/k2-fsa/sherpa-onnéx/releases/download/v1.10.19/sherpa-onnéx-v1.10.19-android.tar.bz2
```

Please always use the latest release.

Hint: This section is originally written for speech-to-text. However, it is also applicable to other folders in <https://github.com/k2-fsa/sherpa-onnéx/tree/master/android>.

For instance, you can replace `SherpaOnnx` in this section with

- `SherpaOnnx2Pass`
- `SherpaOnnxTts` (this is for text-to-speech)
- `SherpaOnnxTtsEngine` (this is for text-to-speech)
- `SherpaOnnxVad`
- `SherpaOnnxVadAsr`
- `SherpaOnnxSpeakerIdentification`
- `SherpaOnnxAudioTagging`
- `SherpaOnnxAudioTaggingWearOs`

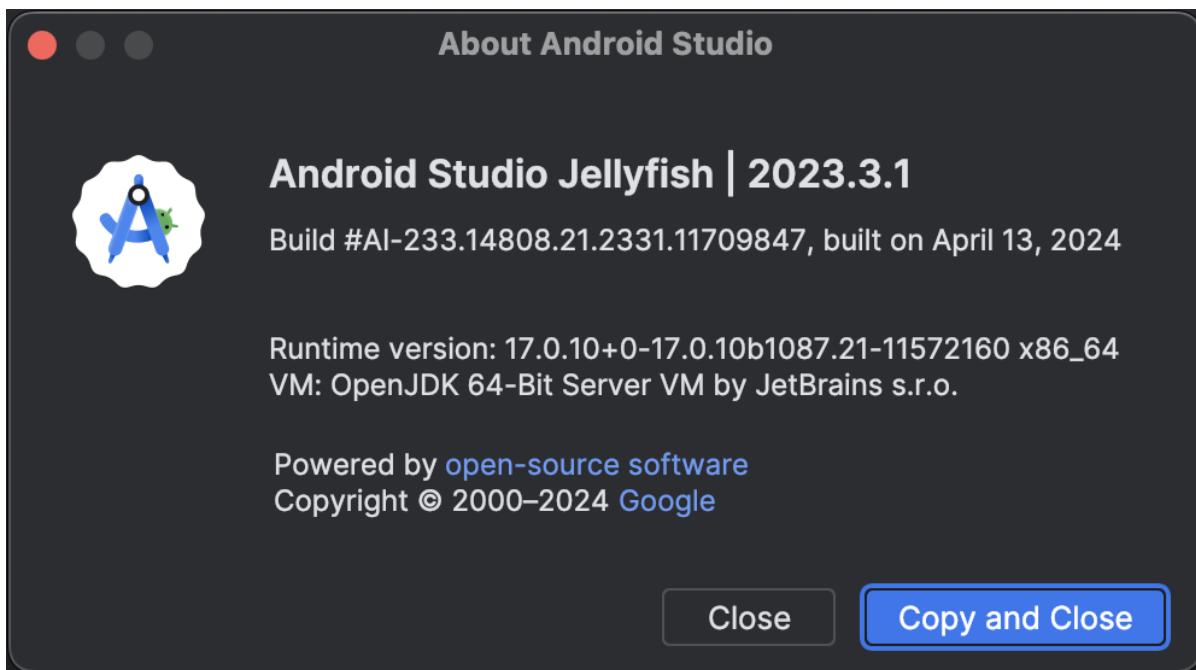
Install Android Studio

The first step is to download and install Android Studio.

Please refer to <https://developer.android.com/studio> for how to install Android Studio.

Hint: Any recent version of Android Studio should work fine. Also, you can use the default settings of Android Studio during installation.

For reference, we post the version we are using below:



Download sherpa-onnx

Next, download the source code of sherpa-onnx:

```
git clone https://github.com/k2-fsa/sherpa-onnx
```

Install NDK

Step 1, start Android Studio.

Step 2, Open sherpa-onnx/android/SherpaOnnx.

Step 3, Select Tools -> SDK Manager.

Step 4, Install NDK.

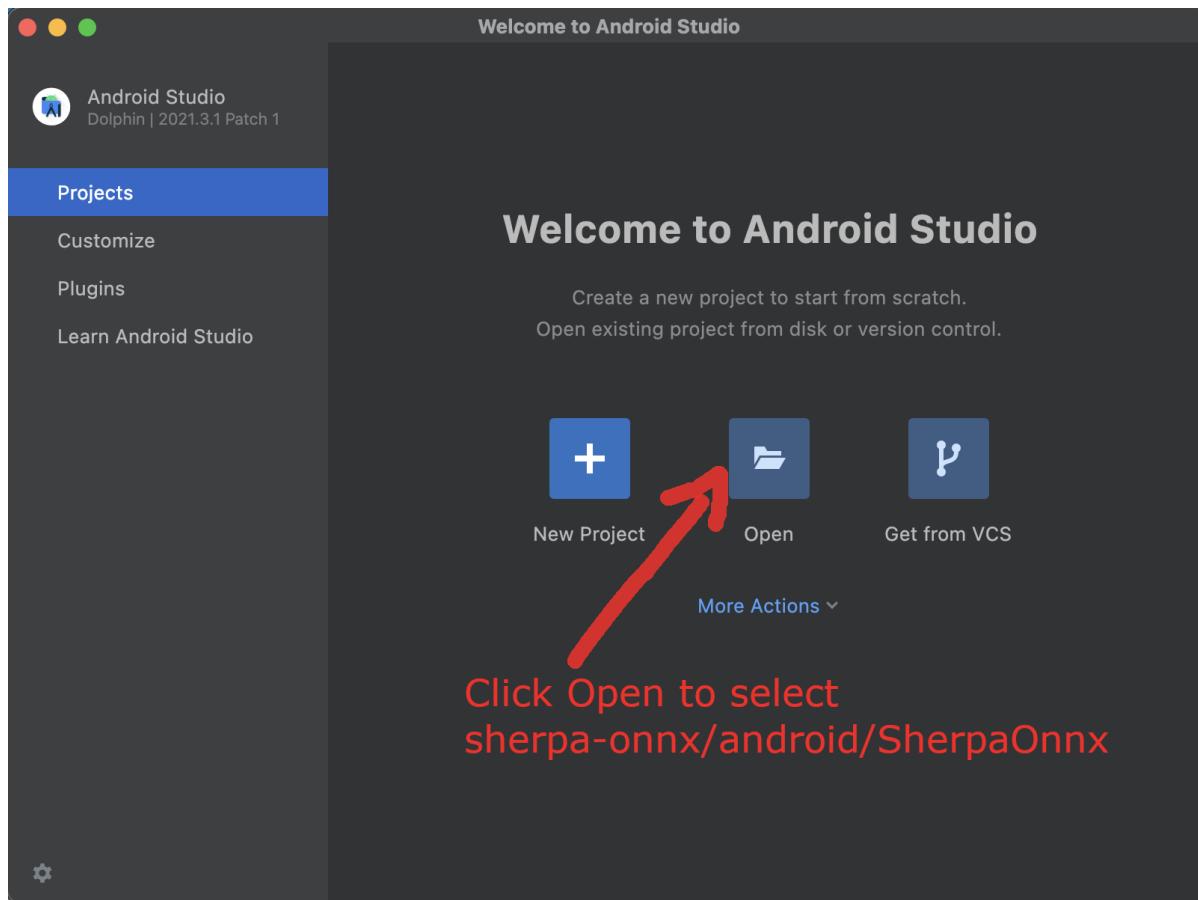


Fig. 8.4: Step 1: Click Open to select `sherpa-onnx/android/SherpaOnnx`

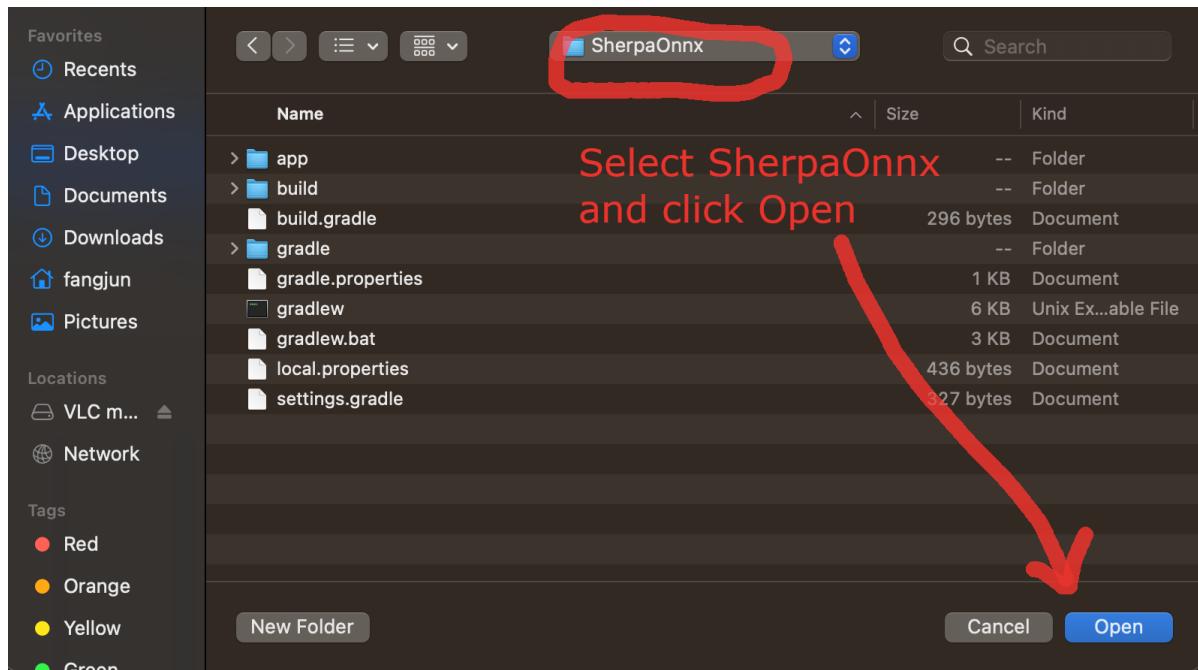


Fig. 8.5: Step 2: Open SherpaOnnx.

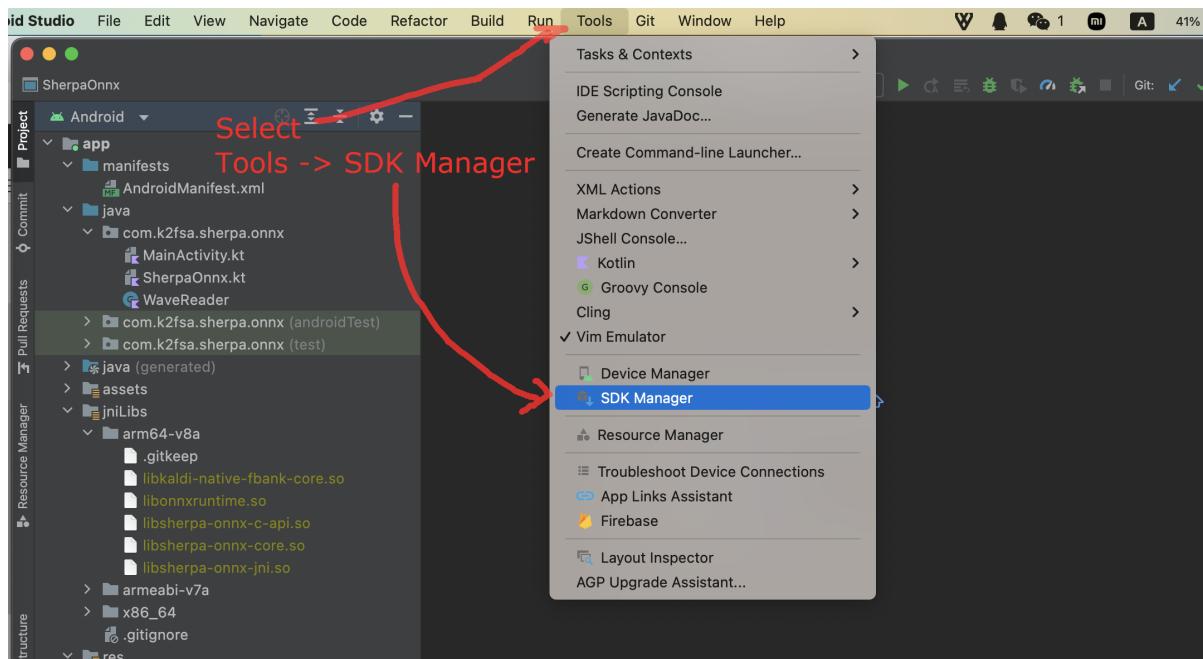


Fig. 8.6: Step 3: Select Tools -> SDK Manager.

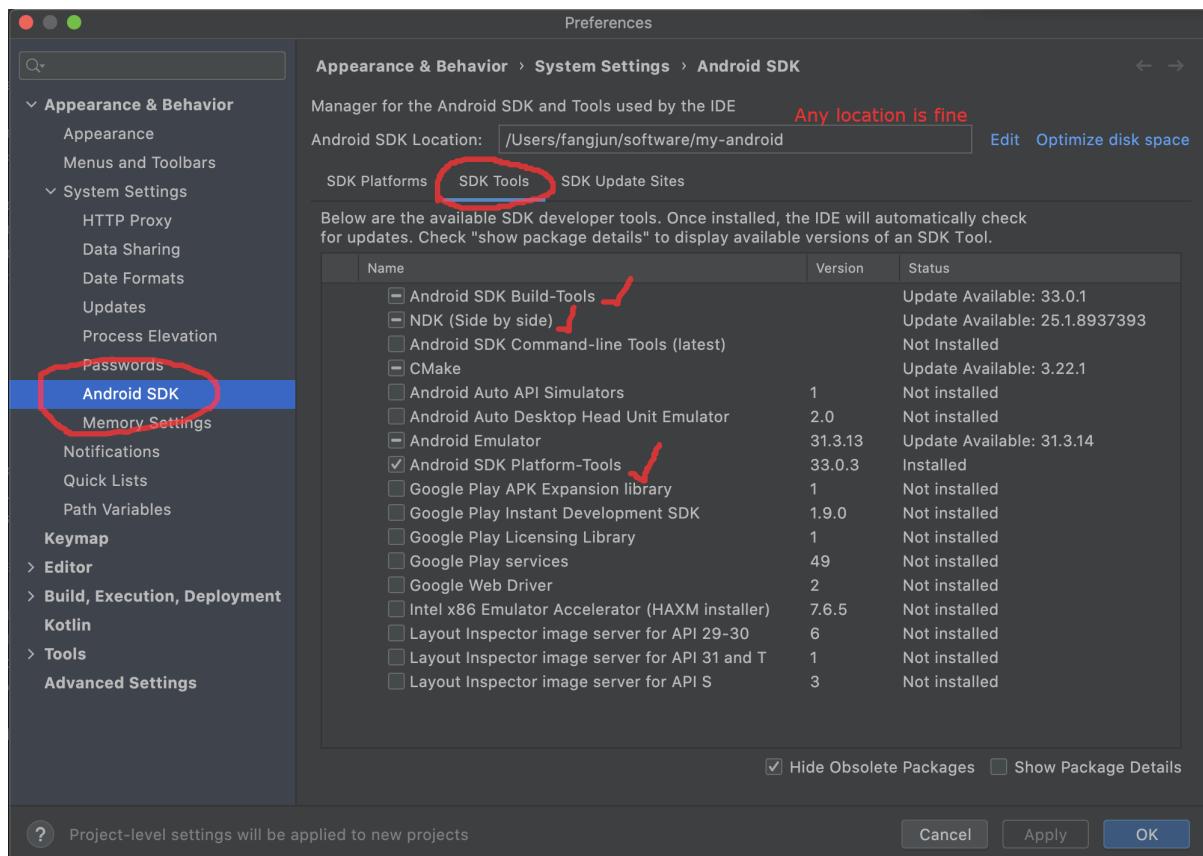


Fig. 8.7: Step 4: Install NDK.

In the following, we assume Android SDK location was set to /Users/fangjun/software/my-android. You can change it accordingly below.

After installing NDK, you can find it in

```
/Users/fangjun/software/my-android/ndk/22.1.7171670
```

Warning: If you selected a different version of NDK, please replace 22.1.7171670 accordingly.

Next, let us set the environment variable ANDROID_NDK for later use.

```
export ANDROID_NDK=/Users/fangjun/software/my-android/ndk/22.1.7171670
```

Note: Note from <https://github.com/Tencent/ncnn/wiki/how-to-build#build-for-android>

(Important) remove the hardcoded debug flag in Android NDK to fix the android-ndk issue: <https://github.com/android/ndk/issues/243>

1. open \$ANDROID_NDK/build/cmake/android.toolchain.cmake for ndk < r23 or \$ANDROID_NDK/build/cmake/android-legacy.toolchain.cmake for ndk >= r23

2. delete the line containing “-g”

```
list(APPEND ANDROID_COMPILER_FLAGS  
-g  
-DANDROID)
```

Build sherpa-onnx (C++ code)

After installing NDK, it is time to build the C++ code of sherpa-onnx.

In the following, we show how to build `sherpa-onnx` for the following Android ABIs:

- arm64-v8a
- armv7-eabi
- x86_64
- x86

Caution: You only need to select one and only one ABI. `arm64-v8a` is probably the most common one.

If you want to test the app on an emulator, you probably need `x86_64`.

Hint: Building scripts for this section are for macOS and Linux. If you are using Windows or if you don't want to build the shared libraries by yourself, you can download pre-compiled shared libraries for this section by visiting

<https://github.com/k2-fsa/sherpa-onnx/releases>

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows or you don't want to setup your local environment to build the C++ libraries, please use the above colab notebook.

Build for arm64-v8a

```
cd sherpa-onnx # Go to the root repo
./build-android-arm64-v8a.sh
```

After building, you will find the following shared libraries:

```
ls -lh build-android-arm64-v8a/install/lib/
-rw-r--r-- 1 fangjun staff 15M Jul 28 12:54 libonnxruntime.so
-rwxr-xr-x 1 fangjun staff 3.7M Jul 28 12:54 libsherpa-onnx-jni.so
```

Please copy them to android/SherpaOnnx/app/src/main/jniLibs/arm64-v8a/:

```
cp build-android-arm64-v8a/install/lib/lib*.so android/SherpaOnnx/app/src/main/jniLibs/
↳arm64-v8a/
```

You should see the following screen shot after running the above copy cp command.

Build for armv7-eabi

```
cd sherpa-onnx # Go to the root repo
./build-android-armv7-eabi.sh
```

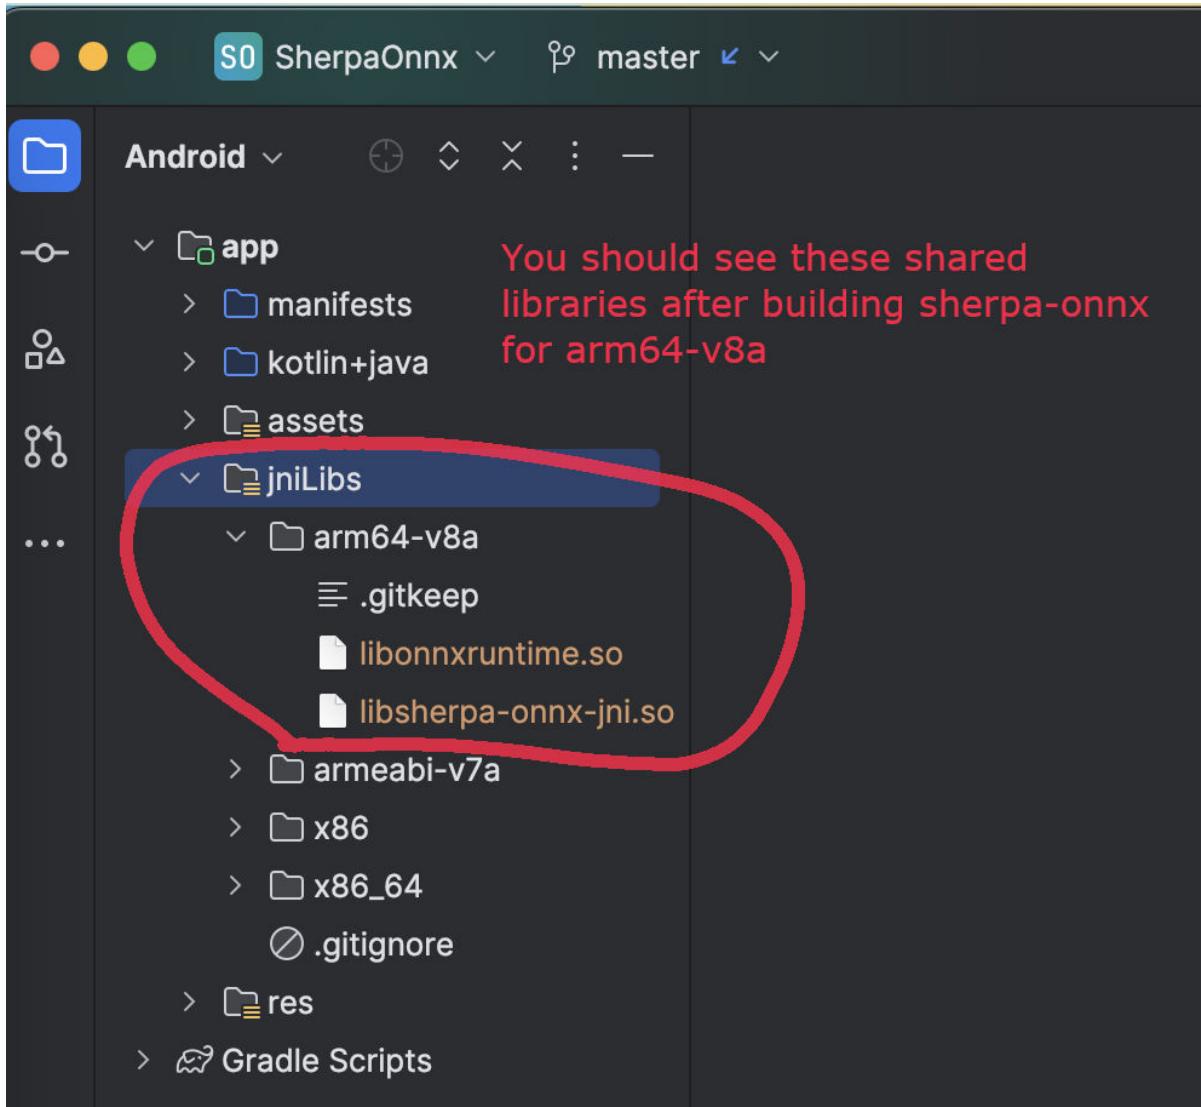
After building, you will find the following shared libraries:

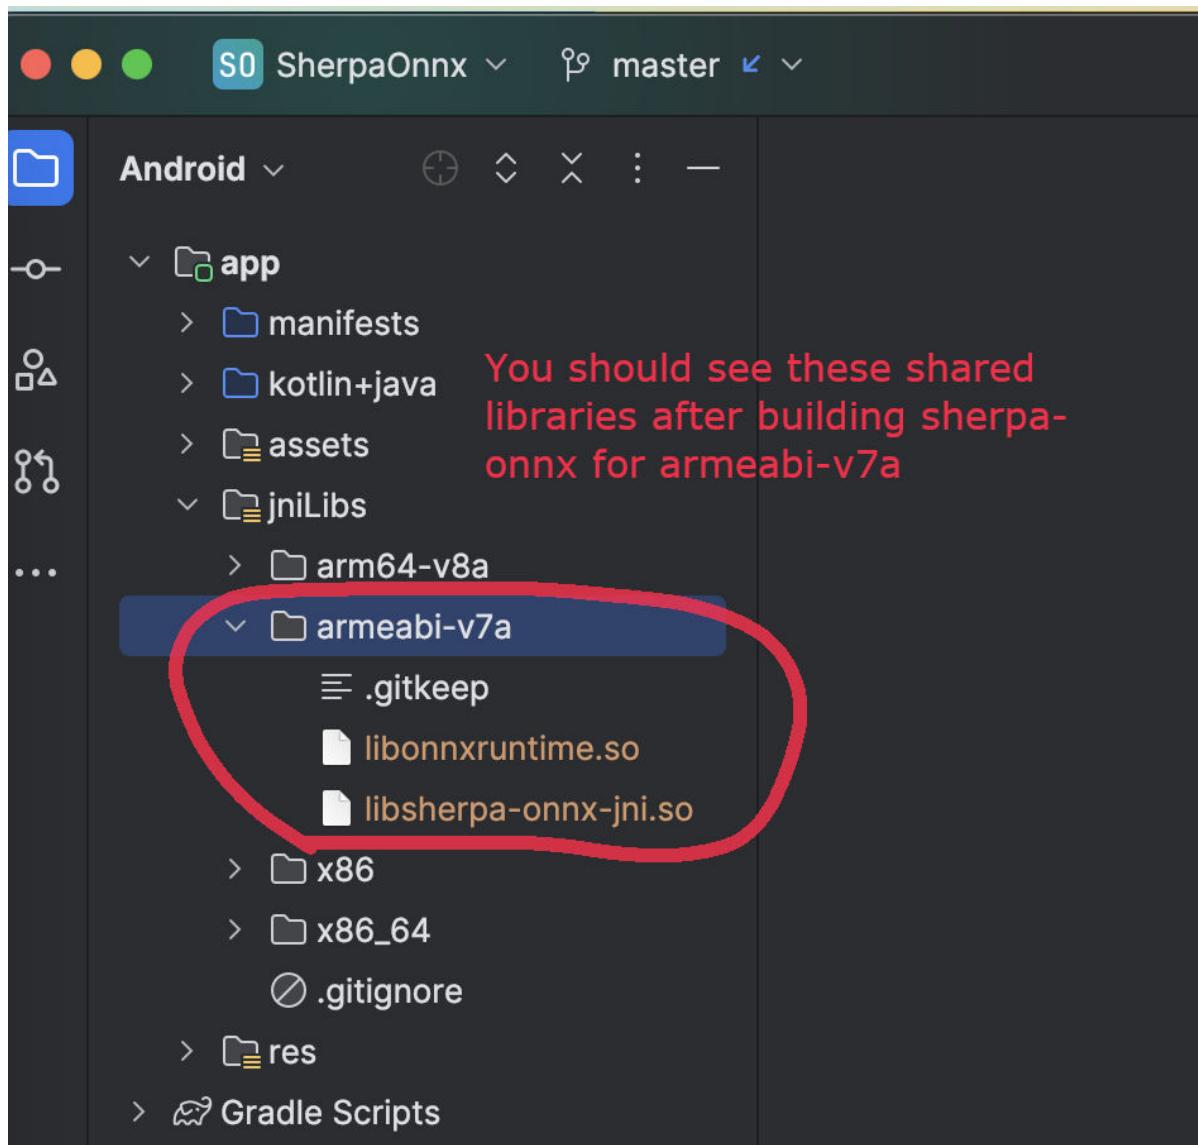
```
ls -lh build-android-armv7-eabi/install/lib/
-rw-r--r-- 1 fangjun staff 10M Jul 28 13:18 libonnxruntime.so
-rwxr-xr-x 1 fangjun staff 2.1M Jul 28 13:18 libsherpa-onnx-jni.so
```

Please copy them to android/SherpaOnnx/app/src/main/jniLibs/armeabi-v7a:

```
cp build-android-armv7-eabi/install/lib/lib*.so android/SherpaOnnx/app/src/main/jniLibs/
↳armeabi-v7a/
```

You should see the following screen shot after running the above copy cp command.





Build for x86_64

```
cd sherpa-onnx # Go to the root repo  
./build-android-x86-64.sh
```

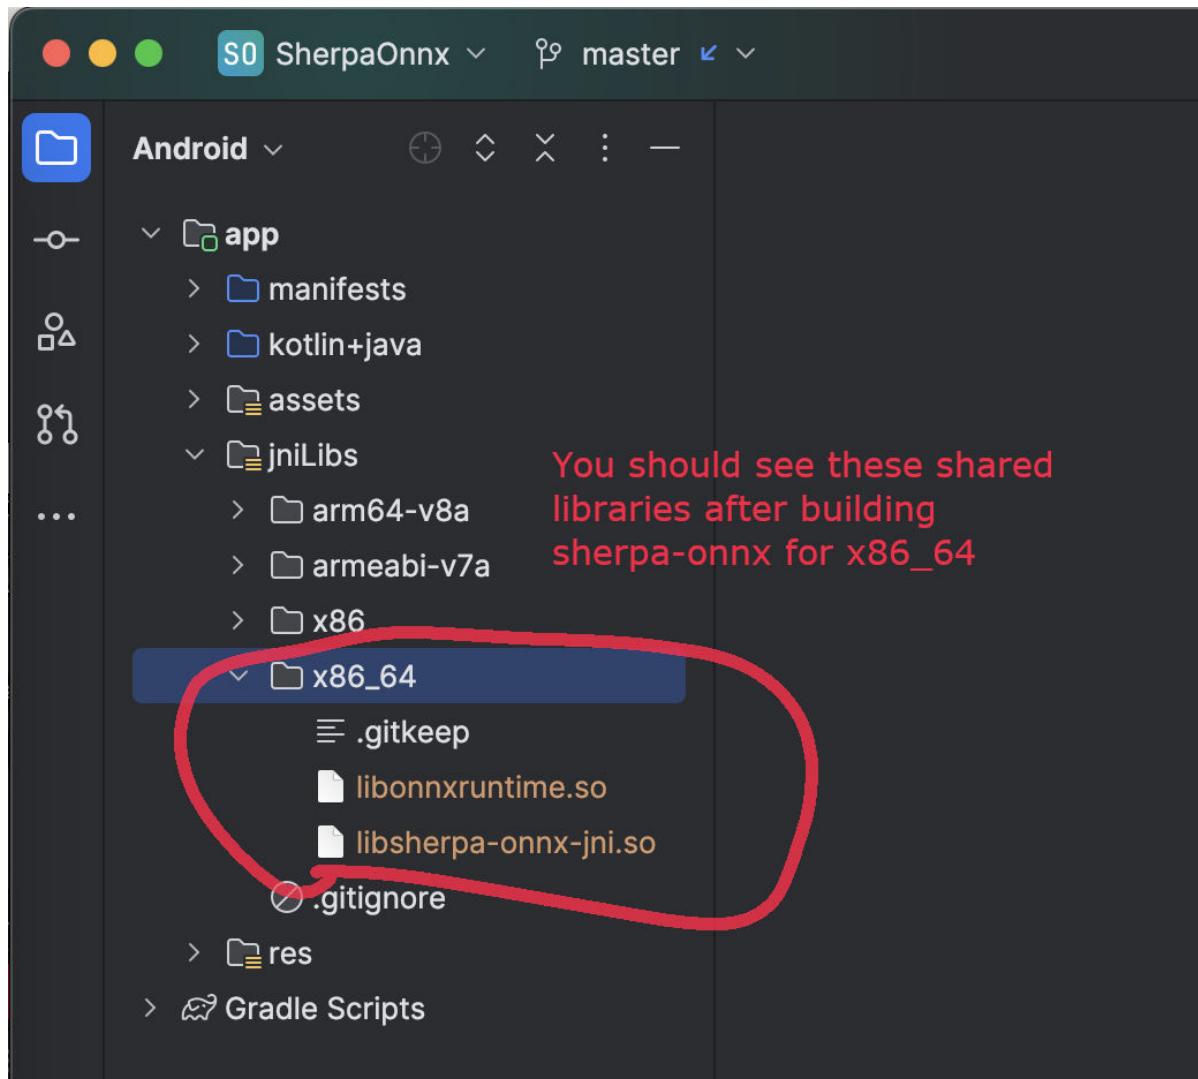
After building, you will find the following shared libraries:

```
ls -lh build-android-x86-64/install/lib/  
  
-rw-r--r-- 1 fangjun staff 17M Jul 28 13:26 libonnxruntime.so  
-rwxr-xr-x 1 fangjun staff 4.0M Jul 28 13:26 libsherpa-onnx-jni.so
```

Please copy them to android/SherpaOnnx/app/src/main/jniLibs/x86_64/:

```
cp build-android-x86-64/install/lib/lib*.so android/SherpaOnnx/app/src/main/jniLibs/x86_64/
```

You should see the following screen shot after running the above copy cp command.



Build for x86

```
cd sherpa-onnx # Go to the root repo
./build-android-x86.sh
```


After building, you will find the following shared libraries:

```
ls -lh build-android-x86/install/lib/
-rw-r--r-- 1 fangjun staff 17M Jul 28 13:28 libonnxruntime.so
-rwxr-xr-x 1 fangjun staff 3.9M Jul 28 13:28 libsherpa-onnx-jni.so
```

Please copy them to android/SherpaOnnx/app/src/main/jniLibs/x86/:

```
cp build-android-x86/install/lib/lib*.so android/SherpaOnnx/app/src/main/jniLibs/x86/
```

You should see the following screen shot after running the above copy cp command.



Download pre-trained models

Please read [Pre-trained models](#) for all available pre-trained models.

In the following, we use a pre-trained model [csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20](https://huggingface.co/csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20) (*Bilingual, Chinese + English*), which supports both Chinese and English.

Hint: The model is trained using `icefall` and the original torchscript model is from <https://huggingface.co/pfluo/k2fsa-zipformer-chinese-english-mixed>.

Use the following command to download the pre-trained model and place it into `android/SherpaOnnx/app/src/main/assets/`:

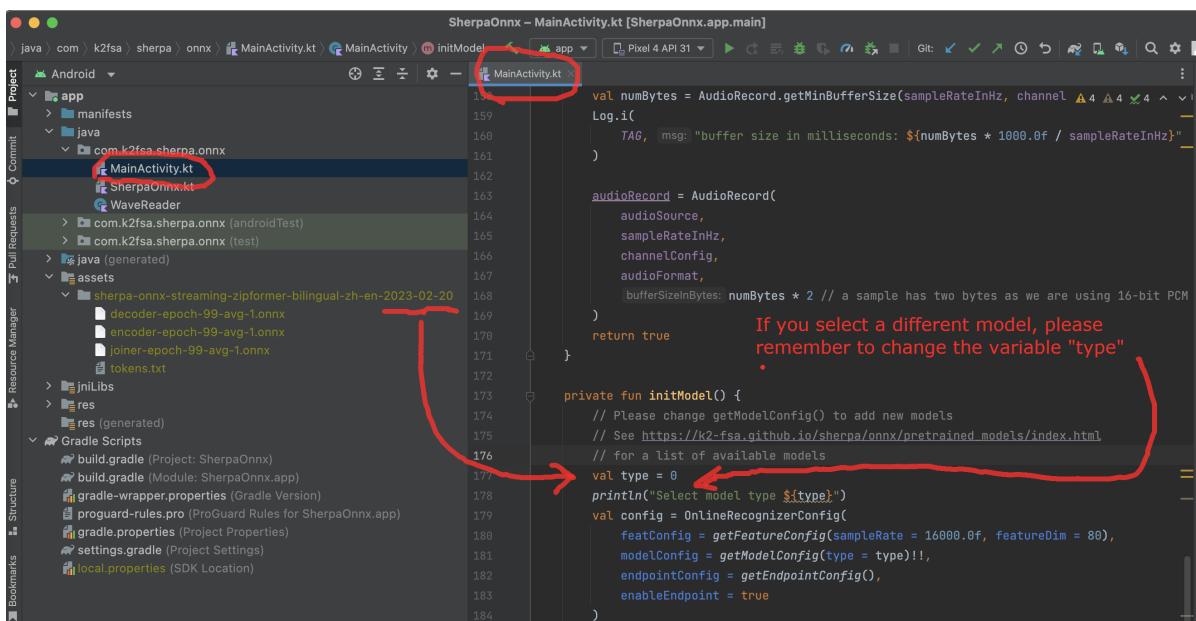
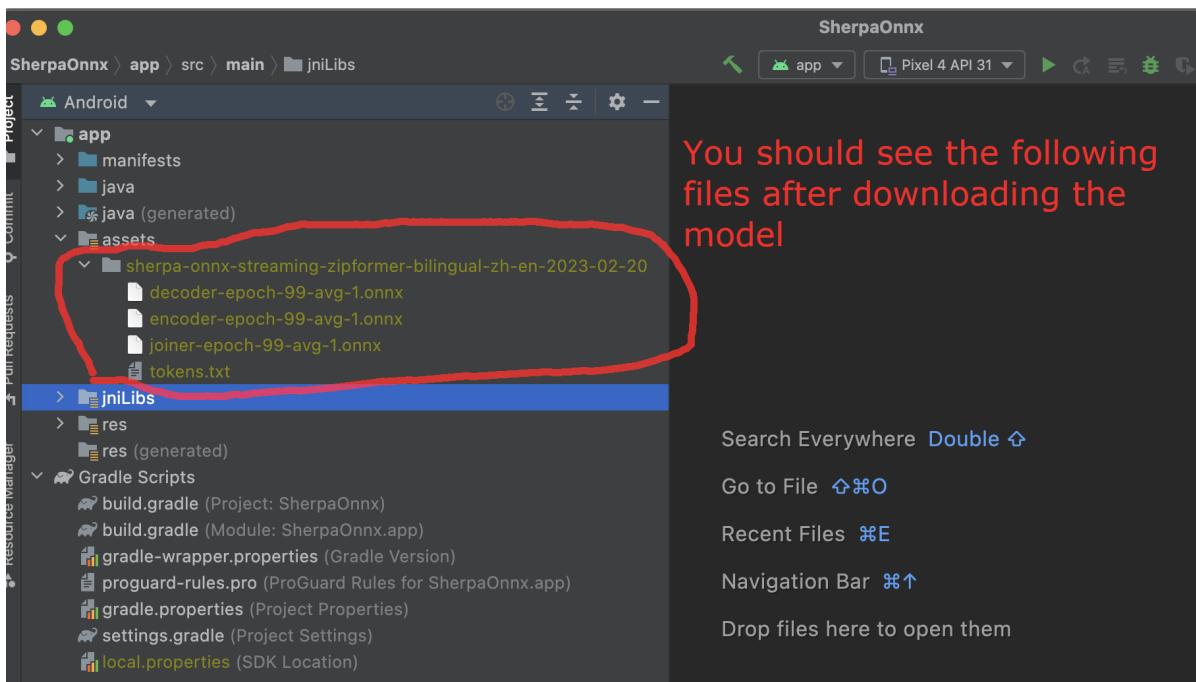
```
cd android/SherpaOnnx/app/src/main/assets/  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
→streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2  
  
tar xvf sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2  
rm sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2  
  
cd sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20  
  
# Now, remove extra files to reduce the file size of the generated apk  
rm -rf test_wavs  
rm -f *.sh README.md  
rm -f bpe.model  
  
rm -f encoder-epoch-99-avg-1.int8.onnx  
rm -f joiner-epoch-99-avg-1.int8.onnx  
rm -f decoder-epoch-99-avg-1.int8.onnx  
rm -f bpe.vocab
```

In the end, you should have the following files:

```
ls -lh  
  
-rw-r--r--@ 1 fangjun staff 13M Jul 28 13:51 decoder-epoch-99-avg-1.onnx  
-rw-r--r--@ 1 fangjun staff 315M Jul 28 13:51 encoder-epoch-99-avg-1.onnx  
-rw-r--r--@ 1 fangjun staff 12M Jul 28 13:51 joiner-epoch-99-avg-1.onnx  
-rw-r--r--@ 1 fangjun staff 55K Nov 21 2023 tokens.txt
```

You should see the following screen shot after downloading the pre-trained model:

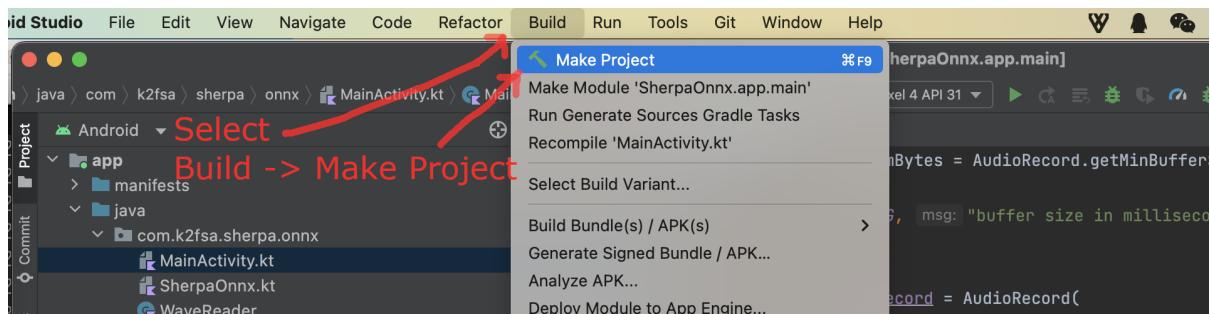
Hint: If you select a different pre-trained model, make sure that you also change the corresponding code listed in the following screen shot:



Generate APK

Finally, it is time to build `sherpa-onnx` to generate an APK package.

Select Build -> Make Project, as shown in the following screen shot.



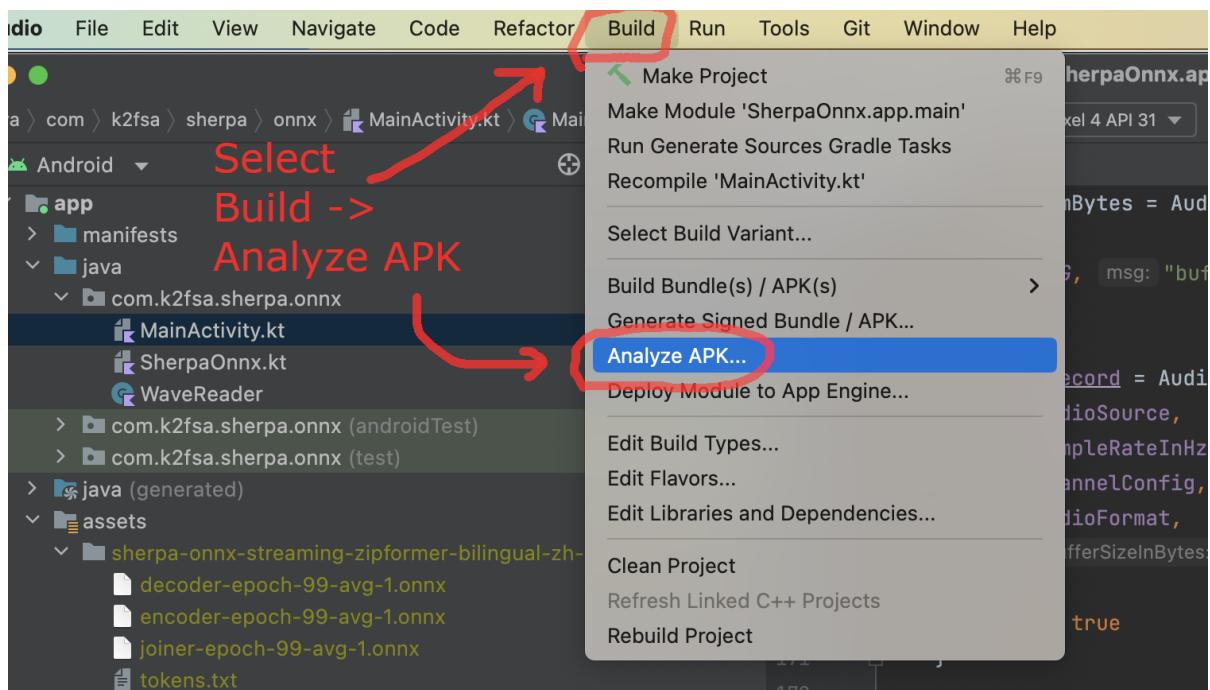
You can find the generated APK in `android/SherpaOnnx/app/build/outputs/apk/debug/app-debug.apk`:

```
ls -lh android/SherpaOnnx/app/build/outputs/apk/debug/app-debug.apk  
-rw-r--r--@ 1 fangjun staff 329M Jul 28 13:56 android/SherpaOnnx/app/build/outputs/  
apk/debug/app-debug.apk
```

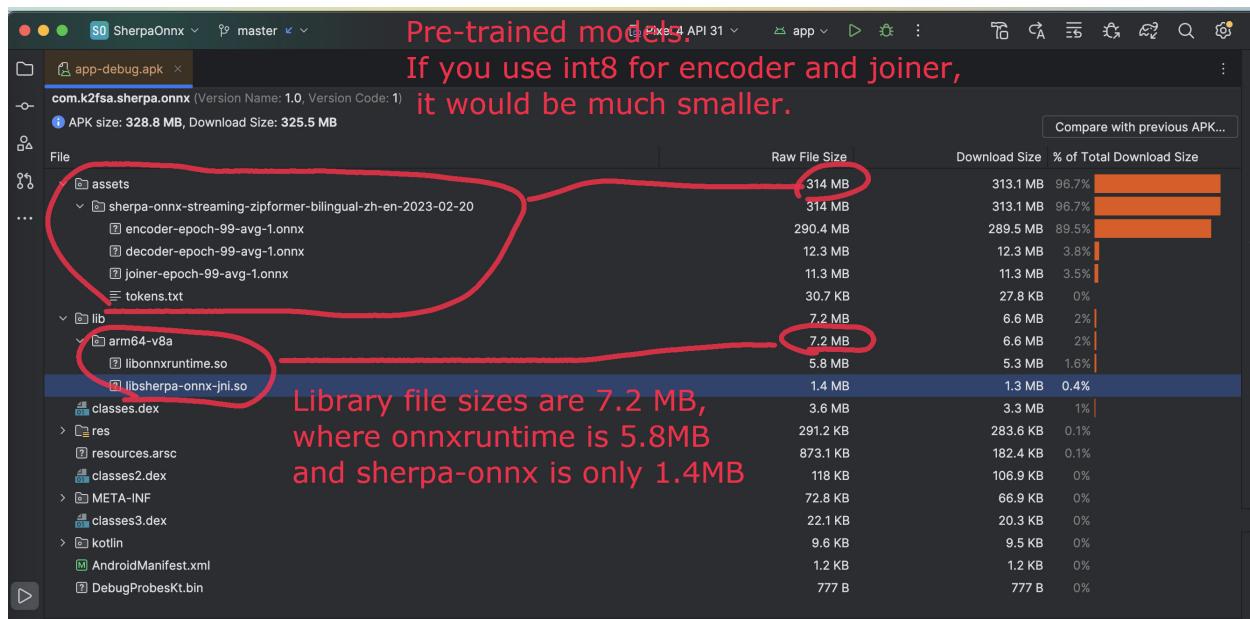
Congratulations! You have successfully built an APK for Android.

Read below to learn more.

Analyze the APK



Select Build -> Analyze APK ... in the above screen shot, in the popped-up dialog select the generated APK `app-debug.apk`, and you will see the following screen shot:



You can see from the above screen shot that most part of the APK is occupied by the pre-trained model, while the runtime, including the shared libraries, is only 7.2 MB.

Caution: You can see that `libonnxruntime.so` alone occupies 5.8MB out of 7.2MB.

We use a so-called Full build instead of Mobile build, so the file size of the library is somewhat a bit larger.

`libonnxruntime.so` is downloaded from

<https://mvnrepository.com/artifact/com.microsoft.onnxruntime/onnxruntime-android/1.17.1>

Please refer to <https://onnxruntime.ai/docs/build/custom.html> for a custom build to reduce the file size of `libonnxruntime.so`.

Note that we are constantly updating the version of `onnxruntime`. By the time you are reading this section, we may be using the latest version of `onnxruntime`.

Hint: We recommend you to use `sherpa-ncnn`. Please see [Analyze the APK for sherpa-ncnn](#). The total runtime of `sherpa-ncnn` is only 1.6 MB, which is much smaller than `sherpa-onnx`.

8.16 iOS

In this section, we describe how to build an iOS app for real-time speech recognition with `sherpa-onnx` and run it within a simulator on your Mac, run it on your iPhone or iPad.

Hint: During speech recognition, it does not need to access the Internet. Everything is processed locally on your device.

8.16.1 Build sherpa-onnx for iOS

This section describes how to build `sherpa-onnx` for iPhone and iPad.

Requirement

Warning: The minimum deployment requires the iOS version >= 13.0.

Before we continue, please make sure the following requirements are satisfied:

- macOS. It won't work on Windows or Linux.
- Xcode. The version 14.2 (14C18) is known to work. Other versions may also work.
- CMake. CMake 3.25.1 is known to work. Other versions may also work.
- (Optional) iPhone or iPad. This is for testing the app on your device. If you don't have a device, you can still run the app within a simulator on your Mac.

Caution:

If you get the following error:

```
CMake Error at toolchains/ios.toolchain.cmake:544 (get_filename_component):
  get_filename_component called with incorrect number of arguments
Call Stack (most recent call first):
  /usr/local/Cellar/cmake/3.29.0/share/cmake/Modules/CMakeDetermineSystem.
  ↗cmake:146 (include)
  CMakeLists.txt:2 (project)
```

please run:

```
sudo xcode-select --install
sudo xcodebuild -license
```

And then delete the build directory `./build-ios` and re-build.

Please see also <https://github.com/k2-fsa/sherpa-onnx/issues/702>.

Download sherpa-onnx

First, let us download the source code of `sherpa-onnx`.

Note: In the following, I will download `sherpa-onnx` to `$HOME/open-source`, i.e., `/Users/fangjun/open-source`, on my Mac.

You can put it anywhere as you like.

```
mkdir -p $HOME/open-source
cd $HOME/open-source
git clone https://github.com/k2-fsa/sherpa-onnx
```

Build sherpa-onnx (in commandline, C++ Part)

After downloading `sherpa-onnx`, let us build the C++ part of `sherpa-onnx`.

```
cd $HOME/open-source/sherpa-onnx/
./build-ios.sh
```

It will generate a directory `$HOME/open-source/sherpa-onnx/build-ios`, which we have already pre-configured for you in Xcode.

Build sherpa-onnx (in Xcode)

Use the following command to open `sherpa-onnx` in Xcode:

```
cd $HOME/open-source/sherpa-onnx/ios-swift/SherpaOnnx
open SherpaOnnx.xcodeproj
```

It will start Xcode and you will see the following screenshot:

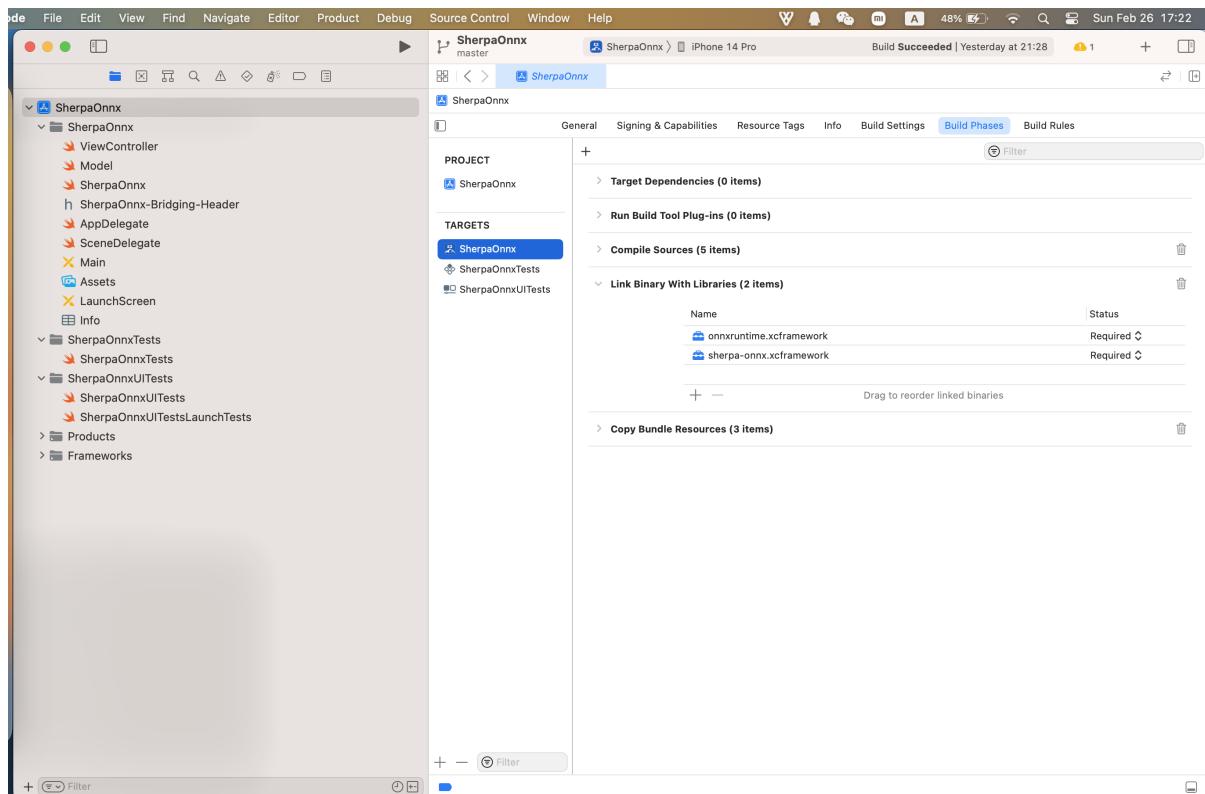


Fig. 8.8: Screenshot after running the command `open SherpaOnnx.xcodeproj`

Please select `Product -> Build` to build the project. See the screenshot below:

After finishing the build, you should see the following screenshot:

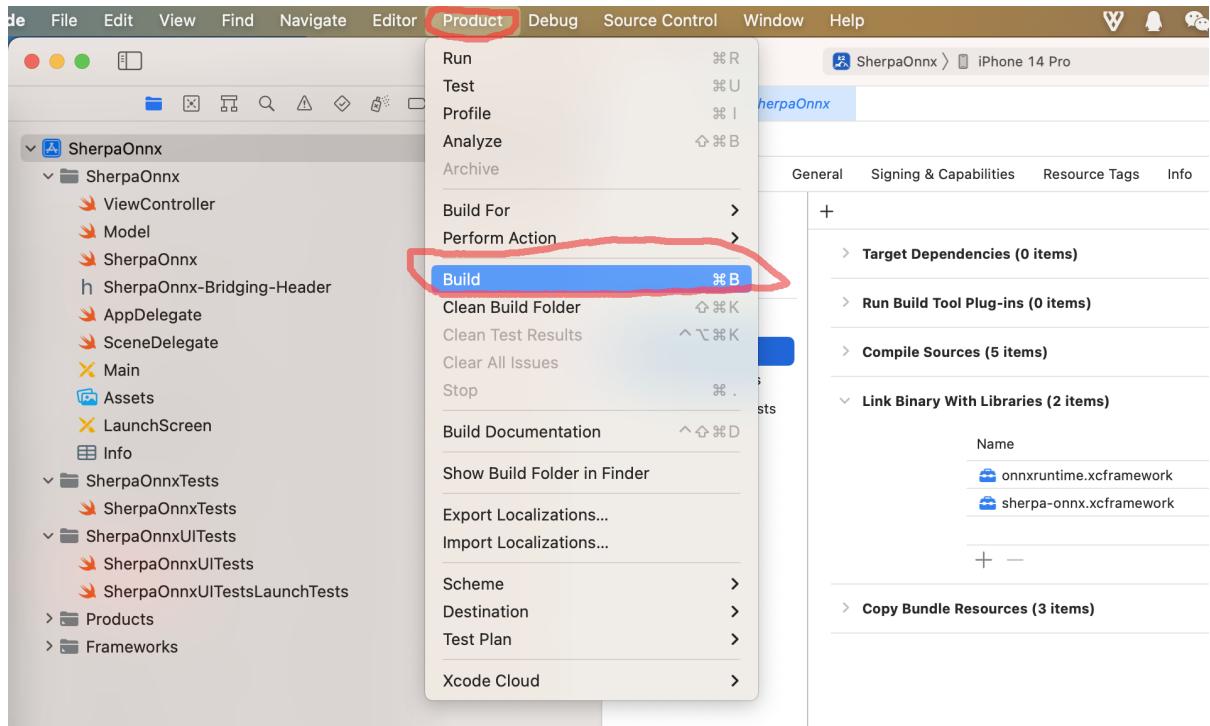


Fig. 8.9: Screenshot for selecting Product -> Build



Fig. 8.10: Screenshot after finishing the build.

Congratulations! You have successfully built the project. Let us run the project by selecting Product -> Run, which is shown in the following screenshot:

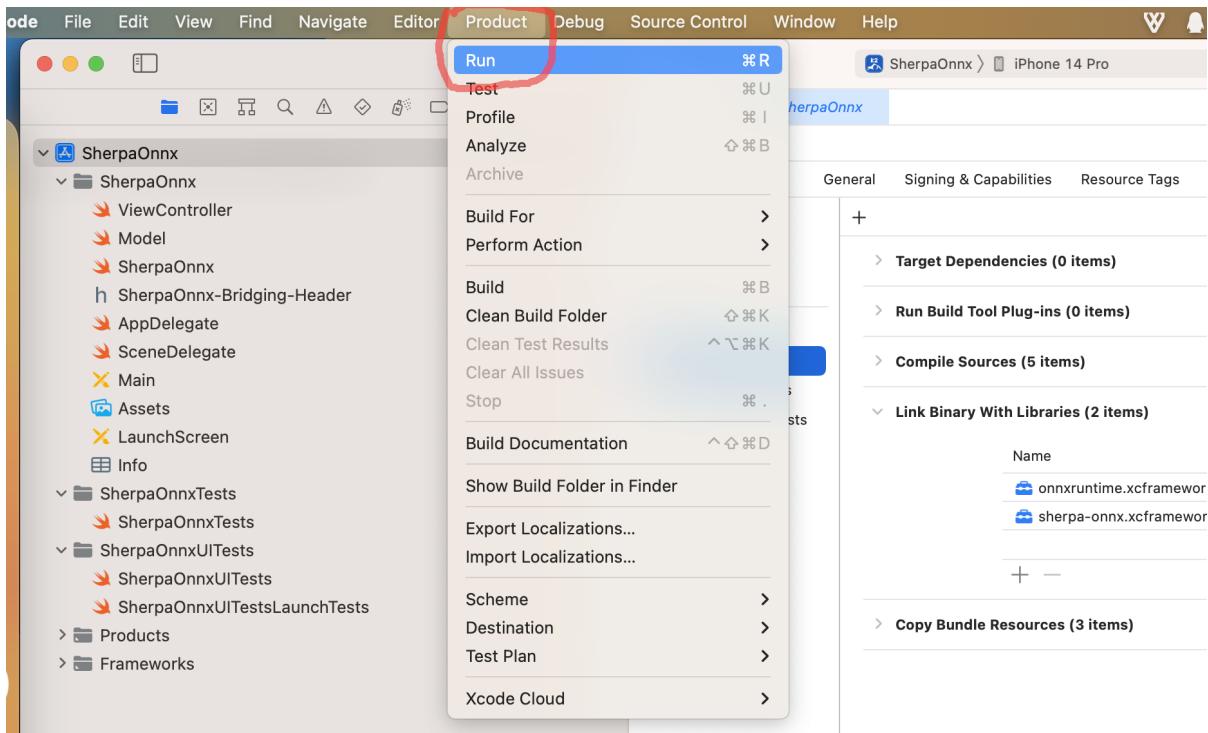


Fig. 8.11: Screenshot for Product -> Run.

Please wait for a few seconds before Xcode starts the simulator.

Unfortunately, it will throw the following error:

The reason for the above error is that we have not provided the pre-trained model yet.

The file `ViewController.swift` pre-selects the pre-trained model to be `csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 (Bilingual, Chinese + English)`, shown in the screenshot below:

Let us add the pre-trained model `csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 (Bilingual, Chinese + English)` to Xcode. Please follow `csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 (Bilingual, Chinese + English)` to download it from huggingface. You can download it to any directory as you like.

Please right click the project SherpaOnnx and select Add Files to "SherpaOnnx" . . . in the popup menu, as is shown in the screenshot below:

In the popup dialog, switch to the folder where you just downloaded the pre-trained model.

In the screenshot below, it is the folder `/Users/fangjun/open-source/icefall-models/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20`:

Select required files and click the button Add:

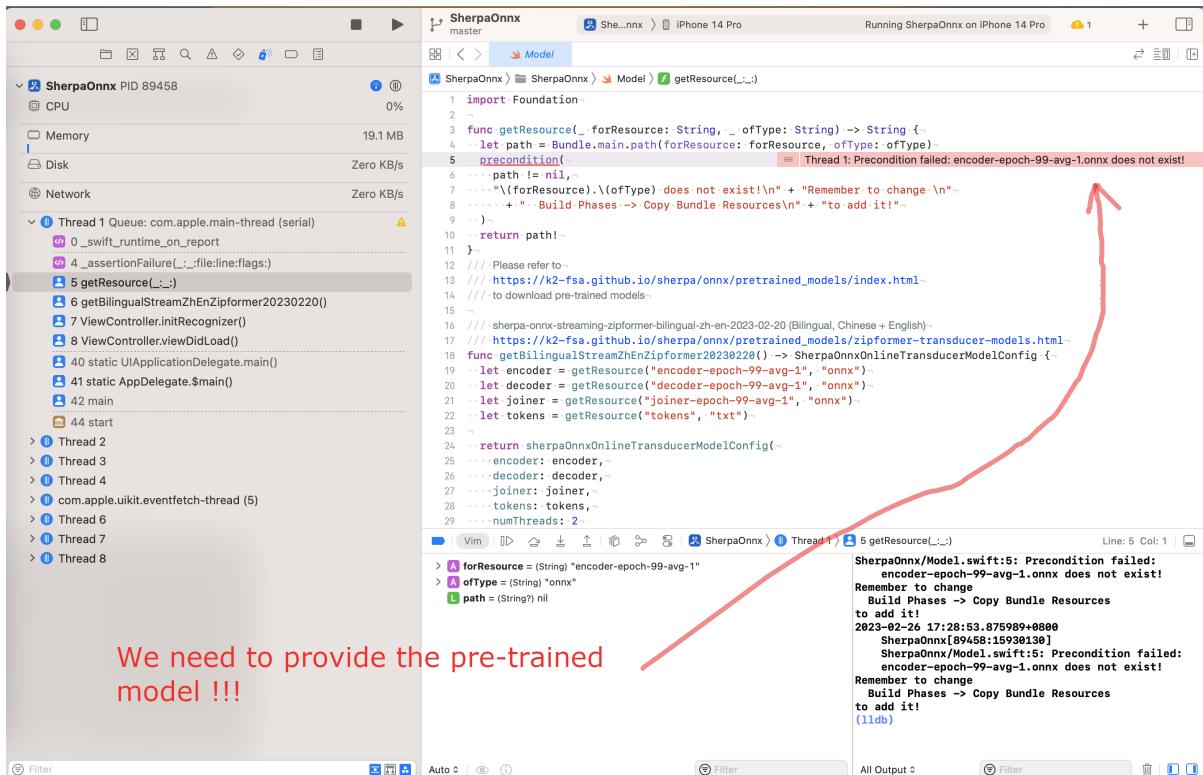


Fig. 8.12: Screenshot for the error




Fig. 8.13: Screenshot for the pre-selected pre-trained model




Fig. 8.14: Screenshot for adding files to SherpaOnnx

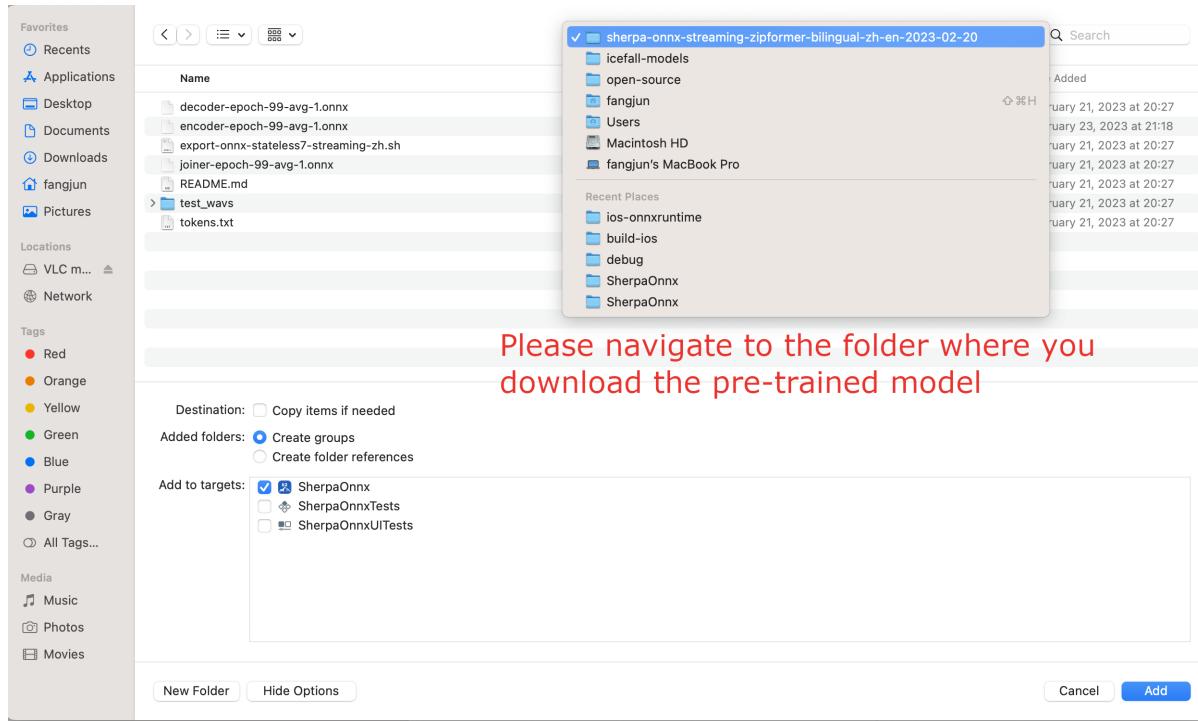


Fig. 8.15: Screenshot for navigating to the folder containing the downloaded pre-trained

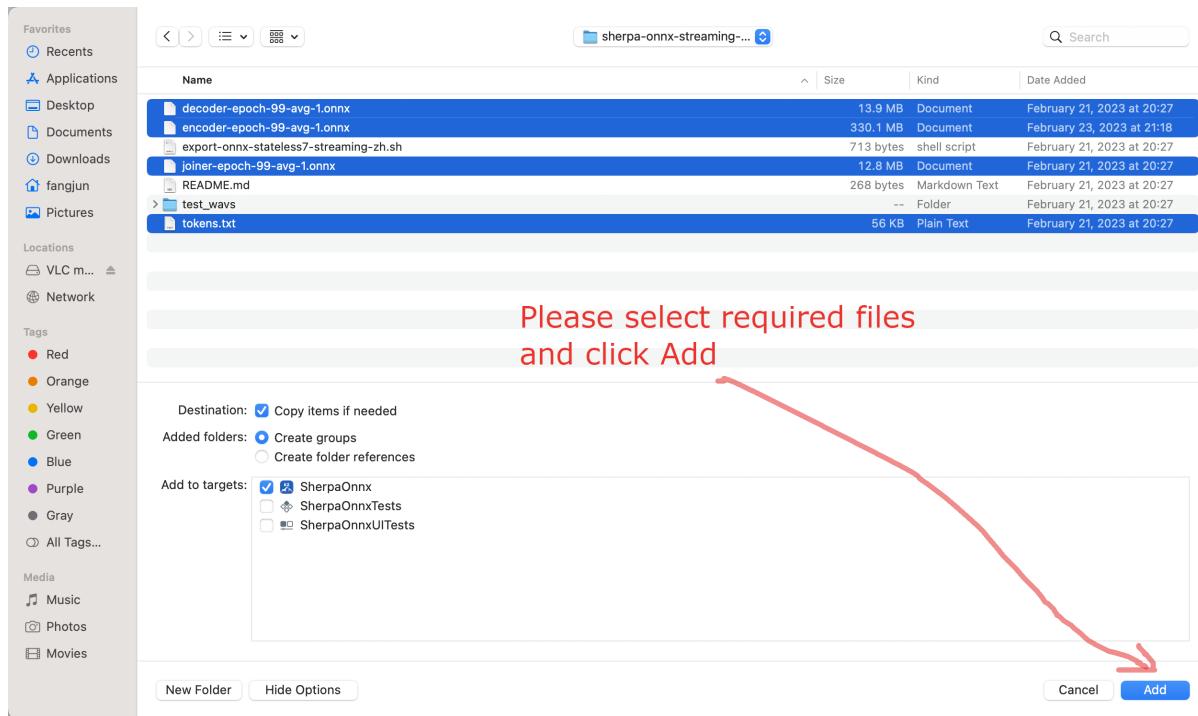


Fig. 8.16: Screenshot for selecting required files

After adding pre-trained model files to Xcode, you should see the following screenshot:

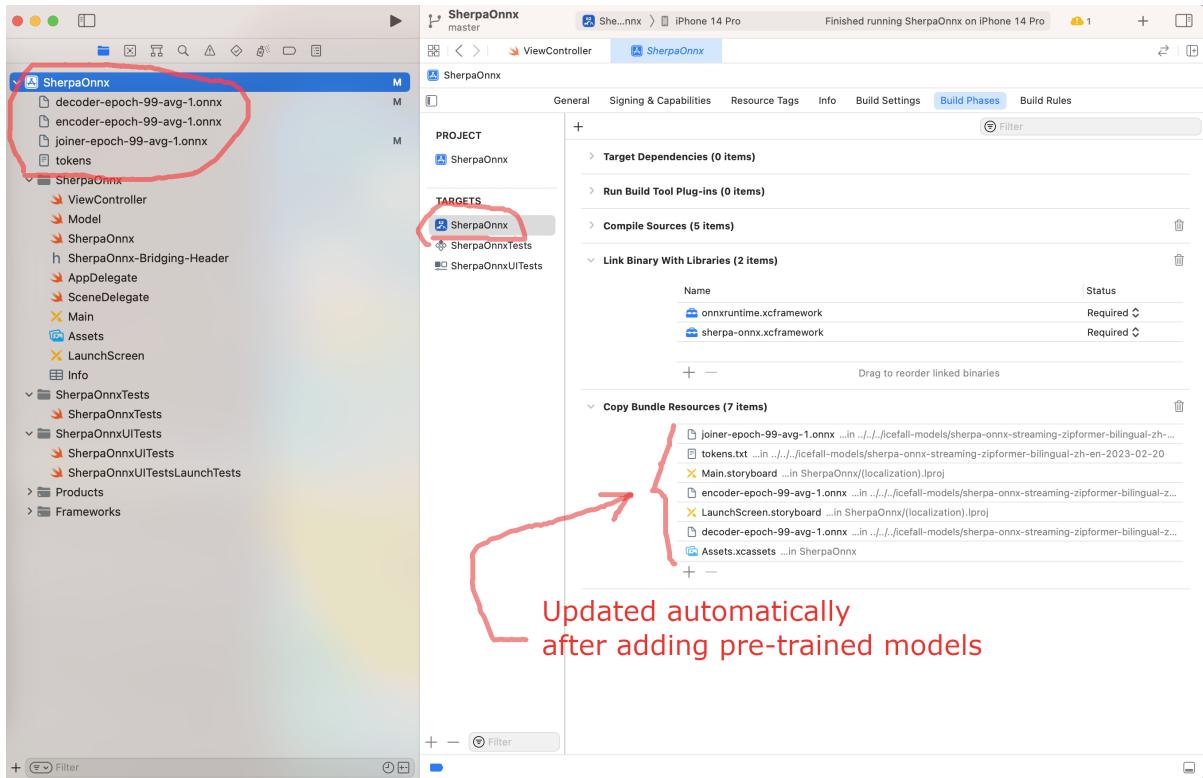


Fig. 8.17: Screenshot after add pre-trained model files

At this point, you should be able to select the menu Product → Run to run the project and you should finally see the following screenshot:

Click the button to start recording! A screenshot is given below:

Congratulations! You have finally succeeded in running `sherpa-onnx` with iOS, though it is in a simulator.

Please read below if you want to run `sherpa-onnx` on your iPhone or iPad.

Run `sherpa-onnx` on your iPhone/iPad

First, please make sure the iOS version of your iPhone/iPad is ≥ 13.0 .

Click the menu Xcode → Settings..., as is shown in the following screenshot:

In the popup dialog, please select Account and click + to add your Apple ID, as is shown in the following screenshots.

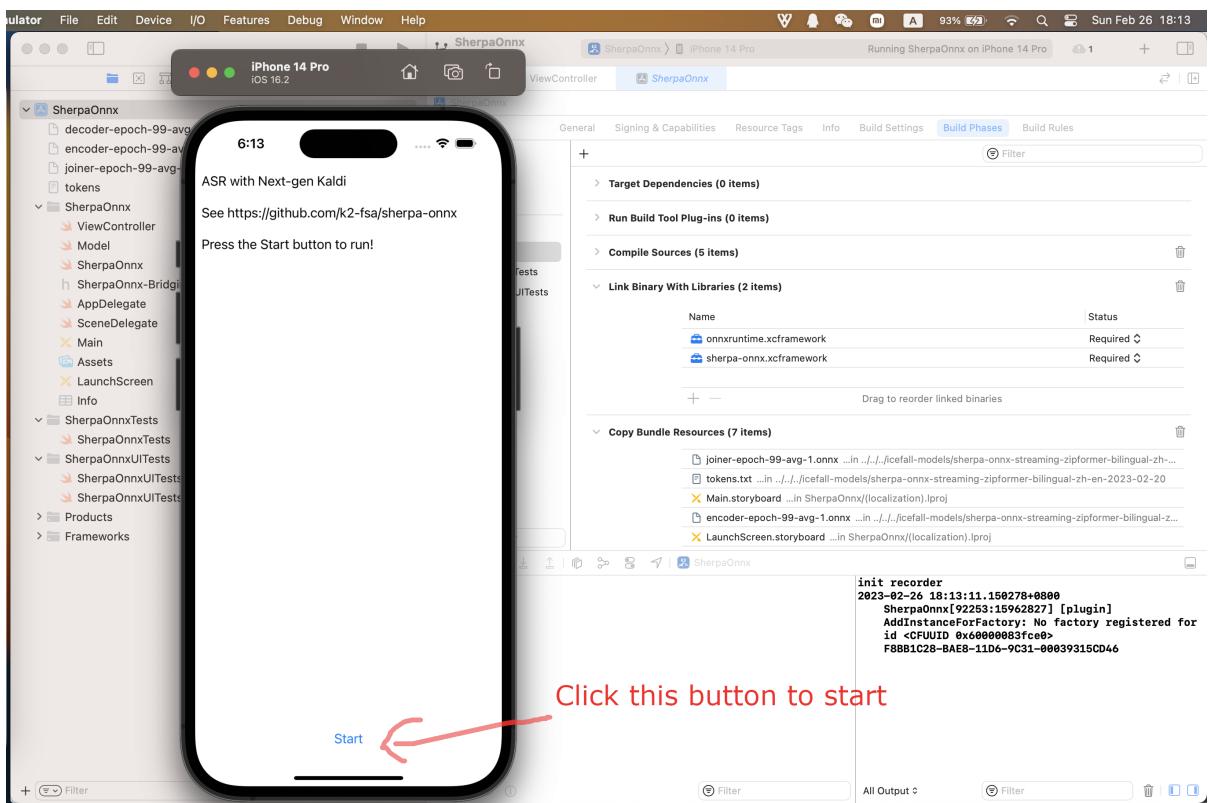


Fig. 8.18: Screenshot for a successful run.

After adding your Apple ID, please connect your iPhone or iPad to your Mac and select your device in Xcode. The following screenshot is an example to select my iPhone.

Now your Xcode should look like below after selecting a device:

Please select Product -> Run again to run `sherpa-onnx` on your selected device, as is shown in the following screenshot:

After a successful build, check your iPhone/iPad and you should see the following screenshot:

At this point, you should be able to run the app on your device. The following is a screenshot about running it on my iPhone:

Congratulations! You have successfully run `sherpa-onnx` on your device!




Fig. 8.19: Screenshot for recording and recognition.




Fig. 8.20: Screenshot for Xcode -> Settings...

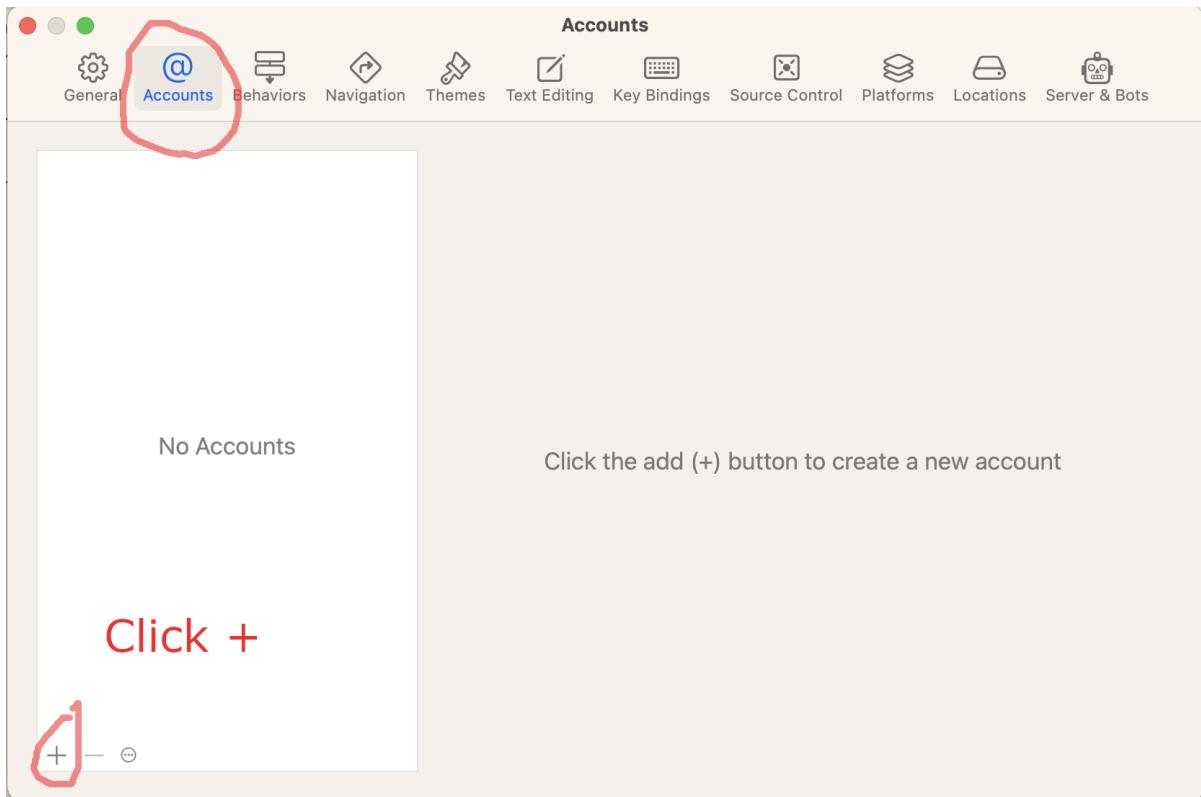


Fig. 8.21: Screenshot for selecting Account and click +.

8.17 Flutter

8.17.1 Pre-built Flutter Apps

Links for pre-built Apps can be found in the following table:

Hint: It runs locally, without internet connection.

Text to speech (TTS, Speech synthesis)

	URL
Android (arm64-v8a, armeabi-v7a, x86_64)	https://k2-fsa.github.io/sherpa/onnx/flutter/tts-android.html
Linux (x64)	https://k2-fsa.github.io/sherpa/onnx/flutter/tts-linux.html
macOS (x64)	https://k2-fsa.github.io/sherpa/onnx/flutter/tts-macos-x64.html
macOS (arm64)	https://k2-fsa.github.io/sherpa/onnx/flutter/tts-macos-arm64.html
Windows (x64)	https://k2-fsa.github.io/sherpa/onnx/flutter/tts-win.html

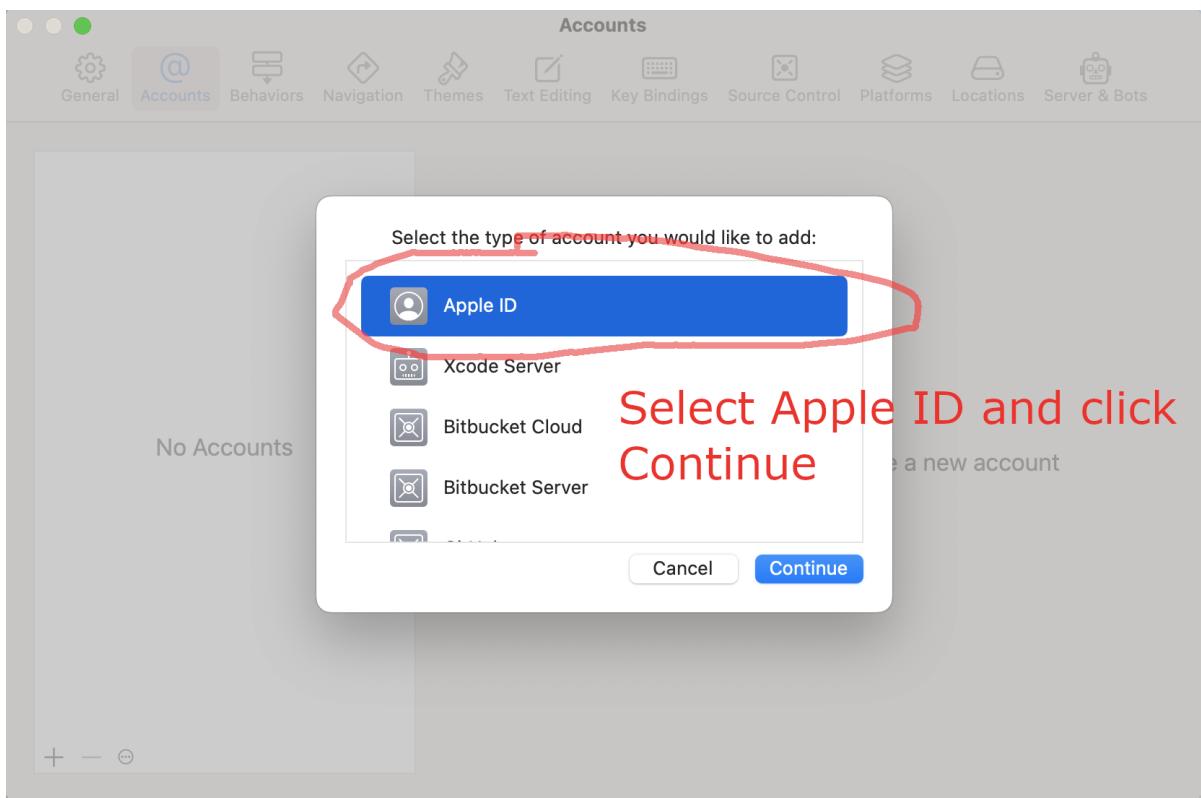


Fig. 8.22: Screenshot for selecting Apple ID and click Continue

Streaming Speech recognition (STT, ASR)

	URL
Streaming speech recognition	https://k2-fsa.github.io/sherpa/onnx/flutter/asr/app.html

8.18 WebSocket

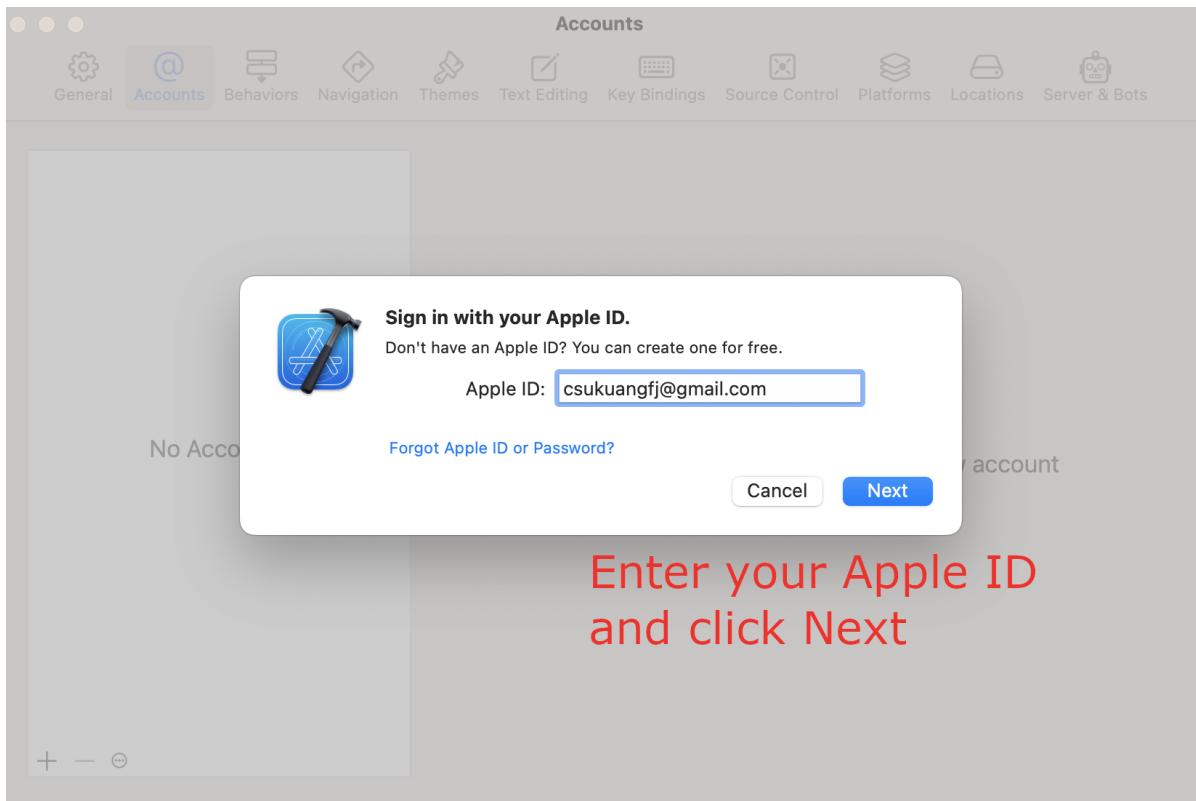
In this section, we describe how to use the `WebSocket` server and client for real-time speech recognition with `sherpa-onnx`.

The `WebSocket` server is implemented in C++ with the help of `websocketpp` and `asio`.

Hint: It does not depend on `boost`.

It does not depend on `boost`.

It does not depend on `boost`.



Enter your Apple ID
and click Next

Fig. 8.23: Screenshot for adding your Apple ID and click Next

8.18.1 Streaming WebSocket server and client

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

Build `sherpa-onnx` with WebSocket support

By default, it will generate the following binaries after [Installation](#):

```
sherpa-onnx fangjun$ ls -lh build/bin/*websocket*
-rwxr-xr-x 1 fangjun staff 1.1M Mar 31 22:09 build/bin/sherpa-onnx-offline-websocket-
server
-rwxr-xr-x 1 fangjun staff 1.0M Mar 31 22:09 build/bin/sherpa-onnx-online-websocket-
client
-rwxr-xr-x 1 fangjun staff 1.2M Mar 31 22:09 build/bin/sherpa-onnx-online-websocket-
server
```

Please refer to [Non-streaming WebSocket server and client](#) for the usage of `sherpa-onnx-offline-websocket-server`.

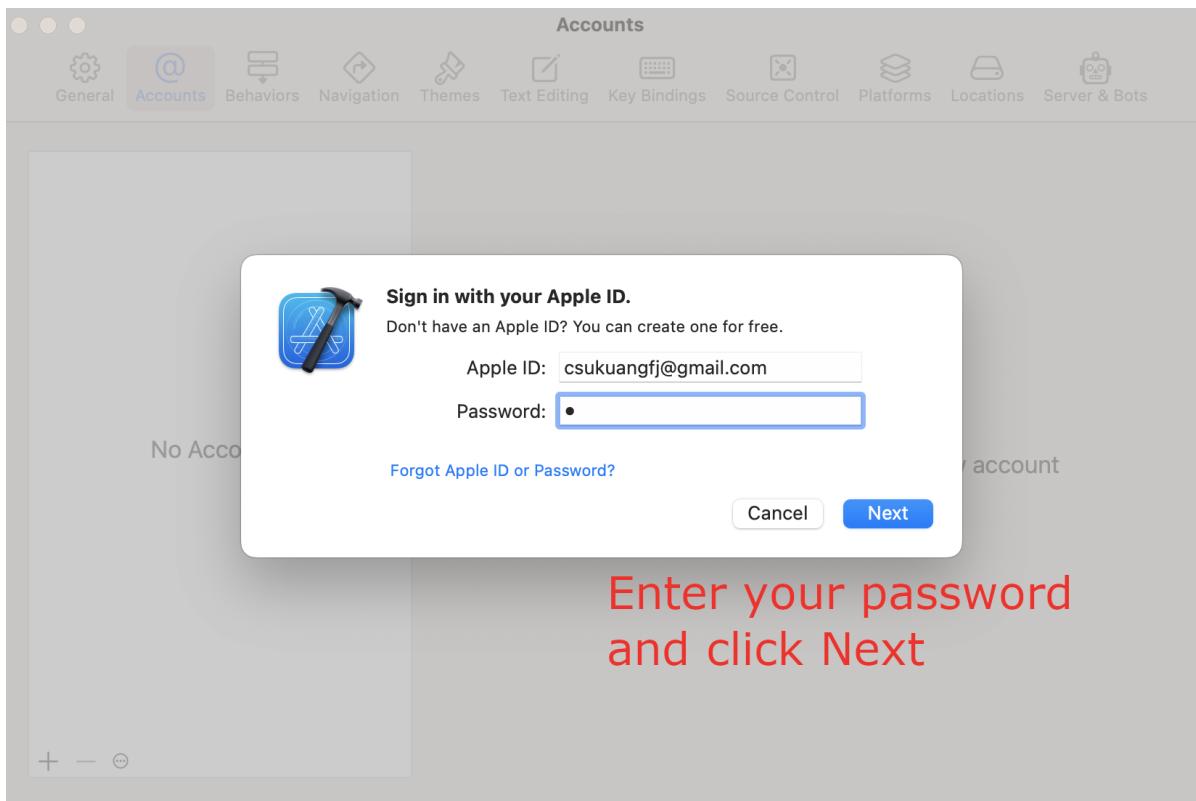


Fig. 8.24: Screenshot for entering your password and click Next

View the server usage

Before starting the server, let us view the help message of `sherpa-onnx-online-websocket-server`:

```
build/bin/sherpa-onnx-online-websocket-server
```

The above command will print the following help information:

```
Automatic speech recognition with sherpa-onnx using websocket.
```

Usage:

```
./bin/sherpa-onnx-online-websocket-server --help
```

```
./bin/sherpa-onnx-online-websocket-server \
--port=6006 \
--num-work-threads=5 \
--tokens=/path/to/tokens.txt \
--encoder=/path/to/encoder.onnx \
--decoder=/path/to/decoder.onnx \
--joiner=/path/to/joiner.onnx \
--log-file=./log.txt \
--max-batch-size=5 \
```

(continues on next page)

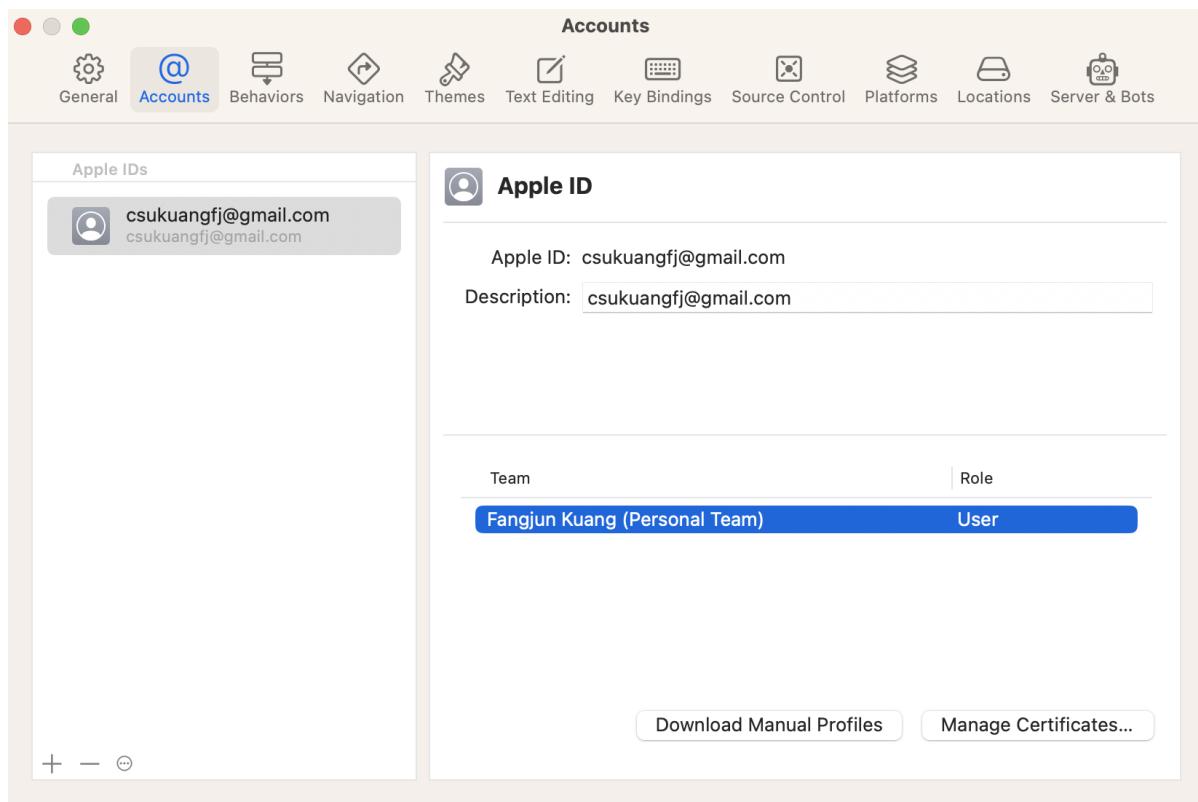


Fig. 8.25: Screenshot after adding your Apple ID

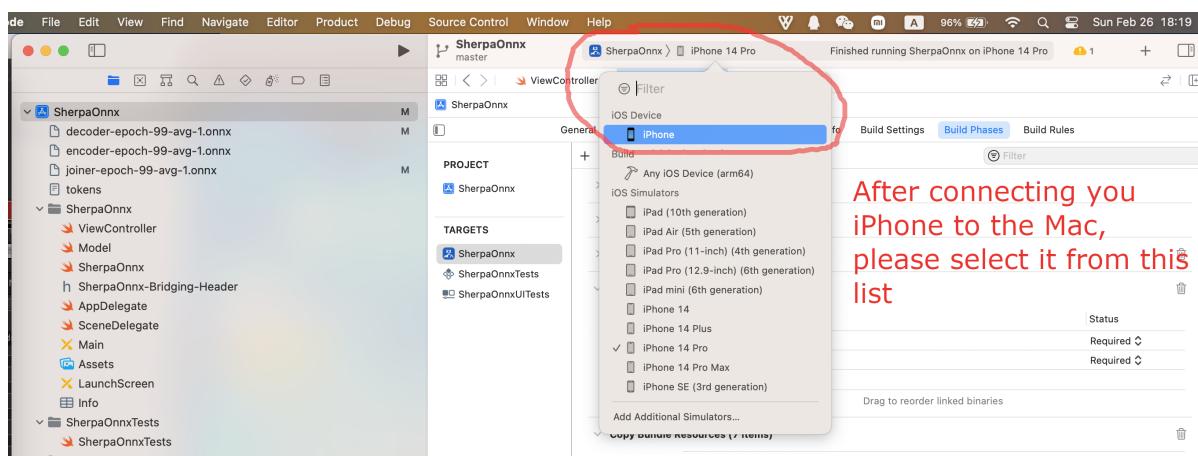


Fig. 8.26: Screenshot for selecting your device

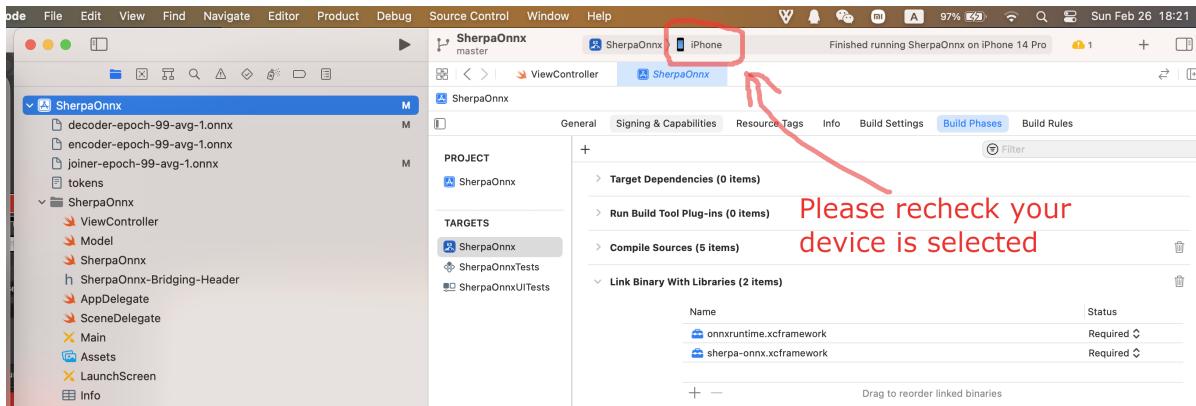


Fig. 8.27: Screenshot after selecting your device

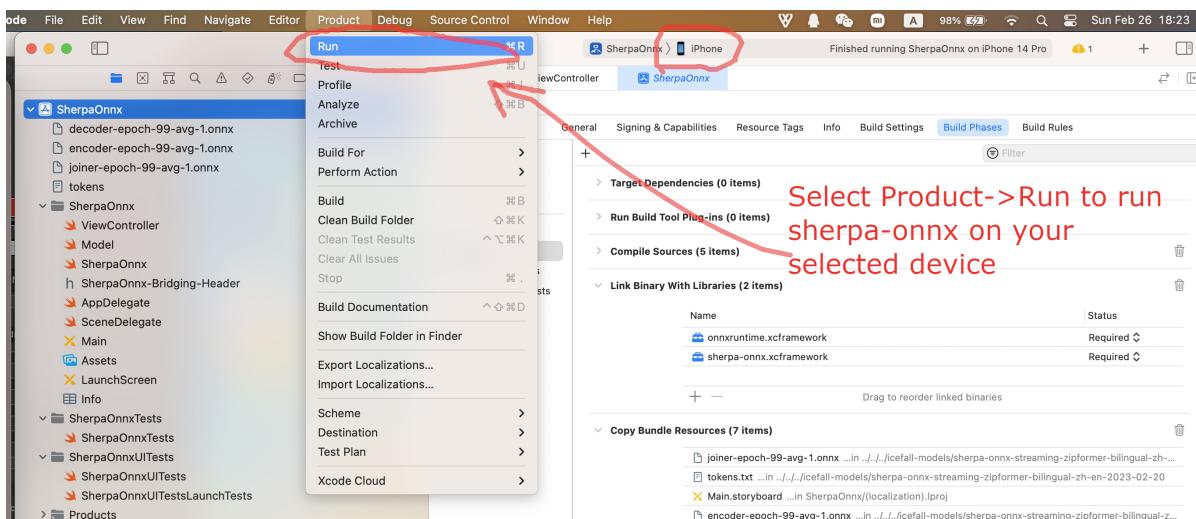


Fig. 8.28: Screenshot for selecting Product -> Run




Fig. 8.29: Screenshot for running sherpa-onnx on your device

··· 中国联通 18:28 100% ⚡

ASR with Next-gen Kaldi

See <https://github.com/k2-fsa/sherpa-onnx>

Press the Start button to run!

Press "Start"




Fig. 8.30: Screenshot for running sherpa-onnx on iPhone

(continued from previous page)

```
--loop-interval-ms=10

Please refer to
https://k2-fsa.github.io/sherpa/onnx/pretrained_models/index.html
for a list of pre-trained models to download.

Options:
--max-batch-size : Max batch size for recognition. (int, default = 5)
--loop-interval-ms : It determines how often the decoder loop runs. (int, default = 10)
--max-active-paths : beam size used in modified beam search. (int, default = 4)
--decoding-method : decoding method, now support greedy_search and modified_beam_search. (string, default = "greedy_search")
--rule3-min-utterance-length : This endpointing rule3 requires utterance-length (in seconds) to be >= this value. (float, default = 20)
--rule3-min-trailing-silence : This endpointing rule3 requires duration of trailing silence in seconds) to be >= this value. (float, default = 0)
--rule3-must-contain-nonsilence : If True, for this endpointing rule3 to apply there must be nonsilence in the best-path traceback. For decoding, a non-blank token is considered as non-silence (bool, default = false)
--rule2-min-utterance-length : This endpointing rule2 requires utterance-length (in seconds) to be >= this value. (float, default = 0)
--rule1-min-trailing-silence : This endpointing rule1 requires duration of trailing silence in seconds) to be >= this value. (float, default = 2.4)
--feat-dim : Feature dimension. Must match the one expected by the model. (int, default = 80)
--rule1-must-contain-nonsilence : If True, for this endpointing rule1 to apply there must be nonsilence in the best-path traceback. For decoding, a non-blank token is considered as non-silence (bool, default = false)
--enable-endpoint : True to enable endpoint detection. False to disable it. (bool, default = true)
--num_threads : Number of threads to run the neural network (int, default = 2)
--debug : true to print model information while loading it. (bool, default = false)
--port : The port on which the server will listen. (int, default = 6006)
--num-io-threads : Thread pool size for network connections. (int, default = 1)
--rule2-must-contain-nonsilence : If True, for this endpointing rule2 to apply there must be nonsilence in the best-path traceback. For decoding, a non-blank token is considered as non-silence (bool, default = true)
--joiner : Path to joiner.onnx (string, default = "")
--tokens : Path to tokens.txt (string, default = "")
--num-work-threads : Thread pool size for neural network computation and decoding. (int, default = 3)
--encoder : Path to encoder.onnx (string, default = "")
--sample-rate : Sampling rate of the input waveform. Note: You can have a different sample rate for the input waveform. We will do resampling inside the feature extractor (int, default = 16000)
--rule2-min-trailing-silence : This endpointing rule2 requires duration of trailing silence in seconds) to be >= this value. (float, default = 1.2)
```

(continues on next page)

(continued from previous page)

```
--log-file : Path to the log file. Logs are appended to this file
--(string, default = "./log.txt")
--rule1-min-utterance-length : This endpointing rule1 requires utterance-length (in
--seconds) to be >= this value. (float, default = 0)
--decoder : Path to decoder.onnx (string, default = "")

Standard options:
--config : Configuration file to read (this option may be repeated)
--(string, default = "")
--help : Print out usage message (bool, default = false)
--print-args : Print the command line arguments (to stderr) (bool,
--default = true)
```

Start the server

Hint: Please refer to [Pre-trained models](#) for a list of pre-trained models.

```
./build/bin/sherpa-onnx-online-websocket-server \
--port=6006 \
--num-work-threads=3 \
--num-io-threads=2 \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-
--99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-
--99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-
--avg-1.onnx \
--log-file=./log.txt \
--max-batch-size=5 \
--loop-interval-ms=20
```

Caution: The arguments are in the form `--key=value`.

It does not support `--key` value.

It does not support `--key` value.

It does not support `--key` value.

Hint: In the above demo, the model files are from [csukuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 \(Bilingual, Chinese + English\)](#).

Note: Note that the server supports processing multiple clients in a batch in parallel. You can use `--max-batch-size`

to limit the batch size.

View the usage of the client (C++)

Let us view the usage of the C++ `WebSocket` client:

```
./build/bin/sherpa-onnx-online-websocket-client
```

The above command will print the following help information:

```
[I] /Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:484:int
↳sherpa_onnx::ParseOptions::Read(int, const char *const *) ./build/bin/sherpa-onnx-
↳online-websocket-client
[I] /Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:525:void
↳sherpa_onnx::ParseOptions::PrintUsage(bool) const
```

Automatic speech recognition **with** `sherpa-onnx` using `websocket`.

Usage:

```
./bin/sherpa-onnx-online-websocket-client --help

./bin/sherpa-onnx-online-websocket-client \
--server-ip=127.0.0.1 \
--server-port=6006 \
--samples-per-message=8000 \
--seconds-per-message=0.2 \
/path/to/foo.wav
```

It support only wave of **with** a single channel, 16kHz, 16-bit samples.

Options:

```
--seconds-per-message      : We will simulate that each message takes this number of
↳seconds to send. If you select a very large value, it will take a long time to send
↳all the samples (float, default = 0.2)
--samples-per-message     : Send this number of samples per message. (int, default =
↳8000)
--sample-rate              : Sample rate of the input wave. Should be the one
↳expected by the server (int, default = 16000)
--server-port               : Port of the websocket server (int, default = 6006)
--server-ip                 : IP address of the websocket server (string, default =
↳"127.0.0.1")
```

Standard options:

```
--help                     : Print out usage message (bool, default = false)
--print-args                : Print the command line arguments (to stderr) (bool,
↳default = true)
--config                   : Configuration file to read (this option may be repeated)
↳(string, default = "")
```

Caution: We only support using IP address for --server-ip.

For instance, please don't use --server-ip=localhost, use --server-ip=127.0.0.1 instead.

For instance, please don't use --server-ip=localhost, use --server-ip=127.0.0.1 instead.

For instance, please don't use --server-ip=localhost, use --server-ip=127.0.0.1 instead.

Start the client (C++)

To start the C++ WebSocket client, use:

```
build/bin/sherpa-onnx-online-websocket-client \
--seconds-per-message=0.1 \
--server-port=6006 \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/0.wav
```

Since the server is able to process multiple clients at the same time, you can use the following command to start multiple clients:

```
for i in $(seq 0 10); do
    k=$((expr $i % 5))
    build/bin/sherpa-onnx-online-websocket-client \
        --seconds-per-message=0.1 \
        --server-port=6006 \
        ./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/${k}.wav &
done

wait

echo "done"
```

View the usage of the client (Python)

Use the following command to view the usage:

```
python3 ./python-api-examples/online-websocket-client-decode-file.py --help
```

Hint: online-websocket-client-decode-file.py is from <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/online-websocket-client-decode-file.py>

It will print:

```
usage: online-websocket-client-decode-file.py [-h] [--server-addr SERVER_ADDR]
                                               [--server-port SERVER_PORT]
                                               [--samples-per-message SAMPLES_PER_MESSAGE]
                                               [--seconds-per-message SECONDS_PER_MESSAGE]
                                               sound_file

positional arguments:
  sound_file            The input sound file. Must be wave with a single
```

(continues on next page)

(continued from previous page)

```
channel, 16kHz sampling rate, 16-bit of each sample.

optional arguments:
  -h, --help            show this help message and exit
  --server-addr SERVER_ADDR
                        Address of the server (default: localhost)
  --server-port SERVER_PORT
                        Port of the server (default: 6006)
  --samples-per-message SAMPLES_PER_MESSAGE
                        Number of samples per message (default: 8000)
  --seconds-per-message SECONDS_PER_MESSAGE
                        We will simulate that the duration of two messages is
                        of this value (default: 0.1)
```

Hint: For the Python client, you can use either a domain name or an IP address for `--server-addr`. For instance, you can use either `--server-addr localhost` or `--server-addr 127.0.0.1`.

For the input argument, you can either use `--key=value` or `--key value`.

Start the client (Python)

```
python3 ./python-api-examples/online-websocket-client-decode-file.py \
  --server-addr localhost \
  --server-port 6006 \
  --seconds-per-message 0.1 \
  ./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/4.wav
```

Start the client (Python, with microphone)

```
python3 ./python-api-examples/online-websocket-client-microphone.py \
  --server-addr localhost \
  --server-port 6006

``online-websocket-client-microphone.py`` is from
`<https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/online-
websocket-client-microphone.py>`_
```

8.18.2 Non-streaming WebSocket server and client

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

Build `sherpa-onnx` with WebSocket support

By default, it will generate the following binaries after [Installation](#):

```
sherpa-onnx fangjun$ ls -lh build/bin/*websocket*
-rwxr-xr-x 1 fangjun staff 1.1M Mar 31 22:09 build/bin/sherpa-onnx-offline-websocket-
 ↵server
-rwxr-xr-x 1 fangjun staff 1.0M Mar 31 22:09 build/bin/sherpa-onnx-online-websocket-
 ↵client
-rwxr-xr-x 1 fangjun staff 1.2M Mar 31 22:09 build/bin/sherpa-onnx-online-websocket-
 ↵server
```

Please refer to [Streaming WebSocket server and client](#) for the usage of `sherpa-onnx-online-websocket-server` and `sherpa-onnx-online-websocket-client`.

View the server usage

Before starting the server, let us view the help message of `sherpa-onnx-offline-websocket-server`:

```
build/bin/sherpa-onnx-offline-websocket-server
```

The above command will print the following help information:

```
Automatic speech recognition with sherpa-onnx using websocket.
```

Usage:

```
./bin/sherpa-onnx-offline-websocket-server --help
```

(1) For transducer models

```
./bin/sherpa-onnx-offline-websocket-server \
--port=6006 \
--num-work-threads=5 \
--tokens=/path/to/tokens.txt \
--encoder=/path/to/encoder.onnx \
--decoder=/path/to/decoder.onnx \
--joiner=/path/to/joiner.onnx \
--log-file=./log.txt \
--max-batch-size=5
```

(2) For Paraformer

```
./bin/sherpa-onnx-offline-websocket-server \
--port=6006 \
--num-work-threads=5 \
--tokens=/path/to/tokens.txt \
--paraformer=/path/to/model.onnx \
--log-file=./log.txt \
--max-batch-size=5
```

Please refer to

https://k2-fsa.github.io/sherpa/onnx/pretrained_models/index.html

(continues on next page)

(continued from previous page)

for a list of pre-trained models to download.

Options:

```
--log-file : Path to the log file. Logs are appended to this file.  
→(string, default = "./log.txt")  
--max-utterance-length : Max utterance length in seconds. If we receive an  
→utterance longer than this value, we will reject the connection. If you have enough  
→memory, you can select a large value for it. (float, default = 300)  
--decoding-method : decoding method,Valid values: greedy_search. (string,  
→default = "greedy_search")  
--num-threads : Number of threads to run the neural network (int,  
→default = 2)  
--feat-dim : Feature dimension. Must match the one expected by the  
→model. (int, default = 80)  
--port : The port on which the server will listen. (int, default  
→= 6006)  
--debug : true to print model information while loading it. (bool,  
→default = false)  
--joiner : Path to joiner.onnx (string, default = "")  
--tokens : Path to tokens.txt (string, default = "")  
--encoder : Path to encoder.onnx (string, default = "")  
--num-work-threads : Thread pool size for neural network computation and  
→decoding. (int, default = 3)  
--paraformer : Path to model.onnx of paraformer. (string, default = "")  
--num-io-threads : Thread pool size for network connections. (int, default  
→= 1)  
--max-batch-size : Max batch size for decoding. (int, default = 5)  
--decoder : Path to decoder.onnx (string, default = "")
```

Standard options:

```
--help : Print out usage message (bool, default = false)  
--print-args : Print the command line arguments (to stderr) (bool,  
→default = true)  
--config : Configuration file to read (this option may be repeated).  
→(string, default = "")
```

Start the server

Hint: Please refer to [Pre-trained models](#) for a list of pre-trained models.

Start the server with a transducer model

```
./build/bin/sherpa-onnx-offline-websocket-server \
--port=6006 \
--num-work-threads=5 \
--tokens=./sherpa-onnx-zipformer-en-2023-03-30/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.onnx \
--log-file=./log.txt \
--max-batch-size=5
```

Caution: The arguments are in the form `--key=value`.

It does not support `--key value`.

It does not support `--key value`.

It does not support `--key value`.

Hint: In the above demo, the model files are from csukuangfj/sherpa-onnx-zipformer-en-2023-03-30 (*English*).

Note: Note that the server supports processing multiple clients in a batch in parallel. You can use `--max-batch-size` to limit the batch size.

Start the server with a paraformer model

```
./build/bin/sherpa-onnx-offline-websocket-server \
--port=6006 \
--num-work-threads=5 \
--tokens=./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2023-03-28/model.onnx \
--log-file=./log.txt \
--max-batch-size=5
```

Hint: In the above demo, the model files are from csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28 (*Chinese + English*).

Start the client (Python)

We provide two clients written in Python:

- `offline-websocket-client-decode-files-paralell.py`: It decodes multiple files in parallel by creating a separate connection for each file
- `offline-websocket-client-decode-files-sequential.py`: It decodes multiple files sequentially by creating only a single connection

`offline-websocket-client-decode-files-paralell.py`

```
python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

`offline-websocket-client-decode-files-sequential.py`

```
python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py \
--server-addr localhost \
--server-port 6006 \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

8.19 Hotwords (Contextual biasing)

In this section, we describe how we implement the hotwords (aka contextual biasing) feature with an Aho-corasick automaton and how to use it in `sherpa-onnx`.

8.19.1 What are hotwords

Current ASR systems work very well for general cases, but they sometimes fail to recognize special words/phrases (aka hotwords) like rare words, personalized information etc. Usually, those words/phrases will be recognized as the words/phrases that pronounce similar to them (for example, recognize LOUIS FOURTEEN as LEWIS FOURTEEN). So we have to provide some kind of contexts information (for example, the LOUIS FOURTEEN) to the ASR systems to boost those words/phrases. Normally, we call this kind of boosting task contextual biasing (aka hotwords recognition).

8.19.2 How do we implement it with an Aho-corasick

We first construct an Aho-corasick automaton on those given hotwords (after tokenizing into tokens). Please refer to https://en.wikipedia.org/wiki/Aho%20Corasick_algorithm for the construction details of Aho-corasick.

The figure below is the aho-corasick on “HE/SHE/SHELL/HIS/THIS” with hotwords-score==1.

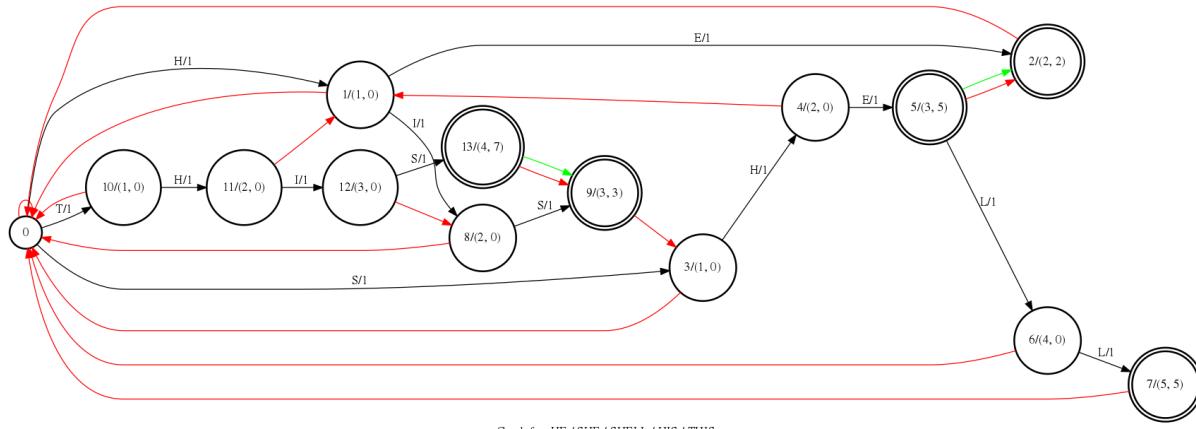


Fig. 8.31: The Aho-corasick for “HE/SHE/SHELL/HIS/THIS”

The black arrows in the graph are the goto arcs, the red arrows are the failure arcs, the green arrows are the output arcs. On each goto arc, there are token and boosting score (Note: we will boost the path when any partial sequence is matched, if the path finally fails to full match any hotwords, the boosted score will be canceled). Currently, the boosting score distributes on the arcs evenly along the path. On each state, there are two scores, the first one is the node score (mainly used to cancel score) the second one is output score, the output score is the total scores of the full matched hotwords of this state.

The following are several matching examples of the graph above.

Note: For simplicity, we assume that the system emits a token each frame.

Hint: We have an extra `finalize` step to force the graph state to go back to the root state.

The path is “SHELF”

Frame	Boost score	Total boost score	Graph state	Matched hotwords
init	0	0	0	
1	1	1	3	
2	1	2	4	
3	1 + 5	8	5	HE, SHE
4	1	9	6	
5	-4	5	0	
finalize	0	5	0	

At frame 3 we reach state 5 and match HE, SHE, so we get a boosting score 1 + 5, the score 1 here because the SHEL still might be the prefix of other hotwords. At frame 5 F can not match any tokens and fail back to root, so we cancel the score for SHEL which is 4 (the node score of state 6).

The path is “HI”

Frame	Boost score	Total boost score	Graph state	Matched hotwords
init	0	0	0	
1	1	1	1	
2	1	2	8	
finalize	-2	0	0	

H and I all match the tokens in the graph, unfortunately, we have to go back to root state when finishing matching a path, so we cancel the boosting score of HI which is 2 (the node score of state 8).

The path is “THE”

Frame	Boost score	Total boost score	Graph state	Matched hotwords
init	0	0	0	
1	1	1	10	
2	1	2	11	
3	0 + 2	4	2	HE
finalize	-2	3	0	

At frame 3 we jump from state 11 to state 2 and get a boosting score of 0 + 2, 0 because the node score of state 2 is the same as state 11 so we don't get score by partial match (the prefix of state 11 is TH has the same length of the prefix of state 2 which is HE), but we do get the output score (at state 2 it outputs HE).

Note: We implement the hotwords feature during inference time, you don't have to re-train the models to use this feature.

8.19.3 How to use hotwords in sherpa-onnx

Caution: Currently, the hotwords feature is only supported in the `modified_beam_search` decoding method of the **transducer models** (both streaming and non-streaming).

The use of the hotwords is no different for streaming and non-streaming models, and in fact it is even no different for all the API supported by sherpa onnx. We add **FOUR** extra arguments for hotwords:

- `hotwords-file`

The file path of the hotwords, one hotwords per line. They could be Chinese words, English words or both according to the modeling units used to train the model. Here are some examples:

For Chinese models trained on `cjkchar` it looks like:

For English like language models trained on `bpe` it looks like:

SPEECH RECOGNITION
DEEP LEARNING

For multilingual models trained on `cjkchar+bpe` (Chinese + English) it looks like:

SPEECH
SPEECH RECOGNITION

You can also specify the boosting score for each hotword, the score should follow the predefined character :, for example:

:3.5
:2.0

It means, hotword will have a boosting score of 3.5, hotword will have a boosting score of 2.0. For those hotwords that don't have specific scores, they will use the global score provided by *hotword-score* below.

Caution: The specific score MUST BE the last item of each hotword (i.e You shouldn't break the hotword into two parts by the score). SPEECH :2.0 # This is invalid

- **hotwords-score**

The boosting score for each matched token.

Note: We match the hotwords at token level, so the **hotwords-score** is applied at token level.

- **modeling-unit**

The modeling unit of the used model, currently support *cjkchar* (for Chinese), *bpe* (for English like languages) and *cjkchar+bpe* (for multilingual models). We need this modeling-unit to select tokenizer to encode words/phrases into tokens, so **do provide correct modeling-unit according to your model**.

- **bpe-vocab**

The bpe vocabulary generated by sentencepiece toolkit, it also can be exported from *bpe.model* (see *script/export_bpe_vocab.py* for details). This vocabulary is used to tokenize words/phrases into bpe units. It is only used when *modeling-unit* is *bpe* or *cjkchar+bpe*.

Hint: We need *bpe.vocab* rather than *bpe.model*, because we don't introduce sentencepiece c++ codebase into *sherpa-onnx* (which has a dependency issue of protobuf), we implement a simple sentencepiece encoder and decoder which takes *bpe.vocab* as input.

The main difference of using hotwords feature is about the modeling units. The following shows how to use it for different modeling units.

Hint: You can use any transducer models here https://k2-fsa.github.io/sherpa/onnx/pretrained_models/index.html, we just choose three of them randomly for the following examples.

Note: In the following example, we use a non-streaming model, if you are using a streaming model, you should use *sherpa-onnx*. *sherpa-onnx-alsa*, *sherpa-onnx-microphone*, *sherpa-onnx-microphone-offline*, *sherpa-onnx-online-websocket-server* and *sherpa-onnx-offline-websocket-server* all support hotwords.

Modeling unit is bpe

Download the model

```
cd /path/to/sherpa-onnx  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
zipformer-en-2023-04-01.tar.bz2  
tar xvf sherpa-onnx-zipformer-en-2023-04-01.tar.bz2  
rm sherpa-onnx-zipformer-en-2023-04-01.tar.bz2  
ln -s sherpa-onnx-zipformer-en-2023-04-01 exp
```

Export bpe.vocab if you can't find bpe.vocab in the model directory.

```
python script/export_bpe_vocab.py --bpe-model exp/bpe.model
```

The hotwords_en.txt contains:

```
QUARTERS  
FOREVER
```

C++ api

Decoding without hotwords

```
./build/bin/sherpa-onnx-offline \  
--encoder=exp/encoder-epoch-99-avg-1.onnx \  
--decoder=exp/decoder-epoch-99-avg-1.onnx \  
--joiner=exp/joiner-epoch-99-avg-1.onnx \  
--decoding-method=modified_beam_search \  
--tokens=exp/tokens.txt \  
exp/test_wavs/0.wav exp/test_wavs/1.wav
```

The output is:

```
/star-kw/kangwei/code/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./build/bin/  
sherpa-onnx-offline --encoder=exp/encoder-epoch-99-avg-1.onnx --decoder=exp/decoder-  
epoch-99-avg-1.onnx --joiner=exp/joiner-epoch-99-avg-1.onnx --decoding-method=modified_  
beam_search --tokens=exp/tokens.txt exp/test_wavs/0.wav exp/test_wavs/1.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_config=OfflineModelConfig(transducer=OfflineTrans-  
ducerModelConfig(encoder_filename="exp/encoder-epoch-99-avg-1.onnx", decoder_filename=  
"exp/decoder-epoch-99-avg-1.onnx", joiner_filename="exp/joiner-epoch-99-avg-1.onnx"),  
paraformer=OfflineParaformerModelConfig(model=""), nemo_=  
ctc=OfflineNemoEncDecCtcModelConfig(model=""),  
whisper=OfflineWhisperModelConfig(encoder="$", decoder="", language="", task="transcribe"  
"$"), tdnn=OfflineTdnmModelConfig(model=""), tokens="exp/tokens.txt", num_threads$2,  
debug=False, provider="cpu", model_type=""), lm_config=OfflineLMConfig(model=""),  
scale=0.5), decoding_method="modified_beam_search", max_active_paths=4, ho$words_file=,  
hotwords_score=1.5)  
Creating recognizer ...  
Started
```

(continues on next page)

(continued from previous page)

Done!

```

exp/test_wavs/0.wav
{"text":"ALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE
→BROTHELS","timestamps":[1.44, 1.48, 1.56, 1.72, 1.88, 1.96, 2.16, 2.28$ 2.36, 2.48, 2.
→60, 2.80, 3.08, 3.28, 3.40, 3.60, 3.80, 4.08, 4.24, 4.32, 4.48, 4.64, 4.84, 4.88, 5.00,
→ 5.08, 5.32, 5.48, 5.60, 5.68, 5.84, 6.04, 6.24],"tokens":["A","LL"," THE"," YE","LL",
→"OW"," LA","M","P","S"," WOULD"," LIGHT"," UP"," HE","RE"," AND"," THERE"," THE"," S",
→"QUA","LI","D"," ","QUA","R","TER"," OF","THE"," B","RO","TH","EL","S"]}

exp/test_wavs/1.wav
{"text":"IN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY CHILD WHOSE PLACE WAS ON THAT
→SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER WITH THE RACE AN
D DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN","timestamps":[2.44, 2.
→64, 2.88, 3.16, 3.28, 3.48, 3.60, 3.80, 3.96, 4.12, 4.36, 4.52, 4.72, 4
.92, 5.16, 5.44, 5.68, 6.04, 6.24, 6.48, 6.84, 7.08, 7.32, 7.56, 7.84, 8.12, 8.24, 8.32,
→8.44, 8.60, 8.76, 8.88, 9.08, 9.28, 9.44, 9.56, 9.64, 9.76, 9.96, 10.0
4, 10.20, 10.40, 10.64, 10.76, 11.04, 11.20, 11.36, 11.60, 11.80, 12.00, 12.12, 12.28,
→12.32, 12.52, 12.72, 12.84, 12.96, 13.04, 13.24, 13.40, 13.60, 13.76, 13
.96, 14.12, 14.24, 14.36, 14.52, 14.68, 14.76, 15.04, 15.28, 15.52, 15.76, 16.00, 16.16,
→16.24, 16.32],"tokens":["IN"," WHICH"," MAN"," TH","US"," P","UN","IS
H","ED"," HAD"," GIVE","N"," HER"," A"," LOVE","LY"," CHILD"," WHO","SE"," PLACE"," WAS",
→" ON"," THAT"," SAME"," DIS","HO","N","OUR","ED"," BO","S","OM"," TO",
" CON","NE","C","T"," HER"," P","AR","ENT"," FOR"," E","VER"," WITH"," THE"," RA","CE",
→"AND"," DE","S","C","ENT"," OF"," MO","R","T","AL","S"," AND"," TO"," B
E"," FI","N","AL","LY"," A"," B","LESS","ED"," SO","UL"," IN"," HE","A","VE","N"]}

num threads: 2
decoding method: modified_beam_search
max active paths: 4
Elapsed seconds: 1.775 s
Real time factor (RTF): 1.775 / 23.340 = 0.076

```

Decoding with hotwords

```

./build/bin/sherpa-onnx-offline \
--encoder=exp/encoder-epoch-99-avg-1.onnx \
--decoder=exp/decoder-epoch-99-avg-1.onnx \
--joiner=exp/joiner-epoch-99-avg-1.onnx \
--decoding-method=modified_beam_search \
--tokens=exp/tokens.txt \
--modeling-unit=bpe \
--bpe-vocab=exp/bpe.vocab \
--hotwords-file=hotwords_en.txt \
--hotwords-score=2.0 \
exp/test_wavs/0.wav exp/test_wavs/1.wav

```

The output is:

```

/star-kw/kangwei/code/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./build/bin/
→sherpa-onnx-offline --encoder=exp/encoder-epoch-99-avg-1.onnx --decoder=exp/decoder-
→epoch-99-avg-1.onnx --joiner=exp/joiner-epoch-99-avg-1.onnx --decoding-method=modified_
→beam_search --tokens=exp/tokens.txt --hotwords-file=hotwords_en.txt --hotwords-score=2.
→0 exp/test_wavs/0.wav exp/test_wavs/1.wav

```

(continues on next page)

(continued from previous page)

```

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
    ↵ feature_dim=80), model_
    ↵ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename=
    ↵ "exp/encoder-epoch-99-avg-1.onnx", decoder_filename="exp/decoder-epoch-99-avg-1.onnx", ↵
    ↵ joiner_filename="exp/joiner-epoch-99-avg-1.onnx"), ↵
    ↵ paraformer=OfflineParaformerModelConfig(model=""), nemo_
    ↵ ctc=OfflineNemoEncDecCtcModelConfig(model=""), ↵
    ↵ whisper=OfflineWhisperModelConfig(encoder="",
    ↵ , decoder="", language="", task="transcribe"), tdnn=OfflineTdnnModelConfig(model=""),
    ↵ tokens="exp/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
    ↵ lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="modified_beam_search",
    ↵ max_active_paths=4, hotwords_file=hotwords_en.txt, hotwords_score=2)
Creating recognizer ...
Started
Done!

exp/test_wavs/0.wav
{"text": "ALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTERS OF THE
    ↵ BROTHELS", "timestamps": "[1.44, 1.48, 1.56, 1.72, 1.88, 1.96, 2.16, 2.28
    , 2.36, 2.48, 2.60, 2.80, 3.08, 3.28, 3.40, 3.60, 3.80, 4.08, 4.24, 4.32, 4.48, 4.
    ↵ 84, 4.88, 5.00, 5.08, 5.12, 5.36, 5.48, 5.60, 5.68, 5.84, 6.04, 6.24]", "tokens": [
    "A", "LL", "THE", "YE", "LL", "OW", "LA", "M", "P", "S", "WOULD", "LIGHT", "UP", "HE
    ↵ ", "RE", "AND", "THERE", "THE", "S", "QUA", "LI", "D", "QUA", "R", "TER", "$
    S", "OF", "THE", "B", "RO", "TH", "EL", "S"]}

exp/test_wavs/1.wav
{"text": "IN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY CHILD WHOSE PLACE WAS ON THAT
    ↵ SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH THE RACE AN$
    DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN", "timestamps": "[2.44, 2.
    ↵ 64, 2.88, 3.16, 3.28, 3.48, 3.60, 3.80, 3.96, 4.12, 4.36, 4.52, 4.72, 4$92, 5.16, 5.44, 5.68, 6.04, 6.24, 6.48, 6.84, 7.08, 7.32, 7.56, 7.84, 8.12, 8.24, 8.32,
    ↵ 8.44, 8.60, 8.76, 8.88, 9.08, 9.28, 9.44, 9.56, 9.64, 9.76, 9.96, 10.0$96, 10.20, 10.40, 10.68, 10.76, 11.04, 11.20, 11.36, 11.60, 11.80, 12.00, 12.12, 12.28, 12.
    ↵ 32, 12.52, 12.72, 12.84, 12.96, 13.04, 13.24, 13.40, 13.60, 13.76, 13$96, 14.12, 14.24, 14.36, 14.52, 14.68, 14.76, 15.04, 15.28, 15.52, 15.76, 16.00, 16.16,
    ↵ 16.24, 16.32]", "tokens": [
    "IN", "WHICH", "MAN", "TH", "US", "P", "UN", "IS$",
    "ED", "HAD", "GIVE", "N", "HER", "A", "LOVE", "LY", "CHILD", "WHO", "SE", "PLACE", "WAS",
    "ON", "THAT", "SAME", "DIS", "HO", "N", "OUR", "ED", "BO", "S", "OM", "TO", "$CON", "NE", "C", "T", "HER", "P", "AR", "ENT", "FOR", "E", "VER", "WITH", "THE", "RA", "CE",
    ↵ "AND", "DE", "S", "C", "ENT", "OF", "MO", "R", "T", "AL", "S", "AND", "TO", "BE$",
    "FI", "N", "AL", "LY", "A", "B", "LESS", "ED", "SO", "UL", "IN", "HE", "A", "VE", "N"]}

num threads: 2
decoding method: modified_beam_search
max active paths: 4
Elapsed seconds: 1.522 s
Real time factor (RTF): 1.522 / 23.340 = 0.065

```

Hint: QUARTER -> QUARTERS

FOR EVER -> FOREVER

Python api

Decoding without hotwords

```
python python-api-examples/offline-decode-files.py \
--encoder exp/encoder-epoch-99-avg-1.onnx \
--decoder exp/decoder-epoch-99-avg-1.onnx \
--joiner exp/joiner-epoch-99-avg-1.onnx \
--decoding modified_beam_search \
--tokens exp/tokens.txt \
exp/test_wavs/0.wav exp/test_wavs/1.wav
```

The output is:

```
Started!
Done!
exp/test_wavs/0.wav
ALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS
-----
exp/test_wavs/1.wav
IN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY CHILD WHOSE PLACE WAS ON THAT SAME_
↪DISHONOURED BOSOM TO CONNECT HER PARENT FOR EVER WITH THE RACE AND DESCENT OF MORTALS_
↪AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
num_threads: 1
decoding_method: modified_beam_search
Wave duration: 23.340 s
Elapsed time: 2.546 s
Real time factor (RTF): 2.546/23.340 = 0.109
```

Decoding with hotwords

```
python python-api-examples/offline-decode-files.py \
--encoder exp/encoder-epoch-99-avg-1.onnx \
--decoder exp/decoder-epoch-99-avg-1.onnx \
--joiner exp/joiner-epoch-99-avg-1.onnx \
--decoding modified_beam_search \
--tokens exp/tokens.txt \
--modeling-unit bpe \
--bpe-vocab exp/bpe.vocab \
--hotwords-file hotwords_en.txt \
--hotwords-score 2.0 \
exp/test_wavs/0.wav exp/test_wavs/1.wav
```

The output is:

```
Started!
Done!
exp/test_wavs/0.wav
ALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTERS OF THE BROTHELS
```

(continues on next page)

(continued from previous page)

```
-----  
exp/test_wavs/1.wav  
IN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY CHILD WHOSE PLACE WAS ON THAT SAME  
DISHONoured BOSOM TO CONNECT HER PARENT FOREVER WITH THE RACE AND DESCENT OF MORTALS  
AND TO BE FINALLY A BLESSED SOUL IN HEAVEN  
-----  
num_threads: 1  
decoding_method: modified_beam_search  
Wave duration: 23.340 s  
Elapsed time: 2.463 s  
Real time factor (RTF): 2.463/23.340 = 0.106
```

Hint: QUARTER -> QUARTERSFOR EVER -> FOREVER

Modeling unit is cjkchar

Download the model

```
cd /path/to/sherpa-onnx  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
conformer-zh-stateless2-2023-05-23.tar.bz2  
tar xvf sherpa-onnx-conformer-zh-stateless2-2023-05-23.tar.bz2  
rm sherpa-onnx-conformer-zh-stateless2-2023-05-23.tar.bz2  
  
ln -s sherpa-onnx-conformer-zh-stateless2-2023-05-23 exp-zh
```

The hotwords_cn.txt contains:

C++ api

Decoding without hotwords

```
./build/bin/sherpa-onnx-offline \  
--encoder=exp-zh/encoder-epoch-99-avg-1.onnx \  
--decoder=exp-zh/decoder-epoch-99-avg-1.onnx \  
--joiner=exp-zh/joiner-epoch-99-avg-1.onnx \  
--tokens=exp-zh/tokens.txt \  
--decoding-method=modified_beam_search \  
exp-zh/test_wavs/3.wav exp-zh/test_wavs/4.wav exp-zh/test_wavs/5.wav exp-zh/test_  
wavs/6.wav
```

The output is:

```

/star-kw/kangwei/code/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./build/bin/
sherpa-onnx-offline --encoder=exp-zh/encoder-epoch-99-avg-1.onnx --decoder=exp-zh/
--decoder-epoch-99-avg-1.onnx --joiner=exp-zh/joiner-epoch-99-avg-1.onnx --tokens=exp-zh/
--tokens.txt --decoding-method=modified_beam_search exp-zh/test_wavs/3.wav exp-zh/test_
wavs/4.wav exp-zh/test_wavs/5.wav exp-zh/test_wavs/6.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename=
"exp-zh/encoder-epoch-99-avg-1.onnx", decoder_filename="exp-zh/decoder-epoch-99-avg-1.
onnx", joiner_filename="exp-zh/joiner-$poch-99-avg-1.onnx"), paraformer=OfflineParaformerModelConfig(model=""),
nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""), whisper=OfflineWhisperModelConfig
$encoder="", decoder="", language="", task="transcribe"), tdnn=OfflineTdnnModelConfig(model=""),
tokens="exp-zh/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="modified_beam_search",
max_active_paths=4, hotwords_file=, hotwords_score=1.5)
Creating recognizer ...
Started
Done!

exp-zh/test_wavs/3.wav
{"text":"","timestamps":[0.00, 0.16, 0.68, 1.32, 1.72, 2.08, 2.60, 2.88, 3.20, 3.52, 3.
92, 4.40, 4.68, 5.12, 5.44, 6.36, .96, 7.32],"tokens":["","","","","","","","","","","","","",
"","","","","","","","",""]}

exp-zh/test_wavs/4.wav
{"text":"","timestamps":[0.00, 0.20, 0.88, 1.36, 1.76, 2.08, 2.28, 2.68, 2.92, 3.16, 3.
44, 3.80],"tokens":["","","","","","","","","$",
"","","","","","",""]}

exp-zh/test_wavs/5.wav
{"text":"","timestamps":[0.00, 0.16, 0.88, 1.24, 1.64, 1.96, 2.76, 3.04, 3.32],"tokens
":["","","","","","","","","","",""]}

exp-zh/test_wavs/6.wav
{"text":"","timestamps":[0.00, 0.16, 0.80, 1.12, 1.44, 1.68, 1.92, 2.16, 2.36, 2.60, 2.
84, 3.12],"tokens":["","","","","","","","","$",
"","","","","","",""]}

num threads: 2
decoding method: modified_beam_search
max active paths: 4
Elapsed seconds: 1.883 s
Real time factor (RTF): 1.883 / 20.328 = 0.093

```

Decoding with hotwords

```

./build/bin/sherpa-onnx-offline \
--encoder=exp-zh/encoder-epoch-99-avg-1.onnx \
--decoder=exp-zh/decoder-epoch-99-avg-1.onnx \
--joiner=exp-zh/joiner-epoch-99-avg-1.onnx \

```

(continues on next page)

(continued from previous page)

```
--tokens=exp-zh/tokens.txt \
--decoding-method=modified_beam_search \
--modeling-unit=cjkchar \
--hotwords-file=hotwords_cn.txt \
--hotwords-score=2.0 \
exp-zh/test_wavs/3.wav exp-zh/test_wavs/4.wav exp-zh/test_wavs/5.wav exp-zh/test_wavs/
6.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename=
"exp-zh/encoder-epoch-99-avg-1.onnx", decoder_filename="exp-zh/decoder-epoch-99-avg-1.
onnx", joiner_filename="exp-zh/joiner-$poch-99-avg-1.onnx"), paraformer=OfflineParaformerModelConfig(model=""),
nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""), whisper=OfflineWhisperModelConfig
$encoder="", decoder="", language="", task="transcribe"), tdnn=OfflineTdn
nModelConfig(model=""), tokens="exp-zh/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
lm_config=OfflineLMConfig(model=""), scale=0.5), decoding_method="modified_beam_search", max_active_paths=4, hotwords_
file=hotwords_cn.txt, hotwords_score=2)
Creating recognizer ...
Started
Done!

exp-zh/test_wavs/3.wav
{"text": "", "timestamps": "[0.00, 0.16, 0.64, 1.28, 1.64, 2.04, 2.60, 2.88, 3.20, 3.52, 3.
92, 4.40, 4.68, 5.12, 5.44, 6.36, $.96, 7.32]", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", ""],
---}
exp-zh/test_wavs/4.wav
{"text": "", "timestamps": "[0.00, 0.12, 0.80, 1.36, 1.76, 2.08, 2.28, 2.68, 2.92, 3.16, 3.
44, 3.80]", "tokens": ["", "", "", "", "", "$", "", "", "", "", "", ""], ---}
exp-zh/test_wavs/5.wav
{"text": "", "timestamps": "[0.00, 0.12, 0.80, 1.24, 1.56, 1.96, 2.68, 3.08, 3.32]", "tokens": ["", "", "", "", "", "", "", "", ""], ---}
exp-zh/test_wavs/6.wav
{"text": "", "timestamps": "[0.00, 0.12, 0.80, 1.12, 1.44, 1.68, 1.92, 2.16, 2.36, 2.60, 2.
84, 3.12]", "tokens": ["", "", "", "", "", "$", "", "", "", "", ""], ---}

num threads: 2
decoding method: modified_beam_search
max active paths: 4
Elapsed seconds: 1.810 s
Real time factor (RTF): 1.810 / 20.328 = 0.089
```

Hint: ->

->

->
->

Python api

Decoding without hotwords

```
python python-api-examples/offline-decode-files.py \
--encoder exp-zh/encoder-epoch-99-avg-1.onnx \
--decoder exp-zh/decoder-epoch-99-avg-1.onnx \
--joiner exp-zh/joiner-epoch-99-avg-1.onnx \
--tokens exp-zh/tokens.txt \
--decoding-method modified_beam_search \
exp-zh/test_wavs/3.wav exp-zh/test_wavs/4.wav exp-zh/test_wavs/5.wav exp-zh/test_wavs/6.
˓wav
```

The output is:

```
Started!
Done!
exp-zh/test_wavs/3.wav

-----
exp-zh/test_wavs/4.wav

-----
exp-zh/test_wavs/5.wav

-----
exp-zh/test_wavs/6.wav

-----
num_threads: 1
decoding_method: modified_beam_search
Wave duration: 20.328 s
Elapsed time: 2.653 s
Real time factor (RTF): 2.653/20.328 = 0.131
```

Decoding with hotwords

```
python python-api-examples/offline-decode-files.py \
--encoder exp-zh/encoder-epoch-99-avg-1.onnx \
--decoder exp-zh/decoder-epoch-99-avg-1.onnx \
--joiner exp-zh/joiner-epoch-99-avg-1.onnx \
--tokens exp-zh/tokens.txt \
--decoding-method modified_beam_search \
--modeling-unit=cjkchar \
--hotwords-file hotwords_cn.txt \
--hotwords-score 2.0 \
exp-zh/test_wavs/3.wav exp-zh/test_wavs/4.wav exp-zh/test_wavs/5.wav exp-zh/test_wavs/6.wav
```

The output is:

```
Started!
Done!
exp-zh/test_wavs/3.wav

-----
exp-zh/test_wavs/4.wav

-----
exp-zh/test_wavs/5.wav

-----
exp-zh/test_wavs/6.wav

-----
num_threads: 1
decoding_method: modified_beam_search
Wave duration: 20.328 s
Elapsed time: 2.636 s
Real time factor (RTF): 2.636/20.328 = 0.130
```

Hint: ->

->
->
->

Modeling unit is cjkchar+bpe

Download the model

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
tar xvf sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
rm sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2

ln -s sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 exp-mixed
```

Export bpe.vocab if you can't find bpe.vocab in the model directory.

```
python script/export_bpe_vocab.py --bpe-model exp/bpe.model
```

The hotwords_mix.txt contains:

C++ api

Decoding without hotwords

```
./build/bin/sherpa-onnx \
--encoder=exp-mixed/encoder-epoch-99-avg-1.onnx \
--decoder=exp-mixed/decoder-epoch-99-avg-1.onnx \
--joiner=exp-mixed/joiner-epoch-99-avg-1.onnx \
--decoding-method=modified_beam_search \
--tokens=exp-mixed/tokens.txt \
exp-mixed/test_wavs/0.wav exp-mixed/test_wavs/2.wav
```

The output is:

```
/star-kw/kangwei/code/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./build/bin/
sherpa-onnx --encoder=exp-mixed/encoder-epoch-99-avg-1.onnx --decoder=exp-mixed/
--decoder-epoch-99-avg-1.onnx --joiner=exp-mixed/joiner-epoch-99-avg-1.onnx --decoding-
--method=modified_beam_search --tokens=exp-mixed/tokens.txt exp-mixed/test_wavs/0.wav
exp-mixed/test_wavs/2.wav
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
"exp-mixed/encoder-epoch-99-avg-1.onnx", decoder="exp-mixed/decoder-epoch-99-avg-1.onnx
", joiner="exp-mixed/joiner-epoch-99-avg-1.onnx"),_
paraformer=OnlineParaformerModelConfig(encoder="", decoder=""), tokens="exp-mixed/
tokens.txt", num_threads=1, debug=False, provider="cpu", model_type=""), lm_
config=OnlineLMConfig(model="", scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, hotwords_score=1.5, hotwords_file="", decoding_
method="modified_beam_search")

exp-mixed/test_wavs/0.wav
Elapsed seconds: 3, Real time factor (RTF): 0.3
MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW
{"is_final":false,"segment":0,"start_time":0.0,"text":" MONDAY TODAY IS LIBR THE DAY_
AFTER TOMORROW","timestamps":[0.64, 1.04, 1.60, 2.08, 2.20, 2.40, 4.16, 4.40, 4.88, 5.
56, 5.80, 6.16, 6.84, 7.12, 7.44, 8.04, 8.16, 8.24, 8.28, 9.04, 9.40, 9.64, 9.88],"_
"tokens":["","","",""," MO","N","DAY"," TO","DAY"," IS"," LI","B","R"," THE"," DAY","_
AFTER"," TO","M","OR","ROW","","","","","",""]}

exp-mixed/test_wavs/2.wav
Elapsed seconds: 1.7, Real time factor (RTF): 0.37
FREQUENTLY
{"is_final":false,"segment":0,"start_time":0.0,"text":" FREQUENTLY","timestamps":[0.00,_
0.40, 0.52, 0.96, 1.08, 1.28, 1.48, 1.68, 1.84, 2.00, 2.24, 2.36, 2.52, 2.68, 2.92, 3.
00, 3.12, 3.32, 3.64, 3.96, 4.36],"tokens":["","","","","","","","","","","","","","",""," F
","RE","QU","ENT","LY","","","",""]}
```

Decoding with hotwords

```
./build/bin/sherpa-onnx \
```

(continues on next page)

(continued from previous page)

```
--encoder=exp-mixed/encoder-epoch-99-avg-1.onnx \
--decoder=exp-mixed/decoder-epoch-99-avg-1.onnx \
--joiner=exp-mixed/joiner-epoch-99-avg-1.onnx \
--decoding-method=modified_beam_search \
--tokens=exp-mixed/tokens.txt \
--modeling-unit=cjkchar+bpe \
--bpe-vocab=exp/bpe.vocab \
--hotwords-file=hotwords_mix.txt \
--hotwords-score=2.0 \
exp-mixed/test_wavs/0.wav exp-mixed/test_wavs/2.wav
```

The output is:

```
/star-kw/kangwei/code/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./build/bin/
sherpa-onnx --encoder=exp-mixed/encoder-epoch-99-avg-1.onnx --decoder=exp-mixed/
--decoder-epoch-99-avg-1.onnx --joiner=exp-mixed/joiner-epoch-99-avg-1.onnx --decoding-
--method=modified_beam_search --tokens=exp-mixed/tokens.txt --tokens-type=cjkchar+bpe --
--bpe-model=exp-mixed/bpe.model --hotwords-file=hotwords_mix.txt --hotwords-score=2.0
exp-mixed/test_wavs/0.wav exp-mixed/test_wavs/2.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
"exp-mixed/encoder-epoch-99-avg-1.onnx", decoder="exp-mixed/decoder-epoch-99-avg-1.onnx
", joiner="exp-mixed/joiner-epoch-99-avg-1.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder=""), tokens="exp-mixed/
tokens.txt", num_threads=1, debug=False, provider="cpu", model_type=""), lm_
config=OnlineLMConfig(model="", scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, hotwords_score=2, hotwords_file="hotwords_mix.txt",
decoding_method="modified_beam_search")

exp-mixed/test_wavs/0.wav
Elapsed seconds: 3.2, Real time factor (RTF): 0.32
MONDAY TODAY IS THE DAY AFTER TOMORROW
{"is_final":false,"segment":0,"start_time":0.0,"text":" MONDAY TODAY IS THE DAY AFTER_
TOMORROW","timestamps":[0.64, 1.04, 1.60, 2.08, 2.20, 2.40, 4.16, 4.40, 4.88, 5.56, 5.
68, 6.00, 6.84, 7.12, 7.44, 8.04, 8.16, 8.24, 8.28, 9.04, 9.40, 9.64, 9.88],"tokens":[
 "", "", "", " MO", "N", "DAY", " TO", "DAY", " IS", "", "", " THE", " DAY", " AFTER", " TO", "M",
 "OR", "ROW", "", "", "", ""}]}

exp-mixed/test_wavs/2.wav
Elapsed seconds: 1.9, Real time factor (RTF): 0.4
FREQUENTLY
{"is_final":false,"segment":0,"start_time":0.0,"text":" FREQUENTLY","timestamps":[0.00,
0.40, 0.52, 0.96, 1.08, 1.28, 1.48, 1.68, 1.84, 2.00, 2.24, 2.36, 2.52, 2.68, 2.92, 3.
00, 3.12, 3.32, 3.64, 3.96, 4.36],"tokens":[" ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " F
", "RE", "QU", "ENT", "LY", " ", " ", ""]} 
```

Hint: LIBR ->

->

Python api

Decoding without hotwords

```
python python-api-examples/online-decode-files.py \
--encoderexp-mixed/encoder-epoch-99-avg-1.onnx \
--decoder exp-mixed/decoder-epoch-99-avg-1.onnx \
--joiner exp-mixed/joiner-epoch-99-avg-1.onnx \
--decoding-method modified_beam_search \
--tokens exp-mixed/tokens.txt
exp-mixed/test_wavs/0.wav exp-mixed/test_wavs/2.wav
```

The output is:

```
Started!
Done!
exp-mixed/test_wavs/0.wav
MONDAY TODAY IS LIBR THE DAY AFTER TOMORROW
-----
exp-mixed/test_wavs/2.wav
FREQUENTLY
-----
num_threads: 1
decoding_method: modified_beam_search
Wave duration: 14.743 s
Elapsed time: 3.052 s
Real time factor (RTF): 3.052/14.743 = 0.207
```

Decoding with hotwords

```
python python-api-examples/online-decode-files.py \
--encoder exp-mixed/encoder-epoch-99-avg-1.onnx \
--decoder exp-mixed/decoder-epoch-99-avg-1.onnx \
--joiner exp-mixed/joiner-epoch-99-avg-1.onnx \
--decoding-method modified_beam_search \
--tokens exp-mixed/tokens.txt \
--modeling-unit cjkchar+bpe \
--bpe-vocab exp-mixed/bpe.vocab \
--hotwords-file hotwords_mix.txt \
--hotwords-score 2.0 \
exp-mixed/test_wavs/0.wav exp-mixed/test_wavs/2.wav
```

The output is:

```
Started!
Done!
exp-mixed/test_wavs/0.wav
MONDAY TODAY IS THE DAY AFTER TOMORROW
-----
exp-mixed/test_wavs/2.wav
```

(continues on next page)

(continued from previous page)

```
FREQUENTLY
-----
num_threads: 1
decoding_method: modified_beam_search
Wave duration: 14.743 s
Elapsed time: 3.060 s
Real time factor (RTF): 3.060/14.743 = 0.208
```

Hint: LIBR ->

->

8.20 Keyword spotting

In this section, we describe how we implement the open vocabulary keyword spotting (aka customized keyword spotting) feature and how to use it in `sherpa-onnx`.

8.20.1 What is open vocabulary keyword spotting

Basically, an open vocabulary keyword spotting system is just like a tiny ASR system, but it can only decode words/phrases in the given keywords. For example, if the given keyword is HELLO WORLD, then the decoded result should be either HELLO WORLD or empty. As for open vocabulary (or customized), it means you can specify any keywords without re-training the model. For building a conventional keyword spotting systems, people need to prepare a lot of audio-text pairs for the selected keywords and the trained model can only be used to detect those selected keywords. While an open vocabulary keyword spotting system allows people using one system to detect different keywords, even the keywords might not be in the training data.

8.20.2 Decoder for open vocabulary keyword spotting

For now, we only implement a beam search decoder to make the system only trigger the given keywords (i.e. the model itself is actually a tiny ASR). To make it is able to balance between the triggered rate and false alarm, we introduce two parameters for each keyword, `boosting score` and `trigger threshold`. The `boosting score` works like the hotwords recognition, it help the paths containing keywords to survive beam search, the larger this score is the easier the corresponding keyword will be triggered, read [Hotwords \(Contextual biasing\)](#) for more details. The `trigger threshold` defines the minimum acoustic probability of decoded sequences (token sequences) that can be triggered, it is a float value between 0 to 1, the lower this threshold is the easier the corresponding keyword will be triggered.

8.20.3 Keywords file

The input keywords looks like (the keywords are HELLO WORLD, HI GOOGLE and HEY SIRI):

```
HE LL O WORLD :1.5 #0.35
HI GO O G LE :1.0 #0.25
HE Y S I RI
```

Each line contains a keyword, the first several tokens (separated by spaces) are encoded tokens of the keyword, the item starts with : is the **boosting score** and the item starts with # is the **trigger threshold**. Note: No spaces between : (or #) and the float value.

To get the tokens you need to use the command line tool in `sherpa-onnx` to convert the original keywords, you can see the usage as follows:

```
sherpa-onnx-cli text2token --help
Usage: sherpa-onnx-cli text2token [OPTIONS] INPUT OUTPUT

Options:

--text TEXT      Path to the input texts. Each line in the texts contains the
                ↵original phrase, it might also contain some extra items,
                ↵for example, the boosting score (startting with :), the triggering
                ↵threshold
                ↵(startting with #, only used in keyword spotting task) and the
                ↵original phrase (startting with @).
                ↵Note: extra items will be kept in the output.

example input 1 (tokens_type = ppinyin):
:2.0 #0.6 @
:3.5 @
#0.6 @
example output 1:
x iǎo ài tóng x ué :2.0 #0.6 @
n ī h ǎo w èn w èn :3.5 @
x iǎo y ì x iǎo y ì #0.6 @

example input 2 (tokens_type = bpe):
HELLO WORLD :1.5 #0.4
HI GOOGLE :2.0 #0.8
HEY SIRI #0.35
example output 2:
HE LL O WORLD :1.5 #0.4
HI GO O G LE :2.0 #0.8
HE Y S I RI #0.35

--tokens TEXT      The path to tokens.txt.
--tokens-type TEXT The type of modeling units, should be cjkchar, bpe, cjkchar+bpe,
                ↵fpinyin or ppinyin.
                ↵ppinyin           fpinyin means full pinyin, each cjkchar has a pinyin(with tone).
                ↵means partial pinyin, it splits pinyin into initial and final,
--bpe-model TEXT   The path to bpe.model. Only required when tokens-type is bpe or
                ↵cjkchar+bpe.
```

(continues on next page)

(continued from previous page)

--help	Show this message and exit.
--------	-----------------------------

Note: If the tokens-type is fpinyin or ppinyin, you MUST provide the original keyword (starting with @).

Note: If you install sherpa-onnx from sources (i.e. not by pip), you can use the alternative script in *scripts*, the usage is almost the same as the command line tool, read the help information by:

python3 scripts/text2token.py --help

8.20.4 How to use keyword spotting in sherpa-onnx

Currently, we provide command-line tool and android app for keyword spotting.

command-line tool

After installing `sherpa-onnx`, type `sherpa-onnx-keyword-spotter --help` for the help message.

8.20.5 Android application

You can find pre-built Android APKs for keyword spotting at

<https://k2-fsa.github.io/sherpa/onnx/kws/apk.html>

Here is a demo video (Note: It is in Chinese).

8.20.6 Pretrained models

You can find the pre-trained models in `sherpa-onnx-kws-pre-trained-models`.

8.21 Punctuation

This section introduces the models that `sherpa-onnx` supports for adding punctuations to text.

Hint: After getting text from speech using speech-to-text, you can use models from this section to add punctuations to text.

8.21.1 Pre-trained models

This section lists pre-trained models for adding punctuations to text.

You can find all models at the following URL:

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/punctuation-models>

sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12

This model is converted from

https://modelscope.cn/models/iic/punc_ct-transformer_zh-cn-common-vocab272727-pytorch/summary

and it supports both Chinese and English.

Hint: If you want to know how the model is converted to `sherpa-onnx`, please download it and you can find related scripts in the downloaded model directory.

In the following, we describe how to download and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/punctuation-models/sherpa-
→onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12.tar.bz2

# For Chinese users, you can also use the following mirror:
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/punctuation-models/
→sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12.tar.bz2

tar xvf sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12.tar.bz2
rm sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12.tar.bz2
```

You will find the following files after unzipping:

-rw-r--r--	1	fangjun	staff	1.4K	Apr 12	12:32	README.md
-rwxr-xr-x	1	fangjun	staff	1.6K	Apr 12	14:40	add-model-metadata.py
-rw-r--r--	1	fangjun	staff	810B	Apr 12	11:56	config.yaml
-rw-r--r--	1	fangjun	staff	42B	Apr 12	11:45	configuration.json
-rw-r--r--	1	fangjun	staff	281M	Apr 12	14:40	model.onnx
-rwxr-xr-x	1	fangjun	staff	745B	Apr 12	11:53	show-model-input-output.py
-rwxr-xr-x	1	fangjun	staff	4.9K	Apr 13	18:45	test.py
-rw-r--r--	1	fangjun	staff	4.0M	Apr 12	11:56	tokens.json

Only `model.onnx` is needed in `sherpa-onnx`. All other files are for your information about how the model is converted to `sherpa-onnx`.

C++ binary examples

After installing `sherpa-onnx`, you can use the following command to add punctuations to text:

```
./bin/sherpa-onnx-offline-punctuation \
--ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.
˓→onnx \
    
```

The output is given below:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./bin/
˓→sherpa-onnx-offline-punctuation --ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-
˓→en-vocab272727-2024-04-12/model.onnx ''  
  
OfflinePunctuationConfig(model=OfflinePunctuationModelConfig(ct_transformer=". /sherpa-
˓→onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.onnx", num_threads=1,_
˓→debug=False, provider="cpu"))  
Creating OfflinePunctuation ...  
Started  
Done  
Num threads: 1  
Elapsed seconds: 0.007 s  
Input text:  
Output text:
```

The second example is for text containing both Chinese and English:

```
./bin/sherpa-onnx-offline-punctuation \
--ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.
˓→onnx \
    "How are youthank you are you ok"
```

Its output is given below:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./bin/
˓→sherpa-onnx-offline-punctuation --ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-
˓→en-vocab272727-2024-04-12/model.onnx 'How are youthank you are you ok'  
  
OfflinePunctuationConfig(model=OfflinePunctuationModelConfig(ct_transformer=". /sherpa-
˓→onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.onnx", num_threads=1,_
˓→debug=False, provider="cpu"))  
Creating OfflinePunctuation ...  
Started  
Done  
Num threads: 1  
Elapsed seconds: 0.005 s  
Input text: How are youthank you are you ok  
Output text: How are youthank youare you ok
```

The last example is for text containing only English:

```
./bin/sherpa-onnx-offline-punctuation \
--ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.
˓→onnx \
```

(continues on next page)

(continued from previous page)

"The African blogosphere is rapidly expanding bringing more voices online in the form of commentaries opinions analyses rants and poetry"

Its output is given below:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./bin/
sherpa-onnx-offline-punctuation --ct-transformer=./sherpa-onnx-punct-ct-transformer-zh-
en-vocab272727-2024-04-12/model.onnx 'The African blogosphere is rapidly expanding,
bringing more voices online in the form of commentaries opinions analyses rants and
poetry'

OfflinePunctuationConfig(model=OfflinePunctuationModelConfig(ct_transformer=".sherpa-
onnx-punct-ct-transformer-zh-en-vocab272727-2024-04-12/model.onnx", num_threads=1,
debug=False, provider="cpu"))
Creating OfflinePunctuation ...
Started
Done
Num threads: 1
Elapsed seconds: 0.003 s
Input text: The African blogosphere is rapidly expanding bringing more voices online in
the form of commentaries opinions analyses rants and poetry
Output text: The African blogosphere is rapidly expanding bringing more voices online in
the form of commentaries opinions analyses rants and poetry
```

Python API examples

Please see

<https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/add-punctuation.py>

Huggingface space examples

Please see

- <https://huggingface.co/spaces/k2-fsa/generate-subtitles-for-videos>
- <https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>

Hint: For Chinese users, please visit the following mirrors:

- <https://hf-mirror.com/spaces/k2-fsa/generate-subtitles-for-videos>
 - <https://hf-mirror.com/spaces/k2-fsa/automatic-speech-recognition>
-

Video demos

The following video is in Chinese.

8.22 Audio tagging

This section introduces the models that `sherpa-onnx` supports for audio tagging, which aims to recognize sound events within an audio clip without its temporal localization.

8.22.1 Pre-trained models

This section lists pre-trained models for audio tagging.

You can find all models at the following URL:

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/audio-tagging-models>

sherpa-onnx-zipformer-small-audio-tagging-2024-04-15

This model is trained by <https://github.com/k2-fsa/icefall/pull/1421> using the dataset `audioset`.

In the following, we describe how to download and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/audio-tagging-models/sherpa-
˓→onnx-zipformer-small-audio-tagging-2024-04-15.tar.bz2

# For Chinese users, you can also use the following mirror:
wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/audio-tagging-models/
˓→sherpa-onnx-zipformer-small-audio-tagging-2024-04-15.tar.bz2

tar xvf sherpa-onnx-zipformer-small-audio-tagging-2024-04-15.tar.bz2
rm sherpa-onnx-zipformer-small-audio-tagging-2024-04-15.tar.bz2
```

You will find the following files after unzipping:

```
-rw-r--r--  1 fangjun  staff   243B Apr 15 16:14 README.md
-rw-r--r--  1 fangjun  staff   14K  Apr 15 16:14 class_labels_indices.csv
-rw-r--r--  1 fangjun  staff   26M  Apr 15 16:14 model.int8.onnx
-rw-r--r--  1 fangjun  staff   88M  Apr 15 16:14 model.onnx
drwxr-xr-x  15 fangjun  staff  480B Apr 15 16:14 test_wavs
```

C++ binary examples

Hint: You can find the binary executable file `sherpa-onnx-offline-audio-tagging` after installing `sherpa-onnx` either from source or using `pip install sherpa-onnx`.

Cat

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/1.wav
```

prints the following:

```
0: AudioEvent(name="Animal", index=72, prob=0.947886)
1: AudioEvent(name="Cat", index=81, prob=0.938876)
2: AudioEvent(name="Domestic animals, pets", index=73, prob=0.931975)
3: AudioEvent(name="Caterwaul", index=85, prob=0.178876)
4: AudioEvent(name="Meow", index=83, prob=0.176177)
Num threads: 1
Wave duration: 10.000
Elapsed seconds: 0.297 s
Real time factor (RTF): 0.297 / 10.000 = 0.030
```

Hint: By default, it outputs the top 5 events. The first event has the largest probability.

Whistle

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/2.wav
```

prints the following:

```
0: AudioEvent(name="Whistling", index=40, prob=0.804928)
1: AudioEvent(name="Music", index=137, prob=0.27548)
2: AudioEvent(name="Piano", index=153, prob=0.135418)
3: AudioEvent(name="Keyboard (musical)", index=152, prob=0.0580414)
4: AudioEvent(name="Musical instrument", index=138, prob=0.0400399)
Num threads: 1
Wave duration: 10.000
Elapsed seconds: 0.289 s
Real time factor (RTF): 0.289 / 10.000 = 0.029
```

Music

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/3.wav
```

prints the following:

```
0: AudioEvent(name="Music", index=137, prob=0.79673)
1: AudioEvent(name="A capella", index=255, prob=0.765521)
2: AudioEvent(name="Singing", index=27, prob=0.473899)
3: AudioEvent(name="Vocal music", index=254, prob=0.459337)
4: AudioEvent(name="Choir", index=28, prob=0.458174)
Num threads: 1
Wave duration: 10.000
Elapsed seconds: 0.279 s
Real time factor (RTF): 0.279 / 10.000 = 0.028
```

Laughter

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/4.wav
```

prints the following:

```

0: AudioEvent(name="Laughter", index=16, prob=0.929239)
1: AudioEvent(name="Snicker", index=19, prob=0.321969)
2: AudioEvent(name="Giggle", index=18, prob=0.149667)
3: AudioEvent(name="Inside, small room", index=506, prob=0.119332)
4: AudioEvent(name="Belly laugh", index=20, prob=0.100728)
Num threads: 1
Wave duration: 10.000
Elapsed seconds: 0.314 s
Real time factor (RTF): 0.314 / 10.000 = 0.031

```

Finger snapping

For the following test wave,

the command:

```

./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/5.wav

```

prints the following:

```

0: AudioEvent(name="Finger snapping", index=62, prob=0.690543)
1: AudioEvent(name="Slap, smack", index=467, prob=0.452133)
2: AudioEvent(name="Clapping", index=63, prob=0.179213)
3: AudioEvent(name="Sound effect", index=504, prob=0.101151)
4: AudioEvent(name="Whack, thwack", index=468, prob=0.0294559)
Num threads: 1
Wave duration: 8.284
Elapsed seconds: 0.225 s
Real time factor (RTF): 0.225 / 8.284 = 0.027

```

Baby cry

For the following test wave,

the command:

```

./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/6.wav

```

prints the following:

```
0: AudioEvent(name="Baby cry, infant cry", index=23, prob=0.912273)
1: AudioEvent(name="Crying, sobbing", index=22, prob=0.670927)
2: AudioEvent(name="Whimper", index=24, prob=0.187221)
3: AudioEvent(name="Inside, small room", index=506, prob=0.0314955)
4: AudioEvent(name="Sound effect", index=504, prob=0.0118726)
Num threads: 1
Wave duration: 8.719
Elapsed seconds: 0.232 s
Real time factor (RTF): 0.232 / 8.719 = 0.027
```

Smoke alarm

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/7.wav
```

prints the following:

```
0: AudioEvent(name="Smoke detector, smoke alarm", index=399, prob=0.781478)
1: AudioEvent(name="Beep, bleep", index=481, prob=0.641056)
2: AudioEvent(name="Buzzer", index=398, prob=0.218576)
3: AudioEvent(name="Fire alarm", index=400, prob=0.140145)
4: AudioEvent(name="Alarm", index=388, prob=0.012525)
Num threads: 1
Wave duration: 2.819
Elapsed seconds: 0.080 s
Real time factor (RTF): 0.080 / 2.819 = 0.028
```

Siren

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/8.wav
```

prints the following:

```

0: AudioEvent(name="Siren", index=396, prob=0.877108)
1: AudioEvent(name="Civil defense siren", index=397, prob=0.732789)
2: AudioEvent(name="Vehicle", index=300, prob=0.0113797)
3: AudioEvent(name="Inside, small room", index=506, prob=0.00537381)
4: AudioEvent(name="Outside, urban or manmade", index=509, prob=0.00261939)
Num threads: 1
Wave duration: 7.721
Elapsed seconds: 0.220 s
Real time factor (RTF): 0.220 / 7.721 = 0.028

```

Stream water

For the following test wave,

the command:

```

./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/10.wav

```

prints the following:

```

0: AudioEvent(name="Stream", index=292, prob=0.247785)
1: AudioEvent(name="Water", index=288, prob=0.231587)
2: AudioEvent(name="Gurgling", index=297, prob=0.170981)
3: AudioEvent(name="Trickle, dribble", index=450, prob=0.108859)
4: AudioEvent(name="Liquid", index=444, prob=0.0693812)
Num threads: 1
Wave duration: 7.837
Elapsed seconds: 0.212 s
Real time factor (RTF): 0.212 / 7.837 = 0.027

```

Meow

For the following test wave,

the command:

```

./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/11.wav

```

prints the following:

```
0: AudioEvent(name="Meow", index=83, prob=0.814944)
1: AudioEvent(name="Cat", index=81, prob=0.698858)
2: AudioEvent(name="Domestic animals, pets", index=73, prob=0.564516)
3: AudioEvent(name="Animal", index=72, prob=0.535303)
4: AudioEvent(name="Music", index=137, prob=0.105332)
Num threads: 1
Wave duration: 11.483
Elapsed seconds: 0.361 s
Real time factor (RTF): 0.361 / 11.483 = 0.031
```

Dog bark

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/12.wav
```

prints the following:

```
0: AudioEvent(name="Animal", index=72, prob=0.688237)
1: AudioEvent(name="Dog", index=74, prob=0.637803)
2: AudioEvent(name="Bark", index=75, prob=0.608597)
3: AudioEvent(name="Bow-wow", index=78, prob=0.515501)
4: AudioEvent(name="Domestic animals, pets", index=73, prob=0.495074)
Num threads: 1
Wave duration: 8.974
Elapsed seconds: 0.261 s
Real time factor (RTF): 0.261 / 8.974 = 0.029
```

Oink (pig)

For the following test wave,

the command:

```
./bin/sherpa-onnx-offline-audio-tagging \
--zipformer-model=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/model.int8.
˓→onnx \
--labels=./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/class_labels_indices.
˓→csv \
./sherpa-onnx-zipformer-small-audio-tagging-2024-04-15/test_wavs/13.wav
```

prints the following:

```

0: AudioEvent(name="Oink", index=94, prob=0.888416)
1: AudioEvent(name="Pig", index=93, prob=0.164295)
2: AudioEvent(name="Animal", index=72, prob=0.160802)
3: AudioEvent(name="Speech", index=0, prob=0.0276513)
4: AudioEvent(name="Snort", index=46, prob=0.0201952)
Num threads: 1
Wave duration: 9.067
Elapsed seconds: 0.261 s
Real time factor (RTF): 0.261 / 9.067 = 0.029

```

Python API examples

Please see

<https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/audio-tagging-from-a-file.py>

Huggingface space

You can try audio tagging with `sherpa-onnx` from within your browser by visiting the following URL:

<https://huggingface.co/spaces/k2-fsa/audio-tagging>

Note: For Chinese users, please use

<https://hf-mirror.com/spaces/k2-fsa/audio-tagging>

8.22.2 Android

You can find Android APKs for each model at the following page

<https://k2-fsa.github.io/sherpa/onnx/audio-tagging/apk.html>

Please follow *Android* to build Android APKs from source.

If you want to run audio tagging on your WearOS watches, please see *WearOS*.

8.22.3 WearOS

You can find APKs for WearOS of each model at the following page

<https://k2-fsa.github.io/sherpa/onnx/audio-tagging/apk-wearos.html>

Please follow *Android* to build APKs for WearOS from source.

If you want to run audio tagging on your Android phones, please see *Android*.

8.23 Spoken language identification

This section describes how to use `sherpa-onnx` for spoken language identification.

8.23.1 Pre-trained models

whisper

Currently, we support whisper multilingual models for spoken language identification.

Model type	Huggingface repo
<code>tiny</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny
<code>base</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-base
<code>small</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-small
<code>medium</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-medium

Hint: You can also download them from

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/asr-models>

In the following, we use the `tiny` model as an example. You can replace `tiny` with `base`, `small`, or `medium` and everything still holds.

Download the model

Please use the following commands to download the `tiny` model:

```
 wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
 ↵whisper-tiny.tar.bz2

# For Chinese users, please use
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
 ↵whisper-tiny.tar.bz2

tar xvf sherpa-onnx-whisper-tiny.tar.bz2
rm sherpa-onnx-whisper-tiny.tar.bz2
```

You should find the following files after unzipping:

```
-rw-r--r-- 1 fangjun staff 427B Jan 31 16:21 README.md
-rw-r--r-- 1 fangjun staff 19K Jan 31 16:21 export-onnx.py
-rw-r--r-- 1 fangjun staff 15B Jan 31 16:21 requirements.txt
-rw-r--r-- 1 fangjun staff 12K Jan 31 16:21 test.py
drwxr-xr-x 6 fangjun staff 192B Jan 31 16:22 test_wavs
-rw-r--r-- 1 fangjun staff 86M Jan 31 16:22 tiny-decoder.int8.onnx
-rw-r--r-- 1 fangjun staff 109M Jan 31 16:22 tiny-decoder.onnx
-rw-r--r-- 1 fangjun staff 12M Jan 31 16:22 tiny-encoder.int8.onnx
-rw-r--r-- 1 fangjun staff 36M Jan 31 16:22 tiny-encoder.onnx
-rw-r--r-- 1 fangjun staff 798K Jan 31 16:22 tiny-tokens.txt
```

Download test waves

Please use the following command to download test data:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/spoken-language-
identification-test-wavs.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/spoken-
language-identification-test-wavs.tar.bz2

tar xvf spoken-language-identification-test-wavs.tar.bz2
rm spoken-language-identification-test-wavs.tar.bz2
```

You can find the following test files after unzipping:

-rw-r--r--	1	fangjun	staff	222K	Mar 24 12:51	ar-arabic.wav
-rw-r--r--@	1	fangjun	staff	137K	Mar 24 13:09	bg-bulgarian.wav
-rw-r--r--	1	fangjun	staff	83K	Mar 24 13:07	cs-czech.wav
-rw-r--r--	1	fangjun	staff	112K	Mar 24 13:07	da-danish.wav
-rw-r--r--	1	fangjun	staff	199K	Mar 24 12:50	de-german.wav
-rw-r--r--	1	fangjun	staff	207K	Mar 24 13:06	el-greek.wav
-rw-r--r--	1	fangjun	staff	31K	Mar 24 12:45	en-english.wav
-rw-r--r--@	1	fangjun	staff	77K	Mar 24 12:23	es-spanish.wav
-rw-r--r--@	1	fangjun	staff	371K	Mar 24 12:21	fa-persian.wav
-rw-r--r--	1	fangjun	staff	136K	Mar 24 13:08	fi-finnish.wav
-rw-r--r--	1	fangjun	staff	112K	Mar 24 12:49	fr-french.wav
-rw-r--r--	1	fangjun	staff	179K	Mar 24 12:47	hi-hindi.wav
-rw-r--r--@	1	fangjun	staff	177K	Mar 24 12:29	hr-croatian.wav
-rw-r--r--	1	fangjun	staff	167K	Mar 24 12:53	id-indonesian.wav
-rw-r--r--	1	fangjun	staff	136K	Mar 24 12:54	it-italian.wav
-rw-r--r--	1	fangjun	staff	46K	Mar 24 12:44	ja-japanese.wav
-rw-r--r--@	1	fangjun	staff	122K	Mar 24 12:52	ko-korean.wav
-rw-r--r--	1	fangjun	staff	85K	Mar 24 12:54	nl-dutch.wav
-rw-r--r--@	1	fangjun	staff	241K	Mar 24 12:38	no-norwegian.wav
-rw-r--r--@	1	fangjun	staff	121K	Mar 24 12:35	po-polish.wav
-rw-r--r--	1	fangjun	staff	166K	Mar 24 12:48	pt-portuguese.wav
-rw-r--r--@	1	fangjun	staff	144K	Mar 24 12:33	ro-romanian.wav
-rw-r--r--	1	fangjun	staff	111K	Mar 24 12:51	ru-russian.wav
-rw-r--r--@	1	fangjun	staff	239K	Mar 24 12:40	sk-slovak.wav
-rw-r--r--	1	fangjun	staff	196K	Mar 24 13:01	sv-swedish.wav
-rw-r--r--	1	fangjun	staff	106K	Mar 24 13:14	ta-tamil.wav
-rw-r--r--	1	fangjun	staff	104K	Mar 24 13:02	tl-tagalog.wav
-rw-r--r--	1	fangjun	staff	76K	Mar 24 13:00	tr-turkish.wav
-rw-r--r--	1	fangjun	staff	188K	Mar 24 13:05	uk-ukrainian.wav
-rw-r--r--	1	fangjun	staff	181K	Mar 24 13:20	zh-chinese.wav

Test with Python APIs

After installing `sherpa-onnx` either from source or from using `pip install sherpa-onnx`, you can run:

```
python3 ./python-api-examples/spoken-language-identification.py \
--whisper-encoder ./sherpa-onnx-whisper-tiny/tiny-encoder.int8.onnx \
--whisper-decoder ./sherpa-onnx-whisper-tiny/tiny-decoder.onnx \
./spoken-language-identification-test-wavs/de-german.wav
```

You should see the following output:

```
2024-04-17 15:53:23,104 INFO [spoken-language-identification.py:158] File: ./spoken-
language-identification-test-wavs/de-german.wav
2024-04-17 15:53:23,104 INFO [spoken-language-identification.py:159] Detected language: de
2024-04-17 15:53:23,104 INFO [spoken-language-identification.py:160] Elapsed seconds: 0.
275
2024-04-17 15:53:23,105 INFO [spoken-language-identification.py:161] Audio duration in
seconds: 6.374
2024-04-17 15:53:23,105 INFO [spoken-language-identification.py:162] RTF: 0.275/6.374 =
0.043
```

Hint: You can find `spoken-language-identification.py` at

<https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/spoken-language-identification.py>

Android APKs

You can find pre-built Android APKs for spoken language identification at the following address:

<https://k2-fsa.github.io/sherpa/onnx/spoken-language-identification/apk.html>

Huggingface space

We provide a huggingface space for spoken language identification.

You can visit the following URL:

<http://huggingface.co/spaces/k2-fsa/spoken-language-identification>

Note: For Chinese users, you can use the following mirror:

<http://hf-mirror.com/spaces/k2-fsa/spoken-language-identification>

8.24 VAD

We support `silero-vad` for voice activity detection.

You can find pre-built Android APKs for VAD at:

<https://k2-fsa.github.io/sherpa/onnx/vad/apk.html>

APKs for VAD + speech recognition can be found at:

<https://k2-fsa.github.io/sherpa/onnx/vad/apk-asr.html>

8.25 Pre-trained models

The following table lists links for all pre-trained models.

Description	URL
Speech recognition (speech to text, ASR)	https://github.com/k2-fsa/sherpa-onné/releases/tag/asr-models
Text to speech (TTS)	https://github.com/k2-fsa/sherpa-onné/releases/tag/tts-models
VAD	https://github.com/k2-fsa/sherpa-onné/releases/download/asr-models/silero_vad.onné
Keyword spotting	https://github.com/k2-fsa/sherpa-onné/releases/tag/kws-models
Speech identification (Speaker ID)	https://github.com/k2-fsa/sherpa-onné/releases/tag/speaker-recognition-models
Spoken language identification (Language ID)	https://github.com/k2-fsa/sherpa-onné/releases/tag/asr-models (multilingual_whisper)
Audio tagging	https://github.com/k2-fsa/sherpa-onné/releases/tag/audio-tagging-models
Punctuation	https://github.com/k2-fsa/sherpa-onné/releases/tag/punctuation-models

In this section, we describe how to download and use all available pre-trained models for speech recognition.

Hint: Please install `git-lfs` before you continue.

Otherwise, you will be SAD later.

8.25.1 Online transducer models

This section lists available online transducer models.

Zipformer-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

sherpa-onnx-streaming-zipformer-korean-2024-06-16 (Korean)

Training code for this model can be found at <https://github.com/k2-fsa/icfall/pull/1651>. It supports only Korean.

PyTorch checkpoint with optimizer state can be found at <https://huggingface.co/johnBamma/icfall-asr-k sponspeech-pruned-transducer-stateless7-streaming-2024-06-12>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-korean-2024-06-16.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-korean-2024-06-16.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-korean-2024-06-16.tar.bz2
rm sherpa-onnx-streaming-zipformer-korean-2024-06-16.tar.bz2
ls -lh sherpa-onnx-streaming-zipformer-korean-2024-06-16
```

The output is given below:

```
$ ls -lh sherpa-onnx-streaming-zipformer-korean-2024-06-16

total 907104
-rw-r--r-- 1 fangjun staff 307K Jun 16 17:36 bpe.model
-rw-r--r-- 1 fangjun staff 2.7M Jun 16 17:36 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 11M Jun 16 17:36 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 121M Jun 16 17:36 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 279M Jun 16 17:36 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 2.5M Jun 16 17:36 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 9.8M Jun 16 17:36 joiner-epoch-99-avg-1.onnx
drwxr-xr-x 7 fangjun staff 224B Jun 16 17:36 test_wavs
-rw-r--r-- 1 fangjun staff 59K Jun 16 17:36 tokens.txt
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/encoder-epoch-99-avg-1.\
.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/decoder-epoch-99-avg-1.\
.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/joiner-epoch-99-avg-1.\
.onnx \
./sherpa-onnx-streaming-zipformer-korean-2024-06-16/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/
.tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/encoder-epoch-
99-avg-1.onnx --decoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/decoder-
epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/
joiner-epoch-99-avg-1.onnx ./sherpa-onnx-streaming-zipformer-korean-2024-06-16/test-
wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80, low_freq=20, high_freq=400, dither=0), model_
config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder="./sherpa-onnx-
streaming-zipformer-korean-2024-06-16/encoder-epoch-99-avg-1.onnx", decoder="./sherpa-
onnx-streaming-zipformer-korean-2024-06-16/decoder-epoch-99-avg-1.onnx", joiner="./
sherpa-onnx-streaming-zipformer-korean-2024-06-16/joiner-epoch-99-avg-1.onnx"), u
paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), wenet_
ctc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4), zipformer2_
ctc=OnlineZipformer2CtcModelConfig(model=""), nemo_ctc=OnlineNeMoCtcModelConfig(model=""
), tokens="./sherpa-onnx-streaming-zipformer-korean-2024-06-16/tokens.txt", threads=1, warm_up=0, debug=False, provider="cpu", model_type="", modeling_unit=
"cjkchar", bpe_vocab=""), lm_config=OnlineLmConfig(model="", scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), ctc_fst_decoder_
8.25. Pre-trained models 343
```

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-korean-2024-06-16/test_wavs/0.wav
Elapsed seconds: 0.56, Real time factor (RTF): 0.16

{ "text": " ", "tokens": [ " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " "],  
  "timestamps": [0.52, 0.96, 1.28, 1.44, 1.52, 1.84, 2.28, 2.48, 2.88, 3.04, 3.20, 3.44],  
  "ys_probs": [-1.701665, -0.208116, -0.527190, -1.777411, -1.853504, -0.768175, -1.  
  029222, -1.657714, -0.514807, -0.360788, -0.842238, -0.218511], "lm_probs": [],  
  "context_scores": [], "segment": 0, "words": [], "start_time": 0.00, "is_final": false}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/encoder-epoch-99-avg-1.  
int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/decoder-epoch-99-avg-1.  
onnx \
--joiner=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/joiner-epoch-99-avg-1.  
int8.onnx \
./sherpa-onnx-streaming-zipformer-korean-2024-06-16/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/  
tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/encoder-epoch-  
99-avg-1.int8.onnx --decoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/  
decoder-epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-korean-2024-06-  
16/joiner-epoch-99-avg-1.int8.onnx ./sherpa-onnx-streaming-zipformer-korean-2024-06-16/  
test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  
dim=80, low_freq=20, high_freq=400, dither=0), model_=  
config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder="./sherpa-onnx-  
streaming-zipformer-korean-2024-06-16/encoder-epoch-99-avg-1.int8.onnx", decoder="./  
sherpa-onnx-streaming-zipformer-korean-2024-06-16/decoder-epoch-99-avg-1.onnx", joiner=  
"./sherpa-onnx-streaming-zipformer-korean-2024-06-16/joiner-epoch-99-avg-1.int8.onnx"),  
paraformer=OnlineParaformerModelConfig(encoder="", decoder=""), wenet_=  
344tc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4) zipformer?  
ctc=OnlineZipformer2CtcModelConfig(model=""), nemo_ctc=OnlineNeMoCtcModelConfig(model="  
"), tokens="./sherpa-onnx-streaming-zipformer-korean-2024-06-16/tokens.txt", num_  
threads=1, warm_up=0, debug=False, provider="cpu", model_type="", modeling_unit=  
"cikchar", bpe_vocab=""), lm_config=OnlineLmConfig(model="", scale=0.5), endpoint
```

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-korean-2024-06-16/test_wavs/0.wav
Elapsed seconds: 0.37, Real time factor (RTF): 0.1

{ "text": " ", "tokens": [ " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " "],  

  ↵ "timestamps": [0.52, 0.96, 1.28, 1.44, 1.52, 1.84, 2.28, 2.48, 2.88, 3.04, 3.20, 3.44],  

  ↵ "ys_probs": [-1.750286, -0.241571, -0.621155, -1.862032, -1.977561, -0.789718, -1.  

  ↵ 002497, -1.627276, -0.554654, -0.443969, -0.852731, -0.218611], "lm_probs": [],  

  ↵ "context_scores": [], "segment": 0, "words": [], "start_time": 0.0, "is_final": false}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/encoder-epoch-99-avg-1.  

  ↵ int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/decoder-epoch-99-avg-1.  

  ↵ onnx \
--joiner=./sherpa-onnx-streaming-zipformer-korean-2024-06-16/joiner-epoch-99-avg-1.  

  ↵ int8.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12 (Chinese)

Training code for this model can be found at <https://github.com/k2-fsa/icefall/pull/1369>. It supports only Chinese.

Please refer to https://github.com/k2-fsa/icefall/tree/master/egs/multi_zh-hans/ASR#included-training-sets for the detailed information about the training data. In total, there are 14k hours of training data.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  

  ↵ streaming-zipformer-multi-zh-hans-2023-12-12.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  

  ↵ streaming-zipformer-multi-zh-hans-2023-12-12.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xf sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12.tar.bz2
rm sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12.tar.bz2
ls -lh sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12
```

The output is given below:

```
$ ls -lh sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12
total 668864
-rw-r--r-- 1 fangjun staff    28B Dec 12 18:59 README.md
-rw-r--r-- 1 fangjun staff   131B Dec 12 18:59 bpe.model
-rw-r--r-- 1 fangjun staff   1.2M Dec 12 18:59 decoder-epoch-20-avg-1-chunk-16-left-
↪ 128.int8.onnx
-rw-r--r-- 1 fangjun staff   4.9M Dec 12 18:59 decoder-epoch-20-avg-1-chunk-16-left-
↪ 128.onnx
-rw-r--r-- 1 fangjun staff   67M Dec 12 18:59 encoder-epoch-20-avg-1-chunk-16-left-
↪ 128.int8.onnx
-rw-r--r-- 1 fangjun staff  249M Dec 12 18:59 encoder-epoch-20-avg-1-chunk-16-left-
↪ 128.onnx
-rw-r--r-- 1 fangjun staff  1.0M Dec 12 18:59 joiner-epoch-20-avg-1-chunk-16-left-128.
↪ int8.onnx
-rw-r--r-- 1 fangjun staff  3.9M Dec 12 18:59 joiner-epoch-20-avg-1-chunk-16-left-128.
↪ onnx
drwxr-xr-x 8 fangjun staff  256B Dec 12 18:59 test_wavs
-rw-r--r-- 1 fangjun staff   18K Dec 12 18:59 tokens.txt
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/encoder-epoch-20-
↪ avg-1-chunk-16-left-128.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/decoder-epoch-20-
↪ avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-
↪ avg-1-chunk-16-left-128.onnx \
./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/test_wavs/DEV_T000000000000.
↪ wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-  
12/tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/  
encoder-epoch-20-avg-1-chunk-16-left-128.onnx --decoder=./sherpa-onnx-streaming-  
zipformer-multi-zh-hans-2023-12-12/decoder-epoch-20-avg-1-chunk-16-left-128.onnx --  
joiner=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-avg-  
1-chunk-16-left-128.onnx ./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/  
test_wavs/DEV_T0000000000.wav  
  
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=  
"./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/encoder-epoch-20-avg-1-  
chunk-16-left-128.onnx", decoder="./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-  
12-12/decoder-epoch-20-avg-1-chunk-16-left-128.onnx", joiner="./sherpa-onnx-streaming-  
zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-avg-1-chunk-16-left-128.onnx"),  
paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), wenet_  
ctc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4), tokens="./  
sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/tokens.txt", num_threads=1,  
debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",  
scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_  
nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),  
rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_  
utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_  
silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,  
hotwords_score=1.5, hotwords_file="", decoding_method="greedy_search")  
./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/test_wavs/DEV_T0000000000.wav  
Elapsed seconds: 0.65, Real time factor (RTF): 0.12  
  
{"is_final":false, "segment":0, "start_time":0.00, "text": " ", "timestamps": [0.32, 0.  
64, 0.76, 0.84, 1.08, 1.24, 1.96, 2.04, 2.24, 2.36, 2.56, 2.68, 2.80, 3.28, 3.40, 3.60,  
3.72, 3.84, 3.96, 4.04, 4.16, 4.28, 4.36, 4.60, 4.72], "tokens":[" ", "", "", "", "",  
"", "", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/encoder-epoch-20-
 avg-1-chunk-16-left-128.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/decoder-epoch-20-
 avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-
 avg-1-chunk-16-left-128.int8.onnx \
./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/test_wavs/DEV_T000000000000.
 wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-
12/tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/
encoder-epoch-20-avg-1-chunk-16-left-128.int8.onnx --decoder=./sherpa-onnx-streaming-
zipformer-multi-zh-hans-2023-12-12/decoder-epoch-20-avg-1-chunk-16-left-128.onnx --
joiner=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-avg-
1-chunk-16-left-128.int8.onnx ./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-
12/test_wavs/DEV_T000000000000.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
"./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/encoder-epoch-20-avg-1-
chunk-16-left-128.int8.onnx", decoder="./sherpa-onnx-streaming-zipformer-multi-zh-hans-
2023-12-12/decoder-epoch-20-avg-1-chunk-16-left-128.onnx", joiner="./sherpa-onnx-
streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-avg-1-chunk-16-left-128.
int8.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder=""), wenet_
ctc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4), tokens="./
sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/tokens.txt", num_threads=1,
debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",
scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
hotwords_score=1.5, hotwords_file="", decoding_method="greedy_search") (continues on next page)
```

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/test_wavs/DEV_T0000000000.wav
Elapsed seconds: 0.5, Real time factor (RTF): 0.088

{"is_final":false, "segment":0, "start_time":0.00, "text": " ", "timestamps": [0.32, 0.
˓→64, 0.76, 0.84, 1.04, 1.24, 1.96, 2.04, 2.24, 2.36, 2.56, 2.68, 2.88, 3.28, 3.40, 3.60,
˓→3.72, 3.84, 3.96, 4.04, 4.16, 4.28, 4.36, 4.60, 4.72], "tokens":[" ", "", "", "", "", "",
˓→"", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""}]}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/encoder-epoch-20-
˓→avg-1-chunk-16-left-128.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/decoder-epoch-20-
˓→avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-multi-zh-hans-2023-12-12/joiner-epoch-20-
˓→avg-1-chunk-16-left-128.int8.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

pkufool/icefall-asr-zipformer-streaming-wenetspeech-20230615 (Chinese)

This model is from

<https://huggingface.co/pkufool/icefall-asr-zipformer-streaming-wenetspeech-20230615>

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/1130>.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-
˓→zipformer-streaming-wenetspeech-20230615.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-
˓→zipformer-streaming-wenetspeech-20230615.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf icefall-asr-zipformer-streaming-wenetspeech-20230615.tar.bz2  
rm icefall-asr-zipformer-streaming-wenetspeech-20230615.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
icefall-asr-zipformer-streaming-wenetspeech-20230615 fangjun$ ls -lh exp/*chunk-16-left-  
↳ 128.*onnx  
-rw-r--r-- 1 fangjun staff 11M Jun 26 15:42 exp/decoder-epoch-12-avg-4-chunk-16-  
↳ left-128.int8.onnx  
-rw-r--r-- 1 fangjun staff 12M Jun 26 15:42 exp/decoder-epoch-12-avg-4-chunk-16-  
↳ left-128.onnx  
-rw-r--r-- 1 fangjun staff 68M Jun 26 15:42 exp/encoder-epoch-12-avg-4-chunk-16-  
↳ left-128.int8.onnx  
-rw-r--r-- 1 fangjun staff 250M Jun 26 15:43 exp/encoder-epoch-12-avg-4-chunk-16-  
↳ left-128.onnx  
-rw-r--r-- 1 fangjun staff 2.7M Jun 26 15:42 exp/joiner-epoch-12-avg-4-chunk-16-left-  
↳ 128.int8.onnx  
-rw-r--r-- 1 fangjun staff 11M Jun 26 15:42 exp/joiner-epoch-12-avg-4-chunk-16-left-  
↳ 128.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx \  
--tokens=./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_char/tokens.  
↳ txt \  
--encoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-  
↳ avg-4-chunk-16-left-128.onnx \  
--decoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-  
↳ avg-4-chunk-16-left-128.onnx \  
--joiner=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-  
↳ avg-4-chunk-16-left-128.onnx \  
.icefall-asr-zipformer-streaming-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx --tokens=./icefall-asr-zipformer-streaming-wenetspeech-20230615/
↳ data/lang_char/tokens.txt --encoder=./icefall-asr-zipformer-streaming-wenetspeech-
↳ 20230615/exp/encoder-epoch-12-avg-4-chunk-16-left-128.onnx --decoder=./icefall-asr-
↳ zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-avg-4-chunk-16-left-128.
↳ onnx --joiner=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-
↳ 12-avg-4-chunk-16-left-128.onnx ./icefall-asr-zipformer-streaming-wenetspeech-20230615/
↳ test_wavs/DEV_T0000000000.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
↳ dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./icefall-asr-
↳ zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-avg-4-chunk-16-left-128.
↳ onnx", decoder_filename="./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/
↳ decoder-epoch-12-avg-4-chunk-16-left-128.onnx", joiner_filename="./icefall-asr-
↳ zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-avg-4-chunk-16-left-128.
↳ onnx", tokens="./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_char/
↳ tokens.txt", num_threads=2, provider="cpu", debug=False), lm_
↳ config=OnlineLMConfig(model="", scale=0.5), endpoint_
↳ config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
↳ silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
↳ min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
↳ nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
↳ endpoint=True, max_active_paths=4, decoding_method="greedy_search")
./icefall-asr-zipformer-streaming-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav
Elapsed seconds: 0.56, Real time factor (RTF): 0.1

{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":[0.36, 0.64, 0.76,
↳ 0.88, 1.08, 1.24, 1.96, 2.04, 2.24, 2.36, 2.56, 2.68, 2.92, 3.32, 3.40, 3.64, 3.72, 3.
↳ 84, 3.96, 4.04, 4.16, 4.32, 4.48, 4.64, 4.72],"tokens":["","","","","","","","","","","","","",
↳ "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_char/tokens.
↳ txt \
--encoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-
↳ avg-4-chunk-16-left-128.int8.onnx \
--decoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-
↳ avg-4-chunk-16-left-128.onnx \
(continues on next page)
```

(continued from previous page)

```
--joiner=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-
→avg-4-chunk-16-left-128.int8.onnx \
./icefall-asr-zipformer-streaming-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./icefall-asr-zipformer-streaming-wenetspeech-20230615/
data/lang_char/tokens.txt --encoder=./icefall-asr-zipformer-streaming-wenetspeech-
20230615/exp/encoder-epoch-12-avg-4-chunk-16-left-128.int8.onnx --decoder=./icefall-
asr-zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-avg-4-chunk-16-left-
128.onnx --joiner=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-
epoch-12-avg-4-chunk-16-left-128.int8.onnx ./icefall-asr-zipformer-streaming-
wenetspeech-20230615/test_wavs/DEV_T0000000000.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./icefall-asr-
zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-avg-4-chunk-16-left-128.
int8.onnx", decoder_filename="./icefall-asr-zipformer-streaming-wenetspeech-20230615/
exp/decoder-epoch-12-avg-4-chunk-16-left-128.onnx", joiner_filename="./icefall-asr-
zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-avg-4-chunk-16-left-128.
int8.onnx", tokens="./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_
char/tokens.txt", num_threads=2, provider="cpu", debug=False), lm_
config=OnlineLMConfig(model="", scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, decoding_method="greedy_search")
./icefall-asr-zipformer-streaming-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav
Elapsed seconds: 0.38, Real time factor (RTF): 0.068

{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":"[0.40, 0.52, 0.72,
→ 0.84, 1.08, 1.24, 1.48, 1.92, 2.00, 2.24, 2.32, 2.48, 2.68, 2.80, 3.00, 3.28, 3.36, 3.
→60, 3.72, 3.84, 3.92, 4.00, 4.16, 4.28, 4.36, 4.64, 4.68, 5.00]","tokens":["","","","","",
→","","","","","","","","","","","","","","","","","","","","","","","","","","","",""]}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./icefall-asr-zipformer-streaming-wenetspeech-20230615/data/lang_char/tokens.
↪txt \
--encoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/encoder-epoch-12-
↪avg-4-chunk-16-left-128.int8.onnx \
--decoder=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/decoder-epoch-12-
↪avg-4-chunk-16-left-128.onnx \
--joiner=./icefall-asr-zipformer-streaming-wenetspeech-20230615/exp/joiner-epoch-12-
↪avg-4-chunk-16-left-128.int8.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-06-26 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-streaming-zipformer-2023-05-17>

which supports only English as it is trained on the *LibriSpeech* corpus.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/1058>.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪streaming-zipformer-en-2023-06-26.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪streaming-zipformer-en-2023-06-26.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-26.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes below.

```
-rw-r--r-- 1 1001 127 240K Apr 23 06:45 bpe.model
-rw-r--r-- 1 1001 127 1.3M Apr 23 06:45 decoder-epoch-99-avg-1-chunk-16-left-128.int8.
↪onnx
-rw-r--r-- 1 1001 127 2.0M Apr 23 06:45 decoder-epoch-99-avg-1-chunk-16-left-128.onnx
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 1001 127 68M Apr 23 06:45 encoder-epoch-99-avg-1-chunk-16-left-128.int8.
→ onnx
-rw-r--r-- 1 1001 127 250M Apr 23 06:45 encoder-epoch-99-avg-1-chunk-16-left-128.onnx
-rwxr-xr-x 1 1001 127 814 Apr 23 06:45 export-onnx-zipformer-online.sh
-rw-r--r-- 1 1001 127 254K Apr 23 06:45 joiner-epoch-99-avg-1-chunk-16-left-128.int8.
→ onnx
-rw-r--r-- 1 1001 127 1003K Apr 23 06:45 joiner-epoch-99-avg-1-chunk-16-left-128.onnx
-rw-r--r-- 1 1001 127 216 Apr 23 06:45 README.md
drwxr-xr-x 2 1001 127 4.0K Apr 23 06:45 test_wavs
-rw-r--r-- 1 1001 127 5.0K Apr 23 06:45 tokens.txt
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-
→ 16-left-128.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-
→ 16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-
→ 16-left-128.onnx \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
→ build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.
→ txt --encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-
→ chunk-16-left-128.onnx --decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/
→ decoder-epoch-99-avg-1-chunk-16-left-128.onnx --joiner=./sherpa-onnx-streaming-
→ zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-16-left-128.onnx ./sherpa-onnx-
→ streaming-zipformer-en-2023-06-26/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
→ dim=80, model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
→ streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.onnx",
→ decoder_filename="./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder(continues on next page)
→ epoch-99-avg-1-chunk-16-left-128.onnx", joiner_filename="./sherpa-onnx-streaming-zipformer-en-2023-
→ 06-26/joiner-epoch-99-avg-1-chunk-16-left-128.onnx", tokens="./sheChapter 8. sherpa-onnx
→ zipformer-en-2023-06-26/tokens.txt", num_threads=2, provider="cpu", debug=False), lm_
→ config=OnlineLMConfig(model="", scale=0.5), endpoint_
→ config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
→ silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True
```

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
Elapsed seconds: 0.51, Real time factor (RTF): 0.077
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS
{"is_final":false,"segment":0,"start_time":0.0,"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS","timestamps": "[0.68, 1.04, 1.16, 1.24, 1.60, 1.76, 1.80, 1.92, 2.04, 2.24, 2.32, 2.36, 2.52, 2.68, 2.72, 2.80, 2.92, 3.12, 3.40, 3.64, 3.76, 3.92, 4.12, 4.48, 4.68, 4.72, 4.84, 5.00, 5.20, 5.24, 5.36, 5.40, 5.64, 5.76, 5.92, 5.96, 6.08, 6.24, 6.52]", "tokens": [" AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE", " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", "L", "IGHT", " UP", " HE", "RE", " AND", " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA", "R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-16-left-128.int8.onnx \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-16-left-128.int8.onnx \
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.int8.onnx", decoder_filename="./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-16-left-128.onnx", joiner_filename="./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-16-left-128.int8.onnx", tokens="./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt", num_threads=2, provider="cpu", debug=False), lm_config=OnlineLMConfig(model="", scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=1.2, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4, decoding_method="greedy_search")
```

(continues on next page)

8.25. Pre-trained models 355

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-en-2023-06-26/test_wavs/0.wav
Elapsed seconds: 0.41, Real time factor (RTF): 0.062
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS
{"is_final":false,"segment":0,"start_time":0.0,"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS","timestamps": "[0.68, 1.04, 1.16, 1.24, 1.60, 1.76, 1.80, 1.92, 2.04, 2.24, 2.32, 2.36, 2.52, 2.68, 2.72, 2.80, 2.92, 3.12, 3.40, 3.64, 3.76, 3.92, 4.12, 4.48, 4.68, 4.72, 4.84, 5.00, 5.20, 5.24, 5.36, 5.44, 5.64, 5.76, 5.92, 5.96, 6.08, 6.24, 6.52]", "tokens": [" AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE", " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA", "R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/encoder-epoch-99-avg-1-chunk-16-left-128.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-26/decoder-epoch-99-avg-1-chunk-16-left-128.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-26/joiner-epoch-99-avg-1-chunk-16-left-128.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

cskuangji/sherpa-onnx-streaming-zipformer-en-2023-06-21 (English)

This model is converted from

<https://huggingface.co/marcoyang/icefall-libri-giga-pruned-transducer-stateless7-streaming-2023-04-04>

which supports only English as it is trained on the LibriSpeech and GigaSpeech corpus.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/984>.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-2023-06-21.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-2023-06-21.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-en-2023-06-21.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-06-21.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-en-2023-06-21 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 1.2M Jun 21 15:34 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Jun 21 15:34 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 179M Jun 21 15:36 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 337M Jun 21 15:37 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 253K Jun 21 15:34 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Jun 21 15:34 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.  
txt --encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.  
onnx --decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.  
onnx --joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.  
onnx ./sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav  
  
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-  
streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.onnx", decoder_filename="./  
sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.onnx", joiner_  
filename="./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.onnx",  
tokens="./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.txt", num_threads=2,  
provider="cpu", debug=False), lm_config=OnlineLMConfig(model="", scale=0.5), endpoint_  
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_  
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,  
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_  
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_  
endpoint=True, max_active_paths=4, decoding_method="greedy_search")  
./sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav  
Elapsed seconds: 0.5, Real time factor (RTF): 0.076  
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID  
QUARTER OF THE BROTHELS  
{"is_final":false,"segment":0,"start_time":0.0,"text":" AFTER EARLY NIGHTFALL THE YELLOW  
LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS","timestamps":  
"[0.64, 1.00, 1.12, 1.20, 1.60, 1.76, 1.84, 1.96, 2.08, 2.24, 2.36, 2.40, 2.60, 2.72,  
2.80, 2.88, 3.00, 3.20, 3.44, 3.68, 3.76, 3.96, 4.24, 4.52, 4.72, 4.76, 4.88, 5.04, 5.  
24, 5.28, 5.36, 5.48, 5.64, 5.76, 5.92, 5.96, 6.04, 6.24, 6.36]", "tokens":[" AFTER", " E  

```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx \  
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.txt \  
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.int8.  
onnx \  
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.onnx \  
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.int8.  
onnx \  
.sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```

/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  

build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.  

txt --encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.  

int8.onnx --decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-  

avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-  

avg-1.int8.onnx ./sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  

dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-  

streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.int8.onnx", decoder_filename=  

"./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.onnx", joiner_  

filename="./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.int8.  

onnx", tokens="./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.txt", num_  

threads=2, provider="cpu", debug=False), lm_config=OnlineLMConfig(model="", scale=0.5),  

endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_  

trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_  

nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0),  

rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_  

utterance_length=20)), enable_endpoint=True, max_active_paths=4, decoding_method=  

"greedy_search")  

./sherpa-onnx-streaming-zipformer-en-2023-06-21/test_wavs/0.wav  

Elapsed seconds: 0.41, Real time factor (RTF): 0.062  

AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID  

QUARTER OF THE BROTHELS  

{"is_final":false,"segment":0,"start_time":0.0,"text":" AFTER EARLY NIGHTFALL THE YELLOW  

LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BROTHELS","timestamps":  

"[0.64, 1.00, 1.12, 1.20, 1.60, 1.76, 1.80, 1.96, 2.08, 2.24, 2.36, 2.40, 2.60, 2.72,  

2.80, 2.88, 3.00, 3.20, 3.44, 3.68, 3.76, 3.96, 4.24, 4.52, 4.72, 4.76, 4.88, 5.04, 5.  

24, 5.28, 5.36, 5.48, 5.64, 5.76, 5.92, 5.96, 6.04, 6.24, 6.36]", "tokens":[" AFTER", " E  

", "AR", "LY", " NIGHT", "F", "A", "LL", " THE", " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", "  

LIGHT", " UP", " HE", "RE", " AND", " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA", "R", "TER",  

" OF", " THE", " B", "RO", "TH", "EL", "S"]}
```

Real-time speech recognition from a microphone

```

cd /path/to/sherpa-onnx  

./build/bin/sherpa-onnx-microphone \  

--tokens=./sherpa-onnx-streaming-zipformer-en-2023-06-21/tokens.txt \  

--encoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/encoder-epoch-99-avg-1.onnx \  

--decoder=./sherpa-onnx-streaming-zipformer-en-2023-06-21/decoder-epoch-99-avg-1.onnx \  

--joiner=./sherpa-onnx-streaming-zipformer-en-2023-06-21/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-streaming-zipformer-en-2023-02-21 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-pruned-transducer-stateless7-streaming-2022-12-29>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless7_streaming

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

GitHub

ModelScope

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-2023-02-21.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-2023-02-21.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-en-2023-02-21.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-2023-02-21.tar.bz2
```

```
cd /path/to/sherpa-onnx

GIT_LFS_SKIP_SMUDGE=1 git clone https://www.modelscope.cn/pkufool/sherpa-onnx-streaming-
˓→zipformer-en-2023-02-21.git
cd sherpa-onnx-streaming-zipformer-en-2023-02-21
git lfs pull --include "*.onnx"
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-en-2023-02-21$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root  1.3M Mar 31 23:06 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root  2.0M Feb 21 20:51 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 180M Mar 31 23:07 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 338M Feb 21 20:51 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 254K Mar 31 23:06 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 1003K Feb 21 20:51 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
˓streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.onnx", decoder_filename="./
˓sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx", joiner_
˓filename="./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.onnx",_
˓tokens="./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt", num_threads=2,_
˓debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
˓nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),_
˓rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
˓utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,_
˓decoding_method="greedy_search")
2023-04-01 06:16:29.128344485 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 604840, index: 15, mask: {16, 52, }, error code: 22_
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
2023-04-01 06:16:29.128346568 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 604841, index: 16, mask: {17, 53, }, error code: 22_
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav
wav duration (s): 6.625
Started
Done!
Recognition result for ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav:
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
˓QUARTER OF THE BROTHELS
```

(continues on next page)

(continued from previous page)

```
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.825 s
Real time factor (RTF): 0.825 / 6.625 = 0.125
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.int8.
˓→onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.int8.
˓→onnx \
./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
˓→streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.int8.onnx", decoder_filename=
˓→"./sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx", joiner_
˓→filename="./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.int8.
˓→onnx", tokens="./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt", num_
˓→threads=2, debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_
˓→contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),_
˓→rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
˓→utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓→silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,_
˓→decoding_method="greedy_search")
2023-04-01 06:18:47.466564998 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓→setaffinity_np failed for thread: 604880, index: 15, mask: {16, 52, }, error code: 22
˓→error msg: Invalid argument. Specify the number of threads explicitly so the affinity
˓→is not set.
2023-04-01 06:18:47.466566863 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓→setaffinity_np failed for thread: 604881, index: 16, mask: {17, 53, }, error code: 22
˓→error msg: Invalid argument. Specify the number of threads explicitly so the affinity
˓→is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav
wav duration (s): 6.625
Started
Done!
```

(continues on next page)

(continued from previous page)

```
Recognition result for ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav:
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
→ QUARTER OF THE BROTHELS
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.633 s
Real time factor (RTF): 0.633 / 6.625 = 0.096
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

cskuangfj/sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20 (Bilingual, Chinese + English)

This model is converted from

<https://huggingface.co/csukuangfj/k2fsa-zipformer-chinese-english-mixed>

which supports both Chinese and English. The model is contributed by the community and is trained on tens of thousands of some internal dataset.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→ streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→ streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
rm sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root 13M Mar 31 21:11 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 14M Feb 20 20:13 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 174M Mar 31 21:11 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 315M Feb 20 20:13 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 3.1M Mar 31 21:11 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 13M Feb 20 20:13 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-
99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-
99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-
avg-1.onnx \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
streaming-zipformer-en-2023-02-21/encoder-epoch-99-avg-1.onnx", decoder_filename="./
sherpa-onnx-streaming-zipformer-en-2023-02-21/decoder-epoch-99-avg-1.onnx", joiner_
filename="./sherpa-onnx-streaming-zipformer-en-2023-02-21/joiner-epoch-99-avg-1.onnx",_
tokens="./sherpa-onnx-streaming-zipformer-en-2023-02-21/tokens.txt", num_threads=2,_
debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=True,_
nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(
must_contain_nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(
must_contain_nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,_
decoding_method="greedy_search")
```

(continued from previous page)

```

2023-04-01 06:22:23.030317206 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 604942, index: 16, mask: {17, 53, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
2023-04-01 06:22:23.030315351 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 604941, index: 15, mask: {16, 52, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav
wav duration (s): 6.625
Started
Done!
Recognition result for ./sherpa-onnx-streaming-zipformer-en-2023-02-21/test_wavs/0.wav:
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID
↳ QUARTER OF THE BROTHELS
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.815 s
Real time factor (RTF): 0.815 / 6.625 = 0.123

```

int8

The following code shows how to use fp32 models to decode a wave file:

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-
↳ 99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-
↳ 99-avg-1.int8.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-
↳ avg-1.int8.onnx \
./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.wav

```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-99-avg-1.int8.onnx",
decoder_filename="./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-
epoch-99-avg-1.int8.onnx", joiner_filename="./sherpa-onnx-streaming-zipformer-
bilingual-zh-en-2023-02-20/joiner-epoch-99-avg-1.int8.onnx", tokens="./sherpa-onnx-
streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt", num_threads=2,
debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
nonsense=False, min_trailing_silence=2.4, min_utterance_length=0),
rule2=EndpointRule(must_contain_nonsense=True, min_trailing_silence=1.2, min_
utterance_length=0), rule3=EndpointRule(must_contain_nonsense=False, min_trailing_
silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
decoding_method="greedy_search")
2023-04-01 06:24:10.503505750 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
setaffinity_np failed for thread: 604982, index: 16, mask: {17, 53, }, error code: 22
error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
is not set.
2023-04-01 06:24:10.503503942 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
setaffinity_np failed for thread: 604981, index: 15, mask: {16, 52, }, error code: 22
error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_wavs/1.
wav
wav duration (s): 5.100
Started
Done!
Recognition result for ./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/test_-
wavs/1.wav:
ALWAYS ALWAYS
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.551 s
Real time factor (RTF): 0.551 / 5.100 = 0.108
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/encoder-epoch-
99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/decoder-epoch-
99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-bilingual-zh-en-2023-02-20/joiner-epoch-99-
avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

shaojieli/sherpa-onnx-streaming-zipformer-fr-2023-04-14 (French)

This model is converted from

<https://huggingface.co/shaojieli/icefall-asr-commonvoice-fr-pruned-transducer-stateless7-streaming-2023-04-02>

which supports only French as it is trained on the [CommonVoice](#) corpus. In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-fr-2023-04-14.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-fr-2023-04-14.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-fr-2023-04-14.tar.bz2
rm sherpa-onnx-streaming-zipformer-fr-2023-04-14.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-fr-2023-04-14 shaojieli$ ls -lh *.bin

-rw-r--r-- 1 lishaojie Students  1.3M 4 14 14:09 decoder-epoch-29-avg-9-with-averaged-
˓→model.int8.onnx
-rw-r--r-- 1 lishaojie Students  2.0M 4 14 14:09 decoder-epoch-29-avg-9-with-averaged-
˓→model.onnx
-rw-r--r-- 1 lishaojie Students 121M 4 14 14:09 encoder-epoch-29-avg-9-with-averaged-
˓→model.int8.onnx
-rw-r--r-- 1 lishaojie Students 279M 4 14 14:09 encoder-epoch-29-avg-9-with-averaged-
˓→model.onnx
-rw-r--r-- 1 lishaojie Students 254K 4 14 14:09 joiner-epoch-29-avg-9-with-averaged-
˓→model.int8.onnx
-rw-r--r-- 1 lishaojie Students 1003K 4 14 14:09 joiner-epoch-29-avg-9-with-averaged-
˓→model.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/encoder-epoch-29-avg-9-with-
averaged-model.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/decoder-epoch-29-avg-9-with-
averaged-model.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/joiner-epoch-29-avg-9-with-
averaged-model.onnx \
./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_19364697.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
streaming-zipformer-fr-2023-04-14/encoder-epoch-29-avg-9-with-averaged-model.onnx",_
decoder_filename="./sherpa-onnx-streaming-zipformer-fr-2023-04-14/decoder-epoch-29-avg-
9-with-averaged-model.onnx", joiner_filename="./sherpa-onnx-streaming-zipformer-fr-
2023-04-14/joiner-epoch-29-avg-9-with-averaged-model.onnx", tokens="./sherpa-onnx-
streaming-zipformer-fr-2023-04-14/tokens.txt", num_threads=2, debug=False), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, decoding_method="greedy_search")
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr-
19364697.wav
wav duration (s): 7.128
Started
Done!
Recognition result for ./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_
voice_fr_19364697.wav:
CE SITE CONTIENT QUATRE TOMBEAUX DE LA DYNASTIE ASHÉMÉNIDE ET SEPT DES SASSANDIDES
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.458 s
Real time factor (RTF): 0.458 / 7.128 = 0.064
```

int8

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/encoder-epoch-29-avg-9-with-
--averaged-model.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/decoder-epoch-29-avg-9-with-
--averaged-model.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/joiner-epoch-29-avg-9-with-
--averaged-model.int8.onnx \
./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr_19364697.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
streaming-zipformer-fr-2023-04-14/encoder-epoch-29-avg-9-with-averaged-model.int8.onnx
", decoder_filename="./sherpa-onnx-streaming-zipformer-fr-2023-04-14/decoder-epoch-29-
avg-9-with-averaged-model.onnx", joiner_filename="./sherpa-onnx-streaming-zipformer-fr-
2023-04-14/joiner-epoch-29-avg-9-with-averaged-model.int8.onnx", tokens="./sherpa-onnx-
streaming-zipformer-fr-2023-04-14/tokens.txt", num_threads=2, debug=False), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, decoding_method="greedy_search")
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_voice_fr-
19364697.wav
wav duration (s): 7.128
Started
Done!
Recognition result for ./sherpa-onnx-streaming-zipformer-fr-2023-04-14/test_wavs/common_
voice_fr_19364697.wav:
CE SITE CONTIENT QUATRE TOMBEAUX DE LA DYNASTIE ASHÉMÉNIDE ET SEPT DES SASSANDIDES
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.485 s
Real time factor (RTF): 0.485 / 7.128 = 0.068
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx
./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/encoder-epoch-29-avg-9-with-
˓→averaged-model.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/decoder-epoch-29-avg-9-with-
˓→averaged-model.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-fr-2023-04-14/joiner-epoch-29-avg-9-with-
˓→averaged-model.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16 (Bilingual, Chinese + English)

Hint: It is a small model.

This model is converted from

<https://huggingface.co/csukuangfj/k2fsa-zipformer-bilingual-zh-en-t>

which supports both Chinese and English. The model is contributed by the community and is trained on tens of thousands of some internal dataset.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2

tar xf sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2
rm sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16 fangjun$ ls -lh *.onnx
total 158M
drwxr-xr-x 2 1001 127 4.0K Mar 20 13:11 64
```

(continues on next page)

(continued from previous page)

```
drwxr-xr-x 2 1001 127 4.0K Mar 20 13:11 96
-rw-r--r-- 1 1001 127 240K Mar 20 13:11 bpe.model
-rw-r--r-- 1 1001 127 3.4M Mar 20 13:11 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 1001 127 14M Mar 20 13:11 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 1001 127 41M Mar 20 13:11 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 1001 127 85M Mar 20 13:11 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 1001 127 3.1M Mar 20 13:11 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 1001 127 13M Mar 20 13:11 joiner-epoch-99-avg-1.onnx
drwxr-xr-x 2 1001 127 4.0K Mar 20 13:11 test_wavs
-rw-r--r-- 1 1001 127 55K Mar 20 13:11 tokens.txt
```

Hint: There are two sub-folders in the model directory: 64 and 96. The number represents chunk size. The larger the number, the lower the RTF. The default chunk size is 32.

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx --tokens=./sherpa-onnx-
˓→streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt --encoder=./sherpa-
˓→onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-epoch-99-avg-1.onnx -
˓→decoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder-
˓→epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-
˓→2023-02-16/joiner-epoch-99-avg-1.onnx ./sherpa-onnx-streaming-zipformer-small-
˓→bilingual-zh-en-2023-02-16/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
˓→"./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-epoch-99-
˓→avg-1.onnx", decoder="./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-
˓→16/decoder-epoch-99-avg-1.onnx", joiner="./sherpa-onnx-streaming-zipformer-small-
˓→bilingual-zh-en-2023-02-16/joiner-epoch-99-avg-1.onnx"),_
˓→paraformer=OnlineParaformerModelConfig(encoder="", decoder=""), wenet_
˓→ctc=OnlineWenetCtcModelConfig(model="", chunk_size=16, num_left_chunks=4), zipformer2_
˓→ctc=OnlineZipformer2CtcModelConfig(model=""), tokens="./sherpa-onnx-streaming-
˓→zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt", num_threads=1, debug=False,_
˓→provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="", scale=0.5),_
˓→endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_
˓→trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
˓→nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0),_
˓→rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
˓→utterance_length=20)), enable_endpoint=True, max_active_paths=4, hotwords_score=1.5,_
˓→hotwords_file="", decoding_method="greedy_search", blank_penalty=0)
./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
Elapsed seconds: 1, Real time factor (RTF): 0.1
MONDAY TODAY IS THEY AFTER TOMORROW
{
  "text": " MONDAY TODAY IS THEY AFTER TOMORROW", "tokens": [ "", "", "", " MO", "N",
˓→"DAY", " TO", "DAY", " IS", " THEY", " AFTER", " TO", "M", "OR", "ROW", "", "", "", "" ],
˓→"timestamps": [ 0.64, 1.08, 1.64, 2.08, 2.20, 2.36, 4.16, 4.36, 5.12, 7.16, 7.44, 8.
˓→00, 8.12, 8.20, 8.44, 9.08, 9.44, 9.64, 9.88 ], "ys_probs": [ -0.000507, -0.056152, -0.
˓→007374, -0.213242, -0.362640, -0.117561, -1.036179, -0.219900, -0.150360, -0.734749,
˓→-0.113281, -0.060974, -0.117775, -0.361603, -0.039993, -0.217766, -0.042011, -0.108857,
˓→-0.135108 ], "lm_probs": [ ], "context_scores": [ ], "segment": 0, "start_time": 0.
˓→00, "is_final": false}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-
epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder-
epoch-99-avg-1.onnx \
```

(continues on next page)

(continued from previous page)

```
--joiner=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner-
˓→epoch-99-avg-1.int8.onnx \
./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx --tokens=./sherpa-onnx-
˓→streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt --encoder=./sherpa-
˓→onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-epoch-99-avg-1.int8.
˓→onnx --decoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/
˓→decoder-epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-small-bilingual-
˓→zh-en-2023-02-16/joiner-epoch-99-avg-1.int8.onnx ./sherpa-onnx-streaming-zipformer-
˓→small-bilingual-zh-en-2023-02-16/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
˓→"./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-epoch-99-
˓→avg-1.int8.onnx", decoder="./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-
˓→2023-02-16/decoder-epoch-99-avg-1.onnx", joiner="./sherpa-onnx-streaming-zipformer-
˓→small-bilingual-zh-en-2023-02-16/joiner-epoch-99-avg-1.int8.onnx"), paraformer=
˓→OnlineParaformerModelConfig(encoder="", decoder ""), wenet_ctc=OnlineWenetCtcModelConfig(model="",
˓→chunk_size=16, num_left_chunks=4), zipformer2_ctc=OnlineZipformer2CtcModelConfig(model=""),
˓→tokens="./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt",
˓→num_threads=1, debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",
˓→scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False,
˓→min_trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
˓→nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0),
˓→rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_
˓→utterance_length=20)), enable_endpoint=True, max_active_paths=4, hotwords_score=1.5,
˓→hotwords_file="", decoding_method="greedy_search", blank_penalty=0)
./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/test_wavs/0.wav
Elapsed seconds: 0.69, Real time factor (RTF): 0.069
MONDAY TODAY IS THEY AFTER TOMORROW
{
  "text": " MONDAY TODAY IS THEY AFTER TOMORROW", "tokens": [ "", "", "", " MO", " N",
˓→" DAY", " TO", " DAY", " IS", " THEY", " AFTER", " TO", " M", " OR", " ROW", " ", " ", " ", " "
˓→], "timestamps": [ 0.64, 1.08, 1.64, 2.08, 2.20, 2.36, 4.20, 4.36, 5.12, 7.16, 7.44, 8.
˓→00, 8.12, 8.20, 8.40, 9.04, 9.44, 9.64, 9.88 ], "ys_probs": [ -0.000305, -0.152557, -0.
˓→007835, -0.156221, -0.622139, -0.081843, -1.140152, -0.418322, -0.198410, -0.939461,
˓→-0.224989, -0.052963, -0.098366, -0.081665, -0.453255, -0.335670, -0.039482, -0.381765,
˓→-0.192475 ], "lm_probs": [ ], "context_scores": [ ], "segment": 0, "start_time": 0.
˓→00, "is_final": false}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/encoder-
epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/decoder-
epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16/joiner-
epoch-99-avg-1.int8.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23 (Chinese)

Hint: It is a small model.

This model is from

<https://huggingface.co/marcoyang/sherpa-ncnn-streaming-zipformer-zh-14M-2023-02-23/>

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
streaming-zipformer-zh-14M-2023-02-23.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
# streaming-zipformer-zh-14M-2023-02-23.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23.tar.bz2
rm sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 1.8M Sep 10 15:31 decoder-epoch-99-avg-1.int8.onnx
```

(continues on next page)

(continued from previous page)

-rw-r--r--	1	fangjun	staff	7.2M	Sep 10 15:31	decoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	21M	Sep 10 15:31	encoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	39M	Sep 10 15:31	encoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	1.7M	Sep 10 15:31	joiner-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	6.8M	Sep 10 15:31	joiner-epoch-99-avg-1.onnx

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-99-avg-1.
˓→onnx \
--decoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/joiner-epoch-99-avg-1.
˓→onnx \
./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

`CHCP 65001`

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
˓→build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/
˓→tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-
˓→99-avg-1.onnx --decoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/decoder-
˓→epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/
˓→joiner-epoch-99-avg-1.onnx ./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/test_
˓→wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
→"./sherpa-onnéx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-99-avg-1.onnéx",
→decoder="./sherpa-onnéx-streaming-zipformer-zh-14M-2023-02-23/decoder-epoch-99-avg-1.
→onnéx", joiner="./sherpa-onnéx-streaming-zipformer-zh-14M-2023-02-23/joiner-epoch-99-avg-
→1.onnéx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), tokens="./
→sherpa-onnéx-streaming-zipformer-zh-14M-2023-02-23/tokens.txt", num_threads=1,
→debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",
→scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
→nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
→rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
→utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
→silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
→context_score=1.5, decoding_method="greedy_search")
./sherpa-onnéx-streaming-zipformer-zh-14M-2023-02-23/test_wavs/0.wav
Elapsed seconds: 0.21, Real time factor (RTF): 0.038

{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":[0.32, 0.64, 0.76,
→0.96, 1.08, 1.16, 1.96, 2.04, 2.24, 2.36, 2.56, 2.68, 2.80, 3.36, 3.52, 3.64, 3.72, 3.
→84, 3.92, 4.00, 4.08, 4.24, 4.48, 4.56, 4.72],"tokens":["","","","","","","","","","","","","",
→"","","","","","","","","","","","","","",""]}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-99-avg-1.
↪int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/decoder-epoch-99-avg-1.
↪onnx \
--joiner=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/joiner-epoch-99-avg-1.
↪int8.onnx \
./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  

build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/  

tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-  

99-avg-1.int8.onnx --decoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/  

decoder-epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-  

23/joiner-epoch-99-avg-1.int8.onnx ./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/  

test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  

dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=  

"./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-99-avg-1.int8.onnx",  

decoder="./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/decoder-epoch-99-avg-1.  

onnx", joiner="./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/joiner-epoch-99-avg-  

1.int8.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), tokens=  

"./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/tokens.txt", num_threads=1,  

debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",  

scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_  

nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),  

rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_  

utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_  

silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,  

context_score=1.5, decoding_method="greedy_search")  

./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/test_wavs/0.wav  

Elapsed seconds: 0.16, Real time factor (RTF): 0.028

{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":"[0.32, 0.64, 0.76,  

0.96, 1.08, 1.16, 1.92, 2.04, 2.24, 2.36, 2.56, 2.68, 2.76, 3.36, 3.52, 3.64, 3.72, 3.  

84, 3.92, 4.00, 4.08, 4.24, 4.48, 4.56, 4.72]","tokens":["","","","","","","","","","","","","","","  

","","","","","","","","","","","","","",""]}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/encoder-epoch-99-avg-1.  

onnx \
--decoder=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/decoder-epoch-99-avg-1.  

onnx \
--joiner=./sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-streaming-zipformer-en-20M-2023-02-17 (English)

Hint: It is a small model.

This model is from

<https://huggingface.co/desh2608/icefall-asr-librispeech-pruned-transducer-stateless7-streaming-small>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-20M-2023-02-17.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-en-20M-2023-02-17.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-en-20M-2023-02-17.tar.bz2
rm sherpa-onnx-streaming-zipformer-en-20M-2023-02-17.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-zipformer-en-20M-2023-02-17 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 527K Sep 10 17:06 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Sep 10 17:06 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 41M Sep 10 17:06 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 85M Sep 10 17:06 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 253K Sep 10 17:06 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Sep 10 17:06 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.
˓→onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.
˓→onnx \
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
˓→build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
˓→tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-
˓→99-avg-1.onnx --decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-
˓→epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
˓→joiner-epoch-99-avg-1.onnx ./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_
˓→wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
˓→"./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.onnx",
˓→decoder="./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
˓→onnx", joiner="./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-
˓→1.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), tokens="./
˓→sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt", num_threads=1,
˓→debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",
˓→scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
˓→nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
˓→rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
˓→utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓→silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
˓→context_score=1.5, decoding_method="greedy_search")
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
Elapsed seconds: 0.32, Real time factor (RTF): 0.049
THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLELS
{"is_final":false,"segment":0,"start_time":0.0,"text":" THE YELLOW LAMPS WOULD LIGHT UP
˓→HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLELS","timestamps":[2.04, 2.16, 2.28,
˓→2.36, 2.52, 2.64, 2.68, 2.76, 2.92, 3.08, 3.40, 3.60, 3.72, 3.88, 4.12, 4.48, 4.64, 4.
˓→68, 4.84, 4.96, 5.16, 5.20, 5.32, 5.36, 5.60, 5.72, 5.92, 5.96, 6.08, 6.24, 6.36, 6.52
˓→,"tokens":[" THE", " YE", " LL", " OW", " LA", " M", " P", " S", " WOULD", " LIGHT", " UP", " HE", " RE
˓→", " AND", " THERE", " THE", " S", " QUA", " LI", " D", " ", " QUA", " R", " TER", " OF", " THE", " B", " RA
˓→", " FF", " L", " EL", " S"]}
```

(continues on next page)

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.
→ int8.onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
→ onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.
→ int8.onnx \
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
→ build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
→ tokens.txt --encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-
→ 99-avg-1.int8.onnx --decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
→ decoder-epoch-99-avg-1.onnx --joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-
→ 17/joiner-epoch-99-avg-1.int8.onnx ./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/
→ test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
→ dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
→ "./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.int8.onnx",
→ decoder="./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
→ onnx", joiner="./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-
→ 1.int8.onnx"), paraformer=OnlineParaformerModelConfig(encoder="", decoder ""), tokens=
→ "./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt", num_threads=1,
→ debug=False, provider="cpu", model_type=""), lm_config=OnlineLMConfig(model="",
→ scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
→ nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
→ rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
→ utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
→ silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
→ context_score=1.5, decoding_method="greedy_search")
./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/test_wavs/0.wav
Elapsed seconds: 0.25, Real time factor (RTF): 0.038
THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLS
{"is_final":false,"segment":0,"start_time":0.0,"text":" THE YELLOW LAMPS WOULD LIGHT UP
→ HERE AND THERE THE SQUALID QUARTER OF THE BRAFFLS","timestamps":[2.04, 2.20, 2.28, 2.
→ 36, 2.52, 2.64, 2.68, 2.76, 2.92, 3.08, 3.40, 3.60, 3.72, 3.88, 4.12, 4.48, 4.64, 4.68,
→ 4.84, 4.96, 5.16, 5.20, 5.32, 5.36, 5.60, 5.72, 5.92, 5.96, 6.08, 6.24, 6.36], "tokens"
→ :[" THE", " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", " LIGHT", " UP", " HE", "RE", " AND",
→ " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA", "R", "TER", " OF", " THE", "B", "RA", "EE", "I
→ , "S"]}
```

(continued from previous page)

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/encoder-epoch-99-avg-1.
˓→onnx \
--decoder=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-streaming-zipformer-en-20M-2023-02-17/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

Conformer-transducer-based Models

Hint: Please refer to *Installation* to install *sherpa-onnx* before you read this section.

csukuangfj/sherpa-onnx-streaming-conformer-zh-2023-05-23 (Chinese)

This model is converted from

https://huggingface.co/luomingshuang/icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/wenetspeech/ASR/pruned_transducer_stateless5

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-conformer-zh-2023-05-23.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-conformer-zh-2023-05-23.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf sherpa-onnx-streaming-conformer-zh-2023-05-23.tar.bz2
rm sherpa-onnx-streaming-conformer-zh-2023-05-23.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-conformer-zh-2023-05-23 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 11M May 23 14:44 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 12M May 23 14:44 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 160M May 23 14:46 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 345M May 23 14:47 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 2.7M May 23 14:44 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 11M May 23 14:44 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-conformer-zh-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
streaming-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx", decoder_filename="./
sherpa-onnx-streaming-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx", joiner_
filename="./sherpa-onnx-streaming-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx",_
tokens=".sherpa-onnx-streaming-conformer-zh-2023-05-23/tokens.txt", num_threads=2,_
debug=False), lm_config=OnlineLMConfig(model="", scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, decoding_method="greedy_search") (continues on next page)
```

(continued from previous page)

```
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav
wav duration (s): 5.611
Started
Done!
Recognition result for ./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav:
{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":"[0.00, 0.48, 0.76,
← 0.88, 1.08, 1.24, 2.00, 2.04, 2.16, 2.36, 2.56, 2.72, 2.92, 3.36, 3.44, 3.64, 3.72, 3.
← 84, 3.96, 4.04, 4.16, 4.28, 4.48, 4.64, 4.84, 5.16]","tokens":["","","","","","","","","","","",
←","","","","","","","","","","","","","","","",""]}
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.559 s
Real time factor (RTF): 0.559 / 5.611 = 0.100
```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-conformer-zh-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.int8.
←onnx \
--decoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.int8.
←onnx \
--joiner=./sherpa-onnx-streaming-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.int8.
←onnx \
./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
←dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-
←streaming-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.int8.onnx", decoder_filename=
←"./sherpa-onnx-streaming-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.int8.onnx",
←joiner_filename="./sherpa-onnx-streaming-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.
←int8.onnx", tokens="./sherpa-onnx-streaming-conformer-zh-2023-05-23/tokens.txt", num_
←threads=2, debug=False), lm_config=OnlineLMConfig(model="", scale=0.5), endpoint_
←config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
←silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, page_
←min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
←nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
8.25. Pre-trained models ←endpoint=True, max_active_paths=4, decoding_method="greedy_search")
```

(continued from previous page)

```
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav
wav duration (s): 5.611
Started
Done!
Recognition result for ./sherpa-onnx-streaming-conformer-zh-2023-05-23/test_wavs/0.wav:
{"is_final":false,"segment":0,"start_time":0.0,"text":"","timestamps":"[0.00, 0.48, 0.72,
← 0.88, 1.08, 1.24, 2.00, 2.04, 2.16, 2.28, 2.56, 2.72, 2.92, 3.36, 3.44, 3.60, 3.72, 3.
← 84, 3.92, 4.04, 4.16, 4.28, 4.48, 4.60, 4.84, 5.16]","tokens":["","","","","","","","","","","",
← "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""}]}
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.493 s
Real time factor (RTF): 0.493 / 5.611 = 0.088
```

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-conformer-zh-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-streaming-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-streaming-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

LSTM-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

csukuangfj/sherpa-onnx-lstm-en-2023-02-17 (English)

This model trained using the [GigaSpeech](#) and the [LibriSpeech](#) dataset.

Please see <https://github.com/k2-fsa/icefall/pull/558> for how the model is trained.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/lstm_transducer_stateless2

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-lstm-
en-2023-02-17.tar.bz2
# For Chinese users, please use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
lstm-en-2023-02-17.tar.bz2
tar xvf sherpa-onnx-lstm-en-2023-02-17.tar.bz2
rm sherpa-onnx-lstm-en-2023-02-17.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
sherpa-onnx-lstm-en-2023-02-17$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root  1.3M Mar 31 22:41 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root  2.0M Mar 31 22:41 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root  80M Mar 31 22:41 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 319M Mar 31 22:41 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 254K Mar 31 22:41 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 1003K Mar 31 22:41 joiner-epoch-99-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-lstm-en-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-lstm-en-2023-02-17/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-lstm-en-2023-02-17/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-lstm-en-2023-02-17/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-lstm-
˓en-2023-02-17/encoder-epoch-99-avg-1.onnx", decoder_filename="./sherpa-onnx-lstm-en-
˓2023-02-17/decoder-epoch-99-avg-1.onnx", joiner_filename="./sherpa-onnx-lstm-en-2023-
˓02-17/joiner-epoch-99-avg-1.onnx", tokens="./sherpa-onnx-lstm-en-2023-02-17/tokens.txt
˓", num_threads=2, debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_
˓contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
˓rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
˓utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓silence=0, min_utterance_length=20)), enable_endpoint=True,
max_active_paths=4, decoding_method="greedy_search")
2023-03-31 22:53:22.120185169 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 576406, index: 16, mask: {17, 53, }, error code: 22
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
2023-03-31 22:53:22.120183162 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 576405, index: 15, mask: {16, 52, }, error code: 22
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav
wav duration (s): 6.625
Started
Done!
Recognition result for ./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav:
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
˓QUARTER OF THE BROTHELS
```

(continues on next page)

(continued from previous page)

```
num threads: 2
decoding method: greedy_search
Elapsed seconds: 2.927 s
Real time factor (RTF): 2.927 / 6.625 = 0.442
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-lstm-en-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-lstm-en-2023-02-17/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-lstm-en-2023-02-17/decoder-epoch-99-avg-1.int8.onnx \
--joiner=./sherpa-onnx-lstm-en-2023-02-17/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-lstm-
˓en-2023-02-17/encoder-epoch-99-avg-1.int8.onnx", decoder_filename="./sherpa-onnx-lstm-
˓en-2023-02-17/decoder-epoch-99-avg-1.int8.onnx", joiner_filename="./sherpa-onnx-lstm-
˓en-2023-02-17/joiner-epoch-99-avg-1.int8.onnx", tokens="./sherpa-onnx-lstm-en-2023-02-
˓17/tokens.txt", num_threads=2, debug=False), endpoint_
˓config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
˓min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
˓nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
˓endpoint=True, max_active_paths=4, decoding_method="greedy_search")
2023-03-31 22:55:46.608941959 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 578689, index: 16, mask: {17, 53, }, error code: 22_
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
2023-03-31 22:55:46.608939862 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
˓setaffinity_np failed for thread: 578688, index: 15, mask: {16, 52, }, error code: 22_
˓error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
˓is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav
wav duration (s): 6.625
Started
Done!
Recognition result for ./sherpa-onnx-lstm-en-2023-02-17/test_wavs/0.wav:
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
˓QUARTER OF THE BROTHELS
```

(continues on next page)

(continued from previous page)

```
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.009 s
Real time factor (RTF): 1.009 / 6.625 = 0.152
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-lstm-en-2023-02-17/tokens.txt \
--encoder=./sherpa-onnx-lstm-en-2023-02-17/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-lstm-en-2023-02-17/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-lstm-en-2023-02-17/joiner-epoch-99-avg-1.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-lstm-zh-2023-02-20 (Chinese)

This is a model trained using the WenetSpeech dataset.

Please see <https://github.com/k2-fsa/icefall/pull/595> for how the model is trained.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-lstm-
zh-2023-02-20.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
lstm-zh-2023-02-20.tar.bz2

tar xvf sherpa-onnx-lstm-zh-2023-02-20.tar.bz2
rm sherpa-onnx-lstm-zh-2023-02-20.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-lstm-zh-2023-02-20$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root 12M Mar 31 20:55 decoder-epoch-11-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 12M Mar 31 20:55 decoder-epoch-11-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 80M Mar 31 20:55 encoder-epoch-11-avg-1.int8.onnx
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 kuangfangjun root 319M Mar 31 20:55 encoder-epoch-11-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 2.8M Mar 31 20:55 joiner-epoch-11-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 11M Mar 31 20:55 joiner-epoch-11-avg-1.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-lstm-zh-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-lstm-zh-2023-02-20/encoder-epoch-11-avg-1.onnx \
--decoder=./sherpa-onnx-lstm-zh-2023-02-20/decoder-epoch-11-avg-1.onnx \
--joiner=./sherpa-onnx-lstm-zh-2023-02-20/joiner-epoch-11-avg-1.onnx \
./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-lstm-
zh-2023-02-20/encoder-epoch-11-avg-1.onnx", decoder_filename="./sherpa-onnx-lstm-zh-
2023-02-20/decoder-epoch-11-avg-1.onnx", joiner_filename="./sherpa-onnx-lstm-zh-2023-
02-20/joiner-epoch-11-avg-1.onnx", tokens="./sherpa-onnx-lstm-zh-2023-02-20/tokens.txt
", num_threads=2, debug=False), endpoint_config=EndpointConfig(rule1=EndpointRule(must_
contain_nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),_
rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,_
decoding_method="greedy_search")
2023-03-31 22:58:59.348229346 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
setaffinity_np failed for thread: 578800, index: 15, mask: {16, 52, }, error code: 22,_
error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
is not set.
```

(continues on next page)

(continued from previous page)

```
2023-03-31 22:58:59.348231417 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 578801, index: 16, mask: {17, 53, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav
wav duration (s): 5.611
Started
Done!
Recognition result for ./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav:

num threads: 2
decoding method: greedy_search
Elapsed seconds: 3.030 s
Real time factor (RTF): 3.030 / 5.611 = 0.540
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-lstm-zh-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-lstm-zh-2023-02-20/encoder-epoch-11-avg-1.int8.onnx \
--decoder=./sherpa-onnx-lstm-zh-2023-02-20/decoder-epoch-11-avg-1.int8.onnx \
--joiner=./sherpa-onnx-lstm-zh-2023-02-20/joiner-epoch-11-avg-1.int8.onnx \
./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
↳ dim=80), model_config=OnlineTransducerModelConfig(encoder_filename="./sherpa-onnx-lstm-
↳ zh-2023-02-20/encoder-epoch-11-avg-1.int8.onnx", decoder_filename="./sherpa-onnx-lstm-
↳ zh-2023-02-20/decoder-epoch-11-avg-1.int8.onnx", joiner_filename="./sherpa-onnx-lstm-
↳ zh-2023-02-20/joiner-epoch-11-avg-1.int8.onnx", tokens="./sherpa-onnx-lstm-zh-2023-02-
↳ 20/tokens.txt", num_threads=2, debug=False), endpoint_
↳ config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
↳ silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True, ↳
↳ min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
↳ nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
↳ endpoint=True, max_active_paths=4, decoding_method="greedy_search")
```

(continues on next page)

(continued from previous page)

```

2023-03-31 23:01:05.737519659 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 578880, index: 15, mask: {16, 52, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
2023-03-31 23:01:05.737521655 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 578881, index: 16, mask: {17, 53, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
sampling rate of input file: 16000
wav filename: ./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav
wav duration (s): 5.611
Started
Done!
Recognition result for ./sherpa-onnx-lstm-zh-2023-02-20/test_wavs/0.wav:

num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.091 s
Real time factor (RTF): 1.091 / 5.611 = 0.194

```

Real-time speech recognition from a microphone

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-lstm-zh-2023-02-20/tokens.txt \
--encoder=./sherpa-onnx-lstm-zh-2023-02-20/encoder-epoch-11-avg-1.onnx \
--decoder=./sherpa-onnx-lstm-zh-2023-02-20/decoder-epoch-11-avg-1.onnx \
--joiner=./sherpa-onnx-lstm-zh-2023-02-20/joiner-epoch-11-avg-1.onnx

```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

8.25.2 Online paraformer models

This section lists available online paraformer models.

Paraformer models

Hint: Please refer to *Installation* to install *sherpa-onnx* before you read this section.

csukuangfj/sherpa-onnx-streaming-parafomer-bilingual-zh-en (Chinese + English)

Note: This model does not support timestamps. It is a bilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/damo/speech_parafomer_asr_nat-zh-cn-16k-common-vocab8404-online/summary

The code for converting can be found at

<https://huggingface.co/csukuangfj/streaming-parafomer-zh>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
      ↪streaming-parafomer-bilingual-zh-en.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
      ↪streaming-parafomer-bilingual-zh-en.tar.bz2

tar xvf sherpa-onnx-streaming-parafomer-bilingual-zh-en.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-parafomer-bilingual-zh-en fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 68M Aug 14 09:53 decoder.int8.onnx
-rw-r--r-- 1 fangjun staff 218M Aug 14 09:55 decoder.onnx
-rw-r--r-- 1 fangjun staff 158M Aug 14 09:54 encoder.int8.onnx
-rw-r--r-- 1 fangjun staff 607M Aug 14 09:57 encoder.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/tokens.txt \
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/encoder.onnx \
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/decoder.onnx \
./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/
tokens.txt --paraformer-encoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/
encoder.onnx --paraformer-decoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/
decoder.onnx ./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='./sherpa-
onnx-streaming-paraformer-bilingual-zh-en/encoder.onnx', decoder='./sherpa-onnx-
streaming-paraformer-bilingual-zh-en/decoder.onnx'), tokens='./sherpa-onnx-streaming-
paraformer-bilingual-zh-en/tokens.txt', num_threads=1, debug=False, provider='cpu',_
model_type=''), lm_config=OnlineLMConfig(model='', scale=0.5), endpoint_
config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,_
min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
endpoint=True, max_active_paths=4, context_score=1.5, decoding_method='greedy_search')
./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav
Elapsed seconds: 2.2, Real time factor (RTF): 0.21
monday today day is the day after tomorrow
{"is_final":false,"segment":0,"start_time":0.0,"text":" monday today day is the day_
after tomorrow ","timestamps":[],"tokens":["","","","","mon@@", "day", "today", "day", "is",
","","","","the", "day", "after", "tom@@", "or@@", "row", "", "", "", ""]}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/tokens.txt \
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/encoder.int8.
˓→onnx \
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/decoder.int8.
˓→onnx \
./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
˓→build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/
˓→tokens.txt --paraformer-encoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/
˓→encoder.int8.onnx --paraformer-decoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-
˓→en/decoder.int8.onnx ./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
˓→'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='./sherpa-
˓→onnx-streaming-paraformer-bilingual-zh-en/encoder.int8.onnx', decoder='./sherpa-onnx-
˓→streaming-paraformer-bilingual-zh-en/decoder.int8.onnx'), tokens='./sherpa-onnx-
˓→streaming-paraformer-bilingual-zh-en/tokens.txt', num_threads=1, debug=False, provider=
˓→'cpu', model_type=''), lm_config=OnlineLMConfig(model='', scale=0.5), endpoint_
˓→config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_trailing_
˓→silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_nonsilence=True,
˓→min_trailing_silence=1.2, min_utterance_length=0), rule3=EndpointRule(must_contain_
˓→nonsilence=False, min_trailing_silence=0, min_utterance_length=20)), enable_
˓→endpoint=True, max_active_paths=4, context_score=1.5, decoding_method="greedy_search")
˓→./sherpa-onnx-streaming-paraformer-bilingual-zh-en/test_wavs/0.wav
Elapsed seconds: 1.6, Real time factor (RTF): 0.15
monday today day is the day after tomorrow
{"is_final":false,"segment":0,"start_time":0.0,"text":" monday today day is the day
˓→after tomorrow ","timestamps":[[],"tokens":["","","","","mon@@", "day", "today", "day", "is",
˓→","","","","the", "day", "after", "tom@@", "or@@", "row", "", "", "", ""}}}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/tokens.txt \
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/encoder.int8.
→onnx \
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/decoder.int8.
→onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

csukuangfj/sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en (Chinese + Cantonese + English)

Note: This model does not support timestamps. It is a trilingual model, supporting both Chinese and English. ()

This model is converted from

https://modelscope.cn/models/dengcunqin/speech_paraformer-large_asr_nat-zh-cantonese-en-16k-vocab8501-online/files

You can find the conversion code after downloading and unzipping the model.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→streaming-paraformer-trilingual-zh-cantonese-en.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→streaming-paraformer-trilingual-zh-cantonese-en.tar.bz2

tar xvf sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 69M Feb 29 19:44 decoder.int8.onnx
-rw-r--r-- 1 fangjun staff 218M Feb 29 19:44 decoder.onnx
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 159M Feb 29 19:44 encoder.int8.onnx  
-rw-r--r-- 1 fangjun staff 607M Feb 29 19:44 encoder.onnx
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx \  
--tokens=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/tokens.txt \  
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/  
encoder.onnx \  
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/  
decoder.onnx \  
.sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-paraformer-trilingual-zh-  
cantonese-en/tokens.txt --paraformer-encoder=./sherpa-onnx-streaming-paraformer-  
trilingual-zh-cantonese-en/encoder.int8.onnx --paraformer-decoder=./sherpa-onnx-  
streaming-paraformer-trilingual-zh-cantonese-en/decoder.int8.onnx ./sherpa-onnx-  
streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav  
  
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_=  
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder_=  
'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='./sherpa-  
onnx-streaming-paraformer-trilingual-zh-cantonese-en/encoder.int8.onnx', decoder='./  
sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/decoder.int8.onnx'), wenet_=  
ctc=OnlineWenetCtcModelConfig(model='', chunk_size=16, num_left_chunks=4), zipformer2_=  
ctc=OnlineZipformer2CtcModelConfig(model=''), tokens='./sherpa-onnx-streaming-  
paraformer-trilingual-zh-cantonese-en/tokens.txt', num_threads=1, debug=False, _  
provider='cpu', model_type=''), lm_config=OnlineLmConfig(model='', scale=0.5), _  
396 endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsilence=False, min_=  
trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_=  
nonsilence=True, min_trailing_silence=1.2, min_utterance_length=0), _  
rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_silence=0, min_=  
utterance_length=20), enable_endpoint=True, max_active_paths=4, hotwords_score=1.5)
```

(continued from previous page)

```
./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav
Elapsed seconds: 0.98, Real time factor (RTF): 0.16

{ "text": "", "tokens": [ "", "", "", "", "", "", "", "", "", "", "" ], "timestamps": [ ], "ys_probs": [ ], "lm_probs": [ ], "context_scores": [ ], "segment": 0, "start_time": 0.00, "is_final": false}
```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--tokens=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/tokens.txt \
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/ \
encoder.int8.onnx \
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/ \
decoder.int8.onnx \
./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx --tokens=./sherpa-onnx-streaming-paraformer-trilingual-zh-
cantonese-en/tokens.txt --paraformer-encoder=./sherpa-onnx-streaming-paraformer-
trilingual-zh-cantonese-en/encoder.int8.onnx --paraformer-decoder=./sherpa-onnx-
streaming-paraformer-trilingual-zh-cantonese-en/decoder.int8.onnx ./sherpa-onnx-
streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='./sherpa-
onnx-streaming-paraformer-trilingual-zh-cantonese-en/encoder.int8.onnx', decoder='./
sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/decoder.int8.onnx'), wenet_-
ctc=OnlineWenetCtcModelConfig(model='', chunk_size=16, num_left_chunks=4), zipformer2_-
ctc=OnlineZipformer2CtcModelConfig(model=''), tokens='./sherpa-onnx-streaming-
paraformer-trilingual-zh-cantonese-en/tokens.txt', num_threads=1, debug=False,_
provider='cpu', model_type=''), lm_config=OnlineLMConfig(model='', scale=0.5),_
endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_nonsense=False, min_
trailing_silence=2.4, min_utterance_length=0), rule2=EndpointRule(must_contain_
nonsense=True, min_trailing_silence=1.2, min_utterance_length=0),_
enable_endpoint=True, max_active_paths=4, hotwords_score=1.5,_
hotwords_file='', decoding_method="greedy_search", blank_penalty=0)
```

(continues on next page)

8.25. Pre-trained Models 397

(continued from previous page)

```
./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav
Elapsed seconds: 0.84, Real time factor (RTF): 0.14

{ "text": "", "tokens": [ "", "", "", "", "", "", "", "", "", "", "" ], "timestamps": [ ], "ys_probs": [ ], "lm_probs": [ ], "context_scores": [ ], "segment": 0, "start_time": 0.00, "is_final": false}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone \
--tokens=./sherpa-onnx-streaming-paraformer-bilingual-zh-en/tokens.txt \
--paraformer-encoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/
encoder.int8.onnx \
--paraformer-decoder=./sherpa-onnx-streaming-paraformer-trilingual-zh-cantonese-en/
decoder.int8.onnx
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

8.25.3 Online CTC models

This section lists available online CTC models.

Zipformer-CTC-based Models

Hint: Please refer to *Installation* to install *sherpa-onnx* before you read this section.

sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13 (Chinese)

Training code for this model can be found at <https://github.com/k2-fsa/icefall/pull/1369>. It supports only Chinese.

Please refer to https://github.com/k2-fsa/icefall/tree/master/egs/multi_zh-hans/ASR#included-training-sets for the detailed information about the training data. In total, there are 14k hours of training data.

In the following, we describe how to download it and use it with *sherpa-onnx*.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-ctc-multi-zh-hans-2023-12-13.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→streaming-zipformer-ctc-multi-zh-hans-2023-12-13.tar.bz2

tar xvf sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13.tar.bz2
rm sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13.tar.bz2
ls -lh sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13
```

The output is given below:

```
$ ls -lh sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13
total 654136
-rw-r--r--@ 1 fangjun  staff   28B Dec 13 16:19 README.md
-rw-r--r--@ 1 fangjun  staff  258K Dec 13 16:19 bpe.model
-rw-r--r--@ 1 fangjun  staff  68M Dec 13 16:19 ctc-epoch-20-avg-1-chunk-16-left-128.
˓→int8.onnx
-rw-r--r--@ 1 fangjun  staff  252M Dec 13 16:19 ctc-epoch-20-avg-1-chunk-16-left-128.
˓→onnx
drwxr-xr-x@ 8 fangjun  staff  256B Dec 13 16:19 test_wavs
-rw-r--r--@ 1 fangjun  staff   18K Dec 13 16:19 tokens.txt
```

Decode a single wave file

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode a wave file:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx \
--zipformer2-ctc-model=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/
˓→ctc-epoch-20-avg-1-chunk-16-left-128.onnx \
--tokens=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt \
./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/test_wavs/DEV_
˓→T0000000000.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx --zipformer2-ctc-model=./sherpa-onnx-streaming-zipformer-ctc-  
multi-zh-hans-2023-12-13/ctc-epoch-20-avg-1-chunk-16-left-128.onnx --tokens=./sherpa-  
onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt ./sherpa-onnx-  
streaming-zipformer-ctc-multi-zh-hans-2023-12-13/test_wavs/DEV_T0000000000.wav  
  
OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  
dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=  
'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='', decoder=  
''), wenet_ctc=OnlineWenetCtcModelConfig(model='', chunk_size=16, num_left_chunks=4),  
zipformer2_ctc=OnlineZipformer2CtcModelConfig(model='./sherpa-onnx-streaming-zipformer-  
ctc-multi-zh-hans-2023-12-13/ctc-epoch-20-avg-1-chunk-16-left-128.onnx'), tokens='./  
sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt', num_  
threads=1, debug=False, provider='cpu', model_type=''), lm_config=OnlineLMConfig(model=  
'', scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_  
nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),  
rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_  
utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_  
silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,  

```

int8

The following code shows how to use int8 models to decode a wave file:

```
cd /path/to/sherpa-onnx  
  
./build/bin/sherpa-onnx \  
--zipformer2-ctc-model=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/  
--ctc-epoch-20-avg-1-chunk-16-left-128.int8.onnx \  
--tokens=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt \  
(continues on next page)
```

(continued from previous page)

```
./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/test_wavs/DEV_
↪T0000000000.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↪build/bin/sherpa-onnx --zipformer2-ctc-model=./sherpa-onnx-streaming-zipformer-ctc-
↪multi-zh-hans-2023-12-13/ctc-epoch-20-avg-1-chunk-16-left-128.int8.onnx --tokens=./
↪sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt ./sherpa-onnx-
↪streaming-zipformer-ctc-multi-zh-hans-2023-12-13/test_wavs/DEV_T0000000000.wav

OnlineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
↪dim=80), model_config=OnlineModelConfig(transducer=OnlineTransducerModelConfig(encoder=
↪'', decoder='', joiner=''), paraformer=OnlineParaformerModelConfig(encoder='', decoder=
↪''), wenet_ctc=OnlineWenetCtcModelConfig(model='', chunk_size=16, num_left_chunks=4),
↪zipformer2_ctc=OnlineZipformer2CtcModelConfig(model='./sherpa-onnx-streaming-zipformer-
↪ctc-multi-zh-hans-2023-12-13/ctc-epoch-20-avg-1-chunk-16-left-128.int8.onnx'), tokens=
↪"./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt", num_
↪threads=1, debug=False, provider="cpu", model_type=''), lm_config=OnlineLMConfig(model=
↪'', scale=0.5), endpoint_config=EndpointConfig(rule1=EndpointRule(must_contain_
↪nonsilence=False, min_trailing_silence=2.4, min_utterance_length=0),
↪rule2=EndpointRule(must_contain_nonsilence=True, min_trailing_silence=1.2, min_
↪utterance_length=0), rule3=EndpointRule(must_contain_nonsilence=False, min_trailing_
↪silence=0, min_utterance_length=20)), enable_endpoint=True, max_active_paths=4,
↪hotwords_score=1.5, hotwords_file='', decoding_method="greedy_search")
./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/test_wavs/DEV_T0000000000.
↪wav
Elapsed seconds: 0.44, Real time factor (RTF): 0.078

{"is_final":false, "segment":0, "start_time":0.00, "text": " ", "timestamps": [0.00, 0.
↪52, 0.76, 0.84, 1.04, 1.24, 1.96, 2.04, 2.24, 2.36, 2.56, 2.68, 2.80, 3.28, 3.40, 3.60,
↪3.72, 3.84, 3.96, 4.04, 4.16, 4.28, 4.36, 4.60, 4.84], "tokens":[" ", "", "", "", "", "",
↪"", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

Real-time speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-microphone \  
--zipformer2-ctc-model=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/  
↪ctc-epoch-20-avg-1-chunk-16-left-128.onnx \  
--tokens=./sherpa-onnx-streaming-zipformer-ctc-multi-zh-hans-2023-12-13/tokens.txt
```

Hint: If your system is Linux (including embedded Linux), you can also use *sherpa-onnx-alsa* to do real-time speech recognition with your microphone if *sherpa-onnx-microphone* does not work for you.

8.25.4 Offline transducer models

This section lists available offline transducer models.

Zipformer-transducer-based Models

Hint: Please refer to *Installation* to install *sherpa-onnx* before you read this section.

sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01 (Japanese,)

This model is from ReazonSpeech and supports only Japanese. It is trained by 35k hours of data.

The code for training the model can be found at <https://github.com/k2-fsa/icefall/tree/master/egs/reasonspeech/ASR>

Paper about the dataset is https://research.reazon.jp/_static/reasonspeech_nlp2023.pdf

In the following, we describe how to download it and use it with *sherpa-onnx*.

Hint: The original onnx model is from

<https://huggingface.co/reazon-research/reasonspeech-k2-v2>

Download the model

Please use the following commands to download it.

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
↪zipformer-ja-reasonspeech-2024-08-01.tar.bz2  
  
# For Chinese users, you can use the following mirror  
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
↪zipformer-ja-reasonspeech-2024-08-01.tar.bz2  
  
tar xvf sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01.tar.bz2
```

(continues on next page)

(continued from previous page)

```
rm sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01.tar.bz2
ls -lh sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01
```

You should see the following output:

-rw-r--r--	1	fangjun	staff	1.2K Aug 1 18:32	README.md
-rw-r--r--	1	fangjun	staff	2.8M Aug 1 18:32	decoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	11M Aug 1 18:32	decoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	148M Aug 1 18:32	encoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	565M Aug 1 18:32	encoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	2.6M Aug 1 18:32	joiner-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	10M Aug 1 18:32	joiner-epoch-99-avg-1.onnx
drwxr-xr-x	8	fangjun	staff	256B Aug 1 18:31	test_wavs
-rw-r--r--	1	fangjun	staff	45K Aug 1 18:32	tokens.txt

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01/encoder-epoch-99-avg-1.
˓→onnx \
--decoder=./sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01/joiner-epoch-99-avg-1.onnx
˓→\
--num-threads=1 \
./sherpa-onnx-zipformer-ja-reasonspeech-2024-08-01/test_wavs/1.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/encoder-epoch-99-avg-1.
˓→int8.onnx \
--decoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/joiner-epoch-99-avg-1.int8.
˓→onnx \
--num-threads=1 \
./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/test_wavs/1.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/encoder-epoch-99-avg-1.
˓→int8.onnx \
--decoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/joiner-epoch-99-avg-1.int8.
˓→onnx
```

Speech recognition from a microphone with VAD

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx

./build/bin/sherpa-onnx-vad-microphone-offline-asr \
--silero-vad-model=./silero_vad.onnx \
--tokens=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/encoder-epoch-99-avg-1.
˓→int8.onnx \
--decoder=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/decoder-epoch-99-avg-1.
˓→onnx \
--joiner=./sherpa-onnx-zipformer-ja-reazonspeech-2024-08-01/joiner-epoch-99-avg-1.int8.
˓→onnx
```

sherpa-onnx-zipformer-korean-2024-06-24 (Korean,)

PyTorch checkpoints of this model can be found at <<https://huggingface.co/johnBamma/icefall-asr-ksponspeech-zipformer-2024-06-24>>.

The training dataset can be found at <https://aihub.or.kr/aihubdata/data/view.do?currMenu=115&topMenu=100&aihubDataSe=realm&dataSetSn=123>.

Paper about the dataset is <https://www.mdpi.com/2076-3417/10/19/6936>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→zipformer-korean-2024-06-24.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→zipformer-korean-2024-06-24.tar.bz2

tar xf sherpa-onnx-zipformer-korean-2024-06-24.tar.bz2
rm sherpa-onnx-zipformer-korean-2024-06-24.tar.bz2

ls -lh sherpa-onnx-zipformer-korean-2024-06-24
```

You should see the following output:

-rw-r--r--	1	fangjun	staff	307K	Jun 24 15:33	bpe.model
-rw-r--r--	1	fangjun	staff	2.7M	Jun 24 15:33	decoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	11M	Jun 24 15:33	decoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	68M	Jun 24 15:33	encoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	249M	Jun 24 15:33	encoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	2.5M	Jun 24 15:33	joiner-epoch-99-avg-1.int8.onnx

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 9.8M Jun 24 15:33 joiner-epoch-99-avg-1.onnx
drwxr-xr-x 7 fangjun staff 224B Jun 24 15:32 test_wavs
-rw-r--r-- 1 fangjun staff 59K Jun 24 15:33 tokens.txt
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-korean-2024-06-24/tokens.txt \
--encoder=./sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-korean-2024-06-24/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:360 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-korean-2024-06-24/
↳ tokens.txt --encoder=./sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.
↳ onnx --decoder=./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx --
↳ joiner=./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.onnx ./sherpa-
↳ onnx-zipformer-korean-2024-06-24/test_wavs/0.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
↳ dim=80, low_freq=20, high_freq=-400, dither=0), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.onnx", decoder_filename=
↳ "./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx", joiner_
↳ filename="./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), whisper=OfflineWhisperModelConfig(encoder="",
decoder="", language="", task="transcribe
406, tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer=Chapter 8. sherpa-onnx
↳ ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
↳ ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc="", tokens="./sherpa-onnx-
zipformer-korean-2024-06-24/tokens.txt", num_threads=2, debug=False, provider="cpu",_
model_type="", modeling_unit="cikchar", bpe_vocab=""), lm_config=OfflineLMConfig(model_
```

(continues on next page)

(continued from previous page)

```
Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-korean-2024-06-24/test_wavs/0.wav
{"text": " ", "timestamps": [0.12, 0.24, 0.56, 1.00, 1.20, 1.32, 2.00, 2.16, 2.32, 2.52, 2.68, 2.80, 3.08, 3.28], "tokens": [" ", "", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ."], "words": []}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.119 s
Real time factor (RTF): 0.119 / 3.526 = 0.034
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-korean-2024-06-24/tokens.txt \
--encoder=./sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-korean-2024-06-24/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:360 ./
build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-korean-2024-06-24/
tokens.txt --encoder=./sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.
int8.onnx --decoder=./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.
onnx --joiner=./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.int8.
onnx ./sherpa-onnx-zipformer-korean-2024-06-24/test_wavs/0.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80, low_freq=20, high_freq=-400, dither=0), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.int8.onnx", decoder_
filename="./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx",
joiner_filename="./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.int8.
onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
```

8.25. Pre-trained models 407
 ctc=OfflineNemoCtcModelConfig(model=""),
 whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
ctc=OfflineWenetCtcModelConfig(model=""), teleSpeech_ctc="" tokens="" /sherpa-onnx

(continued from previous page)

```
Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-korean-2024-06-24/test_wavs/0.wav
{"text": " .", "timestamps": [0.12, 0.24, 0.56, 1.00, 1.20, 1.32, 2.00, 2.16, 2.32, ↵ 2.52, 2.68, 2.84, 3.08, 3.28], "tokens": [" ", "", " ", " ", "", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ", " ."], "words": []}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.092 s
Real time factor (RTF): 0.092 / 3.526 = 0.026
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-korean-2024-06-24/tokens.txt \
--encoder=./sherpa-onnx-zipformer-korean-2024-06-24/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-zipformer-korean-2024-06-24/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-korean-2024-06-24/joiner-epoch-99-avg-1.int8.onnx
```

sherpa-onnx-zipformer-thai-2024-06-20 (Thai,)

PyTorch checkpoints of this model can be found at <https://huggingface.co/yfyeung/icefall-asr-gigaspeech2-th-zipformer-2024-06-20>.

The training dataset can be found at <https://github.com/SpeechColab/GigaSpeech2>.

The paper about the dataset is <https://arxiv.org/pdf/2406.11546>.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪zipformer-thai-2024-06-20.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪zipformer-thai-2024-06-20.tar.bz2

tar xf sherpa-onnx-zipformer-thai-2024-06-20.tar.bz2
rm sherpa-onnx-zipformer-thai-2024-06-20.tar.bz2

ls -lh sherpa-onnx-zipformer-thai-2024-06-20
```

You should see the following output:

```
-rw-r--r-- 1 fangjun staff 277K Jun 20 16:47 bpe.model
-rw-r--r-- 1 fangjun staff 1.2M Jun 20 16:47 decoder-epoch-12-avg-5.int8.onnx
-rw-r--r-- 1 fangjun staff 4.9M Jun 20 16:47 decoder-epoch-12-avg-5.onnx
-rw-r--r-- 1 fangjun staff 148M Jun 20 16:47 encoder-epoch-12-avg-5.int8.onnx
-rw-r--r-- 1 fangjun staff 565M Jun 20 16:47 encoder-epoch-12-avg-5.onnx
-rw-r--r-- 1 fangjun staff 1.0M Jun 20 16:47 joiner-epoch-12-avg-5.int8.onnx
-rw-r--r-- 1 fangjun staff 3.9M Jun 20 16:47 joiner-epoch-12-avg-5.onnx
drwxr-xr-x 6 fangjun staff 192B Jun 20 16:46 test_wavs
-rw-r--r-- 1 fangjun staff 38K Jun 20 16:47 tokens.txt
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-thai-2024-06-20/tokens.txt \
--encoder=./sherpa-onnx-zipformer-thai-2024-06-20/encoder-epoch-12-avg-5.onnx \
--decoder=./sherpa-onnx-zipformer-thai-2024-06-20/decoder-epoch-12-avg-5.onnx \
--joiner=./sherpa-onnx-zipformer-thai-2024-06-20/joiner-epoch-12-avg-5.onnx \
./sherpa-onnx-zipformer-thai-2024-06-20/test_wavs/0.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx  
  
./build/bin/sherpa-onnx-offline \  
--tokens=./sherpa-onnx-zipformer-thai-2024-06-20/tokens.txt \  
--encoder=./sherpa-onnx-zipformer-thai-2024-06-20/encoder-epoch-12-avg-5.int8.onnx \  
--decoder=./sherpa-onnx-zipformer-thai-2024-06-20/decoder-epoch-12-avg-5.onnx \  
--joiner=./sherpa-onnx-zipformer-thai-2024-06-20/joiner-epoch-12-avg-5.int8.onnx \  
.sherpa-onnx-zipformer-thai-2024-06-20/test_wavs/0.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
./build/bin/sherpa-onnx-microphone-offline \  
--tokens=./sherpa-onnx-zipformer-thai-2024-06-20/tokens.txt \  
--encoder=./sherpa-onnx-zipformer-thai-2024-06-20/encoder-epoch-12-avg-5.int8.onnx \  
--decoder=./sherpa-onnx-zipformer-thai-2024-06-20/decoder-epoch-12-avg-5.onnx \  
--joiner=./sherpa-onnx-zipformer-thai-2024-06-20/joiner-epoch-12-avg-5.int8.onnx
```

sherpa-onnx-zipformer-cantonese-2024-03-13 (Cantonese,)

Training code for this model can be found at <https://github.com/k2-fsa/icefall/pull/1537>. It supports only Cantonese since it is trained on a Canatonese dataset. The paper for the dataset can be found at <https://arxiv.org/pdf/2201.02419.pdf>.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓zipformer-cantonese-2024-03-13.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓zipformer-cantonese-2024-03-13.tar.bz2

tar xf sherpa-onnx-zipformer-cantonese-2024-03-13.tar.bz2
rm sherpa-onnx-zipformer-cantonese-2024-03-13.tar.bz2

ls -lh sherpa-onnx-zipformer-cantonese-2024-03-13
```

You should see the following output:

```
total 340M
-rw-r--r-- 1 1001 127 2.7M Mar 13 09:06 decoder-epoch-45-avg-35.int8.onnx
-rw-r--r-- 1 1001 127 11M Mar 13 09:06 decoder-epoch-45-avg-35.onnx
-rw-r--r-- 1 1001 127 67M Mar 13 09:06 encoder-epoch-45-avg-35.int8.onnx
-rw-r--r-- 1 1001 127 248M Mar 13 09:06 encoder-epoch-45-avg-35.onnx
-rw-r--r-- 1 1001 127 2.4M Mar 13 09:06 joiner-epoch-45-avg-35.int8.onnx
-rw-r--r-- 1 1001 127 9.5M Mar 13 09:06 joiner-epoch-45-avg-35.onnx
drwxr-xr-x 2 1001 127 4.0K Mar 13 09:06 test_wavs
-rw-r--r-- 1 1001 127 42K Mar 13 09:06 tokens.txt
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--blank-penalty=1.2 \
--tokens=./sherpa-onnx-zipformer-cantonese-2024-03-13/tokens.txt \
--encoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.onnx \
--decoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx \
--joiner=./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.onnx \
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav \
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --blank-
˓ penalty=1.2 --tokens=./sherpa-onnx-zipformer-cantonese-2024-03-13/tokens.txt --
˓ encoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.onnx --
˓ decoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx --
˓ joiner=./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.onnx ./
˓ sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav ./sherpa-onnx-
˓ zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
˓ feature_dim=80), model_
˓ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
˓ sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.onnx", decoder_
˓ filename="./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx",
˓ joiner_filename="./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.
˓ onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
˓ ctc=OfflineNemoEncDecCtcModelConfig(model=""),
˓ whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓ ", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
˓ ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓ ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-cantonese-
˓ 2024-03-13/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
˓ lm_config=OfflineLmConfig(model="", scale=0.5), ctc_fst_decoder_
˓ config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
˓ search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=1.2)
Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav
{"text": "", "timestamps": [0.00, 0.88, 1.28, 1.52, 1.84, 2.08, 2.32, 2.56, 2.80, 3.04,
˓ 3.20, 3.44, 3.68, 3.92], "tokens":["", "", "", "", "", "", "", "", "", "", "", ""],
˓ ""}
-----
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav
{"text": "", "timestamps": [0.00, 0.64, 0.88, 1.12, 1.28, 1.60, 1.80, 2.16, 2.36, 2.56,
˓ 2.88, 3.08, 3.32, 3.44, 3.60], "tokens":["", "", "", "", "", "", "", "", "", "", ""],
˓ ""}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.349 s
```

(continues on next page)

(continued from previous page)

Real time factor (RTF): 1.349 / 10.320 = 0.131
--

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--blank-penalty=1.2 \
--tokens=./sherpa-onnx-zipformer-cantonese-2024-03-13/tokens.txt \
--encoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.int8.
˓→onnx \
--decoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx \
--joiner=./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.int8.onnx
˓→\
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav \
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --blank-
˓→penalty=1.2 --tokens=./sherpa-onnx-zipformer-cantonese-2024-03-13/tokens.txt --
˓→encoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.int8.onnx
˓→--decoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx --
˓→joiner=./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.int8.onnx ./
˓→sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav ./sherpa-onnx-
˓→zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
˓→feature_dim=80), model_
˓→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
˓→sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.int8.onnx", decoder_
˓→filename="./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx",
˓→joiner_filename="./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.
˓→int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
˓→ctc=OfflineNemoEncDecCtcModelConfig(model=""), whisper=
˓→OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓→", tail_paddings=-1), tdnn=OfflineTdnmModelConfig(model=""), zipformer_
˓→ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓→ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-cantonese-
˓→2024-03-13/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
˓→config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
˓→config=OfflineFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy
˓→search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=1.2)
```

(continues on next page)

8.25. Pre-trained Models 413

(continued from previous page)

```
Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_1.wav
{"text": "", "timestamps": [0.00, 0.88, 1.28, 1.52, 1.84, 2.08, 2.32, 2.56, 2.80, 3.04, 3.20, 3.44, 3.68, 3.92], "tokens": ["", "", "", "", "", "", "", "", "", "", "", ""], "onnx": ""}

-----
./sherpa-onnx-zipformer-cantonese-2024-03-13/test_wavs/test_wavs_2.wav
{"text": "", "timestamps": [0.00, 0.64, 0.88, 1.12, 1.28, 1.60, 1.80, 2.16, 2.36, 2.56, 2.88, 3.08, 3.32, 3.44, 3.60], "tokens": ["", "", "", "", "", "", "", "", "", "", "", ""], "onnx": ""}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.907 s
Real time factor (RTF): 0.907 / 10.320 = 0.088
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-cantonese-2024-03-13/tokens.txt \
--encoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/encoder-epoch-45-avg-35.int8 \
--onnx \
--decoder=./sherpa-onnx-zipformer-cantonese-2024-03-13/decoder-epoch-45-avg-35.onnx \
--joiner=./sherpa-onnx-zipformer-cantonese-2024-03-13/joiner-epoch-45-avg-35.int8.onnx
```

sherpa-onnx-zipformer-gigaspeech-2023-12-12 (English)

Training code for this model is <https://github.com/k2-fsa/icefall/pull/1254>. It supports only English since it is trained on the [GigaSpeech](#) dataset.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-gigaspeech-2023-12-12.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-gigaspeech-2023-12-12.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xf sherpa-onnx-zipformer-gigaspeech-2023-12-12.tar.bz2  
rm sherpa-onnx-zipformer-gigaspeech-2023-12-12.tar.bz2  
ls -lh sherpa-onnx-zipformer-qigaspeech-2023-12-12
```

You should see the following output:

```
$ ls -lh sherpa-onnx-zipformer-gigaspeech-2023-12-12
total 656184
-rw-r--r-- 1 fangjun staff 28B Dec 12 19:00 README.md
-rw-r--r-- 1 fangjun staff 239K Dec 12 19:00 bpe.model
-rw-r--r-- 1 fangjun staff 528K Dec 12 19:00 decoder-epoch-30-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Dec 12 19:00 decoder-epoch-30-avg-1.onnx
-rw-r--r-- 1 fangjun staff 68M Dec 12 19:00 encoder-epoch-30-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 249M Dec 12 19:00 encoder-epoch-30-avg-1.onnx
-rw-r--r-- 1 fangjun staff 253K Dec 12 19:00 joiner-epoch-30-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Dec 12 19:00 joiner-epoch-30-avg-1.onnx
drwxr-xr-x 5 fangjun staff 160B Dec 12 19:00 test_wavs
-rw-r--r-- 1 fangjun staff 4.9K Dec 12 19:00 tokens.txt
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnéx-offline \
--tokens=./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/tokens.txt \
--encoder=./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.onnéx \
--decoder=./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnéx \
--joiner=./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.onnéx \
./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-0001.wav \
./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav \
./sherpa-onnéx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
→ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/  
→ tokens.txt --encoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-  
→ avg-1.onnx --decoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-  
→ avg-1.onnx --joiner=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-average  
→ 1.onnx ./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-0001.wav ./  
→ sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav ./sherpa-  
8.25. Pre-trained models 415  
→ onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav
```

(continued from previous page)

```

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
    ↵ feature_dim=80), model_
    ↵ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
    ↵ sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.onnx", decoder_
    ↵ filename="./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnx",
    ↵ joiner_filename="./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.
    ↵ onnx"), paraformer=OfflineParaformerModelConfig(model ""), nemo_
    ↵ ctc=OfflineNemoEncDecCtcModelConfig(model ""), whisper=
    ↵ OfflineWhisperModelConfig(encoder "", decoder "", language "", task="transcribe
    ↵ ", tail_paddings=-1), tdnm=OfflineTdnmModelConfig(model ""), zipformer_
    ↵ ctc=OfflineZipformerCtcModelConfig(model ""), wenet_
    ↵ ctc=OfflineWenetCtcModelConfig(model ""), tokens="./sherpa-onnx-zipformer-gigaspeech-
    ↵ 2023-12-12/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type ""),
    ↵ config=OfflineLMConfig(model "", scale=0.5), ctc_fst_decoder_
    ↵ config=OfflineCtcFstDecoderConfig(graph "", max_active=3000), decoding_method="greedy_
    ↵ search", max_active_paths=4, hotwords_file "", hotwords_score=1.5)
Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-0001.wav
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
    ↵ SQUALID QUARTER OF THE BROTHELS", "timestamps": [0.00, 0.36, 0.52, 0.68, 0.96, 1.00, 1.
    ↵ 08, 1.28, 1.40, 1.48, 1.60, 1.76, 1.80, 1.88, 1.92, 2.00, 2.20, 2.32, 2.36, 2.48, 2.60,
    ↵ 2.80, 2.84, 2.92, 3.12, 3.32, 3.56, 3.76, 4.04, 4.20, 4.32, 4.40, 4.56, 4.80, 4.92, 5.
    ↵ 08, 5.36, 5.48, 5.64, 5.72, 5.88, 6.04, 6.24], "tokens": [" AFTER", " E", " AR", " LY", "
    ↵ ", " N", " IGH", " F", " AL", " L", " THE", " ", " Y", " E", " LL", " OW", " LA", " M", " P", " S
    ↵ ", " WOULD", " ", " L", " IGH", " UP", " HERE", " AND", " THERE", " THE", " S", " QU",
    ↵ " AL", " ID", " QU", " AR", " TER", " OF", " THE", " B", " RO", " TH", " EL", " S"]}

----
```

```

./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav
{"text": " GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER_
    ↵ A LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR_
    ↵ EVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN",
    "timestamps": [0.00, 0.16, 0.40, 0.68, 0.84, 0.96, 1.04, 1.12, 1.32, 1.52, 1.68, 1.76,
    ↵ 2.00, 2.12, 2.28, 2.40, 2.64, 2.92, 3.20, 3.32, 3.52, 3.64, 3.76, 3.96, 4.12, 4.36, 4.
    ↵ 52, 4.72, 4.92, 5.16, 5.40, 5.64, 5.76, 5.88, 6.12, 6.28, 6.48, 6.84, 7.08, 7.32, 7.60,
    ↵ 7.92, 8.12, 8.24, 8.36, 8.48, 8.64, 8.76, 8.88, 9.12, 9.32, 9.48, 9.56, 9.60, 9.76,
    ↵ 10.00, 10.12, 10.20, 10.44, 10.68, 10.80, 11.00, 11.20, 11.36, 11.52, 11.76, 12.00, 12.
    ↵ 12, 12.24, 12.28, 12.52, 12.72, 12.84, 12.96, 13.04, 13.24, 13.40, 13.64, 13.76, 14.00,
    ↵ 14.08, 14.24, 14.52, 14.68, 14.80, 15.00, 15.04, 15.28, 15.52, 15.76, 16.00, 16.12,
    ↵ 16.20, 16.32], "tokens": [" GO", " D", " AS", " A", " DI", " RE", " C", " T", " CON", " SE",
    ↵ " QU", " ENCE", " OF", " THE", " S", " IN", " WHICH", " MAN", " TH", " US", " P", " UN",
    ↵ " ISH", " ED", " HAD", " GIVE", " N", " HER", " A", " LOVE", " LY", " CHI", " L", " D", " WHO
    ↵ ", " SE", " PLACE", " WAS", " ON", " THAT", " SAME", " DIS", " HO", " N", " OR", " ED", " BO
    ↵ ", " S", " OM", " TO", " CON", " NE", " C", " T", " HER", " PA", " R", " ENT", " FOR", " E",
    ↵ " VER", " WITH", " THE", " RA", " CE", " AND", " DE", " S", " C", " ENT", " OF", " MO", " R",
    ↵ " T", " AL", " S", " AND", " TO", " BE", " F", " IN", " ALLY", " A", " B", " LES", " S", " ED
    ↵ ", " SO", " UL", " IN", " HE", " A", " VE", " N"]}

----
```

```

./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav

```

(continues on next page)

(continued from previous page)

```
{"text": " YET THESE THOUGHTS AFFECTED HESTER PRYNE LESS WITH HOPE THAN APPREHENSION",
  "timestamps": [0.00, 0.04, 0.12, 0.40, 0.68, 0.88, 0.96, 1.12, 1.20, 1.32, 1.44, 1.48,
  ↵ 1.64, 1.76, 1.88, 2.04, 2.16, 2.28, 2.52, 2.68, 2.72, 2.88, 3.12, 3.28, 3.52, 3.80, 4.
  ↵ 00, 4.16, 4.24, 4.40, 4.48], "tokens": [" ", "Y", "ET", " THESE", " THOUGH", "T", "S",
  ↵ " A", "FF", "E", "C", "TED", " HE", "S", "TER", " P", "RY", "NE", " LE", "S", "S", "WITH",
  ↵ " HO", "PE", " THAN", " APP", "RE", "HE", "N", "S", "ION"]}
```

```
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.407 s
Real time factor (RTF): 1.407 / 28.165 = 0.050
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.int8.
↪ onnx \
--decoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.int8.onnx
↪ \
./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-0001.wav \
./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav \
./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↪ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/
↪ tokens.txt --encoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-
↪ avg-1.int8.onnx --decoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-
↪ 30-avg-1.onnx --joiner=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-
↪ avg-1.int8.onnx ./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-
↪ 0001.wav ./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav .
↪ /sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↪ feature_dim=80), model_
↪ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↪ sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.int8.onnx",
decoder_
↪ filename="./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnx",
↪ joiner_filename="./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.
↪ int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
↪ ctc=OfflineNemoEncDecCtcModelConfig(model=""),
↪ whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe"
↪ ", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
```

8.25. Pre-trained models 417

```
↪ ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
↪ ctc=OfflineWenetCtcModelConfig(model=""), tokens=".sherpa-onnx-zipformer-gigaspeech-
↪ 2023-12-12/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
lm_
↪ config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
↪ config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy"
```

(continued from previous page)

```

Creating recognizer ...
Started
Done!

./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1089-134686-0001.wav
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↪SQUALID QUARTER OF THE BROTHELS", "timestamps": [0.00, 0.36, 0.52, 0.68, 0.96, 1.00, 1.
↪08, 1.28, 1.40, 1.48, 1.60, 1.76, 1.80, 1.88, 1.92, 2.00, 2.20, 2.32, 2.36, 2.48, 2.60,
↪2.80, 2.84, 2.92, 3.12, 3.32, 3.56, 3.76, 4.04, 4.24, 4.32, 4.40, 4.56, 4.80, 4.92, 5.
↪08, 5.36, 5.48, 5.64, 5.72, 5.88, 6.04, 6.24], "tokens": ["AFTER", "E", "AR", "LY",
↪", "N", "IGHT", "F", "AL", "L", "THE", " ", "Y", "E", "LL", "OW", "LA", "M", "P", "S
↪", "WOULD", " ", "L", "IGHT", "UP", "HERE", "AND", "THERE", "THE", "S", "QU",
↪"AL", "ID", "QU", "AR", "TER", "OF", "THE", "B", "RO", "TH", "EL", "S"]}

-----
./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0001.wav
{"text": " GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER_
↪A LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR_
↪EVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN",
↪"timestamps": [0.00, 0.16, 0.40, 0.68, 0.84, 0.96, 1.08, 1.12, 1.32, 1.52, 1.68, 1.76,
↪2.00, 2.12, 2.28, 2.40, 2.64, 2.92, 3.20, 3.32, 3.52, 3.64, 3.76, 3.96, 4.12, 4.36, 4.
↪52, 4.72, 4.92, 5.16, 5.40, 5.64, 5.76, 5.88, 6.12, 6.28, 6.52, 6.84, 7.08, 7.32, 7.60,
↪7.92, 8.12, 8.24, 8.36, 8.48, 8.64, 8.76, 8.88, 9.12, 9.32, 9.48, 9.56, 9.60, 9.76,
↪10.00, 10.12, 10.20, 10.44, 10.68, 10.80, 11.00, 11.20, 11.36, 11.52, 11.76, 12.00, 12.
↪12, 12.24, 12.28, 12.52, 12.72, 12.84, 12.96, 13.04, 13.24, 13.44, 13.64, 13.76, 14.00,
↪14.08, 14.24, 14.52, 14.68, 14.80, 15.00, 15.04, 15.28, 15.48, 15.76, 16.00, 16.12,
↪16.16, 16.32], "tokens": ["GO", "D", "AS", "A", "DI", "RE", "C", "T", "CON", "SE",
↪"QU", "ENCE", "OF", "THE", "S", "IN", "WHICH", "MAN", "TH", "US", "P", "UN",
↪"ISH", "ED", "HAD", "GIVE", "N", "HER", "A", "LOVE", "LY", "CHI", "L", "D", "WHO
↪", "SE", "PLACE", "WAS", "ON", "THAT", "SAME", "DIS", "HO", "N", "OR", "ED", "BO
↪", "S", "OM", "TO", "CON", "NE", "C", "T", "HER", "PA", "R", "ENT", "FOR", "E",
↪"VER", "WITH", "THE", "RA", "CE", "AND", "DE", "S", "C", "ENT", "OF", "MO", "R",
↪"T", "AL", "S", "AND", "TO", "BE", "F", "IN", "ALLY", "A", "B", "LES", "S", "ED
↪", "SO", "UL", "IN", "HE", "A", "VE", "N"]}

-----
./sherpa-onnx-zipformer-gigaspeech-2023-12-12/test_wavs/1221-135766-0002.wav
{"text": " YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
↪"timestamps": [0.00, 0.04, 0.12, 0.40, 0.68, 0.88, 0.96, 1.12, 1.24, 1.32, 1.44, 1.48,
↪1.64, 1.76, 1.88, 2.04, 2.16, 2.28, 2.32, 2.52, 2.68, 2.72, 2.88, 3.12, 3.32, 3.52, 3.
↪80, 4.00, 4.16, 4.24, 4.40, 4.48], "tokens": ["", "Y", "ET", "THESE", "THOUGH", "T",
↪"S", "A", "FF", "E", "C", "TED", "HE", "S", "TER", "P", "RY", "N", "NE", "LE", "S",
↪"S", "WITH", "HO", "PE", "THAN", "APP", "RE", "HE", "N", "S", "ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.101 s
Real time factor (RTF): 1.101 / 28.165 = 0.039

```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.int8.
→onnx \
--decoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.int8.onnx
```

Speech recognition from a microphone with VAD

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.
→onnx

./build/bin/sherpa-onnx-vad-microphone-offline-asr \
--silero-vad-model=./silero_vad.onnx \
--tokens=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/tokens.txt \
--encoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/encoder-epoch-30-avg-1.int8.
→onnx \
--decoder=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/decoder-epoch-30-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-gigaspeech-2023-12-12/joiner-epoch-30-avg-1.int8.onnx
```

[zrjin/sherpa-onnx-zipformer-multi-zh-hans-2023-9-2 \(Chinese\)](#)

This model is from

<https://huggingface.co/zrjin/sherpa-onnx-zipformer-multi-zh-hans-2023-9-2>

which supports Chinese as it is trained on whatever datasets involved in the `multi-zh_hans` recipe.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/1238>.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
→zipformer-multi-zh-hans-2023-9-2.tar.bz2
```

(continues on next page)

(continued from previous page)

```
# For Chinese users, you can use the following mirror  
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
zipformer-multi-zh-hans-2023-9-2.tar.bz2
```

```
tar xvf sherpa-onnx-zipformer-multi-zh-hans-2023-9-2.tar.bz2  
rm sherpa-onnx-zipformer-multi-zh-hans-2023-9-2.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-zipformer-multi-zh-hans-2023-9-2 zengruijin$ ls -lh *.onnx  
-rw-rw-r--@ 1 zengruijin staff 1.2M Sep 18 07:04 decoder-epoch-20-avg-1.int8.onnx  
-rw-rw-r--@ 1 zengruijin staff 4.9M Sep 18 07:04 decoder-epoch-20-avg-1.onnx  
-rw-rw-r--@ 1 zengruijin staff 66M Sep 18 07:04 encoder-epoch-20-avg-1.int8.onnx  
-rw-rw-r--@ 1 zengruijin staff 248M Sep 18 07:05 encoder-epoch-20-avg-1.onnx  
-rw-rw-r--@ 1 zengruijin staff 1.0M Sep 18 07:05 joiner-epoch-20-avg-1.int8.onnx  
-rw-rw-r--@ 1 zengruijin staff 3.9M Sep 18 07:05 joiner-epoch-20-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-offline \  
--tokens=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/tokens.txt \  
--encoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.onnx \  
--decoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-1.onnx \  
--joiner=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.onnx \  
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav \  
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/1.wav \  
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361  
sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/tokens.txt  
--encoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.onnx --  
decoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-1.onnx --  
joiner=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.onnx ./  
sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav ./sherpa-onnx-zipformer-  
multi-zh-hans-2023-9-2/test_wavs/1.wav ./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/8k.wav
```

(continued from previous page)

```

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
    ↵ feature_dim=80), model_
    ↵ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
    ↵ sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.onnx", decoder_
    ↵ filename="./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-1.onnx", ↵
    ↵ joiner_filename="./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.
    ↵ onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
    ↵ ctc=OfflineNemoEncDecCtcModelConfig(model=""), ↵
    ↵ whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
    ↵ "), tdnn=OfflineTdnmModelConfig(model=""), tokens="./sherpa-onnx-zipformer-multi-zh-
    ↵ hans-2023-9-2/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
    ↵ lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
    ↵ active_paths=4, hotwords_file="", hotwords_score=1.5)
Creating recognizer ...
Started
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
    ↵ cc:AcceptWaveformImpl:117 Creating a resampler:
        in_sample_rate: 8000
        output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav
{"text":" ","timestamps":[0.00, 0.16, 0.40, 0.60, 0.84, 1.08, 1.60, 1.72, 1.88, 2.04, 2.
    ↵ 24, 2.44, 2.60, 2.96, 3.12, 3.32, 3.40, 3.60, 3.72, 3.84, 4.00, 4.16, 4.32, 4.52, 4.68]
    ↵ , "tokens":[" ","","","","","","","","","","","","","","","","","","","","","","","","","","","",
    ↵ ""]}
---
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/1.wav
{"text":"<0xE8><0x8D><0xA1>","timestamps":[0.00, 0.12, 0.48, 0.68, 0.92, 1.12, 1.28, 1.
    ↵ 48, 1.80, 2.04, 2.40, 2.56, 2.76, 2.96, 3.08, 3.32, 3.48, 3.68, 3.84, 4.00, 4.20, 4.24,
    ↵ 4.28, 4.40, 4.60, 4.84],"tokens":[" ","","","","","","","","","","","","","","","","","",
    ↵ "","","","","","<0xE8>","<0x8D>","<0xA1>","","","",""]}
---
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/8k.wav
{"text":"<0xE8><0x8D><0xA1>","timestamps":[0.00, 0.04, 0.24, 0.52, 0.76, 1.00, 1.40, 1.
    ↵ 64, 1.80, 2.12, 2.32, 2.64, 2.80, 3.00, 3.20, 3.24, 3.28, 3.44, 3.64, 3.76, 3.96, 4.20]
    ↵ , "tokens":[" ","","","","","","","","","","","","","<0xE8>","<0x8D>","<0xA1>","","",
    ↵ "","","","",""]}
---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.362 s
Real time factor (RTF): 0.362 / 15.289 = 0.024

```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/tokens.txt \
--encoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.int8.
˓→onnx \
--decoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.int8.
˓→onnx \
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav \
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/1.wav \
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361
˓→sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/tokens.txt \
˓→--encoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.int8.
˓→onnx --decoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-1.
˓→onnx --joiner=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.
˓→int8.onnx ./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav ./sherpa-onnx-
˓→zipformer-multi-zh-hans-2023-9-2/test_wavs/1.wav ./sherpa-onnx-zipformer-multi-zh-hans-
˓→2023-9-2/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
˓→feature_dim=80), model_
˓→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
˓→sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.int8.onnx",
˓→decoder_filename="./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-20-avg-
˓→1.onnx", joiner_filename="./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-
˓→20-avg-1.int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""),
˓→nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),
˓→whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓→"), tdnn=OfflineTdnnModelConfig(model=""), tokens="./sherpa-onnx-zipformer-multi-zh-
˓→hans-2023-9-2/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
˓→lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
˓→active_paths=4, hotwords_file="", hotwords_score=1.5)
Creating recognizer ...
Started
/Users/runner/work/sherpa-onnx/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
˓→cc:AcceptWaveformImpl:117 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/0.wav
{"text":" ","timestamps":[0.00, 0.16, 0.40, 0.60, 0.84, 1.08, 1.60, 1.72, 1.88, 2.04, 2.28, 2.44, 2.60, 2.96, 3.12, 3.32, 3.40, 3.60, 3.76, 3.84, 4.00, 4.16, 4.32, 4.52, 4.56]
,"tokens":[" ","","","","","","","","","","","","","","","","","","","","","","","","","","","","",""]
]}
-----
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/1.wav
{"text":"<0xE8><0x8D><0xA1>","timestamps":[0.00, 0.12, 0.48, 0.68, 0.92, 1.12, 1.28, 1.48, 1.80, 2.04, 2.40, 2.56, 2.76, 2.96, 3.08, 3.32, 3.48, 3.68, 3.84, 4.00, 4.20, 4.24, 4.28, 4.40, 4.60, 4.84],"tokens":[" ","","","","","","","","","","","","","","","","","","","","","","",""]
,"tokens":["<0xE8>","<0x8D>","<0xA1>","","","",""]
]}
-----
./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/test_wavs/8k.wav
{"text":"<0xE8><0x8D><0xA1>","timestamps":[0.00, 0.04, 0.24, 0.52, 0.76, 1.00, 1.40, 1.64, 1.80, 2.12, 2.36, 2.64, 2.80, 3.04, 3.16, 3.20, 3.24, 3.44, 3.64, 3.76, 3.96, 4.20]
,"tokens":[" ","","","","","","","","","","","","","","","<0xE8>","<0x8D>","<0xA1>","",""
,"","","","",""]
]}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.305 s
Real time factor (RTF): 0.305 / 15.289 = 0.020
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/tokens.txt \
--encoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/encoder-epoch-20-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/decoder-epoch-0-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-multi-zh-hans-2023-9-2/joiner-epoch-20-avg-1.onnx
```

[yfyeung/icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17](https://huggingface.co/yfyeung/icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17) (English)

This model is from

<https://huggingface.co/yfyeung/icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17>
which supports only English as it is trained on the [CommonVoice](#) English dataset.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/997>.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-cv-
↪corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-
↪cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17.tar.bz2

tar xvf icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17.
↪tar.bz2
rm icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17.tar.
↪bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17 fangjun
↪$ ls -lh exp/*epoch-60-avg-20*.onnx
-rw-r--r-- 1 fangjun staff 1.2M Jun 27 09:53 exp/decoder-epoch-60-avg-20.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Jun 27 09:54 exp/decoder-epoch-60-avg-20.onnx
-rw-r--r-- 1 fangjun staff 121M Jun 27 09:54 exp/encoder-epoch-60-avg-20.int8.onnx
-rw-r--r-- 1 fangjun staff 279M Jun 27 09:55 exp/encoder-epoch-60-avg-20.onnx
-rw-r--r-- 1 fangjun staff 253K Jun 27 09:53 exp/joiner-epoch-60-avg-20.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Jun 27 09:53 exp/joiner-epoch-60-avg-20.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↪04-17/data/lang_bpe_500/tokens.txt \
--encoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↪04-17/exp/encoder-epoch-60-avg-20.onnx \
--decoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↪04-17/exp/decoder-epoch-60-avg-20.onnx \
--joiner=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↪04-17/exp/joiner-epoch-60-avg-20.onnx \
```

(continues on next page)

(continued from previous page)

```
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1089-134686-0001.wav \
↳ ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1221-135766-0001.wav \
↳ ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1221-135766-0002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./icefall-asr-cv-corpus-13.0-2023-03-09-en-
↳ pruned-transducer-stateless7-2023-04-17/data/lang_bpe_500/tokens.txt --encoder=./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ encoder-epoch-60-avg-20.onnx --decoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-
↳ pruned-transducer-stateless7-2023-04-17/exp/decoder-epoch-60-avg-20.onnx --joiner=./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ joiner-epoch-60-avg-20.onnx ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-
↳ transducer-stateless7-2023-04-17/test_wavs/1089-134686-0001.wav ./icefall-asr-cv-
↳ corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_wavs/1221-
↳ 135766-0001.wav ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-
↳ stateless7-2023-04-17/test_wavs/1221-135766-0002.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ encoder-epoch-60-avg-20.onnx", decoder_filename="./icefall-asr-cv-corpus-13.0-2023-03-
↳ 09-en-pruned-transducer-stateless7-2023-04-17/exp/decoder-epoch-60-avg-20.onnx",
↳ joiner_filename="./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-
↳ stateless7-2023-04-17/exp/joiner-epoch-60-avg-20.onnx"),
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-cv-corpus-13.0-
↳ 2023-03-09-en-pruned-transducer-stateless7-2023-04-17/data/lang_bpe_500/tokens.txt",
↳ num_threads=2, debug=False, provider="cpu"), lm_config=OfflineLMConfig(model="",
↳ scale=0.5), decoding_method="greedy_search", max_active_paths=4, context_score=1.5)
Creating recognizer ...
Started
Done!

./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_
↳ wavs/1089-134686-0001.wav
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↳ SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.64, 0.76, 0.84, 1.04, 1.08, 1.
↳ 16, 1.32, 1.44, 1.56, 1.72, 1.84, 1.88, 1.92, 1.96, 2.04, 2.16, 2.32, 2.48, 2.56, 2.76,
↳ 2.80, 2.84, 3.08, 3.28, 3.40, 3.52, 3.68, 4.00, 4.24, 4.28, 4.52, 4.68, 4.84, 4.88, 4.
↳ 96, 5.04, 5.28, 5.40, 5.52, 5.72, 5.88, 6.08]", "tokens": ["AFTER", "E", "AR", "LY", " ", "N
↳ ", "IGHT", "F", "AL", "L", "THE", " ", "Y", "E", "LL", "OW", "LA", "MP", "S", "WOULD", " ", "L",
↳ "IGHT", "UP", "HE", "RE", "AND", "THERE", "THE", "S", "QUA", "LI", "D", " ", "QUA", "R", "TER",
↳ "OF", "THE", "BRO", "TH", "EL", "S"]}
```

(continues on next page)

(continued from previous page)

```
-----
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_
↳ wavs/1221-135766-0001.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
↳ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT_
↳ FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
↳ ","timestamps":[0.04, 0.44, 0.64, 0.84, 0.96, 1.32, 1.52, 1.68, 1.84, 1.88, 2.04, 2.
↳ 16, 2.32, 2.40, 2.64, 2.88, 3.12, 3.24, 3.44, 3.52, 3.72, 3.88, 4.20, 4.40, 4.48, 4.60,
↳ 4.76, 4.96, 5.08, 5.24, 5.36, 5.56, 5.80, 6.20, 6.32, 6.52, 6.92, 7.16, 7.36, 7.60, 7.
↳ 76, 7.92, 8.16, 8.28, 8.40, 8.48, 8.60, 8.76, 8.84, 9.08, 9.24, 9.44, 9.48, 9.72, 9.88,
↳ 10.04, 10.12, 10.52, 10.76, 10.84, 11.08, 11.24, 11.36, 11.60, 11.76, 11.96, 12.08,_
↳ 12.24, 12.28, 12.48, 12.72, 12.84, 12.92, 13.00, 13.20, 13.52, 13.76, 13.88, 14.08, 14.
↳ 28, 14.52, 14.64, 14.76, 14.96, 15.04, 15.24, 15.48, 15.68, 15.84, 16.00, 16.04],
↳ "tokens":[" GO", "D", " AS", " A", " DIRECT", " CON", "SE", "QUE", "N", "CE", " OF", " THE", " S",
↳ "IN", " WHICH", " MAN", " TH", "US", " P", "UN", "ISH", "ED", " HAD", " G", "IVE", "N", " HER", " A",
↳ " LO", "VE", "LY", " CHI", "LD", " WHO", "SE", " PLACE", " WAS", " ON", " THAT", " SA", "ME", " DIS
↳ ", "HO", "N", "OR", "ED", " BO", "S", "OM", " TO", " CON", "N", "ECT", " HER", " PA", "R", "ENT", " FOR
↳ ", "E", "VER", " WITH", " THE", " RA", "CE", " AND", " DE", "S", "C", "ENT", " OF", " MO", "R", "T",
↳ "AL", "S", " AND", " TO", " BE", " FIN", "ALLY", " A", " B", "LES", "S", "ED", " SO", "UL", " IN", "HE",
↳ "A", "VEN"]}

-----
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_
↳ wavs/1221-135766-0002.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRIN LESS WITH HOPE THAN APPREHENSION",
↳ "timestamps":[0.00, 0.04, 0.12, 0.56, 0.80, 0.88, 1.00, 1.04, 1.12, 1.20, 1.28, 1.40,
↳ 1.52, 1.64, 1.76, 1.84, 2.04, 2.24, 2.40, 2.64, 2.68, 2.84, 3.04, 3.24, 3.44, 3.52, 3.
↳ 72, 3.92, 4.00, 4.16, 4.24, 4.36], "tokens":[" ", "Y", "ET", " THESE", " TH", "O", "UGH", "T",
↳ "S", " A", "FF", "ECT", "ED", " HE", "S", "TER", " PRI", "N", " LE", "S", "S", " WITH", " HO", "PE",
↳ "TH", "AN", " APP", "RE", "HE", "N", "S", "ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.611 s
Real time factor (RTF): 1.611 / 28.165 = 0.057
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/data/lang_bpe_500/tokens.txt \
--encoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/encoder-epoch-60-avg-20.int8.onnx \
--decoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/decoder-epoch-60-avg-20.onnx \
--joiner=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/joiner-epoch-60-avg-20.int8.onnx \
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1089-134686-0001.wav \
```

(continues on next page)

(continued from previous page)

```
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1221-135766-0001.wav \
↳ ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/
↳ test_wavs/1221-135766-0002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./icefall-asr-cv-corpus-13.0-2023-03-09-en-
↳ pruned-transducer-stateless7-2023-04-17/data/lang_bpe_500/tokens.txt --encoder=./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ encoder-epoch-60-avg-20.int8.onnx --decoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-
↳ pruned-transducer-stateless7-2023-04-17/exp/decoder-epoch-60-avg-20.onnx --joiner=./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ joiner-epoch-60-avg-20.int8.onnx ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-
↳ transducer-stateless7-2023-04-17/test_wavs/1089-134686-0001.wav ./icefall-asr-cv-
↳ corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_wavs/1221-
↳ 135766-0001.wav ./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-
↳ stateless7-2023-04-17/test_wavs/1221-135766-0002.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/
↳ encoder-epoch-60-avg-20.int8.onnx", decoder_filename="./icefall-asr-cv-corpus-13.0-
↳ 2023-03-09-en-pruned-transducer-stateless7-2023-04-17/exp/decoder-epoch-60-avg-20.onnx
↳ ", joiner_filename="./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-
↳ stateless7-2023-04-17/exp/joiner-epoch-60-avg-20.int8.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-cv-corpus-13.0-
↳ 2023-03-09-en-pruned-transducer-stateless7-2023-04-17/data/lang_bpe_500/tokens.txt",
↳ num_threads=2, debug=False, provider="cpu"), lm_config=OfflineLMConfig(model="",
↳ scale=0.5), decoding_method="greedy_search", max_active_paths=4, context_score=1.5)
Creating recognizer ...
Started
Done!

./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_
↳ wavs/1089-134686-0001.wav
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↳ SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.64, 0.76, 0.84, 1.04, 1.08, 1.
↳ 16, 1.36, 1.44, 1.56, 1.72, 1.84, 1.88, 1.92, 1.96, 2.04, 2.20, 2.32, 2.48, 2.56, 2.76,
↳ 2.80, 2.84, 3.08, 3.28, 3.40, 3.52, 3.68, 4.00, 4.24, 4.28, 4.52, 4.68, 4.84, 4.88, 4.
↳ 96, 5.04, 5.28, 5.36, 5.52, 5.72, 5.88, 6.08]", "tokens": ["AFTER", "E", "AR", "LY", "N", "IGHT", "F", "AL", "L", "THE", "Y", "E", "LL", "OW", "LA", "MP", "S", "WOULD", "L", "IGHT", "UP", "HE", "RE", "AND", "THERE", "THE", "S", "QUA", "LI", "D", "QUA", "R", "TER", "OF", "THE", "BRO", "TH", "EL", "S"]}

...
./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_
↳ wavs/1221-135766-0001.wav
```

(continues on next page)

(continued from previous page)

```
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_\n↳ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT_\n↳ FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN_\n", "timestamps": "[0.04, 0.44, 0.64, 0.84, 0.96, 1.32, 1.52, 1.68, 1.84, 1.88, 2.04, 2_\n↳ 16, 2.32, 2.40, 2.64, 2.88, 3.12, 3.24, 3.44, 3.52, 3.72, 3.88, 4.20, 4.40, 4.48, 4.60,\n↳ 4.76, 4.96, 5.08, 5.24, 5.36, 5.56, 5.80, 6.20, 6.32, 6.52, 6.92, 7.16, 7.32, 7.60, 7_\n↳ 76, 7.92, 8.16, 8.28, 8.40, 8.48, 8.60, 8.76, 8.84, 9.08, 9.24, 9.44, 9.48, 9.72, 9.88,\n↳ 10.04, 10.12, 10.52, 10.76, 10.84, 11.08, 11.24, 11.36, 11.60, 11.76, 11.96, 12.08,\n↳ 12.24, 12.28, 12.48, 12.72, 12.84, 12.92, 13.00, 13.20, 13.52, 13.76, 13.88, 14.08, 14_\n↳ 28, 14.52, 14.64, 14.76, 14.96, 15.04, 15.24, 15.48, 15.68, 15.84, 16.00, 16.04]",\n↳ "tokens": ["GO", "D", "AS", "A", "DIRECT", "CON", "SE", "QUE", "N", "CE", "OF", "THE", "S",\n↳ "IN", "WHICH", "MAN", "TH", "US", "P", "UN", "ISH", "ED", "HAD", "G", "IVE", "N", "HER", "A",\n↳ "LO", "VE", "LY", "CHI", "LD", "WHO", "SE", "PLACE", "WAS", "ON", "THAT", "SA", "ME", "DIS",\n↳ "HO", "N", "OR", "ED", "BO", "S", "OM", "TO", "CON", "N", "ECT", "HER", "PA", "R", "ENT", "FOR",\n↳ "E", "VER", "WITH", "THE", "RA", "CE", "AND", "DE", "S", "C", "ENT", "OF", "MO", "R", "T",\n↳ "AL", "S", "AND", "TO", "BE", "FIN", "ALLY", "A", "B", "LES", "S", "ED", "SO", "UL", "IN", "HE",\n↳ "A", "VEN"]}\n----\n./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-04-17/test_\n↳ wavs/1221-135766-0002.wav\n{"text":" YET THESE THOUGHTS AFFECTED HESTER PRIN LESS WITH HOPE THAN APPREHENSION",\n↳ "timestamps": "[0.00, 0.04, 0.12, 0.56, 0.80, 0.88, 1.00, 1.04, 1.12, 1.20, 1.28, 1.40,\n↳ 1.52, 1.64, 1.76, 1.84, 2.04, 2.24, 2.40, 2.64, 2.68, 2.84, 3.04, 3.24, 3.44, 3.52, 3.\n↳ 72, 3.92, 4.00, 4.16, 4.24, 4.36]", "tokens": [" ", "Y", "ET", "THESE", "TH", "O", "UGH", "T",\n↳ "S", "A", "FF", "ECT", "ED", "HE", "S", "TER", "PRI", "N", "LE", "S", "S", "WITH", "HO", "PE", "\n↳ TH", "AN", "APP", "RE", "HE", "N", "S", "ION"]}\n----\nnum threads: 2\ndecoding method: greedy_search\nElapsed seconds: 1.368 s\nReal time factor (RTF): 1.368 / 28.165 = 0.049
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx\n\n./build/bin/sherpa-onnx-microphone-offline \
--tokens=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/data/lang_bpe_500/tokens.txt \
--encoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/encoder-epoch-60-avg-20.onnx \
--decoder=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/decoder-epoch-60-avg-20.onnx \
--joiner=./icefall-asr-cv-corpus-13.0-2023-03-09-en-pruned-transducer-stateless7-2023-
↳ 04-17/exp/joiner-epoch-60-avg-20.onnx
```

pkufool/icefall-asr-zipformer-wenetspeech-20230615 (Chinese)

This model is from

<https://huggingface.co/pkufool/icefall-asr-zipformer-wenetspeech-20230615>

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

If you are interested in how the model is trained, please refer to <https://github.com/k2-fsa/icefall/pull/1130>.

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-
zipformer-wenetspeech-20230615.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/icefall-asr-
zipformer-wenetspeech-20230615.tar.bz2

tar xvf icefall-asr-zipformer-wenetspeech-20230615.tar.bz2
rm icefall-asr-zipformer-wenetspeech-20230615.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
icefall-asr-zipformer-wenetspeech-20230615 fangjun$ ls -lh exp/*.onnx
-rw-r--r-- 1 fangjun staff 11M Jun 26 14:31 exp/decoder-epoch-12-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 12M Jun 26 14:31 exp/decoder-epoch-12-avg-4.onnx
-rw-r--r-- 1 fangjun staff 66M Jun 26 14:32 exp/encoder-epoch-12-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 248M Jun 26 14:34 exp/encoder-epoch-12-avg-4.onnx
-rw-r--r-- 1 fangjun staff 2.7M Jun 26 14:31 exp/joiner-epoch-12-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 11M Jun 26 14:31 exp/joiner-epoch-12-avg-4.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-zipformer-wenetspeech-20230615/data/lang_char/tokens.txt \
--encoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/encoder-epoch-12-avg-4.onnx \
--decoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/decoder-epoch-12-avg-4.onnx \
--joiner=./icefall-asr-zipformer-wenetspeech-20230615/exp/joiner-epoch-12-avg-4.onnx \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000001.wav \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx-offline --tokens=./icefall-asr-zipformer-wenetspeech-20230615/
data/lang_char/tokens.txt --encoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/
encoder-epoch-12-avg-4.onnx --decoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/
decoder-epoch-12-avg-4.onnx --joiner=./icefall-asr-zipformer-wenetspeech-20230615/exp/
joiner-epoch-12-avg-4.onnx ./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_
T0000000000.wav ./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000001.
wav ./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000002.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
icefall-asr-zipformer-wenetspeech-20230615/exp/encoder-epoch-12-avg-4.onnx", decoder_
filename="./icefall-asr-zipformer-wenetspeech-20230615/exp/decoder-epoch-12-avg-4.onnx
", joiner_filename="./icefall-asr-zipformer-wenetspeech-20230615/exp/joiner-epoch-12-
avg-4.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-zipformer-
wenetspeech-20230615/data/lang_char/tokens.txt", num_threads=2, debug=False, provider=
"cpu"), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search
", max_active_paths=4, context_score=1.5)
Creating recognizer ...
Started
Done!
```

(continues on next page)

(continued from previous page)

```
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000000.wav
{"text":"","timestamps":[0.00, 0.12, 0.48, 0.64, 0.88, 1.16, 1.64, 1.76, 1.92, 2.08, 2.
↪32, 2.48, 2.64, 3.08, 3.20, 3.40, 3.48, 3.64, 3.76, 3.88, 3.96, 4.12, 4.28, 4.52, 4.72,
↪4.84],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","","","",
↪"","",""]}

-----
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000001.wav
{"text":"","timestamps":[0.00, 0.16, 0.48, 0.72, 0.92, 1.08, 1.28, 1.52, 1.92, 2.08, 2.
↪52, 2.64, 2.88, 3.04, 3.20, 3.40, 3.56, 3.76, 3.84, 4.00, 4.16, 4.32, 4.56, 4.84],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","","",""]}

-----
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000002.wav
{"text":"","timestamps":[0.00, 0.32, 0.56, 0.84, 1.12, 1.44, 1.68, 1.84, 2.28, 2.48, 2.
↪76, 2.92, 3.12, 3.28, 3.44, 3.60, 3.72, 3.92, 4.20],"tokens":["","","","","","","","","","","","","","","","","","","","",""]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.458 s
Real time factor (RTF): 0.458 / 15.289 = 0.030
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-zipformer-wenetspeech-20230615/data/lang_char/tokens.txt \
--encoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/encoder-epoch-12-avg-4.int8.
↪onnx \
--decoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/decoder-epoch-12-avg-4.onnx_
↪ \
--joiner=./icefall-asr-zipformer-wenetspeech-20230615/exp/joiner-epoch-12-avg-4.int8.
↪onnx \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000000.wav \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000001.wav \
./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T000000000002.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx-offline --tokens=./icefall-asr-zipformer-wenetspeech-20230615/  
data/lang_char/tokens.txt --encoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/  
encoder-epoch-12-avg-4.int8.onnx --decoder=./icefall-asr-zipformer-wenetspeech-  
20230615/exp/decoder-epoch-12-avg-4.onnx --joiner=./icefall-asr-zipformer-wenetspeech-  
20230615/exp/joiner-epoch-12-avg-4.int8.onnx ./icefall-asr-zipformer-wenetspeech-  
20230615/test_wavs/DEV_T0000000000.wav ./icefall-asr-zipformer-wenetspeech-20230615/  
test_wavs/DEV_T0000000001.wav ./icefall-asr-zipformer-wenetspeech-20230615/test_wavs/  
DEV_T0000000002.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./  
icefall-asr-zipformer-wenetspeech-20230615/exp/encoder-epoch-12-avg-4.int8.onnx",  
decoder_filename="./icefall-asr-zipformer-wenetspeech-20230615/exp/decoder-epoch-12-  
avg-4.onnx", joiner_filename="./icefall-asr-zipformer-wenetspeech-20230615/exp/joiner-  
epoch-12-avg-4.int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_  
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-zipformer-  
wenetspeech-20230615/data/lang_char/tokens.txt", num_threads=2, debug=False, provider=  
"cpu"), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search  
", max_active_paths=4, context_score=1.5)  
Creating recognizer ...  
Started  
Done!  
  
.icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000000.wav  
{"text": "", "timestamps": "[0.00, 0.12, 0.48, 0.60, 0.80, 1.08, 1.64, 1.76, 1.92, 2.08, 2.  
32, 2.48, 2.64, 3.08, 3.20, 3.28, 3.44, 3.60, 3.72, 3.84, 3.92, 4.12, 4.28, 4.48, 4.72,  
4.84]", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "",  
""]}  
----  
.icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000001.wav  
{"text": "", "timestamps": "[0.00, 0.16, 0.48, 0.68, 0.84, 1.08, 1.20, 1.48, 1.64, 2.08, 2.  
36, 2.52, 2.64, 2.84, 3.00, 3.16, 3.40, 3.52, 3.72, 3.84, 4.00, 4.16, 4.32, 4.56, 4.84]  
", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "",  
""]}  
----  
.icefall-asr-zipformer-wenetspeech-20230615/test_wavs/DEV_T0000000002.wav  
{"text": "", "timestamps": "[0.00, 0.12, 0.48, 0.84, 1.08, 1.44, 1.60, 1.84, 2.24, 2.48, 2.  
76, 2.88, 3.12, 3.24, 3.28, 3.36, 3.60, 3.72, 3.84, 4.16]", "tokens": ["", "", "", "", "",  
""]}  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 0.338 s  
Real time factor (RTF): 0.338 / 15.289 = 0.022
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./icefall-asr-zipformer-wenetspeech-20230615/data/lang_char/tokens.txt \
--encoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/encoder-epoch-12-avg-4.onnx \
--decoder=./icefall-asr-zipformer-wenetspeech-20230615/exp/decoder-epoch-12-avg-4.onnx \
--joiner=./icefall-asr-zipformer-wenetspeech-20230615/exp/joiner-epoch-12-avg-4.onnx
```

csukuangfj/sherpa-onnx-zipformer-large-en-2023-06-26 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-zipformer-large-2023-05-16>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

<https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/zipformer>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-large-en-2023-06-26.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-large-en-2023-06-26.tar.bz2

tar xvf sherpa-onnx-zipformer-large-en-2023-06-26.tar.bz2
rm sherpa-onnx-zipformer-large-en-2023-06-26.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
sherpa-onnx-zipformer-large-en-2023-06-26 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 1.2M Jun 26 13:19 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Jun 26 13:19 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 145M Jun 26 13:20 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 564M Jun 26 13:22 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 253K Jun 26 13:19 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Jun 26 13:19 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-large-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-large-en-2023-06-26/
↳ tokens.txt --encoder=./sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-
↳ 1.onnx --decoder=./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.
↳ onnx --joiner=./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.onnx ./
↳ sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-zipformer-
↳ large-en-2023-06-26/test_wavs/1.wav ./sherpa-onnx-zipformer-large-en-2023-06-26/test_
↳ wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-1.onnx", decoder_
↳ filename="./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.onnx",
↳ joiner_filename="./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.onnx
↳ "), paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-large-
↳ en-2023-06-26/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
↳ config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
↳ active_paths=4, context_score=1.5)
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
↳ cc:AcceptWaveformImpl:108 Creating a resampler:
    in_sample_rate: 8000
```

(continues on next page)

(continued from previous page)

```

output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↪SQUALID QUARTER OF THE BROTHELS","timestamps":[0.00, 0.48, 0.60, 0.72, 1.04, 1.28, 1.
↪36, 1.48, 1.60, 1.84, 1.96, 2.00, 2.16, 2.32, 2.40, 2.48, 2.60, 2.80, 3.04, 3.28, 3.40,
↪3.56, 3.76, 4.04, 4.24, 4.28, 4.48, 4.64, 4.80, 4.84, 5.00, 5.04, 5.28, 5.40, 5.56, 5.
↪60, 5.76, 5.96, 6.16],"tokens":[" AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE",
↪" YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", "THE",
↪" S", "QUA", "LI", "D", " ", "QUA", "R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}

-----
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
↪LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR_
↪EVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN",
↪"timestamps":[0.00, 0.20, 0.48, 0.72, 0.88, 1.04, 1.12, 1.20, 1.36, 1.52, 1.68, 1.84,
↪1.88, 2.00, 2.12, 2.32, 2.36, 2.60, 2.84, 3.12, 3.24, 3.48, 3.56, 3.76, 3.92, 4.12, 4.
↪36, 4.56, 4.72, 4.96, 5.16, 5.44, 5.68, 6.12, 6.28, 6.48, 6.88, 7.12, 7.36, 7.56, 7.92,
↪8.16, 8.28, 8.40, 8.48, 8.60, 8.76, 8.88, 9.08, 9.28, 9.44, 9.52, 9.60, 9.72, 9.92,
↪10.00, 10.12, 10.48, 10.68, 10.76, 11.00, 11.20, 11.36, 11.56, 11.76, 12.00, 12.12, 12.
↪28, 12.32, 12.52, 12.72, 12.84, 12.92, 13.04, 13.20, 13.44, 13.64, 13.76, 14.00, 14.12,
↪14.24, 14.36, 14.52, 14.72, 14.80, 15.04, 15.28, 15.52, 15.76, 16.00, 16.20, 16.24,
↪16.32],"tokens":[" GO", "D", " AS", " A", " DI", "RE", "C", "T", " CON", "SE", "QUE", "N", "CE", "O
↪F", " THE", " S", "IN", " WHICH", " MAN", " TH", "US", " P", "UN", "ISH", "ED", " HAD", " GIVE", "N
↪", " HER", " A", " LOVE", "LY", " CHILD", " WHO", "SE", " PLACE", " WAS", " ON", " THAT", " SAME",
↪" DIS", "HO", "N", "OR", "ED", " BO", "S", "OM", " TO", " CON", "NE", "C", "T", " HER", " P", "AR",
↪"ENT", " FOR", " E", "VER", " WITH", " THE", " RA", "CE", " AND", " DE", "S", "C", "ENT", " OF", " MO
↪", "R", "T", "AL", "S", " AND", " TO", " BE", " FI", "N", "AL", "LY", " A", " B", "LESS", "ED", " SO",
↪"UL", " IN", " HE", "A", "VE", "N"]}

-----
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/8k.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
↪"timestamps":[0.00, 0.12, 0.36, 0.48, 0.76, 0.96, 1.12, 1.24, 1.32, 1.44, 1.48, 1.68,
↪1.76, 1.88, 2.04, 2.12, 2.24, 2.28, 2.48, 2.56, 2.80, 3.08, 3.28, 3.52, 3.80, 3.92, 4.
↪00, 4.16, 4.24, 4.36, 4.44],"tokens":[" YE", "T", " THE", "SE", " THOUGHT", "S", " A", "FF",
↪"E", "C", "TED", " HE", "S", "TER", " P", "RY", "N", "NE", " ", "LESS", " WITH", " HO", "PE", " THAN",
↪" A", "PP", "RE", "HE", "N", "S", "ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.843 s
Real time factor (RTF): 1.843 / 28.165 = 0.065

```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-large-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-large-en-2023-06-26/
tokens.txt --encoder=./sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-
1.int8.onnx --decoder=./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-
1.onnx --joiner=./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.int8.
onnx ./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-
zipformer-large-en-2023-06-26/test_wavs/1.wav ./sherpa-onnx-zipformer-large-en-2023-06-
26/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx", decoder_
filename="./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.onnx",
joiner_filename="./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.
int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-large-
en-2023-06-26/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
active_paths=4, context_score=1.5)
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
cc:AcceptWaveformImpl:108 Creating a resampler:
in_sample_rate: 8000
output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
SQUALID QUARTER OF THE BROTHELS","timestamps":[0.00, 0.48, 0.60, 0.72, 1.04, 1.28, 1.
36, 1.48, 1.60, 1.84, 1.96, 2.00, 2.16, 2.32, 2.40, 2.48, 2.60, 2.80, 3.04, 3.28, 3.40]
3.56, 3.76, 4.04, 4.24, 4.28, 4.48, 4.64, 4.80, 4.84, 5.00, 5.04, 5.28, 5.40, 5.56, 5.
60, 5.76, 5.96, 6.16],"tokens":["AFTER","E","AR","LY","NIGHT","F","A","LL","THE",
"YE","LL","OW","LA","M","P","S","WOULD","LIGHT","UP","HE","RE","AND","THERE",
"THE","S","QUA","LI","D","","QUA","R","TER","OF","THE","B","RO","TH","EL","S"]}
```

(continued from previous page)

```
-----
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
↳ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR_
↳ EVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN",
↳ "timestamps":[0.00, 0.20, 0.48, 0.72, 0.88, 1.04, 1.12, 1.20, 1.36, 1.52, 1.64, 1.84,
↳ 1.88, 2.00, 2.12, 2.32, 2.36, 2.60, 2.84, 3.12, 3.24, 3.48, 3.56, 3.76, 3.92, 4.12, 4.
↳ 36, 4.52, 4.72, 4.96, 5.16, 5.44, 5.68, 6.12, 6.28, 6.48, 6.88, 7.12, 7.36, 7.56, 7.92,
↳ 8.16, 8.28, 8.40, 8.48, 8.60, 8.76, 8.88, 9.08, 9.28, 9.44, 9.52, 9.60, 9.72, 9.92,_
↳ 10.00, 10.12, 10.48, 10.68, 10.76, 11.00, 11.20, 11.36, 11.56, 11.76, 12.00, 12.12, 12.
↳ 28, 12.32, 12.52, 12.72, 12.84, 12.92, 13.04, 13.20, 13.44, 13.64, 13.76, 14.00, 14.08,
↳ 14.24, 14.36, 14.52, 14.72, 14.76, 15.04, 15.28, 15.52, 15.76, 16.00, 16.20, 16.24,
↳ 16.32]", "tokens": ["GO", "D", "AS", "A", "DI", "RE", "C", "T", "CON", "SE", "QUE", "N", "CE",
↳ "OF", "THE", "S", "IN", "WHICH", "MAN", "TH", "US", "P", "UN", "ISH", "ED", "HAD", "GIVE", "N
↳ ", "HER", "A", "LOVE", "LY", "CHILD", "WHO", "SE", "PLACE", "WAS", "ON", "THAT", "SAME",
↳ "DIS", "HO", "N", "OR", "ED", "BO", "S", "OM", "TO", "CON", "NE", "C", "T", "HER", "P", "AR",
↳ "ENT", "FOR", "E", "VER", "WITH", "THE", "RA", "CE", "AND", "DE", "S", "C", "ENT", "OF", "MO
↳ ", "R", "T", "AL", "S", "AND", "TO", "BE", "FI", "N", "AL", "LY", "A", "B", "LESS", "ED", "SO",
↳ "UL", "IN", "HE", "A", "VE", "N"]}

-----
./sherpa-onnx-zipformer-large-en-2023-06-26/test_wavs/8k.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
↳ "timestamps":[0.00, 0.12, 0.36, 0.48, 0.76, 0.96, 1.12, 1.24, 1.32, 1.44, 1.48, 1.68,
↳ 1.76, 1.88, 2.04, 2.12, 2.28, 2.32, 2.48, 2.52, 2.80, 3.08, 3.28, 3.52, 3.76, 3.92, 4.
↳ 00, 4.16, 4.24, 4.36, 4.44]", "tokens": ["YE", "T", "THE", "SE", "THOUGHT", "S", "A", "FF",
↳ "E", "C", "TED", "HE", "S", "TER", "P", "RY", "N", "NE", "LESS", "WITH", "HO", "PE", "THAN",
↳ "A", "PP", "RE", "HE", "N", "S", "ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.490 s
Real time factor (RTF): 1.490 / 28.165 = 0.053
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-large-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-large-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-large-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-large-en-2023-06-26/joiner-epoch-99-avg-1.onnx
```

csukuangfj/sherpa-onnx-zipformer-small-en-2023-06-26 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-zipformer-small-2023-05-16>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

<https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/zipformer>

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪zipformer-small-en-2023-06-26.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪zipformer-small-en-2023-06-26.tar.bz2

tar xvf sherpa-onnx-zipformer-small-en-2023-06-26.tar.bz2
rm sherpa-onnx-zipformer-small-en-2023-06-26.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-zipformer-small-en-2023-06-26 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff  1.2M Jun 26 13:04 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff  2.0M Jun 26 13:04 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff   25M Jun 26 13:04 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff   87M Jun 26 13:04 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff  253K Jun 26 13:04 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff   1.0M Jun 26 13:04 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-small-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-small-en-2023-06-26/
↳ tokens.txt --encoder=./sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-
↳ 1.onnx --decoder=./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.
↳ onnx --joiner=./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.onnx ./
↳ sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-zipformer-
↳ small-en-2023-06-26/test_wavs/1.wav ./sherpa-onnx-zipformer-small-en-2023-06-26/test_
↳ wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-1.onnx", decoder_
↳ filename="./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.onnx",
↳ joiner_filename="./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.onnx
↳ "), paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-small-
↳ en-2023-06-26/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
↳ config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
↳ active_paths=4, context_score=1.5)
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
↳ cc:AcceptWaveformImpl:108 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↳ SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.64, 0.76, 0.84, 1.12, 1.36, 1.
↳ 44, 1.56, 1.72, 1.84, 1.96, 2.04, 2.20, 2.32, 2.36, 2.44, 2.60, 2.76, 3.04, 3.24, 3.40,
↳ 3.52, 3.72, 4.04, 4.20, 4.28, 4.48, 4.64, 4.80, 4.84, 4.96, 5.00, 5.28, (604 bytes on page)
↳ 60, 5.76, 5.92, 6.08]", "tokens": [" "AFTER", " "E", " "AR", " "LY", " "NIGHT", " "F", " "A", " "LL", " "THE",
↳ " "YE", " "LL", " "OW", " "LA", " "M", " "P", " "S", " "WOULD", " "LIGHT", " "UP", " "HE", " "RE", " "AND", " "THERE",
↳ " "THE", " "S", " "QUA", " "LI", " "D", " ", " "QUA", " "R", " "TER", " "OF", " "THE", " "B", " "RO", " "TH", " "EL", " "S"]}
```

(continued from previous page)

```
-----
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
↳ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONoured BOSOM TO CONNECT HER PARENT_
↳ FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
↳ ","timestamps":[0.00, 0.32, 0.64, 0.80, 0.96, 1.08, 1.16, 1.20, 1.32, 1.52, 1.68, 1.
↳ 80, 1.88, 2.04, 2.16, 2.32, 2.40, 2.64, 2.88, 3.16, 3.20, 3.44, 3.52, 3.72, 3.88, 4.16,
↳ 4.44, 4.60, 4.76, 4.96, 5.16, 5.36, 5.60, 6.16, 6.32, 6.52, 6.88, 7.16, 7.32, 7.60, 7.
↳ 96, 8.16, 8.28, 8.36, 8.48, 8.64, 8.76, 8.84, 9.04, 9.28, 9.44, 9.52, 9.60, 9.68, 9.88,
↳ 9.92, 10.12, 10.52, 10.76, 10.80, 11.08, 11.20, 11.36, 11.56, 11.76, 11.96, 12.08, 12.
↳ 24, 12.28, 12.48, 12.68, 12.80, 12.92, 13.00, 13.20, 13.48, 13.72, 13.84, 14.04, 14.20,
↳ 14.28, 14.40, 14.56, 14.68, 14.76, 15.00, 15.24, 15.48, 15.68, 15.92, 16.08, 16.12,_
↳ 16.20],"tokens":[" GO","D"," AS"," A"," DI","RE","C","T"," CON","SE","QUE","N","CE","_
↳ OF"," THE"," S","IN"," WHICH"," MAN"," TH","US"," P","UN","ISH","ED"," HAD"," GIVE","N
↳ "," HER"," A"," LOVE","LY"," CHILD"," WHO","SE"," PLACE"," WAS"," ON"," THAT"," SAME",
↳ " DIS","HO","N","OUR","ED"," BO","S","OM"," TO"," CON","NE","C","T"," HER"," P","AR",
↳ "ENT"," FOR","E","VER"," WITH"," THE"," RA","CE"," AND"," DE","S","C","ENT"," OF"," MO
↳ ","R","T","AL","S"," AND"," TO"," BE"," FI","N","AL","LY"," A"," B","LESS","ED"," SO",
↳ "UL"," IN"," HE","A","VE","N"]}

-----
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/8k.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
"timestamps":[0.00, 0.32, 0.48, 0.64, 0.84, 1.08, 1.20, 1.32, 1.36, 1.44, 1.48, 1.64,
1.76, 1.88, 2.08, 2.12, 2.24, 2.28, 2.44, 2.48, 2.80, 3.04, 3.24, 3.48, 3.72, 3.88, 3.
92, 4.08, 4.16, 4.24, 4.36],"tokens":[" YE","T"," THE","SE"," THOUGHT","S"," A","FF",
"E","C","TED"," HE","S","TER"," P","RY","N","NE"," ","LESS"," WITH"," HO","PE"," THAN",
" A","PP","RE","HE","N","S","ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.953 s
Real time factor (RTF): 0.953 / 28.165 = 0.034
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-small-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx \
\
--decoder=./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
→build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-small-en-2023-06-26/  
→tokens.txt --encoder=./sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-  
→1.int8.onnx --decoder=./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-  
→1.onnx --joiner=./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.int8.  
→onnx ./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-  
→zipformer-small-en-2023-06-26/test_wavs/1.wav ./sherpa-onnx-zipformer-small-en-2023-06-  
→26/test_wavs/8k.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
→feature_dim=80), model_=  
→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./  
→sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx", decoder_=  
→filename="./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.onnx",  
→joiner_filename="./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.  
→int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_=  
→ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-small-  
→en-2023-06-26/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_=  
→config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_=  
→active_paths=4, context_score=1.5)  
Creating recognizer ...  
Started  
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.  
→cc:AcceptWaveformImpl:108 Creating a resampler:  
    in_sample_rate: 8000  
    output sample rate: 16000
```

Done!

```
./sherpa-onnéx-zipformer-small-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE  

→ SQUALID QUARTER OF THE BROTHELS","timestamps":[0.00, 0.64, 0.76, 0.84, 1.08, 1.36, 1.  

→ 44, 1.56, 1.72, 1.84, 1.96, 2.04, 2.20, 2.32, 2.36, 2.44, 2.60, 2.76, 3.04, 3.24, 3.40,  

→ 3.52, 3.72, 4.00, 4.20, 4.28, 4.48, 4.64, 4.80, 4.84, 4.96, 5.00, 5.28, 5.40, 5.52, 5.  

→ 60, 5.76, 5.92, 6.08]","tokens":[" AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE",  

→ " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", "  

→ THE", " S", "QUA", "LI", "D", " ", "QUA", "R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}  

----  

./sherpa-onnéx-zipformer-small-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A  

→ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONoured BOSOM TO CONNECT HER PARENT  

→ FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN  

→ ","timestamps":[0.00, 0.32, 0.64, 0.80, 0.96, 1.08, 1.16, 1.20, 1.32, 1.52, 1.68, 1.  

→ 80, 1.88, 2.04, 2.16, 2.32, 2.40, 2.64, 2.88, 3.16, 3.20, 3.44, 3.52, 3.72, 3.88, 4.16,  

→ 4.44, 4.60, 4.76, 4.96, 5.16, 5.36, 5.60, 6.16, 6.32, 6.52, 6.88, 7.16, 7.32, 7.60, 7.  

→ 96, 8.16, 8.28, 8.36, 8.48, 8.64, 8.76, 8.84, 9.04, 9.28, 9.44, 9.52, 9.60, 9.68, 9.88,  

→ 9.92, 10.12, 10.52, 10.76, 10.80, 11.08, 11.20, 11.36, 11.56, 11.76, 11.96, 12.08, 12.  

→ 24, 12.28, 12.48, 12.68, 12.80, 12.92, 13.04, 13.16, 13.48, 13.72, 13.84, 14.04, 14.20,  

→ 14.28, 14.40, 14.56, 14.68, 14.76, 15.00, 15.28, 15.48, 15.68, 15.92, 16.08, 16.24, 16.40, 16.56, 16.72, 16.88, 16.94, 17.10, 17.26, 17.42, 17.58, 17.74, 17.90, 18.06, 18.22, 18.38, 18.54, 18.70, 18.86, 18.92, 19.08, 19.24, 19.40, 19.56, 19.72, 19.88, 19.94, 20.10, 20.26, 20.42, 20.58, 20.74, 20.90, 20.96, 21.12, 21.28, 21.44, 21.60, 21.76, 21.92, 21.98, 22.14, 22.30, 22.46, 22.62, 22.78, 22.94, 22.96, 23.12, 23.28, 23.44, 23.60, 23.76, 23.92, 23.98, 24.14, 24.30, 24.46, 24.62, 24.78, 24.94, 24.96, 25.12, 25.28, 25.44, 25.60, 25.76, 25.92, 25.98, 26.14, 26.30, 26.46, 26.62, 26.78, 26.94, 26.96, 27.12, 27.28, 27.44, 27.60, 27.76, 27.92, 27.98, 28.14, 28.30, 28.46, 28.62, 28.78, 28.94, 28.96, 29.12, 29.28, 29.44, 29.60, 29.76, 29.92, 29.98, 30.14, 30.30, 30.46, 30.62, 30.78, 30.94, 30.96, 31.12, 31.28, 31.44, 31.60, 31.76, 31.92, 31.98, 32.14, 32.30, 32.46, 32.62, 32.78, 32.94, 32.96, 33.12, 33.28, 33.44, 33.60, 33.76, 33.92, 33.98, 34.14, 34.30, 34.46, 34.62, 34.78, 34.94, 34.96, 35.12, 35.28, 35.44, 35.60, 35.76, 35.92, 35.98, 36.14, 36.30, 36.46, 36.62, 36.78, 36.94, 36.96, 37.12, 37.28, 37.44, 37.60, 37.76, 37.92, 37.98, 38.14, 38.30, 38.46, 38.62, 38.78, 38.94, 38.96, 39.12, 39.28, 39.44, 39.60, 39.76, 39.92, 39.98, 40.14, 40.30, 40.46, 40.62, 40.78, 40.94, 40.96, 41.12, 41.28, 41.44, 41.60, 41.76, 41.92, 41.98, 42.14, 42.30, 42.46, 42.62, 42.78, 42.94, 42.96, 43.12, 43.28, 43.44, 43.60, 43.76, 43.92, 43.98, 44.14, 44.30, 44.46, 44.62, 44.78, 44.94, 44.96, 45.12, 45.28, 45.44, 45.60, 45.76, 45.92, 45.98, 46.14, 46.30, 46.46, 46.62, 46.78, 46.94, 46.96, 47.12, 47.28, 47.44, 47.60, 47.76, 47.92, 47.98, 48.14, 48.30, 48.46, 48.62, 48.78, 48.94, 48.96, 49.12, 49.28, 49.44, 49.60, 49.76, 49.92, 49.98, 50.14, 50.30, 50.46, 50.62, 50.78, 50.94, 50.96, 51.12, 51.28, 51.44, 51.60, 51.76, 51.92, 51.98, 52.14, 52.30, 52.46, 52.62, 52.78, 52.94, 52.96, 53.12, 53.28, 53.44, 53.60, 53.76, 53.92, 53.98, 54.14, 54.30, 54.46, 54.62, 54.78, 54.94, 54.96, 55.12, 55.28, 55.44, 55.60, 55.76, 55.92, 55.98, 56.14, 56.30, 56.46, 56.62, 56.78, 56.94, 56.96, 57.12, 57.28, 57.44, 57.60, 57.76, 57.92, 57.98, 58.14, 58.30, 58.46, 58.62, 58.78, 58.94, 58.96, 59.12, 59.28, 59.44, 59.60, 59.76, 59.92, 59.98, 60.14, 60.30, 60.46, 60.62, 60.78, 60.94, 60.96, 61.12, 61.28, 61.44, 61.60, 61.76, 61.92, 61.98, 62.14, 62.30, 62.46, 62.62, 62.78, 62.94, 62.96, 63.12, 63.28, 63.44, 63.60, 63.76, 63.92, 63.98, 64.14, 64.30, 64.46, 64.62, 64.78, 64.94, 64.96, 65.12, 65.28, 65.44, 65.60, 65.76, 65.92, 65.98, 66.14, 66.30, 66.46, 66.62, 66.78, 66.94, 66.96, 67.12, 67.28, 67.44, 67.60, 67.76, 67.92, 67.98, 68.14, 68.30, 68.46, 68.62, 68.78, 68.94, 68.96, 69.12, 69.28, 69.44, 69.60, 69.76, 69.92, 69.98, 70.14, 70.30, 70.46, 70.62, 70.78, 70.94, 70.96, 71.12, 71.28, 71.44, 71.60, 71.76, 71.92, 71.98, 72.14, 72.30, 72.46, 72.62, 72.78, 72.94, 72.96, 73.12, 73.28, 73.44, 73.60, 73.76, 73.92, 73.98, 74.14, 74.30, 74.46, 74.62, 74.78, 74.94, 74.96, 75.12, 75.28, 75.44, 75.60, 75.76, 75.92, 75.98, 76.14, 76.30, 76.46, 76.62, 76.78, 76.94, 76.96, 77.12, 77.28, 77.44, 77.60, 77.76, 77.92, 77.98, 78.14, 78.30, 78.46, 78.62, 78.78, 78.94, 78.96, 79.12, 79.28, 79.44, 79.60, 79.76, 79.92, 79.98, 80.14, 80.30, 80.46, 80.62, 80.78, 80.94, 80.96, 81.12, 81.28, 81.44, 81.60, 81.76, 81.92, 81.98, 82.14, 82.30, 82.46, 82.62, 82.78, 82.94, 82.96, 83.12, 83.28, 83.44, 83.60, 83.76, 83.92, 83.98, 84.14, 84.30, 84.46, 84.62, 84.78, 84.94, 84.96, 85.12, 85.28, 85.44, 85.60, 85.76, 85.92, 85.98, 86.14, 86.30, 86.46, 86.62, 86.78, 86.94, 86.96, 87.12, 87.28, 87.44, 87.60, 87.76, 87.92, 87.98, 88.14, 88.30, 88.46, 88.62, 88.78, 88.94, 88.96, 89.12, 89.28, 89.44, 89.60, 89.76, 89.92, 89.98, 90.14, 90.30, 90.46, 90.62, 90.78, 90.94, 90.96, 91.12, 91.28, 91.44, 91.60, 91.76, 91.92, 91.98, 92.14, 92.30, 92.46, 92.62, 92.78, 92.94, 92.96, 93.12, 93.28, 93.44, 93.60, 93.76, 93.92, 93.98, 94.14, 94.30, 94.46, 94.62, 94.78, 94.94, 94.96, 95.12, 95.28, 95.44, 95.60, 95.76, 95.92, 95.98, 96.14, 96.30, 96.46, 96.62, 96.78, 96.94, 96.96, 97.12, 97.28, 97.44, 97.60, 97.76, 97.92, 97.98, 98.14, 98.30, 98.46, 98.62, 98.78, 98.94, 98.96, 99.12, 99.28, 99.44, 99.60, 99.76, 99.92, 99.98, 100.14, 100.30, 100.46, 100.62, 100.78, 100.94, 100.96, 101.12, 101.28, 101.44, 101.60, 101.76, 101.92, 101.98, 102.14, 102.30, 102.46, 102.62, 102.78, 102.94, 102.96, 103.12, 103.28, 103.44, 103.60, 103.76, 103.92, 103.98, 104.14, 104.30, 104.46, 104.62, 104.78, 104.94, 104.96, 105.12, 105.28, 105.44, 105.60, 105.76, 105.92, 105.98, 106.14, 106.30, 106.46, 106.62, 106.78, 106.94, 106.96, 107.12, 107.28, 107.44, 107.60, 107.76, 107.92, 107.98, 108.14, 108.30, 108.46, 108.62, 108.78, 108.94, 108.96, 109.12, 109.28, 109.44, 109.60, 109.76, 109.92, 109.98, 110.14, 110.30, 110.46, 110.62, 110.78, 110.94, 110.96, 111.12, 111.28, 111.44, 111.60, 111.76, 111.92, 111.98, 112.14, 112.30, 112.46, 112.62, 112.78, 112.94, 112.96, 113.12, 113.28, 113.44, 113.60, 113.76, 113.92, 113.98, 114.14, 114.30, 114.46, 114.62, 114.78, 114.94, 114.96, 115.12, 115.28, 115.44, 115.60, 115.76, 115.92, 115.98, 116.14, 116.30, 116.46, 116.62, 116.78, 116.94, 116.96, 117.12, 117.28, 117.44, 117.60, 117.76, 117.92, 117.98, 118.14, 118.30, 118.46, 118.62, 118.78, 118.94, 118.96, 119.12, 119.28, 119.44, 119.60, 119.76, 119.92, 119.98, 120.14, 120.30, 120.46, 120.62, 120.78, 120.94, 120.96, 121.12, 121.28, 121.44, 121.60, 121.76, 121.92, 121.98, 122.14, 122.30, 122.46, 122.62, 122.78, 122.94, 122.96, 123.12, 123.28, 123.44, 123.60, 123.76, 123.92, 123.98, 124.14, 124.30, 124.46, 124.62, 124.78, 124.94, 124.96, 125.12, 125.28, 125.44, 125.60, 125.76, 125.92, 125.98, 126.14, 126.30, 126.46, 126.62, 126.78, 126.94, 126.96, 127.12, 127.28, 127.44, 127.60, 127.76, 127.92, 127.98, 128.14, 128.30, 128.46, 128.62, 128.78, 128.94, 128.96, 129.12, 129.28, 129.44, 129.60, 129.76, 129.92, 129.98, 130.14, 130.30, 130.46, 130.62, 130.78, 130.94, 130.96, 131.12, 131.28, 131.44, 131.60, 131.76, 131.92, 131.98, 132.14, 132.30, 132.46, 132.62, 132.78, 132.94, 132.96, 133.12, 133.28, 133.44, 133.60, 133.76, 133.92, 133.98, 134.14, 134.30, 134.46, 134.62, 134.78, 134.94, 134.96, 135.12, 135.28, 135.44, 135.60, 135.76, 135.92, 135.98, 136.14, 136.30, 136.46, 136.62, 136.78, 136.94, 136.96, 137.12, 137.28, 137.44, 137.60, 137.76, 137.92, 137.98, 138.14, 138.30, 138.46, 138.62, 138.78, 138.94, 138.96, 139.12, 139.28, 139.44, 139.60, 139.76, 139.92, 139.98, 140.14, 140.30, 140.46, 140.62, 140.78, 140.94, 140.96, 141.12, 141.28, 141.44, 141.60, 141.76, 141.92, 141.98, 142.14, 142.30, 142.46, 142.62, 142.78, 142.94, 142.96, 143.12, 143.28, 143.44, 143.60, 143.76, 143.92, 143.98, 144.14, 144.30, 144.46, 144.62, 144.78, 144.94, 144.96, 145.12, 145.28, 145.44, 145.60, 145.76, 145.92, 145.98, 146.14, 146.30, 146.46, 146.62, 146.78, 146.94, 146.96, 147.12, 147.28, 147.44, 147.60, 147.76, 147.92, 147.98, 148.14, 148.30, 148.46, 148.62, 148.78, 148.94, 148.96, 149.12, 149.28, 149.44, 149.60, 149.76, 149.92, 149.98, 150.14, 150.30, 150.46, 150.62, 150.78, 150.94, 150.96, 151.12, 151.28, 151.44, 151.60, 151.76, 151.92, 151.98, 152.14, 152.30, 152.46, 152.62, 152.78, 152.94, 152.96, 153.12, 153.28, 153.44, 153.60, 153.76, 153.92, 153.98, 154.14, 154.30, 154.46, 154.62, 154.78, 154.94, 154.96, 155.12, 155.28, 155.44, 155.60, 155.76, 155.92, 155.98, 156.14, 156.30, 156.46, 156.62, 156.78, 156.94, 156.96, 157.12, 157.28, 157.44, 157.60, 157.76, 157.92, 157.98, 158.14, 158.30, 158.46, 158.62, 158.78, 158.94, 158.96, 159.12, 159.28, 159.44, 159.60, 159.76, 159.92, 159.98, 160.14, 160.30, 160.46, 160.62, 160.78, 160.94, 160.96, 161.12, 161.28, 161.44, 161.60, 161.76, 161.92, 161.98, 162.14, 162.30, 162.46, 162.62, 162.78, 162.94, 162.96, 163.12, 163.28, 163.44, 163.60, 163.76, 163.92, 163.98, 164.14, 164.30, 164.46, 164.62, 164.78, 164.94, 164.96, 165.12, 165.28, 165.44, 165.60, 165.76, 165.92, 165.98, 166.14, 166.30, 166.46, 166.62, 166.78, 166.94, 166.96, 167.12, 167.28, 167.44, 167.60, 167.76, 167.92, 167.98, 168.14, 168.30, 168.46, 168.62, 168.78, 168.94, 168.96, 169.12, 169.28, 169.44, 169.60, 169.76, 169.92, 169.98, 170.14, 170.30, 170.46, 170.62, 170.78, 170.94, 170.96, 171.12, 171.28, 171.44, 171.60, 171.76, 171.92, 171.98, 172.14, 172.30, 172.46, 172.62, 172.78, 172.94, 172.96, 173.12, 173.28, 173.44, 173.60, 173.76, 173.92, 173.98, 174.14, 174.30, 174.46, 174.62, 174.78, 174.94, 174.96, 175.12, 175.28, 175.44, 175.60, 175.76, 175.92, 175.98, 176.14, 176.30, 176.46, 176.62, 176.78, 176.94, 176.96, 177.12, 177.28, 177.44, 177.60, 177.76, 177.92, 177.98, 178.14, 178.30, 178.46, 178.62, 178.78, 178.94, 178.96, 179.12, 179.28, 179.44, 179.60, 179.76, 179.92, 179.98, 180.14, 180.30, 180.46, 180.62, 180.78, 180.94, 180.96, 181.12, 181.28, 181.44, 181.60, 181.76, 181.92, 181.98, 182.14, 182.30, 182.46, 182.62, 182.78, 182.94, 182.96, 183.12, 183.28, 183.44, 183.60, 183.76, 183.92, 183.98, 184.14, 184.30, 184.46, 184.62, 184.78, 184.94, 184.96, 185.12, 185.28, 185.44, 185.60, 185.76, 185.92, 185.98, 186.14, 186.30, 186.46, 186.62, 186.78, 186.94, 186.96, 187.12, 187.28, 187.44, 187.60, 187.76, 187.92, 187.98, 188.14, 188.30, 188.46, 188.62, 188.78, 188.94, 188.96, 189.12, 189.28, 189.44, 189.60, 189.76, 189.92, 189.98, 190.14, 190.30, 190.46, 190.62, 190.78, 190.94, 190.96, 191.12, 191.28, 191.44, 191.60, 191.76, 191.92, 191.98, 192.14, 192.30, 192.46, 192.62, 192.78, 192.94, 192.96, 193.12, 193.28, 193.44, 193.60, 193.76, 193.92, 193.98, 194.14, 194.30, 194.46, 194.62, 194.78, 194.94, 194.96, 195.12, 195.28, 195.44, 195.60, 195.76, 195.92, 195.98, 196.14, 196.30, 196.46, 196.62, 196.78, 196.94, 196.96, 197.12, 197.28, 197.44, 197.60, 197.76, 197.92, 197.98, 198.14, 198.30, 198.46, 198.62, 198.78, 198.94, 198.96, 199.12, 199.28, 199.44, 199.60, 199.76, 199.92, 199.98, 200.14, 200.30, 200.46, 200.62, 200.78, 200.94, 200.96, 201.12, 201.28, 201.44, 201.60, 201.76, 201.92, 201.98, 202.14, 202.30, 202.46, 202.62, 202.78, 202.94, 202.96, 203.12, 203.28, 203.44, 203.60, 203.76, 203.92, 203.98, 204.14, 204.30, 204.46, 204.62, 204.78, 204.94, 204.96, 205.12, 205.28, 205.44, 205.60, 205.76, 205.92, 205.98, 206.14, 206.30, 206.46, 206.62, 206.78, 206.94, 206.96, 207.12, 207.28, 207.44, 207.60, 207.76, 207.92, 207.98, 208.14, 208.30, 208.46, 208.62, 208.78, 208.94, 208.96, 209.12, 209.28, 209.44, 209.60, 209.76, 209.92, 209.98, 210.14, 210.30, 210.46, 210.62, 210.78, 210.94, 210.96, 211.12, 211.28, 211.44, 211.60, 211.76, 211.92, 211.98, 212.14, 212.30, 212.46, 212.62, 212.78, 212.94, 212.96, 213.12, 213.28, 213.44, 213.60, 213.76, 213.92, 213.98, 214.14, 214.30, 214.46, 214.62, 214.78, 214.94, 214.96, 215.12, 215.28, 215.44, 215.60, 215.76, 215.92, 215.98, 216.14, 216.30, 216.46, 216.62, 216.78, 216.94, 216.96, 217.12, 217.28, 217.44, 217.60, 217.76, 217.92, 217.98, 218.14, 218.30, 218.46, 218.62, 218.78, 218.94, 218.96, 219.12, 219.28, 219.44, 219.60, 219.76, 219.92, 219.98, 220.14, 220.30, 220.46, 220.62, 220.78, 220.94, 220.96, 221.12, 221.28, 221.44, 221.60, 221.76, 221.92, 221.98, 222.14, 222.30, 222.46, 222.62, 222.78, 222.94, 222.96, 223.12, 223.28, 223.44, 223.60, 223.76, 223.92, 223.98, 224.14, 224.30, 224.46, 224.62, 224.78, 224.94, 224.96, 225.12, 225.28, 225.44, 225.60, 225.76, 225.92, 225.98, 226.14, 226.30, 226.46, 226.62, 
```

(continued from previous page)

```
----  
./sherpa-onnx-zipformer-small-en-2023-06-26/test_wavs/8k.wav  
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",  
 "timestamps":[0.00, 0.32, 0.48, 0.64, 0.84, 1.08, 1.20, 1.32, 1.36, 1.44, 1.48, 1.64,  
 1.76, 1.88, 2.08, 2.12, 2.24, 2.28, 2.44, 2.48, 2.80, 3.04, 3.24, 3.48, 3.72, 3.88, 3.  
 92, 4.08, 4.16, 4.24, 4.36],"tokens":["YE","T","THE","SE","THOUGHT","S","A","FF",  
 "E","C","TED","HE","S","TER","P","RY","N","NE","","LESS","WITH","HO","PE","THAN",  
 "A","PP","RE","HE","N","S","ION"]}  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 0.891 s  
Real time factor (RTF): 0.891 / 28.165 = 0.032
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-microphone-offline \  
--tokens=./sherpa-onnx-zipformer-small-en-2023-06-26/tokens.txt \  
--encoder=./sherpa-onnx-zipformer-small-en-2023-06-26/encoder-epoch-99-avg-1.onnx \  
--decoder=./sherpa-onnx-zipformer-small-en-2023-06-26/decoder-epoch-99-avg-1.onnx \  
--joiner=./sherpa-onnx-zipformer-small-en-2023-06-26/joiner-epoch-99-avg-1.onnx
```

csukuangfj/sherpa-onnx-zipformer-en-2023-06-26 (English)

This model is converted from

<https://huggingface.co/Zengwei/icefall-asr-librispeech-zipformer-2023-05-15>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

<https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/zipformer>

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
zipformer-en-2023-06-26.tar.bz2  
  
# For Chinese users, you can use the following mirror  
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
zipformer-en-2023-06-26.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf sherpa-onnx-zipformer-en-2023-06-26.tar.bz2
rm sherpa-onnx-zipformer-en-2023-06-26.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-zipformer-en-2023-06-26 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 1.2M Jun 26 12:45 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M Jun 26 12:45 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 66M Jun 26 12:45 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 248M Jun 26 12:46 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 253K Jun 26 12:45 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M Jun 26 12:45 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-en-2023-06-26/tokens.
↳ txt --encoder=./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx --
↳ decoder=./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx --joiner=./
↳ sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx ./sherpa-onnx-zipformer-
↳ en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav ./
↳ sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_                                            (continues on next page)
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx", decoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx", joiner_filename="./
8.25. Pre-trained models                                              443
↳ sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx"), paraformer=
↳ OfflineParaformerModelConfig(model=""), nemo=
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-en-2023-
```

(continued from previous page)

```

Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
→cc:AcceptWaveformImpl:108 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
→SQUALID QUARTER OF THE BROTHELS","timestamps":[0.00, 0.56, 0.64, 0.80, 1.08, 1.36, 1.
→40, 1.52, 1.68, 1.84, 1.96, 2.04, 2.20, 2.32, 2.40, 2.48, 2.60, 2.80, 3.04, 3.28, 3.40,
→3.56, 3.76, 4.08, 4.24, 4.32, 4.48, 4.64, 4.80, 4.84, 5.00, 5.04, 5.28, 5.40, 5.56, 5.
→60, 5.76, 5.96, 6.12],"tokens":[" AFTER"," E","AR","LY"," NIGHT","F","A","LL"," THE",
→" YE","LL","OW"," LA","M","P","S"," WOULD"," LIGHT"," UP"," HE","RE"," AND"," THERE",
→" THE"," S","QUA","LI","D"," ","QUA","R","TER"," OF"," THE"," B","RO","TH","EL","S"]}

---  

./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
→LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT_
→FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
→","timestamps":[0.00, 0.24, 0.56, 0.76, 0.92, 1.04, 1.16, 1.20, 1.36, 1.52, 1.64, 1.
→80, 1.88, 2.00, 2.16, 2.32, 2.40, 2.64, 2.88, 3.12, 3.24, 3.48, 3.56, 3.72, 3.92, 4.12,
→4.40, 4.52, 4.72, 4.96, 5.16, 5.36, 5.64, 6.12, 6.28, 6.52, 6.88, 7.12, 7.32, 7.56, 7.
→92, 8.16, 8.28, 8.40, 8.48, 8.64, 8.76, 8.88, 9.04, 9.28, 9.44, 9.52, 9.60, 9.72, 9.92,
→9.96, 10.16, 10.48, 10.72, 10.80, 11.04, 11.20, 11.36, 11.56, 11.76, 12.00, 12.12, 12.
→28, 12.32, 12.52, 12.72, 12.84, 12.92, 13.04, 13.20, 13.44, 13.68, 13.84, 14.00, 14.16,
→14.28, 14.40, 14.56, 14.72, 14.76, 15.00, 15.28, 15.48, 15.68, 15.96, 16.16, 16.20,
→16.28],"tokens":[" GO","D"," AS"," A"," DI","RE","C","T"," CON","SE","QUE","N","CE",
→"OF"," THE"," S","IN"," WHICH"," MAN"," TH","US"," P","UN","ISH","ED"," HAD"," GIVE","N
→"," HER"," A"," LOVE","LY"," CHILD"," WHO","SE"," PLACE"," WAS"," ON"," THAT"," SAME",
→" DIS","HO","N","OR","ED"," BO","S","OM"," TO"," CON","NE","C","T"," HER"," P","AR",
→"ENT"," FOR","E","VER"," WITH"," THE"," RA","CE"," AND"," DE","S","C","ENT"," OF"," MO
→","R","T","AL","S"," AND"," TO"," BE"," FI","N","AL","LY"," A"," B","LESS","ED"," SO",
→"UL"," IN"," HE","A","VE","N"]}

---  

./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
→"timestamps":[0.00, 0.24, 0.40, 0.60, 0.80, 1.04, 1.16, 1.28, 1.36, 1.44, 1.48, 1.68,
→1.76, 1.88, 2.00, 2.12, 2.24, 2.28, 2.48, 2.52, 2.80, 3.08, 3.28, 3.52, 3.68, 3.84, 3.
→96, 4.12, 4.20, 4.32, 4.44],"tokens":[" YE","T"," THE","SE"," THOUGHT","S"," A","FF",
→"E","C","TED"," HE","S","TER"," P","RY","N","NE"," ","LESS"," WITH"," HO","PE"," THAN",
→" A","PP","RE","HE","N","S","ION"]}

---  

num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.301 s
Real time factor (RTF): 1.301 / 28.165 = 0.046

```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-zipformer-en-2023-06-26/tokens.
↳ txt --encoder=./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx --
↳ decoder=./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx --joiner=/
↳ ./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.int8.onnx ./sherpa-onnx-
↳ zipformer-en-2023-06-26/test_wavs/0.wav ./sherpa-onnx-zipformer-en-2023-06-26/test_-
↳ wavs/1.wav ./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.int8.onnx", decoder_
↳ filename="./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx", joiner_
↳ filename="./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.int8.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-zipformer-en-2023-
↳ 06-26/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
↳ config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
↳ active_paths=4, context_score=1.5)
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
↳ cc:AcceptWaveformImpl:108 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/0.wav
{"text":" AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↳ SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.56, 0.64, 0.80, 1.08, 1.36, 1.
↳ 40, 1.52, 1.68, 1.84, 1.96, 2.04, 2.20, 2.32, 2.40, 2.48, 2.60, 2.76, 3.04, 3.28, 3.40,
↳ 3.56, 3.76, 4.08, 4.24, 4.32, 4.48, 4.64, 4.80, 4.84, 5.00, 5.04, 5.28, 5.40, 5.56, 5.
↳ 60, 5.76, 5.96, 6.12]", "tokens": [" AFTER", " E", "AR", "LY", " NIGHT", "F", "A"("çönlüüs ön TEK page"
↳ " YE", "LL", "OW", " LA", "M", "P", "S", " WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", "_
↳ THE", " S", "QUA", "LT", "D", " ", "QUA", "R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}
```

(continued from previous page)

```
-----
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/1.wav
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A_
↳ LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT_
↳ FOREVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
↳ ","timestamps":[0.00, 0.24, 0.56, 0.76, 0.92, 1.04, 1.16, 1.20, 1.36, 1.52, 1.64, 1.
↳ 80, 1.88, 2.00, 2.16, 2.32, 2.40, 2.64, 2.88, 3.12, 3.24, 3.48, 3.56, 3.72, 3.92, 4.12,
↳ 4.40, 4.52, 4.72, 4.96, 5.12, 5.40, 5.64, 6.12, 6.28, 6.52, 6.88, 7.12, 7.32, 7.60, 7.
↳ 92, 8.16, 8.28, 8.40, 8.48, 8.64, 8.76, 8.88, 9.04, 9.28, 9.44, 9.52, 9.60, 9.72, 9.92,
↳ 9.96, 10.16, 10.48, 10.72, 10.80, 11.04, 11.20, 11.36, 11.56, 11.76, 12.00, 12.12, 12.
↳ 28, 12.32, 12.52, 12.72, 12.84, 12.92, 13.04, 13.20, 13.44, 13.68, 13.84, 14.00, 14.16,
↳ 14.28, 14.40, 14.56, 14.72, 14.76, 15.00, 15.28, 15.48, 15.68, 15.96, 16.16, 16.20,
↳ 16.28]","tokens":[" GO","D"," AS"," A"," DI","RE","C","T"," CON","SE","QUE","N","CE",
↳ "OF"," THE"," S","IN"," WHICH"," MAN"," TH","US"," P","UN","ISH","ED"," HAD"," GIVE","N
↳ "," HER"," A"," LOVE","LY"," CHILD"," WHO","SE"," PLACE"," WAS"," ON"," THAT"," SAME",
↳ "DIS","HO","N","OR","ED"," BO","S","OM"," TO"," CON","NE","C","T"," HER"," P","AR",
↳ "ENT"," FOR","E","VER"," WITH"," THE"," RA","CE"," AND"," DE","S","C","ENT"," OF"," MO
↳ ","R","T","AL","S"," AND"," TO"," BE"," FI","N","AL","LY"," A"," B","LESS","ED"," SO",
↳ "UL"," IN"," HE","A","VE","N"]}

-----
./sherpa-onnx-zipformer-en-2023-06-26/test_wavs/8k.wav
{"text":" YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
↳ "timestamps":[0.00, 0.24, 0.40, 0.60, 0.80, 1.04, 1.16, 1.28, 1.36, 1.44, 1.48, 1.68,
↳ 1.76, 1.88, 2.00, 2.08, 2.24, 2.28, 2.48, 2.52, 2.80, 3.08, 3.28, 3.52, 3.68, 3.84, 3.
↳ 96, 4.12, 4.20, 4.32, 4.44]","tokens":[" YE","T"," THE","SE"," THOUGHT","S"," A","FF",
↳ "E","C","TED"," HE","S","TER"," P","RY","N","NE"," ", "LESS"," WITH"," HO","PE"," THAN",
↳ " A","PP","RE","HE","N","S","ION"]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.106 s
Real time factor (RTF): 1.106 / 28.165 = 0.039
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-06-26/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-06-26/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-06-26/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-06-26/joiner-epoch-99-avg-1.onnx
```

icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04 (English)

This model is trained using GigaSpeech + LibriSpeech + Common Voice 13.0 with zipformer

See <https://github.com/k2-fsa/icefall/pull/1010> if you are interested in how it is trained.

In the following, we describe how to download it and use it with `sherpa-onnix`.

Download the model

Please use the following commands to download it.

```
 wget https://github.com/k2-fsa/sherpa-onné/releases/download/asr-models/icefall-asr-  
 ↵multidataset-pruned_transducer_stateless7-2023-05-04.tar.bz2  
  
 # For Chinese users, you can use the following mirror  
 # wget https://hub.nuaa.cf/k2-fsa/sherpa-onné/releases/download/asr-models/icefall-asr-  
 ↵multidataset-pruned_transducer_stateless7-2023-05-04.tar.bz2  
  
 tar xvf icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04.tar.bz2  
 rm icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff  1.2M May 15 11:11 decoder-epoch-30-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 2.0M May 15 11:11 decoder-epoch-30-avg-4.onnx
-rw-r--r-- 1 fangjun staff 121M May 15 11:12 encoder-epoch-30-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 279M May 15 11:13 encoder-epoch-30-avg-4.onnx
-rw-r--r-- 1 fangjun staff 253K May 15 11:11 joiner-epoch-30-avg-4.int8.onnx
-rw-r--r-- 1 fangjun staff 1.0M May 15 11:11 joiner-epoch-30-avg-4.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/data/lang_
↪bpe_500/tokens.txt \
--encoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪encoder-epoch-30-avg-4.onnx \
--decoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪decoder-epoch-30-avg-4.onnx \
```

(continued from previous page)

```
--joiner=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/joiner-
↪epoch-30-avg-4.onnx \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1089-
↪134686-0001.wav \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-
↪135766-0001.wav \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-
↪135766-0002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↪build/bin/sherpa-onnx-offline --tokens=./icefall-asr-multidataset-pruned_transducer_
↪stateless7-2023-05-04/data/lang_bpe_500/tokens.txt --encoder=./icefall-asr-
↪multidataset-pruned_transducer_stateless7-2023-05-04/exp/encoder-epoch-30-avg-4.onnx --
↪decoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/decoder-
↪epoch-30-avg-4.onnx --joiner=./icefall-asr-multidataset-pruned_transducer_stateless7-
↪2023-05-04/exp/joiner-epoch-30-avg-4.onnx ./icefall-asr-multidataset-pruned_transducer_
↪stateless7-2023-05-04/test_wavs/1089-134686-0001.wav ./icefall-asr-multidataset-pruned_
↪transducer_stateless7-2023-05-04/test_wavs/1221-135766-0001.wav ./icefall-asr-
↪multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-0002.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↪feature_dim=80), model_
↪config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↪icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/encoder-epoch-30-
↪avg-4.onnx", decoder_filename="./icefall-asr-multidataset-pruned_transducer_stateless7-
↪2023-05-04/exp/decoder-epoch-30-avg-4.onnx", joiner_filename="./icefall-asr-
↪multidataset-pruned_transducer_stateless7-2023-05-04/exp/joiner-epoch-30-avg-4.onnx"), 
↪paraformer=OfflineParaformerModelConfig(model=""), nemo_
↪ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-multidataset-
↪pruned_transducer_stateless7-2023-05-04/data/lang_bpe_500/tokens.txt", num_threads=2,
↪debug=False, provider="cpu"), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_
↪method="greedy_search", max_active_paths=4)
Creating recognizer ...
Started
Done!

./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1089-134686-
↪0001.wav
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE_
↪SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.40, 0.56, 0.64, 0.96, 1.24, 1.32, 1.
↪44, 1.56, 1.76, 1.88, 1.96, 2.16, 2.32, 2.36, 2.48, 2.60, 2.80, 3.08, 3.28, 3.36, 3.56, 3.80, 4.04, 4.
↪24, 4.32, 4.48, 4.64, 4.84, 4.88, 5.00, 5.08, 5.32, 5.44, 5.56, 5.64, 5.80, 5.96, 6.20]", "tokens": [
↪"AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE", " YE", "LL", "OW", " LA", "M", "P", "S", " ",
↪"WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA",
↪"R", "TER", " OF", " THE", " B", "RO", "TH", "EL", "S"]}

...
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-
↪0001.wav
```

(continues on next page)

(continued from previous page)

```
{"text":" GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A  
→LOVELY CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR  
→EVER WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN",  
→"timestamps": "[0.00,0.16,0.44,0.68,0.84,1.00,1.12,1.16,1.32,1.48,1.64,1.80,1.84,2.00,2.  
→12,2.28,2.40,2.64,2.88,3.16,3.28,3.56,3.60,3.76,3.92,4.12,4.36,4.52,4.72,4.92,5.16,5.  
→44,5.72,6.12,6.24,6.48,6.84,7.08,7.28,7.56,7.88,8.12,8.28,8.36,8.48,8.60,8.76,8.88,9.  
→12,9.28,9.48,9.56,9.64,9.80,10.00,10.04,10.20,10.44,10.68,10.80,11.04,11.20,11.40,11.  
→56,11.80,12.00,12.12,12.28,12.32,12.52,12.72,12.84,12.96,13.04,13.24,13.40,13.64,13.80,  
→14.00,14.16,14.24,14.36,14.56,14.72,14.80,15.08,15.32,15.52,15.76,16.04,16.16,16.24,16.  
→36]", "tokens": [" GO", "D", " AS", " A", " DI", "RE", "C", "T", " CON", "SE", "QUE", "N", "CE", " OF  
→", " THE", " S", "IN", " WHICH", " MAN", " TH", "US", " P", "UN", "ISH", "ED", " HAD", " GIVE", "N",  
→" HER", " A", " LOVE", "LY", " CHILD", " WHO", "SE", " PLACE", " WAS", " ON", " THAT", " SAME", " u  
→DIS", "HO", "N", "OR", "ED", " BO", "S", "OM", " TO", " CON", "NE", "C", "T", " HER", " P", "AR", "ENT  
→", " FOR", " E", "VER", " WITH", " THE", " RA", "CE", " AND", " DE", "S", "C", "ENT", " OF", " MO", "R  
→", "T", "AL", "S", " AND", " TO", " BE", " FI", "N", "AL", "LY", " A", " B", "LESS", "ED", " SO", "UL",  
→" IN", " HE", "A", "VE", "N"]}  
----  
.icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-  
→0002.wav  
{ "text": " YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",  
→"timestamps": "[0.00,0.08,0.32,0.48,0.68,0.92,1.08,1.20,1.28,1.40,1.44,1.64,1.76,1.88,2.  
→04,2.12,2.24,2.32,2.48,2.56,2.88,3.12,3.32,3.52,3.76,3.92,4.00,4.20,4.28,4.40,4.52]",  
→"tokens": [" YE", "T", " THE", "SE", " THOUGHT", "S", " A", "FF", "E", "C", "TED", " HE", "S", "TER",  
→" P", "RY", "N", "NE", " ", "LESS", " WITH", " HO", "PE", " THAN", " A", "PP", "RE", "HE", "N", "S",  
→"ION"]}  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 1.662 s  
Real time factor (RTF): 1.662 / 28.165 = 0.059
```

int8

The following code shows how to use `int8` models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/data/lang-
↪_bpe_500/tokens.txt \
--encoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪encoder-epoch-30-avg-4.int8.onnx \
--decoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪decoder-epoch-30-avg-4.onnx \
--joiner=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/joiner-
↪epoch-30-avg-4.int8.onnx \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1089-
↪134686-0001.wav \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-
↪135766-0001.wav \
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-
↪135766-0002.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-ONNX/sherpa-ONNX/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-ONNX-offline --tokens=./icefall-asr-multidataset-pruned_transducer_  
stateless7-2023-05-04/data/lang_bpe_500/tokens.txt --encoder=./icefall-asr-  
multidataset-pruned_transducer_stateless7-2023-05-04/exp/encoder-epoch-30-avg-4.int8.  
onnx --decoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/  
decoder-epoch-30-avg-4.onnx --joiner=./icefall-asr-multidataset-pruned_transducer_  
stateless7-2023-05-04/exp/joiner-epoch-30-avg-4.int8.onnx ./icefall-asr-multidataset-  
pruned_transducer_stateless7-2023-05-04/test_wavs/1089-134686-0001.wav ./icefall-asr-  
multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-0001.wav ./  
icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-  
0002.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_=  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./  
icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/encoder-epoch-30-  
avg-4.int8.onnx", decoder_filename="./icefall-asr-multidataset-pruned_transducer_  
stateless7-2023-05-04/exp/decoder-epoch-30-avg-4.onnx", joiner_filename="./icefall-asr-  
multidataset-pruned_transducer_stateless7-2023-05-04/exp/joiner-epoch-30-avg-4.int8.  
onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_=  
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./icefall-asr-multidataset-  
pruned_transducer_stateless7-2023-05-04/data/lang_bpe_500/tokens.txt", num_threads=2,  
debug=False, provider="cpu"), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_=  
method="greedy_search", max_active_paths=4)  
Creating recognizer ...  
Started  
Done!  
  
.icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1089-134686-  
0001.wav  
{"text": " AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE  
SQUALID QUARTER OF THE BROTHELS", "timestamps": "[0.00, 0.40, 0.56, 0.64, 0.96, 1.24, 1.32, 1.  
44, 1.56, 1.76, 1.88, 1.96, 2.16, 2.32, 2.36, 2.48, 2.60, 2.80, 3.08, 3.28, 3.36, 3.56, 3.80, 4.04, 4.  
24, 4.32, 4.48, 4.64, 4.84, 4.88, 5.00, 5.08, 5.32, 5.44, 5.56, 5.64, 5.80, 5.96, 6.20]", "tokens": ["  
AFTER", " E", "AR", "LY", " NIGHT", "F", "A", "LL", " THE", " YE", "LL", "OW", " LA", "M", "P", "S", "  
WOULD", " LIGHT", " UP", " HE", "RE", " AND", " THERE", " THE", " S", "QUA", "LI", "D", " ", "QUA",  

```

(continued from previous page)

```
-----
./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/test_wavs/1221-135766-
↪0002.wav
>{"text": " YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION",
↪"timestamps": "[0.00, 0.08, 0.32, 0.48, 0.68, 0.92, 1.08, 1.20, 1.28, 1.40, 1.44, 1.64, 1.76, 1.88, 2.
↪04, 2.12, 2.28, 2.32, 2.52, 2.56, 2.88, 3.12, 3.32, 3.52, 3.76, 3.92, 4.00, 4.20, 4.28, 4.40, 4.52]", 
↪"tokens": ["YE", "T", "THE", "SE", "THOUGHT", "S", "A", "FF", "E", "C", "TED", "HE", "S", "TER",
↪"P", "RY", "N", "NE", "LESS", "WITH", "HO", "PE", "THAN", "A", "PP", "RE", "HE", "N", "S",
↪"ION"]}
```

num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.424 s
Real time factor (RTF): 1.424 / 28.165 = 0.051

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/data/lang_
↪bpe_500/tokens.txt \
--encoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪encoder-epoch-30-avg-4.onnx \
--decoder=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/
↪decoder-epoch-30-avg-4.onnx \
--joiner=./icefall-asr-multidataset-pruned_transducer_stateless7-2023-05-04/exp/joiner-
↪epoch-30-avg-4.onnx
```

csukuangfj/sherpa-onnx-zipformer-en-2023-04-01 (English)

This model is converted from

<https://huggingface.co/WeijiZhuang/icefall-asr-librispeech-pruned-transducer-stateless8-2022-12-02>

which supports only English as it is trained on the LibriSpeech and GigaSpeech corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless8

In the following, we describe how to download it and use it with sherpa-onnx.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-en-2023-04-01.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
# zipformer-en-2023-04-01.tar.bz2

tar xvf sherpa-onnx-zipformer-en-2023-04-01.tar.bz2
rm sherpa-onnx-zipformer-en-2023-04-01.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-zipformer-en-2023-04-01$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root  1.3M Apr  1 14:34 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root  2.0M Apr  1 14:34 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 180M Apr  1 14:34 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 338M Apr  1 14:34 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 254K Apr  1 14:34 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 1003K Apr  1 14:34 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-04-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
    ↪feature_dim=80), model_
    ↪config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename='./
    ↪sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.onnx', decoder_filename='./
    ↪sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx', joiner_filename='./
    ↪sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.onnx'), ↪
    ↪paraformer=OfflineParaformerModelConfig(model="", tokens="./sherpa-onnx-zipformer-en-
    ↪2023-04-01/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
2023-04-01 14:40:56.353883875 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
    ↪setaffinity_np failed for thread: 638155, index: 16, mask: {17, 53, }, error code: 22
    ↪error msg: Invalid argument. Specify the number of threads explicitly so the affinity
    ↪is not set.
2023-04-01 14:40:56.353881478 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
    ↪setaffinity_np failed for thread: 638154, index: 15, mask: {16, 52, }, error code: 22
    ↪error msg: Invalid argument. Specify the number of threads explicitly so the affinity
    ↪is not set.
Started
Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID
    ↪QUARTER OF THE BROTHELS
---
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY
    ↪CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER
    ↪WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
---
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 2.151 s
Real time factor (RTF): 2.151 / 28.165 = 0.076
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
    --tokens=./sherpa-onnx-zipformer-en-2023-04-01/tokens.txt \
    --encoder=./sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.int8.onnx \
```

(continues on next page)

(continued from previous page)

```
--decoder=./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
    ↵ feature_dim=80), model_
    ↵ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
    ↵ sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.int8.onnx", decoder_
    ↵ filename="./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx", joiner_
    ↵ filename="./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.int8.onnx"), ↵
    ↵ paraformer=OfflineParaformerModelConfig(model=""), tokens="./sherpa-onnx-zipformer-en-
    ↵ 2023-04-01/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
2023-04-01 14:42:00.407939001 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
    ↵ setaffinity_np failed for thread: 638195, index: 15, mask: {16, 52, }, error code: 22
    ↵ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
    ↵ is not set.
2023-04-01 14:42:00.407940827 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
    ↵ setaffinity_np failed for thread: 638196, index: 16, mask: {17, 53, }, error code: 22
    ↵ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
    ↵ is not set.
Started
Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID
    ↵ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY
    ↵ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER
    ↵ WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-zipformer-en-2023-04-01/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.478 s
Real time factor (RTF): 1.478 / 28.165 = 0.052
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-04-01/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-04-01/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-04-01/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-04-01/joiner-epoch-99-avg-1.onnx
```

csukuangfj/sherpa-onnx-zipformer-en-2023-03-30 (English)

This model is converted from

<https://huggingface.co/csukuangfj/icefall-asr-librispeech-pruned-transducer-stateless7-2022-11-11>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless7

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
zipformer-en-2023-03-30.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
# zipformer-en-2023-03-30.tar.bz2

tar xvf sherpa-onnx-zipformer-en-2023-03-30.tar.bz2
rm sherpa-onnx-zipformer-en-2023-03-30.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-zipformer-en-2023-03-30$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root  1.3M Mar 31 00:37 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root  2.0M Mar 30 20:10 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 180M Mar 31 00:37 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 338M Mar 30 20:10 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root  254K Mar 31 00:37 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 1003K Mar 30 20:10 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-03-30/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.onnx", decoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx", joiner_filename="./
↳ sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), tokens="./sherpa-onnx-zipformer-en-
↳ 2023-03-30/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
2023-04-01 06:47:56.620698024 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 607690, index: 15, mask: {16, 52, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
↳ is not set.
2023-04-01 06:47:56.620700026 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 607691, index: 16, mask: {17, 53, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
↳ is not set.
Started
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
↳ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY_
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER_
↳ WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTION HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.950 s
Real time factor (RTF): 1.950 / 28.165 = 0.069
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-03-30/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/0.wav \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/1.wav \
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.int8.onnx", decoder_
↳ filename="./sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx", joiner_
↳ filename="./sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.int8.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), tokens="./sherpa-onnx-zipformer-en-
↳ 2023-03-30/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
2023-04-01 06:49:34.370117205 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 607732, index: 16, mask: {17, 53, }, error code: 22_
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
↳ is not set.
2023-04-01 06:49:34.370115197 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 607731, index: 15, mask: {16, 52, }, error code: 22_
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity_
↳ is not set.
```

(continued from previous page)

```
Started
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
→ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY_
→ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOR EVER_
→ WITH THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-zipformer-en-2023-03-30/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.710 s
Real time factor (RTF): 1.710 / 28.165 = 0.061
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-zipformer-en-2023-03-30/tokens.txt \
--encoder=./sherpa-onnx-zipformer-en-2023-03-30/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-zipformer-en-2023-03-30/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-zipformer-en-2023-03-30/joiner-epoch-99-avg-1.onnx
```

Conformer-transducer-based Models

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

csukuangfj/sherpa-onnx-conformer-zh-stateless2-2023-05-23 (Chinese)

This model is converted from

https://huggingface.co/luomingshuang/icefall_asr_wenetspeech_pruned_transducer_stateless2

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/wenetspeech/ASR/pruned_transducer_stateless2

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪conformer-zh-stateless2-2023-05-23.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪conformer-zh-stateless2-2023-05-23.tar.bz2

tar xvf sherpa-onnx-conformer-zh-stateless2-2023-05-23.tar.bz2
rm sherpa-onnx-conformer-zh-stateless2-2023-05-23.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-conformer-zh-stateless2-2023-05-23 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 11M May 23 15:29 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 12M May 23 15:29 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 122M May 23 15:30 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 315M May 23 15:31 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 fangjun staff 2.7M May 23 15:29 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 fangjun staff 11M May 23 15:29 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.wav \
./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav \
./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/2.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-conformer-zh-stateless2-2023-05-
23/tokens.txt --encoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-
99-avg-1.onnx --decoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-
99-avg-1.onnx --joiner=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-epoch-
99-avg-1.onnx ./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.wav ./
sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav ./sherpa-onnx-conformer-
zh-stateless2-2023-05-23/test_wavs/2.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-99-avg-1.onnx", decoder_
filename="./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-99-avg-1.onnx
", joiner_filename="./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-epoch-99-
avg-1.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-conformer-zh-
stateless2-2023-05-23/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
active_paths=4)
Creating recognizer ...
Started
Done!

./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.wav
{"text": "", "timestamps": "[0.00, 0.12, 0.44, 0.64, 0.84, 1.04, 1.64, 1.72, 1.88, 2.08, 2.
28, 2.44, 2.56, 2.76, 3.08, 3.20, 3.32, 3.48, 3.64, 3.76, 3.88, 4.00, 4.16, 4.24, 4.44,
4.60, 4.84]", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], }
-----
./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav
```

(continues on next page)

(continued from previous page)

```
{"text":"","timestamps":[0.00, 0.12, 0.48, 0.64, 0.88, 1.08, 1.28, 1.48, 1.80, 2.12, 2.  
↳ 40, 2.56, 2.68, 2.88, 3.04, 3.16, 3.36, 3.56, 3.68, 3.84, 4.00, 4.16, 4.32, 4.56, 4.76]  
↳ ","tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","","",""]}  
----  
.sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/2.wav  
{"text":"","timestamps":[0.00, 0.16, 0.60, 0.88, 1.08, 1.36, 1.64, 1.84, 2.24, 2.52, 2.  
↳ 72, 2.92, 3.08, 3.24, 3.40, 3.56, 3.72, 3.88, 4.12],"tokens":["","","","","","","","","","","","","","","","","","","","","","",""]}  
----  
num threads: 2  
decoding method: greedy_search  

```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-offline \  
--tokens=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/tokens.txt \  
--encoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-99-avg-1.int8.  
↳ onnx \  
--decoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-99-avg-1.onnx.  
↳ \  
--joiner=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-epoch-99-avg-1.int8.  
↳ onnx \  
.sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.wav \  
.sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav \  
.sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/2.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: We did not use int8 for the decoder model above.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-conformer-zh-stateless2-2023-05-  
23/tokens.txt --encoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-  
99-avg-1.int8.onnx --decoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-  
epoch-99-avg-1.onnx --joiner=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-  
epoch-99-avg-1.int8.onnx ./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.  
wav ./sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav ./sherpa-onnx-  
conformer-zh-stateless2-2023-05-23/test_wavs/2.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./  
sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-99-avg-1.int8.onnx",  
decoder_filename="./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-99-  
avg-1.onnx", joiner_filename="./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-  
epoch-99-avg-1.int8.onnx"), paraformer=OfflineParaformerModelConfig(model=""), nemo_  
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-conformer-zh-  
stateless2-2023-05-23/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_  
config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_  
active_paths=4)  
Creating recognizer ...  
Started  
Done!  
  
../sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/0.wav  
{"text": "", "timestamps": "[0.00, 0.12, 0.44, 0.64, 0.84, 1.08, 1.64, 1.72, 1.88, 2.08, 2.  
28, 2.44, 2.56, 2.76, 3.08, 3.20, 3.32, 3.48, 3.64, 3.76, 3.88, 4.00, 4.16, 4.24, 4.48,  
4.60, 4.84]", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "",  
"", "", "", "", "", ""]}  
---  
../sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/1.wav  
{"text": "", "timestamps": "[0.00, 0.08, 0.48, 0.64, 0.88, 1.08, 1.28, 1.48, 1.80, 2.08, 2.  
40, 2.56, 2.68, 2.88, 3.04, 3.16, 3.36, 3.56, 3.68, 3.84, 4.00, 4.16, 4.32, 4.56, 4.76]  
", "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "",  
""]}  
---  
../sherpa-onnx-conformer-zh-stateless2-2023-05-23/test_wavs/2.wav  
{"text": "", "timestamps": "[0.00, 0.12, 0.56, 0.84, 1.08, 1.40, 1.64, 1.84, 2.24, 2.52, 2.  
72, 2.92, 3.08, 3.24, 3.40, 3.56, 3.72, 3.88, 4.12]", "tokens": ["", "", "", "", "", "",  
"", "", "", "", "", "", "", "", ""]}  
---  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 0.439 s  
Real time factor (RTF): 0.439 / 15.289 = 0.029
```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-zh-stateless2-2023-05-23/joiner-epoch-99-avg-1.onnx
```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

csukuangfj/sherpa-onnx-conformer-zh-2023-05-23 (Chinese)

This model is converted from

https://huggingface.co/luomingshuang/icefall_asr_wenetspeech_pruned_transducer_stateless5_offline

which supports only Chinese as it is trained on the [WenetSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/wenetspeech/ASR/pruned_transducer_stateless5

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
--conformer-zh-2023-05-23.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
--conformer-zh-2023-05-23.tar.bz2

tar xvf sherpa-onnx-conformer-zh-2023-05-23.tar.bz2
rm sherpa-onnx-conformer-zh-2023-05-23.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
sherpa-onnx-conformer-zh-2023-05-23 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 11M May 23 13:45 decoder-epoch-99-avg-1.int8.onnx
```

(continues on next page)

(continued from previous page)

-rw-r--r--	1	fangjun	staff	12M	May 23 13:45	decoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	129M	May 23 13:47	encoder-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	345M	May 23 13:48	encoder-epoch-99-avg-1.onnx
-rw-r--r--	1	fangjun	staff	2.7M	May 23 13:45	joiner-epoch-99-avg-1.int8.onnx
-rw-r--r--	1	fangjun	staff	11M	May 23 13:45	joiner-epoch-99-avg-1.onnx

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-conformer-zh-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/0.wav \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/2.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-conformer-zh-2023-05-23/tokens.
↳ txt --encoder=./sherpa-onnx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx --
↳ decoder=./sherpa-onnx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx --joiner=./
↳ sherpa-onnx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx ./sherpa-onnx-conformer-
↳ zh-2023-05-23/test_wavs/0.wav ./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav ./
↳ sherpa-onnx-conformer-zh-2023-05-23/test_wavs/2.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
↳ sherpa-onnx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnx", decoder_filename="./
↳ sherpa-onnx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx", joiner_filename="./
↳ sherpa-onnx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnx"),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-conformer-zh-2023-
↳ 05-23/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_
↳ config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_
↳ active_paths=4)
```

(continued from previous page)

```

Creating recognizer ...
Started
Done!

./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/0.wav
{"text":"","timestamps":[0.00, 0.12, 0.52, 0.64, 0.84, 1.04, 1.68, 1.80, 1.92, 2.12, 2.
↪32, 2.48, 2.64, 2.76, 3.08, 3.20, 3.44, 3.52, 3.64, 3.76, 3.88, 4.00, 4.16, 4.32, 4.48,
↪4.64, 4.84],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","",
↪","","","","","","","",""]}

-----
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav
{"text":"","timestamps":[0.04, 0.16, 0.36, 0.48, 0.68, 0.92, 1.08, 1.24, 1.44, 1.84, 2.
↪08, 2.36, 2.52, 2.68, 2.88, 3.04, 3.16, 3.40, 3.56, 3.72, 3.84, 4.04, 4.16, 4.32, 4.56,
↪4.76],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","",
↪","","","",""]}

-----
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/2.wav
{"text":"","timestamps":[0.00, 0.12, 0.60, 0.84, 1.04, 1.44, 1.68, 1.84, 2.28, 2.52, 2.
↪80, 2.92, 3.08, 3.24, 3.40, 3.60, 3.72, 3.84, 4.12],"tokens":["","","","","","","","","",
↪","","","","","","","","","","","","","",""]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.706 s
Real time factor (RTF): 0.706 / 15.289 = 0.046

```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

int8

The following code shows how to use int8 models to decode wave files:

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-conformer-zh-2023-05-23/tokens.txt \
--encoder=./sherpa-onnx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/0.wav \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav \
./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/2.wav

```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: We did not use int8 for the decoder model above.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./  
build/bin/sherpa-onnx-offline --decoding-method=greedy_search --tokens=./sherpa-onnx-  
conformer-zh-2023-05-23/tokens.txt --encoder=./sherpa-onnx-conformer-zh-2023-05-23/  
encoder-epoch-99-avg-1.int8.onnx --decoder=./sherpa-onnx-conformer-zh-2023-05-23/  
decoder-epoch-99-avg-1.onnx --joiner=./sherpa-onnx-conformer-zh-2023-05-23/joiner-  
epoch-99-avg-1.int8.onnx ./sherpa-onnx-conformer-zh-2023-05-23/test_wavs/0.wav ./  
sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav ./sherpa-onnx-conformer-zh-2023-05-  
23/test_wavs/2.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./  
sherpa-onnx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.int8.onnx", decoder_  
filename="./sherpa-onnx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnx", joiner_  
filename="./sherpa-onnx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.int8.onnx"),  
paraformer=OfflineParaformerModelConfig(model=""), nemo_  
ctc=OfflineNemoEncDecCtcModelConfig(model=""), tokens="./sherpa-onnx-conformer-zh-2023-  
05-23/tokens.txt", num_threads=2, debug=False, provider="cpu"), lm_  
config=OfflineLMConfig(model="", scale=0.5), decoding_method="greedy_search", max_  
active_paths=4)  
Creating recognizer ...  
Started  
Done!  
  
../sherpa-onnx-conformer-zh-2023-05-23/test_wavs/0.wav  
{"text":"","timestamps":[0.00, 0.12, 0.52, 0.64, 0.84, 1.04, 1.68, 1.80, 1.92, 2.08, 2.  
32, 2.48, 2.64, 2.76, 3.08, 3.20, 3.44, 3.52, 3.64, 3.76, 3.88, 4.00, 4.16, 4.32, 4.48,  
4.60, 4.84],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","  
","","","","","","","",""]}  
---  
../sherpa-onnx-conformer-zh-2023-05-23/test_wavs/1.wav  
{"text":"","timestamps":[0.04, 0.16, 0.36, 0.48, 0.68, 0.92, 1.08, 1.24, 1.44, 1.88, 2.  
08, 2.36, 2.52, 2.64, 2.88, 3.00, 3.16, 3.40, 3.56, 3.72, 3.84, 4.04, 4.20, 4.32, 4.56,  
4.76],"tokens":["","","","","","","","","","","","","","","","","","","","","","","","","","","  
","","","","","",""]}  
---  
../sherpa-onnx-conformer-zh-2023-05-23/test_wavs/2.wav  
{"text":"","timestamps":[0.00, 0.12, 0.60, 0.84, 1.04, 1.44, 1.64, 1.84, 2.28, 2.52, 2.  
80, 2.92, 3.08, 3.28, 3.36, 3.60, 3.72, 3.84, 4.12],"tokens":["","","","","","","","","  
","","","","","","","","","","","","","",""]}  
---  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 0.502 s  
Real time factor (RTF): 0.502 / 15.289 = 0.033
```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
./build/bin/sherpa-onnéx-microphone-offline \  
--tokens=./sherpa-onnéx-conformer-zh-2023-05-23/tokens.txt \  
--encoder=./sherpa-onnéx-conformer-zh-2023-05-23/encoder-epoch-99-avg-1.onnéx \  
--decoder=./sherpa-onnéx-conformer-zh-2023-05-23/decoder-epoch-99-avg-1.onnéx \  
--joiner=./sherpa-onnéx-conformer-zh-2023-05-23/joiner-epoch-99-avg-1.onnéx
```

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

[cskuangfj/sherpa-onnéx-conformer-en-2023-03-18 \(English\)](#)

This model is converted from

<https://huggingface.co/csukuangfj/icefall-asr-librispeech-pruned-transducer-stateless3-2022-05-13>

which supports only English as it is trained on the [LibriSpeech](#) corpus.

You can find the training code at

https://github.com/k2-fsa/icefall/tree/master/egs/librispeech/ASR/pruned_transducer_stateless3

In the following, we describe how to download it and use it with `sherpa-onnix`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onné/releases/download/asr-models/sherpa-onné-
→conformer-en-2023-03-18.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onné/releases/download/asr-models/sherpa-onné-
→conformer-en-2023-03-18.tar.bz2

tar xvf sherpa-onné-conformer-en-2023-03-18.tar.bz2
rm sherpa-onné-conformer-en-2023-03-18.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-en-2023-03-18$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root 1.3M Apr 1 07:02 decoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 2.0M Apr 1 07:02 decoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 122M Apr 1 07:02 encoder-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 315M Apr 1 07:02 encoder-epoch-99-avg-1.onnx
-rw-r--r-- 1 kuangfangjun root 254K Apr 1 07:02 joiner-epoch-99-avg-1.int8.onnx
-rw-r--r-- 1 kuangfangjun root 1003K Apr 1 07:02 joiner-epoch-99-avg-1.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-conformer-en-2023-03-18/tokens.txt \
--encoder=./sherpa-onnx-conformer-en-2023-03-18/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-conformer-en-2023-03-18/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-en-2023-03-18/joiner-epoch-99-avg-1.onnx \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/0.wav \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/1.wav \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
sherpa-onnx-conformer-en-2023-03-18/encoder-epoch-99-avg-1.onnx", decoder_filename="./
sherpa-onnx-conformer-en-2023-03-18/decoder-epoch-99-avg-1.onnx", joiner_filename="./
sherpa-onnx-conformer-en-2023-03-18/joiner-epoch-99-avg-1.onnx"),
paraformer=OfflineParaformerModelConfig(model=""), tokens="./sherpa-onnx-conformer-en-
2023-03-18/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
2023-04-01 07:11:51.666456713 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
setaffinity_np failed for thread: 608379, index: 15, mask: {16, 52, }, error code: 22
error msg: Invalid argument. Specify the number of threads explicitly so the affinity
is not set.
2023-04-01 07:11:51.666458525 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
setaffinity_np failed for thread: 608380, index: 16, mask: {17, 53, }, error code: 22
(error msg: Invalid argument. Specify the number of threads explicitly so the affinity
is not set.)
```

(continued from previous page)

```

Started
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-conformer-en-2023-03-18/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID_
→ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY_
→ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH_
→ THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 2.264 s
Real time factor (RTF): 2.264 / 28.165 = 0.080

```

int8

The following code shows how to use int8 models to decode wave files:

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-conformer-en-2023-03-18/tokens.txt \
--encoder=./sherpa-onnx-conformer-en-2023-03-18/encoder-epoch-99-avg-1.int8.onnx \
--decoder=./sherpa-onnx-conformer-en-2023-03-18/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-en-2023-03-18/joiner-epoch-99-avg-1.int8.onnx \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/0.wav \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/1.wav \
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/8k.wav

```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
→ feature_dim=80), model_
→ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="./
→ sherpa-onnx-conformer-en-2023-03-18/encoder-epoch-99-avg-1.int8.onnx", decoder_
→ filename="./sherpa-onnx-conformer-en-2023-03-18/decoder-epoch-99-avg-1.onnx", joiner_
→ filename="./sherpa-onnx-conformer-en-2023-03-18/joiner-epoch-99-avg-1.int8.onnx"),_
→ paraformer=OfflineParaformerModelConfig(model=""), tokens="./sherpa-onnx-conformer-en-
→ 2023-03-18/tokens.txt", num_threads=2, debug=False), decoding_method="greedy_search")

```

(continues on next page)

(continued from previous page)

```

Creating recognizer ...
2023-04-01 07:13:26.514109433 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 608419, index: 15, mask: {16, 52, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
2023-04-01 07:13:26.514112711 [E:onnxruntime:, env.cc:251 ThreadMain] pthread_
↳ setaffinity_np failed for thread: 608420, index: 16, mask: {17, 53, }, error code: 22
↳ error msg: Invalid argument. Specify the number of threads explicitly so the affinity
↳ is not set.
Started
Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-conformer-en-2023-03-18/test_wavs/0.wav
AFTER EARLY NIGHTFALL THE YELLOW LAMPS WOULD LIGHT UP HERE AND THERE THE SQUALID
↳ QUARTER OF THE BROTHELS
-----
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/1.wav
GOD AS A DIRECT CONSEQUENCE OF THE SIN WHICH MAN THUS PUNISHED HAD GIVEN HER A LOVELY
↳ CHILD WHOSE PLACE WAS ON THAT SAME DISHONORED BOSOM TO CONNECT HER PARENT FOREVER WITH
↳ THE RACE AND DESCENT OF MORTALS AND TO BE FINALLY A BLESSED SOUL IN HEAVEN
-----
./sherpa-onnx-conformer-en-2023-03-18/test_wavs/8k.wav
YET THESE THOUGHTS AFFECTED HESTER PRYNNE LESS WITH HOPE THAN APPREHENSION
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 1.370 s
Real time factor (RTF): 1.370 / 28.165 = 0.049

```

Speech recognition from a microphone

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-conformer-en-2023-03-18/tokens.txt \
--encoder=./sherpa-onnx-conformer-en-2023-03-18/encoder-epoch-99-avg-1.onnx \
--decoder=./sherpa-onnx-conformer-en-2023-03-18/decoder-epoch-99-avg-1.onnx \
--joiner=./sherpa-onnx-conformer-en-2023-03-18/joiner-epoch-99-avg-1.onnx

```

8.25.5 Offline paraformer models

This section lists available offline paraformer models.

Paraformer models

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

csukuangfj/sherpa-onnx-paraformer-trilingual-zh-cantonese-en (Chinese + English + Cantonese)

Note: This model does not support timestamps. It is a trilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/dengcunqin/speech_seaco_paraformer_large_asr_nat-zh-cantonese-en-16k-common-vocab11666-pytorch/summary

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-trilingual-zh-cantonese-en.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-trilingual-zh-cantonese-en.tar.bz2

tar xvf sherpa-onnx-paraformer-trilingual-zh-cantonese-en.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-paraformer-trilingual-zh-cantonese-en$ ls -lh *.onnx

-rw-r--r-- 1 1001 127 234M Mar 10 02:12 model.int8.onnx
-rw-r--r-- 1 1001 127 831M Mar 10 02:12 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/tokens.txt \
--paraformer=./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/model.onnx \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/1.wav \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/2.wav \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/3-sichuan.wav \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/4-tianjin.wav \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav \
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./
sherpa-onnx-parafomer-trilingual-zh-cantonese-en/tokens.txt --paraformer=./sherpa-
onnx-parafomer-trilingual-zh-cantonese-en/model.onnx ./sherpa-onnx-parafomer-
trilingual-zh-cantonese-en/test_wavs/1.wav ./sherpa-onnx-parafomer-trilingual-zh-
cantonese-en/test_wavs/2.wav ./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_
wavs/3-sichuan.wav ./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/4-
tianjin.wav ./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav .
./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
decoder_filename="", joiner_filename=""),
paraformer=OfflineParaformerModelConfig(model="./sherpa-onnx-parafomer-trilingual-zh-
cantonese-en/model.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),
whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_ (continues on next page)
ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-parafomer-trilingual-
zh-cantonese-en/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type="472
"), lm_config=OfflineLmConfig(model="", scale=0.5), ctc fst_decoder_
config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
```

(continued from previous page)

```
Creating recognizer ...
Started
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 0, 13 vs -1
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 1, 15 vs -1
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 2, 40 vs -1
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 3, 41 vs -1
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 4, 37 vs -1
/project/sherpa-onnx/csrc/offline-paraformer-greedy-search-decoder.cc:Decode:65 time_
→stamp for batch: 5, 16 vs -1
Done!

./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "tokens_int":[]}
-----
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/2.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "tokens_int":[]}
-----
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/3-sichuan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "tokens_int":[]}
-----
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/4-tianjin.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "tokens_int":[]}
-----
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "tokens_int":[]}
-----
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav
{"text": " yesterday was today is tuesday ", "timestamps": [], "tokens":["yesterday", "was", "", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", ""], "tokens_int":[]}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 6.871 s
Real time factor (RTF): 6.871 / 42.054 = 0.163
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/model.int8.onnx \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/1.wav \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/2.wav \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/3-sichuan.wav \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/4-tianjin.wav \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav \
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./
sherpa-onnx-paraformer-trilingual-zh-cantonese-en/tokens.txt --paraformer=./sherpa-
onnx-paraformer-trilingual-zh-cantonese-en/model.int8.onnx ./sherpa-onnx-paraformer-
trilingual-zh-cantonese-en/test_wavs/1.wav ./sherpa-onnx-paraformer-trilingual-zh-
cantonese-en/test_wavs/2.wav ./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_
wavs/3-sichuan.wav ./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/4-
tianjin.wav ./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav .
./sherpa-onnx-paraformer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
feature_dim=80), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
decoder_filename="", joiner_filename=""),
paraformer=OfflineParaformerModelConfig(model="./sherpa-onnx-paraformer-trilingual-zh-
cantonese-en/model.int8.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),
whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe",
tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
ctc=OfflineWenetCtcModelConfig(model=""), tokens=".sherpa-onnx-paraformer-trilingual-
zh-cantonese-en/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
Creating recognizer ...
Started
```

(continues on next page)

(continued from previous page)

```

/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 0, 13 vs -1
/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 1, 15 vs -1
/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 2, 40 vs -1
/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 3, 41 vs -1
/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 4, 37 vs -1
/project/sherpa-onnx/csrc/offline-parafomer-greedy-search-decoder.cc:Decode:65 time_
↳ stamp for batch: 5, 16 vs -1
Done!

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/1.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ ""}
----  

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/2.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", ""]}
----  

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/3-sichuan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", "", "", ""}]
----  

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/4-tianjin.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", "", "", ""}]
----  

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/5-henan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", "", "", "", "", "", "", "", "", "", "", ""], 
↳ "", "", ""}]
----  

./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/test_wavs/6-zh-en.wav
{"text": " yesterday was today is tuesday ", "timestamps": [], "tokens":["yesterday", 
↳ "was", "", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", ""]}
----  

num threads: 2
decoding method: greedy_search
Elapsed seconds: 6.290 s
Real time factor (RTF): 6.290 / 42.054 = 0.150

```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-microphone-offline \  
--tokens=./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/tokens.txt \  
--parafomer=./sherpa-onnx-parafomer-trilingual-zh-cantonese-en/model.int8.onnx
```

csukuangfj/sherpa-onnx-parafomer-en-2024-03-09 (English)

Note: This model does not support timestamps. It supports only English.

This model is converted from

https://www.modelscope.cn/models/iic/speech_parafomer_asr-en-16k-vocab4199-pytorch/summary

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
-parafomer-en-2024-03-09.tar.bz2  
  
# For Chinese users  
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
-parafomer-en-2024-03-09.tar.bz2  
  
tar xvf sherpa-onnx-parafomer-en-2024-03-09.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-parafomer-en-2024-03-09$ ls -lh *.onnx  
-rw-r--r-- 1 1001 127 220M Mar 10 02:12 model.int8.onnx  
-rw-r--r-- 1 1001 127 817M Mar 10 02:12 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-paraformer-en-2024-03-09/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-en-2024-03-09/model.onnx \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=\
./sherpa-onnx-paraformer-en-2024-03-09/tokens.txt --paraformer=./sherpa-onnx-paraformer-\
en-2024-03-09/model.onnx ./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav ./\
sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav ./sherpa-onnx-paraformer-en-2024-\
03-09/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000, \
feature_dim=80), model_\
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", \
decoder_filename="", joiner_filename=""), \
paraformer=OfflineParaformerModelConfig(model="./sherpa-onnx-paraformer-en-2024-03-09/\
model.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""), \
whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe", \
tail_paddings=-1), tdnn=OfflineTdnmModelConfig(model=""), zipformer_\
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_\
ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-paraformer-en-2024-03-\
09/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""), lm_\
config=OfflineLmConfig(model="", scale=0.5), ctc_fst_decoder_\
config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_\
search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
Creating recognizer ...
Started
/project/sherpa-onnx/csrc/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav
{"text": " after early nightfall the yellow lamps would light up here and there the_\
squalid quarter of the brothels", "timestamps": [], "tokens": ["after", "early", "ni@@", \
"ght@@", "fall", "the", "yel@@", "low", "la@@", "mp@@", "s", "would", "light", "up", \
"here", "and", "there", "the", "squ@@", "al@@", "id", "quarter", "of", "the", "bro@@", \
"the@@", "ls"]}
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav
{"text": " god as a direct consequence of the sin which man thus punished had given her\u202e
\u202e a lovely child whose place was 'on' that same dishonoured bosom to connect her parent\u202e
\u202e for ever with the race and descent of mortals and to be finally a blessed soul in\u202e
\u202e heaven", "timestamps": [], "tokens": ["god", "as", "a", "direct", "con@@", "sequence",
\u202e "of", "the", "sin", "which", "man", "thus", "p@@", "uni@@", "shed", "had", "given",
\u202e "her", "a", "lo@@", "vely", "child", "whose", "place", "was", "'on'", "that", "same",
\u202e "di@@", "sh@@", "on@@", "ou@@", "red", "bo@@", "so@@", "m", "to", "connect", "her",
\u202e "paren@@", "t", "for", "ever", "with", "the", "race", "and", "des@@", "cent", "of",
\u202e "mor@@", "tal@@", "s", "and", "to", "be", "finally", "a", "bl@@", "essed", "soul", "in
\u202e ", "hea@@", "ven"]}

---
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/8k.wav
{"text": " yet these thoughts affected hester prynne less with hope than apprehension",
\u202e "timestamps": [], "tokens": ["yet", "these", "thoughts", "aff@@", "ected", "he@@", "ster",
\u202e "pr@@", "y@@", "n@@", "ne", "less", "with", "hope", "than", "ap@@", "pre@@", "hen@@"
\u202e ", "sion"]}

---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 7.173 s
Real time factor (RTF): 7.173 / 28.165 = 0.255
```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-paraformer-en-2024-03-09/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-en-2024-03-09/model.int8.onnx \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav \
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./
\u202e sherpa-onnx-paraformer-en-2024-03-09/tokens.txt --paraformer=./sherpa-onnx-paraformer-
\u202e en-2024-03-09/model.int8.onnx ./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav ./
\u202e sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav ./sherpa-onnx-paraformer-en-2024-
\u202e 03-09/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
\u202e feature_dim=80), model_
\u202e config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
\u202e decoder_filename="", joiner_filename=""),
\u202e paraformer=OfflineParaformerModelConfig(model=".sherpa-onnx-paraformer-en-2024-03-09")
\u202e model.int8.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),
\u202e whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="",
\u202e task="transcribe",
\u202e tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""),
zipformer_
\u202e ctc=OfflineZipformerCtcModelConfig(model=""),
\u202e wenet_ctc=OfflineWenetCtcModelConfig(model=""),
tokens=".sherpa-onnx-paraformer-en-2024-03-09/tokens.txt",
\u202e num_threads=2, debug=False, provider="cpu", model_type=""),
lm_
```

(continued from previous page)

```

Creating recognizer ...
Started
/project/sherpa-onnx/csrc/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/0.wav
{"text": " after early nightfall the yellow lamps would light up here and there the_
↳ squalid quarter of the brothels", "timestamps": [], "tokens": ["after", "early", "ni@@",
↳ "ght@@", "fall", "the", "yel@@", "low", "la@@", "mp@@", "s", "would", "light", "up",
↳ "here", "and", "there", "the", "squ@@", "al@@", "id", "quarter", "of", "the", "bro@@",
↳ "the@@", "ls"]}

---
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/1.wav
{"text": " god as a direct consequence of the sin which man thus punished had given her_
↳ a lovely child whose place was 'on' that same dishonoured bosom to connect her parent_
↳ for ever with the race and descent of mortals and to be finally a blessed soul in_
↳ heaven", "timestamps": [], "tokens": ["god", "as", "a", "direct", "con@@", "sequence",
↳ "of", "the", "sin", "which", "man", "thus", "p@@", "uni@@", "shed", "had", "given",
↳ "her", "a", "lo@@", "vely", "child", "whose", "place", "was", "'on'", "that", "same",
↳ "di@@", "sh@@", "on@@", "ou@@", "red", "bo@@", "so@@", "m", "to", "connect", "her",
↳ "paren@@", "t", "for", "ever", "with", "the", "race", "and", "des@@", "cent", "of",
↳ "mor@@", "tal@@", "s", "and", "to", "be", "finally", "a", "bl@@", "essed", "soul", "in
↳ ", "hea@@", "ven"]}

---
./sherpa-onnx-paraformer-en-2024-03-09/test_wavs/8k.wav
{"text": " yet these thoughts affected hester prynne less with hope than apprehension",
↳ "timestamps": [], "tokens": ["yet", "these", "thoughts", "aff@@", "ected", "he@@", "ster",
↳ "pr@@", "y@@", "n@@", "ne", "less", "with", "hope", "than", "ap@@", "pre@@", "hen@@",
↳ "sion"]}

---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 5.492 s
Real time factor (RTF): 5.492 / 28.165 = 0.195

```

Speech recognition from a microphone

```

cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-paraformer-en-2024-03-09/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-en-2024-03-09/model.int8.onnx

```

[csukuangfj/sherpa-onnx-parafomer-zh-small-2024-03-09 \(Chinese + English\)](#)

Note: This model does not support timestamps. It is a bilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/crazyant/speech_parafomer_asr_nat-zh-cn-16k-common-vocab8358-onnx/summary

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→parafomer-zh-small-2024-03-09.tar.bz2

# For Chinese users
wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→parafomer-zh-small-2024-03-09.tar.bz2

tar xvf sherpa-onnx-parafomer-zh-small-2024-03-09.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-parafomer-zh-small-2024-03-09$ ls -lh *.onnx
-rw-r--r-- 1 1001 127 79M Mar 10 00:48 model.int8.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

int8

The following code shows how to use `int8` models to decode wave files:

```
cd /path/to/sherpa-onnx
./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-parafomer-zh-small-2024-03-09/tokens.txt \
--parafomer=./sherpa-onnx-parafomer-zh-small-2024-03-09/model.int8.onnx \
./sherpa-onnx-parafomer-zh-small-2024-03-09/test_wavs/0.wav \
./sherpa-onnx-parafomer-zh-small-2024-03-09/test_wavs/1.wav \
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/8k.wav \
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/2-zh-en.wav \
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/3-sichuan.wav \
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/4-tianjin.wav \
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/5-henan.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

`CHCP 65001`

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./
˓→sherpa-onnx-paraformer-zh-small-2024-03-09/tokens.txt --paraformer=./sherpa-onnx-
˓→paraformer-zh-small-2024-03-09/model.int8.onnx ./sherpa-onnx-paraformer-zh-small-2024-
˓→03-09/test_wavs/0.wav ./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/1.wav ./
˓→sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/8k.wav ./sherpa-onnx-paraformer-
˓→zh-small-2024-03-09/test_wavs/2-zh-en.wav ./sherpa-onnx-paraformer-zh-small-2024-03-09/
˓→test_wavs/3-sichuan.wav ./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/4-
˓→tianjin.wav ./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/5-henan.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
˓→feature_dim=80), model_
˓→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
˓→decoder_filename="", joiner_filename=""),
˓→paraformer=OfflineParaformerModelConfig(model="../sherpa-onnx-paraformer-zh-small-2024-
˓→03-09/model.int8.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),
˓→whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe",
˓→tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
˓→ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓→ctc=OfflineWenetCtcModelConfig(model=""), tokens="../sherpa-onnx-paraformer-zh-small-
˓→2024-03-09/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""),
˓→config=OfflineLmConfig(model="", scale=0.5), ctc_fst_decoder_
˓→config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
˓→search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
Creating recognizer ...
Started
/project/sherpa-onnx/csrc/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/0.wav
{"text": "", "timestamps": [], "tokens": ["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

(continues on next page)

(continued from previous page)

```
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/1.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/8k.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/2-zh-en.wav
{"text": " yesterday was today is tuesday ", "timestamps": [], "tokens":["ye@@", "ster@@", "day", "was", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/3-sichuan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/4-tianjin.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-paraformer-zh-small-2024-03-09/test_wavs/5-henan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 3.562 s
Real time factor (RTF): 3.562 / 47.023 = 0.076
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-paraformer-zh-small-2024-03-09/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-small-2024-03-09/model.int8.onnx
```

csukuangfj/sherpa-onnx-parafomer-zh-2024-03-09 (Chinese + English)

Note: This model does not support timestamps. It is a bilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/crazyant/speech_parafomer_asr_nat-zh-cn-16k-common-vocab8358-onnx/summary

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→parafomer-zh-2024-03-09.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→parafomer-zh-2024-03-09.tar.bz2

tar xvf sherpa-onnx-parafomer-zh-2024-03-09.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-parafomer-zh-2024-03-09$ ls -lh *.onnx
-rw-r--r-- 1 1001 127 217M Mar 10 02:22 model.int8.onnx
-rw-r--r-- 1 1001 127 785M Mar 10 02:22 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use fp32 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-parafomer-zh-2024-03-09/tokens.txt \
--parafomer=./sherpa-onnx-parafomer-zh-2024-03-09/model.onnx \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/0.wav \
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/1.wav \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/8k.wav \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/2-zh-en.wav \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/3-sichuan.wav \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/4-tianjin.wav \
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/5-henan.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

`CHCP 65001`

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./
˓sherpa-onnx-parafomer-zh-2024-03-09/tokens.txt --parafomer=./sherpa-onnx-parafomer-
˓zh-2024-03-09/model.onnx ./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/0.wav ./
˓sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/1.wav ./sherpa-onnx-parafomer-zh-2024-
˓03-09/test_wavs/8k.wav ./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/2-zh-en.wav ./
˓sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/3-sichuan.wav ./sherpa-onnx-parafomer-
˓zh-2024-03-09/test_wavs/4-tianjin.wav ./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/
˓5-henan.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,_
˓feature_dim=80), model_
˓config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",_
˓decoder_filename="", joiner_filename=""),_
˓paraformer=OfflineParaformerModelConfig(model="./sherpa-onnx-parafomer-zh-2024-03-09/
˓model.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),_
˓whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
˓ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-parafomer-zh-2024-03-
˓09/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""), lm_
˓config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
˓config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
˓search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
Creating recognizer ...
Started
/project/sherpa-onnx/csrc/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!
./sherpa-onnx-parafomer-zh-2024-03-09/test_wavs/0.wav
```

(continues on next page)

(continued from previous page)

```
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],  
---  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/1.wav  
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],  
---  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/8k.wav  
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],  
---  

```

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-offline \  
--tokens=./sherpa-onnx-paraformer-zh-2024-03-09/tokens.txt \  
--paraformer=./sherpa-onnx-paraformer-zh-2024-03-09/model.int8.onnx \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/0.wav \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/1.wav \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/8k.wav \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/2-zh-en.wav \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/3-sichuan.wav \  
.sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/4-tianjin.wav \  
"
```

(continues on next page)

(continued from previous page)

./sherpa-onnéx-parformer-zh-2024-03-09/test_wavs/5-henan.wav

Note: Please use `./build/bin/Release/sherpa-ONNX-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/project/sherpa-ONNX/csrc/parse-options.cc:Read:361 sherpa-ONNX-offline --tokens=./  
sherpa-ONNX-paraformer-zh-2024-03-09/tokens.txt --paraformer=./sherpa-ONNX-paraformer-  
zh-2024-03-09/model.onnx ./sherpa-ONNX-paraformer-zh-2024-03-09/test_wavs/0.wav ./  
sherpa-ONNX-paraformer-zh-2024-03-09/test_wavs/1.wav ./sherpa-ONNX-paraformer-zh-2024-  
03-09/test_wavs/8k.wav ./sherpa-ONNX-paraformer-zh-2024-03-09/test_wavs/2-zh-en.wav ./  
sherpa-ONNX-paraformer-zh-2024-03-09/test_wavs/3-sichuan.wav ./sherpa-ONNX-paraformer-  
zh-2024-03-09/test_wavs/4-tianjin.wav ./sherpa-ONNX-paraformer-zh-2024-03-09/test_wavs/  
5-henan.wav
```

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000, ↴  
↳ feature_dim=80), model_  
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", ↴  
↳ decoder_filename="", joiner_filename=""), ↴  
↳ paraformer=OfflineParaformerModelConfig(model=".sherpa-ONNX-paraformer-zh-2024-03-09/", ↴  
↳ model.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""), ↴  
↳ whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe", ↴  
↳ tail_paddings=-1), tdnn=OfflineTdnNModelConfig(model=""), zipformer_  
↳ ctc=OfflineZipformerCtcModelConfig(model=""), wenet_  
↳ ctc=OfflineWenetCtcModelConfig(model=""), tokens=".sherpa-ONNX-paraformer-zh-2024-03-09/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type=""), lm_  
↳ config=OfflineLmConfig(model="", scale=0.5), ctc_fst_decoder_  
↳ config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_  
↳ search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)
```

Creating recognizer ...

Started

```
/project/sherpa-ONNX/cs/src/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:  
    in_sample_rate: 8000  
    output sample rate: 16000
```

Done!

(continues on next page)

(continued from previous page)

```
-----
./sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/8k.wav
{"text": "", "timestamps": [], "tokens":["","","","","","","","","","","","","","","","","","","","","","",""]}

-----
./sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/2-zh-en.wav
{"text": " yesterday was today is tuesday ", "timestamps": [], "tokens":["ye@@", "ster@@", "day", "was", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", """]}

-----
./sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/3-sichuan.wav
{"text": "", "timestamps": [], "tokens":["","","","","","","","","","","","","","","","","","","","","","",""]}

-----
./sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/4-tianjin.wav
{"text": "", "timestamps": [], "tokens":["","","","","","","","","","","","","","","","","","","","",""]}

-----
./sherpa-onnx-paraformer-zh-2024-03-09/test_wavs/5-henan.wav
{"text": "", "timestamps": [], "tokens":["","","","","","","","","","","","","","","","","","","","",""]}

-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 6.829 s
Real time factor (RTF): 6.829 / 47.023 = 0.145
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-paraformer-zh-2024-03-09/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2024-03-09/model.int8.onnx
```

csukuangfj/sherpa-onnx-paraformer-zh-2023-03-28 (Chinese + English)

Note: This model does not support timestamps. It is a bilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/damo/speech_paraformer-large_asr_nat-zh-cn-16k-common-vocab8404-pytorch

The code for converting can be found at

<https://huggingface.co/csukuangfj/paraformer-onnxruntime-python-example/tree/main>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-zh-2023-03-28.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-zh-2023-03-28.tar.bz2

tar xvf sherpa-onnx-paraformer-zh-2023-03-28.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-paraformer-zh-2023-03-28$ ls -lh *.onnx
-rw-r--r-- 1 kuangfangjun root 214M Apr  1 07:28 model.int8.onnx
-rw-r--r-- 1 kuangfangjun root 824M Apr  1 07:28 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

fp32

The following code shows how to use `fp32` models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-paraformer-zh-2023-03-28/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2023-03-28/model.onnx \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/3-sichuan.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/4-tianjin.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/5-henan.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/6-zh-en.wav \
./sherpa-onnx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

(continued from previous page)

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx  
  
./build/bin/sherpa-onnéx-offline \  
--tokens=./sherpa-onnéx-paraformer-zh-2023-03-28/tokens.txt \  
--paraformer=./sherpa-onnéx-paraformer-zh-2023-03-28/model.int8.onnéx \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/0.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/1.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/2.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/3-sichuan.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/4-tianjin.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/5-henan.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/6-zh-en.wav \  
./sherpa-onnéx-paraformer-zh-2023-03-28/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/4-tianjin.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/5-henan.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/6-zh-en.wav
{"text": " yesterday was today is tuesday ", "timestamps": [], "tokens":["ye@@", "ster@@", "day", "was", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", "", ""], "score": 0.0}
-----
./sherpa-onnx-parafomer-zh-2023-03-28/test_wavs/8k.wav
{"text": "", "timestamps": [], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""], "score": 0.0}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 6.439 s
Real time factor (RTF): 6.439 / 51.236 = 0.126
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-microphone-offline \
--tokens=./sherpa-onnx-parafomer-zh-2023-03-28/tokens.txt \
--parafomer=./sherpa-onnx-parafomer-zh-2023-03-28/model.int8.onnx
```

csukuangfj/sherpa-onnx-parafomer-zh-2023-09-14 (Chinese + English))

Note: This model supports timestamps. It is a bilingual model, supporting both Chinese and English. ()

This model is converted from

https://www.modelscope.cn/models/iic/speech_parafomer-large-vad-punc_asr_nat-zh-cn-16k-common-vocab8404-onnx/summary

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-zh-2023-09-14.tar.bz2

# For Chinese users
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→paraformer-zh-2023-09-14.tar.bz2

tar xvf sherpa-onnx-paraformer-zh-2023-09-14.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-paraformer-zh-2023-09-14$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 232M Sep 14 13:46 model.int8.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

int8

The following code shows how to use int8 models to decode wave files:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-paraformer-zh-2023-09-14/tokens.txt \
--paraformer=./sherpa-onnx-paraformer-zh-2023-09-14/model.int8.onnx \
--model-type=paraformer \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/0.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/1.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/2.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/3-sichuan.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/4-tianjin.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/5-henan.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/6-zh-en.wav \
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

```
CHCP 65001
```

in your commandline.

You should see the following output:

```
/project/sherpa-onnx/csrc/parse-options.cc:Read:361 sherpa-onnx-offline --tokens=./  
sherpa-onnx-paraformer-zh-2023-09-14/tokens.txt --paraformer=./sherpa-onnx-paraformer-  
zh-2023-09-14/model.int8.onnx --model-type=paraformer ./sherpa-onnx-paraformer-zh-2023-  
09-14/test_wavs/0.wav ./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/1.wav ./sherpa-  
onnx-paraformer-zh-2023-09-14/test_wavs/2.wav ./sherpa-onnx-paraformer-zh-2023-09-14/  
test_wavs/3-sichuan.wav ./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/4-tianjin.wav  
./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/5-henan.wav ./sherpa-onnx-paraformer-  
zh-2023-09-14/test_wavs/6-zh-en.wav ./sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/  
8k.wav  
  
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,  
feature_dim=80), model_  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",  
decoder_filename="", joiner_filename ""),  
paraformer=OfflineParaformerModelConfig(model="./sherpa-onnx-paraformer-zh-2023-09-14/  
model.int8.onnx"), nemo_ctc=OfflineNemoEncDecCtcModelConfig(model=""),  
whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe",  
tail_paddings=-1), tdnm=OfflineTdnmModelConfig(model=""), zipformer_  
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_  
ctc=OfflineWenetCtcModelConfig(model=""), tokens="./sherpa-onnx-paraformer-zh-2023-09-  
14/tokens.txt", num_threads=2, debug=False, provider="cpu", model_type="paraformer"),  
lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_  
config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_  
search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0)  
Creating recognizer ...  
Started  
/project/sherpa-onnx/csrc/offline-stream.cc:AcceptWaveformImpl:119 Creating a resampler:  
    in_sample_rate: 8000  
    output_sample_rate: 16000  
  
Done!  
  
.sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/0.wav  
{"text": "", "timestamps": [0.36, 0.48, 0.62, 0.72, 0.86, 1.02, 1.32, 1.74, 1.90, 2.12,  
2.20, 2.38, 2.50, 2.62, 2.74, 3.18, 3.32, 3.52, 3.62, 3.74, 3.82, 3.90, 3.98, 4.08, 4.  
20, 4.34, 4.56, 4.74, 5.10], "tokens": ["", "", "", "", "", "", "", "", "", "", "",  
 "", "", "", "", "", "", "", "", "", "", "", ""]}  
---  
.sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/1.wav  
{"text": "", "timestamps": [0.16, 0.30, 0.42, 0.56, 0.72, 0.96, 1.08, 1.20, 1.30, 2.08,  
2.26, 2.44, 2.58, 2.72, 2.98, 3.14, 3.26, 3.46, 3.62, 3.80, 3.88, 4.02, 4.12, 4.20, 4.  
36, 4.56], "tokens": ["", "", "", "", "", "", "", "", "", "", "", ""]}  
---  
.sherpa-onnx-paraformer-zh-2023-09-14/test_wavs/2.wav
```

(continues on next page)

(continued from previous page)

```
{"text": "", "timestamps": [0.34, 0.54, 0.66, 0.80, 1.08, 1.52, 1.72, 1.90, 2.40, 2.68, 2.86, 2.96, 3.16, 3.26, 3.46, 3.54, 3.66, 3.80, 3.90], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

```
./sherpa-onnx-parafomer-zh-2023-09-14/test_wavs/3-sichuan.wav
```

```
{"text": "", "timestamps": [0.16, 0.30, 0.56, 0.72, 0.92, 1.18, 1.32, 1.88, 2.24, 2.40, 3.16, 3.28, 3.40, 3.54, 3.76, 3.88, 4.06, 4.24, 4.36, 4.56, 4.66, 4.88, 5.14, 5.30, 5.44, 5.60, 5.72, 5.84, 5.96, 6.14, 6.24, 6.38, 6.56, 6.78, 6.98, 7.08, 7.22, 7.38, 7.50, 7.62], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

```
./sherpa-onnx-parafomer-zh-2023-09-14/test_wavs/4-tianjin.wav
```

```
{"text": "", "timestamps": [0.08, 0.24, 0.36, 0.56, 0.66, 0.78, 1.04, 1.14, 1.26, 1.38, 1.50, 1.58, 1.70, 1.84, 2.28, 2.38, 2.64, 2.74, 3.08, 3.28, 3.66, 3.80, 3.94, 4.14, 4.34, 4.64, 4.84, 4.94, 5.12, 5.24, 5.84, 6.10, 6.24, 6.44, 6.54, 6.66, 6.86, 7.02, 7.14, 7.24, 7.44], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

```
./sherpa-onnx-parafomer-zh-2023-09-14/test_wavs/5-henan.wav
```

```
{"text": "", "timestamps": [0.08, 0.20, 0.30, 0.42, 0.94, 1.14, 1.26, 1.46, 1.66, 2.28, 2.50, 2.62, 2.70, 2.82, 2.98, 3.14, 3.28, 3.52, 3.70, 3.86, 4.94, 5.06, 5.18, 5.30, 5.42, 5.66, 5.76, 5.94, 6.08, 6.24, 6.38, 6.60, 6.78, 6.96, 7.10, 7.30, 7.50, 7.62], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

```
./sherpa-onnx-parafomer-zh-2023-09-14/test_wavs/6-zh-en.wav
```

```
{"text": " yesterday was today is tuesday ", "timestamps": [0.36, 0.60, 0.84, 1.22, 2.24, 2.44, 2.74, 3.52, 4.06, 4.68, 5.00, 5.12, 5.76, 5.96, 6.24, 6.82, 7.02, 7.26], "tokens":["ye@@", "ster@@", "day", "was", "", "", "", "today", "is", "tu@@", "es@@", "day", "", "", "", "", "", ""]}
```

```
./sherpa-onnx-parafomer-zh-2023-09-14/test_wavs/8k.wav
```

```
{"text": "", "timestamps": [0.48, 0.78, 1.04, 1.18, 1.52, 1.78, 2.06, 2.18, 2.50, 2.66, 2.88, 3.10, 3.30], "tokens":["", "", "", "", "", "", "", "", "", "", "", "", ""]}
```

```
num threads: 2
decoding method: greedy_search
Elapsed seconds: 9.206 s
Real time factor (RTF): 9.206 / 51.236 = 0.180
```

Speech recognition from a microphone

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-microphone-offline \  
--tokens=./sherpa-onnx-paraformer-zh-2023-09-14/tokens.txt \  
--paraformer=./sherpa-onnx-paraformer-zh-2023-09-14/model.int8.onnx \  
--model-type=paraformer
```

8.25.6 Offline CTC models

This section lists available offline CTC models.

NeMo

This page lists all offline CTC models from NeMo.

Hint: Please refer to https://catalog.ngc.nvidia.com/orgs/nvidia/collections/nemo_asr for a list of pre-trained NeMo models.

How to export models from NeMo to sherpa-onnx

This section describes how to export CTC models from NeMo to `sherpa-onnx`.

Hint: Please refer to https://catalog.ngc.nvidia.com/orgs/nvidia/collections/nemo_asr for a list of pre-trained NeMo models.

You can use method described in this section to convert more models to `sherpa-onnx`.

Let us take the following model as an example:

https://ngc.nvidia.com/models/nvidia:nemo:stt_en_conformer_ctc_small.

Hint: You can find the exported files in this example by visiting

<https://huggingface.co/csukuangfj/sherpa-onnx-nemo-ctc-en-conformer-small>

The steps to export it to `sherpa-onnx` are given below.

Step 1: Export model.onnx

The first step is to obtain `model.onnx`.

```
import nemo.collections.asr as nemo_asr
m = nemo_asr.models.EncDecCTCModelBPE.from_pretrained('stt_en_conformer_ctc_small')
m.export('model.onnx')
```

Step 2: Add metadata

To be usable in `sherpa-onnx`, we have to use `add-model-metadata.py` to add metadata to `model.onnx`.

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-nemo-ctc-en-conformer-small/resolve/
└─main/add-model-metadata.py

# The following command changes model.onnx in-place
python3 add-model-metadata.py
```

Step 3: Obtain model.int8.onnx

We can use `quantize-model.py` to obtain a quantized version of `model.onnx`:

```
wget https://huggingface.co/csukuangfj/sherpa-onnx-nemo-ctc-en-conformer-small/resolve/
└─main/quantize-model.py

# The following command will generate model.int8.onnx
python3 ./quantize-model.py
```

Step 4: Obtain tokens.txt

Use the following command to obtain `tokens.txt`:

```
import nemo.collections.asr as nemo_asr
m = nemo_asr.models.EncDecCTCModelBPE.from_pretrained('stt_en_conformer_ctc_small')

with open('tokens.txt', 'w') as f:
    for i, s in enumerate(m.decoder.vocabulary):
        f.write(f'{s} {i}\n')
        f.write(f'<blk> {i+1}\n')
```

English

Hint: Please refer to [Installation](#) to install `sherpa-onnx` before you read this section.

Note: We use `./build/bin/sherpa-offline` as an example in this section. You can use other scripts such as

- `./build/bin/sherpa-onnx-microphone-offline`
 - `./build/bin/sherpa-onnx-offline-websocket-server`
 - `python-api-examples/offline-decode-files.py`
-

This page lists offline CTC models from NeMo for English.

stt_en_citrinet_512

This model is converted from

https://catalog.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_citrinet_512

Citrinet-512 model which has been trained on the ASR Set dataset with over 7000 hours of english speech.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-
      ↪ctc-en-citrinet-512.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
      ↪nemo-ctc-en-citrinet-512.tar.bz2

tar xvf sherpa-onnx-nemo-ctc-en-citrinet-512.tar.bz2
rm sherpa-onnx-nemo-ctc-en-citrinet-512.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-nemo-ctc-en-citrinet-512 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 36M Apr  7 16:10 model.int8.onnx
-rw-r--r-- 1 fangjun staff 142M Apr  7 14:24 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

The following code shows how to use fp32 models to decode wave files. Please replace `model.onnx` with `model.int8.onnx` to use int8 quantized model.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-nemo-ctc-en-citrinet-512/tokens.txt \
--nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-citrinet-512/model.onnx \
--num-threads=2 \
--decoding-method=greedy_search \
--debug=false \
./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/8k.wav
```

Note: Please use `./build/bin/Release/sherpa-onnx-offline.exe` for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-nemo-ctc-en-citrinet-512/tokens.
↳ txt --nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-citrinet-512/model.onnx --num-threads=2
↳ --decoding-method=greedy_search --debug=false ./sherpa-onnx-nemo-ctc-en-citrinet-512/
↳ test_wavs/0.wav ./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/1.wav ./sherpa-onnx-
↳ nemo-ctc-en-citrinet-512/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
↳ decoder_filename="", joiner_filename=""),
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model="./sherpa-onnx-nemo-ctc-en-citrinet-512/
↳ model.onnx"), tokens="./sherpa-onnx-nemo-ctc-en-citrinet-512/tokens.txt", num_
↳ threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
↳ cc:AcceptWaveformImpl:105 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid
↳ quarter of the brothels
```

(continues on next page)

(continued from previous page)

```
-----
./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/1.wav
god as a direct consequence of the sin which man thus punished had given her a lovely_
_child whose place was on that same dishonoured bosom to connect her parent for ever_
_with the race and descent of mortals and to be finally a blessed soul in heaven
-----
./sherpa-onnx-nemo-ctc-en-citrinet-512/test_wavs/8k.wav
yet these thoughts affected hester prynne less with hope than apprehension
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 4.963 s
Real time factor (RTF): 4.963 / 28.165 = 0.176
```

stt_en_conformer_ctc_small

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_ctc_small

It contains small size versions of Conformer-CTC (13M parameters) trained on NeMo ASRSet with around 16000 hours of english speech. The model transcribes speech in lower case english alphabet along with spaces and apostrophes.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-
_ctc-en-conformer-small.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
_nemo-ctc-en-conformer-small.tar.bz2

tar xvf sherpa-onnx-nemo-ctc-en-conformer-small.tar.bz2
rm sherpa-onnx-nemo-ctc-en-conformer-small.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-nemo-ctc-en-conformer-small fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 44M Apr  7 20:24 model.int8.onnx
-rw-r--r-- 1 fangjun staff 81M Apr  7 18:56 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

The following code shows how to use fp32 models to decode wave files. Please replace model.onnx with model.int8.onnx to use int8 quantized model.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-nemo-ctc-en-conformer-small/tokens.txt \
--nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-small/model.onnx \
--num-threads=2 \
--decoding-method=greedy_search \
--debug=false \
./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
~/build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-nemo-ctc-en-conformer-small/
~/tokens.txt --nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-small/model.onnx --num-
~/threads=2 --decoding-method=greedy_search --debug=false ./sherpa-onnx-nemo-ctc-en-
~/conformer-small/test_wavs/0.wav ./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/1.
~/wav ./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
~/feature_dim=80), model_
~/config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
~/decoder_filename="", joiner_filename=""),
~/paraformer=OfflineParaformerModelConfig(model=""), nemo_
~/ctc=OfflineNemoEncDecCtcModelConfig(model="./sherpa-onnx-nemo-ctc-en-conformer-small/
~/model.onnx"), tokens="./sherpa-onnx-nemo-ctc-en-conformer-small/tokens.txt", num_
~/threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
~/cc:AcceptWaveformImpl:105 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
~/quarter of the brothels
```

(continues on next page)

(continued from previous page)

```
----  
./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/1.wav  
god as a direct consequence of the sin which man thus punished had given her a lovely  
child whose place was on that same dishonoured bosom to connect her parent for ever  
with the race and descent of mortals and to be finally a blessed soul in heaven  
----  
./sherpa-onnx-nemo-ctc-en-conformer-small/test_wavs/8k.wav  
yet these thoughts affected hester prin less with hope than apprehension  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 0.665 s  
Real time factor (RTF): 0.665 / 28.165 = 0.024
```

stt_en_conformer_ctc_medium

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_ctc_medium

It contains medium size versions of Conformer-CTC (around 30M parameters) trained on NeMo ASRSet with around 16000 hours of english speech. The model transcribes speech in lower case english alphabet along with spaces and apostrophes.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-  
-ctc-en-conformer-medium.tar.bz2  
  
# For Chinese users, please use the following mirror  
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
-nemo-ctc-en-conformer-medium.tar.bz2  
  
tar xvf sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2  
rm sherpa-onnx-nemo-ctc-en-conformer-medium.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-nemo-ctc-en-conformer-medium fangjun$ ls -lh *.onnx  
-rw-r--r-- 1 fangjun staff 64M Apr 7 20:44 model.int8.onnx  
-rw-r--r-- 1 fangjun staff 152M Apr 7 20:43 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

The following code shows how to use fp32 models to decode wave files. Please replace model.onnx with model.int8.onnx to use int8 quantized model.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt \
--nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx \
--num-threads=2 \
--decoding-method=greedy_search \
--debug=false \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
~/build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-nemo-ctc-en-conformer-medium/
~/tokens.txt --nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-medium/model.onnx --
~/num-threads=2 --decoding-method=greedy_search --debug=false ./sherpa-onnx-nemo-ctc-en-
~/conformer-medium/test_wavs/0.wav ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/
~/1.wav ./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
~/feature_dim=80), model_
~/config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
~/decoder_filename="", joiner_filename=""),
~/paraformer=OfflineParaformerModelConfig(model=""), nemo_
~/ctc=OfflineNemoEncDecCtcModelConfig(model="./sherpa-onnx-nemo-ctc-en-conformer-medium/
~/model.onnx"), tokens="./sherpa-onnx-nemo-ctc-en-conformer-medium/tokens.txt", num_
~/threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
~/cc:AcceptWaveformImpl:105 Creating a resampler:
  in_sample_rate: 8000
  output_sample_rate: 16000

Done!

./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
~/quarter of the brothels
```

(continues on next page)

(continued from previous page)

```
----  
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/1.wav  
god as a direct consequence of the sin which man thus punished had given her a lovely  
→ child whose place was on that same dishonored bosom to connect her parent for ever  
→ with the race and descent of mortals and to be finally a blessed soul in heaven  
----  
./sherpa-onnx-nemo-ctc-en-conformer-medium/test_wavs/8k.wav  
yet these thoughts affected hester pryne less with hope than apprehension  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 1.184 s  
Real time factor (RTF): 1.184 / 28.165 = 0.042
```

stt_en_conformer_ctc_large

This model is converted from

https://registry.ngc.nvidia.com/orgs/nvidia/teams/nemo/models/stt_en_conformer_ctc_large

It contains large size versions of Conformer-CTC (around 120M parameters) trained on NeMo ASRSet with around 24500 hours of english speech. The model transcribes speech in lower case english alphabet along with spaces and apostrophes

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-nemo-  
→ ctc-en-conformer-large.tar.bz2  
  
# For Chinese users, please use the following mirror  
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
→ nemo-ctc-en-conformer-large.tar.bz2  
  
tar xvf sherpa-onnx-nemo-ctc-en-conformer-large.tar.bz2  
rm sherpa-onnx-nemo-ctc-en-conformer-large.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
sherpa-onnx-nemo-ctc-en-conformer-large fangjun$ ls -lh *.onnx  
-rw-r--r-- 1 fangjun staff 162M Apr 7 22:01 model.int8.onnx  
-rw-r--r-- 1 fangjun staff 508M Apr 7 22:01 model.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

The following code shows how to use fp32 models to decode wave files. Please replace model.onnx with model.int8.onnx to use int8 quantized model.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-nemo-ctc-en-conformer-large/tokens.txt \
--nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-large/model.onnx \
--num-threads=2 \
--decoding-method=greedy_search \
--debug=false \
./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/0.wav \
./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/1.wav \
./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/8k.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-nemo-ctc-en-conformer-large/
↳ tokens.txt --nemo-ctc-model=./sherpa-onnx-nemo-ctc-en-conformer-large/model.onnx --num-
↳ threads=2 --decoding-method=greedy_search --debug=false ./sherpa-onnx-nemo-ctc-en-
↳ conformer-large/test_wavs/0.wav ./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/1.
↳ wav ./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/8k.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,_
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",_
↳ decoder_filename="", joiner_filename=""),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model="./sherpa-onnx-nemo-ctc-en-conformer-large/
↳ model.onnx"), tokens="./sherpa-onnx-nemo-ctc-en-conformer-large/tokens.txt", num_
↳ threads=2, debug=False), decoding_method="greedy_search")
Creating recognizer ...
Started
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/offline-stream.
↳ cc:AcceptWaveformImpl:105 Creating a resampler:
    in_sample_rate: 8000
    output_sample_rate: 16000

Done!

./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/0.wav
after early nightfall the yellow lamps would light up here and there the squalid_
↳ quarter of the brothels
```

(continues on next page)

(continued from previous page)

```
----  
./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/1.wav  
god as a direct consequence of the sin which man thus punished had given her a lovely  
→ child whose place was on that same dishonored bosom to connect her parent for ever  
→ with the race and descent of mortals and to be finally a blesed soul in heaven  
----  
./sherpa-onnx-nemo-ctc-en-conformer-large/test_wavs/8k.wav  
yet these thoughts afected hester pryne les with hope than apprehension  
----  
num threads: 2  
decoding method: greedy_search  
Elapsed seconds: 3.553 s  
Real time factor (RTF): 3.553 / 28.165 = 0.126
```

yesno

This section describes how to use the `tdnn` model of the `yesno` dataset from `icefall` in `sherpa-onnx`.

Note: It is a **non-streaming** model and it can only recognize two words in Hebrew: yes and no.

To download the model, please use:

```
cd /path/to/sherpa-onnx  
  
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-tdnn-  
→yesno.tar.bz2  
  
# For Chinese users, please use the following mirror  
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-  
→tdnn-yesno.tar.bz2  
  
tar xvf sherpa-onnx-tdnn-yesno.tar.bz2  
rm sherpa-onnx-tdnn-yesno.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
sherpa-onnx-tdnn-yesno fangjun$ ls -lh *.onnx  
-rw-r--r-- 1 fangjun staff 55K Aug 12 17:02 model-epoch-14-avg-2.int8.onnx  
-rw-r--r-- 1 fangjun staff 54K Aug 12 17:02 model-epoch-14-avg-2.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

The following code shows how to use fp32 models to decode wave files. Please replace `model-epoch-14-avg-2.int8.onnx` with `model-epoch-14-avg-2.int8.onnx` to use the int8 quantized model.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--sample-rate=8000 \
--feat-dim=23 \
--tokens=./sherpa-onnx-tdnn-yesno/tokens.txt \
--tdnn-model=./sherpa-onnx-tdnn-yesno/model-epoch-14-avg-2.onnx \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_0_1_0_0_0_1.wav \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_0_0_1_0.wav \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_0_1_1_1.wav \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_1_0_0_1.wav \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_1_0_0_0_1.wav \
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_1_0_1_1_0.wav
```

The output is given below:

```
OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=8000,
→ feature_dim=23), model_
→ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
→ decoder_filename="", joiner_filename=""),
→ paraformer=OfflineParaformerModelConfig(model=""), nemo_
→ ctc=OfflineNemoEncDecCtcModelConfig(model=""),
→ whisper=OfflineWhisperModelConfig(encoder="", decoder=""),
→ tdnn=OfflineTdnnModelConfig(model="./sherpa-onnx-tdnn-yesno/model-epoch-14-avg-2.onnx
→ "), tokens="./sherpa-onnx-tdnn-yesno/tokens.txt", num_threads=2, debug=False, provider=
→ "cpu", model_type=""), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method=
→ "greedy_search", max_active_paths=4, context_score=1.5)
Creating recognizer ...
Started
Done!

./sherpa-onnx-tdnn-yesno/test_wavs/0_0_0_1_0_0_0_1.wav
{"text": "NNNNNNNN", "timestamps": "[]", "tokens": ["N", "N", "N", "Y", "N", "N", "N", "Y"]}
---
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_0_0_1_0.wav
{"text": "NNNNNNYN", "timestamps": "[]", "tokens": ["N", "N", "Y", "N", "N", "N", "Y", "N"]}
---
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_0_1_1_1.wav
{"text": "NNYNNYYY", "timestamps": "[]", "tokens": ["N", "N", "Y", "N", "N", "Y", "Y"]}
---
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_0_1_0_1.wav
{"text": "NNYNYNYY", "timestamps": "[]", "tokens": ["N", "N", "Y", "N", "Y", "N", "Y"]}
---
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_1_0_0_1.wav
{"text": "NNYYNNNY", "timestamps": "[]", "tokens": ["N", "N", "Y", "Y", "N", "N", "Y"]}
---
./sherpa-onnx-tdnn-yesno/test_wavs/0_0_1_1_0_1_1_0.wav
{"text": "NNYYNYYN", "timestamps": "[]", "tokens": ["N", "N", "Y", "Y", "N", "Y", "N"]}
---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 0.071 s
Real time factor (RTF): 0.071 / 38.530 = 0.002
```

Note: In the above output, N represents NO, while Y is YES. So for the last wave, NNYYNYYN means NO NO YES YES NO YES YES NO.

In the filename of the last wave `0_0_1_1_0_1_1_0.wav`, 0 means NO and 1 means YES. So the ground truth of the last wave is NO NO YES YES NO YES YES NO.

8.25.7 TeleSpeech

This page lists all offline CTC models from <https://github.com/Tele-AI/TeleSpeech-ASR>.

Hint: Please see the license at [TeleSpeech.pdf](#)

Models from <https://github.com/Tele-AI/TeleSpeech-ASR> are subject to the above license if you want to use them for commercial purpose.

How to export models from Tele-AI/TeleSpeech-ASR to sherpa-onnx

This section describes how to export CTC models from *Tele-AI/TeleSpeech-ASR* to `sherpa-onnx`.

Step 1: Export model.onnx

The first step is to obtain `model.onnx`.

Please see https://github.com/lovemefan/telespeech-asr-python/blob/main/telespeechasr/onnx/onnx_export.py for details.

Step 2: Add metadata

To be usable in `sherpa-onnx`, we have to use `add-metadata.py` to add metadata to `model.onnx`.

Please see <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/tele-speech/run.sh> for details.

Step 3: Obtain tokens.txt

Please also see <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/tele-speech/add-metadata.py>

Models

sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04 ()

Hint:

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪ telespeech-ctc-int8-zh-2024-06-04.tar.bz2

# For Chinese users, please use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪ telespeech-ctc-int8-zh-2024-06-04.tar.bz2

tar xvf sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04.tar.bz2
rm sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 325M Jun 4 11:56 model.int8.onnx
```

Decode wave files

Hint: It supports decoding only wave files of a single channel with 16-bit encoded samples, while the sampling rate does not need to be 16 kHz.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--telespeech-ctc=./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/model.int8.onnx \
--tokens=./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/tokens.txt \
--model-type=telespeech_ctc \
--num-threads=1 \
./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/3-sichuan.wav \
./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/4-tianjin.wav \
./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/5-henan.wav
```

Note: Please use ./build/bin/Release/sherpa-onnx-offline.exe for Windows.

Caution: If you use Windows and get encoding issues, please run:

CHCP 65001

in your commandline.

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↪ build/bin/sherpa-onnx-offline --telespeech-ctc=./sherpa-onnx-telespeech-ctc-int8-zh-
↪ 2024-06-04/model.int8.onnx --tokens=./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/
↪ tokens.txt --model-type=telespeech_ctc --num-threads=1 ./sherpa-onnx-telespeech-ctc-
↪ int8-zh-2024-06-04/test_wavs/3-sichuan.wav ./sherpa-onnx-telespeech-ctc-int8-zh-2024-
↪ 06-04/test_wavs/4-tianjin.wav ./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_
↪ wavs/5-henan.wav
```

(continues on next page)

(continued from previous page)

```

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
    ↴dim=80, low_freq=20, high_freq=-400, dither=0), model_
    ↴config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", ↴
    ↴decoder_filename="", joiner_filename=""), ↴
    ↴paraformer=OfflineParaformerModelConfig(model=""), nemo_
    ↴ctc=OfflineNemoEncDecCtcModelConfig(model=""), ↴
    ↴whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe"
    ↴", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
    ↴ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
    ↴ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc=".sherpa-onnx-telespeech-ctc-
    ↴int8-zh-2024-06-04/model.int8.onnx", tokens=".sherpa-onnx-telespeech-ctc-int8-zh-2024-
    ↴06-04/tokens.txt", num_threads=1, debug=False, provider="cpu", model_type="telespeech_
    ↴ctc", modeling_unit="cjkchar", bpe_vocab=""), lm_config=OfflineLMConfig(model="",
    ↴scale=0.5), ctc_fst_decoder_config=OfflineCtcFstDecoderConfig(graph="", max_
    ↴active=3000), decoding_method="greedy_search", max_active_paths=4, hotwords_file="",
    ↴hotwords_score=1.5, blank_penalty=0)
Creating recognizer ...
Started
Done!

./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/3-sichuan.wav
>{"text": "", "timestamps": [0.08, 0.36, 0.52, 0.72, 0.92, 1.16, 1.36, 1.88, 2.20, 2.36,
    ↴3.16, 3.28, 3.40, 3.60, 3.80, 3.92, 4.08, 4.24, 4.40, 4.56, 4.76, 5.16, 5.32, 5.44, 5.
    ↴64, 5.76, 5.88, 6.04, 6.16, 6.28, 6.40, 6.60, 6.88, 7.12, 7.40, 7.52, 7.64], "tokens": [
    ↴"", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],
    ↴[], ""]
----  

./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/4-tianjin.wav
>{"text": "", "timestamps": [0.36, 0.56, 1.04, 1.16, 1.24, 1.64, 1.88, 2.24, 2.40, 2.60,
    ↴2.80, 3.12, 3.32, 3.64, 3.80, 3.96, 4.16, 4.44, 4.68, 4.80, 5.00, 5.16, 5.28, 6.12, 6.
    ↴28, 6.44, 6.60, 6.72, 6.88, 7.04, 7.12, 7.32, 7.52], "tokens": [
    ↴[], "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],
    ↴[], ""]
----  

./sherpa-onnx-telespeech-ctc-int8-zh-2024-06-04/test_wavs/5-henan.wav
>{"text": "", "timestamps": [0.04, 0.12, 0.24, 0.40, 1.00, 1.24, 1.44, 1.68, 2.32, 2.48,
    ↴2.60, 2.64, 2.80, 3.00, 3.16, 3.32, 3.52, 3.68, 3.92, 5.00, 5.16, 5.28, 5.32, 5.44, 5.
    ↴84, 6.00, 6.12, 6.48, 6.68, 6.84, 7.00, 7.16, 7.32, 7.56, 7.68], "tokens": [
    ↴[], "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", "", ""],
    ↴[], ""]
----  

num threads: 1
decoding method: greedy_search
Elapsed seconds: 3.406 s
Real time factor (RTF): 3.406 / 23.634 = 0.144

```

Note: The feature_dim=80 is incorrect in the above logs. The actual value is 40.

Hint: There is also a float32 model. Please see <https://github.com/k2-fsa/sherpa-onnx/releases/download/>

asr-models/sherpa-onnx-telespeech-ctc-zh-2024-06-04.tar.bz2

8.25.8 Whisper

This section describes how to use models from [Whisper](#) with [sherpa-onnx](#) for non-streaming speech recognition.

Export Whisper to ONNX

This section describes how to export [Whisper](#) models to [onnx](#).

Available models

Note that we have already exported [Whisper](#) models to [onnx](#) and they are available from the following [huggingface](#) repositories:

Model type	Huggingface repo
<code>tiny.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny.en
<code>base.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-base.en
<code>small.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-small.en
<code>distil-small.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-distil-small.en
<code>medium.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-medium.en
<code>distil-medium.en</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-distil-medium.en
<code>tiny</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny
<code>base</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-base
<code>small</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-small
<code>medium</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-medium
<code>large</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-large
<code>large-v1</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-large-v1
<code>large-v2</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-large-v2
<code>large-v3</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-large-v3
<code>distil-large-v2</code>	https://huggingface.co/csukuangfj/sherpa-onnx-whisper-distil-large-v2

Note: You can also download them from

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/asr-models>

Models end with `.en` support only English and all other models are multilingual.

If you want to export the models by yourself or/and want to learn how the models are exported, please read below.

Export to onnx

We use

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/whisper/export-onnx.py>

to export Whisper models to onnx.

First, let us install dependencies and download the export script

```
pip install torch openai-whisper onnxruntime onnx librosa soundfile

git clone https://github.com/k2-fsa/sherpa-onnx/
cd sherpa-onnx/scripts/whisper
python3 ./export-onnx.py --help
```

It will print the following message:

```
usage: export-onnx.py [-h] --model
                      {tiny,tiny.en,base,base.en,small,small.en,medium,medium.en,large,
                     large-v1,large-v2,large-v3,distil-medium.en,distil-small.en,distil-large-v2,medium-
                     aishell}

optional arguments:
  -h, --help            show this help message and exit
  --model {tiny,tiny.en,base,base.en,small,small.en,medium,medium.en,large,large-v1,
           large-v2,large-v3,distil-medium.en,distil-small.en,distil-large-v2,medium-aishell}
```

Example 1: Export tiny.en

To export `tiny.en`, we can use:

```
python3 ./export-onnx.py --model tiny.en
```

It will generate the following files:

```
(py38) fangjuns-MacBook-Pro:whisper fangjun$ ls -lh tiny.en-*
-rw-r--r--  1 fangjun  staff   105M Aug  7 15:43 tiny.en-decoder.int8.onnx
-rw-r--r--  1 fangjun  staff   185M Aug  7 15:43 tiny.en-decoder.onnx
-rw-r--r--  1 fangjun  staff    12M Aug  7 15:43 tiny.en-encoder.int8.onnx
-rw-r--r--  1 fangjun  staff    36M Aug  7 15:43 tiny.en-encoder.onnx
-rw-r--r--  1 fangjun  staff   816K Aug  7 15:43 tiny.en-tokens.txt
```

`tiny.en-encoder.onnx` is the encoder model and `tiny.en-decoder.onnx` is the decoder model.

`tiny.en-encoder.int8.onnx` is the quantized encoder model and `tiny.en-decoder.onnx` is the quantized decoder model.

`tiny.en-tokens.txt` contains the token table, which maps an integer to a token and vice versa.

To check whether the exported model works correctly, we can use

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/whisper/test.py>

We use https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny.en/resolve/main/test_wavs/0.wav as the test wave.

```

pip install kaldinative-fbank
wget https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny.en/resolve/main/test_
↪wavs/0.wav

python3 ./test.py \
--encoder ./tiny.en-encoder.onnx \
--decoder ./tiny.en-decoder.onnx \
--tokens ./tiny.en-tokens.txt \
./0.wav

```

To test int8 quantized models, we can use:

```

python3 ./test.py \
--encoder ./tiny.en-encoder.int8.onnx \
--decoder ./tiny.en-decoder.int8.onnx \
--tokens ./tiny.en-tokens.txt \
./0.wav

```

Example 2: Export large-v3

To export large-v3, we can use:

```
python3 ./export-onnx.py --model large-v3
```

It will generate the following files:

```
(py38) fangjuns-MacBook-Pro:whisper fangjun$ ls -lh large-v3-*
-rw-r--r-- 1 fangjun staff 2.7M Jul 12 20:38 large-v3-decoder.onnx
-rw-r--r-- 1 fangjun staff 3.0G Jul 12 20:38 large-v3-decoder.weights
-rw-r--r-- 1 fangjun staff 744K Jul 12 20:35 large-v3-encoder.onnx
-rw-r--r-- 1 fangjun staff 2.8G Jul 12 20:35 large-v3-encoder.weights
-rw-r--r-- 1 fangjun staff 798K Jul 12 20:32 large-v3-tokens.txt
```

large-v3-encoder.onnx is the encoder model and large-v3-decoder.onnx is the decoder model.

Note that for large models, there will also be two additional weights files.

large-v3-tokens.txt contains the token table, which maps an integer to a token and vice versa.

To check whether the exported model works correctly, we can use

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/whisper/test.py>

We use https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny.en/resolve/main/test_wavs/0.wav as the test wave.

```

pip install kaldinative-fbank
wget https://huggingface.co/csukuangfj/sherpa-onnx-whisper-tiny.en/resolve/main/test_
↪wavs/0.wav

python3 ./test.py \
--encoder ./large-v3-encoder.onnx \
--decoder ./large-v3-decoder.onnx \
--tokens ./large-v3-tokens.txt \
./0.wav

```

Hint: We provide a colab notebook for you to try the exported large-v3 onnx model with sherpa-onnx on CPU as well as on GPU.

You will find the RTF on GPU (Tesla T4) is less than 1.

tiny.en

You can use the following commands to download the exported `onnx` models of `tiny.en`:

Hint: Please replace `tiny.en` with `base.en`, `small.en`, `medium.en`, `distil-small.en`, `tiny`, `base`, `small`, and `medium` if you want to try a different type of model.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
˓→whisper-tiny.en.tar.bz2
tar xvf sherpa-onnx-whisper-tiny.en.tar.bz2
```

Please check that the file sizes of the downloaded models are correct. See the file size of `*.onnx` files below.

```
(py38) fangjuns-MacBook-Pro:sherpa-onnx-whisper-tiny.en fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 105M Aug  7 16:22 tiny.en-decoder.int8.onnx
-rw-r--r-- 1 fangjun staff 185M Aug  7 16:23 tiny.en-decoder.onnx
-rw-r--r-- 1 fangjun staff 12M Aug  7 16:22 tiny.en-encoder.int8.onnx
-rw-r--r-- 1 fangjun staff 36M Aug  7 16:22 tiny.en-encoder.onnx
```

To use the downloaded files to decode waves, please run:

Hint: Please first follow [Installation](#) to build `sherpa-onnx` before you continue.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--whisper-encoder=./sherpa-onnx-whisper-tiny.en/tiny.en-encoder.onnx \
--whisper-decoder=./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.onnx \
--tokens=./sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt \
./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
```

To use `int8` quantized models, please use:

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline \
--whisper-encoder=./sherpa-onnx-whisper-tiny.en/tiny.en-encoder.int8.onnx \
--whisper-decoder=./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.int8.onnx \
--tokens=./sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt \
./sherpa-onnx-whisper-tiny.en/test_wavs/0.wav \
```

(continues on next page)

(continued from previous page)

```
./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav \
./sherpa-onnx-whisper-tiny.en/test_wavs/8k.wav
```

Real-time factor (RTF) on Raspberry Pi 4 Model B

One of the test command is given below:

```
./sherpa-onnx-offline \
--num-threads=1 \
--whisper-encoder=./sherpa-onnx-whisper-tiny.en/tiny.en-encoder.onnx \
--whisper-decoder=./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.onnx \
--tokens=./sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt \
./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav
```

And its output is:

```
/root/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:361 ./
↳ sherpa-onnx-offline --num-threads=1 --whisper-encoder=./sherpa-onnx-whisper-tiny.en/
↳ tiny.en-encoder.onnx --whisper-decoder=./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.
↳ onnx --tokens=./sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt ./sherpa-onnx-whisper-
↳ tiny.en/test_wavs/1.wav

OfflineRecognizerConfig(feat_config=OfflineFeatureExtractorConfig(sampling_rate=16000,
↳ feature_dim=80), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",
↳ decoder_filename="", joiner_filename=""),
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""),
↳ whisper=OfflineWhisperModelConfig(encoder="./sherpa-onnx-whisper-tiny.en/tiny.en-
↳ encoder.onnx", decoder="./sherpa-onnx-whisper-tiny.en/tiny.en-decoder.onnx"), tokens="",
↳ /sherpa-onnx-whisper-tiny.en/tiny.en-tokens.txt", num_threads=1, debug=False, provider=
↳ "cpu", model_type=""), lm_config=OfflineLMConfig(model="", scale=0.5), decoding_method=
↳ "greedy_search", max_active_paths=4, context_score=1.5)
Creating recognizer ...
Started
Done!

./sherpa-onnx-whisper-tiny.en/test_wavs/1.wav
{"text":" God, as a direct consequence of the sin which man thus punished, had given her,
↳ a lovely child, whose place was on that same dishonored bosom to connect her parent,
↳ forever with the race and descent of mortals, and to be finally a blessed soul in
↳ heaven.", "timestamps": "[]", "tokens": [" God", "", " as", " a", " direct", " consequence", " "
↳ of", " the", " sin", " which", " man", " thus", " punished", "", " had", " given", " her", " a",
↳ " lovely", " child", "", " whose", " place", " was", " on", " that", " same", " dishon", " ored",
↳ " bos", " om", " to", " connect", " her", " parent", " forever", " with", " the", " race", " and",
↳ " descent", " of", " mortals", "", " and", " to", " be", " finally", " a", " blessed", " soul",
↳ " in", " heaven", "."]}

---
num threads: 1
decoding method: greedy_search
Elapsed seconds: 11.454 s
```

(continues on next page)

(continued from previous page)

Real time factor (RTF): 11.454 / 16.715 = 0.685

The following table compares the RTF between different number of threads and types of `onnx` models:

Model type	Number of threads	RTF
float32	1	0.685
float32	2	0.559
float32	3	0.526
float32	4	0.520
int8	1	0.547
int8	2	0.431
int8	3	0.398
int8	4	0.386

large-v3

Before we start, let us follow [Linux](#) to build a CUDA-enabled version of `sherpa-onnx`.

In the following, we assume you have run

```
cd /content

git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

mkdir -p build
cd build
cmake \
-DBUILD_SHARED_LIBS=ON \
-DSHERPA_ONNX_ENABLE_GPU=ON ..

make -j2 sherpa-onnx-offline
```

You can use the following commands to download the exported `onnx` models of `large-v3`:

Hint: Please replace `large-v3` with `large`, `large-v1`, `large-v2`, and `distil-large-v2` if you want to try a different type of model.

```
cd /content

git lfs install
git clone https://huggingface.co/csukuangfj/sherpa-onnx-whisper-large-v3

ls -lh sherpa-onnx-whisper-large-v3
```

The logs of the above commands are given below:

```
Git LFS initialized.
Cloning into 'sherpa-onnx-whisper-large-v3'...
remote: Enumerating objects: 26, done.
```

(continues on next page)

(continued from previous page)

```

remote: Counting objects: 100% (22/22), done.
remote: Compressing objects: 100% (21/21), done.
remote: Total 26 (delta 2), reused 0 (delta 0), pack-reused 4 (from 1)
Unpacking objects: 100% (26/26), 1.00 MiB | 9.10 MiB/s, done.
Filtering content: 100% (6/6), 7.40 GiB | 34.50 MiB/s, done.
total 7.5G
-rw-r--r-- 1 root root 962M Jul 13 14:19 large-v3-decoder.int8.onnx
-rw-r--r-- 1 root root 2.8M Jul 13 14:18 large-v3-decoder.onnx
-rw-r--r-- 1 root root 3.0G Jul 13 14:22 large-v3-decoder.weights
-rw-r--r-- 1 root root 732M Jul 13 14:19 large-v3-encoder.int8.onnx
-rw-r--r-- 1 root root 745K Jul 13 14:18 large-v3-encoder.onnx
-rw-r--r-- 1 root root 2.8G Jul 13 14:21 large-v3-encoder.weights
-rw-r--r-- 1 root root 798K Jul 13 14:18 large-v3-tokens.txt
drwxr-xr-x 2 root root 4.0K Jul 13 14:18 test_wavs

```

Caution: Please remember to run `git lfs install` before you run `git clone`. If you have any issues about `git lfs install`, please follow <https://git-lfs.com/> to install git-lfs.

Caution: Please check the file sizes are correct before proceeding. Otherwise, you would be SAD later.

Caution: Please check the file sizes are correct before proceeding. Otherwise, you would be SAD later.

Caution: Please check the file sizes are correct before proceeding. Otherwise, you would be SAD later.

Run with CPU (float32)

```

cd /content

exe=$PWD/sherpa-onnx/build/bin/sherpa-onnx-offline

cd sherpa-onnx-whisper-large-v3

time $exe \
    --whisper-encoder=./large-v3-encoder.onnx \
    --whisper-decoder=./large-v3-decoder.onnx \
    --tokens=./large-v3-tokens.txt \
    --num-threads=2 \
    ./test_wavs/0.wav

```

The logs are given below:

```

/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 /content/sherpa-onnx/
↳build/bin/sherpa-onnx-offline --whisper-encoder=./large-v3-encoder.onnx --whisper-
↳decoder=./large-v3-decoder.onnx --tokens=./large-v3-tokens.txt --num-threads=2 ./test_
↳wavs/0.wav

```

(continues on next page)

(continued from previous page)

```

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
    ↴dim=80, low_freq=20, high_freq=-400, dither=0), model_
    ↴config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", ↴
    ↴decoder_filename="", joiner_filename=""), ↴
    ↴paraformer=OfflineParaformerModelConfig(model=""), nemo_
    ↴ctc=OfflineNemoEncDecCtcModelConfig(model=""), ↴
    ↴whisper=OfflineWhisperModelConfig(encoder="./large-v3-encoder.onnx", decoder="./large-
    ↴v3-decoder.onnx", language="", task="transcribe", tail_paddings=-1), ↴
    ↴tdnn=OfflineTdnnModelConfig(model=""), zipformer_
    ↴ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
    ↴ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc="", tokens="./large-v3-tokens.
    ↴txt", num_threads=2, debug=False, provider="cpu", model_type="", modeling_unit="cjkchar
    ↴", bpe_vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
    ↴config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
    ↴search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0, ↴
    ↴rule fstts="", rule_fars="")
Creating recognizer ...
Started
Done!

./test_wavs/0.wav
{"text": " after early nightfall the yellow lamps would light up here and there the_
    ↴squalid quarter of the brothels", "timestamps": [], "tokens": ["after", "early", "_
    ↴night", "fall", "the", "yellow", "lamps", "would", "light", "up", "here", "and
    ↴", "there", "the", "squ", "alid", "quarter", "of", "the", "broth", "els"], ↴
    ↴"words": []}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 54.070 s
Real time factor (RTF): 54.070 / 6.625 = 8.162

real      1m32.107s
user      1m39.877s
sys  0m10.405s

```

Run with CPU (int8)

```

cd /content

exe=$PWD/sherpa-onnx/build/bin/sherpa-onnx-offline

cd sherpa-onnx-whisper-large-v3

time $exe \
    --whisper-encoder=./large-v3-encoder.int8.onnx \
    --whisper-decoder=./large-v3-decoder.int8.onnx \
    --tokens=./large-v3-tokens.txt \
    --num-threads=2 \

```

(continues on next page)

(continued from previous page)

```
./test_wavs/0.wav
```

The logs are given below:

```
/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 /content/sherpa-onnx/
↳ build/bin/sherpa-onnx-offline --whisper-encoder=./large-v3-encoder.int8.onnx --whisper-
↳ decoder=./large-v3-decoder.int8.onnx --tokens=./large-v3-tokens.txt --num-threads=2 ./
↳ test_wavs/0.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
↳ dim=80, low_freq=20, high_freq=400, dither=0), model_
↳ config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",_
↳ decoder_filename="", joiner_filename=""),_
↳ paraformer=OfflineParaformerModelConfig(model=""), nemo_
↳ ctc=OfflineNemoEncDecCtcModelConfig(model=""),_
↳ whisper=OfflineWhisperModelConfig(encoder="./large-v3-encoder.int8.onnx", decoder="./
↳ large-v3-decoder.int8.onnx", language="", task="transcribe", tail_paddings=-1),_
↳ tdnn=OfflineTdnmModelConfig(model=""), zipformer_
↳ ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
↳ ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc="", tokens="./large-v3-tokens.
↳ txt", num_threads=2, debug=False, provider="cpu", model_type="", modeling_unit="cjkchar
↳ ", bpe_vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
↳ config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
↳ search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0,_
↳ rule fst="", rule_fars="")
Creating recognizer ...
Started
Done!

./test_wavs/0.wav
{"text": " after early nightfall the yellow lamps would light up here and there the_
↳ squalid quarter of the brothels", "timestamps": [], "tokens": ["after", "early", "_"
↳ night", "fall", "the", "yellow", "lamps", "would", "light", "up", "here", "and"
↳ ", "there", "the", "squ", "alid", "quarter", "of", "the", "broth", "els"],_
↳ "words": []}
---
num threads: 2
decoding method: greedy_search
Elapsed seconds: 49.991 s
Real time factor (RTF): 49.991 / 6.625 = 7.546

real 1m15.555s
user 1m41.488s
sys 0m9.156s
```

Run with GPU (float32)

```
cd /content
exe=$PWD/sherpa-onnx/build/bin/sherpa-onnx-offline

cd sherpa-onnx-whisper-large-v3

time $exe \
--whisper-encoder=./large-v3-encoder.onnx \
--whisper-decoder=./large-v3-decoder.onnx \
--tokens=./large-v3-tokens.txt \
--provider=cuda \
--num-threads=2 \
./test_wavs/0.wav
```

The logs are given below:

```
/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 /content/sherpa-onnx/
build/bin/sherpa-onnx-offline --whisper-encoder=./large-v3-encoder.onnx --whisper-
decoder=./large-v3-decoder.onnx --tokens=./large-v3-tokens.txt --provider=cuda --num-
threads=2 ./test_wavs/0.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80, low_freq=20, high_freq=-400, dither=0), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", u
decoder_filename "", joiner_filename ""), u
paraformer=OfflineParaformerModelConfig(model=""), nemo_
ctc=OfflineNemoEncDecCtcModelConfig(model=""), u
whisper=OfflineWhisperModelConfig(encoder="./large-v3-encoder.onnx", decoder="./large-
v3-decoder.onnx", language="", task="transcribe", tail_paddings=-1), u
tdnn=OfflineTdnndModelConfig(model=""), zipformer_
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc="", tokens="./large-v3-tokens.
txt", num_threads=2, debug=False, provider="cuda", model_type="", modeling_unit=
"cjkchar", bpe_vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_
decoder_config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method=
"greedy_search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_
penalty=0, rule_fsts="", rule_fars="")
Creating recognizer ...
Started
Done!

./test_wavs/0.wav
{"text": " after early nightfall the yellow lamps would light up here and there the_
squallid quarter of the brothels", "timestamps": [], "tokens": ["after", "early", "n
ight", "fall", "the", "yellow", "lamps", "would", "light", "up", "here", "and", "there", "the", "squ
allid", "quarter", "of", "the", "broth", "els"], "words": []}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 5.910 s
Real time factor (RTF): 5.910 / 6.625 = 0.892
```

(continues on next page)

(continued from previous page)

```
real 0m26.996s
user 0m12.854s
sys 0m4.486s
```

Note: The above command is run within a colab notebook using Tesla T4 GPU. You can see the RTF is less than 1. If you have some more performant GPU, you would get an even lower RTF.

Run with GPU (int8)

```
cd /content
exe=$PWD/sherpa-onnx/build/bin/sherpa-onnx-offline

cd sherpa-onnx-whisper-large-v3

time $exe \
--whisper-encoder=./large-v3-encoder.int8.onnx \
--whisper-decoder=./large-v3-decoder.int8.onnx \
--tokens=./large-v3-tokens.txt \
--provider=cuda \
--num-threads=2 \
./test_wavs/0.wav
```

The logs are given below:

```
/content/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 /content/sherpa-onnx/
build/bin/sherpa-onnx-offline --whisper-encoder=./large-v3-encoder.int8.onnx --whisper-
decoder=./large-v3-decoder.int8.onnx --tokens=./large-v3-tokens.txt --provider=cuda --
num-threads=2 ./test_wavs/0.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
dim=80, low_freq=20, high_freq=400, dither=0), model_
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", 
decoder_filename="", joiner_filename=""), 
paraformer=OfflineParaformerModelConfig(model=""), nemo_
ctc=OfflineNemoEncDecCtcModelConfig(model=""), 
whisper=OfflineWhisperModelConfig(encoder="./large-v3-encoder.int8.onnx", decoder="./
large-v3-decoder.int8.onnx", language="", task="transcribe", tail_paddings=-1), 
tdnn=OfflineTdnModelConfig(model=""), zipformer_
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
ctc=OfflineWenetCtcModelConfig(model=""), telespeech_ctc="", tokens="./large-v3-tokens.
txt", num_threads=2, debug=False, provider="cuda", model_type="", modeling_unit=
"cjkchar", bpe_vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_
decoder_config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method=
"greedy_search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_
penalty=0, rule_fsts="", rule_fars="")
Creating recognizer ...
Started
```

(continues on next page)

(continued from previous page)

Done!

```
./test_wavs/0.wav
>{"text": " after early nightfall the yellow lamps would light up here and there the
↳ squalid quarter of the brothels", "timestamps": [], "tokens":[" after", " early", "_
↳ night", " fall", " the", " yellow", " lamps", " would", " light", " up", " here", " and
↳ ", " there", " the", " squ", " alid", " quarter", " of", " the", " broth", " els"],_
↳ "words": []}
-----
num threads: 2
decoding method: greedy_search
Elapsed seconds: 19.190 s
Real time factor (RTF): 19.190 / 6.625 = 2.897

real  0m46.850s
user   0m50.007s
sys    0m8.013s
```

Fix issues about running on GPU

If you get errors like below:

```
what(): /onnxruntime_src/onnxruntime/core/session/provider_bridge_ort.cc:1426
onnxruntime::Provider& onnxruntime::ProviderLibrary::Get()
[ONNXRuntimeError] : 1 : FAIL :
Failed to load library libonnxruntime_providers_cuda.so with error:
libcublasLt.so.11: cannot open shared object file: No such file or directory
```

please follow <https://www.google.com/url?q=https%3A%2F%2Fk2-fsa.github.io%2Fk2%2Finstallation%2Fcuda-cudnn.html> to install CUDA toolkit.

To determine which version of CUDA toolkit to install, please read <https://onnxruntime.ai/docs/execution-providers/CUDA-ExecutionProvider.html> to figure it out.

For instance, if onnxruntime v1.18.1 is used in `sherpa-onnx`, we have to install CUDA 11.8 according to <https://onnxruntime.ai/docs/execution-providers/CUDA-ExecutionProvider.html>

colab

Please see the following colab notebook .

It walks you step by step to try the exported large-v3 onnx model with `sherpa-onnx` on CPU as well as on GPU.

colab

Non-large models

We provide a colab notebook for you to try `Whisper` models with `sherpa-onnx` step by step.

The screenshot shows a GitHub repository page for the 'sherpa-onnx' project. The URL in the address bar is `github.com/k2-fsa/colab/blob/master/sherpa-onnx/sherpa_onnx_whisper_models.ipynb`. The repository name is 'k2-fsa / colab'. The main navigation menu includes Code, Issues, Pull requests, Actions, Projects, Wiki, Security, Insights, and Settings. A dropdown menu shows 'master'. Below the menu, there's a commit history entry from 'csukuangfj' with the message 'Fix typos'. At the bottom, there are buttons for Preview, Code, Blame, and a file size of 2240 lines (2240 loc) · 147 KB. A prominent 'Open in Colab' button is located at the bottom left of the preview area.

Introduction

This colab notebook shows how to use `sherpa-onnx` to run `whisper` models.

Real time factors (RTF) of each type of model are listed below:

model	CPU or CUDA	RTF
float32 tiny.en	CPU	0.208
float32 tiny.en	CUDA	0.186
int8 tiny.en	CPU	0.142
int8 tiny.en	CUDA	0.146
float32 base.en	CPU	0.430
float32 base.en	CUDA	0.129
int8 base.en	CPU	0.332
int8 base.en	CUDA	0.259
float32 small.en	CPU	1.633
float32 small.en	CUDA	0.268

Large models

For large models of whisper, please see the following colab notebook . It walks you step by step to try the exported large-v3 onnx model with `sherpa-onnx` on CPU as well as on GPU.

You will find the RTF on GPU (Tesla T4) is less than 1.

Huggingface space

You can try `Whisper` models from within your browser without installing anything.

Please visit

<https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>

The screenshot shows a web browser interface for a Huggingface space. At the top, the URL bar contains the address: `https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition`. Below the URL, the page title is "Automatic Speech Recognition with Next-gen Kaldi". The main form has several sections:

- Language:** A row of radio buttons with "English" selected (circled in red). Other options include Chinese, Chinese+English, Tibetan, Arabic, German, and French.
- Select a model:** A dropdown menu showing "whisper-base.en" (circled in red), with a large red text overlay "Please select whisper" placed over it.
- Decoding method:** Radio buttons for "greedy_search" (selected) and "modified_beam_search".
- Number of active paths for modified_beam_search:** A horizontal slider with a blue handle.
- Upload from disk:** A button to upload files, showing a progress bar at 0:00 / 0:06. Other options are "Record from microphone" and "From URL".
- Submit for recognition:** A button to start the recognition process.
- Recognized speech from uploaded file:** A text area displaying the transcription: "after early nightfall the yellow lamps would light up here and there the squalid quarter of the brothels."

8.25.9 WeNet

This page lists all CTC models from [WeNet](#).

How to export models from WeNet to sherpa-onnx

Suppose you have the following files from [WeNet](#):

- `final.pt`
- `train.yaml`
- `global_cmvn`
- `units.txt`

We describe below how to use scripts from [sherpa-onnx](#) to export your files.

Hint: Both streaming and non-streaming models are supported.

Export for non-streaming inference

You can use the following script

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/wenet/export-onnx.py>

to export your model to [sherpa-onnx](#). After running it, you should get two files:

- `model.onnx`
- `model.int8.onnx`.

Next, we rename `units.txt` to `tokens.txt` to follow the convention used in [sherpa-onnx](#):

```
mv units.txt tokens.txt
```

Now you can use the following command for speech recognition with the exported models:

```
# with float32 models
./build/bin/sherpa-onnx-offline \
--wenet-ctc-model=./model.onnx \
--tokens=./tokens.txt \
/path/to/some.wav

# with int8 models
./build/bin/sherpa-onnx-offline \
--wenet-ctc-model=./model.int8.onnx \
--tokens=./tokens.txt \
/path/to/some.wav
```

Export for streaming inference

You can use the following script

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/wenet/export-onnx-streaming.py>

to export your model to `sherpa-onnx`. After running it, you should get two files:

- `model-streaming.onnx`
- `model-streaming.int8.onnx`.

Next, we rename `units.txt` to `tokens.txt` to follow the convention used in `sherpa-onnx`:

```
mv units.txt tokens.txt
```

Now you can use the following command for speech recognition with the exported models:

```
# with float32 models
./build/bin/sherpa-onnx \
--wenet-ctc-model=./model-streaming.onnx
--tokens=./tokens.txt \
/path/to/some.wav

# with int8 models
./build/bin/sherpa-onnx \
--wenet-ctc-model=./model-streaming.int8.onnx
--tokens=./tokens.txt \
/path/to/some.wav
```

FAQs

sherpa-onnx/csrc/online-wenet-ctc-model.cc:Init:144 head does not exist in the metadata

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/online-wenet-ctc-model.cc:Init:144 head does not exist in the metadata
```

To fix the above error, please check the following two items:

- Make sure you are using `model-streaming.onnx` or `model-streaming.int8.onnx`. The executable you are running requires a streaming model as input.
- Make sure you use the script from `sherpa-onnx` to export your model.

All models from WeNet

https://github.com/wenet-e2e/wenet/blob/main/docs/pretrained_models.en.md lists all pre-trained models from WeNet and we have converted all of them to `sherpa-onnx` using the following script:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/wenet/run.sh>.

We have uploaded the exported models to huggingface and you can find them from the following figure:

To make it easier to copy the links, we list them below:

The pretrained model in WeNet follows the license of it's corresponding dataset. For example, the pretrained model on LibriSpeech follows CC BY 4.0 , since it is used as license of the LibriSpeech dataset, see <http://openslr.org/12/>.

Model List

Here is a list of the pretrained models on different datasets.

For non-Chinese users, please visit [Pretrained Models\(En\)](#) to download.

Datasets	Languages	Checkpoint Model	Runtime Model	Contributor	exported for sherpa-onnx
aishell	CN	Conformer	Conformer	 出门问问	https://huggingface.co/csukuangfj/-/sherpa-onnx-zh-wenet-aishell
aishell2	CN	Conformer	Conformer	 出门问问	https://huggingface.co/csukuangfj/sherpa-onnx-zh-wenet-aishell2
gigaspeech	EN	Conformer	Conformer	 出门问问	https://huggingface.co/csukuangfj/-/sherpa-onnx-en-wenet-gigaspeech
librispeech	EN	Conformer	Conformer	 出门问问	https://huggingface.co/csukuangfj/-/sherpa-onnx-en-wenet-librispeech
multi_cn	CN	Conformer	Conformer	 京东	https://huggingface.co/csukuangfj/-/sherpa-onnx-zh-wenet-multi-cn
wenetspeech	CN	Conformer	Conformer	 地平线	https://huggingface.co/csukuangfj/-/sherpa-onnx-zh-wenet-wenetspeech

Fig. 8.32: All pre-trained models from WeNet.

- <https://huggingface.co/csukuangfj/sherpa-onnx-zh-wenet-aishell>
- <https://huggingface.co/csukuangfj/sherpa-onnx-zh-wenet-aishell2>
- <https://huggingface.co/csukuangfj/sherpa-onnx-en-wenet-gigaspeech>
- <https://huggingface.co/csukuangfj/sherpa-onnx-en-wenet-librispeech>
- <https://huggingface.co/csukuangfj/sherpa-onnx-zh-wenet-multi-cn>
- <https://huggingface.co/csukuangfj/sherpa-onnx-zh-wenet-wenetspeech>

Colab

We provide a colab notebook for you to try the exported WeNet models with sherpa-onnx.

8.25.10 Small models

In this section, we list online/streaming models with fewer parameters that are suitable for resource constrained embedded systems.

Hint: You can use them as a first pass model in a two-pass system, where the second pass uses a non-streaming model.

Hint: If you are using Raspberry Pi 4, this section is not so helpful for you since all models in sherpa-onnx are able to run in real-time on it.

This page is especially useful for systems with less resource than Raspberry Pi 4.

- *csukuangfj/sherpa-onnx-streaming-zipformer-zh-14M-2023-02-23 (Chinese)*
- *csukuangfj/sherpa-onnx-streaming-zipformer-en-20M-2023-02-17 (English)*
- *sherpa-onnx-streaming-zipformer-small-bilingual-zh-en-2023-02-16 (Bilingual, Chinese + English)*

8.26 SenseVoice

This section describes how to use models from <https://github.com/FunAudioLLM/SenseVoice>.

A single model from SenseVoice supports the following languages

- Chinese (Mandarin,)
- Cantonese (,)
- English
- Japanese
- Korean

which is similar to what multilingual Whisper is doing.

We have converted SenseVoice to onnx and provided APIs for the following programming languages

- 1. C++
- 2. C

- 3. Python
- 4. C#
- 5. Go
- 6. Kotlin
- 7. Java
- 8. JavaScript (Support WebAssembly and Node)
- 9. Swift
- 10. Dart (Support Flutter)

Note that you can use SenseVoice with `sherpa-onnx` on the following platforms:

- Linux (x64, aarch64, arm, riscv64)
- macOS (x64, arm64)
- Windows (x64, x86, arm64)
- Android (arm64-v8a, armv7-eabi, x86, x86_64)
- iOS (arm64)

In the following, we describe how to download pre-trained SenseVoice models and use them in `sherpa-onnx`.

8.26.1 Huggingface space

You can try SenseVoice with `sherpa-onnx` with the following huggingface space

<https://huggingface.co/spaces/k2-fsa/automatic-speech-recognition>

Hint: You don't need to install anything. All you need is a browser.

You can even run it on your phone or tablet.

8.26.2 Export SenseVoice to sherpa-onnx

This page describes how to export SenseVoice to onnx so that you can use it with `sherpa-onnx`.

The code

Please refer to `export-onnx.py`

The entry point is `run.sh`

After executing `run.sh`, you should get the following files

- `model.onnx`, the float32 onnx model
- `model.int8.onnx`, the 8-bit quantized model
- `tokens.txt`, for converting integer token IDs to strings
- `test_wavs/zh.wav`, test wave for Chinese
- `test_wavs/en.wav`, test wave for English

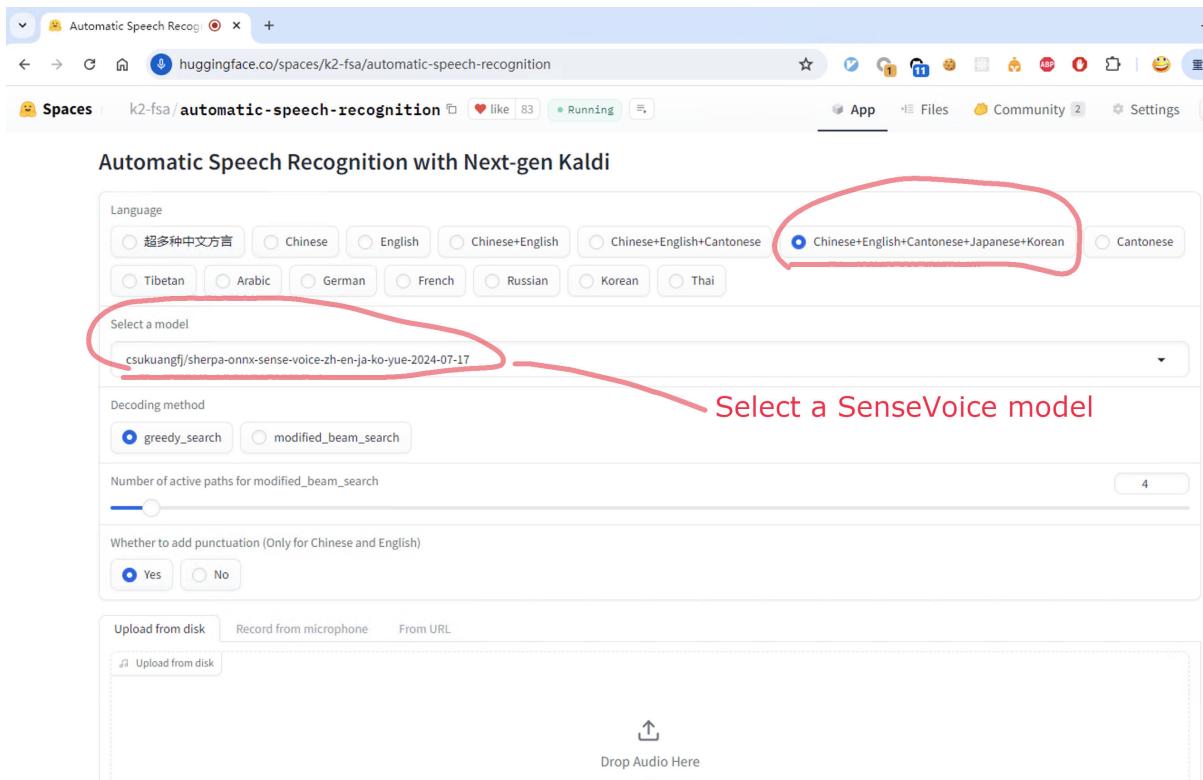


Fig. 8.33: Try SenseVoice in our Huggingface space with sherpa-onnx

- `test_wavs/ko.wav`, test wave for Korean
- `test_wavs/ja.wav`, test wave for Japanese
- `test_wavs/yue.wav`, test wave for Cantonese

Test the exported model

You can use `test.py` to test the exported model.

Note that `test.py` does not depend on `sherpa-onnx`. It uses onnxruntime Python API.

Where to find exported models

You can find the exported SenseVoice models at

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/asr-models>

The following is an example about how to download an exported SenseVoice model:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuua.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
```

(continues on next page)

(continued from previous page)

```
tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
```

To view the downloaded files, please use:

```
ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17

total 1.1G
-rw-r--r-- 1 runner docker 71 Jul 18 13:06 LICENSE
-rw-r--r-- 1 runner docker 104 Jul 18 13:06 README.md
-rwxr-xr-x 1 runner docker 5.8K Jul 18 13:06 export-onnx.py
-rw-r--r-- 1 runner docker 229M Jul 18 13:06 model.int8.onnx
-rw-r--r-- 1 runner docker 895M Jul 18 13:06 model.onnx
drwxr-xr-x 2 runner docker 4.0K Jul 18 13:06 test_wavs
-rw-r--r-- 1 runner docker 309K Jul 18 13:06 tokens.txt

ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs

total 940K
-rw-r--r-- 1 runner docker 224K Jul 18 13:06 en.wav
-rw-r--r-- 1 runner docker 226K Jul 18 13:06 ja.wav
-rw-r--r-- 1 runner docker 145K Jul 18 13:06 ko.wav
-rw-r--r-- 1 runner docker 161K Jul 18 13:06 yue.wav
-rw-r--r-- 1 runner docker 175K Jul 18 13:06 zh.wav
```

8.26.3 Pre-trained Models

This page describes how to download pre-trained SenseVoice models.

sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17

This model is converted from <https://www.modelscope.cn/models/iic/SenseVoiceSmall> using the script `export-onnx.py`.

It supports the following 5 languages:

- Chinese (Mandarin,)
- Cantonese (,)
- English
- Japanese
- Korean

In the following, we describe how to download it.

Download

Please use the following commands to download it:

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

# For Chinese users, you can use the following mirror
# wget https://hub.nuaa.cf/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-onnx-
↪sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
```

After downloading, you should find the following files:

```
ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17

total 1.1G
-rw-r--r-- 1 runner docker 71 Jul 18 13:06 LICENSE
-rw-r--r-- 1 runner docker 104 Jul 18 13:06 README.md
-rwxr-xr-x 1 runner docker 5.8K Jul 18 13:06 export-onnx.py
-rw-r--r-- 1 runner docker 229M Jul 18 13:06 model.int8.onnx
-rw-r--r-- 1 runner docker 895M Jul 18 13:06 model.onnx
drwxr-xr-x 2 runner docker 4.0K Jul 18 13:06 test_wavs
-rw-r--r-- 1 runner docker 309K Jul 18 13:06 tokens.txt

ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs

total 940K
-rw-r--r-- 1 runner docker 224K Jul 18 13:06 en.wav
-rw-r--r-- 1 runner docker 226K Jul 18 13:06 ja.wav
-rw-r--r-- 1 runner docker 145K Jul 18 13:06 ko.wav
-rw-r--r-- 1 runner docker 161K Jul 18 13:06 yue.wav
-rw-r--r-- 1 runner docker 175K Jul 18 13:06 zh.wav
```

Decode a file with `model.onnx`

Without inverse text normalization

To decode a file without inverse text normalization, please use:

```
./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/sense-voice/model.
↪onnx \
--num-threads=1 \
--debug=0 \
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav \
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
```

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 ./  
build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-  
07-17/tokens.txt --sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/  
sense-voice/model.onnx --num-threads=1 --debug=0 ./sherpa-onnx-sense-voice-zh-en-ja-ko-  
yue-2024-07-17/test_wavs/zh.wav ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/  
test_wavs/en.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_  
dim=80, low_freq=20, high_freq=-400, dither=0), model_=  
config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="",  
decoder_filename="", joiner_filename=""),  
paraformer=OfflineParaformerModelConfig(model=""), nemo_=  
ctc=OfflineNemoEncDecCtcModelConfig(model=""),  
whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe",  
tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_=  
ctc=OfflineZipformerCtcModelConfig(model=""), wenet_=  
ctc=OfflineWenetCtcModelConfig(model=""), sense_=  
voice=OfflineSenseVoiceModelConfig(model="/Users/fangjun/open-source/sherpa-onnx/  
scripts/sense-voice/model.onnx", language="auto", use_itn=False), telespeech_ctc="",  
tokens="../sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt", num_=  
threads=1, debug=False, provider="cpu", model_type="", modeling_unit="cjkchar", bpe_=  
vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_=  
config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_"  
search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0,  
rule_fsts="", rule_fars="")
Creating recognizer ...
Started
Done!

./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
{"text": "", "timestamps": [0.72, 0.96, 1.26, 1.44, 1.92, 2.10, 2.58, 2.82, 3.30, 3.90,  
4.20, 4.56, 4.74], "tokens": ["\"", "", "", "", "", "", "", "", "", "", "", ""],  
"words": []}
----  

./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
{"text": "the tribal chieftain called for the boy and presented him with fifty pieces of  
gold", "timestamps": [0.90, 1.26, 1.56, 1.80, 2.16, 2.46, 2.76, 2.94, 3.12, 3.60, 3.96,  
4.50, 4.74, 5.10, 5.52, 5.88, 6.18], "tokens": ["the", " tri", "bal", " chief", "tain",  
" called", " for", " the", " boy", " and", " presented", " him", " with", " fifty", "  
pieces", " of", " gold"], "words": []}
----  

num threads: 1
decoding method: greedy_search
Elapsed seconds: 2.320 s
Real time factor (RTF): 2.320 / 12.744 = 0.182
```

With inverse text normalization

To decode a file with inverse text normalization, please use:

```
./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/sense-voice/model.
˓→onnx \
--num-threads=1 \
--sense-voice-use-itn=1 \
--debug=0 \
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav \
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
```

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 ./
˓→build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-
˓→07-17/tokens.txt --sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/
˓→sense-voice/model.onnx --num-threads=1 --sense-voice-use-itn=1 --debug=0 ./sherpa-onnx-
˓→sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav ./sherpa-onnx-sense-voice-zh-
˓→en-ja-ko-yue-2024-07-17/test_wavs/en.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80, low_freq=20, high_freq=-400, dither=0), model_
˓→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", u
˓→decoder_filename="", joiner_filename=""), u
˓→paraformer=OfflineParaformerModelConfig(model=""), nemo_
˓→ctc=OfflineNemoEncDecCtcModelConfig(model=""), u
˓→whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓→", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
˓→ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓→ctc=OfflineWenetCtcModelConfig(model=""), sense_
˓→voice=OfflineSenseVoiceModelConfig(model="/Users/fangjun/open-source/sherpa-onnx/
˓→scripts/sense-voice/model.onnx", language="auto", use_itn=True), telespeech_ctc="", u
˓→tokens="../sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt", num_
˓→threads=1, debug=False, provider="cpu", model_type="", modeling_unit="cjkchar", bpe_
˓→vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
˓→config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
˓→search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0,
˓→rule_fsts="", rule_fars="")
Creating recognizer ...
Started
Done!

./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
{"text": "95", "timestamps": [0.72, 0.96, 1.26, 1.44, 1.92, 2.10, 2.58, 2.82, 3.30, 3.90,
˓→ 4.20, 4.56, 4.74, 5.46], "tokens": ["", "", "", "", "", "", "9", "", "", "", "5",
˓→ "", ""], "words": []}
---
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
{"text": "The tribal chieftain called for the boy and presented him with 50 pieces of
˓→ gold.", "timestamps": [0.90, 1.26, 1.56, 1.80, 2.16, 2.46, 2.76, 2.94, 3.12, 3.60, 3.
˓→ 96, 4.50, 4.74, 4.92, 5.10, 5.28, 5.52, 5.88, 6.18, 7.02], "tokens": ["The", "tri",
˓→ "bal", "chief", "tain", "called", "for", "the", "boy", "and", "presented", "him
˓→ ", "with", "5", "0", "pieces", "of", "gold", "."], "words": []}
```

(continues on next page)

(continued from previous page)

```
---
num threads: 1
decoding method: greedy_search
Elapsed seconds: 1.543 s
Real time factor (RTF): 1.543 / 12.744 = 0.121
```

Hint: When inverse text normalization is enabled, the results also punctuations.

Specify a language

If you don't provide a language when decoding, it uses auto.

To specify the language when decoding, please use:

```
./build/bin/sherpa-onnx-offline \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/sense-voice/model.
˓→onnx \
--num-threads=1 \
--sense-voice-language=zh \
--debug=0 \
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
```

You should see the following output:

```
/Users/fangjun/open-source/sherpa-onnx/sherpa-onnx/csrc/parse-options.cc:Read:375 ./
˓→build/bin/sherpa-onnx-offline --tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-
˓→07-17/tokens.txt --sense-voice-model=/Users/fangjun/open-source/sherpa-onnx/scripts/
˓→sense-voice/model.onnx --num-threads=1 --sense-voice-language=zh --debug=0 ./sherpa-
˓→onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav

OfflineRecognizerConfig(feat_config=FeatureExtractorConfig(sampling_rate=16000, feature_
˓→dim=80, low_freq=20, high_freq=-400, dither=0), model_
˓→config=OfflineModelConfig(transducer=OfflineTransducerModelConfig(encoder_filename="", u
˓→decoder_filename="", joiner_filename=""), u
˓→paraformer=OfflineParaformerModelConfig(model=""), nemo_
˓→ctc=OfflineNemoEncDecCtcModelConfig(model=""), u
˓→whisper=OfflineWhisperModelConfig(encoder="", decoder="", language="", task="transcribe
˓→", tail_paddings=-1), tdnn=OfflineTdnnModelConfig(model=""), zipformer_
˓→ctc=OfflineZipformerCtcModelConfig(model=""), wenet_
˓→ctc=OfflineWenetCtcModelConfig(model=""), sense_
˓→voice=OfflineSenseVoiceModelConfig(model="/Users/fangjun/open-source/sherpa-onnx/
˓→scripts/sense-voice/model.onnx", language="zh", use_itn=False), telespeech_ctc="", u
˓→tokens=".sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt", num_
˓→threads=1, debug=False, provider="cpu", model_type="", modeling_unit="cjkchar", bpe_
˓→vocab=""), lm_config=OfflineLMConfig(model="", scale=0.5), ctc_fst_decoder_
˓→config=OfflineCtcFstDecoderConfig(graph="", max_active=3000), decoding_method="greedy_
˓→search", max_active_paths=4, hotwords_file="", hotwords_score=1.5, blank_penalty=0, u
˓→rule_ftsts="", rule_fars=""")
Creating recognizer ...
```

(continues on next page)

(continued from previous page)

Started

Done!

```
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
{"text": "", "timestamps": [0.72, 0.96, 1.26, 1.44, 1.92, 2.10, 2.58, 2.82, 3.30, 3.90, 4.20, 4.56, 4.74], "tokens": ["", "", "", "", "", "", "", "", "", "", "", ""], "words": []}
-----
num threads: 1
decoding method: greedy_search
Elapsed seconds: 0.625 s
Real time factor (RTF): 0.625 / 5.592 = 0.112
```

Hint: Valid values for --sense-voice-language are auto, zh, en, ko, ja, and yue. where zh is for Chinese, en for English, ko for Korean, ja for Japanese, and yue for Cantonese.

8.26.4 C API for SenseVoice

This page describes how to use the C API for SenseVoice.

Please refer to [C API](#) for how to build `sherpa-onnx`.

The following is a very quick introduction for using the C API of `sherpa-onnx` in the form of shared libraries on macOS and Linux.

Hint: We do support static libraries and also support Windows.

If you copy, paste, and run the following commands in your terminal, you should be able to see the following recognition result:

```
Decoded text: The tribal chieftain called for the boy and presented him with 50 pieces
of gold.
```

```
cd /tmp

git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17
echo "----"
ls -lh sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs
```

(continues on next page)

(continued from previous page)

```

mkdir build
cd build
cmake \
-D CMAKE_BUILD_TYPE=Release \
-D BUILD_SHARED_LIBS=ON \
-D CMAKE_INSTALL_PREFIX=./install \
-D SHERPA_ONNX_ENABLE_BINARY=OFF \
..

make -j2 install

ls -lh install/lib
ls -lh install/include

cd ..

gcc -o sense-voice-c-api ./c-api-examples/sense-voice-c-api.c \
-I ./build/install/include \
-L ./build/install/lib/ \
-l sherpa-onnx-c-api \
-l onnxruntime

ls -lh sense-voice-c-api

export LD_LIBRARY_PATH=$PWD/build/install/lib:$LD_LIBRARY_PATH
export DYLD_LIBRARY_PATH=$PWD/build/install/lib:$DYLD_LIBRARY_PATH

./sense-voice-c-api

```

Note that we have hard-coded the file paths inside `sense-voice-c-api.c`

Hint: Since we are using shared libraries in the above example, you have to set the environment variable `LD_LIBRARY_PATH` for Linux and `DYLD_LIBRARY_PATH` for macOS. Otherwise, you would get runtime errors when running `./sense-voice-c-api`.

Explanations

1. Download `sherpa-onnx`

```

cd /tmp

git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

```

In this example, we download `sherpa-onnx` and place it inside the directory `/tmp/`. You can replace `/tmp/` with any directory you like.

Please always download the latest master of `sherpa-onnx`.

2. Download the model

```
curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-  
→onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2  
tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2  
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
```

Note that we have placed the model in the directory /tmp/sherpa-onnx.

3. Build sherpa-onnx

```
mkdir build  
cd build  
cmake \  
→ -D CMAKE_BUILD_TYPE=Release \  
→ -D BUILD_SHARED_LIBS=ON \  
→ -D CMAKE_INSTALL_PREFIX=./install \  
→ -D SHERPA_ONNX_ENABLE_BINARY=OFF \  
→ ..  
make -j2 install
```

We build a Release version of `sherpa-onnx`. Also, we use shared libraries here. The header file `c-api.h` and shared libraries are installed into the directory `./build/install`.

If you are using Linux, you should see the following content:

```
Install the project...  
-- Install configuration: "Release"  
-- Installing: /tmp/sherpa-onnx/build/install/lib/libonnxruntime.so  
-- Installing: /tmp/sherpa-onnx/build/install/.sherpa-onnx.pc  
-- Installing: /tmp/sherpa-onnx/build/install/lib/libsherpa-onnx-c-api.so  
-- Set non-toolchain portion of runtime path of "/tmp/sherpa-onnx/build/install/lib/  
→ libsherpa-onnx-c-api.so" to "$ORIGIN"  
-- Installing: /tmp/sherpa-onnx/build/install/include/sherpa-onnx/c-api/c-api.h
```

If you are using macOS, you should see:

```
Install the project...  
-- Install configuration: "Release"  
-- Installing: /tmp/sherpa-onnx/build/install/lib/libonnxruntime.1.17.1.dylib  
-- Installing: /tmp/sherpa-onnx/build/install/lib/libonnxruntime.dylib  
-- Installing: /tmp/sherpa-onnx/build/install/.sherpa-onnx.pc  
-- Installing: /tmp/sherpa-onnx/build/install/lib/libsherpa-onnx-c-api.dylib  
-- Installing: /tmp/sherpa-onnx/build/install/include/sherpa-onnx/c-api/c-api.h
```

4. View the build result

```
ls -lh install/lib
ls -lh install/include
```

If you are using Linux, you should see the following content:

```
total 19M
-rw-r--r-- 1 runner docker 15M Jul 22 08:47 libbonnrxruntime.so
-rw-r--r-- 1 runner docker 4.1M Jul 22 08:47 libsherpa-onnx-c-api.so
drwxr-xr-x 2 runner docker 4.0K Jul 22 08:47 pkgconfig
total 4.0K
drwxr-xr-x 3 runner docker 4.0K Jul 22 08:47 sherpa-onnx
```

If you are using macOS, you should see the following content:

```
total 53976
-rw-r--r-- 1 runner staff 23M Jul 22 08:48 libbonnrxruntime.1.17.1.dylib
lrwxr-xr-x 1 runner staff 27B Jul 22 08:48 libbonnrxruntime.dylib -> libbonnrxruntime.1.
→17.1.dylib
-rwxr-xr-x 1 runner staff 3.5M Jul 22 08:48 libsherpa-onnx-c-api.dylib
drwxr-xr-x 3 runner staff 96B Jul 22 08:48 pkgconfig
total 0
drwxr-xr-x 3 runner staff 96B Jul 22 08:48 sherpa-onnx
```

5. Build the C API example

```
cd ..
gcc -o sense-voice-c-api ./c-api-examples/sense-voice-c-api.c \
-I ./build/install/include \
-L ./build/install/lib/ \
-l sherpa-onnx-c-api
ls -lh sense-voice-c-api
```

Note that:

- `-I ./build/install/include` is to add the directory `./build/install/include` to the header search path so that `#include "sherpa-onnx/c-api/c-api.h"` won't throw an error.
- `-L ./build/install/lib/` is to add the directory `./build/install/lib` to the library search path so that it can find `-l sherpa-onnx-c-api`
- `-l sherpa-onnx-c-api` is to link the library `libsherpa-onnx-c-api.so` for Linux and `libsherpa-onnx-c-api.dylib` for macOS.

6. Run it

```
export LD_LIBRARY_PATH=$PWD/build/install/lib:$LD_LIBRARY_PATH
export DYLD_LIBRARY_PATH=$PWD/build/install/lib:$DYLD_LIBRARY_PATH

./sense-voice-c-api
```

Note that we have to use:

```
# For Linux
export LD_LIBRARY_PATH=$PWD/build/install/lib:$LD_LIBRARY_PATH
```

and:

```
# for macOS
export DYLD_LIBRARY_PATH=$PWD/build/install/lib:$DYLD_LIBRARY_PATH
```

Otherwise, it cannot find `libsherpa-onnx-c-api.so` for Linux and `libsherpa-onnx-c-api.dylib` at runtime.

7. Where to find `sense-voice-c-api.c`

You can find `sense-voice-c-api.c` at the following address:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/c-api-examples/sense-voice-c-api.c>

8.26.5 Dart API for SenseVoice

This page describes how to use the Dart API to run `SenseVoice` models in `sherpa-onnx`

Note that we have published the package `sherpa_onnx` at https://pub.dev/packages/sherpa_onnx.

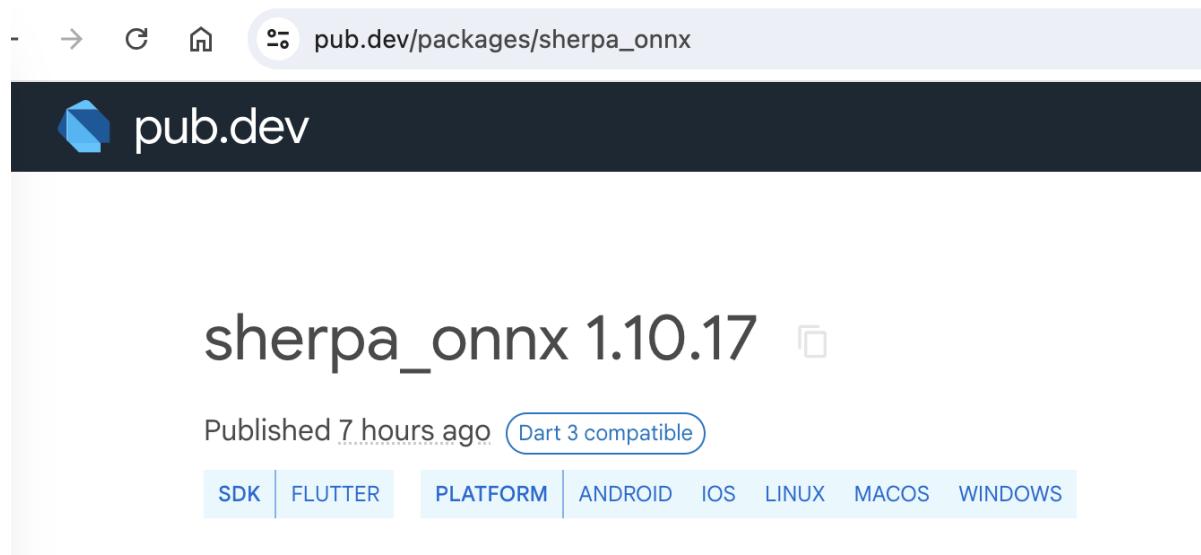


Fig. 8.34: Screenshot of `sherpa-onnx` on `pub.dev`.

Note that the package supports the following platforms:

- Android
- iOS
- Linux
- macOS
- Windows

In the following, we show how to use the pure Dart API to decode files with SenseVoice models.

```
cd /tmp

git clone http://github.com/k2-fsa/sherpa-onnx

cd sherpa-onnx
cd dart-api-examples
cd non-streaming-asr
dart pub get
./run-sense-voice.sh
```

You should see the following recognition result:

95

Explanations

1. Download the code

```
cd /tmp

git clone http://github.com/k2-fsa/sherpa-onnx
```

In this example, we download `sherpa-onnx` and place it inside the directory `/tmp/`. You can replace `/tmp/` with any directory you like.

2. Download the sherpa-onnx package

```
cd sherpa-onnx
cd dart-api-examples
cd non-streaming-asr
dart pub get
```

The command `dart pub get` will download the `sherpa_onnx` package automagically from `pub.dev`.

You should see something like below after running `dart pub get`:

```
(py38) fangjuns-MacBook-Pro:non-streaming-asr fangjun$ dart pub get
Resolving dependencies... (1.2s)
Downloading packages... (33.3s)
  collection 1.18.0 (1.19.0 available)
  lints 3.0.0 (4.0.0 available)
  material_color_utilities 0.8.0 (0.12.0 available)
```

(continues on next page)

(continued from previous page)

```
meta 1.12.0 (1.15.0 available)
> sherpa_onnx 1.10.17 (was 1.9.29)
+ sherpa_onnx_android 1.10.17
+ sherpa_onnx_ios 1.10.17
+ sherpa_onnx_linux 1.10.17
+ sherpa_onnx_macos 1.10.17
+ sherpa_onnx_windows 1.10.17
Changed 6 dependencies!
4 packages have newer versions incompatible with dependency constraints.
Try `dart pub outdated` for more information.
```

3. Run it

```
./run-sense-voice.sh
```

The above script downloads models and run the code automatically.

You can find `run-sense-voice.sh` at the following address:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/dart-api-examples/non-streaming-asr/run-sense-voice.sh>

The Dart API example code can be found at:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/dart-api-examples/non-streaming-asr/bin/sense-voice.dart>

8.26.6 Python API for SenseVoice

This page describes how to use the Python API for `SenseVoice`.

Please refer to *Install the Python Package* for how to install the Python package of `sherpa-onnx`.

The following is a quick way to do that:

```
pip install sherpa-onnx
```

Decode a file

After installing the Python package, you can download the Python example code and run it with the following commands:

```
cd /tmp
git clone https://github.com/k2-fsa/sherpa-onnx.git/
cd sherpa-onnx

curl -SL -O https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/sherpa-
˓→onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
tar xvf sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2
rm sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17.tar.bz2

python3 ./python-api-examples/offline-sense-voice-ctc-decode-files.py
```

You should see something like below:

```
./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
{"text": "95", "timestamps": [0.72, 0.96, 1.26, 1.44, 1.92, 2.10, 2.58, 2.82, 3.30, 3.90,
  ↵ 4.20, 4.56, 4.74, 5.46], "tokens": ["", "", "", "", "", "", "9", "", "", "", "", "5", "",
  ↵ "", ""], "words": []}

(py38) fangjuns-MacBook-Pro:sherpa-onnx fangjun$ #python3 ./python-api-examples/offline-
  ↵ sense-voice-ctc-decode-files.py
```

You can find `offline-sense-voice-ctc-decode-files.py` at the following address:

<https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/offline-sense-voice-ctc-decode-files.py>

Speech recognition from a microphone

The following example shows how to use a microphone with `SenseVoice` and `silero-vad` for speech recognition:

```
cd /tmp/sherpa-onnx

# Assume you have downloaded the SenseVoice model

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx

python3 ./python-api-examples/vad-with-non-streaming-asr.py \
--silero-vad-model=./silero_vad.onnx \
--sense-voice=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/model.int8.onnx \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--num-threads=2
```

You should see something like below:

```
0 Background Music, Core Audio (2 in, 2 out)
1 Background Music (UI Sounds), Core Audio (2 in, 2 out)
> 2 MacBook Pro Microphone, Core Audio (1 in, 0 out)
< 3 MacBook Pro Speakers, Core Audio (0 in, 2 out)
  4 WeMeet Audio Device, Core Audio (2 in, 2 out)
Use default device: MacBook Pro Microphone
Creating recognizer. Please wait...
Started! Please speak
```

If you start speaking, you should see some output after you stop speaking.

Hint: It starts speech recognition after `silero-vad` detects a pause.

Generate subtitles

This section describes how to use SenseVoice and silero-vad to generate subtitles.

Chinese

Test with a wave file containing Chinese:

```
cd /tmp/sherpa-onnx

# Assuem you have downloaded the SenseVoice model

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/lei-jun-test.wav

python3 ./python-api-examples/generate-subtitles.py \
--silero-vad-model=./silero_vad.onnx \
--sense-voice=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/model.onnx \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--num-threads=2 \
./lei-jun-test.wav
```

It will generate a text file lei-jun-test.srt, which is given below:

Click to see [lei-jun-test.srt](#).

```
1
0:00:28,934 --> 0:00:36,006

2
0:00:42,118 --> 0:00:46,374

3
0:00:46,918 --> 0:00:50,118

4
0:00:50,406 --> 0:00:55,750

5
0:00:56,134 --> 0:00:57,574

6
0:00:58,182 --> 0:01:06,854

7
0:01:07,718 --> 0:01:10,886
```

(continues on next page)

(continued from previous page)

8
0:01:11,654 --> 0:01:15,142

9
0:01:15,430 --> 0:01:19,526
3500

10
0:01:19,942 --> 0:01:22,278

11
0:01:22,694 --> 0:01:25,798

12
0:01:26,342 --> 0:01:30,886

13
0:01:38,470 --> 0:01:39,910

14
0:01:40,358 --> 0:01:44,486

15
0:01:45,926 --> 0:01:47,334

16
0:01:47,910 --> 0:01:50,694
130

17
0:01:51,750 --> 0:01:52,838

18
0:01:53,350 --> 0:01:54,886

19
0:01:55,206 --> 0:01:57,222

20
0:01:58,054 --> 0:01:59,558

(continues on next page)

(continued from previous page)

21
0:01:59,814 --> 0:02:02,598

22
0:02:03,654 --> 0:02:05,670

23
0:02:06,246 --> 0:02:08,614

24
0:02:09,190 --> 0:02:11,462

25
0:02:11,686 --> 0:02:14,182

26
0:02:14,982 --> 0:02:17,670

27
0:02:18,278 --> 0:02:19,494

28
0:02:21,830 --> 0:02:23,814
36

29
0:02:25,926 --> 0:02:27,654
1987

30
0:02:28,678 --> 0:02:31,622

31
0:02:32,678 --> 0:02:35,174

32
0:02:35,398 --> 0:02:36,710

33
0:02:37,574 --> 0:02:38,630

(continues on next page)

(continued from previous page)

34
0:02:39,334 --> 0:02:41,638

35
0:02:43,302 --> 0:02:44,454

36
0:02:45,286 --> 0:02:46,438

37
0:02:47,590 --> 0:02:49,318

38
0:02:50,406 --> 0:02:51,238

39
0:02:52,006 --> 0:02:52,966

40
0:02:53,318 --> 0:02:54,662

41
0:02:55,398 --> 0:02:57,670

42
0:02:58,342 --> 0:02:59,782

43
0:03:00,774 --> 0:03:02,470

44
0:03:02,950 --> 0:03:05,222

45
0:03:06,470 --> 0:03:07,750

46
0:03:08,934 --> 0:03:10,310

(continues on next page)

(continued from previous page)

47
0:03:10,598 --> 0:03:11,814

48
0:03:13,958 --> 0:03:14,822

49
0:03:17,606 --> 0:03:18,822

50
0:03:19,270 --> 0:03:22,502

51
0:03:29,766 --> 0:03:30,534

52
0:03:30,758 --> 0:03:32,550

53
0:03:40,326 --> 0:03:42,726

54
0:03:43,814 --> 0:03:47,046

55
0:03:48,134 --> 0:03:50,630

56
0:03:51,526 --> 0:03:56,326

57
0:03:57,574 --> 0:04:00,102

58
0:04:00,358 --> 0:04:02,278

59
0:04:03,846 --> 0:04:04,934

(continues on next page)

(continued from previous page)

60
0:04:05,190 --> 0:04:06,918

61
0:04:07,974 --> 0:04:08,966

62
0:04:10,342 --> 0:04:13,798

63
0:04:14,982 --> 0:04:16,102

64
0:04:16,774 --> 0:04:18,022

65
0:04:18,342 --> 0:04:19,878

66
0:04:20,422 --> 0:04:21,382

67
0:04:22,150 --> 0:04:22,982

68
0:04:23,366 --> 0:04:25,670

69
0:04:26,662 --> 0:04:27,174

70
0:04:28,486 --> 0:04:31,398

English

Test with a wave file containing English:

```
cd /tmp/sherpa-onnx

# Assuem you have downloaded the SenseVoice model

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/silero_vad.onnx
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/asr-models/Obama.wav

python3 ./python-api-examples/generate-subtitles.py \
--silero-vad-model=./silero_vad.onnx \
--sense-voice=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/model.onnx \
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt \
--num-threads=2 \
./Obama.wav
```

It will generate a text file `Obama.srt`, which is given below:

Click to see `Obama.srt`.

```
1
0:00:09,286 --> 0:00:12,486
Everybody, all right, everybody, go ahead and have a seat.

2
0:00:13,094 --> 0:00:15,014
How's everybody doing today?

3
0:00:18,694 --> 0:00:20,742
How about Tim Sper.

4
0:00:25,894 --> 0:00:31,942
I am here with students at Wakefield High School in Arlington, Virginia.

5
0:00:32,710 --> 0:00:48,326
And we've got students tuning in from all across America, from kindergarten through 12thu
grade. And I am just so glad that all could join us today. And I want to thanku
Wakefield for being such an outstanding host. Give yourselves a big line.

6
0:00:54,406 --> 0:00:59,238
Now I know that for many of you, today is the first day of school.

7
0:00:59,590 --> 0:01:09,798
And for those of you in kindergarten or starting middle or high school, it's your firstu
day in a new school, so it's understandable if you're a little nervous.

8
```

(continues on next page)

(continued from previous page)

0:01:10,630 --> 0:01:16,006

I imagine there's some seniors out there who are feeling pretty good right now with just ↵one more year to go.

9

0:01:18,790 --> 0:01:27,142

And no matter what grade you're in, some of you are probably wishing it we're still ↵summer, and you could have stayed in bed just a little bit longer this morning.

10

0:01:27,942 --> 0:01:29,414

I know that field.

11

0:01:31,654 --> 0:01:51,708

When I was young, my family lived overseas, I lived in Indonesia for a few years and my ↵mother, she didn't have the money to send me where all the American kids went to ↵school, but she thought it was important for me to keep up with American education, so ↵she decided to teach me extra lessons herself.

12

0:01:52,230 --> 0:01:58,790

Monday through Friday, but because she had to go to work, the only time she could do it ↵was at 4:30 in the morning.

13

0:02:00,038 --> 0:02:03,750

Now as you might imagine, I wasn't too happy about getting up that early.

14

0:02:04,102 --> 0:02:07,302

A lot of times I'd fall asleep right there at the kitchen table.

15

0:02:08,262 --> 0:02:15,014

But whenever I'd complained, my mother would just give me one of those looks and she'd ↵say, this is no picnic for me either, Buster.

16

0:02:17,094 --> 0:02:25,382

So I know that some of you are still adjusting to being back at school, but I'm here ↵today because I have something important to discuss with you.

17

0:02:25,798 --> 0:02:33,798

I'm here because I want to talk with you about your education and what's expected of all ↵of you in this new school year.

18

0:02:34,470 --> 0:02:40,422

I've given a lot of speeches about education, and I've talked about responsibility a lot.

(continues on next page)

(continued from previous page)

19
0:02:40,806 --> 0:02:47,174
I've talked about teachers responsibility for inspiring students and pushing you to
→learn.

20
0:02:47,430 --> 0:02:58,726
I talked about your parents' responsibility for making sure you stay on track and you
→get your homework done and don't spend every waking hour in front of the TV or with
→the Xbox.

21
0:02:59,078 --> 0:03:00,774
I've talked a lot about.

22
0:03:01,350 --> 0:03:13,286
Your government's responsibility for setting high standards and supporting teachers and
→principals and turning around schools that aren't working where students aren't
→getting the opportunities that they deserve.

23
0:03:13,990 --> 0:03:15,366
But at the end of the day.

24
0:03:16,006 --> 0:03:26,054
We can have the most dedicated teachers, the most supportive parents, the best schools
→in the world, and none of it will make a difference, none of it will matter.

25
0:03:26,694 --> 0:03:30,694
Unless all of you fulfill your responsibilities.

26
0:03:31,238 --> 0:03:43,814
Unless you show up to those schools, unless you pay attention to those teachers, unless
→you listen to your parents and grandparents and other adults and put in the hard work
→it takes to succeed.

27
0:03:44,646 --> 0:03:46,598
That's what I want to focus on today.

28
0:03:47,110 --> 0:03:50,918
The responsibility each of you has for your education.

29
0:03:51,718 --> 0:03:54,854
I want to start with the responsibility you have to yourself.

(continues on next page)

(continued from previous page)

30
0:03:55,654 --> 0:03:59,078
Every single one of you has something that you're good at.

31
0:03:59,782 --> 0:04:02,406
Every single one of you has something to offer.

32
0:04:02,982 --> 0:04:07,590
And you have a responsibility to yourself to discover what that is.

33
0:04:08,326 --> 0:04:11,494
That's the opportunity an education can provide.

34
0:04:12,326 --> 0:04:22,598
Maybe you could be a great writer, maybe even good enough to write a book or articles in
a newspaper, but you might not know it until you write that English paper.

35
0:04:23,078 --> 0:04:25,894
That English class paper that's assigned to you.

36
0:04:26,694 --> 0:04:38,726
Maybe you could be an innovator or an inventor, maybe even good enough to come up with
the next iPhone or the new medicine or a vaccine, but you might not know it until you
do your project for your science class.

37
0:04:39,814 --> 0:04:44,838
Maybe you could be a mayor or a senator or a Supreme Court justice.

38
0:04:45,350 --> 0:04:50,182
But you might not know that until you join student government or the debate team.

39
0:04:51,558 --> 0:04:56,774
And no matter what you want to do with your life, I guarantee that you'll need an
education to do it.

40
0:04:57,318 --> 0:05:00,710
You want to be a doctor or a teacher or a police officer?

41
0:05:00,998 --> 0:05:09,702
You want to be a nurse or an architect, a lawyer or a member of our military, you're
going to need a good education for every single one of those careers.

(continues on next page)

(continued from previous page)

```
42  
0:05:10,054 --> 0:05:14,278  
You cannot drop out of school and just drop into a good job.  
  
43  
0:05:15,174 --> 0:05:19,846  
You've got to train for it and work for it and learn for it.  
  
44  
0:05:20,518 --> 0:05:23,654  
And this isn't just important for your own life and your own future.  
  
45  
0:05:24,678 --> 0:05:29,670  
What you make of your education will decide nothing less than the future of this country.  
  
46  
0:05:29,958 --> 0:05:32,998  
The future of America depends on you.
```

WebSocket server and client example

This example shows how to use a WebSocket server with SenseVoice for speech recognition.

1. Start the server

Please run

```
cd /tmp/sherpa-onnx  
  
# Assuem you have downloaded the SenseVoice model  
  
python3 ./python-api-examples/non_streaming_server.py \  
--sense-voice=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/model.int8.onnx \  
--tokens=./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt
```

You should see the following output after starting the server:

```
2024-07-28 20:22:38,389 INFO [non_streaming_server.py:1001] {'encoder': '', 'decoder': ''  
..., 'joiner': '', 'paraformer': '', 'sense_voice': './sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/model.int8.onnx', 'nemo_ctc': '', 'wenet_ctc': ''  
..., 'tdnn_model': '', 'whisper_encoder': '', 'whisper_decoder': '', 'whisper_language': '', 'whisper_task': 'transcribe', 'whisper_tail_paddings': -1  
..., 'tokens': './sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/tokens.txt', 'num_threads': 2, 'provider': 'cpu', 'sample_rate': 16000, 'feat_dim': 80, 'decoding_method': 'greedy_search', 'max_active_paths': 4, 'hotwords_file': '', 'hotwords_score': 1.5, 'blank_penalty': 0.0, 'port': 6006, 'max_batch_size': 3, 'max_wait_ms': 5, 'nn_pool_size': 1, 'max_m': 1}
```

(continues on next page)

(continued from previous page)

```
essage_size': 1048576, 'max_queue_size': 32, 'max_active_connections': 200, 'certificate
˓→': None, 'doc_root': './python-api-examples/web'}
2024-07-28 20:22:41,861 INFO [non_streaming_server.py:647] started
2024-07-28 20:22:41,861 INFO [non_streaming_server.py:659] No certificate provided
2024-07-28 20:22:41,866 INFO [server.py:707] server listening on 0.0.0.0:6006
2024-07-28 20:22:41,866 INFO [server.py:707] server listening on [::]:6006
2024-07-28 20:22:41,866 INFO [non_streaming_server.py:679] Please visit one of the
˓→following addresses:
```

<http://localhost:6006>

You can either visit the address <http://localhost:6006> or write code to interact with the server.

In the following, we describe possible approaches for interacting with the WebSocket server.

Hint: The WebSocket server is able to handle multiple clients/connections at the same time.

2. Start the client (decode files sequentially)

The following code sends the files in sequential one by one to the server for decoding.

```
cd /tmp/sherpa-onnx

python3 ./python-api-examples/offline-websocket-client-decode-files-sequential.py ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
```

You should see something like below on the server side:

```
2024-07-28 20:28:15,749 INFO [server.py:642] connection open
2024-07-28 20:28:15,749 INFO [non_streaming_server.py:835] Connected: ('::1', 53252, 0, 0). Number of connections: 1/200
2024-07-28 20:28:15,933 INFO [non_streaming_server.py:851] result: 95
2024-07-28 20:28:16,194 INFO [non_streaming_server.py:851] result: The tribal chieftain
˓→called for the boy and presented him with 50 pieces of gold.
2024-07-28 20:28:16,195 INFO [non_streaming_server.py:819] Disconnected: ('::1', 53252, 0, 0). Number of connections: 0/200
2024-07-28 20:28:16,196 INFO [server.py:260] connection closed
```

You should see something like below on the client side:

```
2024-07-28 20:28:15,750 INFO [offline-websocket-client-decode-files-sequential.py:114]
˓→Sending ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
95
2024-07-28 20:28:15,934 INFO [offline-websocket-client-decode-files-sequential.py:114]
˓→Sending ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
The tribal chieftain called for the boy and presented him with 50 pieces of gold.
```

3. Start the client (decode files in parallel)

The following code sends the files in parallel at the same time to the server for decoding.

```
cd /tmp/sherpa-onnx

python3 ./python-api-examples/offline-websocket-client-decode-files-paralell.py ./sherpa-
˓→onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav ./sherpa-onnx-sense-
˓→voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
```

You should see something like below on the server side:

```
2024-07-28 20:31:10,147 INFO [server.py:642] connection open
2024-07-28 20:31:10,148 INFO [non_streaming_server.py:835] Connected: ('::1', 53436, 0, 0). Number of connections: 1/200
2024-07-28 20:31:10,149 INFO [server.py:642] connection open
2024-07-28 20:31:10,149 INFO [non_streaming_server.py:835] Connected: ('::1', 53437, 0, 0). Number of connections: 2/200
2024-07-28 20:31:10,353 INFO [non_streaming_server.py:851] result: 95
2024-07-28 20:31:10,354 INFO [non_streaming_server.py:819] Disconnected: ('::1', 53436, 0, 0). Number of connections: 1/200
2024-07-28 20:31:10,356 INFO [server.py:260] connection closed
2024-07-28 20:31:10,541 INFO [non_streaming_server.py:851] result: The tribal chieftain
˓→called for the boy and presented him with 50 pieces of gold.
2024-07-28 20:31:10,542 INFO [non_streaming_server.py:819] Disconnected: ('::1', 53437, 0, 0). Number of connections: 0/200
2024-07-28 20:31:10,544 INFO [server.py:260] connection closed
```

You should see something like below on the client side:

```
2024-07-28 20:31:10,112 INFO [offline-websocket-client-decode-files-paralell.py:139] {
˓→'server_addr': 'localhost', 'server_port': 6006, 'sound_files': ['./sherpa-onnx-sense-
˓→voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav', './sherpa-onnx-sense-voice-zh-en-
˓→ja-ko-yue-2024-07-17/test_wavs/en.wav']}
2024-07-28 20:31:10,148 INFO [offline-websocket-client-decode-files-paralell.py:113]
˓→Sending ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
2024-07-28 20:31:10,191 INFO [offline-websocket-client-decode-files-paralell.py:113]
˓→Sending ./sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
2024-07-28 20:31:10,353 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
˓→sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/zh.wav
95
2024-07-28 20:31:10,542 INFO [offline-websocket-client-decode-files-paralell.py:131] ./
˓→sherpa-onnx-sense-voice-zh-en-ja-ko-yue-2024-07-17/test_wavs/en.wav
The tribal chieftain called for the boy and presented him with 50 pieces of gold.
```

4. Start the Web browser client

You can also start a browser to interact with the WebSocket server.

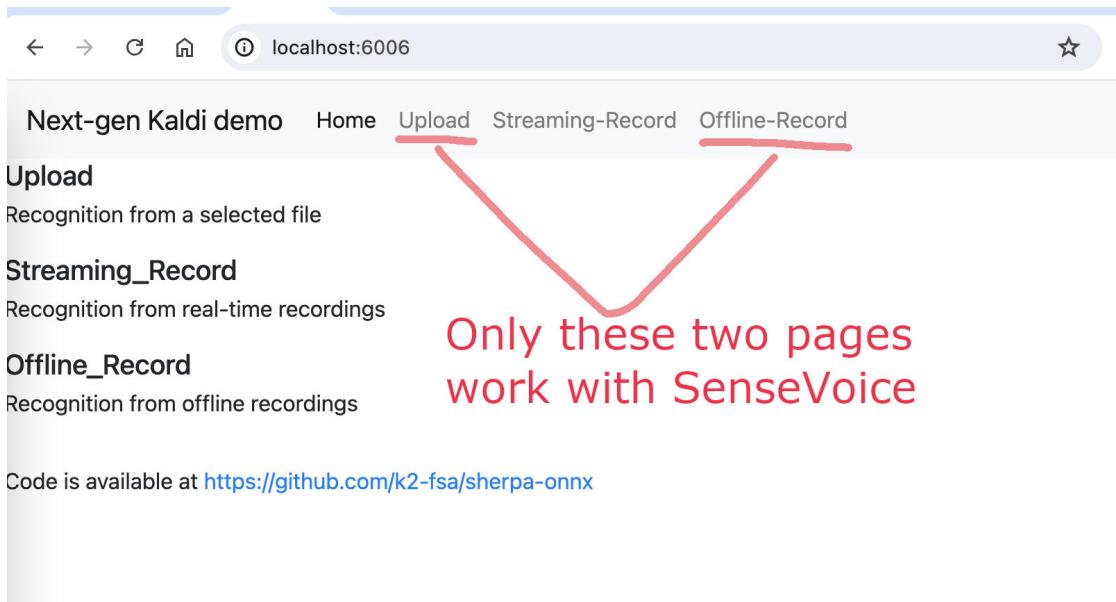
Please visit <http://localhost:6006>.

Warning: We are not using a certificate to start the server, so the only correct URL is <http://localhost:6006>.

All of the following addresses are **incorrect**:

- Incorrect/Wrong address: <https://localhost:6006>
- Incorrect/Wrong address: <http://127.0.0.1:6006>
- Incorrect/Wrong address: <https://127.0.0.1:6006>
- Incorrect/Wrong address: <http://a.b.c.d:6006>
- Incorrect/Wrong address: <https://a.b.c.d:6006>

After starting the browser, you should see the following page:



Upload a file for recognition

If we click Upload, we will see the following page:

localhost:6006/upload.html

Next-gen Kaldi demo Home **Upload** Streaming-Record Offline-Record

Recognition from a selected file

Click me to connect ws://localhost:6006

Select file

Choose File No file chosen

Recognition results

Please upload only a single channel, 16-bit, 16000 Hz wav file.
All other formats are NOT supported

Clear results

After clicking Click me to connect and Choose File, you will see the recognition result returned from the server:

localhost:6006/upload.html

Next-gen Kaldi demo Home Upload Streaming-Record Offline-Record

Recognition from a selected file

Click me to connect! ws://localhost:6006

Select file

Choose File en.wav

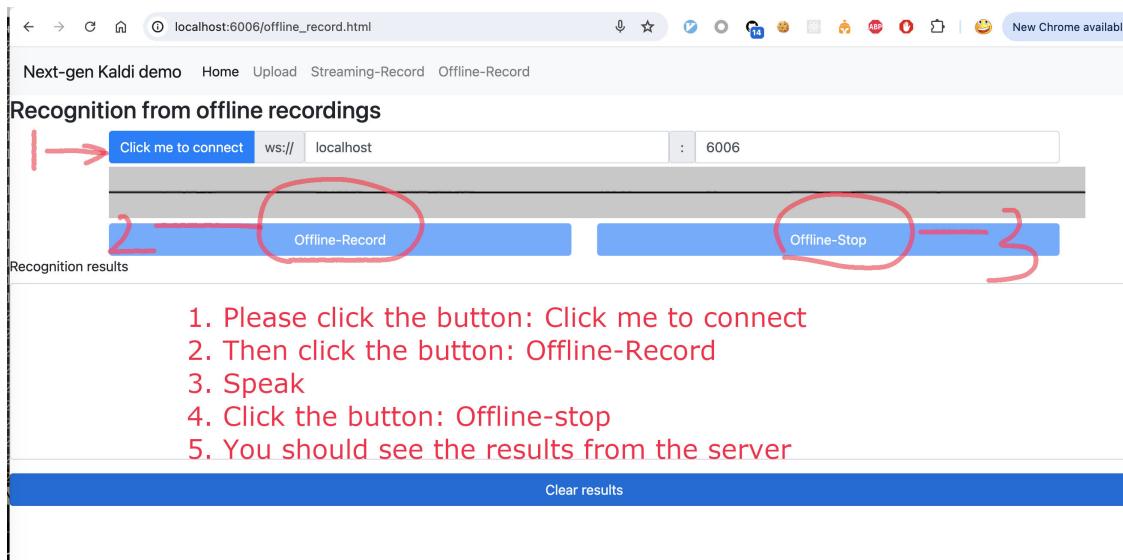
Recognition results

The tribal chieftain called for the boy and presented him with 50 pieces of gold.

Clear results


Record your speech with a microphone for recognition

If you click Offline-Record, you should see the following page:



Please click the button `Click me to connect`, and then click the button `Offline-Record`, then speak, finally, click the button `Offline-Stop`;

you should see the results from the server. A screenshot is given below:



Note that you can save the recorded audio into a wave file for debugging.

The recorded audio from the above screenshot is saved to `test.wav` and is given below:

Input File	:	'test.wav'
Channels	:	1
Sample Rate	:	16000

(continues on next page)

(continued from previous page)

Precision	:	16-bit
Duration	:	00:00:07.00 = 112012 samples ~ 525.056 CDDA sectors
File Size	:	224k
Bit Rate	:	256k
Sample Encoding:	16-bit Signed Integer PCM	

8.27 Speaker Identification

This page describes how to use `sherpa-onnx` for speaker identification.

Please first follow [Installation](#) and/or [Install the Python Package](#) to install `sherpa-onnx` before you continue.

Pre-trained models can be found at <https://github.com/k2-fsa/sherpa-onnx/releases/tag/speaker-recognition-models>

Hint: You can find Android APKs for each model at the following page

<https://k2-fsa.github.io/sherpa/onnx/speaker-identification/apk.html>

Please refer to <https://github.com/k2-fsa/sherpa-onnx/tree/master/python-api-examples> for usage examples.

8.28 Text-to-speech (TTS)

This page describes how to use `sherpa-onnx` for text-to-speech (TTS).

Please first follow [Installation](#) and/or [Install the Python Package](#) to install `sherpa-onnx` before you continue.

8.28.1 Huggingface space

We provide a huggingface space where you can try text-to-speech with `sherpa-onnx` from within your browser without installing anything.

Hint: We also have spaces using [WebAssembly](#) for text-to-speech. Please see [Huggingface Spaces \(WebAssembly\)](#).

All you need is a browser, either running on your desk computer, your phone, or your iPad, etc.

Please visit

<https://huggingface.co/spaces/k2-fsa/text-to-speech>

1. Select a language

2. Select a model

3. Input your text

4. Select a speaker

5. Click it

8.28.2 Pre-trained models

This page lists pre-trained models for text-to-speech.

Hint: Please install `git-lfs` before you continue.

Otherwise, you will be SAD later.

vits

This page lists pre-trained **vits** models.

All models in a single table

The following table summarizes the information of all models in this page.

Note: Since there are more than 100 pre-trained models for over 40 languages, we don't list all of them on this page. Please find them at <https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models>.

You can try all the models at the following huggingface space. <https://huggingface.co/spaces/k2-fsa/text-to-speech>.

Hint: You can find Android APKs for each model at the following page

<https://k2-fsa.github.io/sherpa/onnx/tts/apk.html>

Model	Language	# Speakers	Dataset	Model file-size (MB)	Sample rate (Hz)
<i>vits-melo-tts-zh_en (Chinese + English, 1 speaker)</i>	Chinese + English	1	N/A	163	44100
<i>vits-piper-en_US-libritts_r-medium (English, 904 speakers)</i>	English	904	LibriTTS-R	75	22050
<i>vits-piper-en_US-glados (English, 1 speaker)</i>	English	1	N/A	61	22050
<i>csukuangfj/sherpa-onnx-vits-zh-l1 (Chinese, 5 speakers)</i>	Chinese	5	N/A	115	16000
<i>csukuangfj/vits-zh-hf-fanchen-C (Chinese, 187 speakers)</i>	Chinese	187	N/A	116	16000
<i>csukuangfj/vits-zh-hf-fanchen-wnj (Chinese, 1 male)</i>	Chinese	1	N/A	116	16000
<i>csukuangfj/vits-zh-hf-theresa (Chinese, 804 speakers)</i>	Chinese	804	N/A	117	22050
<i>csukuangfj/vits-zh-hf-eula (Chinese, 804 speakers)</i>	Chinese	804	N/A	117	22050
<i>aishell3 (Chinese, multi-speaker, 174 speakers)</i>	Chinese	174	aishell3	116	8000
<i>ljspeech (English, single-speaker)</i>	English (US)	1 (Female)	LJ Speech	109	22050
<i>VCTK (English, multi-speaker, 109 speakers)</i>	English	109	VCTK	116	22050
<i>en_US-lessac-medium (English, single-speaker)</i>	English (US)	1 (Male)	lessac_blizzard2613	2613	22050

vits-melo-tts-zh_en (Chinese + English, 1 speaker)

This model is converted from <https://huggingface.co/myshell-ai/MeloTTS-Chinese> and it supports only 1 speaker. It supports both Chinese and English.

Note that if you input English words, only those that are present in the `lexicon.txt` can be pronounced. Please refer to <https://github.com/k2-fsa/sherpa-onnx/pull/1209> for how to add new words.

Hint: The converting script is available at <https://github.com/k2-fsa/sherpa-onnx/tree/master/scripts/melo-tts>

You can convert more models from <https://github.com/myshell-ai/MeloTTS> by yourself.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-melo-tts-zh_
en.tar.bz2
tar xvf vits-melo-tts-zh_en.tar.bz2
rm vits-melo-tts-zh_en.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
ls -lh vits-melo-tts-zh_en/
total 346848
-rw-r--r-- 1 fangjun staff 1.0K Jul 16 13:38 LICENSE
-rw-r--r-- 1 fangjun staff 156B Jul 16 13:38 README.md
-rw-r--r-- 1 fangjun staff 58K Jul 16 13:38 date.fst
drwxr-xr-x 9 fangjun staff 288B Apr 19 20:42 dict
-rw-r--r-- 1 fangjun staff 6.5M Jul 16 13:38 lexicon.txt
-rw-r--r-- 1 fangjun staff 163M Jul 16 13:38 model.onnx
-rw-r--r-- 1 fangjun staff 63K Jul 16 13:38 number.fst
-rw-r--r-- 1 fangjun staff 87K Jul 16 13:38 phone.fst
-rw-r--r-- 1 fangjun staff 655B Jul 16 13:38 tokens.txt
```

Generate speech with executable compiled from C++

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-melo-tts-zh_en/model.onnx \
--vits-lexicon=./vits-melo-tts-zh_en/lexicon.txt \
--vits-tokens=./vits-melo-tts-zh_en/tokens.txt \
--vits-dict-dir=./vits-melo-tts-zh_en/dict \
--output-filename=./zh-en-0.wav \
"This is a text to speech "
```

(continues on next page)

(continued from previous page)

```

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-melo-tts-zh_en/model.onnx \
--vits-lexicon=./vits-melo-tts-zh_en/lexicon.txt \
--vits-tokens=./vits-melo-tts-zh_en/tokens.txt \
--vits-dict-dir=./vits-melo-tts-zh_en/dict \
--output-filename=./zh-en-1.wav \
"machine learningartificial intelligence"

./build/bin/sherpa-onnx-offline-tts-play \
--vits-model=./vits-melo-tts-zh_en/model.onnx \
--vits-lexicon=./vits-melo-tts-zh_en/lexicon.txt \
--vits-tokens=./vits-melo-tts-zh_en/tokens.txt \
--tts-rule-fsts=". ./vits-melo-tts-zh_en/date.fst,./vits-melo-tts-zh_en/number.fst" \
--vits-dict-dir=./vits-melo-tts-zh_en/dict \
--output-filename=./zh-en-2.wav \
"Are you ok 20154 I am very happy to be in China."

```

After running, it will generate three files zh-en-1.wav, zh-en-2.wav, and zh-en-3.wav in the current directory.

```
soxi zh-en-*.wav
```

```

Input File      : 'zh-en-0.wav'
Channels        : 1
Sample Rate     : 44100
Precision       : 16-bit
Duration        : 00:00:03.54 = 156160 samples = 265.578 CDDA sectors
File Size       : 312k
Bit Rate        : 706k
Sample Encoding: 16-bit Signed Integer PCM

```

```

Input File      : 'zh-en-1.wav'
Channels        : 1
Sample Rate     : 44100
Precision       : 16-bit
Duration        : 00:00:05.98 = 263680 samples = 448.435 CDDA sectors
File Size       : 527k
Bit Rate        : 706k
Sample Encoding: 16-bit Signed Integer PCM

```

```

Input File      : 'zh-en-2.wav'
Channels        : 1
Sample Rate     : 44100
Precision       : 16-bit
Duration        : 00:00:18.92 = 834560 samples = 1419.32 CDDA sectors
File Size       : 1.67M
Bit Rate        : 706k
Sample Encoding: 16-bit Signed Integer PCM

```

```
Total Duration of 3 files: 00:00:28.44
```

Generate speech with Python script

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts-play.py \
--vits-model=./vits-melo-tts-zh_en/model.onnx \
--vits-lexicon=./vits-melo-tts-zh_en/lexicon.txt \
--vits-tokens=./vits-melo-tts-zh_en/tokens.txt \
--vits-dict-dir=./vits-melo-tts-zh_en/dict \
--output-filename=./zh-en-3.wav \
". Genius is one percent inspiration and ninety-nine percent perspiration. "
```

After running, it will generate a file zh-en-3.wav in the current directory.

```
soxi zh-en-3.wav

Input File      : 'zh-en-3.wav'
Channels        : 1
Sample Rate     : 44100
Precision       : 16-bit
Duration        : 00:00:09.83 = 433664 samples = 737.524 CDDA sectors
File Size       : 867k
Bit Rate        : 706k
Sample Encoding: 16-bit Signed Integer PCM
```

vits-piper-en_US-glados (English, 1 speaker)

This model is converted from <https://github.com/dnhkng/Glados/raw/main/models/glados.onnx> and it supports only English.

See also <https://github.com/dnhkng/GlaDOS>.

If you are interested in how the model is converted to sherpa-onnx, please see the following colab notebook:

<https://colab.research.google.com/drive/1m3Zr8H1RJaoZu4Y7hpQlav5vhtw3A513?usp=sharing>

In the following, we describe how to download it and use it with sherpa-onnx.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-piper-en_US-
glados.tar.bz2
tar xvf vits-piper-en_US-glados.tar.bz2
rm vits-piper-en_US-glados.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
ls -lh vits-piper-en_US-glados/  
  
-rw-r--r--    1 fangjun  staff   242B Dec 13 2023 README.md  
-rw-r--r--    1 fangjun  staff   61M Dec 13 2023 en_US-glados.onnx  
drwxr-xr-x 122 fangjun  staff  3.8K Dec 13 2023 espeak-ng-data  
-rw-r--r--    1 fangjun  staff  940B Dec 13 2023 tokens.txt
```

Generate speech with executable compiled from C++

```
cd /path/to/sherpa-onnx  
  
.build/bin/sherpa-onnx-offline-tts \  
  --vits-model=./vits-piper-en_US-glados/en_US-glados.onnx\  
  --vits-tokens=./vits-piper-en_US-glados/tokens.txt \  
  --vits-data-dir=./vits-piper-en_US-glados/espeak-ng-data \  
  --output-filename=./glados-liliana.wav \  
  "liliana, the most beautiful and lovely assistant of our team!"  
  
.build/bin/sherpa-onnx-offline-tts \  
  --vits-model=./vits-piper-en_US-glados/en_US-glados.onnx\  
  --vits-tokens=./vits-piper-en_US-glados/tokens.txt \  
  --vits-data-dir=./vits-piper-en_US-glados/espeak-ng-data \  
  --output-filename=./glados-code.wav \  
  "Talk is cheap. Show me the code."  
  
.build/bin/sherpa-onnx-offline-tts \  
  --vits-model=./vits-piper-en_US-glados/en_US-glados.onnx\  
  --vits-tokens=./vits-piper-en_US-glados/tokens.txt \  
  --vits-data-dir=./vits-piper-en_US-glados/espeak-ng-data \  
  --output-filename=./glados-men.wav \  
  "Today as always, men fall into two groups: slaves and free men. Whoever does not  
  ↪ have two-thirds of his day for himself, is a slave, whatever he may be: a statesman, a  
  ↪ businessman, an official, or a scholar."
```

After running, it will generate 3 files glados-liliana.wav, glados-code.wav, and glados-men.wav in the current directory.

```
soxi glados*.wav  
  
Input File      : 'glados-code.wav'  
Channels        : 1  
Sample Rate     : 22050  
Precision       : 16-bit  
Duration        : 00:00:02.18 = 48128 samples ~ 163.701 CDDA sectors  
File Size       : 96.3k  
Bit Rate        : 353k  
Sample Encoding: 16-bit Signed Integer PCM  
  
Input File      : 'glados-liliana.wav'  
Channels        : 1
```

(continues on next page)

(continued from previous page)

```

Sample Rate      : 22050
Precision       : 16-bit
Duration        : 00:00:03.97 = 87552 samples ~ 297.796 CDDA sectors
File Size       : 175k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Input File      : 'glados-men.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:15.31 = 337664 samples ~ 1148.52 CDDA sectors
File Size       : 675k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Total Duration of 3 files: 00:00:21.47

```

Generate speech with Python script

```

cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-piper-en_US-glados/en_US-glados.onnx \
--vits-tokens=./vits-piper-en_US-glados/tokens.txt \
--vits-data-dir=./vits-piper-en_US-glados/espeak-ng-data \
--output-filename=./glados-ship.wav \
"A ship in port is safe, but that's not what ships are built for."

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-piper-en_US-glados/en_US-glados.onnx \
--vits-tokens=./vits-piper-en_US-glados/tokens.txt \
--vits-data-dir=./vits-piper-en_US-glados/espeak-ng-data \
--output-filename=./glados-bug.wav \
"Given enough eyeballs, all bugs are shallow."

```

After running, it will generate two files `glados-ship.wav` and `glados-bug.wav` in the current directory.

```

soxi ./glados-{ship,bug}.wav

Input File      : './glados-ship.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.74 = 82432 samples ~ 280.381 CDDA sectors
File Size       : 165k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

```

(continues on next page)

(continued from previous page)

```
Input File      : './glados-bug.wav'
Channels       : 1
Sample Rate    : 22050
Precision      : 16-bit
Duration       : 00:00:02.67 = 58880 samples ~ 200.272 CDDA sectors
File Size      : 118k
Bit Rate       : 353k
Sample Encoding: 16-bit Signed Integer PCM

Total Duration of 2 files: 00:00:06.41
```

vits-piper-en_US-libritts_r-medium (English, 904 speakers)

This model is converted from https://huggingface.co/rhasspy/piper-voices/tree/main/en/en_US/libritts_r/medium and it supports 904 speakers. It supports only English.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-piper-en_US-
↪libritts_r-medium.tar.bz2
tar xvf vits-piper-en_US-libritts_r-medium.tar.bz2
rm vits-piper-en_US-libritts_r-medium.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
ls -lh vits-piper-en_US-libritts_r-medium/
total 153552
-rw-r--r--    1 fangjun  staff   279B Nov 29 2023 MODEL_CARD
-rw-r--r--    1 fangjun  staff   75M Nov 29 2023 en_US-libritts_r-medium.onnx
-rw-r--r--    1 fangjun  staff   20K Nov 29 2023 en_US-libritts_r-medium.onnx.json
drwxr-xr-x  122 fangjun  staff   3.8K Nov 28 2023 espeak-ng-data
-rw-r--r--    1 fangjun  staff   954B Nov 29 2023 tokens.txt
-rwxr-xr-x    1 fangjun  staff   1.8K Nov 29 2023 vits-piper-en_US.py
-rwxr-xr-x    1 fangjun  staff   730B Nov 29 2023 vits-piper-en_US.sh
```

Generate speech with executable compiled from C++

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-piper-en_US-libritts_r-medium/en_US-libritts_r-medium.onnx \
--vits-tokens=./vits-piper-en_US-libritts_r-medium/tokens.txt \
--vits-data-dir=./vits-piper-en_US-libritts_r-medium/espeak-ng-data \
--output-filename=./libritts-liliana-109.wav \
--sid=109 \
"liliana, the most beautiful and lovely assistant of our team!"

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-piper-en_US-libritts_r-medium/en_US-libritts_r-medium.onnx \
--vits-tokens=./vits-piper-en_US-libritts_r-medium/tokens.txt \
--vits-data-dir=./vits-piper-en_US-libritts_r-medium/espeak-ng-data \
--output-filename=./libritts-liliana-900.wav \
--sid=900 \
"liliana, the most beautiful and lovely assistant of our team!"
```

After running, it will generate two files `libritts-liliana-109.wav` and `libritts-liliana-900.wav` in the current directory.

```
soxi libritts-liliana-*.wav

Input File      : 'libritts-liliana-109.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:02.73 = 60160 samples ~ 204.626 CDDA sectors
File Size       : 120k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Input File      : 'libritts-liliana-900.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.36 = 73984 samples ~ 251.646 CDDA sectors
File Size       : 148k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Total Duration of 2 files: 00:00:06.08
```

Generate speech with Python script

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-piper-en_US-libritts_r-medium/en_US-libritts_r-medium.onnx \
--vits-tokens=./vits-piper-en_US-libritts_r-medium/tokens.txt \
--vits-data-dir=./vits-piper-en_US-libritts_r-medium/espeak-ng-data \
--sid=200 \
--output-filename=./libritts-armstrong-200.wav \
"That's one small step for a man, a giant leap for mankind."

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-piper-en_US-libritts_r-medium/en_US-libritts_r-medium.onnx \
--vits-tokens=./vits-piper-en_US-libritts_r-medium/tokens.txt \
--vits-data-dir=./vits-piper-en_US-libritts_r-medium/espeak-ng-data \
--sid=500 \
--output-filename=./libritts-armstrong-500.wav \
"That's one small step for a man, a giant leap for mankind."
```

After running, it will generate two files libritts-armstrong-200.wav and libritts-armstrong-500.wav in the current directory.

```
soxi ./libritts-armstrong*.wav

Input File      : './libritts-armstrong-200.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.11 = 68608 samples ~ 233.361 CDDA sectors
File Size       : 137k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Input File      : './libritts-armstrong-500.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.42 = 75520 samples ~ 256.871 CDDA sectors
File Size       : 151k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM

Total Duration of 2 files: 00:00:06.54
```

ljspeech (English, single-speaker)

This model is converted from `pretrained_ljspeech.pth`, which is trained by the `vits` author Jaehyeon Kim on the [LJ Speech](#) dataset. It supports only English and is a single-speaker model.

Note: If you are interested in how the model is converted, please see <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/vits/export-onnx-ljs.py>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-ljs.tar.bz2
tar xvf vits-ljs.tar.bz2
rm vits-ljs.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
-rw-r--r-- 1 1001 127 109M Apr 22 02:38 vits-ljs/vits-ljs.onnx
```

Generate speech with executable compiled from C++

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-ljs/vits-ljs.onnx \
--vits-lexicon=./vits-ljs/lexicon.txt \
--vits-tokens=./vits-ljs/tokens.txt \
--output-filename=./liliana.wav \
"liliana, the most beautiful and lovely assistant of our team!"
```

After running, it will generate a file `liliana.wav` in the current directory.

```
soxi ./liliana.wav

Input File      : './liliana.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:04.39 = 96768 samples ~ 329.143 CDDA sectors
File Size       : 194k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM
```

Generate speech with Python script

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-ljs/vits-ljs.onnx \
--vits-lexicon=./vits-ljs/lexicon.txt \
--vits-tokens=./vits-ljs/tokens.txt \
--output-filename=./armstrong.wav \
"That's one small step for a man, a giant leap for mankind."
```

After running, it will generate a file `armstrong.wav` in the current directory.

```
soxi ./armstrong.wav

Input File      : './armstrong.wav'
Channels       : 1
Sample Rate    : 22050
Precision      : 16-bit
Duration       : 00:00:04.81 = 105984 samples ~ 360.49 CDDA sectors
File Size      : 212k
Bit Rate       : 353k
Sample Encoding: 16-bit Signed Integer PCM
```

VCTK (English, multi-speaker, 109 speakers)

This model is converted from `pretrained_vctk.pth`, which is trained by the `vits` author Jaehyeon Kim on the VCTK dataset. It supports only English and is a multi-speaker model. It contains 109 speakers.

Note: If you are interested in how the model is converted, please see <https://github.com/k2-fsa/sherpa-onnx/blob/master/scripts/vits/export-onnx-vctk.py>

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-vctk.tar.bz2
tar xvf vits-vctk.tar.bz2
rm vits-vctk.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of `*.onnx` files below.

```
vits-vctk fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 37M Oct 16 10:57 vits-vctk.int8.onnx
-rw-r--r-- 1 fangjun staff 116M Oct 16 10:57 vits-vctk.onnx
```

Generate speech with executable compiled from C++

Since there are 109 speakers available, we can choose a speaker from 0 to 198. The default speaker ID is 0.

We use speaker ID 0, 10, and 108 below to generate audio for the same text.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=0 \
--output-filename=./kennedy-0.wav \
"Ask not what your country can do for you; ask what you can do for your country."

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=10 \
--output-filename=./kennedy-10.wav \
"Ask not what your country can do for you; ask what you can do for your country."

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=108 \
--output-filename=./kennedy-108.wav \
"Ask not what your country can do for you; ask what you can do for your country."
```

It will generate 3 files: kennedy-0.wav, kennedy-10.wav, and kennedy-108.wav.

Generate speech with Python script

We use speaker ID 30, 66, and 99 below to generate audio for different transcripts.

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=30 \
--output-filename=./einstein-30.wav \
"Life is like riding a bicycle. To keep your balance, you must keep moving."

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=66 \
```

(continues on next page)

(continued from previous page)

```
--output-filename=./franklin-66.wav \
"Three can keep a secret, if two of them are dead." 

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-vctk/vits-vctk.onnx \
--vits-lexicon=./vits-vctk/lexicon.txt \
--vits-tokens=./vits-vctk/tokens.txt \
--sid=99 \
--output-filename=./martin-99.wav \
"Darkness cannot drive out darkness: only light can do that. Hate cannot drive out
hate: only love can do that"
```

It will generate 3 files: einstein-30.wav, franklin-66.wav, and martin-99.wav.

csukuangfj/sherpa-onnx-vits-zh-II (Chinese, 5 speakers)

You can download the model using the following commands:

```
 wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/sherpa-onnx-vits-
zh-II.tar.bz2
 tar xvf sherpa-onnx-vits-zh-II.tar.bz2
 rm sherpa-onnx-vits-zh-II.tar.bz2
```

Hint: This model is trained with the following framework

<https://github.com/Plachtaa/VITS-fast-fine-tuning>

Please check the file sizes of the downloaded model:

```
ls -lh sherpa-onnx-vits-zh-II/
-rw-r--r-- 1 fangjun staff 2.3K Apr 25 17:58 G_multisperaker_latest.json
-rw-r-----@ 1 fangjun staff 2.2K Apr 25 17:22 G_multisperaker_latest_low.json
-rw-r--r-- 1 fangjun staff 127B Apr 25 17:58 README.md
-rw-r--r-- 1 fangjun staff 58K Apr 25 17:58 date.fst
drwxr-xr-x 9 fangjun staff 288B Jun 21 16:32 dict
-rw-r--r-- 1 fangjun staff 368K Apr 25 17:58 lexicon.txt
-rw-r--r-- 1 fangjun staff 115M Apr 25 17:58 model.onnx
-rw-r--r-- 1 fangjun staff 21K Apr 25 17:58 new_heteronym.fst
-rw-r--r-- 1 fangjun staff 63K Apr 25 17:58 number.fst
-rw-r--r-- 1 fangjun staff 87K Apr 25 17:58 phone.fst
-rw-r--r-- 1 fangjun staff 331B Apr 25 17:58 tokens.txt
```

usage:

```
sherpa-onnx-offline-tts \
--vits-model=./sherpa-onnx-vits-zh-II/model.onnx \
--vits-dict-dir=./sherpa-onnx-vits-zh-II/dict \
--vits-lexicon=./sherpa-onnx-vits-zh-II/lexicon.txt \
--vits-tokens=./sherpa-onnx-vits-zh-II/tokens.txt \
--vits-length-scale=0.5 \
```

(continues on next page)

(continued from previous page)

```
--sid=0 \
--output-filename="./0-value-2x.wav" \
"""

sherpa-onnx-offline-tts \
--vits-model=./sherpa-onnx-vits-zh-ll/model.onnx \
--vits-dict-dir=./sherpa-onnx-vits-zh-ll/dict \
--vits-lexicon=./sherpa-onnx-vits-zh-ll/lexicon.txt \
--vits-tokens=./sherpa-onnx-vits-zh-ll/tokens.txt \
--sid=1 \
--tts-rule-fsts=./sherpa-onnx-vits-zh-ll/number.fst \
--output-filename="./1-numbers.wav" \
"14"

sherpa-onnx-offline-tts \
--vits-model=./sherpa-onnx-vits-zh-ll/model.onnx \
--vits-dict-dir=./sherpa-onnx-vits-zh-ll/dict \
--vits-lexicon=./sherpa-onnx-vits-zh-ll/lexicon.txt \
--vits-tokens=./sherpa-onnx-vits-zh-ll/tokens.txt \
--tts-rule-fsts=./sherpa-onnx-vits-zh-ll/phone.fst,./sherpa-onnx-vits-zh-ll/number.fst \
--sid=2 \
--output-filename="./2-numbers.wav" \
"110 18601200909"

sherpa-onnx-offline-tts \
--vits-model=./sherpa-onnx-vits-zh-ll/model.onnx \
--vits-dict-dir=./sherpa-onnx-vits-zh-ll/dict \
--vits-lexicon=./sherpa-onnx-vits-zh-ll/lexicon.txt \
--vits-tokens=./sherpa-onnx-vits-zh-ll/tokens.txt \
--sid=3 \
--output-filename="./3-wo-mi.wav" \
"""

sherpa-onnx-offline-tts \
--vits-model=./sherpa-onnx-vits-zh-ll/model.onnx \
--vits-dict-dir=./sherpa-onnx-vits-zh-ll/dict \
--vits-lexicon=./sherpa-onnx-vits-zh-ll/lexicon.txt \
--vits-tokens=./sherpa-onnx-vits-zh-ll/tokens.txt \
--tts-rule-fsts=./sherpa-onnx-vits-zh-ll/number.fst \
--sid=4 \
--output-filename="./4-heteronym.wav" \
"35, 9"
```

csukuangfj/vits-zh-hf-fanchen-C (Chinese, 187 speakers)

You can download the model using the following commands:

```
 wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-zh-hf-
 ↵fanchen-C.tar.bz2
 tar xvf vits-zh-hf-fanchen-C.tar.bz2
 rm vits-zh-hf-fanchen-C.tar.bz2
```

Hint: This model is converted from https://huggingface.co/spaces/lkz99/tts_model/tree/main/zh

```
# information about model files

total 291M
-rw-r--r-- 1 1001 127 58K Apr 21 05:40 date.fst
drwxr-xr-x 3 1001 127 4.0K Apr 19 12:42 dict
-rw xr-xr-x 1 1001 127 4.0K Apr 21 05:40 export-onnx-zh-hf-fanchen-models.py
-rw xr-xr-x 1 1001 127 2.5K Apr 21 05:40 generate-lexicon-zh-hf-fanchen-models.py
-rw-r--r-- 1 1001 127 2.4M Apr 21 05:40 lexicon.txt
-rw-r--r-- 1 1001 127 22K Apr 21 05:40 new_heteronym.fst
-rw-r--r-- 1 1001 127 63K Apr 21 05:40 number.fst
-rw-r--r-- 1 1001 127 87K Apr 21 05:40 phone.fst
-rw-r--r-- 1 1001 127 173M Apr 21 05:40 rule.far
-rw-r--r-- 1 1001 127 331 Apr 21 05:40 tokens.txt
-rw-r--r-- 1 1001 127 116M Apr 21 05:40 vits-zh-hf-fanchen-C.onnx
-rw xr-xr-x 1 1001 127 2.0K Apr 21 05:40 vits-zh-hf-fanchen-models.sh
```

usage:

```
sherpa-onnx-offline-tts \
  --vits-model=./vits-zh-hf-fanchen-C/vits-zh-hf-fanchen-C.onnx \
  --vits-dict-dir=./vits-zh-hf-fanchen-C/dict \
  --vits-lexicon=./vits-zh-hf-fanchen-C/lexicon.txt \
  --vits-tokens=./vits-zh-hf-fanchen-C/tokens.txt \
  --vits-length-scale=0.5 \
  --output-filename="./value-2x.wav" \
  ...

sherpa-onnx-offline-tts \
  --vits-model=./vits-zh-hf-fanchen-C/vits-zh-hf-fanchen-C.onnx \
  --vits-dict-dir=./vits-zh-hf-fanchen-C/dict \
  --vits-lexicon=./vits-zh-hf-fanchen-C/lexicon.txt \
  --vits-tokens=./vits-zh-hf-fanchen-C/tokens.txt \
  --vits-length-scale=1.0 \
  --tts-rule-fsts=./vits-zh-hf-fanchen-C/number.fst \
  --output-filename="./numbers.wav" \
  "14"

sherpa-onnx-offline-tts \
  --sid=100 \
```

(continues on next page)

(continued from previous page)

```
--vits-model=./vits-zh-hf-fanchen-C/vits-zh-hf-fanchen-C.onnx \
--vits-dict-dir=./vits-zh-hf-fanchen-C/dict \
--vits-lexicon=./vits-zh-hf-fanchen-C/lexicon.txt \
--vits-tokens=./vits-zh-hf-fanchen-C/tokens.txt \
--vits-length-scale=1.0 \
--tts-rule-fsts=./vits-zh-hf-fanchen-C/phone.fst,./vits-zh-hf-fanchen-C/number.fst \
--output-filename="./numbers-100.wav" \
"110 18601200909"

sherpa-onnx-offline-tts \
--sid=14 \
--vits-model=./vits-zh-hf-fanchen-C/vits-zh-hf-fanchen-C.onnx \
--vits-dict-dir=./vits-zh-hf-fanchen-C/dict \
--vits-lexicon=./vits-zh-hf-fanchen-C/lexicon.txt \
--vits-tokens=./vits-zh-hf-fanchen-C/tokens.txt \
--vits-length-scale=1.0 \
--output-filename="./wo-mi-14.wav" \
"""

sherpa-onnx-offline-tts \
--sid=102 \
--vits-model=./vits-zh-hf-fanchen-C/vits-zh-hf-fanchen-C.onnx \
--vits-dict-dir=./vits-zh-hf-fanchen-C/dict \
--vits-lexicon=./vits-zh-hf-fanchen-C/lexicon.txt \
--vits-tokens=./vits-zh-hf-fanchen-C/tokens.txt \
--tts-rule-fsts=./vits-zh-hf-fanchen-C/number.fst \
--vits-length-scale=1.0 \
--output-filename="./heteronym-102.wav" \
"35, 91"
```

csukuangfi/vits-zh-hf-fanchen-wnj (Chinese, 1 male)

You can download the model using the following commands:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-zh-hf-
-fanchen-wnj.tar.bz2
tar xvf vits-zh-hf-fanchen-wnj.tar.bz2
rm vits-zh-hf-fanchen-wnj.tar.bz2
```

Hint: This model is converted from https://huggingface.co/spaces/lkz99/tts_model/blob/main/G_wnj_latest.pth

```
# information about model files
total 594760
-rw-r--r-- 1 fangjun staff 58K Apr 21 13:40 date.fst
drwxr-xr-x 9 fangjun staff 288B Apr 19 20:42 dict
-rwrxr-xr-x 1 fangjun staff 3.9K Apr 21 13:40 export-onnx-zh-hf-fanchen-models.py
-rwrxr-xr-x 1 fangjun staff 2.4K Apr 21 13:40 generate-lexicon-zh-hf-fanchen-models.py
-rw-r--r-- 1 fangjun staff 2.3M Apr 21 13:40 lexicon.txt
-rw-r--r-- 1 fangjun staff 21K Apr 21 13:40 new_heteronym.fst
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 63K Apr 21 13:40 number.fst
-rw-r--r-- 1 fangjun staff 87K Apr 21 13:40 phone.fst
-rw-r--r-- 1 fangjun staff 172M Apr 21 13:40 rule.far
-rw-r--r-- 1 fangjun staff 331B Apr 21 13:40 tokens.txt
-rwxr-xr-x 1 fangjun staff 1.9K Apr 21 13:40 vits-zh-hf-fanchen-models.sh
-rw-r--r-- 1 fangjun staff 115M Apr 21 13:40 vits-zh-hf-fanchen-wnj.onnx
```

usage:

```
sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-fanchen-wnj/vits-zh-hf-fanchen-wnj.onnx \
--vits-dict-dir=./vits-zh-hf-fanchen-wnj/dict \
--vits-lexicon=./vits-zh-hf-fanchen-wnj/lexicon.txt \
--vits-tokens=./vits-zh-hf-fanchen-wnj/tokens.txt \
--output-filename="./kuayue.wav" \
"""

sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-fanchen-wnj/vits-zh-hf-fanchen-wnj.onnx \
--vits-dict-dir=./vits-zh-hf-fanchen-wnj/dict \
--vits-lexicon=./vits-zh-hf-fanchen-wnj/lexicon.txt \
--vits-tokens=./vits-zh-hf-fanchen-wnj/tokens.txt \
--tts-rule-fsts=./vits-zh-hf-fanchen-wnj/number.fst \
--output-filename="./os.wav" \
"14"
```

csukuangfj/vits-zh-hf-theresa (Chinese, 804 speakers)

You can download the model with the following commands:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-zh-hf-
theresa.tar.bz2
tar xvf vits-zh-hf-theresa.tar.bz2
rm vits-zh-hf-theresa.tar.bz2
```

Hint: This model is converted from https://huggingface.co/spaces/zomehwh/vits-models-genshin-bh3/tree/main/pretrained_models/theresa

```
# information about model files

total 596992
-rw-r--r-- 1 fangjun staff 58K Apr 21 13:39 date.fst
drwxr-xr-x 9 fangjun staff 288B Apr 19 20:42 dict
-rw-r--r-- 1 fangjun staff 2.6M Apr 21 13:39 lexicon.txt
-rw-r--r-- 1 fangjun staff 21K Apr 21 13:39 new_heteronym.fst
-rw-r--r-- 1 fangjun staff 63K Apr 21 13:39 number.fst
-rw-r--r-- 1 fangjun staff 87K Apr 21 13:39 phone.fst
-rw-r--r-- 1 fangjun staff 172M Apr 21 13:39 rule.far
-rw-r--r-- 1 fangjun staff 116M Apr 21 13:39 theresa.onnx
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 268B Apr 21 13:39 tokens.txt
-rwxr-xr-x 1 fangjun staff 5.3K Apr 21 13:39 vits-zh-hf-models.py
-rwxr-xr-x 1 fangjun staff 571B Apr 21 13:39 vits-zh-hf-models.sh
```

usage:

```
sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-theresa/theresa.onnx \
--vits-dict-dir=./vits-zh-hf-theresa/dict \
--vits-lexicon=./vits-zh-hf-theresa/lexicon.txt \
--vits-tokens=./vits-zh-hf-theresa/tokens.txt \
--sid=0 \
--output-filename="./reai-0.wav" \
"""

sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-theresa/theresa.onnx \
--vits-dict-dir=./vits-zh-hf-theresa/dict \
--vits-lexicon=./vits-zh-hf-theresa/lexicon.txt \
--vits-tokens=./vits-zh-hf-theresa/tokens.txt \
--tts-rule-fsts=./vits-zh-hf-theresa/number.fst \
--debug=1 \
--sid=88 \
--output-filename="./mi14-88.wav" \
"141000000"
```

csukuangfj/vits-zh-hf-eula (Chinese, 804 speakers)

You can download the model using the following commands:

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-zh-hf-eula.
tar.bz2
tar xvf vits-zh-hf-eula.tar.bz2
rm vits-zh-hf-eula.tar.bz2
```

Hint: This model is converted from https://huggingface.co/spaces/zomehwh/vits-models-genshin-bh3/tree/main/pretrained_models/eula

```
# information about model files

total 596992
-rw-r--r-- 1 fangjun staff 58K Apr 21 13:39 date.fst
drwxr-xr-x 9 fangjun staff 288B Apr 19 20:42 dict
-rw-r--r-- 1 fangjun staff 116M Apr 21 13:39 eula.onnx
-rw-r--r-- 1 fangjun staff 2.6M Apr 21 13:39 lexicon.txt
-rw-r--r-- 1 fangjun staff 21K Apr 21 13:39 new_heteronym.fst
-rw-r--r-- 1 fangjun staff 63K Apr 21 13:39 number.fst
-rw-r--r-- 1 fangjun staff 87K Apr 21 13:39 phone.fst
-rw-r--r-- 1 fangjun staff 172M Apr 21 13:39 rule.far
```

(continues on next page)

(continued from previous page)

```
-rw-r--r-- 1 fangjun staff 268B Apr 21 13:39 tokens.txt
-rwxr-xr-x 1 fangjun staff 5.3K Apr 21 13:39 vits-zh-hf-models.py
-rwxr-xr-x 1 fangjun staff 571B Apr 21 13:39 vits-zh-hf-models.sh
```

usage:

```
sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-eula/eula.onnx \
--vits-dict-dir=./vits-zh-hf-eula/dict \
--vits-lexicon=./vits-zh-hf-eula/lexicon.txt \
--vits-tokens=./vits-zh-hf-eula/tokens.txt \
--debug=1 \
--sid=666 \
--output-filename="./news-666.wav" \
"""

sherpa-onnx-offline-tts \
--vits-model=./vits-zh-hf-eula/eula.onnx \
--vits-dict-dir=./vits-zh-hf-eula/dict \
--vits-lexicon=./vits-zh-hf-eula/lexicon.txt \
--vits-tokens=./vits-zh-hf-eula/tokens.txt \
--tts-rule-fsts=./vits-zh-hf-eula/number.fst \
--sid=99 \
--output-filename="./news-99.wav" \
"925"
```

aishell3 (Chinese, multi-speaker, 174 speakers)

This model is trained on the [aishell3](#) dataset using [icefall](#).

It supports only Chinese and it's a multi-speaker model and contains 174 speakers.

Hint: You can download the Android APK for this model at

<https://k2-fsa.github.io/sherpa/onnx/tts/apk-engine.html>

(Please search for `vits-icefall-zh-aishell3` in the above Android APK page)

Note: If you are interested in how the model is converted, please see the documentation of [icefall](#).

If you are interested in training your own model, please also refer to [icefall](#).

[icefall](#) is also developed by us.

In the following, we describe how to download it and use it with [sherpa-onnx](#).

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-icefall-zh-
↪aishell3.tar.bz2
tar xvf vits-icefall-zh-aishell3.tar.bz2
rm vits-icefall-zh-aishell3.tar.bz2
```

Please check that the file sizes of the pre-trained models are correct. See the file sizes of *.onnx files below.

```
vits-icefall-zh-aishell3 fangjun$ ls -lh *.onnx
-rw-r--r-- 1 fangjun staff 29M Mar 20 22:50 model.onnx
```

Generate speech with executable compiled from C++

Since there are 174 speakers available, we can choose a speaker from 0 to 173. The default speaker ID is 0.

We use speaker ID 10, 33, and 99 below to generate audio for the same text.

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
↪fst,./vits-icefall-zh-aishell3/number.fst \
--sid=10 \
--output-filename=./liliana-10.wav \
"""

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
↪fst,./vits-icefall-zh-aishell3/number.fst \
--sid=33 \
--output-filename=./liliana-33.wav \
"""

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
↪fst,./vits-icefall-zh-aishell3/number.fst \
--sid=99 \
--output-filename=./liliana-99.wav \
""
```

It will generate 3 files: `liliana-10.wav`, `liliana-33.wav`, and `liliana-99.wav`.

We also support rule-based text normalization, which is implemented with [OpenFst](#). Currently, only number normalization is supported.

Hint: We will support other normalization rules later.

The following is an example:

```
./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
fst,./vits-icefall-zh-aishell3/number.fst \
--sid=66 \
--output-filename=./rule-66.wav \
"35, 91"
```

Generate speech with Python script

We use speaker ID 21, 41, and 45 below to generate audio for different transcripts.

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
fst,./vits-icefall-zh-aishell3/number.fst \
--sid=21 \
--output-filename=./liubei-21.wav \
"""

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
fst,./vits-icefall-zh-aishell3/number.fst \
--sid=41 \
--output-filename=./demokelite-41.wav \
"""

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
fst,./vits-icefall-zh-aishell3/number.fst \
--sid=45 \
```

(continues on next page)

(continued from previous page)

```
--output-filename=./zhugeliang-45.wav \
...''
```

It will generate 3 files: liubei-21.wav, demokelite-41.wav, and zhugeliang-45.wav.

The Python script also supports rule-based text normalization.

```
python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-icefall-zh-aishell3/model.onnx \
--vits-lexicon=./vits-icefall-zh-aishell3/lexicon.txt \
--vits-tokens=./vits-icefall-zh-aishell3/tokens.txt \
--tts-rule-fsts=./vits-icefall-zh-aishell3/phone.fst,./vits-icefall-zh-aishell3/date.
↪fst,./vits-icefall-zh-aishell3/number.fst \
--sid=103 \
--output-filename=./rule-103.wav \
"7144349737831141177872411013812345678"
```

en_US-lessac-medium (English, single-speaker)

This model is converted from https://huggingface.co/rhasspy/piper-voices/tree/main/en/en_US/lessac/medium.

The dataset used to train the model is `lessac_blizzard2013`.

Hint: The model is from `piper`.

In the following, we describe how to download it and use it with `sherpa-onnx`.

Download the model

Please use the following commands to download it.

```
cd /path/to/sherpa-onnx

wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-piper-en_US-
↪lessac-medium.tar.bz2
tar xf vits-piper-en_US-lessac-medium.tar.bz2
```

Hint: You can find a lot of pre-trained models for over 40 languages at <<https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models>>.

Generate speech with executable compiled from C++

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts \
--vits-model=./vits-piper-en_US-lessac-medium/en_US-lessac-medium.onnx \
--vits-data-dir=./vits-piper-en_US-lessac-medium/espeak-ng-data \
--vits-tokens=./vits-piper-en_US-lessac-medium/tokens.txt \
--output-filename=./liliana-piper-en_US-lessac-medium.wav \
"liliana, the most beautiful and lovely assistant of our team!"
```

Hint: You can also use

```
cd /path/to/sherpa-onnx

./build/bin/sherpa-onnx-offline-tts-play \
--vits-model=./vits-piper-en_US-lessac-medium/en_US-lessac-medium.onnx \
--vits-data-dir=./vits-piper-en_US-lessac-medium/espeak-ng-data \
--vits-tokens=./vits-piper-en_US-lessac-medium/tokens.txt \
--output-filename=./liliana-piper-en_US-lessac-medium.wav \
"liliana, the most beautiful and lovely assistant of our team!"
```

which will play the audio as it is generating.

After running, it will generate a file `liliana-piper.wav` in the current directory.

```
soxi ./liliana-piper-en_US-lessac-medium.wav

Input File      : './liliana-piper-en_US-lessac-medium.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.48 = 76800 samples ~ 261.224 CDDA sectors
File Size       : 154k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM
```

Generate speech with Python script

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts.py \
--vits-model=./vits-piper-en_US-lessac-medium/en_US-lessac-medium.onnx \
--vits-data-dir=./vits-piper-en_US-lessac-medium/espeak-ng-data \
--vits-tokens=./vits-piper-en_US-lessac-medium/tokens.txt \
--output-filename=./armstrong-piper-en_US-lessac-medium.wav \
"That's one small step for a man, a giant leap for mankind."
```

Hint: You can also use

```
cd /path/to/sherpa-onnx

python3 ./python-api-examples/offline-tts-play.py \
--vits-model=./vits-piper-en_US-lessac-medium/en_US-lessac-medium.onnx \
--vits-data-dir=./vits-piper-en_US-lessac-medium/espeak-ng-data \
--vits-tokens=./vits-piper-en_US-lessac-medium/tokens.txt \
--output-filename=./armstrong-piper-en_US-lessac-medium.wav \
"That's one small step for a man, a giant leap for mankind."
```

which will play the audio as it is generating.

After running, it will generate a file `armstrong-piper-en_US-lessac-medium.wav` in the current directory.

```
soxi ./armstrong-piper-en_US-lessac-medium.wav

Input File      : './armstrong-piper-en_US-lessac-medium.wav'
Channels        : 1
Sample Rate     : 22050
Precision       : 16-bit
Duration        : 00:00:03.74 = 82432 samples ~ 280.381 CDDA sectors
File Size       : 165k
Bit Rate        : 353k
Sample Encoding: 16-bit Signed Integer PCM
```

8.28.3 WebAssembly

In this section, we describe how to build text-to-speech from `sherpa-onnx` for `WebAssembly` so that you can run text-to-speech with `WebAssembly`.

Please follow the steps below to build and run `sherpa-onnx` for `WebAssembly`.

Hint: We provide a colab notebook for you to try this section step by step.

If you are using Windows or you don't want to setup your local environment to build WebAssembly support, please use the above colab notebook.

Install Emscripten

We need to compile the C/C++ files in `sherpa-onnx` with the help of `emscripten`.

Please refer to https://emscripten.org/docs/getting_started/downloads for detailed installation instructions.

The following is an example to show you how to install it on Linux/macOS.

```
git clone https://github.com/emscripten-core/emsdk.git
cd emsdk
git pull
./emsdk install latest
./emsdk activate latest
source ./emsdk_env.sh
```

To check that you have installed `emscripten` successfully, please run:

```
emcc -v
```

The above command should print something like below:

```
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
  ↳(e967e20b4727956a30592165a3c1cde5c67fa0a8)
shared:INFO: (Emscripten: Running sanity checks)
(py38) fangjuns-MacBook-Pro:open-source fangjun$ emcc -v
emcc (Emscripten gcc/clang-like replacement + linker emulating GNU ld) 3.1.48
  ↳(e967e20b4727956a30592165a3c1cde5c67fa0a8)
clang version 18.0.0 (https://github.com/llvm/llvm-project
  ↳a54545ba6514802178cf7cf1c1dd9f7efbf3cde7)
Target: wasm32-unknown-emscripten
Thread model: posix
InstalledDir: /Users/fangjun/open-source/emsdk/upstream/bin
```

Congratulations! You have successfully installed `emscripten`.

Build

After installing `emscripten`, we can build text-to-speech from `sherpa-onnx` for WebAssembly now.

Please use the following command to build it:

```
git clone https://github.com/k2-fsa/sherpa-onnx
cd sherpa-onnx

cd wasm/tts/assets

wget -q https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/vits-piper-en_
  ↳US-libritts_r-medium.tar.bz2
tar xf vits-piper-en_US-libritts_r-medium.tar.bz2
rm vits-piper-en_US-libritts_r-medium.tar.bz2
mv vits-piper-en_US-libritts_r-medium/en_US-libritts_r-medium.onnx ./model.onnx
mv vits-piper-en_US-libritts_r-medium/tokens.txt ./
mv vits-piper-en_US-libritts_r-medium/espeak-ng-data ./
rm -rf vits-piper-en_US-libritts_r-medium

cd ../../..
./build-wasm-simd-tts.sh
```

Hint: You can visit <https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models> to download a different model.

After building, you should see the following output:

```
Install the project...
-- Install configuration: "Release"
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
  ↳libkaldi-native-fbank-core.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
  ↳libkaldi-decoder-core.a
```

(continues on next page)

(continued from previous page)

```
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libsherpa-onnx-kaldifst-core.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libsherpa-onnx-fst.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libonnxruntime.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libespeak-ng.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libucd.a
-- Up-to-date: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libucd.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libpiper_phonemize.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/./
→ sherpa-onnx.pc
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ pkgconfig/espeak-ng.pc
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/share/
→ vim/vimfiles/ftdetect
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/share/
→ vim/vimfiles/ftdetect/espeak filetype.vim
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/share/
→ vim/vimfiles/syntax
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/share/
→ vim/vimfiles/syntax/espeakrules.vim
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/share/
→ vim/vimfiles/syntax/espeaklist.vim
-- Up-to-date: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libucd.a
-- Up-to-date: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libespeak-ng.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libsherpa-onnx-core.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/lib/
→ libsherpa-onnx-c-api.a
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/
→ include/sherpa-onnx/c-api/c-api.h
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/sherpa-onnx-wasm-main.js
-- Up-to-date: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/sherpa-onnx-wasm-main.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/index.html
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/sherpa-onnx.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/app.js
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/sherpa-onnx-wasm-main.wasm
-- Installing: /Users/fangjun/open-source/sherpa-onnx/build-wasm-simd-tts/install/bin/
→ wasm/tts/sherpa-onnx-wasm-main.data
```

(continues on next page)

(continued from previous page)

```
+ ls -lh install/bin/wasm/tts
total 211248
-rw-r--r-- 1 fangjun staff 5.3K Feb 22 09:18 app.js
-rw-r--r-- 1 fangjun staff 1.3K Feb 22 09:18 index.html
-rw-r--r-- 1 fangjun staff 92M Feb 22 10:35 sherpa-onnx-wasm-main.data
-rw-r--r-- 1 fangjun staff 117K Feb 22 10:39 sherpa-onnx-wasm-main.js
-rw-r--r-- 1 fangjun staff 11M Feb 22 10:39 sherpa-onnx-wasm-main.wasm
-rw-r--r-- 1 fangjun staff 4.5K Feb 22 09:18 sherpa-onnx.js
```

Now you can use the following command to run it:

```
cd build-wasm-simd-tts/install/bin/wasm/tts
python3 -m http.server 6008
```

Start your browser and visit <http://localhost:6008/>; you should see the following page:



Now you can enter some text and click **Generate**

A screenshot is given below:

Congratulations! You have successfully run text-to-speech with [WebAssembly](#) in your browser.

Huggingface Spaces (WebAssembly)


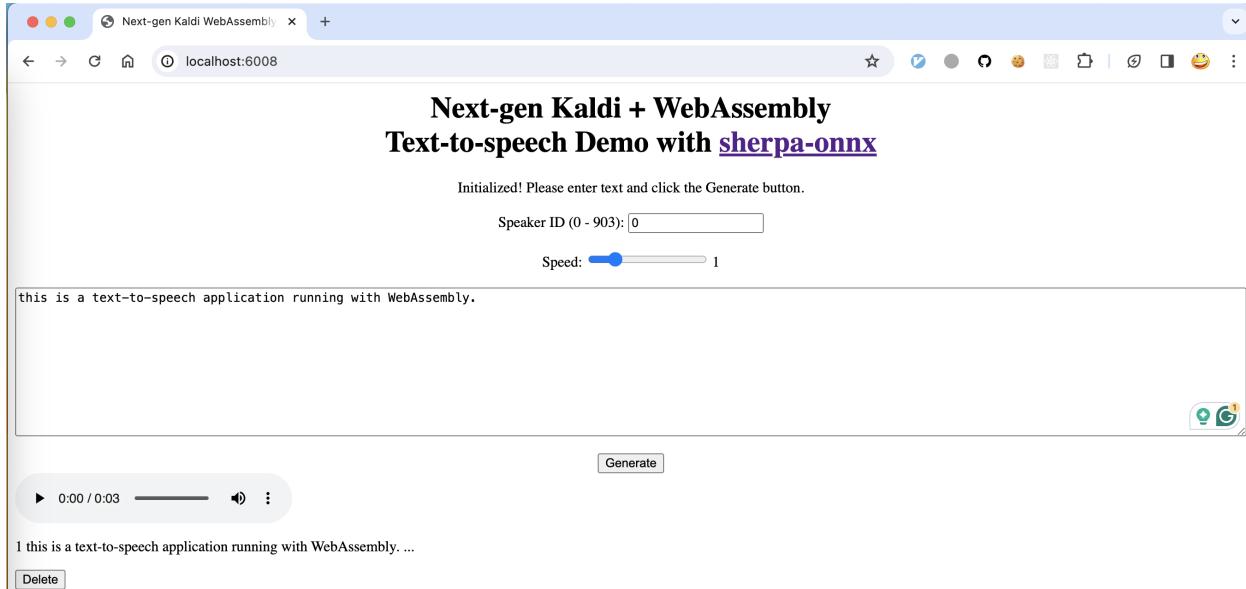
We provide two [Huggingface](#) spaces so that you can try text-to-speech with [WebAssembly](#) in your browser.

English TTS

<https://huggingface.co/spaces/k2-fsa/web-assembly-tts-sherpa-onnx-en>

Hint: If you don't have access to [Huggingface](#), please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-tts-sherpa-onnx-en/summary>



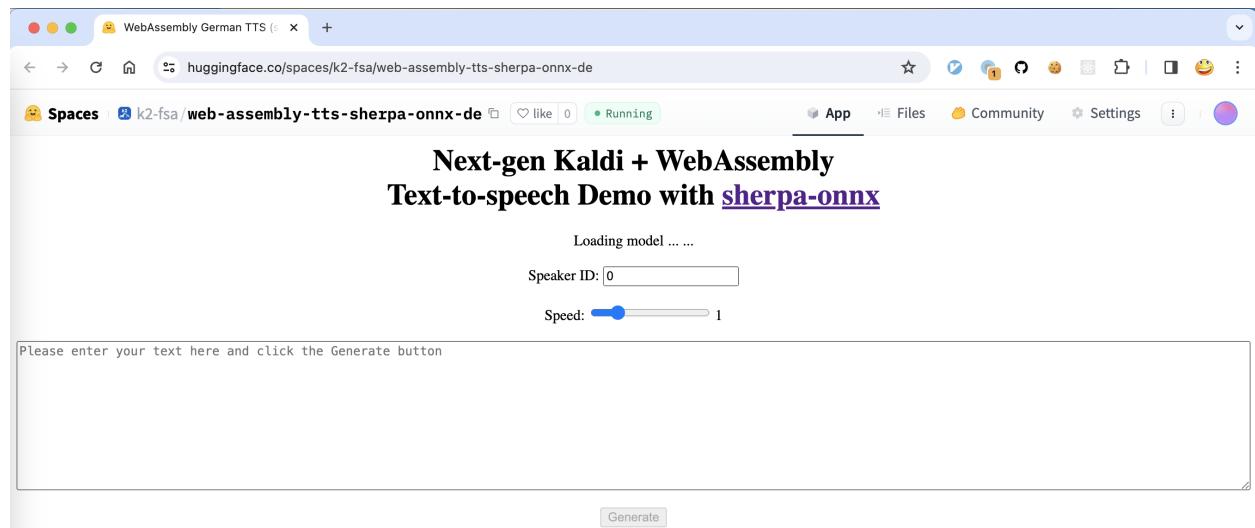
Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-en-tts.yaml>

German TTS

<https://huggingface.co/spaces/k2-fsa/web-assembly-tts-sherpa-onnx-de>

Hint: If you don't have access to Huggingface, please visit the following mirror:

<https://modelscope.cn/studios/k2-fsa/web-assembly-tts-sherpa-onnx-de/summary>



Note: The script for building this space can be found at <https://github.com/k2-fsa/sherpa-onnx/blob/master/.github/workflows/wasm-simd-hf-space-de-tts.yaml>

8.28.4 Piper

In this section, we describe how to convert `piper` pre-trained models from <https://huggingface.co/rhasspy/piper-voices>.

Hint:

You can find all of the converted models from `piper` in the following address:

<https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models>

If you want to convert your own pre-trained `piper` models or if you want to learn how the conversion works, please read on.

Otherwise, you only need to download the converted models from the above link.

Note that there are pre-trained models for over 30 languages from `piper`. All models share the same converting method, so we use an American English model in this section as an example.

Install dependencies

```
pip install onnx onnxruntime
```

Hint: We suggest that you always use the latest version of onnxruntime.

Find the pre-trained model from piper

All American English models from piper can be found at https://huggingface.co/rhasspy/piper-voices/tree/main/en/en_US.

We use https://huggingface.co/rhasspy/piper-voices/tree/main/en/en_US/amy/low as an example in this section.

Download the pre-trained model

We need to download two files for each model:

```
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/amy/low/en_US-amy-
˓→low.onnx
wget https://huggingface.co/rhasspy/piper-voices/resolve/main/en/en_US/amy/low/en_US-amy-
˓→low.onnx.json
```

Add meta data to the onnx model

Please use the following code to add meta data to the downloaded onnx model.

```
#!/usr/bin/env python3

import json
import os
from typing import Any, Dict

import onnx

def add_meta_data(filename: str, meta_data: Dict[str, Any]):
    """Add meta data to an ONNX model. It is changed in-place.

    Args:
        filename:
            Filename of the ONNX model to be changed.
        meta_data:
            Key-value pairs.
    """
    model = onnx.load(filename)
    for key, value in meta_data.items():
        meta = model.metadata_props.add()
        meta.key = key
        meta.value = str(value)
```

(continues on next page)

(continued from previous page)

```
onnx.save(model, filename)

def load_config(model):
    with open(f"{model}.json", "r") as file:
        config = json.load(file)
    return config

def generate_tokens(config):
    id_map = config["phoneme_id_map"]
    with open("tokens.txt", "w", encoding="utf-8") as f:
        for s, i in id_map.items():
            f.write(f"{s} {i[0]}\n")
    print("Generated tokens.txt")

def main():
    # Caution: Please change the filename
    filename = "en_US-amy-low.onnx"

    # The rest of the file should not be changed.
    # You only need to change the above filename = "xxx.onnx" in this file

    config = load_config(filename)

    print("generate tokens")
    generate_tokens(config)

    print("add model metadata")
    meta_data = {
        "model_type": "vits",
        "comment": "piper", # must be piper for models from piper
        "language": config["language"]["name_english"],
        "voice": config["espeak"]["voice"], # e.g., en-us
        "has_espeak": 1,
        "n_speakers": config["num_speakers"],
        "sample_rate": config["audio"]["sample_rate"],
    }
    print(meta_data)
    add_meta_data(filename, meta_data)

main()
```

After running the above script, your `en_US-amy-low.onnx` is updated with meta data and it also generates a new file `tokens.txt`.

From now on, you don't need the config json file `en_US-amy-low.onnx.json` any longer.

Download espeak-ng-data

```
wget https://github.com/k2-fsa/sherpa-onnx/releases/download/tts-models/espeak-ng-data.  
tar.bz2  
tar xf espeak-ng-data.tar.bz2
```

Note that `espeak-ng-data.tar.bz2` is shared by all models from `piper`, no matter which language your are using for your model.

Test your converted model

To have a quick test of your converted model, you can use

```
pip install sherpa-onnx
```

to install `sherpa-onnx` and then use the following commands to test your model:

```
# The command "pip install sherpa-onnx" will install several binaries,  
# including the following one  
  
which sherpa-onnx-offline-tts  
  
sherpa-onnx-offline-tts \  
  --vits-model=./en_US-amy-low.onnx \  
  --vits-tokens=./tokens.txt \  
  --vits-data-dir=./espeak-ng-data \  
  --output-filename=./test.wav \  
  "How are you doing? This is a text-to-speech application using next generation Kaldi."
```

The above command should generate a wave file `test.wav`.

Congratulations! You have successfully converted a model from `piper` and run it with `sherpa-onnx`.

8.28.5 MMS

This section describes how to convert models from <https://huggingface.co/facebook/mms-tts/tree/main> to `sherpa-onnx`.

Note that `facebook/mms-tts` supports more than 1000 languages. You can try models from `facebook/mms-tts` at the huggingface space <https://huggingface.co/spaces/mms-meta/MMS>.

You can try the converted models by visiting <https://huggingface.co/spaces/k2-fsa/text-to-speech>. To download the converted models, please visit <https://github.com/k2-fsa/sherpa-onnx/releases/tag/tts-models>. If a filename contains `vits-mms`, it means the model is from `facebook/mms-tts`.

Install dependencies

```
pip install -qq onnx scipy Cython
pip install -qq torch==1.13.0+cpu -f https://download.pytorch.org/wheel/torch_stable.html
```

Download the model file

Suppose that we want to convert the English model, we need to use the following commands to download the model:

```
name=eng
wget -q https://huggingface.co/facebook/mms-tts/resolve/main/models/$name/G_100000.pth
wget -q https://huggingface.co/facebook/mms-tts/resolve/main/models/$name/config.json
wget -q https://huggingface.co/facebook/mms-tts/resolve/main/models/$name/vocab.txt
```

Download MMS source code

```
git clone https://huggingface.co/spaces/mms-meta/MMS
export PYTHONPATH=$PWD/MMS:$PYTHONPATH
export PYTHONPATH=$PWD/MMS/vits:$PYTHONPATH

pushd MMS/vits/monotonic_align

python3 setup.py build

ls -lh build/
ls -lh build/lib*/
ls -lh build/lib*/*/

cp build/lib*/vits/monotonic_align/core*.so .

sed -i.bak s/.monotonic_align.core/.core/g ./__init__.py
popd
```

Convert the model

Please save the following code into a file with name `./vits-mms.py`:

```
#!/usr/bin/env python3

import collections
import os
from typing import Any, Dict

import onnx
import torch
from vits import commons, utils
from vits.models import SynthesizerTrn
```

(continues on next page)

(continued from previous page)

```

class OnnxModel(torch.nn.Module):
    def __init__(self, model: SynthesizerTrn):
        super().__init__()
        self.model = model

    def forward(
        self,
        x,
        x_lengths,
        noise_scale=0.667,
        length_scale=1.0,
        noise_scale_w=0.8,
    ):
        return self.model.infer(
            x=x,
            x_lengths=x_lengths,
            noise_scale=noise_scale,
            length_scale=length_scale,
            noise_scale_w=noise_scale_w,
        )[0]

def add_meta_data(filename: str, meta_data: Dict[str, Any]):
    """Add meta data to an ONNX model. It is changed in-place.

    Args:
        filename:
            Filename of the ONNX model to be changed.
        meta_data:
            Key-value pairs.
    """
    model = onnx.load(filename)
    for key, value in meta_data.items():
        meta = model.metadata_props.add()
        meta.key = key
        meta.value = str(value)

    onnx.save(model, filename)

def load_vocab():
    return [
        x.replace("\n", "") for x in open("vocab.txt", encoding="utf-8").readlines()
    ]

@torch.no_grad()
def main():
    hps = utils.get_hparams_from_file("config.json")
    is_uroman = hps.data.training_files.split(".")[-1] == "uroman"
    if is_uroman:
        raise ValueError("We don't support uroman!")

```

(continues on next page)

(continued from previous page)

```

symbols = load_vocab()

# Now generate tokens.txt
all_upper_tokens = [i.upper() for i in symbols]
duplicate = set(
    [
        item
        for item, count in collections.Counter(all_upper_tokens).items()
        if count > 1
    ]
)

print("generate tokens.txt")

with open("tokens.txt", "w", encoding="utf-8") as f:
    for idx, token in enumerate(symbols):
        f.write(f"{token} {idx}\n")

        # both upper case and lower case correspond to the same ID
        if (
            token.lower() != token.upper()
            and len(token.upper()) == 1
            and token.upper() not in duplicate
        ):
            f.write(f"{token.upper()} {idx}\n")

net_g = SynthesizerTrn(
    len(symbols),
    hps.data.filter_length // 2 + 1,
    hps.train.segment_size // hps.data.hop_length,
    **hps.model,
)
net_g.cpu()
_ = net_g.eval()

_ = utils.load_checkpoint("G_100000.pth", net_g, None)

model = OnnxModel(net_g)

x = torch.randint(low=1, high=10, size=(50,), dtype=torch.int64)
x = x.unsqueeze(0)

x_length = torch.tensor([x.shape[1]], dtype=torch.int64)
noise_scale = torch.tensor([1], dtype=torch.float32)
length_scale = torch.tensor([1], dtype=torch.float32)
noise_scale_w = torch.tensor([1], dtype=torch.float32)

opset_version = 13

filename = "model.onnx"

```

(continues on next page)

(continued from previous page)

```

torch.onnx.export(
    model,
    (x, x_length, noise_scale, length_scale, noise_scale_w),
    filename,
    opset_version=opset_version,
    input_names=[
        "x",
        "x_length",
        "noise_scale",
        "length_scale",
        "noise_scale_w",
    ],
    output_names=["y"],
    dynamic_axes={
        "x": {0: "N", 1: "L"}, # n_audio is also known as batch_size
        "x_length": {0: "N"},
        "y": {0: "N", 2: "L"},
    },
)
meta_data = {
    "model_type": "vits",
    "comment": "mms",
    "url": "https://huggingface.co/facebook/mms-tts/tree/main",
    "add_blank": int(hps.data.add_blank),
    "language": os.environ.get("language", "unknown"),
    "frontend": "characters",
    "n_speakers": int(hps.data.n_speakers),
    "sample_rate": hps.data.sampling_rate,
}
print("meta_data", meta_data)
add_meta_data(filename=filename, meta_data=meta_data)

main()

```

The you can run it with:

```

export PYTHONPATH=$PWD/MMS:$PYTHONPATH
export PYTHONPATH=$PWD/MMS/vits:$PYTHONPATH
export lang=eng
python3 ./vits-mms.py

```

It will generate the following two files:

- `model.onnx`
- `tokens.txt`

Use the converted model

We can use the converted model with the following command after installing `sherpa-onnx`.

```
./build/bin/sherpa-onnx-offline-tts \
--vits-model=./model.onnx \
--vits-tokens=./tokens.txt \
--debug=1 \
--output-filename=./mms-eng.wav \
"How are you doing today? This is a text-to-speech application using models from \
facebook with next generation Kaldi"
```

The above command should generate a wave file `mms-eng.wav`.

Congratulations! You have successfully converted a model from `MMS` and run it with `sherpa-onnx`.

We are using `eng` in this section as an example, you can replace it with other languages, such as `deu` for German, `fra` for French, etc.

8.28.6 Frequently Asked Question (FAQs)

Is there a colab notebook

Yes, we have one. Please see

https://github.com/k2-fsa/colab/blob/master/sherpa-onnx/Text_to_speech_with_sherpa_onnx.ipynb

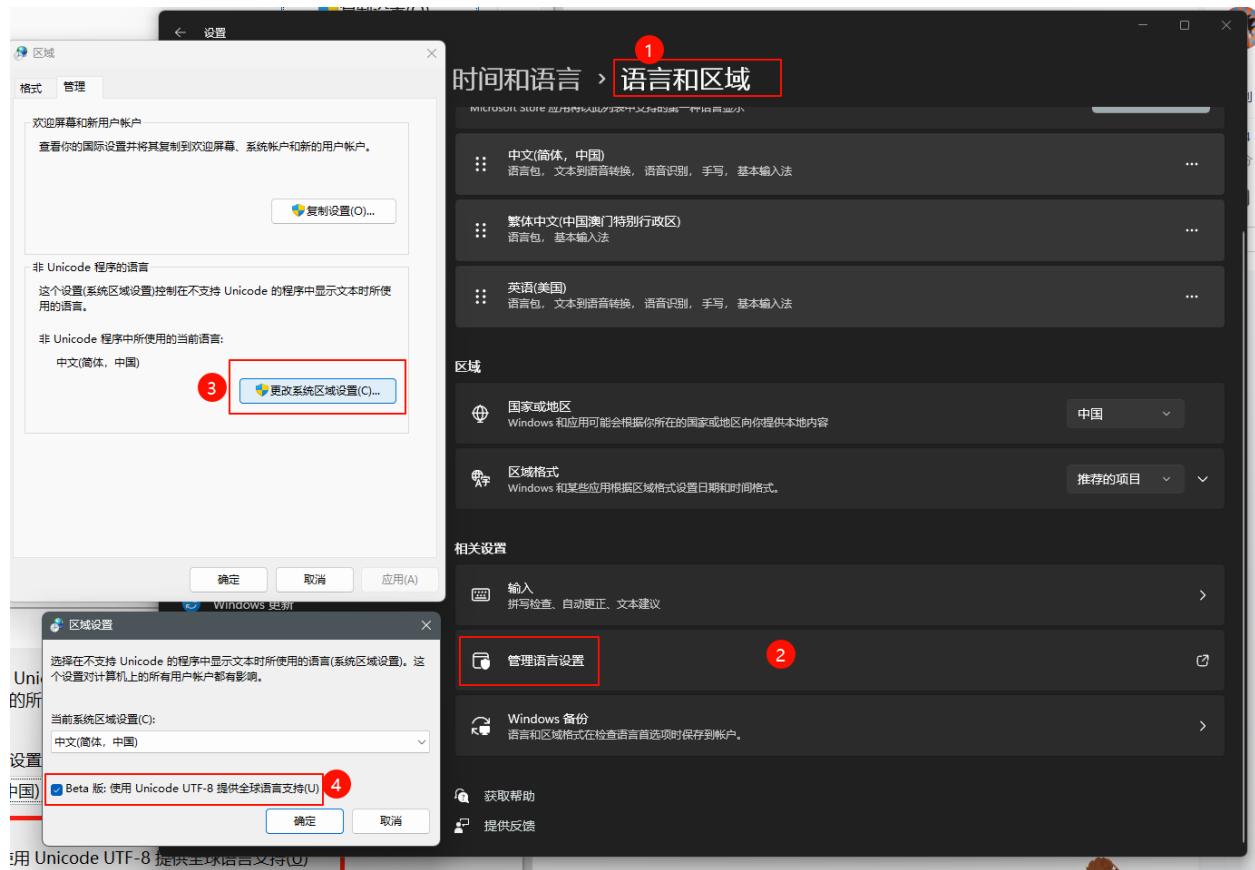
It shows you

- How to install `sherpa-onnx`
- How to download pre-trained text-to-speech (TTS) models
- How to use `sherpa-onnx` with pre-trained models for TTS

How to enable UTF-8 on Windows

For Chinese Users: [issue](#)

Please see [win11 cmdutf-8](#)



How to install sherpa-onnx for TTS

For Python users

The fastest way to install sherpa-onnx for TTS is:

```
pip install sherpa-onnx
```

The above command does NOT require you to install a C++ compiler and it supports a variety of platforms, such as:

- Linux
 - x64
 - arm, e.g., 32-bit Raspberry Pi
 - arm64, e.g., 64-bit Raspberry Pi
- Windows
 - x64, e.g., 64-bit Windows
 - x86, e.g., 32-bit Windows

- macOS
 - x64
 - arm64, e.g., M1 and M2 chips

If you want to build the `sherpa-onnx` Python package from source, please refer to [Install the Python Package](#).

After installation, please refer to <https://github.com/k2-fsa/sherpa-onnx/blob/master/python-api-examples/offline-tts.py> for example usage.

Hint: `pip install sherpa-onnx` also installs an executable `sherpa-onnx-offline-tts`. The directory where it is installed should be already on your PATH after you activate your Python virtual environment.

You can run

```
sherpa-onnx-offline-tts --help
```

in your terminal to get the help information about it.

Build from source

Please refer to [Installation](#).

Where to get pre-trained TTS models

Please refer to [Pre-trained models](#).

How to handle OOVs

Please add them to `lexicon.txt`.

TRITON

Nvidia [Triton](#) Inference Server provides a cloud and edge inferencing solution optimized for both CPUs and GPUs.

The following content describes how to deploy ASR models trained by [icefall](#) using Triton.

9.1 Installation

We prepare a dockerfile based on official triton docker containers. The customized dockerfile intergrates [Triton-server](#), [Triton-client](#) and sherpa-related requirements into a single image. You need to install [Docker](#) first before starting installation.

Hint: For your production environment, you could build triton manually to reduce the size of container.

9.1.1 Build Triton Image

```
git clone https://github.com/k2-fsa/sherpa
cd sherpa/triton
docker build . -f Dockerfile/Dockerfile.server -t sherpa_triton_server:latest
```

Note: It may take a lot of time since we build k2 from source. If you only need to use greedy search scorer, you could comment k2-related lines.

9.1.2 Launch a inference container

```
docker run --gpus all --name sherpa_server --net host --shm-size=1g -it sherpa_triton_
server:latest
```

Now, you should enter into the container successfully.

9.2 Triton-server

This page gives several examples to deploy streaming and offline ASR pretrained models with Triton server.

9.2.1 Deploy streaming ASR models with Onnx

First, we need to export pretrained models with Onnx.

```
export SHERPA_SRC=../sherpa
export ICEFALL_SRC=/workspace/icefall
# copy essentials
cp $SHERPA_SRC/triton/scripts/*onnx*.py $ICEFALL_DIR/egs/wenetspeech/ASR/pruned_
˓→stateless_transducer5/
cd $ICEFALL_SRC/egs/wenetspeech/ASR/
# download pretrained models
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/luomingshuang/icefall_asr_
˓→wenetspeech_pruned_transducer_stateless5_streaming
cd ./icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming
git lfs pull --include "exp/pretrained_epoch_7_avg_1.pt"
cd -
# export to onnx fp16
ln -s ./icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming/exp/pretrained_
˓→epoch_7_avg_1.pt ./icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming/exp/
˓→epoch-999.pt
./pruned_transducer_stateless5/export_onnx.py \
--exp-dir ./icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming/exp \
--tokenizer-file ./icefall_asr_wenetspeech_pruned_transducer_stateless5_streaming/
˓→data/lang_char \
--epoch 999 \
--avg 1 \
--streaming-model 1 \
--causal-convolution 1 \
--onnx 1 \
--left-context 64 \
--right-context 4 \
--fp16
```

Note: For Chinese models, `--tokenizer-file` points to `<pretrained_dir>/data/lang_char`. While for English models, it points to `<pretrained_dir>/data/lang_bpe_500/bpe.model` file.

Then, in the docker container, you could start the service with:

```
cd sherpa/triton/
bash scripts/start_streaming_server.sh
```

9.2.2 Deploy offline ASR models with torchscript

Caution: Currently, we only support FP32 offline ASR inference for torchscript backend. Streaming ASR and FP16 inference are not supported.

First, we need to export pretrained models using jit.

```
export SHERPA_SRC=./sherpa
export ICEFALL_SRC=/workspace/icefall
# Download pretrained models
GIT_LFS_SKIP_SMUDGE=1 git clone https://huggingface.co/csukuangfj/icefall-asr-
↪librispeech-pruned-transducer-stateless3-2022-04-29 $ICEFALL_DIR/egs/librispeech/ASR/
↪pruned_stateless_transducer3/
cd icefall-asr-librispeech-pruned-transducer-stateless3-2022-04-29
git lfs pull --include "exp/pretrained-epoch-25-avg-7.pt"
# export them to three jit models: encoder_jit.pt, decoder_jit.pt, joiner_jit.pt
cp $SHERPA_SRC/triton/scripts/conformer_triton.py $ICEFALL_DIR/egs/librispeech/ASR/
↪pruned_stateless_transducer3/
cp $SHERPA_SRC/triton/scripts/export_jit.py $ICEFALL_DIR/egs/librispeech/ASR/pruned_
↪stateless_transducer3/
cd $ICEFALL_DIR/egs/librispeech/ASR/pruned_stateless_transducer3
python3 export_jit.py \
    --pretrained-model $ICEFALL_DIR/egs/librispeech/ASR/pruned_stateless_transducer3/
↪icefall-asr-librispeech-pruned-transducer-stateless3-2022-04-29 \
    --output-dir <jit_model_dir> --bpe-model <bpe_model_path>
# copy bpe file to <jit_model_dir>, later we would mount <jit_model_dir> to the triton_
↪docker container
cp <bpe_model_path> <jit_model_dir>
```

Note: If you export models outside the docker container, you could mount the exported <jit_model_dir> with -v <host_dir>:<container_dir> when lauching the container.

Then, in the docker container, you could start the service with:

```
cd sherpa/triton/
bash scripts/start_offline_server_jit.sh
```

9.3 Triton-client

9.3.1 Send requests using client

In the docker container, run the client script to do ASR inference:

```
cd sherpa/triton/client
# Test one audio using offline ASR
python3 client.py --audio_file=./test_wavs/1089-134686-0001.wav --url=localhost:8001
```

(continues on next page)

(continued from previous page)

```
# Test one audio using streaming ASR
python3 client.py --audio_file=./test_wavs/1089-134686-0001.wav --url=localhost:8001 --
→streaming
```

The above command sends a single audio `1089-134686-0001.wav` to the server and get the result. `--url` option specifies the IP and port of the server, in this example, we set the server and client on the same machine, therefore IP is `localhost`, and we use port `8001` since it is the default port for gRPC in Triton.

You can also test a bunch of audios together with the client. Just specify the path of `wav.scp` with `--wavscp` option, set the path of test set directory with `--data_dir` option, and set the path of ground-truth transcript file with `--trans` option, the client will infer all the audios in test set and calculate the WER upon the test set.

9.3.2 Decode manifests

You could also decode a whole dataset to benchmark metrics e.g. RTF, WER.

Caution: Decode manifests in simulation streaming mode would be supported in the future.

```
cd sherpa/triton/client

# For aishell manifests:
git lfs install
git clone https://huggingface.co/csukuangfj/aishell-test-dev-manifests
sudo mkdir -p /root/fangjun/open-source/icefall-aishell/egs/aishell/ASR/download/aishell
tar xf ./aishell-test-dev-manifests/data_aishell.tar.gz -C /root/fangjun/open-source/
→icefall-aishell/egs/aishell/ASR/download/aishell/
# dev set: ./aishell-test-dev-manifests/data/fbank/aishell_cuts_test.jsonl.gz
# test set: ./aishell-test-dev-manifests/data/fbank/aishell_cuts_test.jsonl.gz

python3 decode_manifest_triton.py \
    --server-addr localhost \
    --num-tasks 300 \
    --log-interval 20 \
    --model-name transducer \
    --manifest-filename ./aishell-test-dev-manifests/data/fbank/aishell_cuts_test.jsonl.
→gz \
    --compute-cer
```

9.4 Perf Analyzer

We can use `perf_analyzer` provided by Triton to test the performance of the service.

9.4.1 Generate Input Data from Audio Files

For offline ASR server:

```
cd sherpa/triton/client
# en
python3 generate_perf_input.py --audio_file=test_wavs/1089-134686-0001.wav
# zh
python3 generate_perf_input.py --audio_file=test_wavs/zh/mid.wav
```

It will generate a `offline_input.json` file in `sherpa/triton/client`.

For streaming ASR server, you need to add a `--streaming` option:

```
python3 generate_perf_input.py --audio_file=test_wavs/1089-134686-0001.wav --streaming
```

A `online_input.json` file would be generated.

9.4.2 Test Throughput using Perf Analyzer

```
# Offline ASR Test with grpc
perf_analyzer -m transducer -b 1 -a -p 20000 --concurrency-range 100:200:50 -i gRPC --
    ↵input-data=offline_input.json -u localhost:8001

# Streaming ASR Test with grpc
perf_analyzer -m transducer -b 1 -a -p 20000 --concurrency-range 100:200:50 -i gRPC --
    ↵input-data=online_input.json -u localhost:8001 --streaming
```

You could save the below results with a `-f log.txt` option.

Concurrency	Inferences/Send	Client Send	Net-work+Server Queue	Server Send/Recv	Server Compute Input	Server Infer	Server Compute Output	Client Recv	p50 latency	p90 latency	p95 latency	p99 latency
300	226.24	109	230434	1	9314	1068792	14512	1	1254206616224958246551406			

Note: Please refer to https://github.com/triton-inference-server/server/blob/main/docs/user_guide/perf_analyzer.md for advanced usage.

9.5 TensorRT acceleration

This page shows how to use TensorRT engine to accelerate inference speed for K2 models

9.5.1 Preparation

First of all, you have to install the TensorRT. Here we suggest you to use docker container to run TRT. Just run the following command:

```
docker run --gpus '"device=0"' -it --rm --net host -v $PWD:/k2 nvcr.io/nvidia/  
→tensorrt:22.12-py3
```

You can also see [here](#) to build TRT on your machine.

Note: Please pay attention that, the TRT version must have to >= 8.5.3!!!

If your TRT version is < 8.5.3, you can download the desired TRT version and then run the following command inside the docker container to use the TRT you just download:

```
# inside the container  
bash tools/install.sh
```

9.5.2 Model export

You have to prepare the ONNX model by referring [here](#) to export your models into ONNX format. Assume you have put your ONNX model in the \$model_dir directory. Then, just run the command:

```
bash tools/build.sh $model_dir  
cp $model_dir/encoder.trt model_repo_offline_fast_beam_trt/encoder/1
```

The generated TRT model will be saved into \$model_dir/encoder.trt. We also give an example of model_repo of TRT model. You can follow the same procedure as described [here](#) to deploy the pipeline using triton.

9.5.3 Benchmark for Conformer TRT encoder vs ONNX

Model	Batch size	Avg latency(ms)	QPS
ONNX	1	7.44	134.48
	8	14.92	536.09
	16	22.84	700.67
	32	41.62	768.84
	64	80.48	795.27
	128	171.97	744.32
TRT	1	5.21834	193.93
	8	11.7826	703.49
	16	20.4444	815.79
	32	37.583	893.56
	64	69.8312	965.40
	128	139.702	964.57